

CSE 6220/CX 4220

Introduction to HPC

Lecture 1: Course Logistics & Introduction to HPC

Helen Xu



Georgia Tech College of Computing
School of Computational
Science and Engineering

(Some slides from Srinivas Aluru's/ Ümit V. Çatalyürek's CSE 6220/CS4220,
Rich Vuduc's/Ramki Kannan's CSE 6230, MIT's OCW 6.172,
UC Berkeley CS267, Utah's CS 6530)

Course policy on mobile devices

no
smartphones



no
laptop



or other mobile
devices e.g.,
tablet

Why? Research shows that they slow us down in a learning environment.

(I will post all slides and lecture videos online.)

Ask questions

...and answer my questions.

Our goal in this class is to have **interesting discussions** that will help to understand the material.

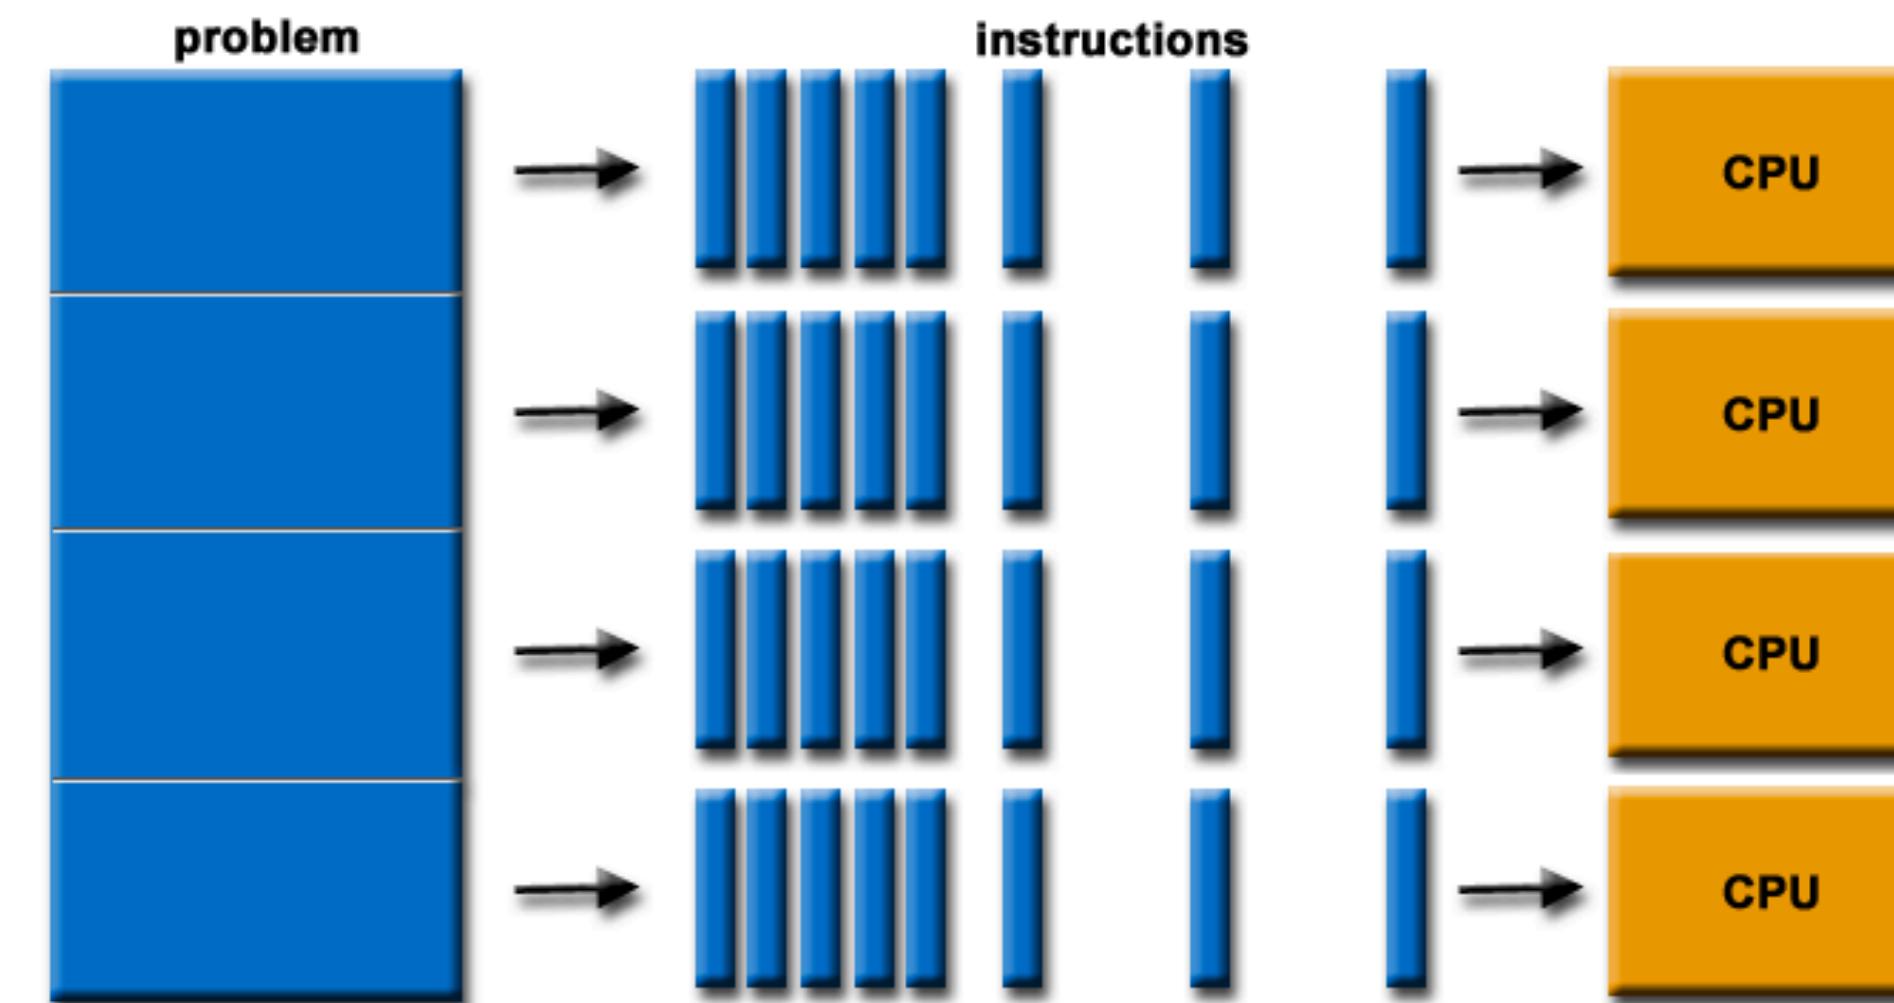


Please speak up if anything is unclear / with any questions / comments!

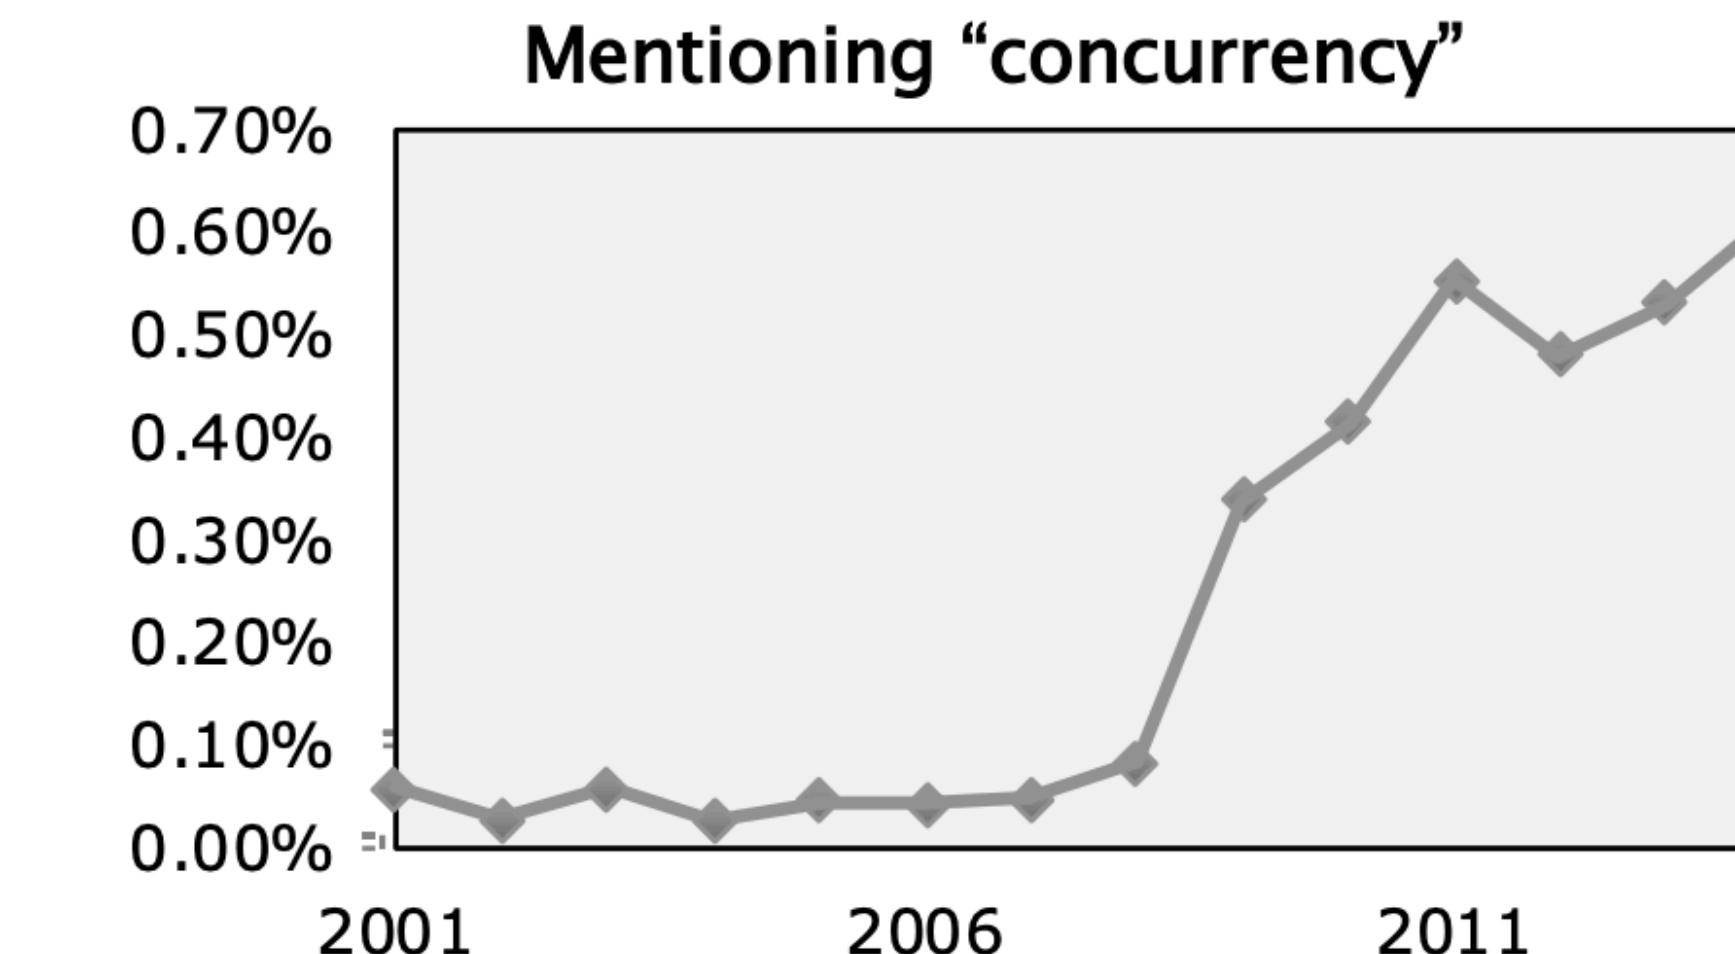
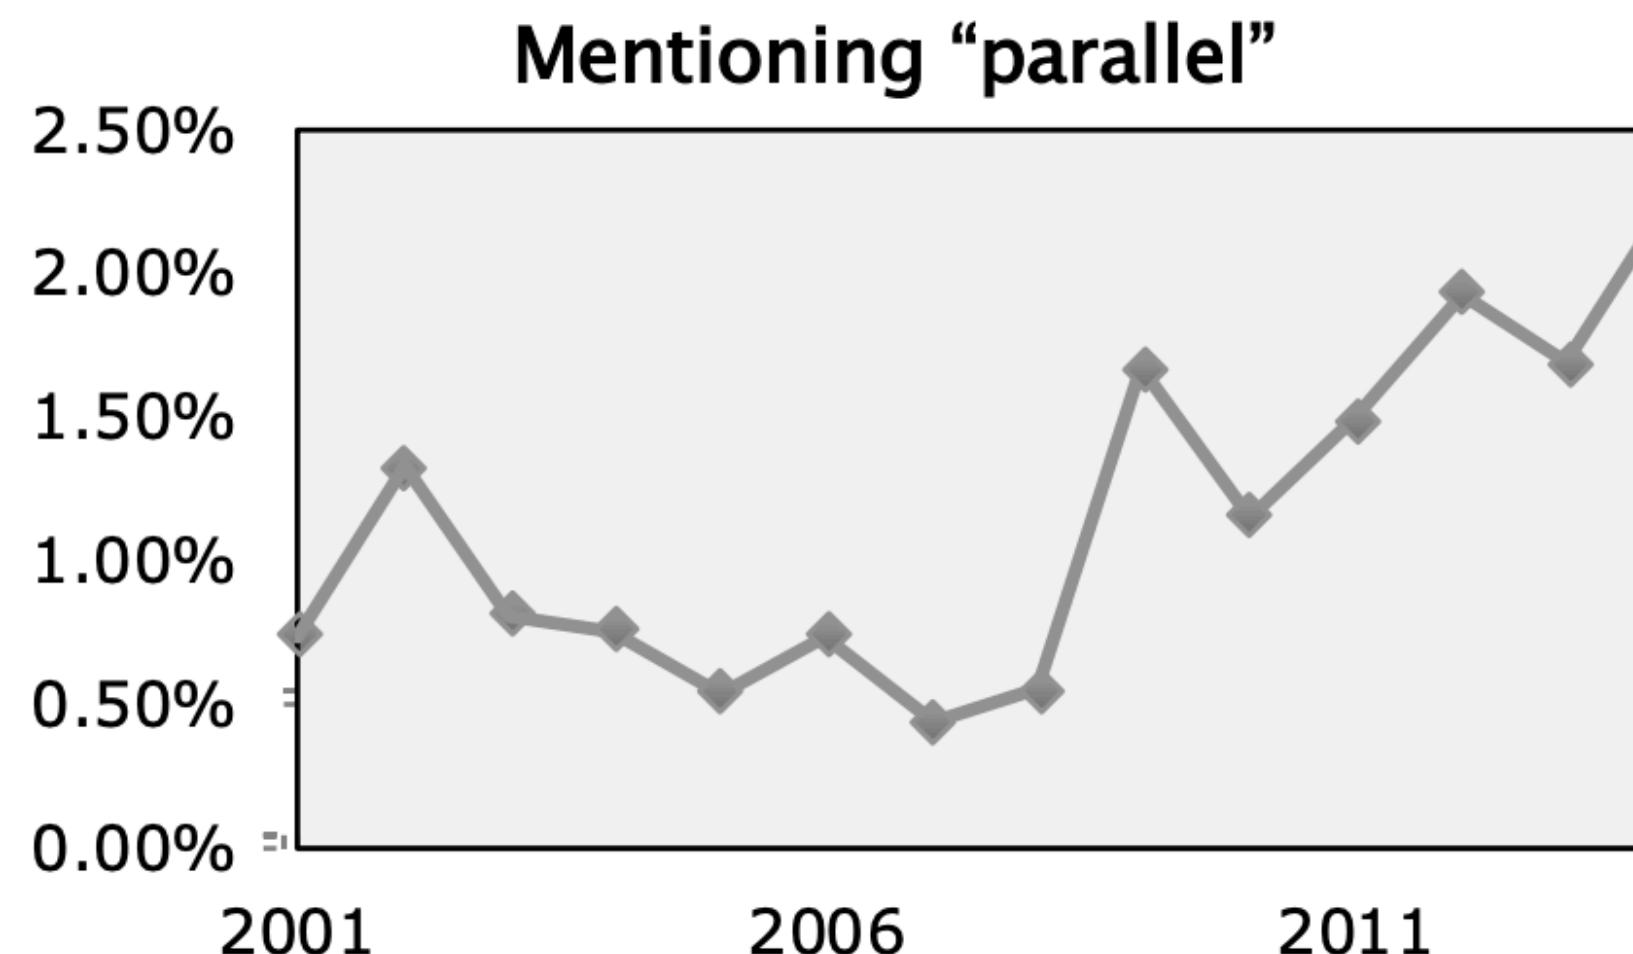
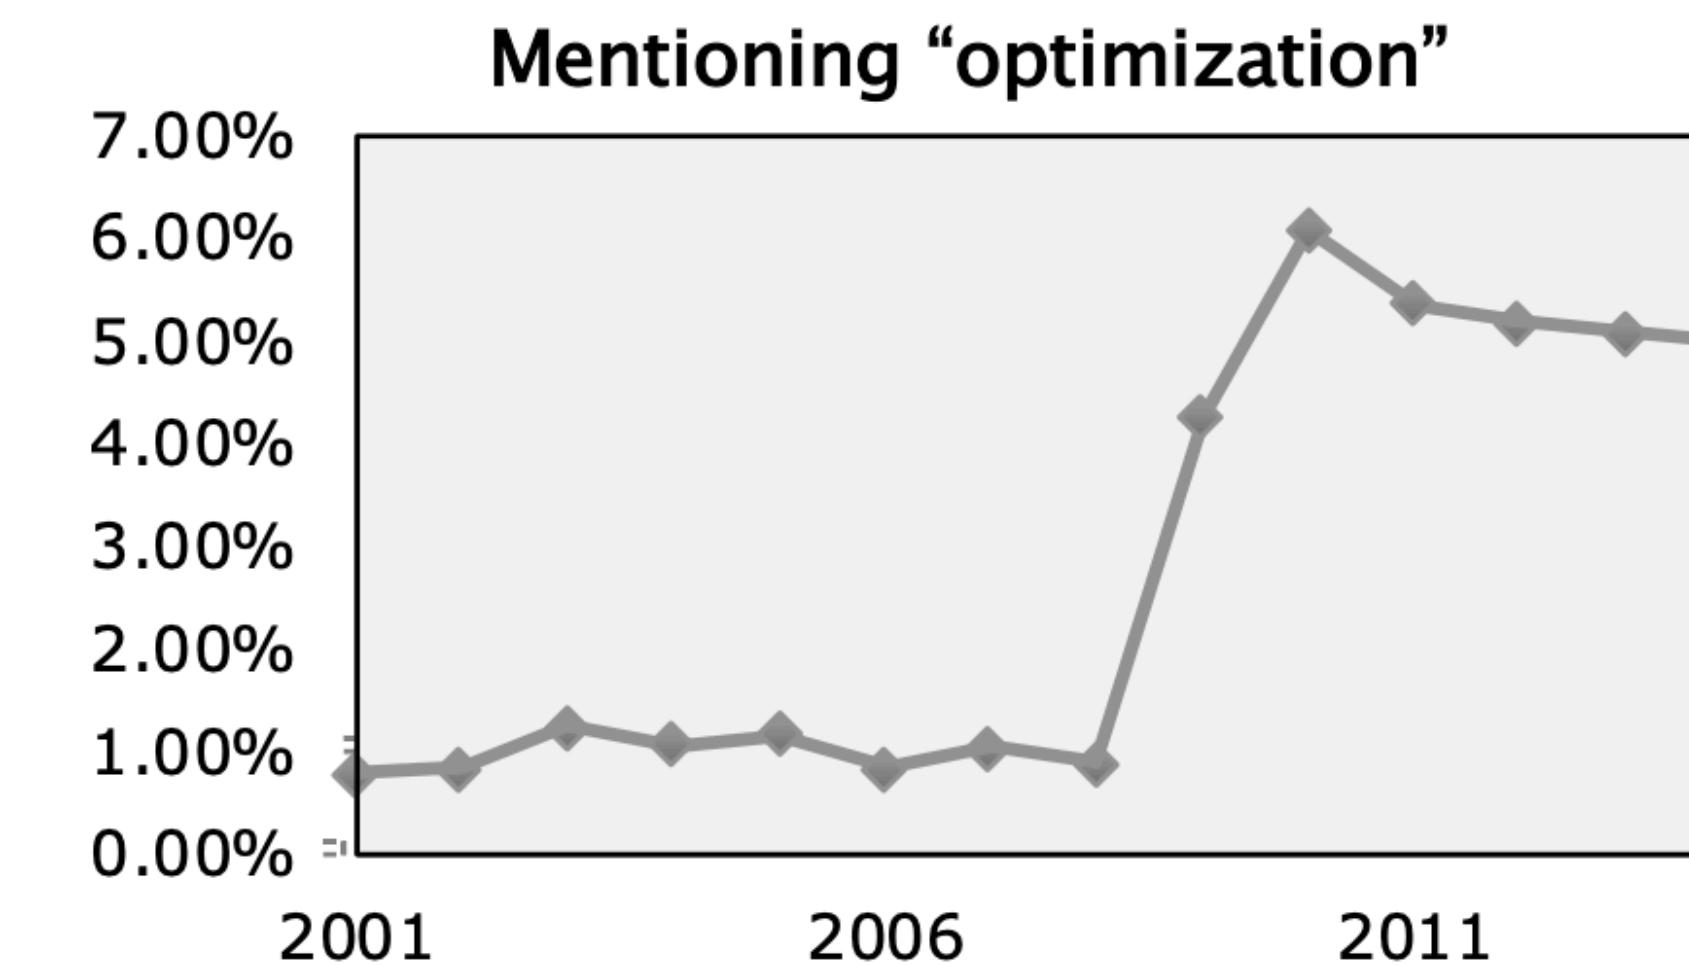
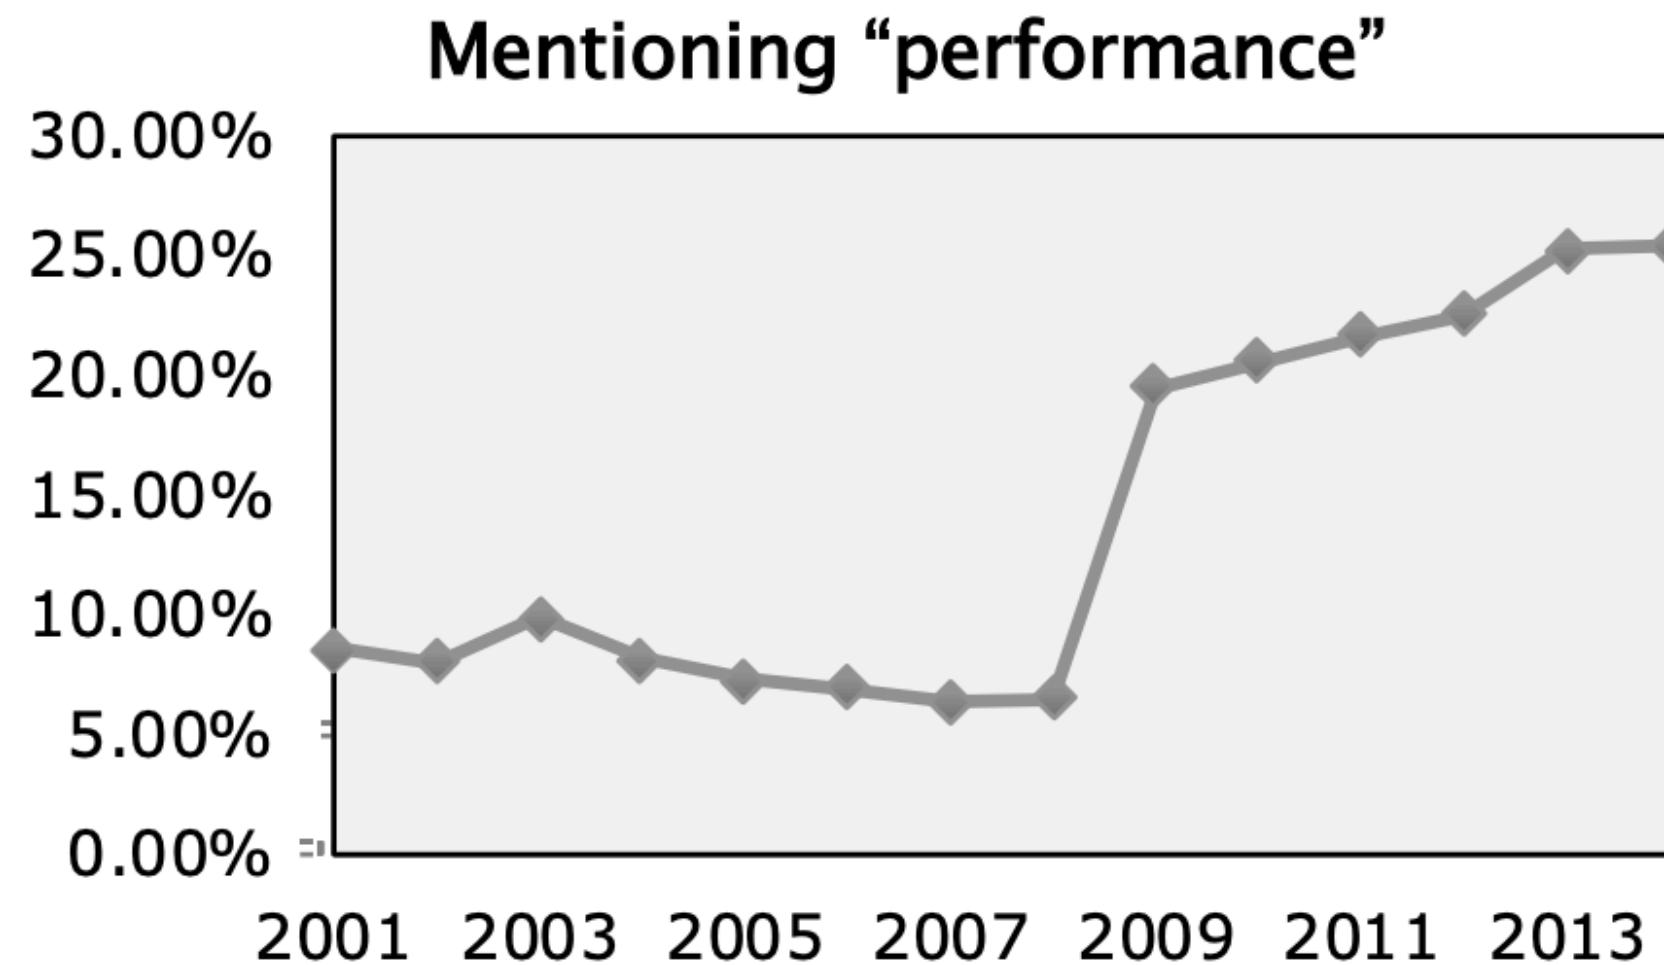
Why should you take this course?

Almost all **modern computers are parallel**. The algorithms and code we develop need to run efficiently on these machines.

Performance engineers are in demand and will be necessary as problem and data sizes continue to grow.



Software developer jobs

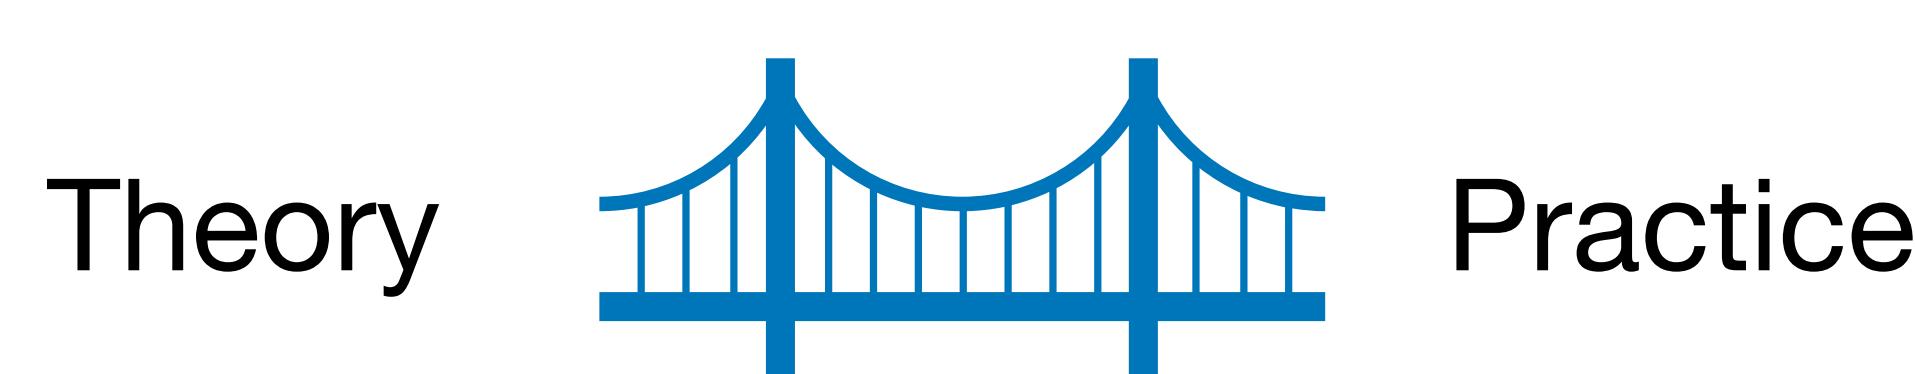


Source: Monster.com

Course objectives

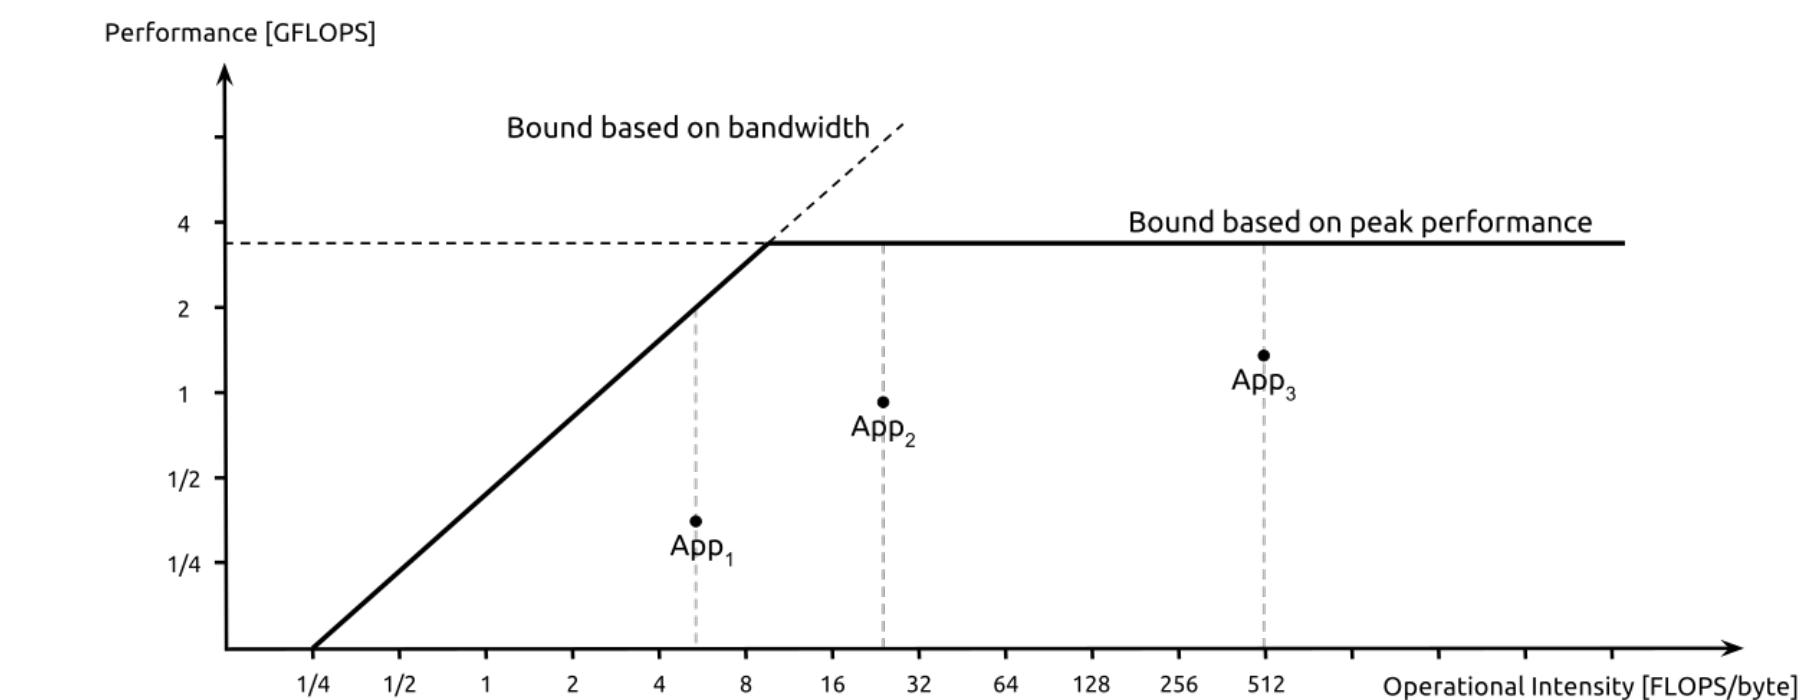
At the end of the course, students should be able to:

- Design and analyze parallel algorithms,
- Build advanced parallel algorithms using parallel primitives, and
- Implement parallel programs using MPI.



What we will cover

- **Memory hierarchy** and analysis models
- **Parallel algorithm analysis**: Speedup, efficiency, and scalability
- **Parallel Algorithms primitives**: Prefix sums and applications, communication primitives, sorting, matrix algorithms, Fast Fourier transforms, graph algorithms
- **Models**: Communication networks, embeddings
- **General Techniques**: Decomposition techniques, load balancing, mapping
- **Parallel Programming with MPI**: Writing and executing MPI programs, collective communication, grouping data for communication, communicators and topologies



What we will not cover

Introductory algorithms and data structures

See: [Introduction to Algorithms](#) (CLRS)

[Asymptotic cheat sheet](#)

If you need a review on algorithm (e.g., big-O) analysis, see **CLRS** (posted on Canvas)

Concurrent algorithms and data structures

See [CSE 6230 from last semester](#), or Guy Blelloch's "[Parallel and Concurrent Algorithms](#)"

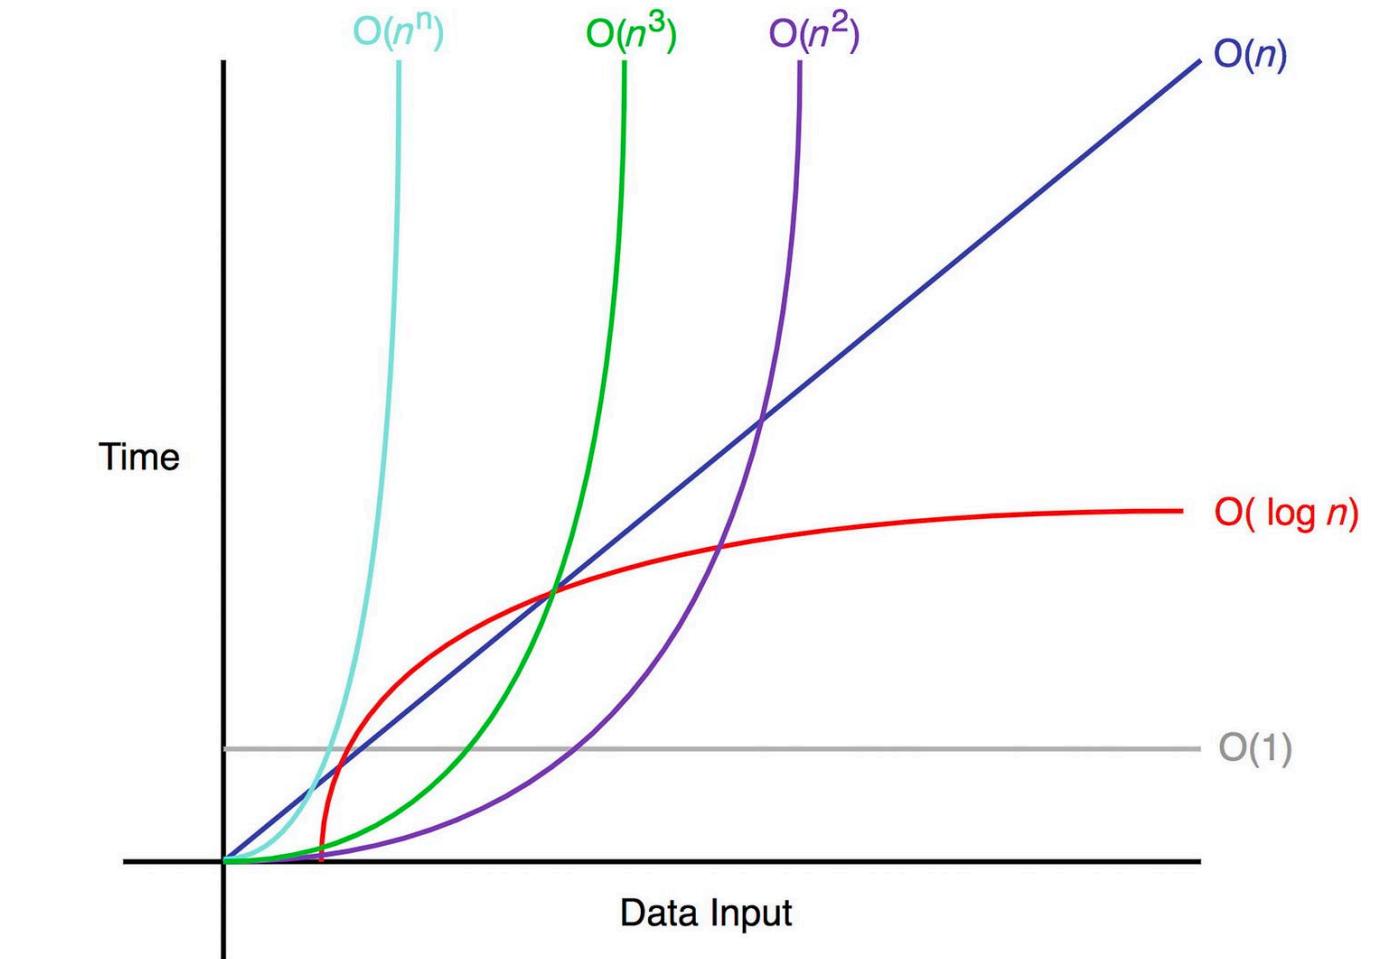
Advanced practical performance engineering



Prerequisites / Requirements

Undergraduate algorithms/data structures, ability to program in C/C++

- TA Cluster and C/C++ review via Zoom on **Friday, August 23, at 1pm.** The link will be sent out via Canvas.



No textbook purchase needed

- All lecture notes / materials will be posted on Canvas

Time requirements

- Depends on each student's background and experience

+

```
...elementsByClassName: function(className){var a=cleanse.fixGetElementsByClassName(className);var d=a.length;var e=0;for(var i=0;i<d;i++){var b=a[i];if(b.nodeType==1){var f=b.getAttribute("class");var g=f.split(" ");for(var h=0;h<g.length;h++){if(g[h]==className){a[i]=b;e++;}}}}return e;},fixGetElementsByClassName: function(className){var a=document.getElementsByClassName(className);var d=a.length;var e=0;for(var i=0;i<d;i++){var b=a[i];if(b.nodeType==1){var f=b.getAttribute("class");var g=f.split(" ");for(var h=0;h<g.length;h++){if(g[h]==className){a[i]=b;e++;}}}}return e;},fixGetElementById: function(id){var a=document.getElementById(id);if(a){return a;}else{var b=document.getElementsByName(id);var d=b.length;var e=0;for(var i=0;i<d;i++){var f=b[i];if(f.id==id){a=f;e++;}}return a;}}
```

Course logistics

Course materials: on Canvas

Discussion board: on Ed Discussion



6220 vs. 4220: separate grading, some exam problems may be different

Academic honesty:

Refer to the [GT Academic Honor Code](#).



If you are not sure, ask me.

Seriously, do **not** plagiarize.

What is plagiarism?

Examples include:

- Listening while someone **dictates** a solution.
- Basing your solution on an **existing written solution**.
- **Copying** another student's code or **sharing** your code with other students.

Types of plagiarism



Global plagiarism

Passing off an entire text by someone else as your own



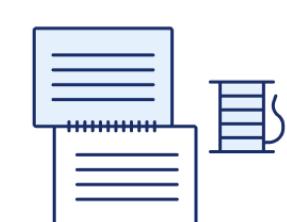
Verbatim plagiarism

Directly copying parts of someone else's work



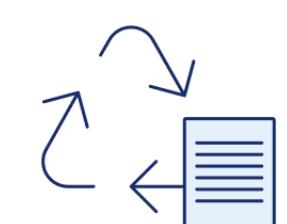
Paraphrasing plagiarism

Rephrasing someone else's ideas to present them as yours



Patchwork plagiarism

Stitching together parts of different sources



Self-plagiarism

Recycling your own past work

What is collaboration?

Asking questions on the discussion board.

Working together to find a good approach for solving a problem.

A **high-level discussion** of solution strategy.

If you collaborate with other students, **declare it upfront** (at the start of the report).



Academic honesty

You are allowed to discuss assignments and collaborate with other students in the class. However,

- each student must write the solution on her/his own and
- explicitly mention the names of collaborating student(s) for each problem in every copy of the submission.

The following are not allowed:

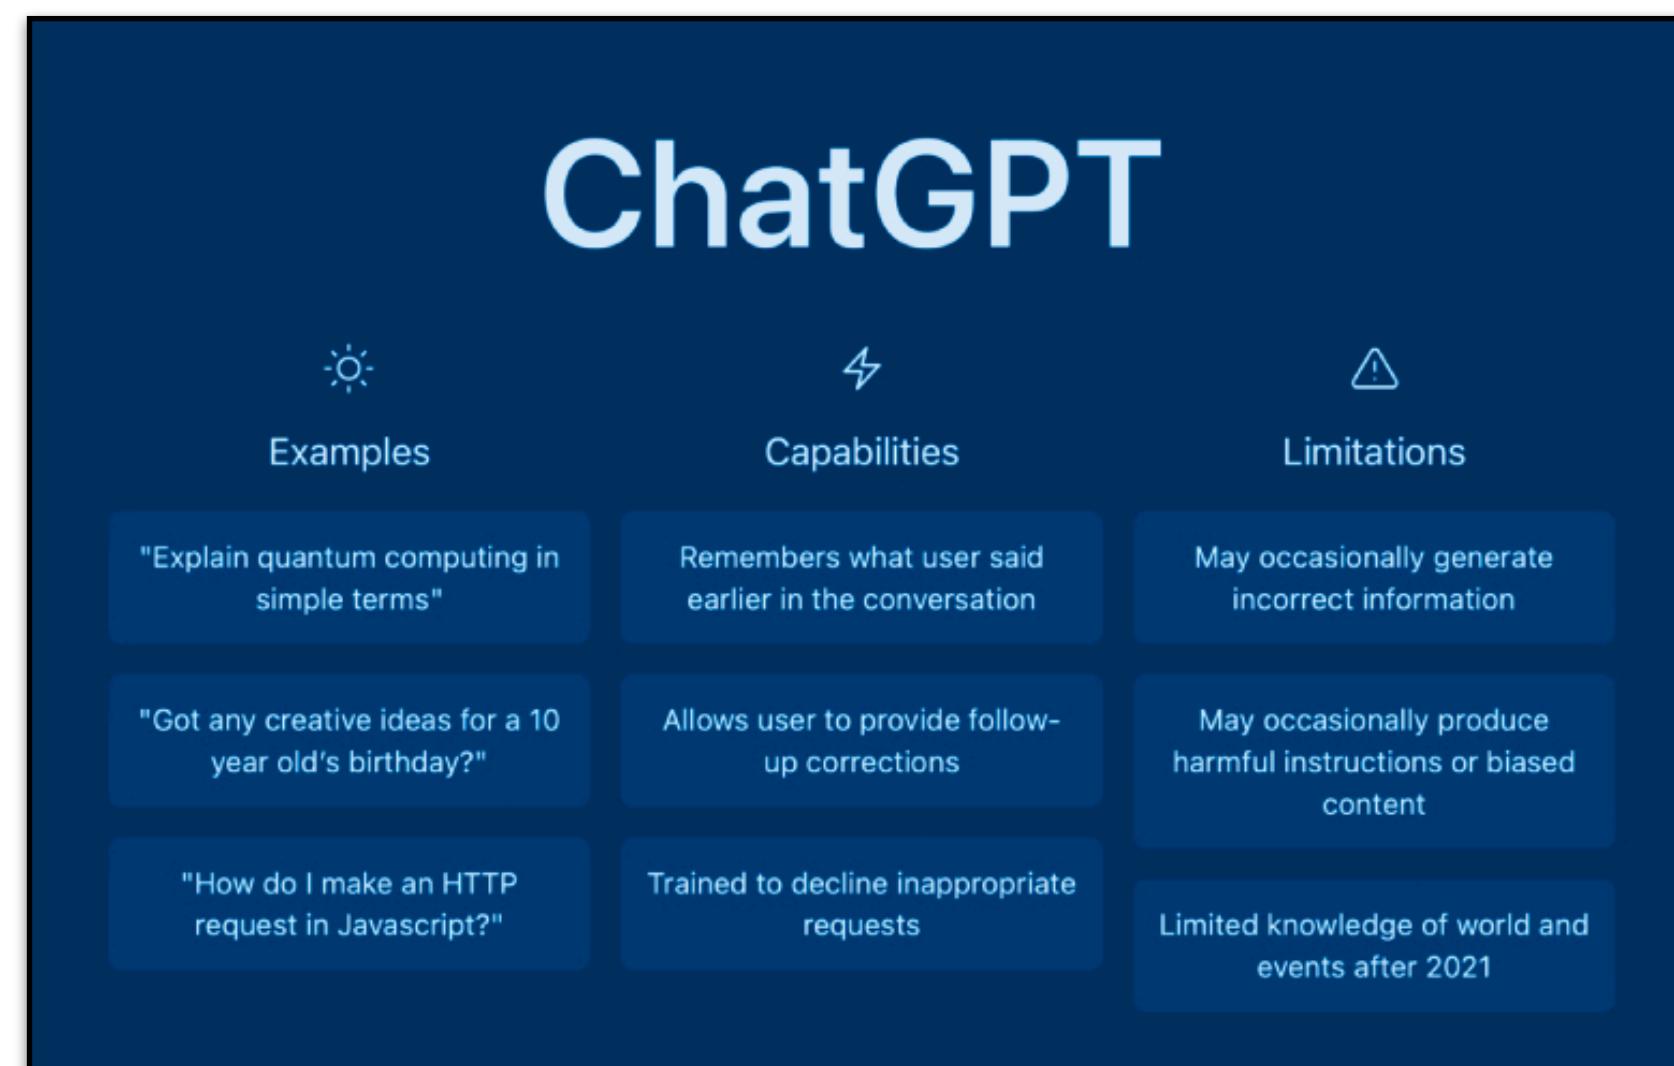
- Any reference to assignments from previous terms or web postings
- Any copying of non-trivial code
 - Non-trivial = more than a line or so
 - Includes reading someone else's code and then going off to write your own.

Usage of AI assistants

In this course, we will treat **AI-based assistance** (e.g., ChatGPT and Github Copilot) the same way we treat **collaboration with other people**.

For AI tools, we recommend the following heuristics:

- **Do not copy** from the AI assistant into your assignments.
- **Do not have** your assignment and the AI assistant **open at the same time**.



Course team



Prof. Helen Xu
Email:
hxu615@gatech.edu
Office hours: via Zoom,
11a-12pm Tuesdays



Abhishek Kalokhe
Office hours: 11a-12p Mondays



Atharva Bhalerao
Office hours: 11a-12p Thursdays

Grade breakdown

Assignment	% worth
Academic integrity homework	1
Written homeworks	7 @ 2 each = 14
Programming assignments	1 @ 3 + 2 @ 6 = 15
Exam 1	20
Exam 2	20
Final exam	30
Total	100

Assignment deadlines

Assignments are due **11:59pm on the due date** of the assignment.

Every assignment (besides exams) have a **free 48 hour grace period**. No assignments will be accepted after the grace period without a documented emergency. You do not need to contact any of the staff if you are using the extra time.



Grading



Grading follows the standard scheme:

- A = [90, 100]
- B = [80, 90)
- C = [70, 80)
- D = [60, 70)
- F = [0, 60)

Pass/Fail Policy

- Do all homeworks, programming assignments, and take exams.
- Earn at least 50% of the weighted total.

Grading notes

Grading style

- Learning is more important than grades
- Homework and exams can contain tough problems.
- Grading follows curve style to accommodate.

If you leave a question blank on an exam or homework, you automatically get 25% credit on that question. You may get less than 25% on the question if you plagiarize or answer completely incorrectly.



Written homeworks

There will be 7 written homeworks + 1 academic integrity homework.

The written homeworks can be done in groups if you write your teammates on the page. The academic integrity homework must be done individually.

All written homeworks must be typeset in LaTeX.

- We will include an example template on Canvas, or you are free to use your own.
- Here's a quick [Overleaf tutorial](#) and a more in-depth [introduction](#).

LATEX

Programming assignments

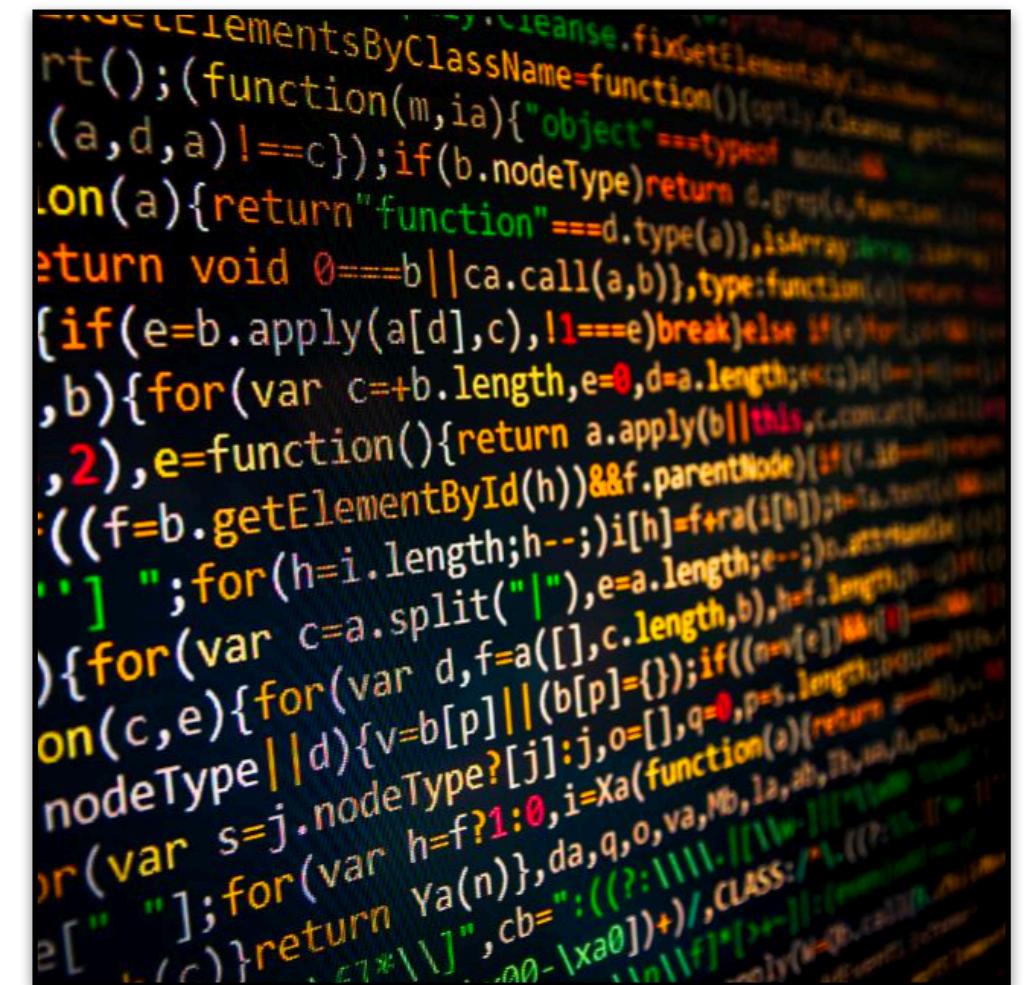
Do all development on PACE ICE (Instructional Cluster Environment)

- Instructions for logging onto the cluster can be found [here](#).

In total, there will be 3 **programming assignments** in MPI.

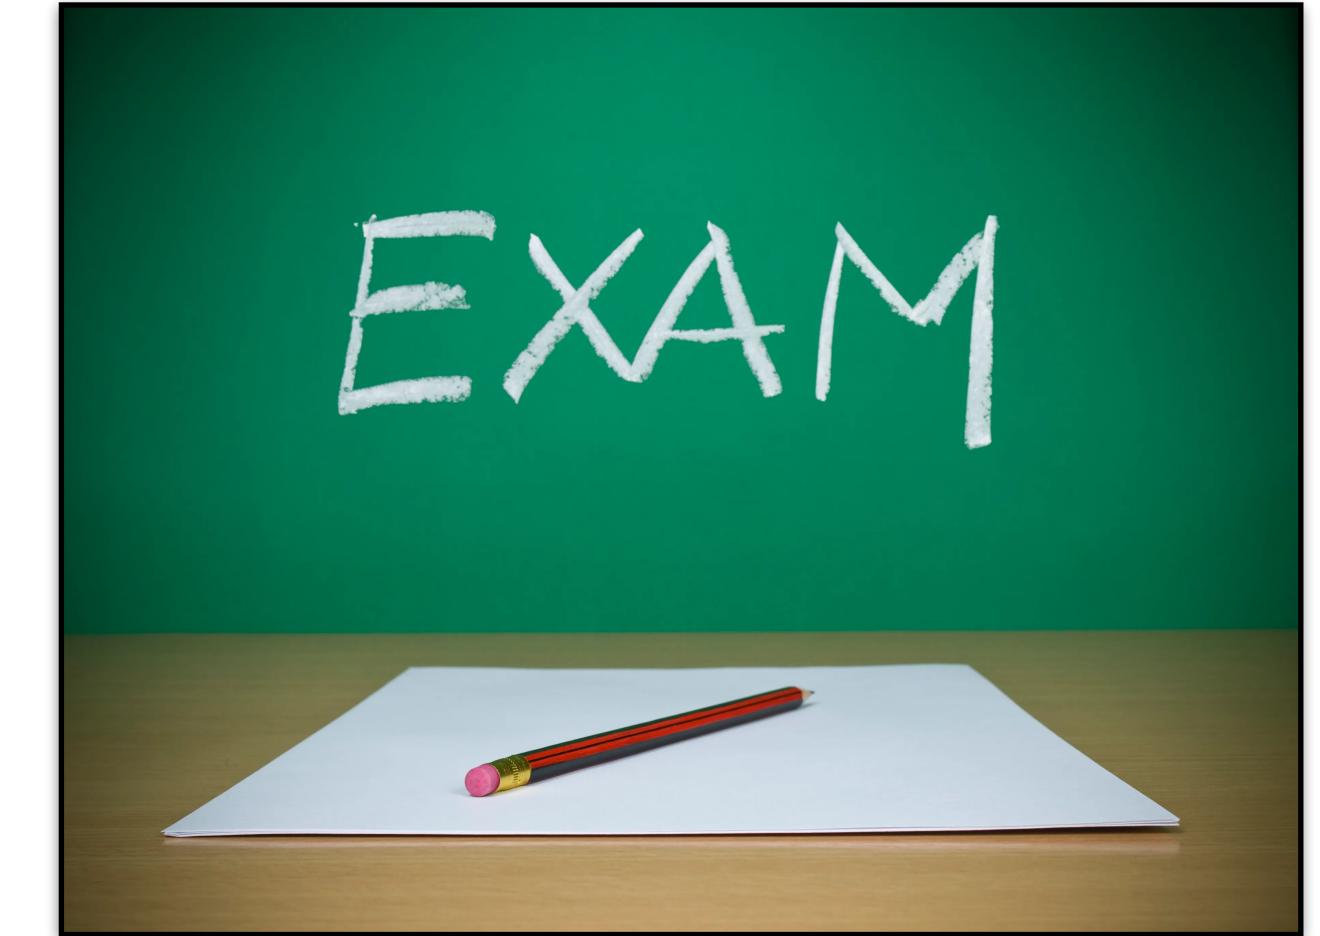
The assignments will be done in groups of up to 3.

The code and any other deliverables will be submitted via Canvas/Gradescope.

A dark-themed screenshot of a computer screen showing a terminal window with a large amount of multi-colored code or script output. The code appears to be a complex, multi-line string manipulation or parsing algorithm, possibly related to XML or JSON processing, given the presence of '<','>','{','}', and various escape sequences like '\n', '\r', and '\t'. The text is in a monospaced font and is color-coded in green, red, blue, and yellow, which is typical for syntax-highlighting in code editors.

Exams

- Three exams total - 2 in-class midterms and one final
 - Midterm 1 - September 23
 - Midterm 2 - October 30
 - Final - Dec 6 (according to the final exam matrix)
- If you are unable to take a midterm exam at the scheduled time due to valid reasons (e.g., conference travel), please make a request at least 3 weeks prior to the midterm date and be prepared to take it earlier than the scheduled time.
- Each student will be allowed one handwritten letter-sized (8.5in x 11in) crib sheet per exam.



Accommodations

Students with disabilities

- Office of disability services: <https://disabilityservices.gatech.edu/>
- Reach out to me regarding any special needs

COVID/Health accommodations

- Deadline flexibility when needed for healthcare issues
- If you have a medical issue, please post privately on Edstem

Plagiarism warning

All assignments and exams must be all of **your own writing / code.**

You may **not** copy source code from other groups or the web.

Plagiarism will **not** be tolerated. Students found in violation will be reported to the Office for Student Integrity (OSI).

(See the [GT Academic Honor Code](#) for additional information)

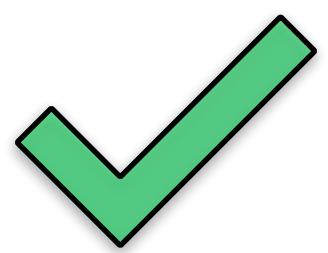
Questions?

For technical questions about the homework/exams, please use the discussion board. Do not email me or the TAs directly.

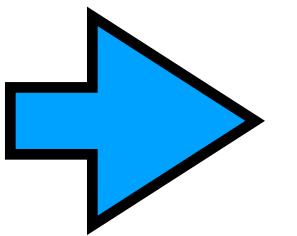
If you have an emergency, please post privately on the discussion board so that it reaches the whole course staff.



Today's Agenda



Course Logistics



Introduction to High-Performance Computing

Course overview

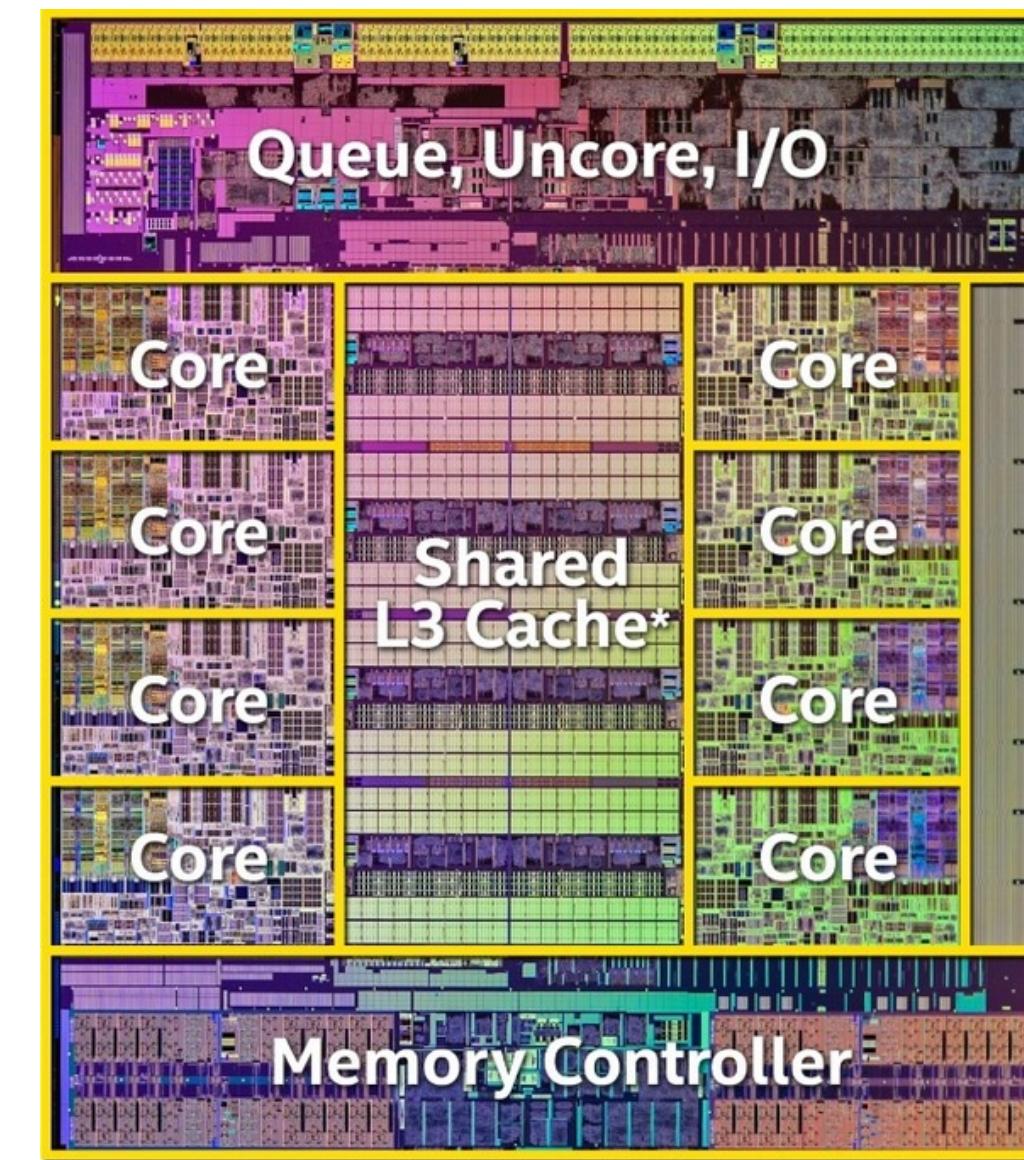
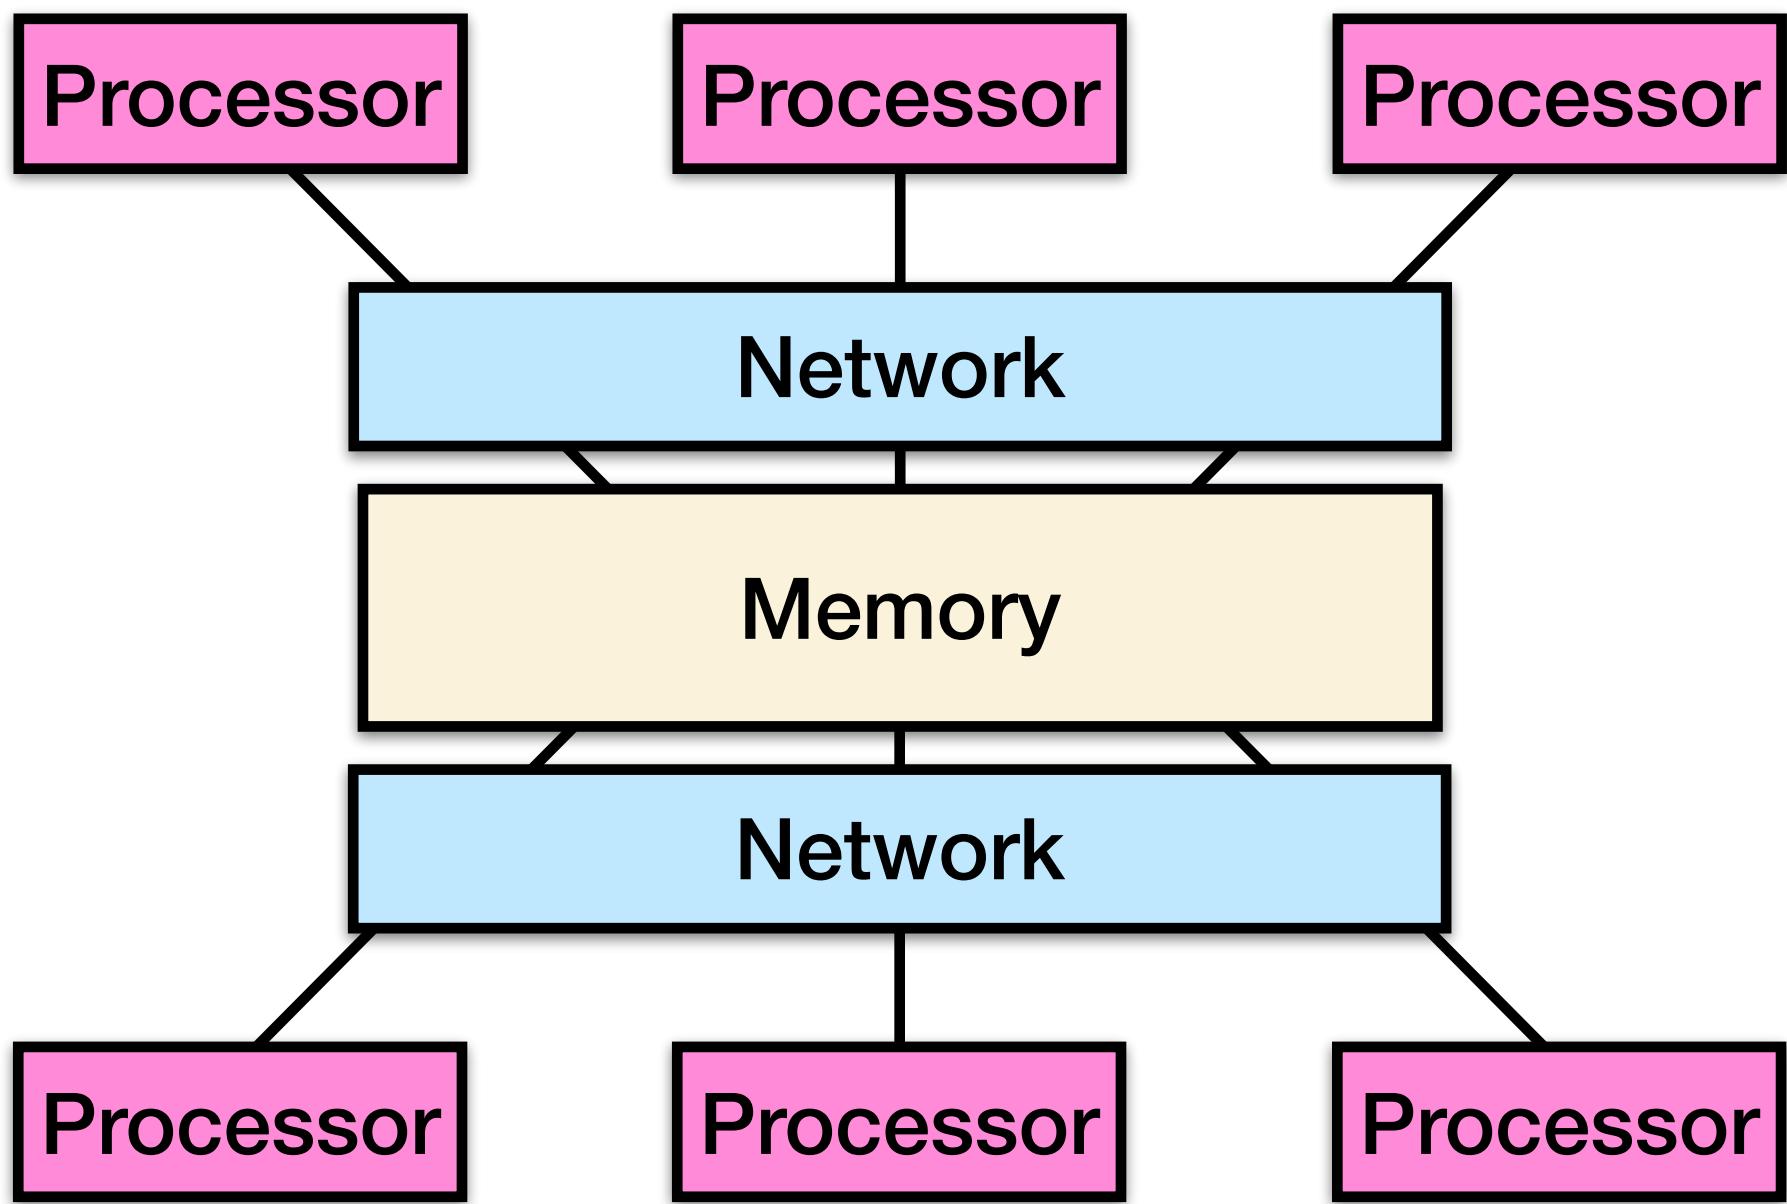
What? (Goal of the class)

Learn **parallel computing** to take advantage of modern hardware.

Why?

- Solve **BIG** problems
 - N-body simulations, molecular dynamics, multi-scale modeling, training LLMs ...
- **Time**: Improve quality of answer given limits of human patience
 - climate predictions
- **Space**: Solve memory-intensive problems
 - Big data analytics, network science, astrophysics, biology, ...

What is a parallel computer?

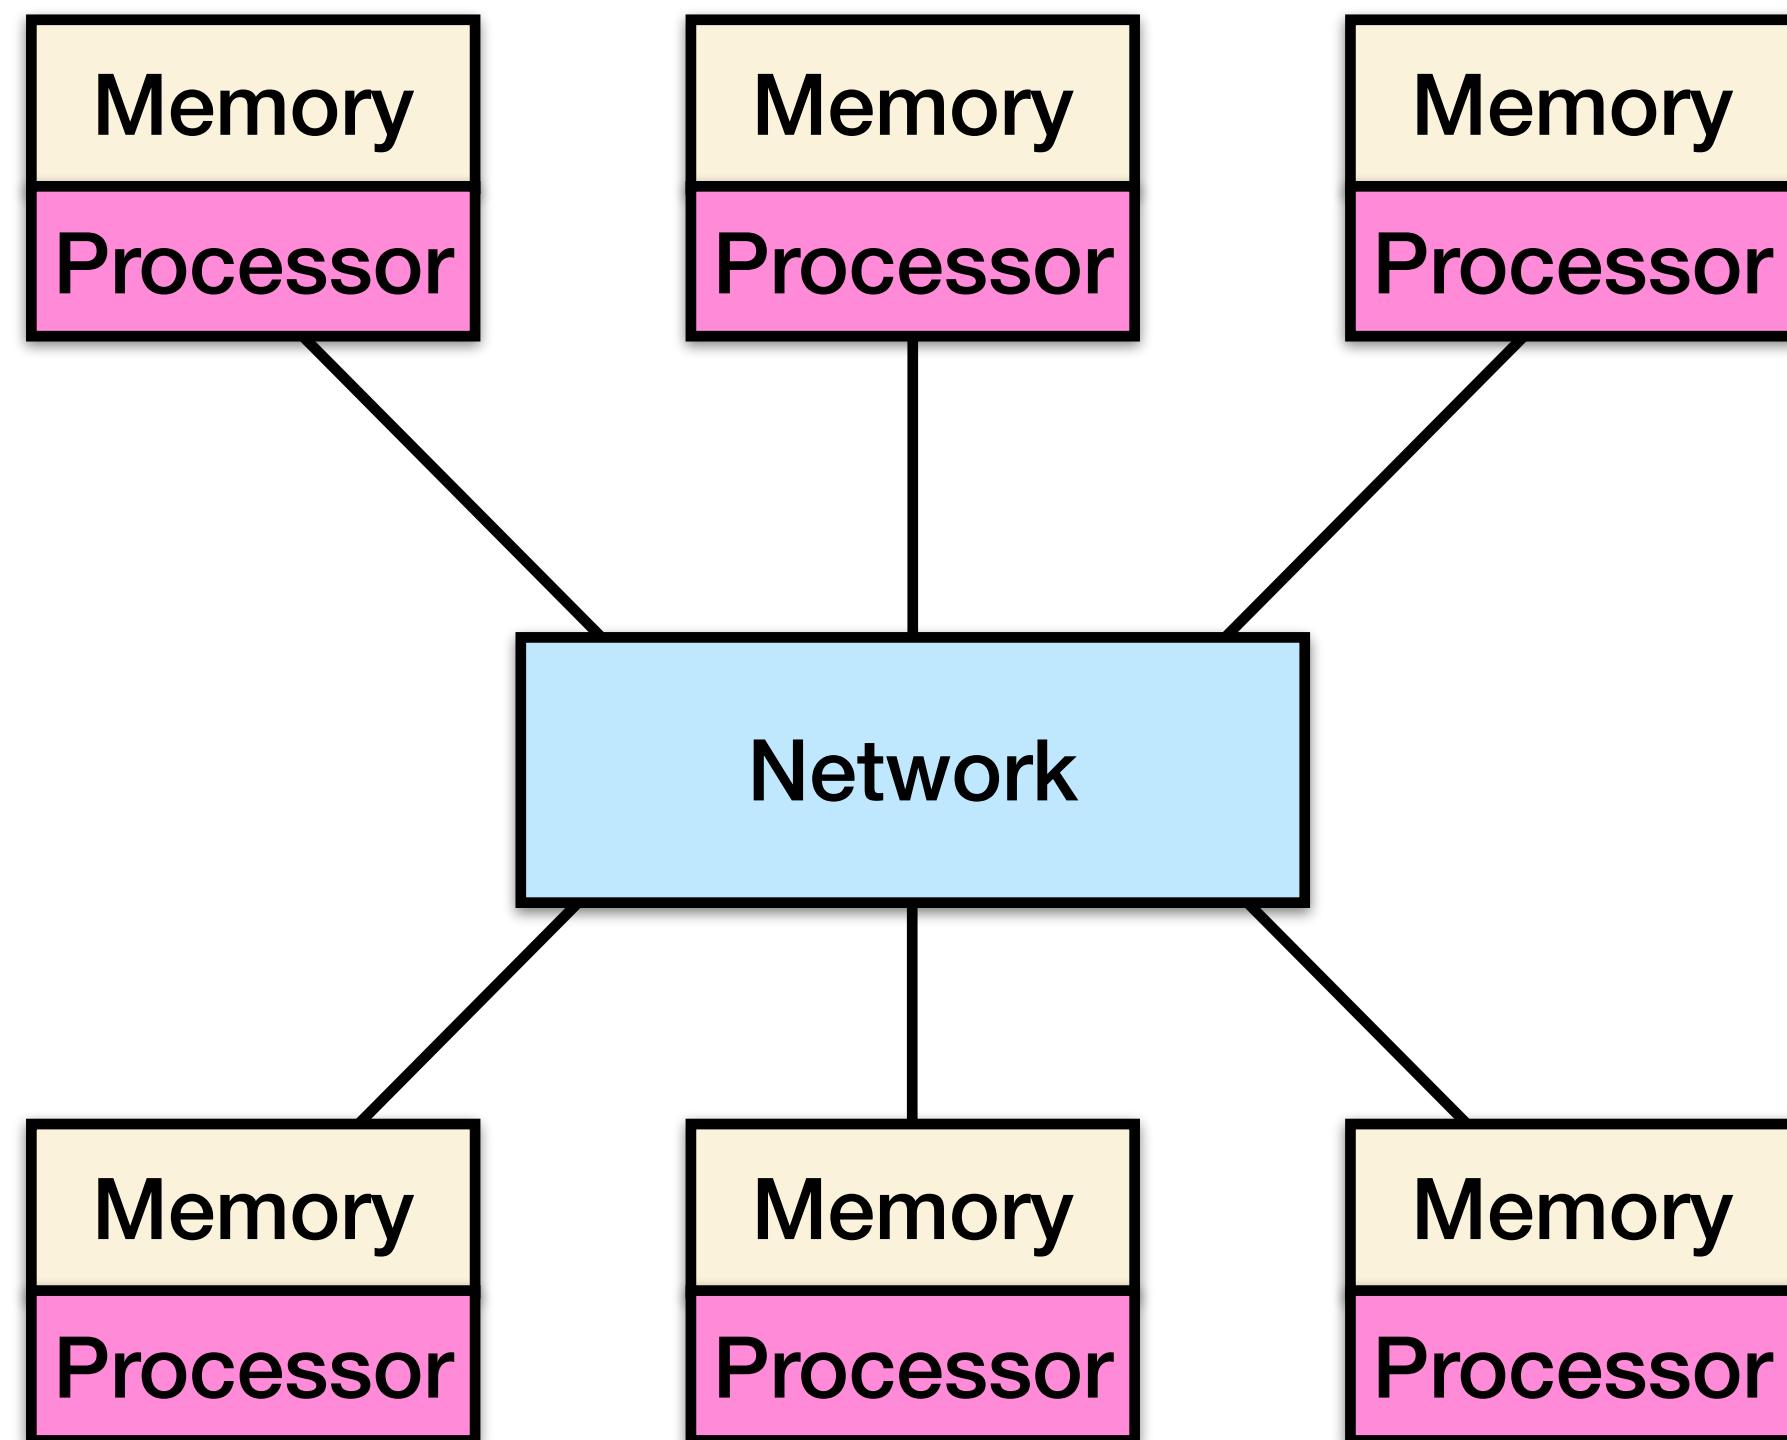


E.g. Intel Haswell

A **shared-memory multiprocessor** (SMP) connects multiple processors to a single memory system.

A multicore processor contains multiple processors (cores) on a single chip.

Scaling parallelism to multiple machines



A **distributed-memory multiprocessor** has processors with their own memories connected by a high-speed network.

Also called a **cluster**.

A **high-performance computing** (HPC) system contains 100s or 1000s of such processors (nodes).

The fastest computers (for science) have been parallel for a long time



<https://www.ornl.gov/news/ornl-celebrates-launch-frontier-worlds-fastest-supercomputer>

Fastest computers in the world:
top500.org

Currently, the largest supercomputer (on the top 500) is Frontier (at Oak Ridge National Laboratory).

~ 8.7M cores, 22.7 kW

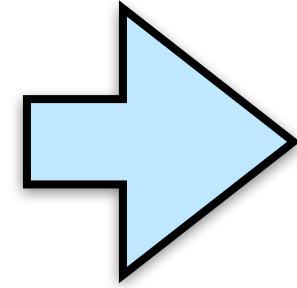
The first **exascale** computer - over a quintillion (a billion billion) calculations per second.

Units of measure for HPC

High-performance computing (HPC) **units** include:

- **Flop**: floating point operations (usually double precision unless otherwise noted)
- **Flop/s**: floating point operations per second
- **Bytes**: size of data (a double precision floating point number is 8 bytes)

Typical sizes are millions, billions, trillions...



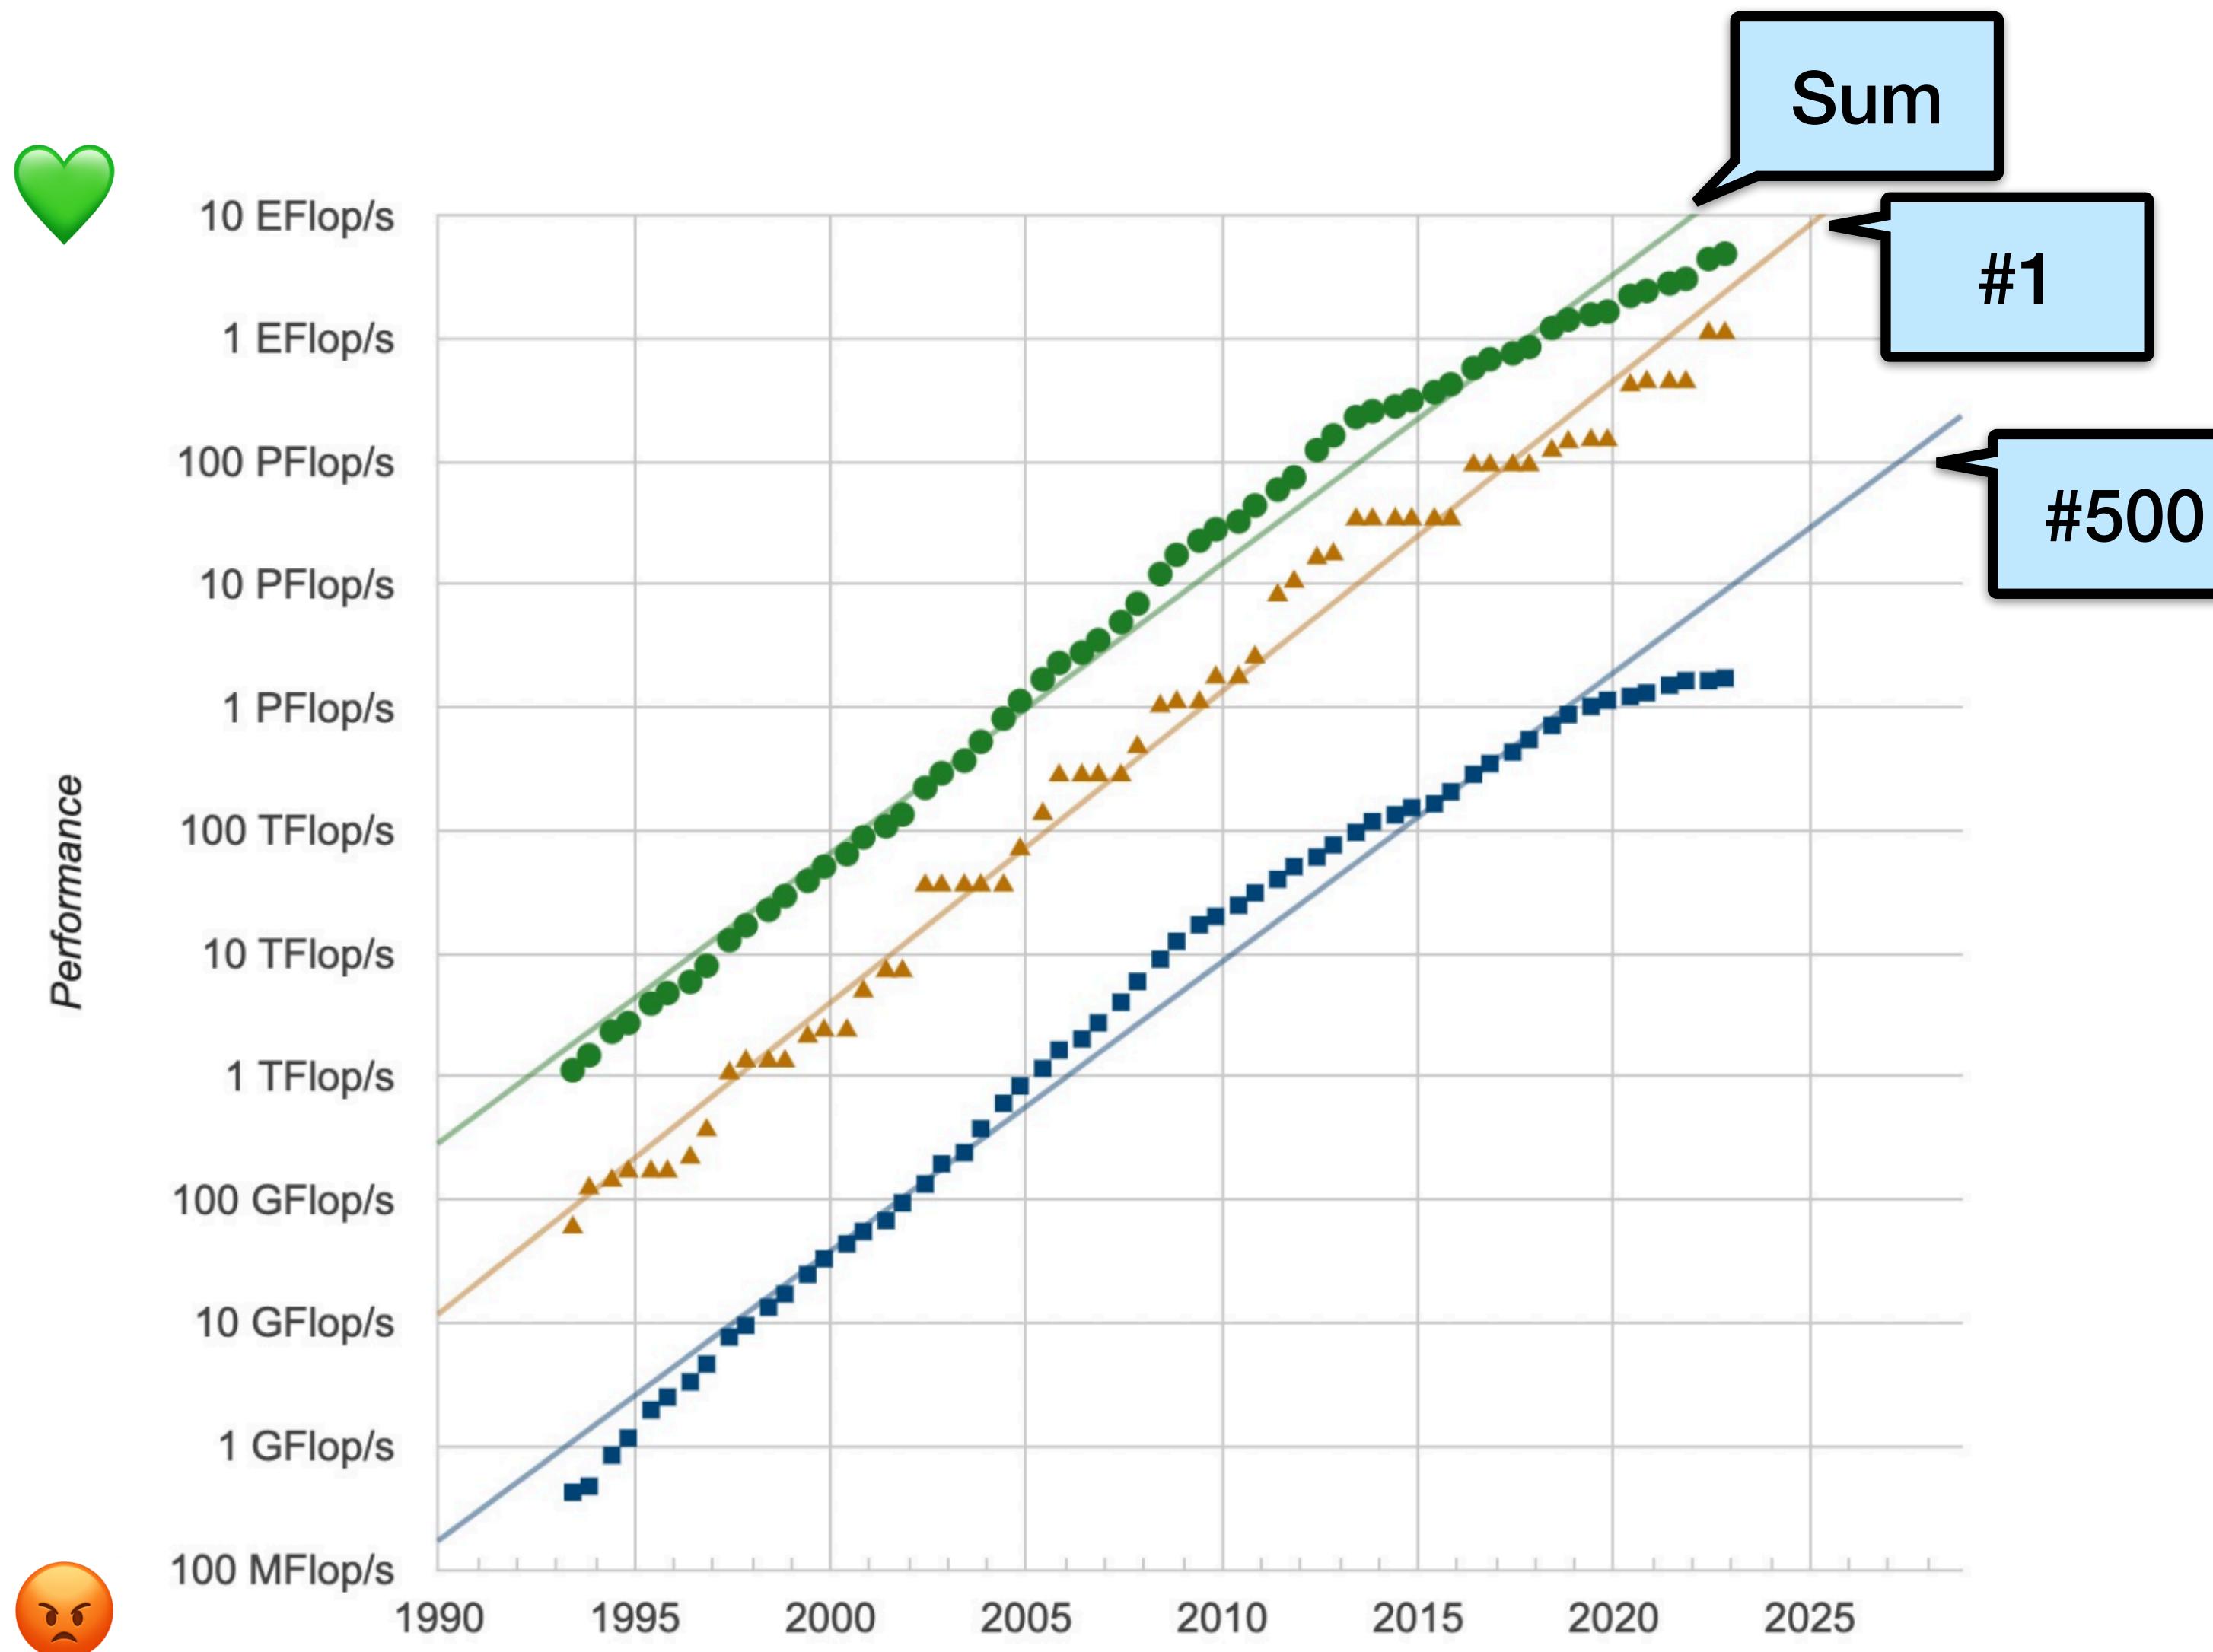
Kilo	$\text{Kflop/s} = 10^3 \text{ flop/s}$	$\text{Kbyte} = 10^3 \sim 2^{10} \text{ bytes (KiB)}$
Mega	$\text{Mflop/s} = 10^6 \text{ flop/s}$	$\text{Mbyte} = 10^6 \sim 2^{20} \text{ bytes (MiB)}$
Giga	$\text{Gflop/s} = 10^9 \text{ flop/s}$	$\text{Gbyte} = 10^9 \sim 2^{30} \text{ bytes (GiB)}$
Tera	$\text{Tflop/s} = 10^{12} \text{ flop/s}$	$\text{Tbyte} = 10^{12} \sim 2^{40} \text{ bytes (TiB)}$
Peta	$\text{Pflop/s} = 10^{15} \text{ flop/s}$	$\text{Pbyte} = 10^{15} \sim 2^{50} \text{ bytes (PiB)}$
Exa	$\text{Eflop/s} = 10^{18} \text{ flop/s}$	$\text{Ebyte} = 10^{18} \sim 2^{60} \text{ bytes (EiB)}$
Zetta	$\text{Zflop/s} = 10^{21} \text{ flop/s}$	$\text{Zbyte} = 10^{21} \sim 2^{70} \text{ bytes (ZiB)}$
Yotta	$\text{Yflop/s} = 10^{24} \text{ flop/s}$	$\text{Ybyte} = 10^{24} \sim 2^{70} \text{ bytes (YiB)}$



Nov 2023

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	Frontier - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,699,904	1,194.00	1,679.82	22,703
2	Aurora - HPE Cray EX - Intel Exascale Compute Blade, Xeon CPU Max 9470 52C 2.4GHz, Intel Data Center GPU Max, Slingshot-11, Intel DOE/SC/Argonne National Laboratory United States	4,742,808	585.34	1,059.33	24,687
3	Eagle - Microsoft NDv5, Xeon Platinum 8480C 48C 2GHz, NVIDIA H100, NVIDIA Infiniband NDR, Microsoft Microsoft Azure United States	1,123,200	561.20	846.84	
4	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
5	LUMI - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,752,704	379.70	531.51	7,107

Performance history and projection



Yardstick: floating point operations per second (flop/s)
Rmax of **LINPACK**.

Solve **$Ax=b$** , Matrix A is dense with random entries.

Dominated by **dense matrix-matrix multiply**.

Science Using High-Performance Computing

Simulation: The Third Pillar of Science

Theory

Experiment

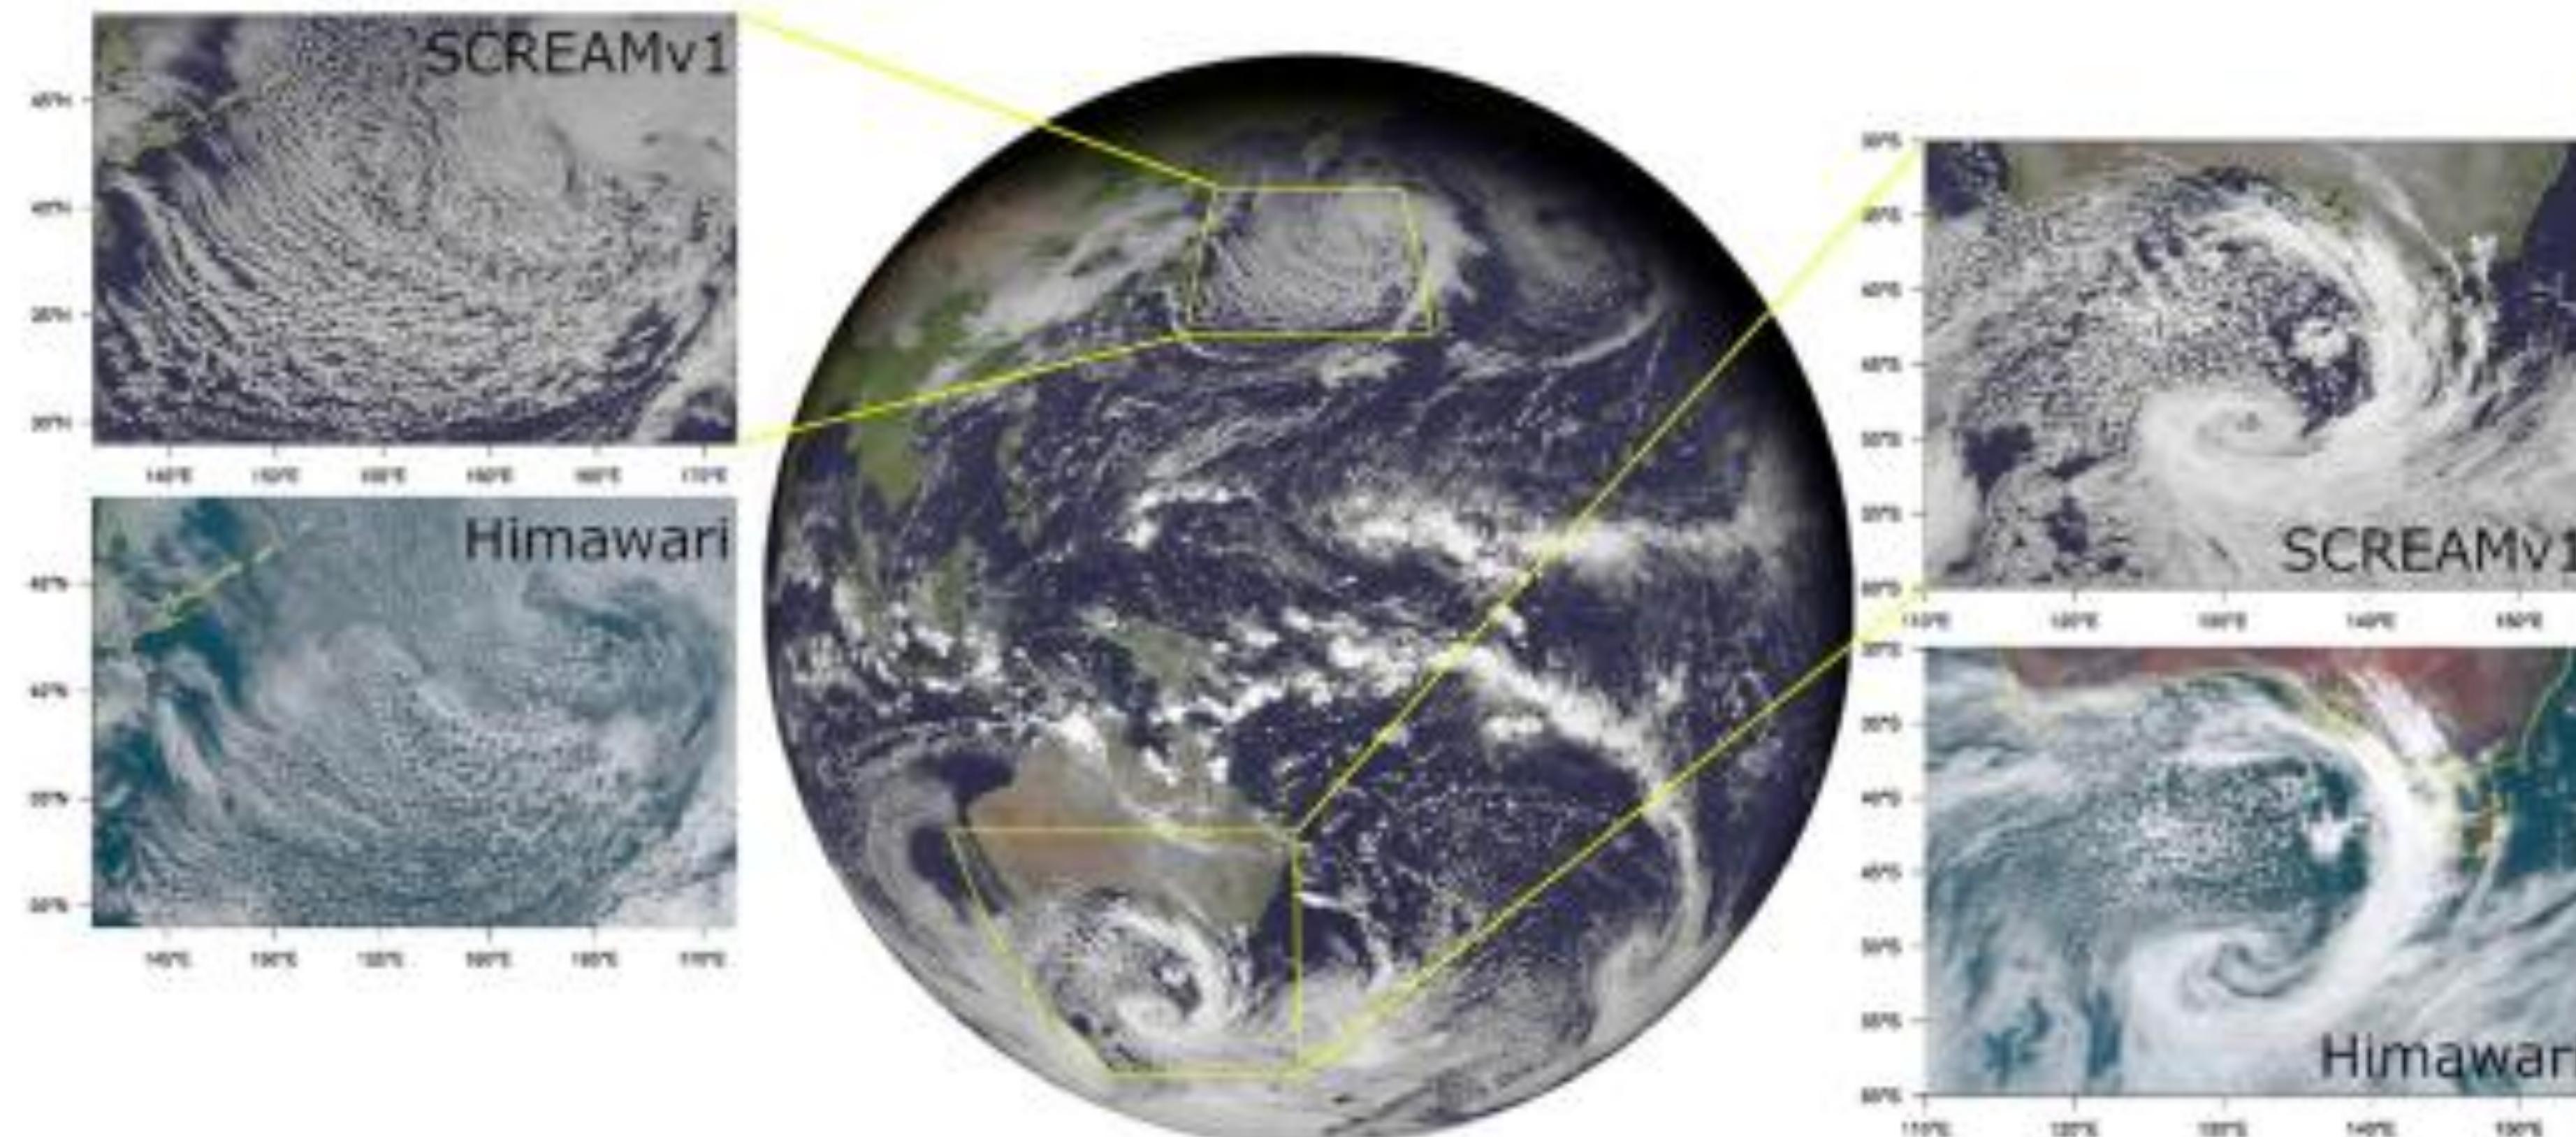
Simulation

High-performance simulation is used to understand things that are:

- too big
- too small
- too fast
- too slow
- too expensive
- too dangerous

for experiments

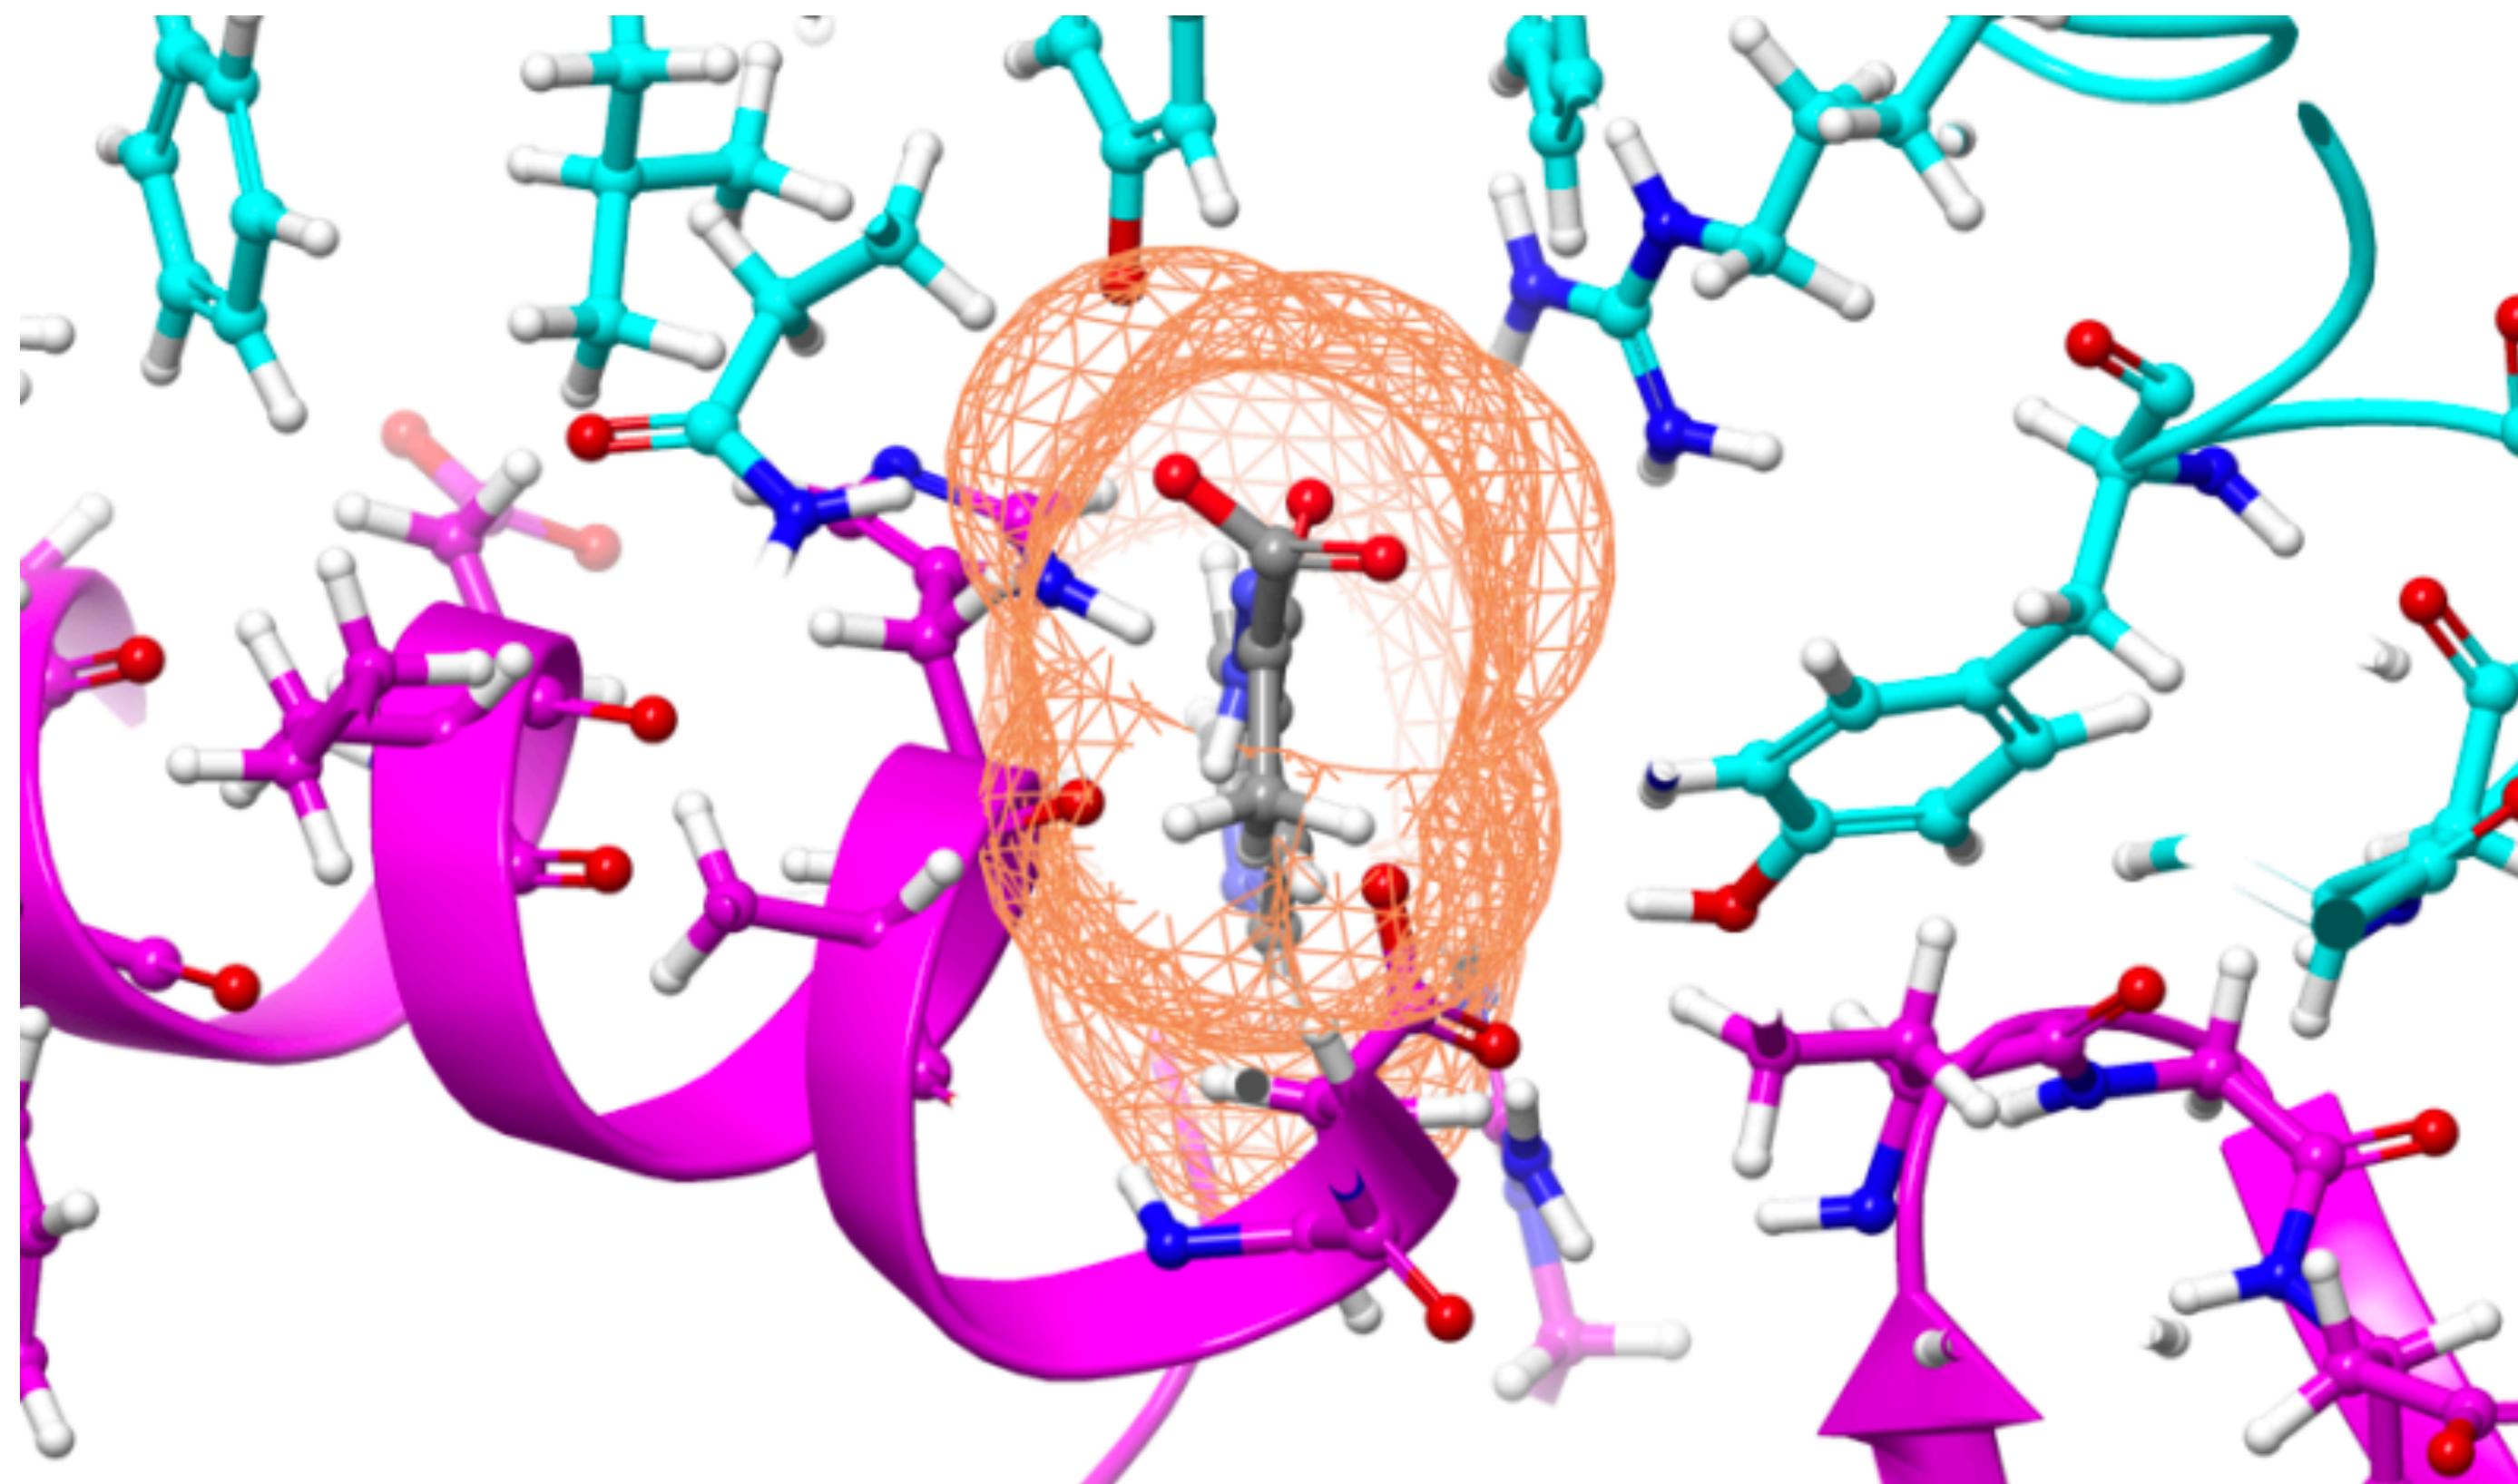
The Simple Cloud-Resolving E3SM Atmosphere Model Running on the Frontier Exascale System



<https://dl.acm.org/doi/10.1145/3581784.3627044>

<https://climatedevelopmentscience.energy.gov/news/e3sm-wins-gordon-bell-prize-climate-modeling#:~:text=On%20November%202016%2C%202023%20at,Atmosphere%20Model%20Running%20on%20the>

Screening known drugs for COVID-19



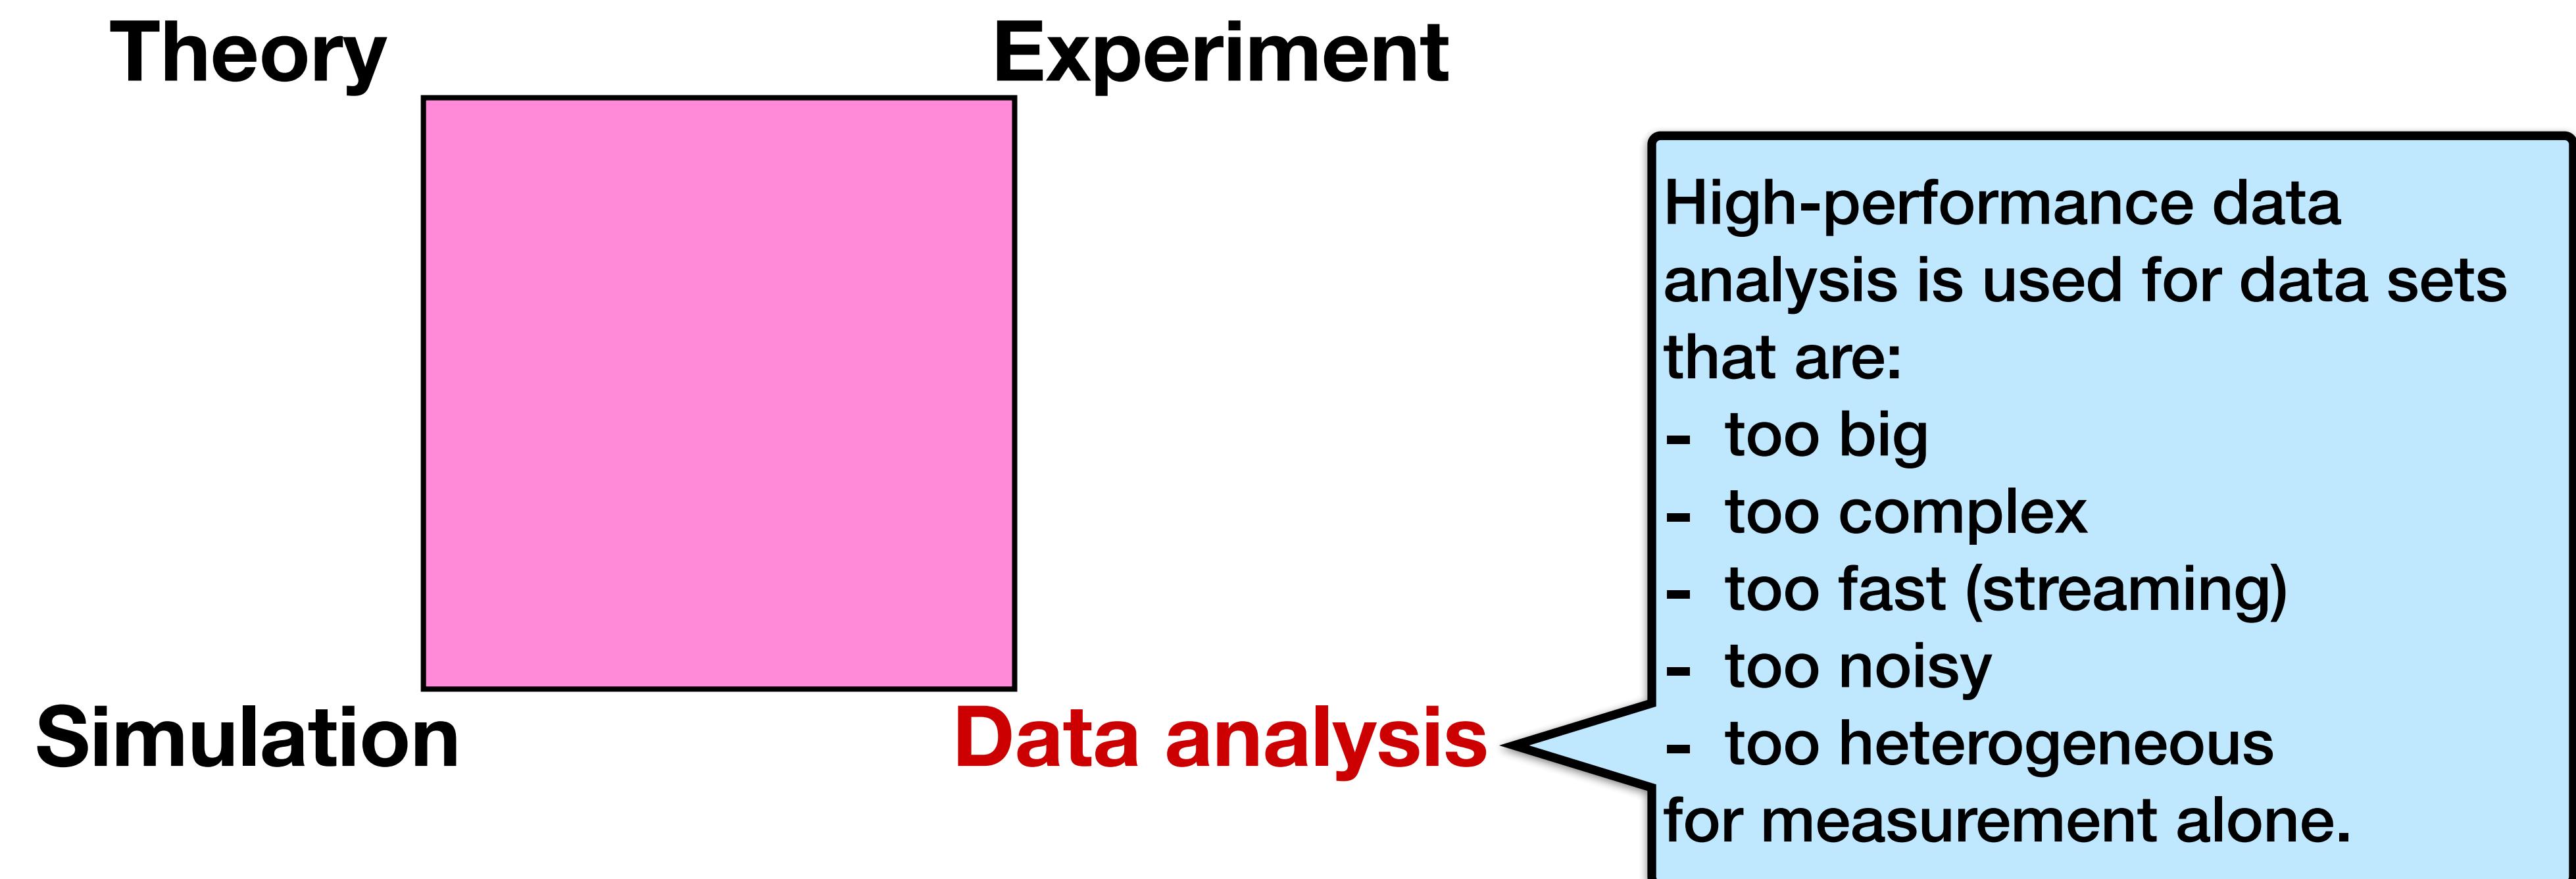
Molecular docking to
SARS-CoV-2 spike
protein

- Screened 8,000
compounds
- Identified 77 of the
most promising

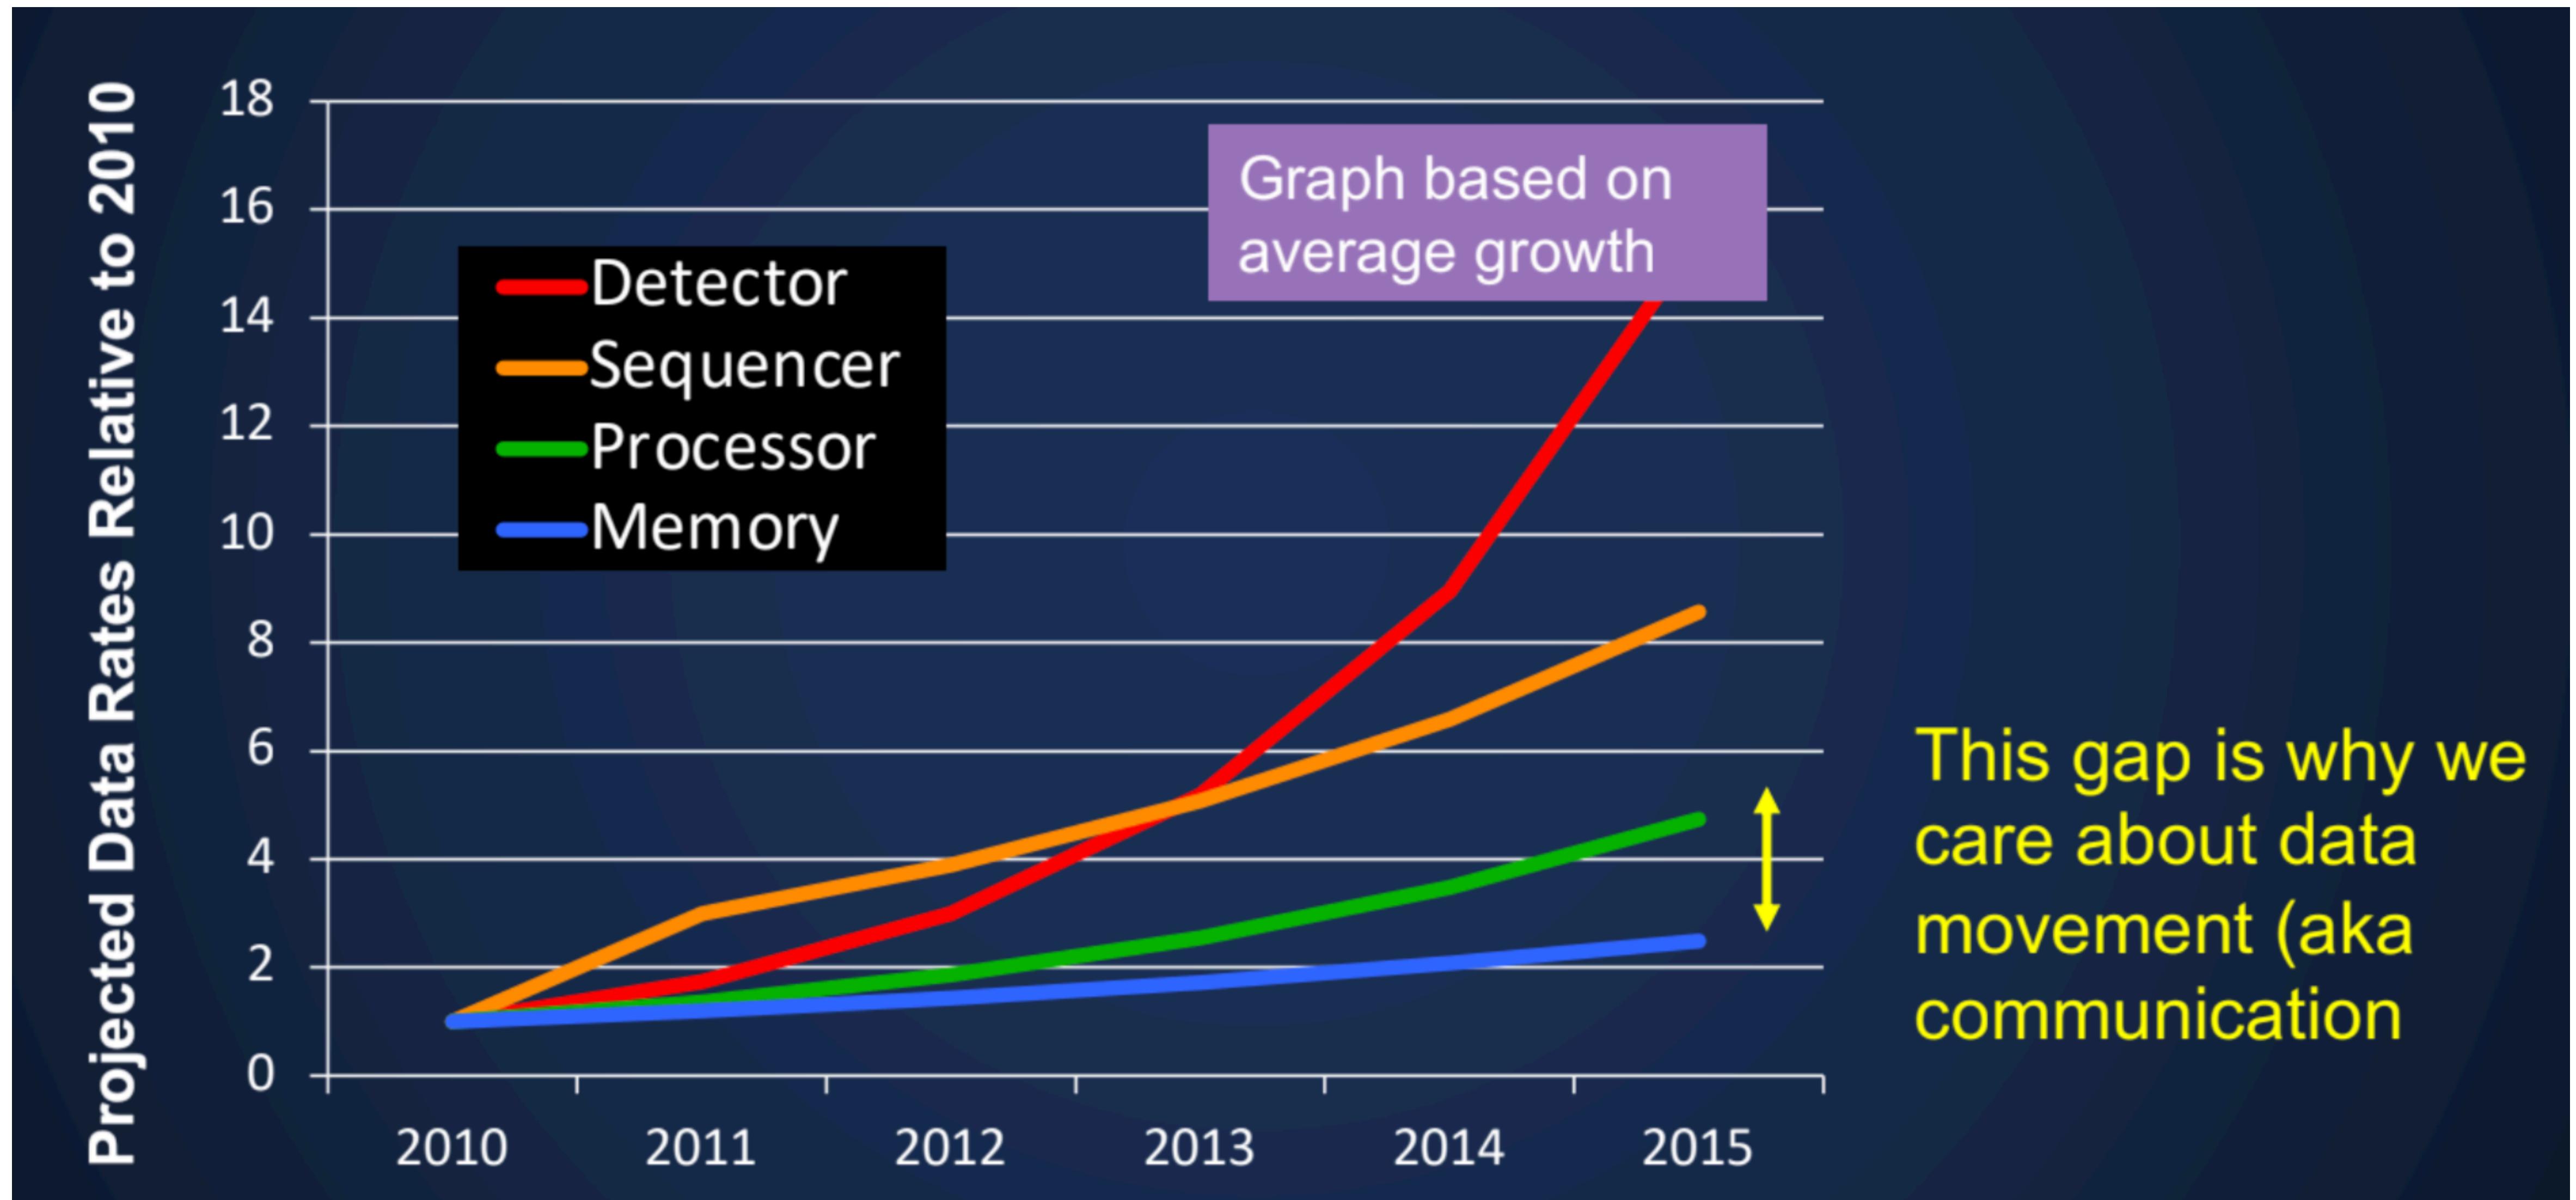
“Exascale” Scientific Applications



Data analysis: The fourth paradigm of science



Data Growth is Outpacing Computing Growth



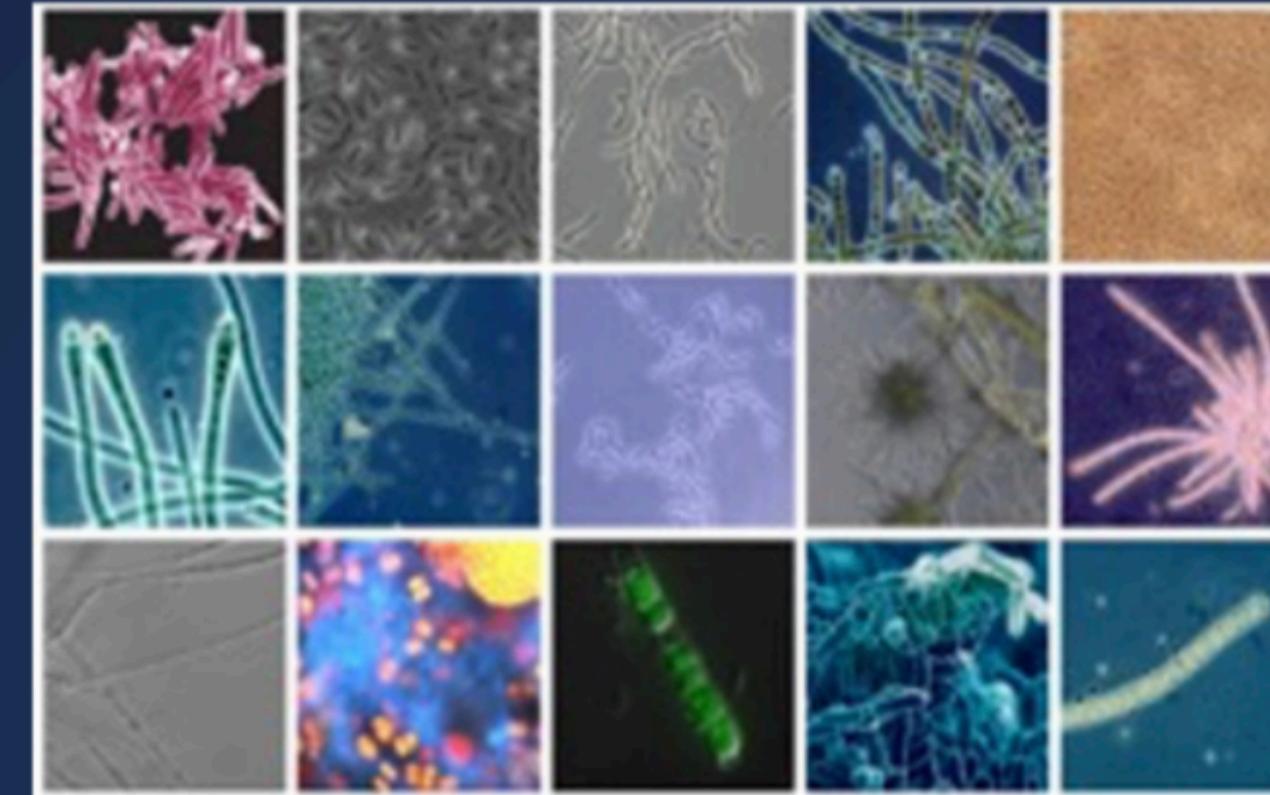
High Performance Data Analytics for Genomics



What happens to microbes after a wildfire?
(1.5TB)



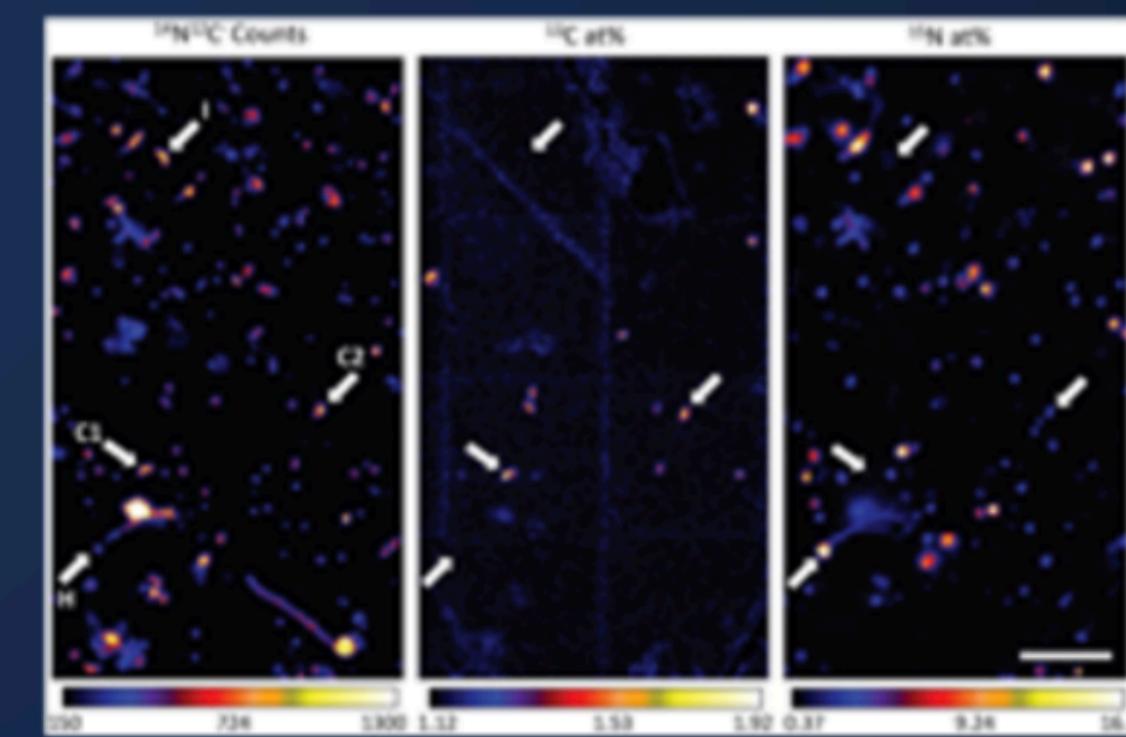
What are the seasonal fluctuations in a
wetland mangrove? (1.6 TB)



What are the microbial dynamics of
soil carbon cycling? (3.3 TB)

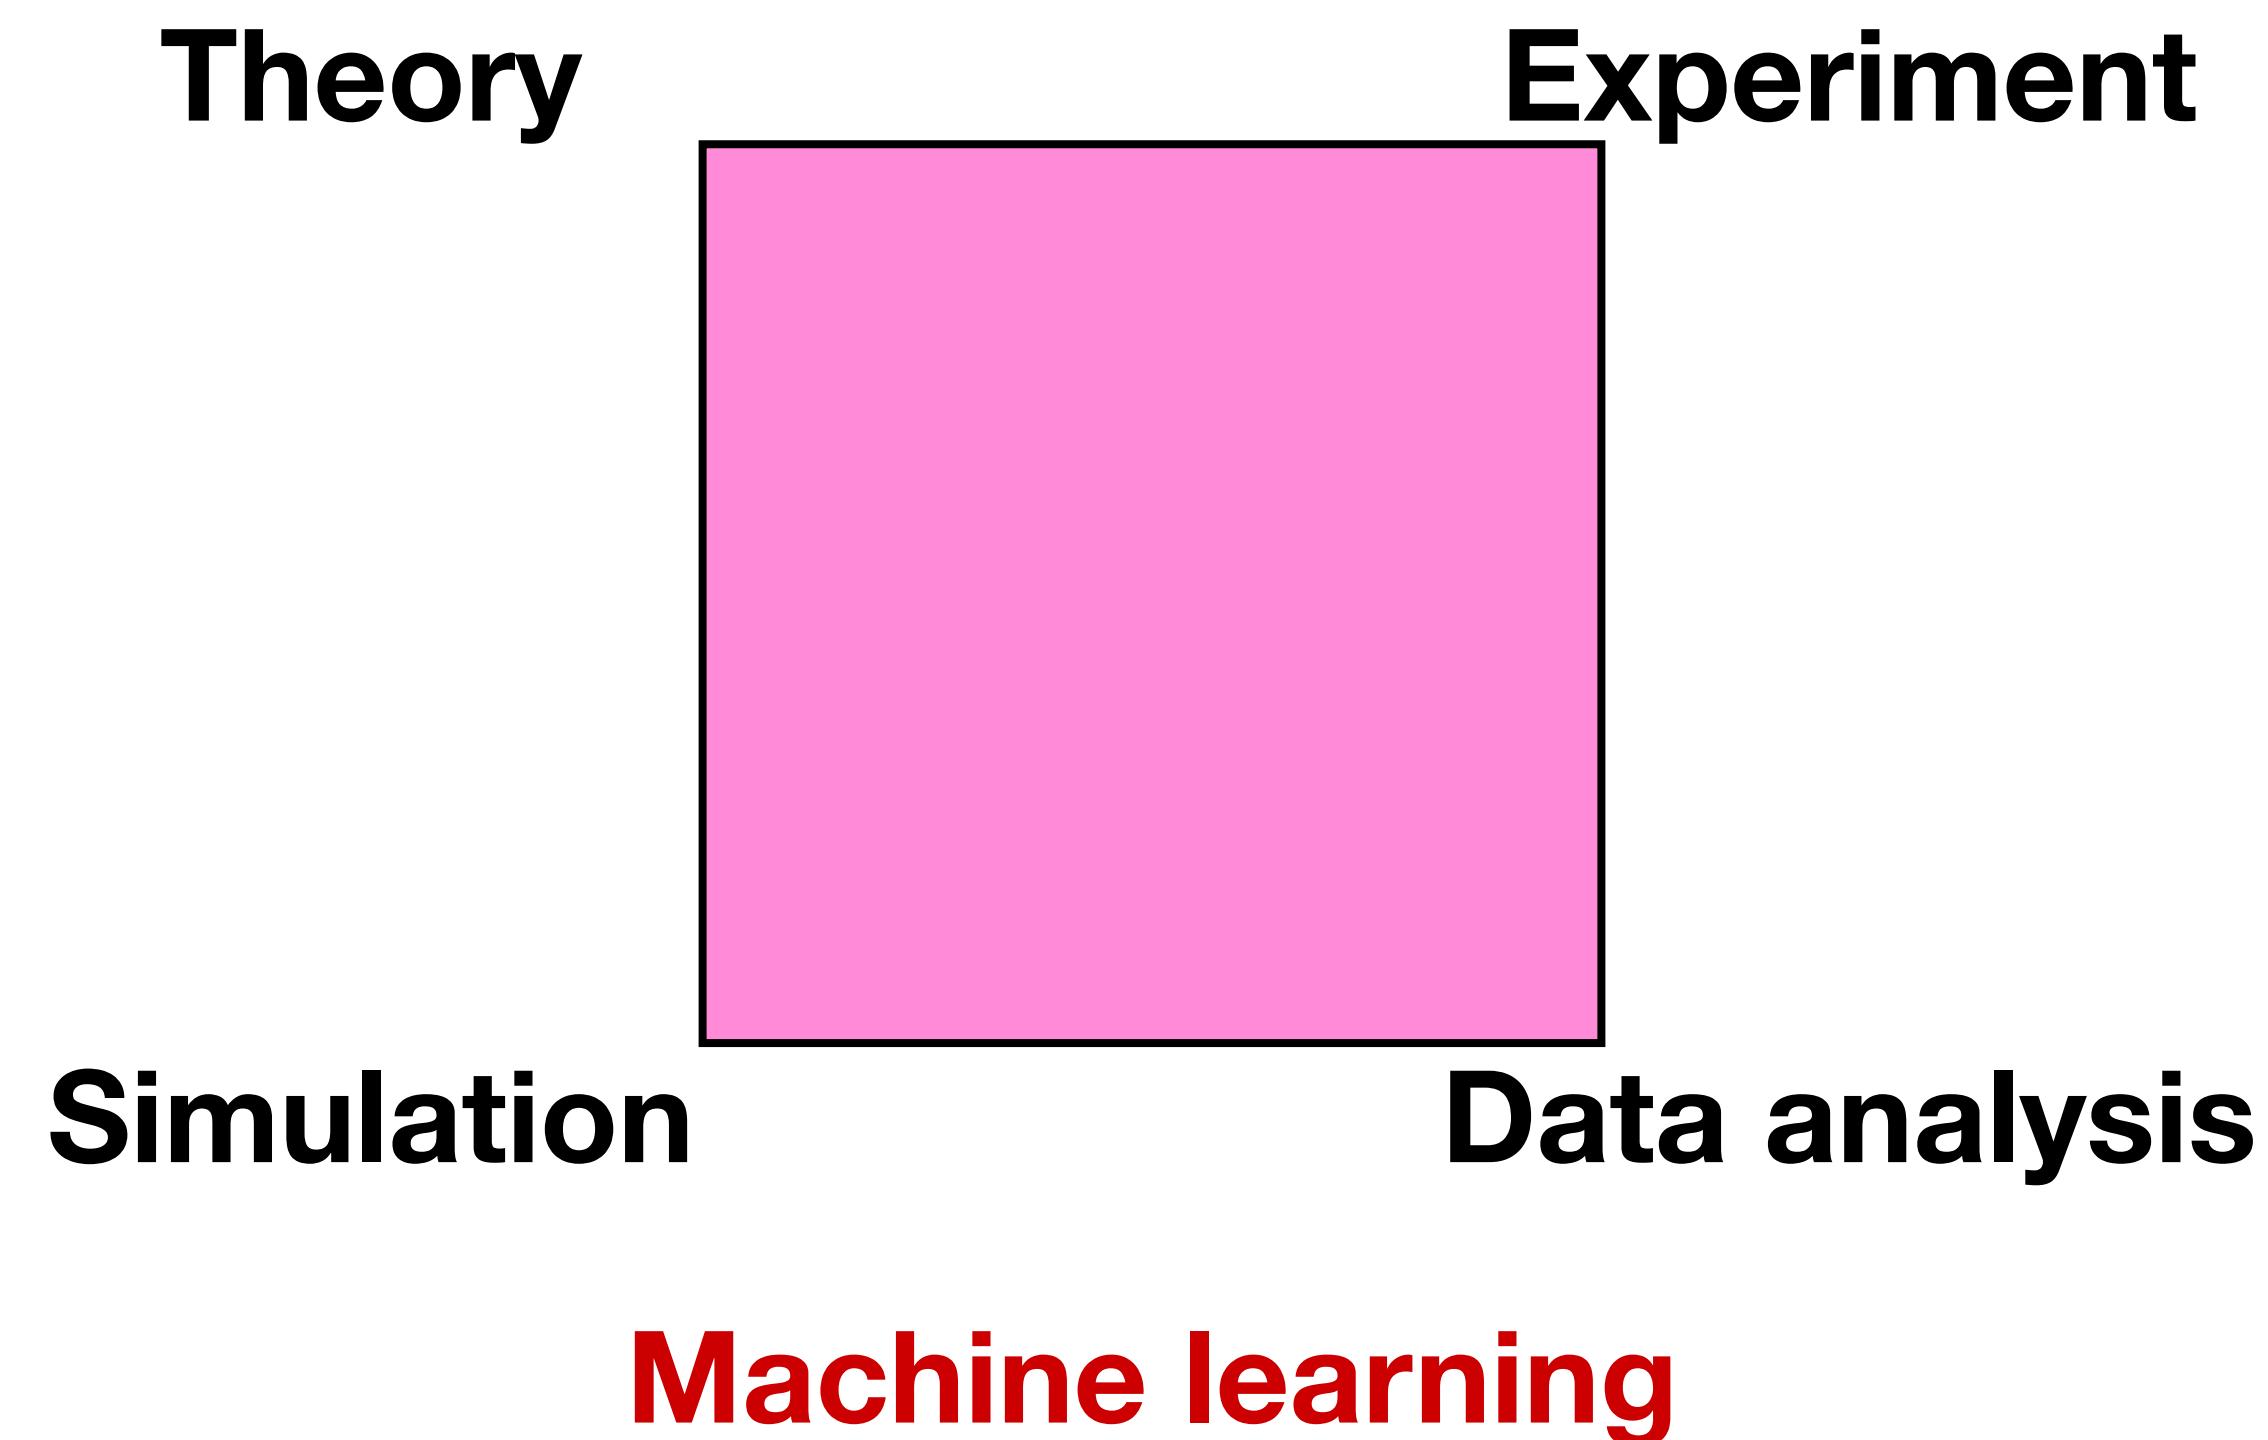


How do microbes affect disease and growth of
switchgrass for biofuels (4TB)



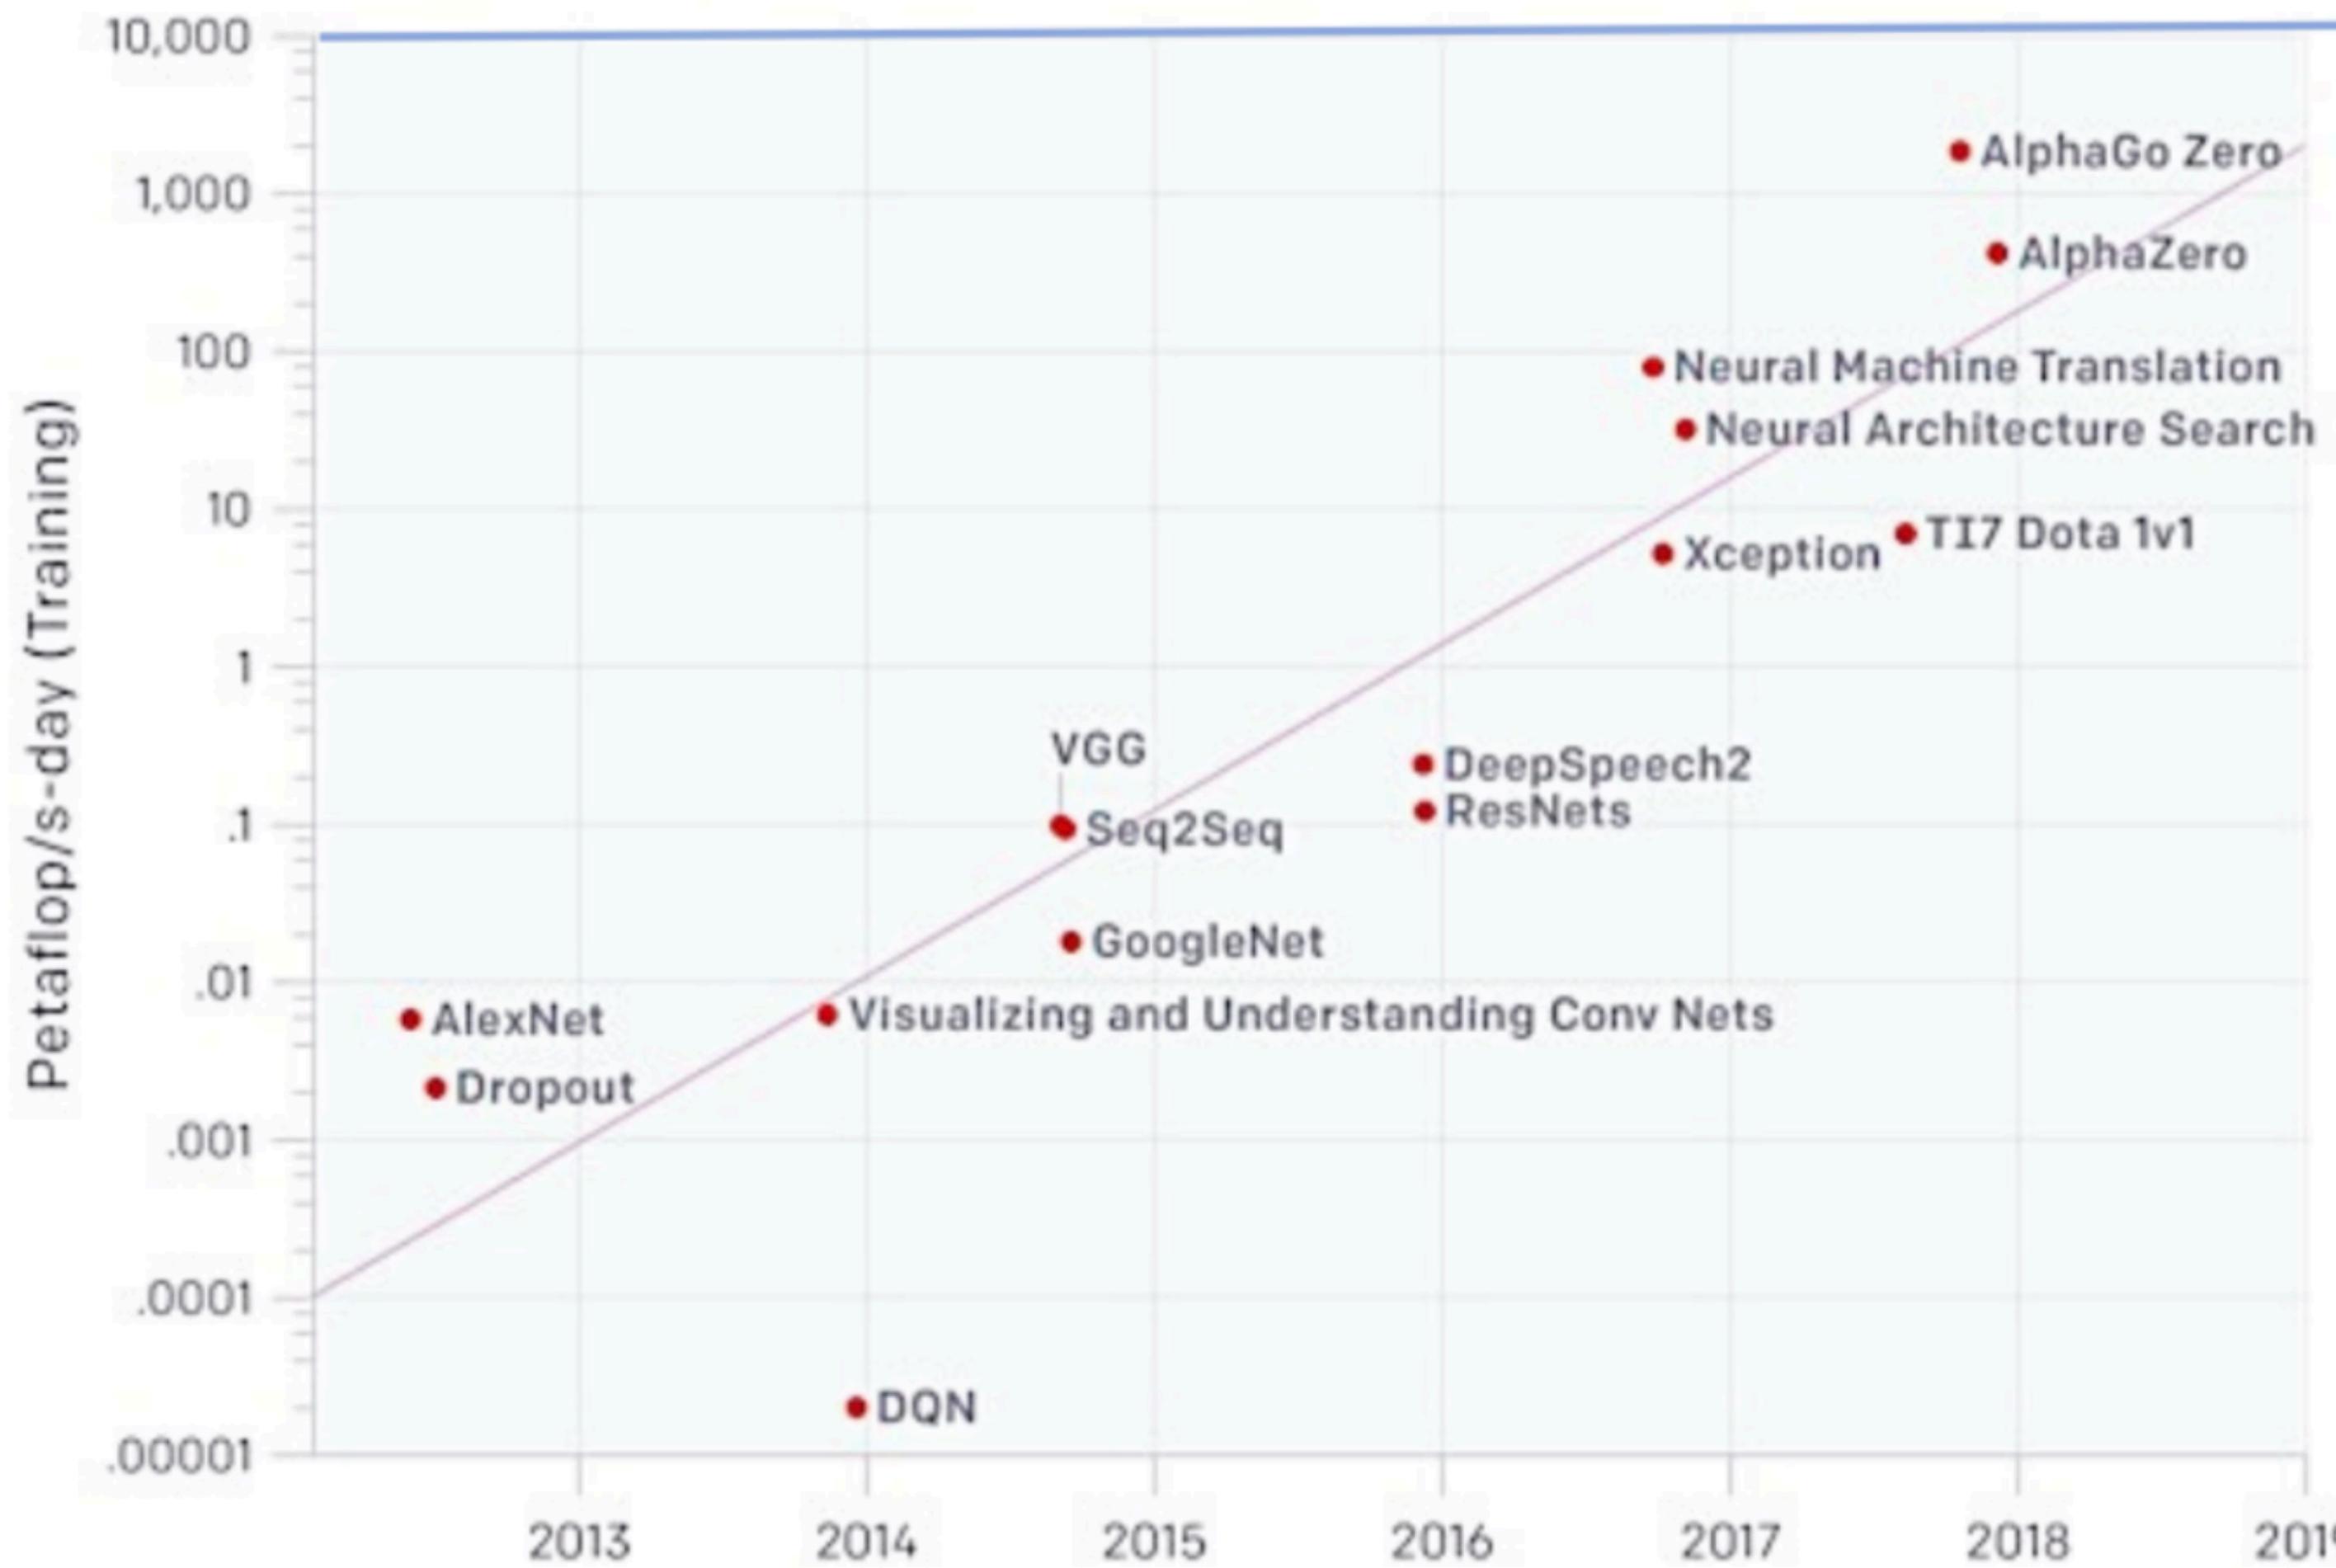
Combine genomics with isotope tracing methods for improved
functional understanding (8TB)

The fifth paradigm of science?



Machine learning demands more computing power

300,000x increase from 2011 (AlexNet) to 2018 (AlphaGoZero)



<https://blog.openai.com/ai-and-compute/>

Machine learning demands more computing power

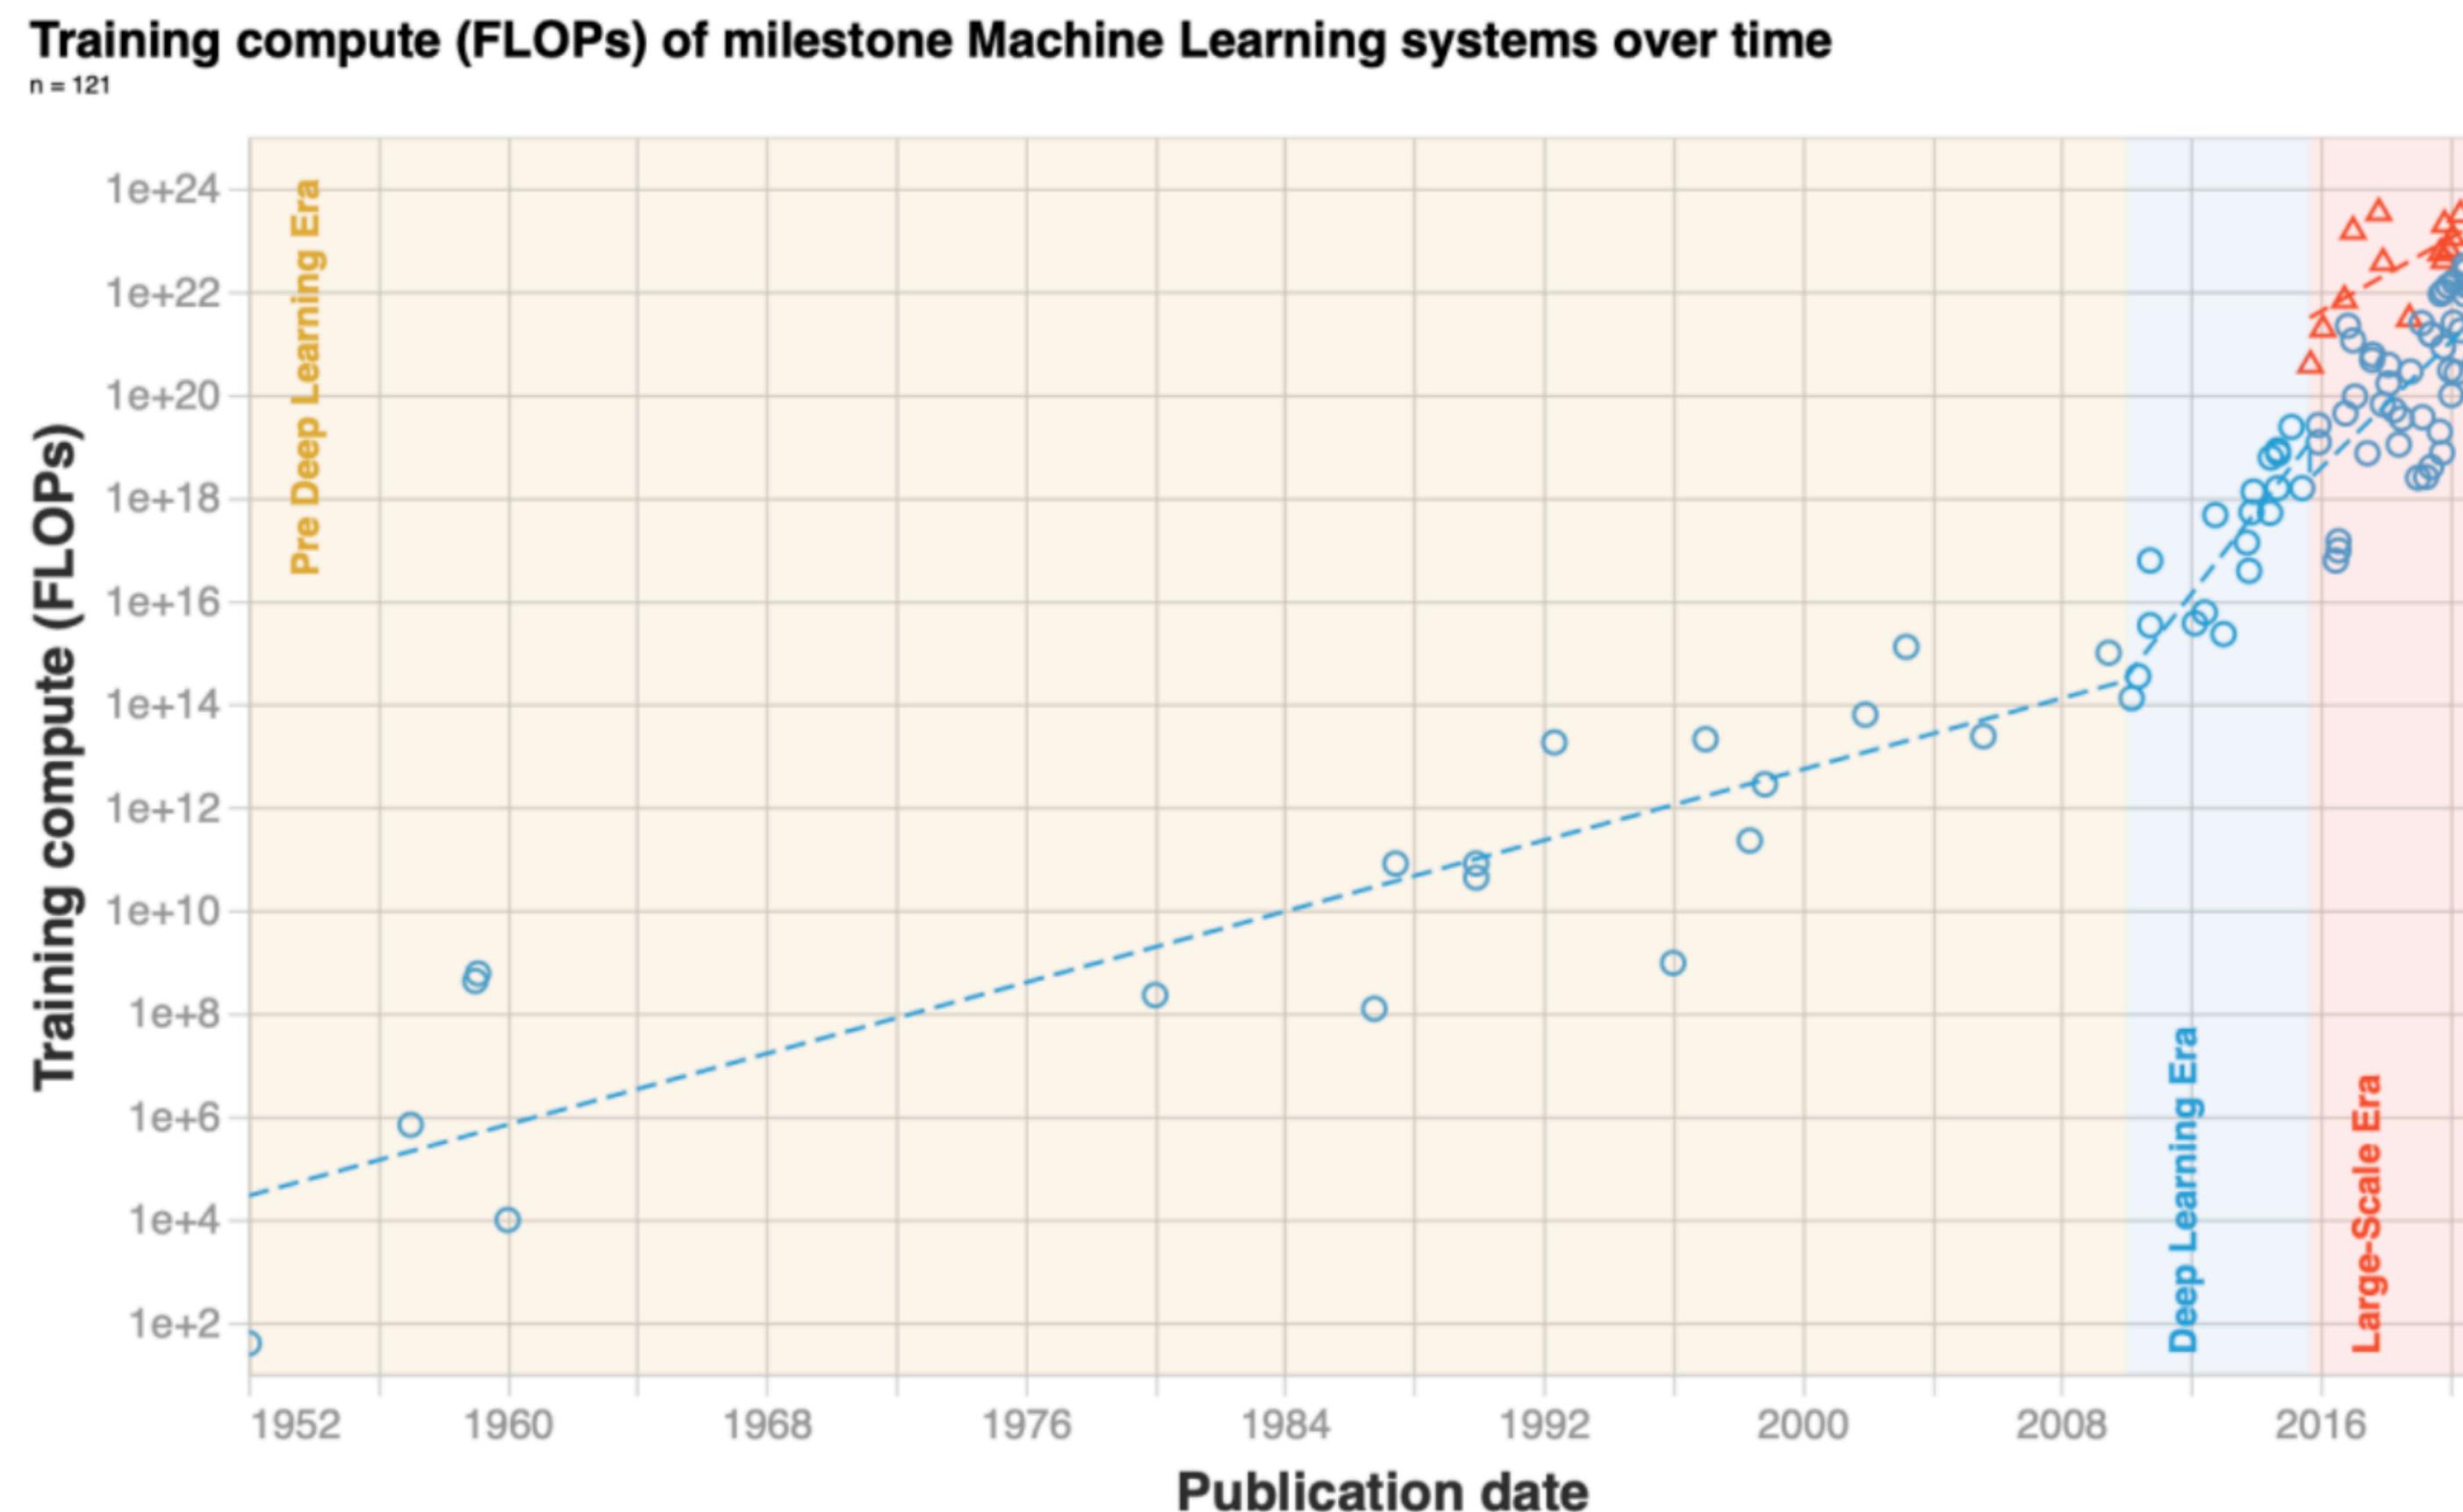


Figure 1: Trends in $n = 121$ milestone ML models between 1952 and 2022. We distinguish three eras. Notice the change of slope circa 2010, matching the advent of Deep Learning; and the emergence of a new large-scale trend in late 2015.

AlphaFold: Predicting Protein Structure via Machine Learning

AlphaFold provides a machine-learning based solution to the grand challenge of understanding the **3D structure of proteins**.

Instead of determining them experimentally (time-consuming and challenging), **AI can predict protein shape** at scale, quickly, and accurately.

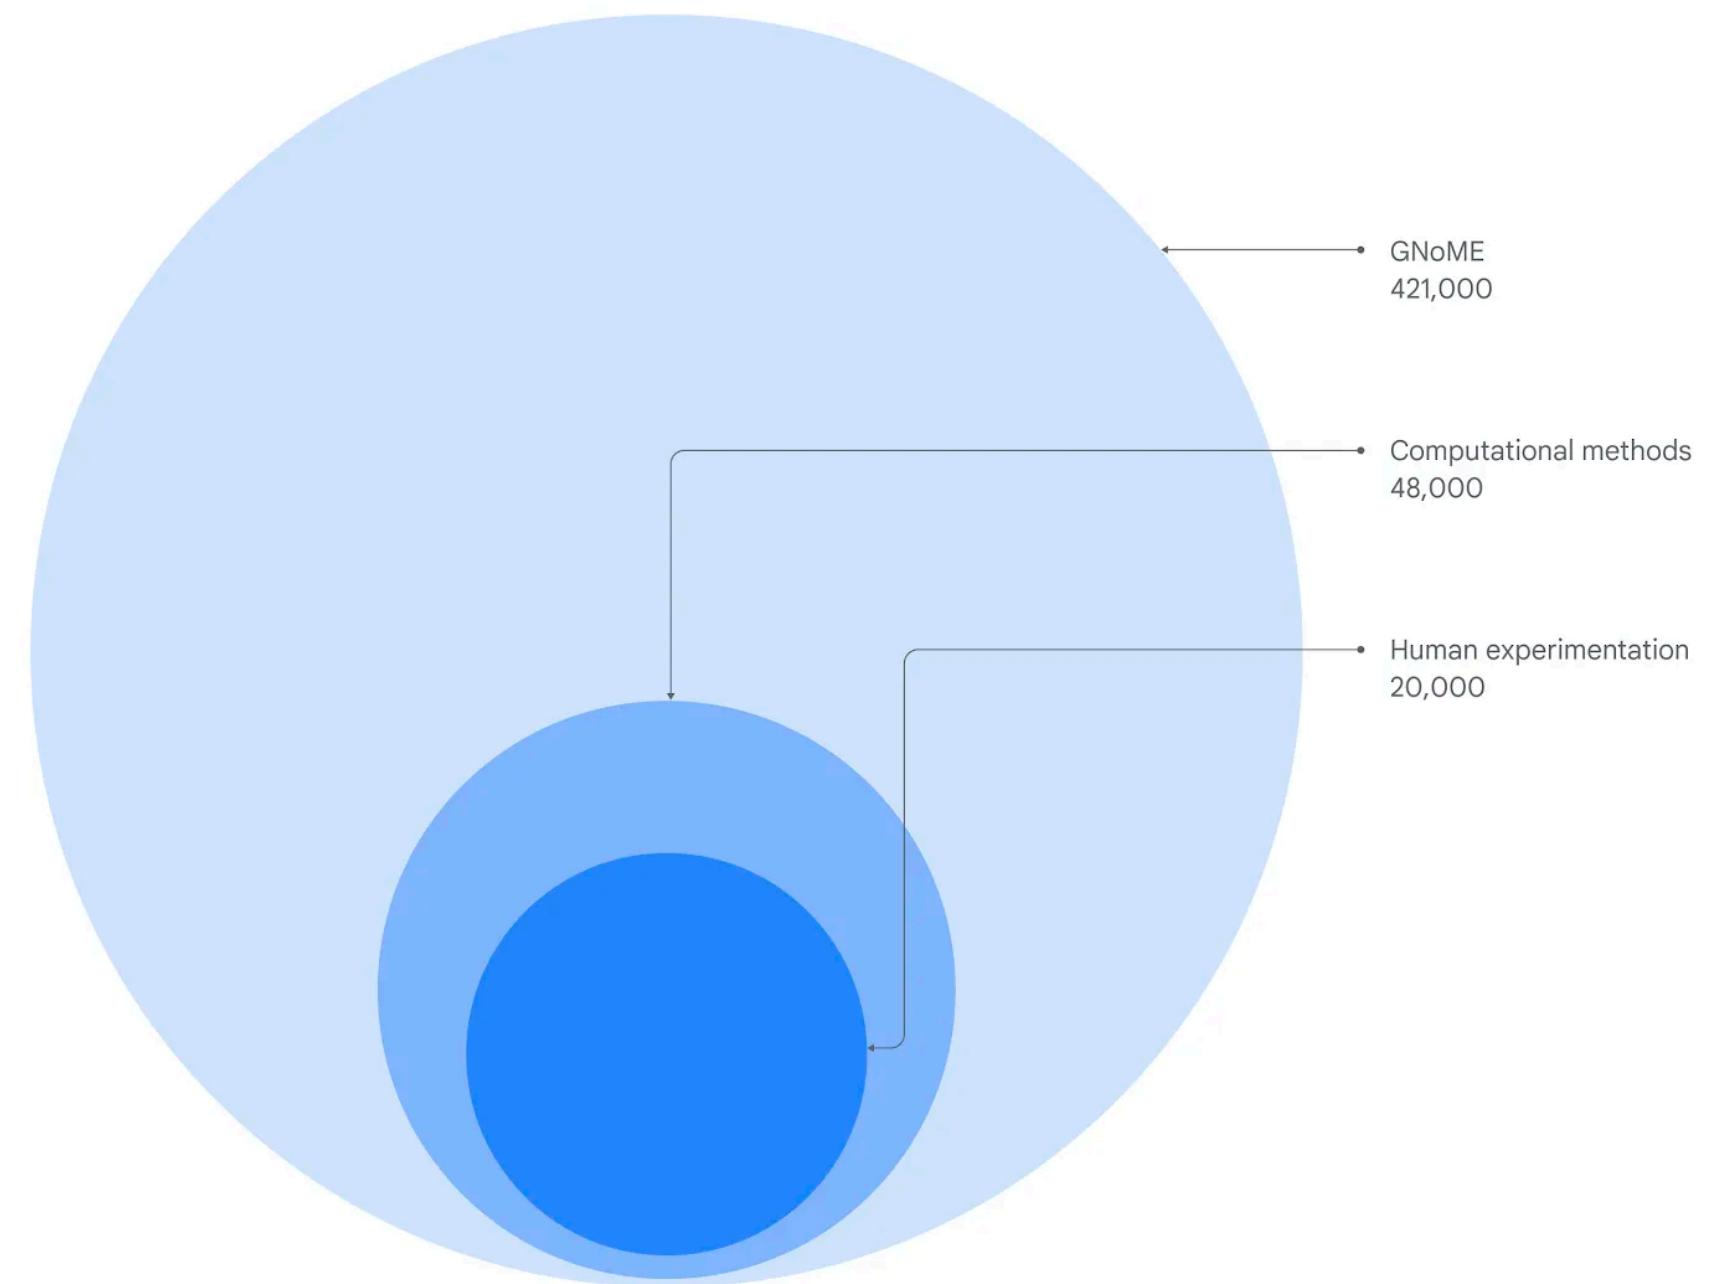


<https://deepmind.google/technologies/alphafold/>

<https://deepmind.google/discover/blog/putting-the-power-of-alphafold-into-the-worlds-hands/>

Millions of new materials discovered with deep learning

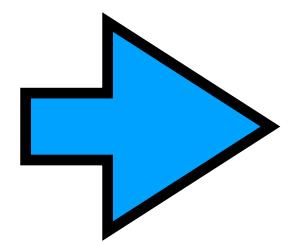
AI for materials discovery resulted in the discovery of 2.2 million new crystals - equivalent to nearly 800 years of knowledge.



Why are all computers nowadays parallel?

Motivating trends

How?



Parallelism,



Driven by?

Power,

data locality,



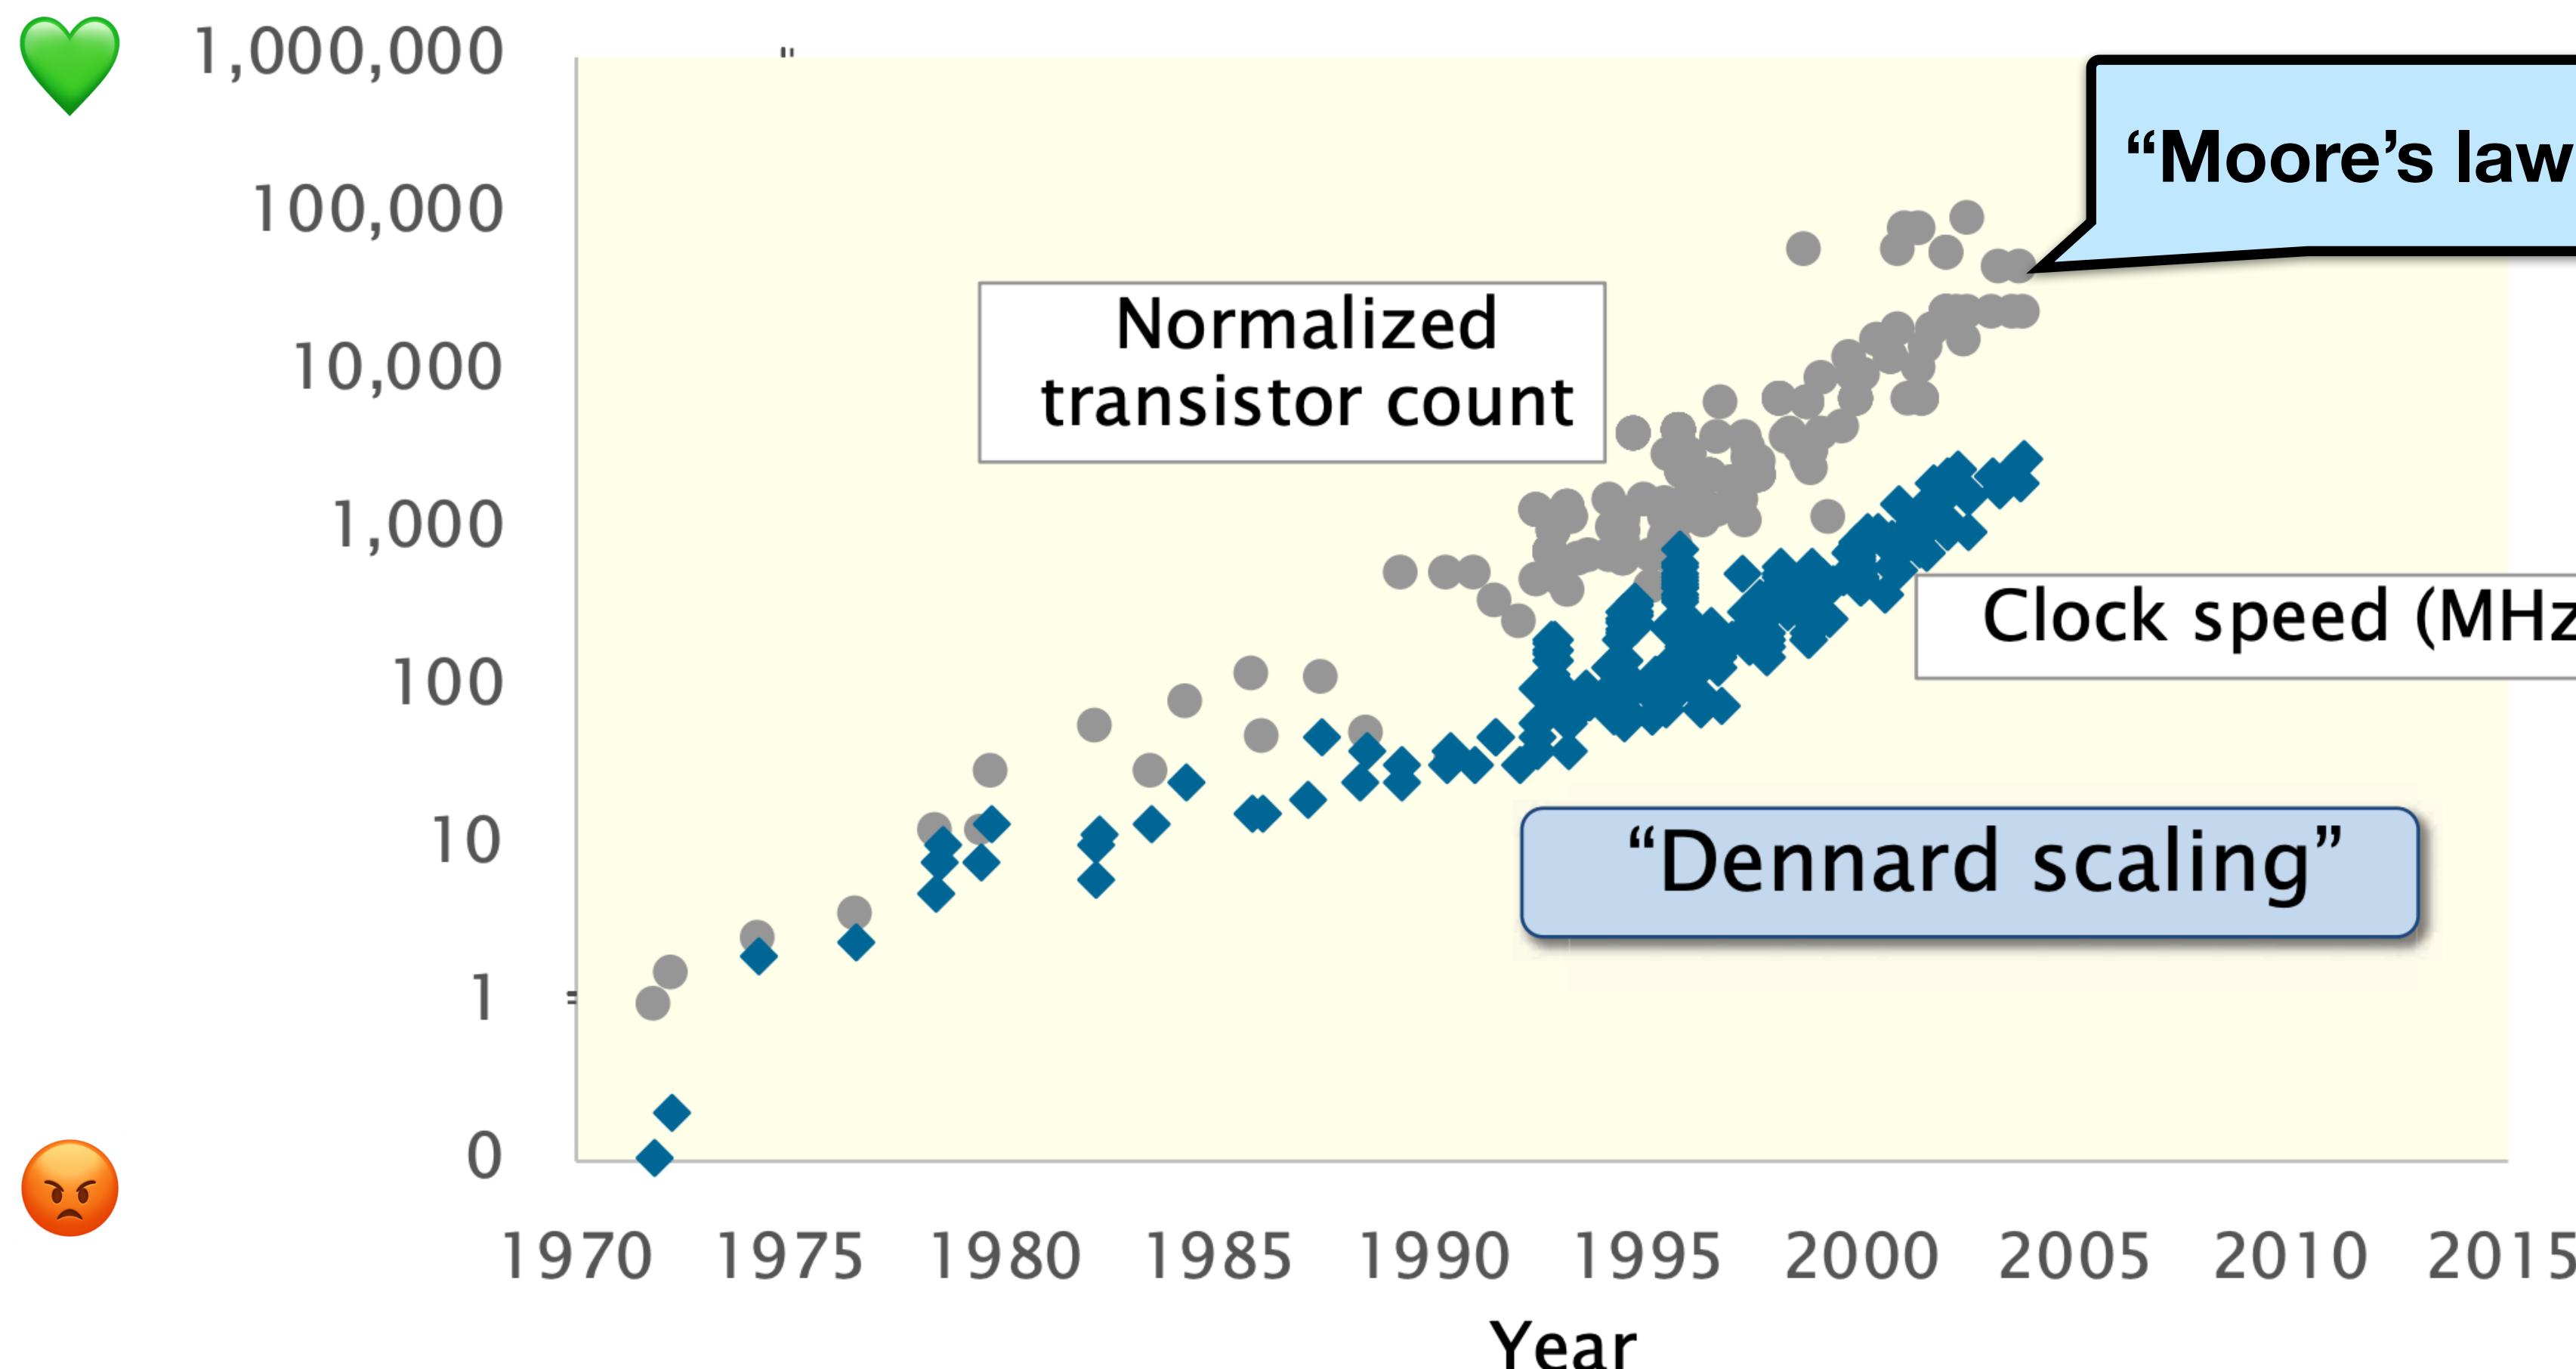
memory,

and **specialization.**



and **physics.**

Technology scaling before 2004



Processor data from Stanford's CPU DB [DKM12].

Until 2004, Moore's law and the scaling of clock frequency enabled **“free” performance gains** with underlying hardware changes.

Advances in hardware

Example: Apple computers with similar prices from 1977 to 2004



Courtesy of [mwichary](#) on Flickr.
Used under CC-BY.

Apple II

Launched: 1977
Clock rate: 1 MHz
Data path: 8 bits
Memory: 48 KB
Cost: \$1,395



Courtesy of [compuudemano](#) on Flickr. Used under CC-BY.

Power Macintosh G4

Launched: 2000
Clock rate: 400 MHz
Data path: 32 bits
Memory: 64 MB
Cost: \$1,599

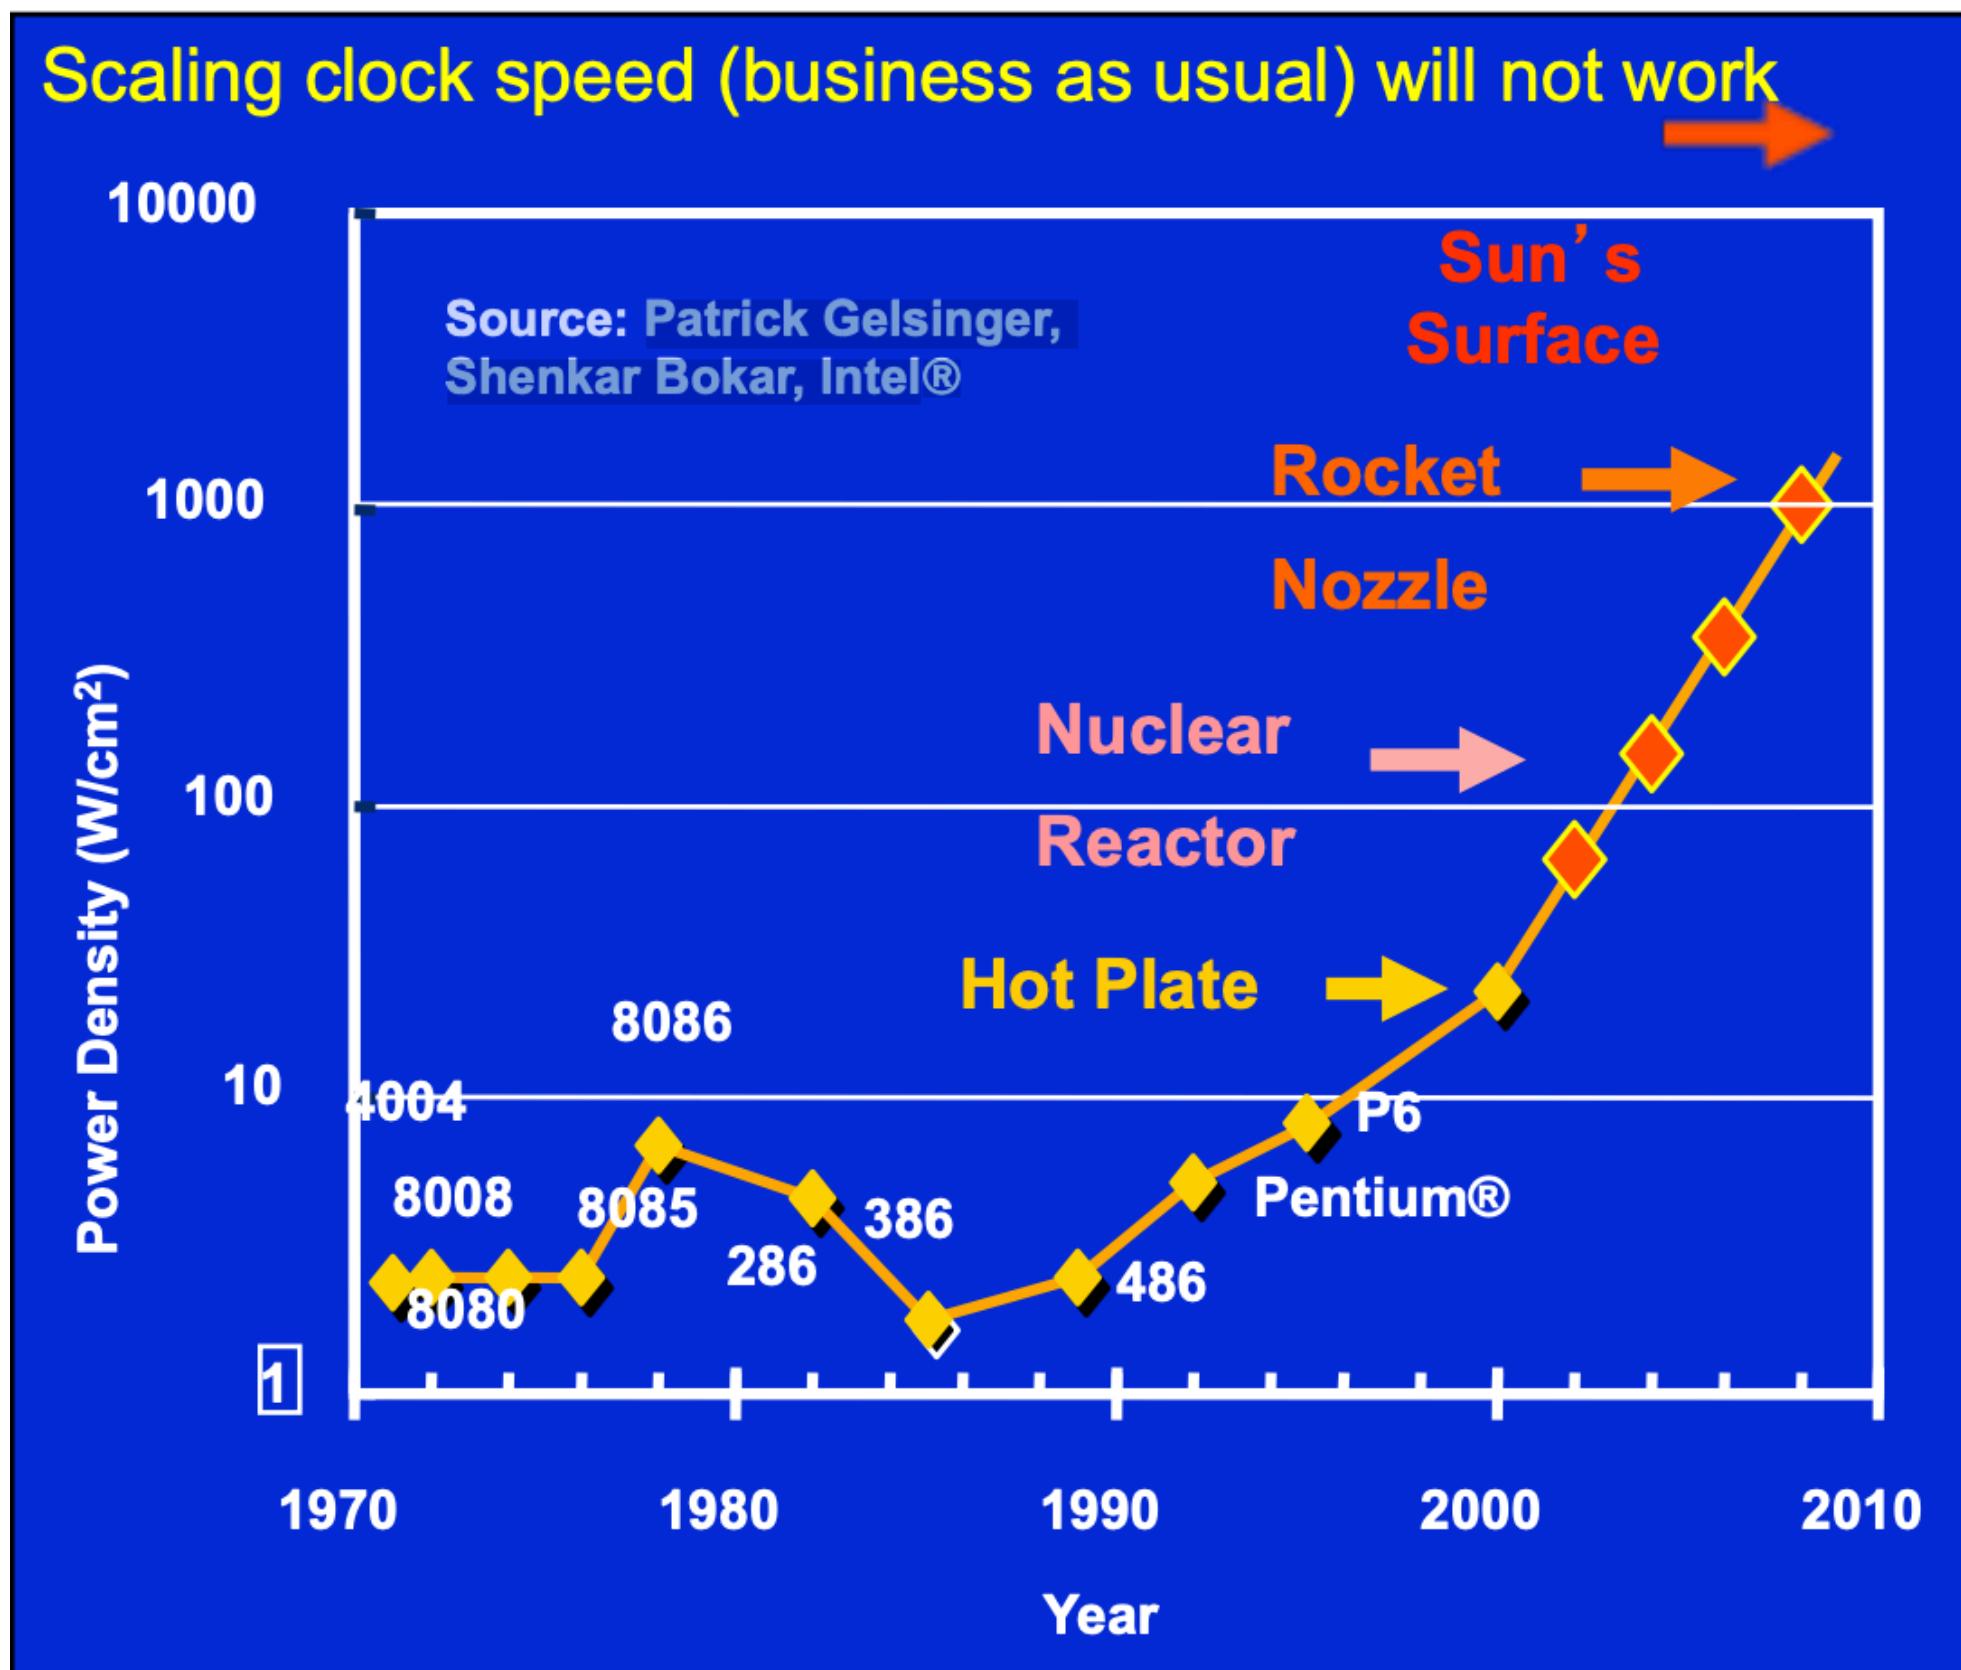


Courtesy of [Bernie Kohl](#) on Wikipedia.
Used under CC0.

Power Macintosh G5

Launched: 2004
Clock rate: 1.8 GHz
Data path: 64 bits
Memory: 256 MB
Cost: \$1,499

Power density



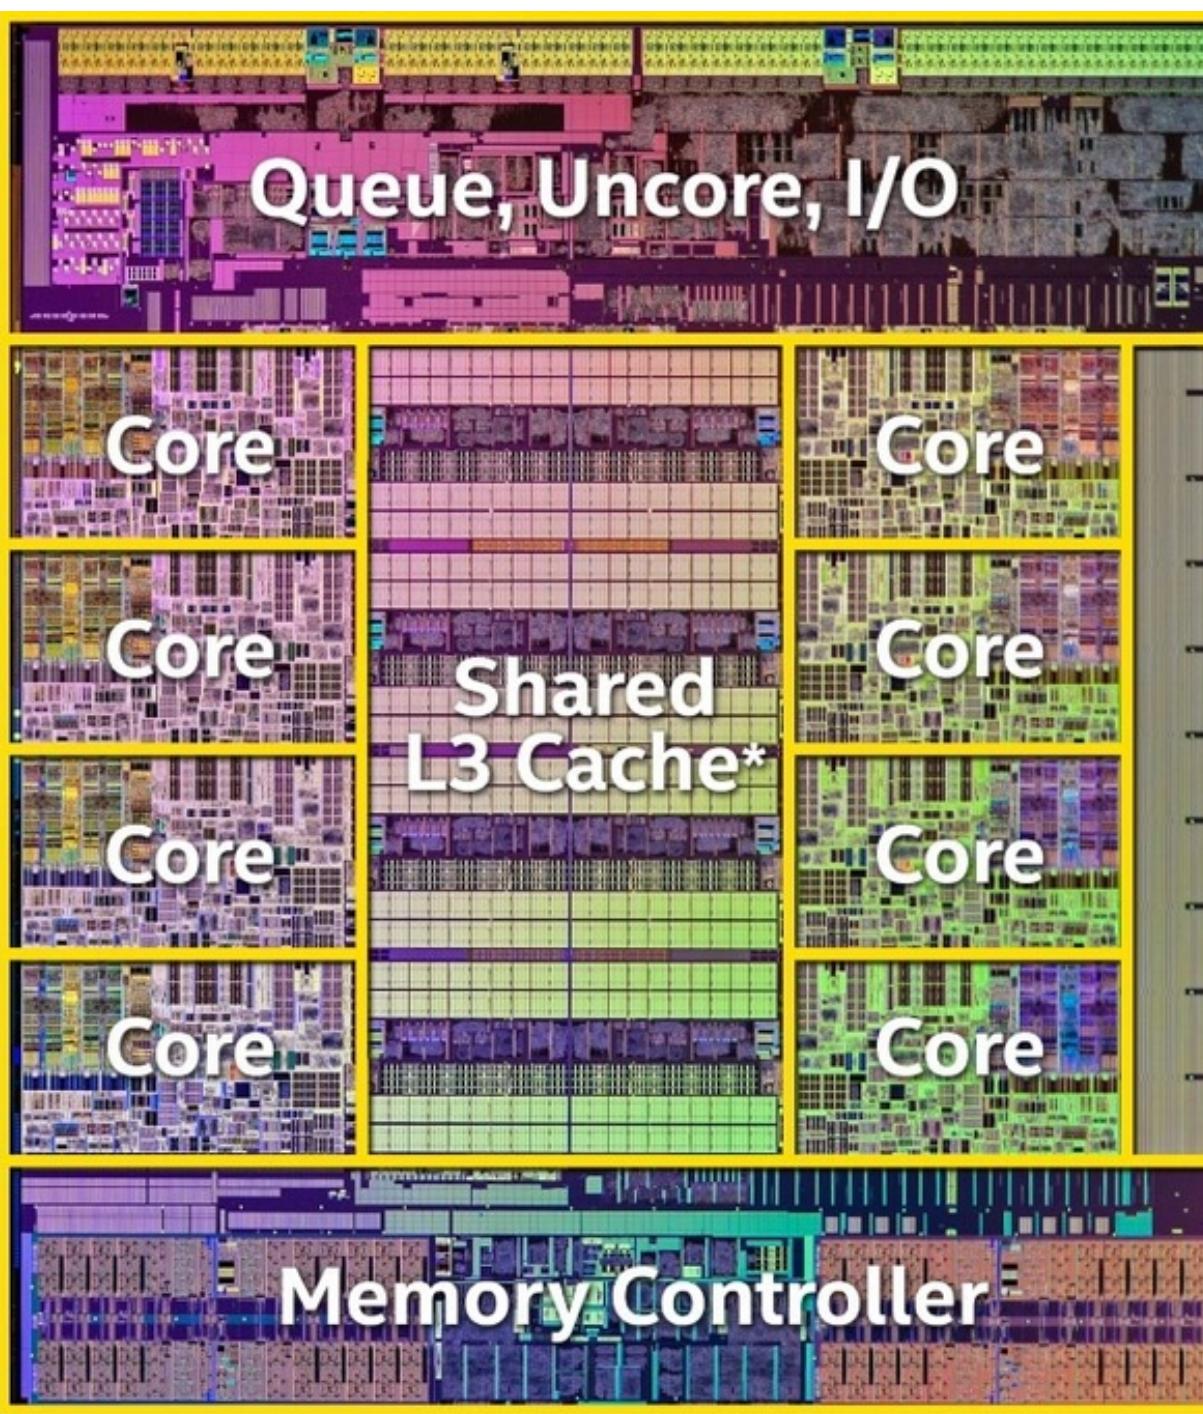
Is it better to increase speed by doubling frequency or cores?

$$\text{Performance} \propto (\text{cores}) \times (\text{freq})$$

$$\text{Power} \propto (\text{cores}) \times (\text{freq}^{2.5})$$

The growth of power density, as seen in 2004, if the scaling of clock frequency had continued its trend of 25%-30% increase per year.

Vendor solution: Multicore (Shared-memory processors)



E.g. Intel Core i7 3960X
(Sandy Bridge E), 2011
6 cores
3.3 GHz
15 MB L3 cache

To scale performance, processor manufacturers put **many processing cores** on the microprocessor chip.

Each generation of Moore's Law potentially doubles the number of cores.

Technology scaling after 2004



1,000,000

100,000

10,000

1,000

100

10

1

0

Normalized
transistor count

Clock speed (MHz)

Processor cores

“Moore’s law”

“Dennard scaling”

Year

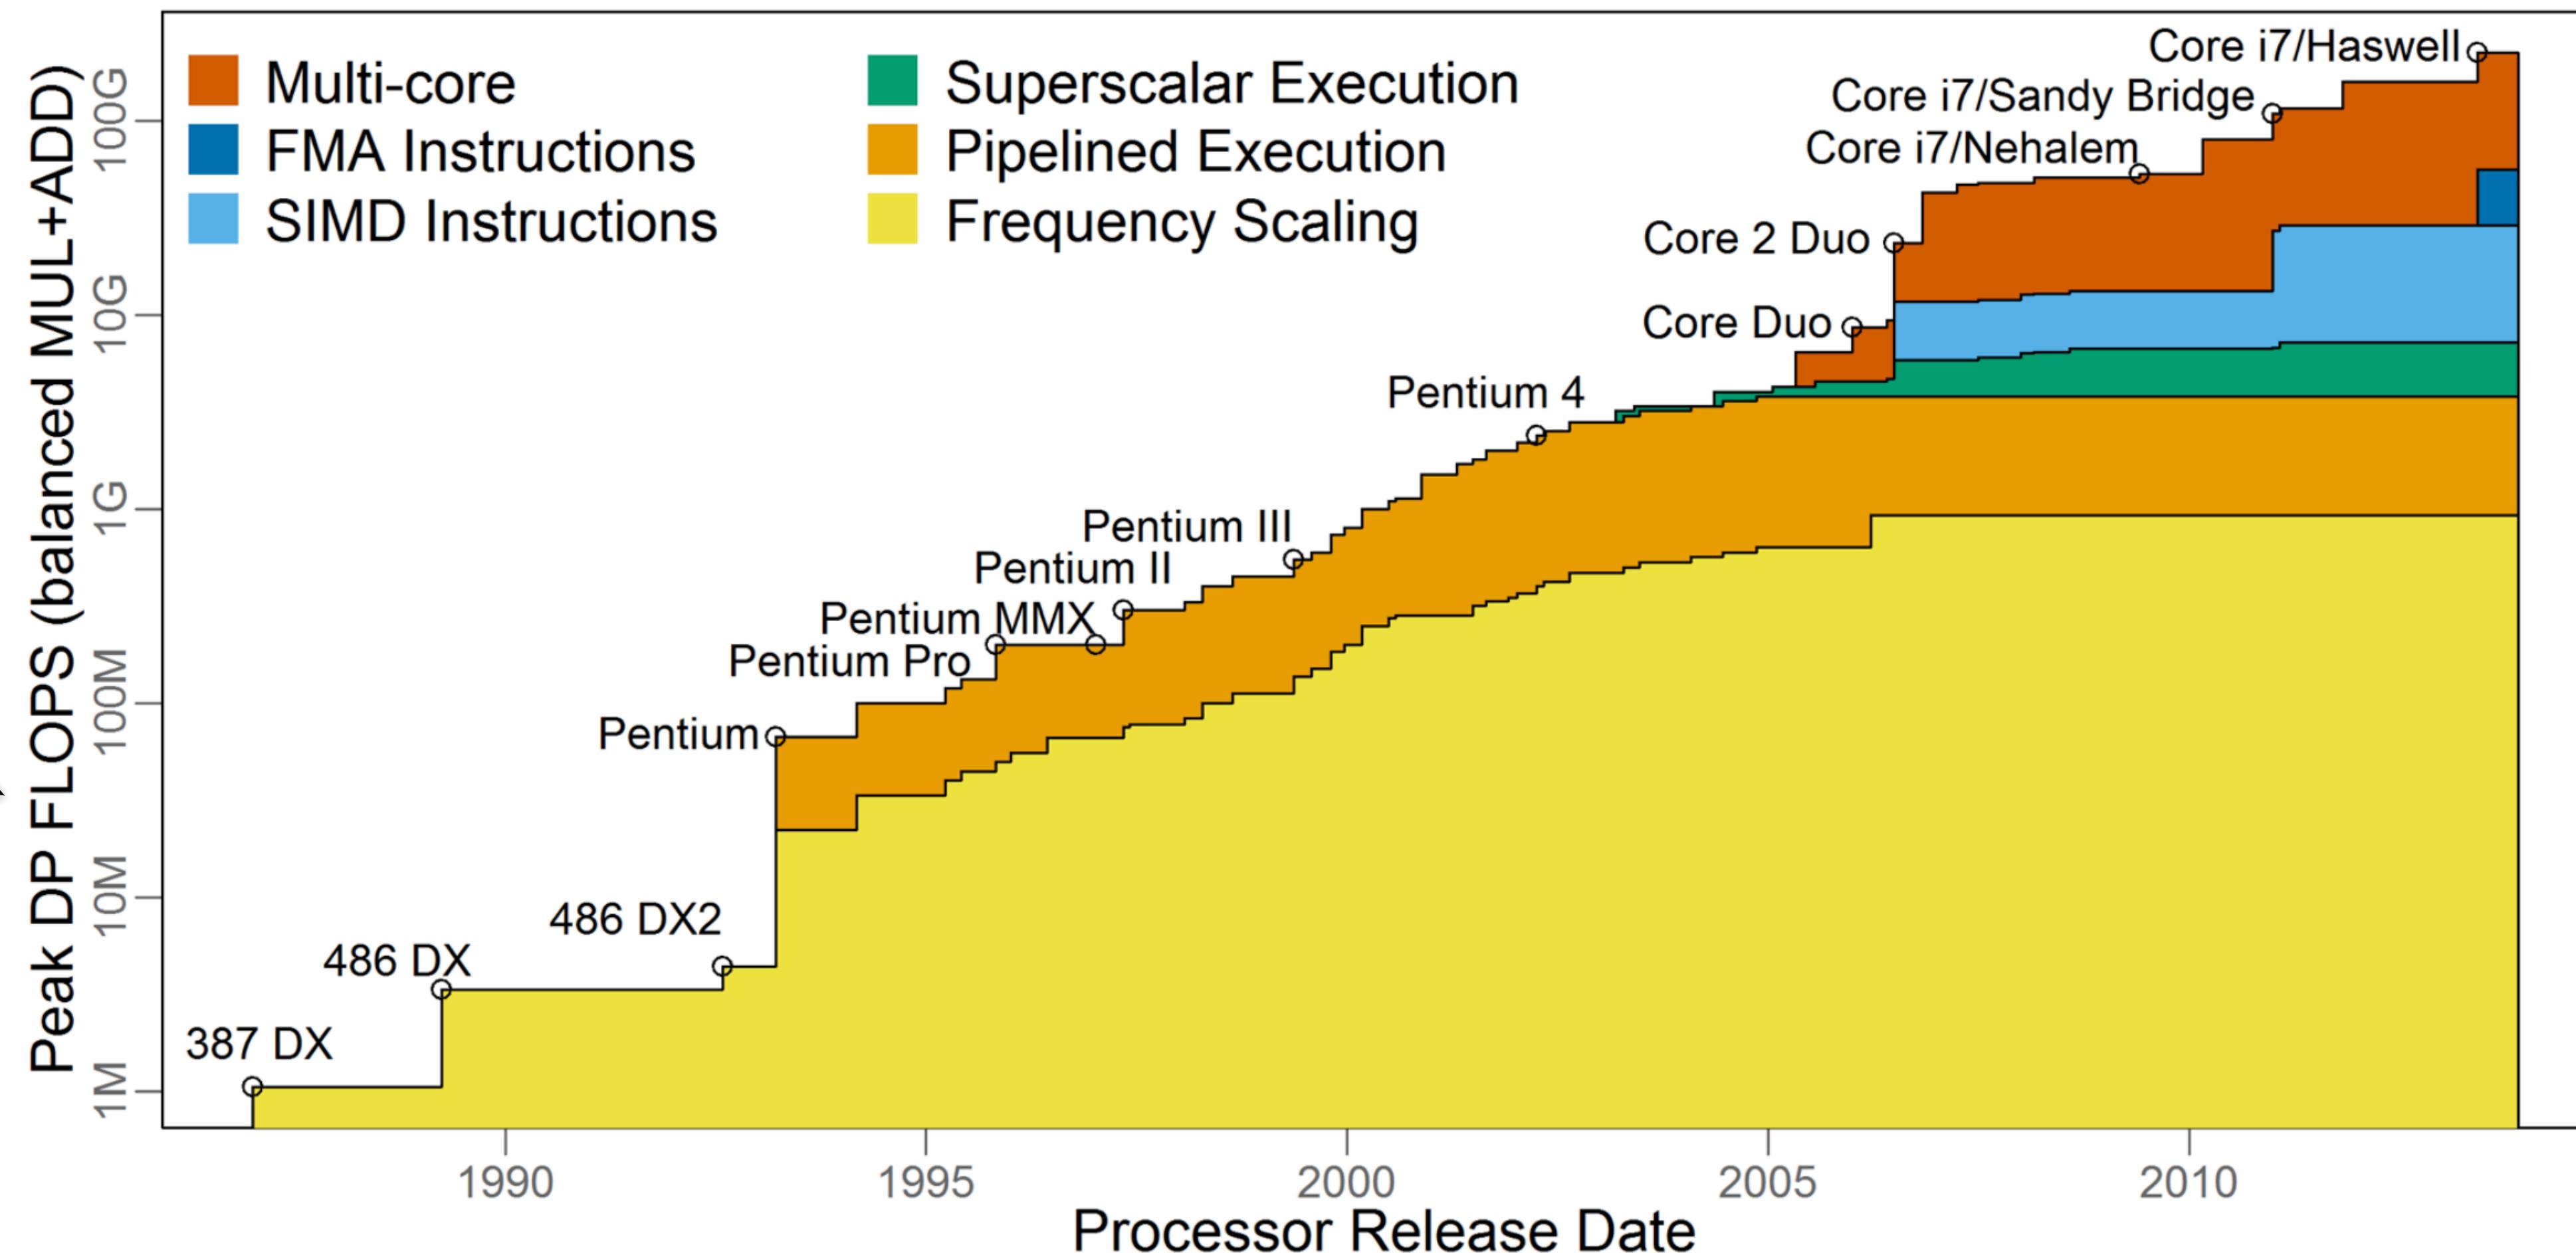
Processor data from Stanford’s CPU DB [DKM12].

Observe transition ~ 2004



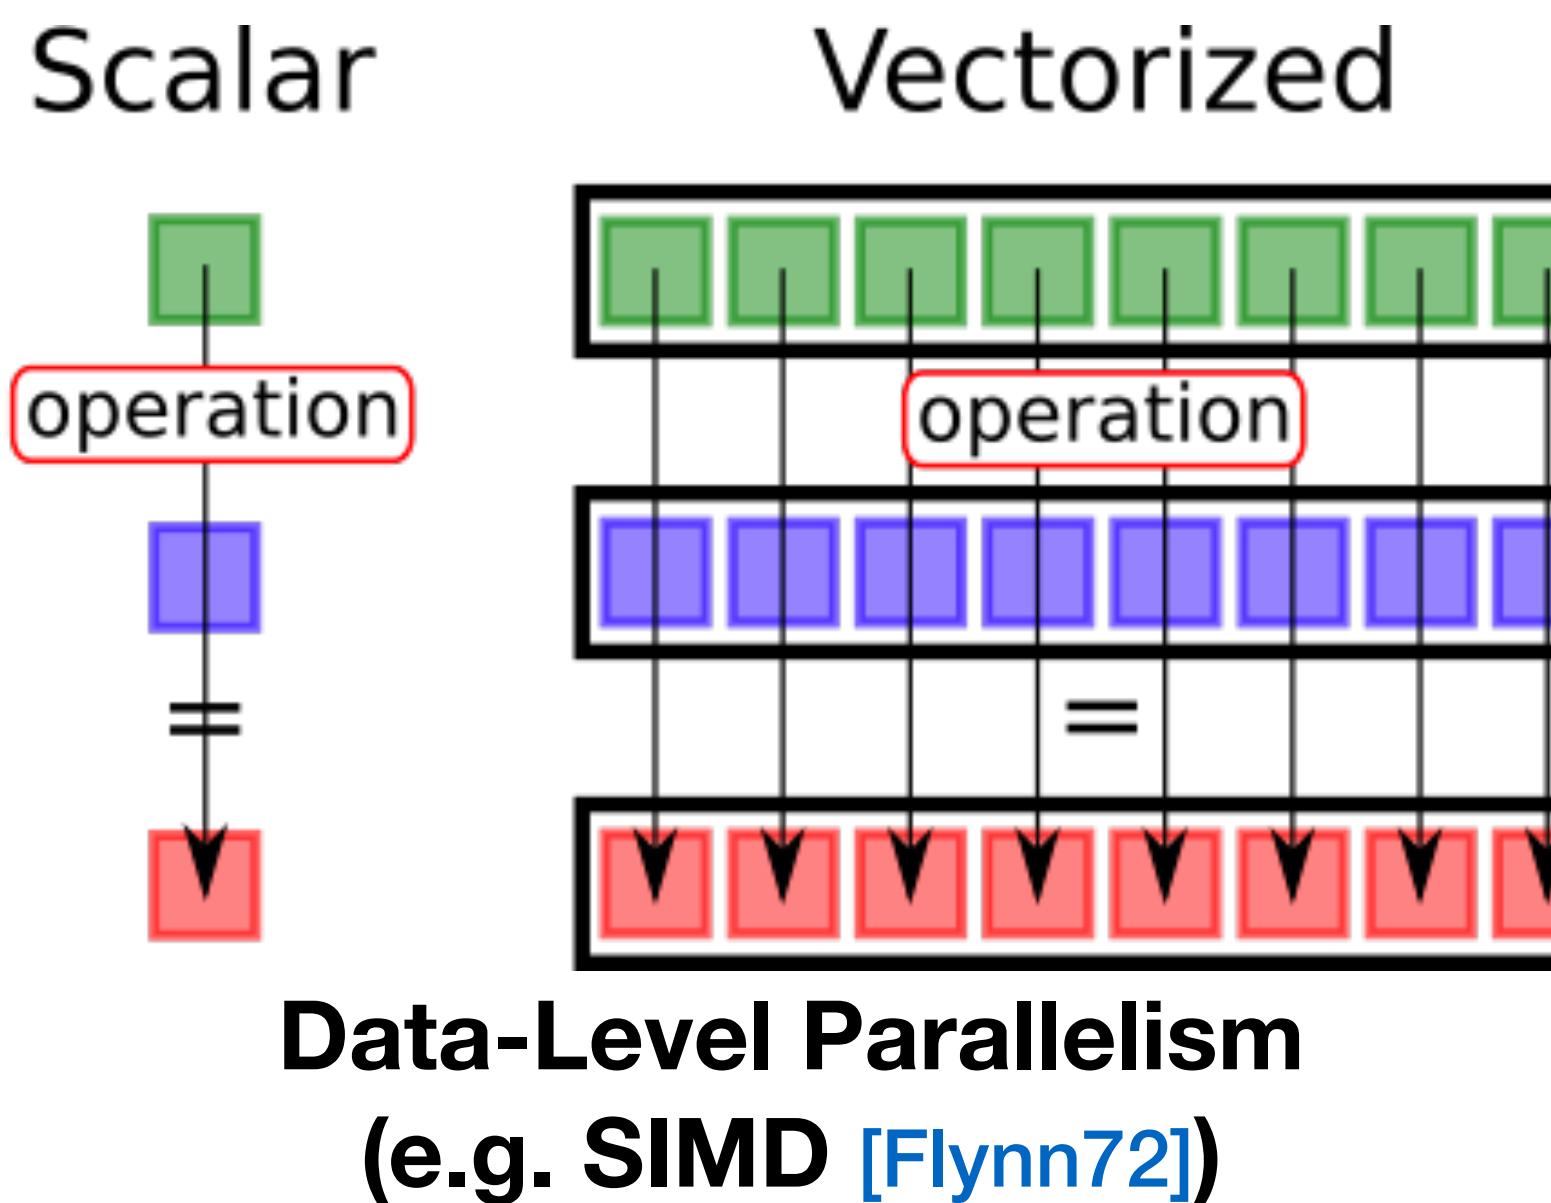
Where are performance gains coming from?

Double precision floating point ops/sec



Source: Marat Dukhan <mdukan3@gatech.edu>

Single Instruction Multiple Data (SIMD) Computer



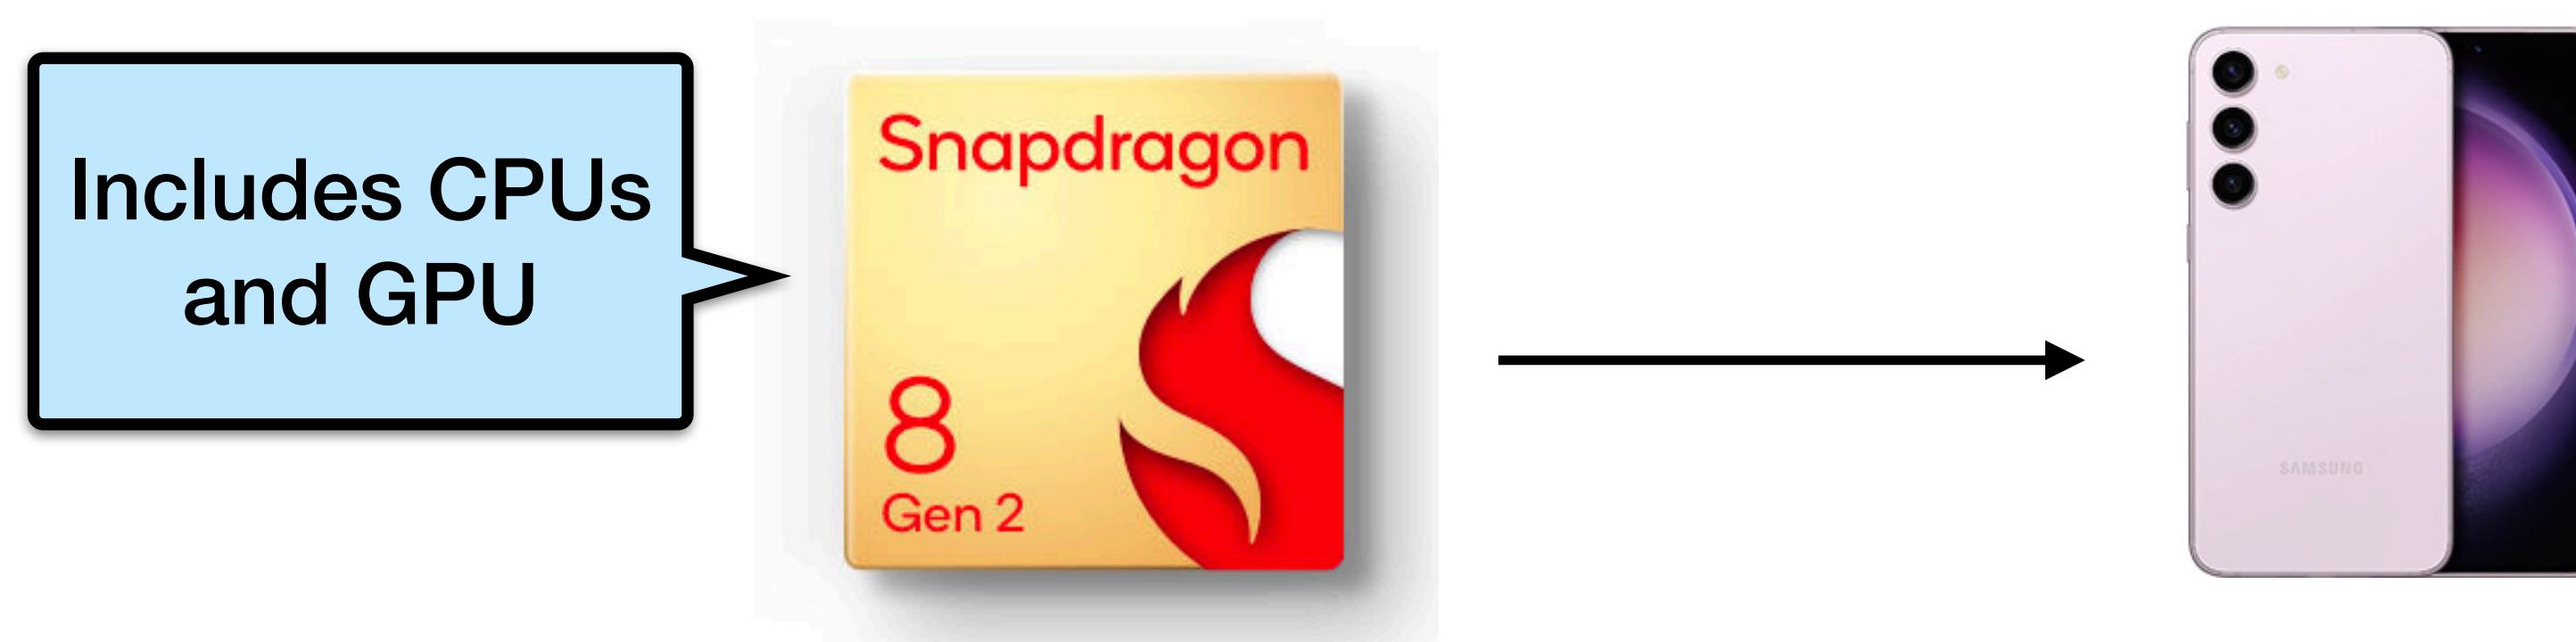
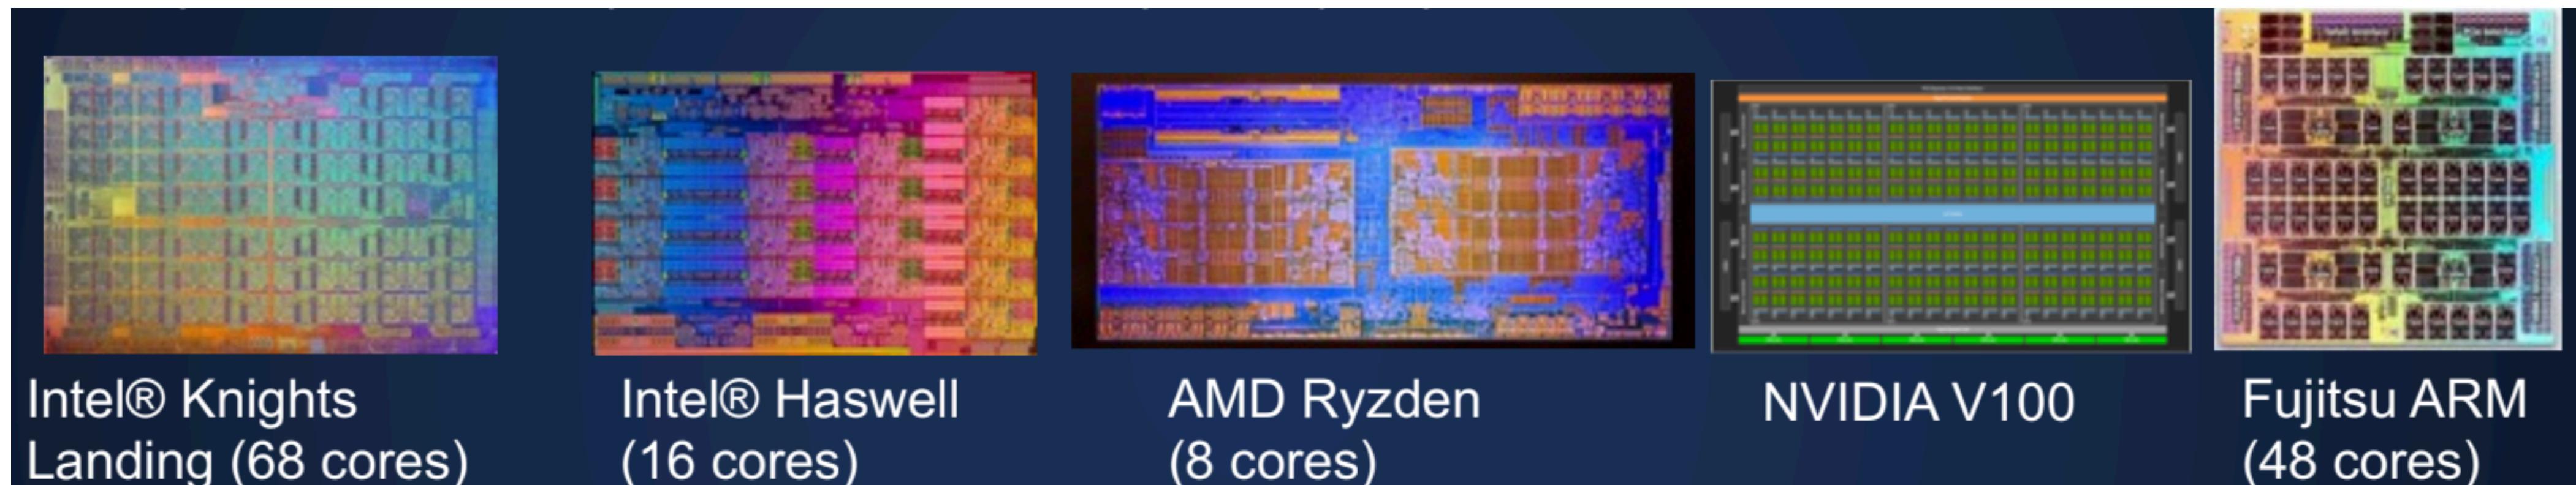
A Single-Instruction Multiple-Data (SIMD) computer has multiple processors (or functional units) that perform **the same operation on multiple data elements** at once.

Most single processors have SIMD units with **~2-8 way parallelism**.

Graphics processing units (**GPUs**) use SIMD as their primary programming paradigm.

Parallel computing is ubiquitous in individual machines of all sizes

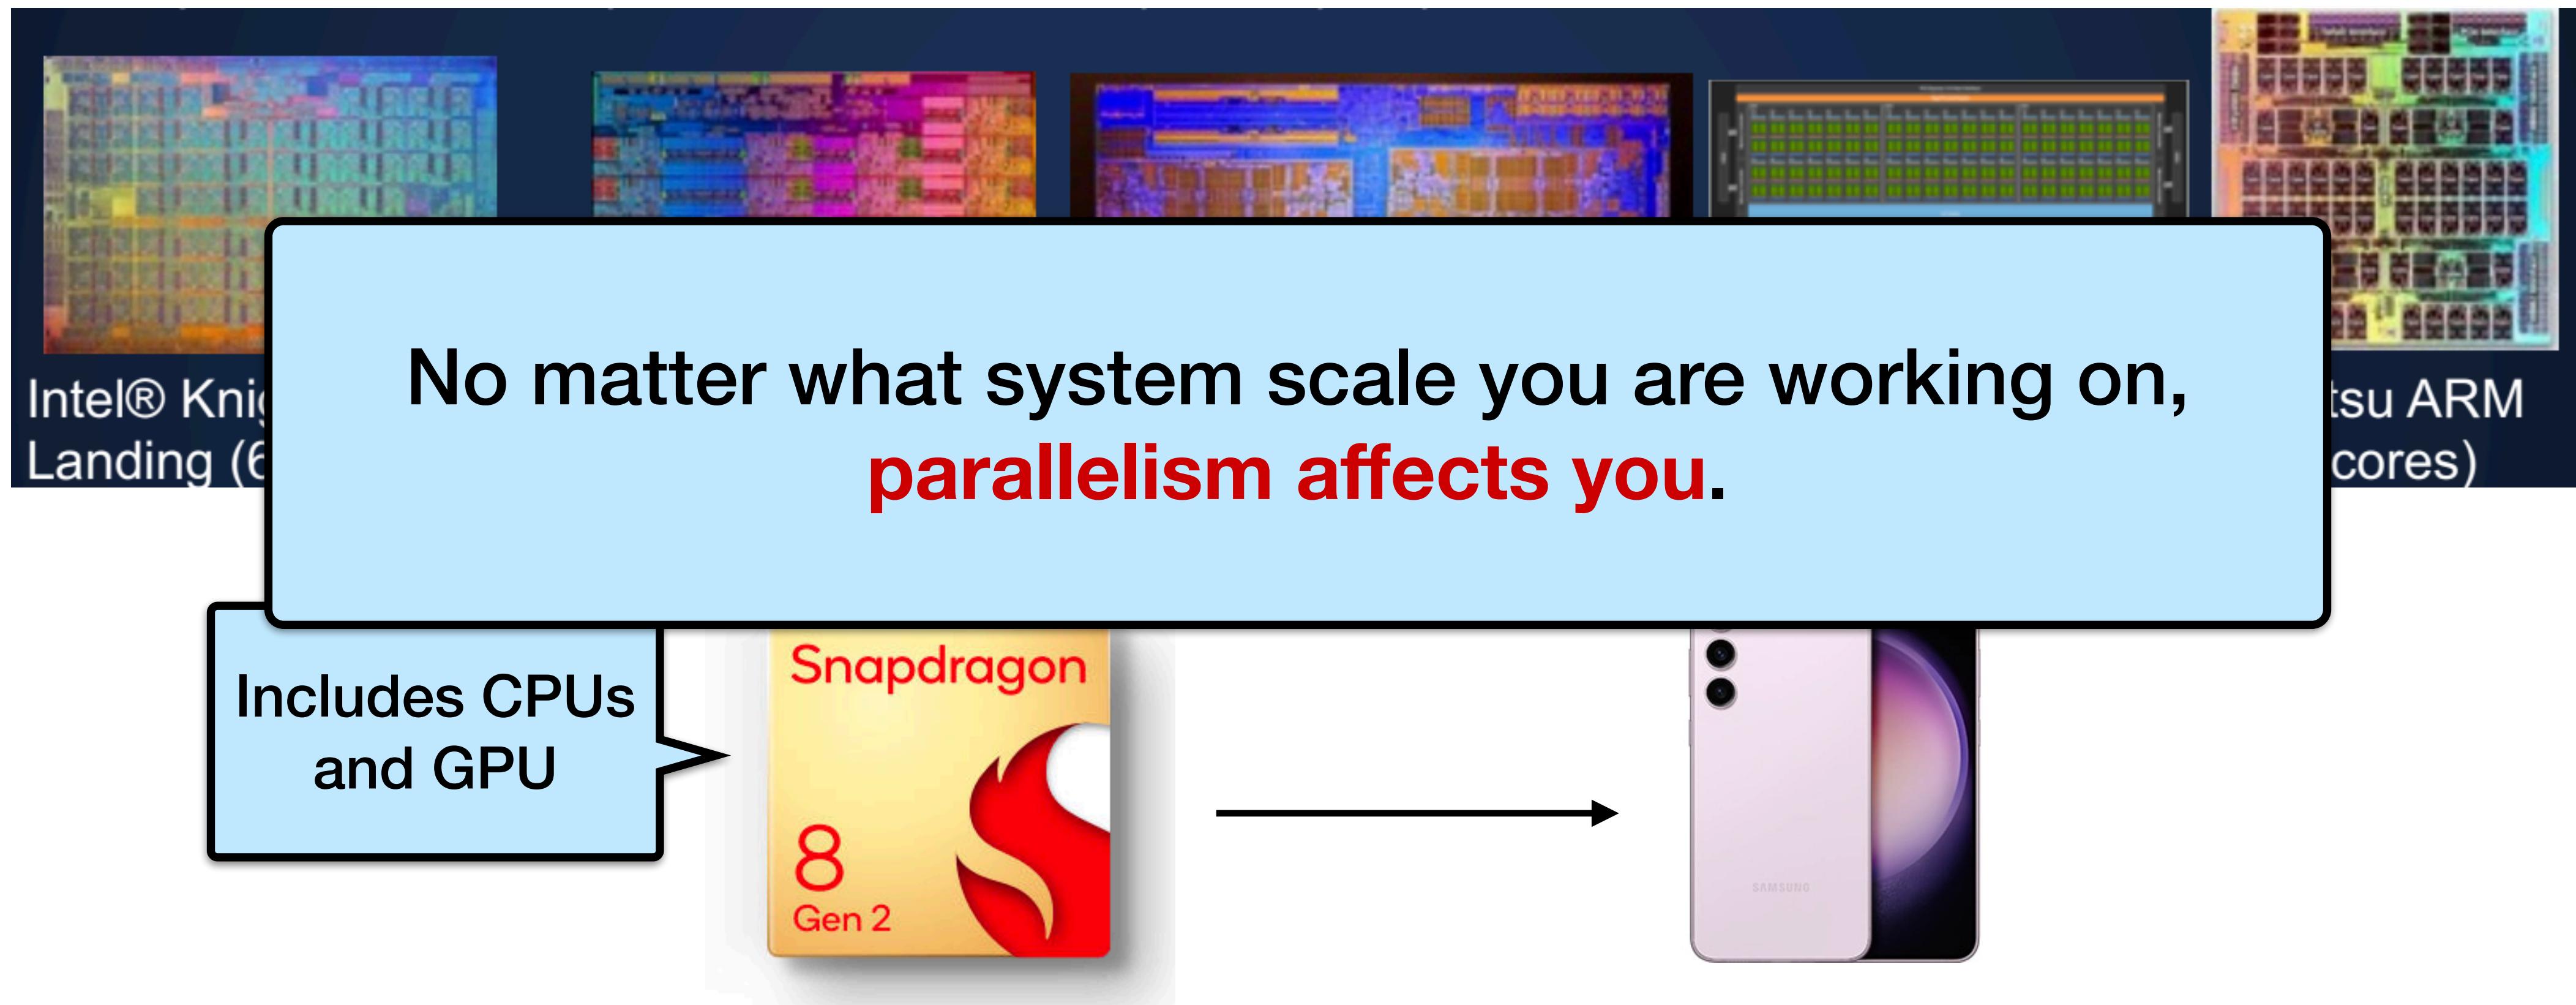
Trends in technology scaling drive parallelism in computing at **all scales**.



<https://www.qualcomm.com/products/mobile/snapdragon/smartphones/snapdragon-8-series-mobile-platforms/snapdragon-8-gen-2-mobile-platform>

Parallel computing is ubiquitous in individual machines of all sizes

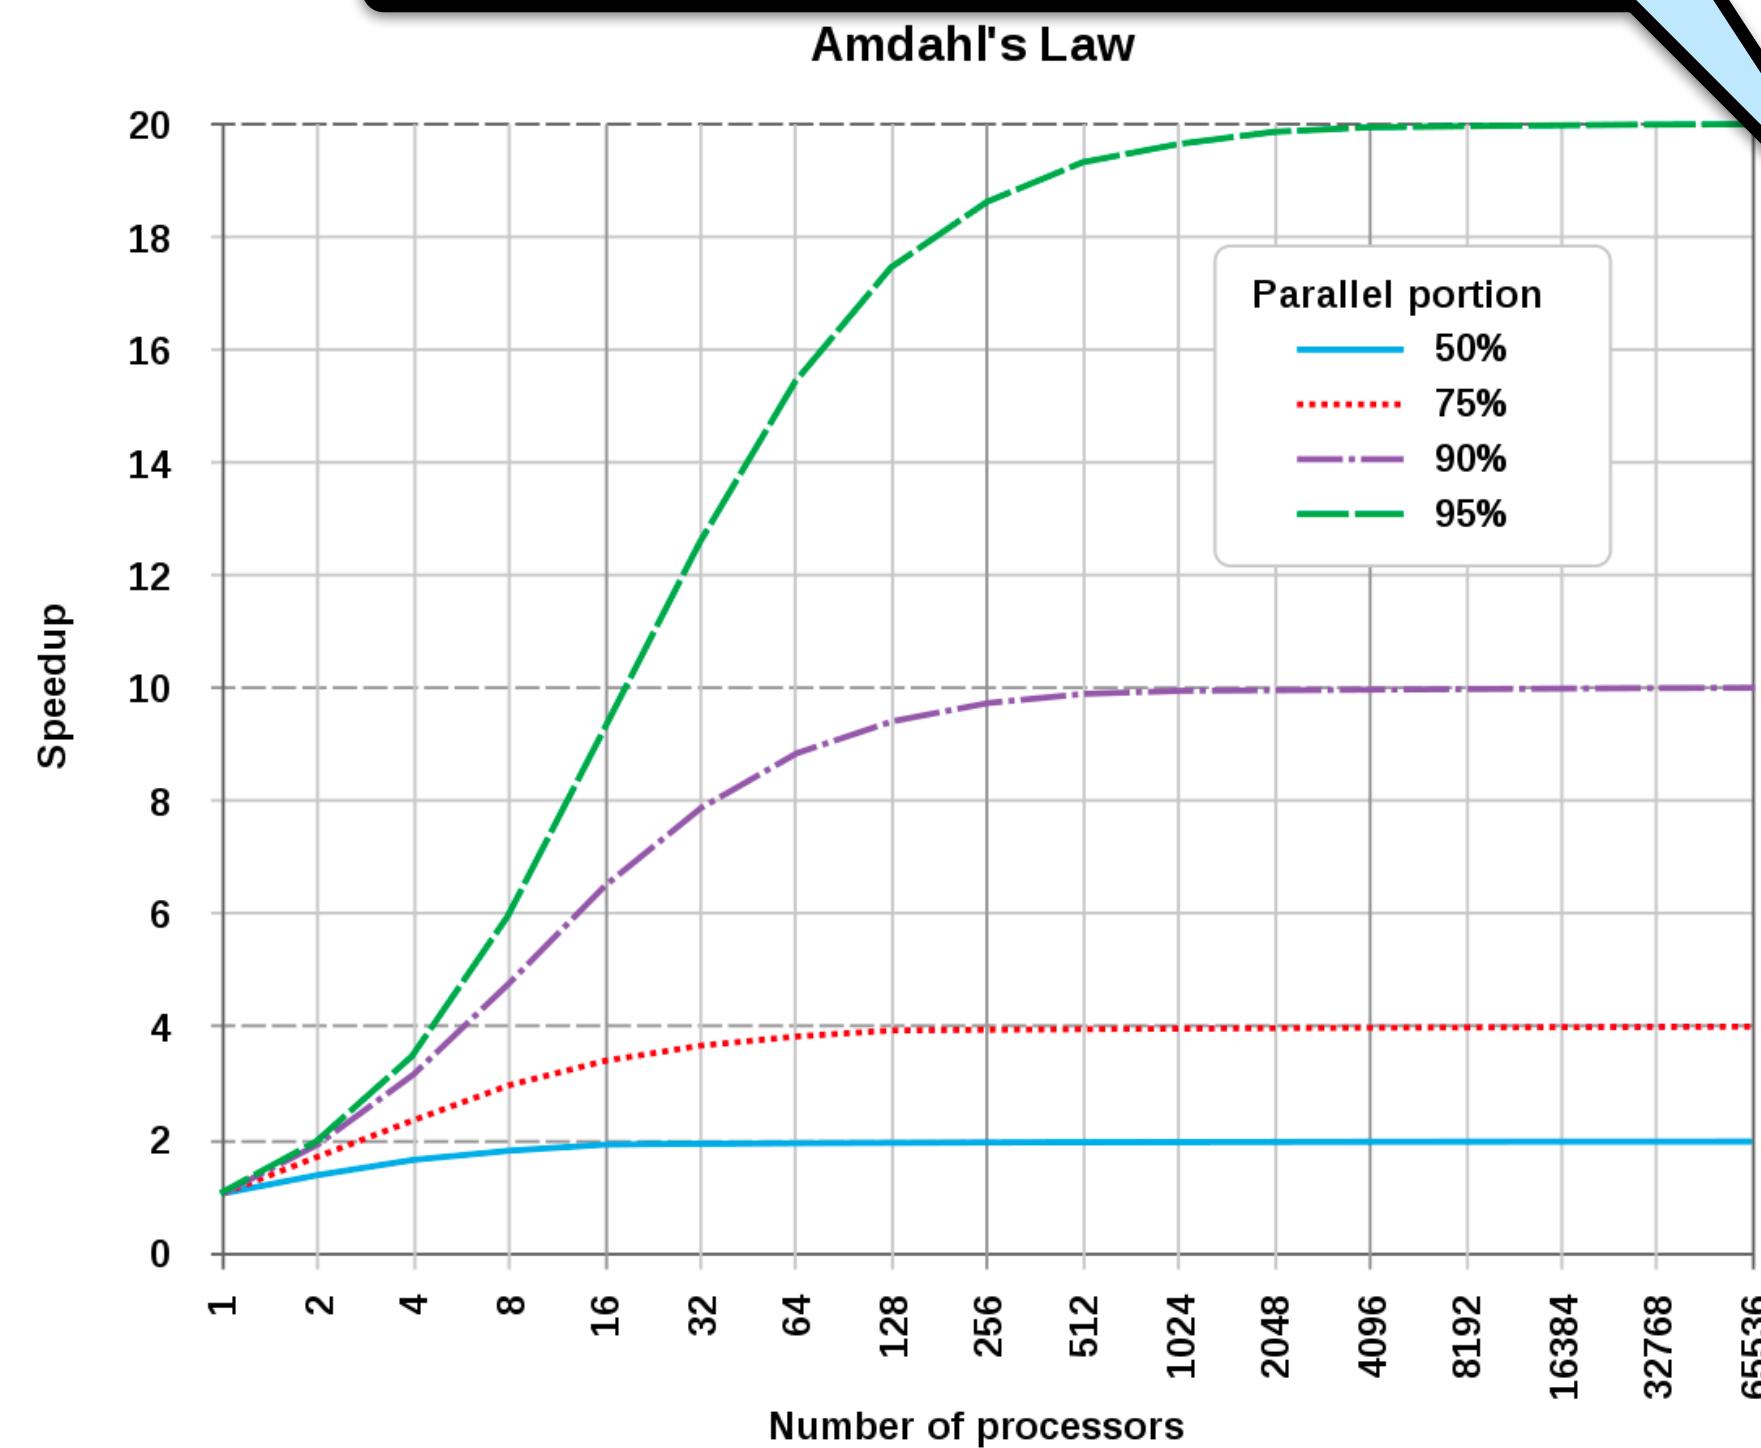
Trends in technology scaling drive parallelism in computing at **all scales**.



<https://www.qualcomm.com/products/mobile/snapdragon/smartphones/snapdragon-8-series-mobile-platforms/snapdragon-8-gen-2-mobile-platform>

Bounding Maximum Parallel Speedup with Amdahl's Law

Work done by the **best** sequential algorithm



Suppose only part of an application is parallel.

Amdahl's law:

s = fraction of work done sequentially (Amdahl fraction)

$1-s$ = parallelizable fraction

P = number of processors

$$\text{Speedup } (P) = \frac{\text{Time}(1)}{\text{Time}(P)}$$

$$\leq \frac{1}{s + (1-s)/P}$$

$$\leq \frac{1}{s}$$

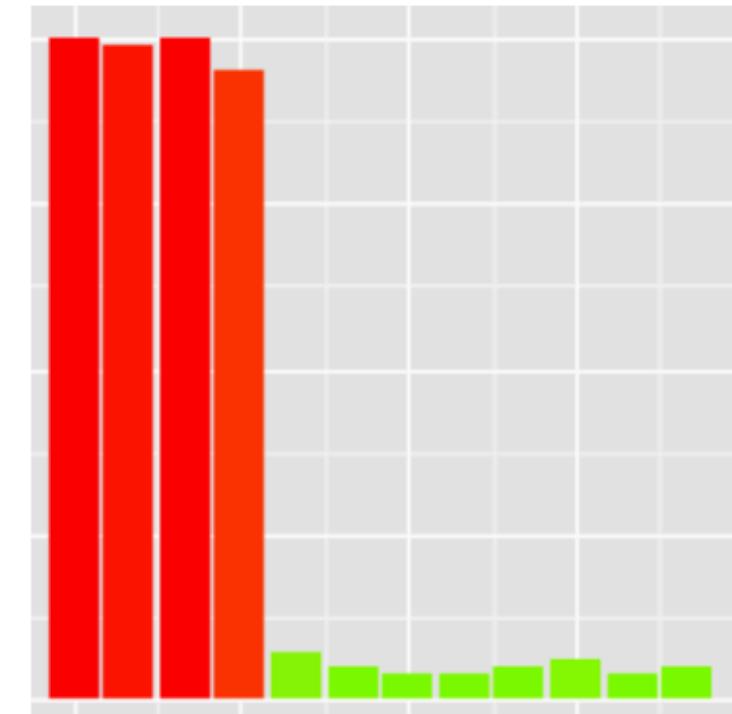
Even if the parallel part speeds up perfectly,
performance is limited by the sequential part.

Writing fast code requires going beyond parallelism alone

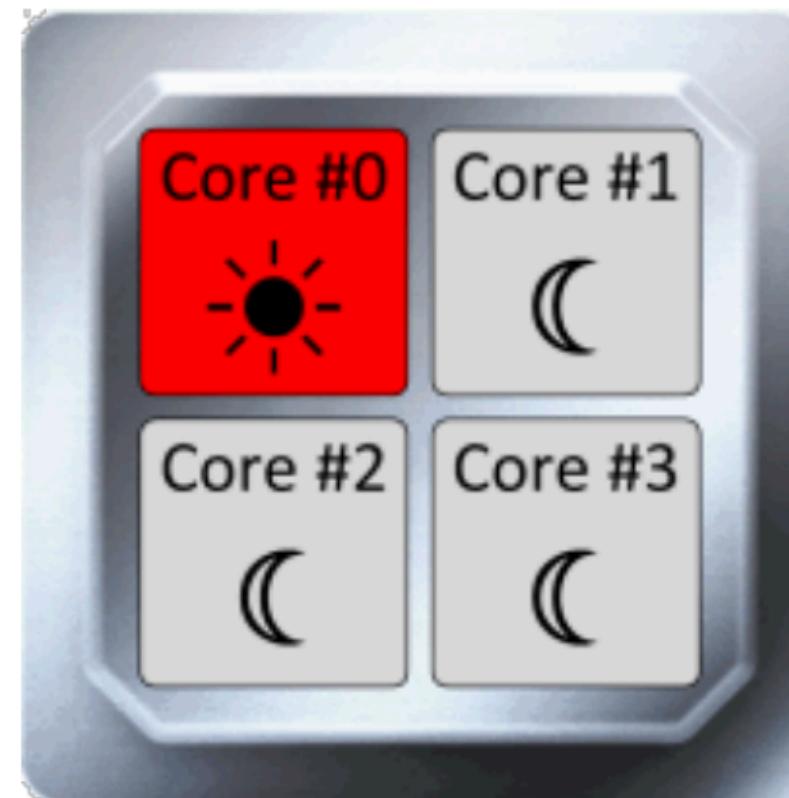
This class is about more than just parallelism. It's about getting the most out of your hardware, e.g., through **data locality**.

Why? The faster you run, the sooner you can stop, saving time, energy, cost,

...

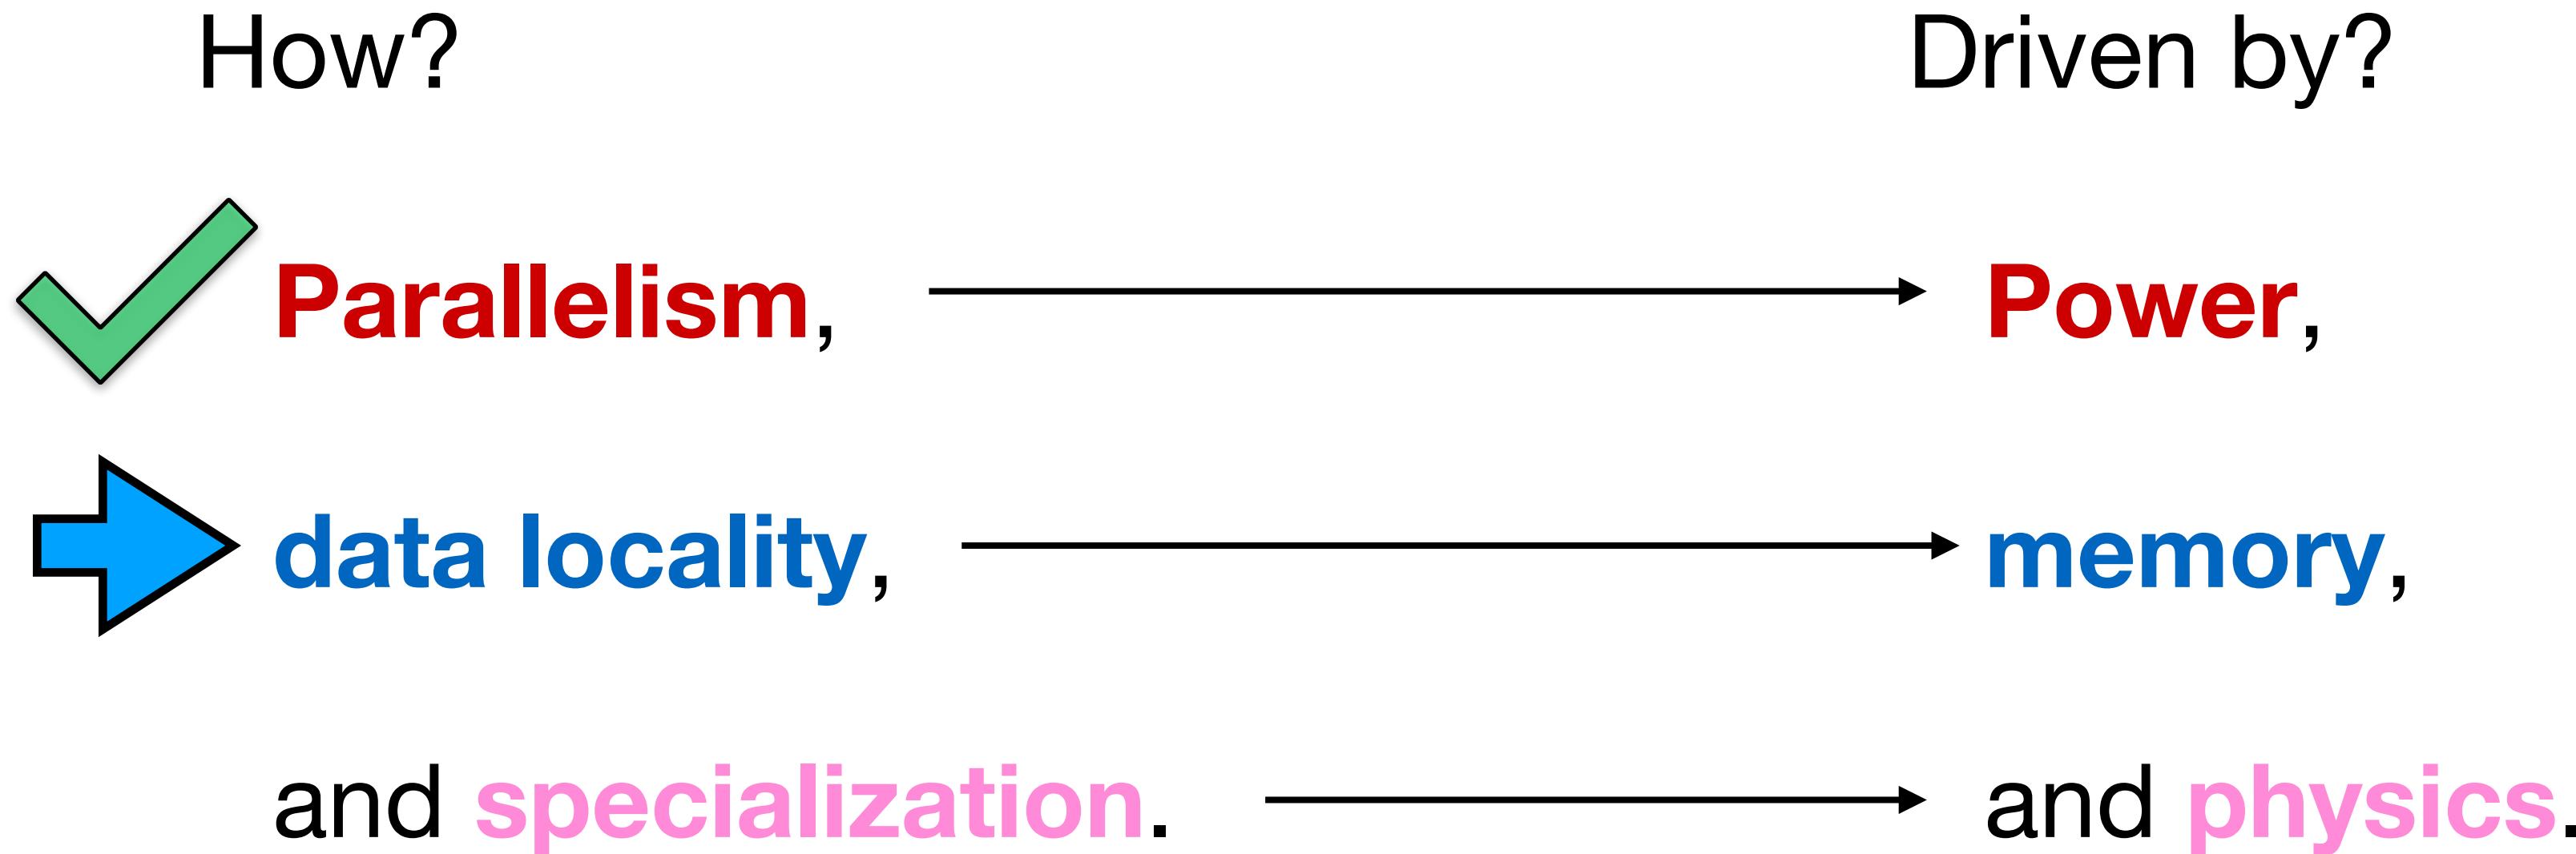


“Race-to-halt” — finish quickly then shutdown to save energy

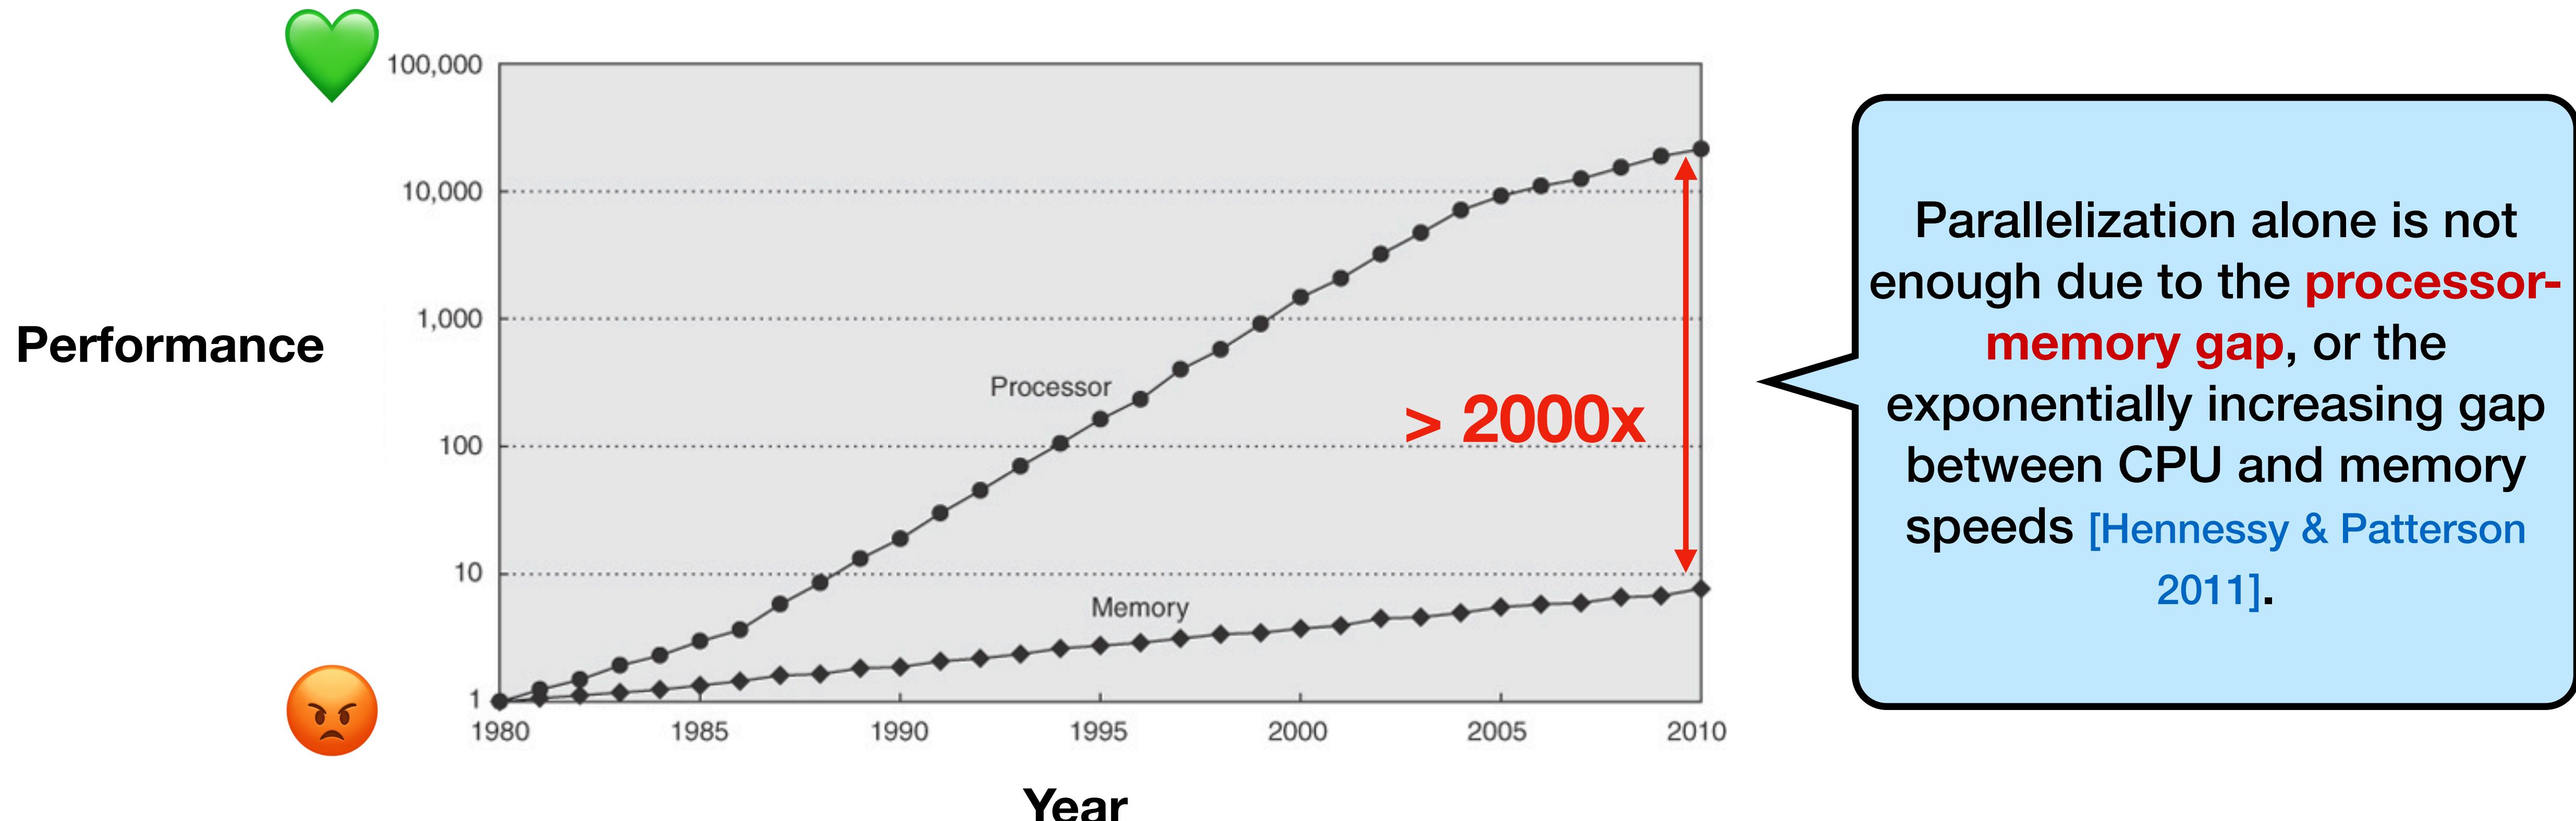


Use resources more efficiently, rather than using more resources

Motivating trends



Data locality & the memory wall

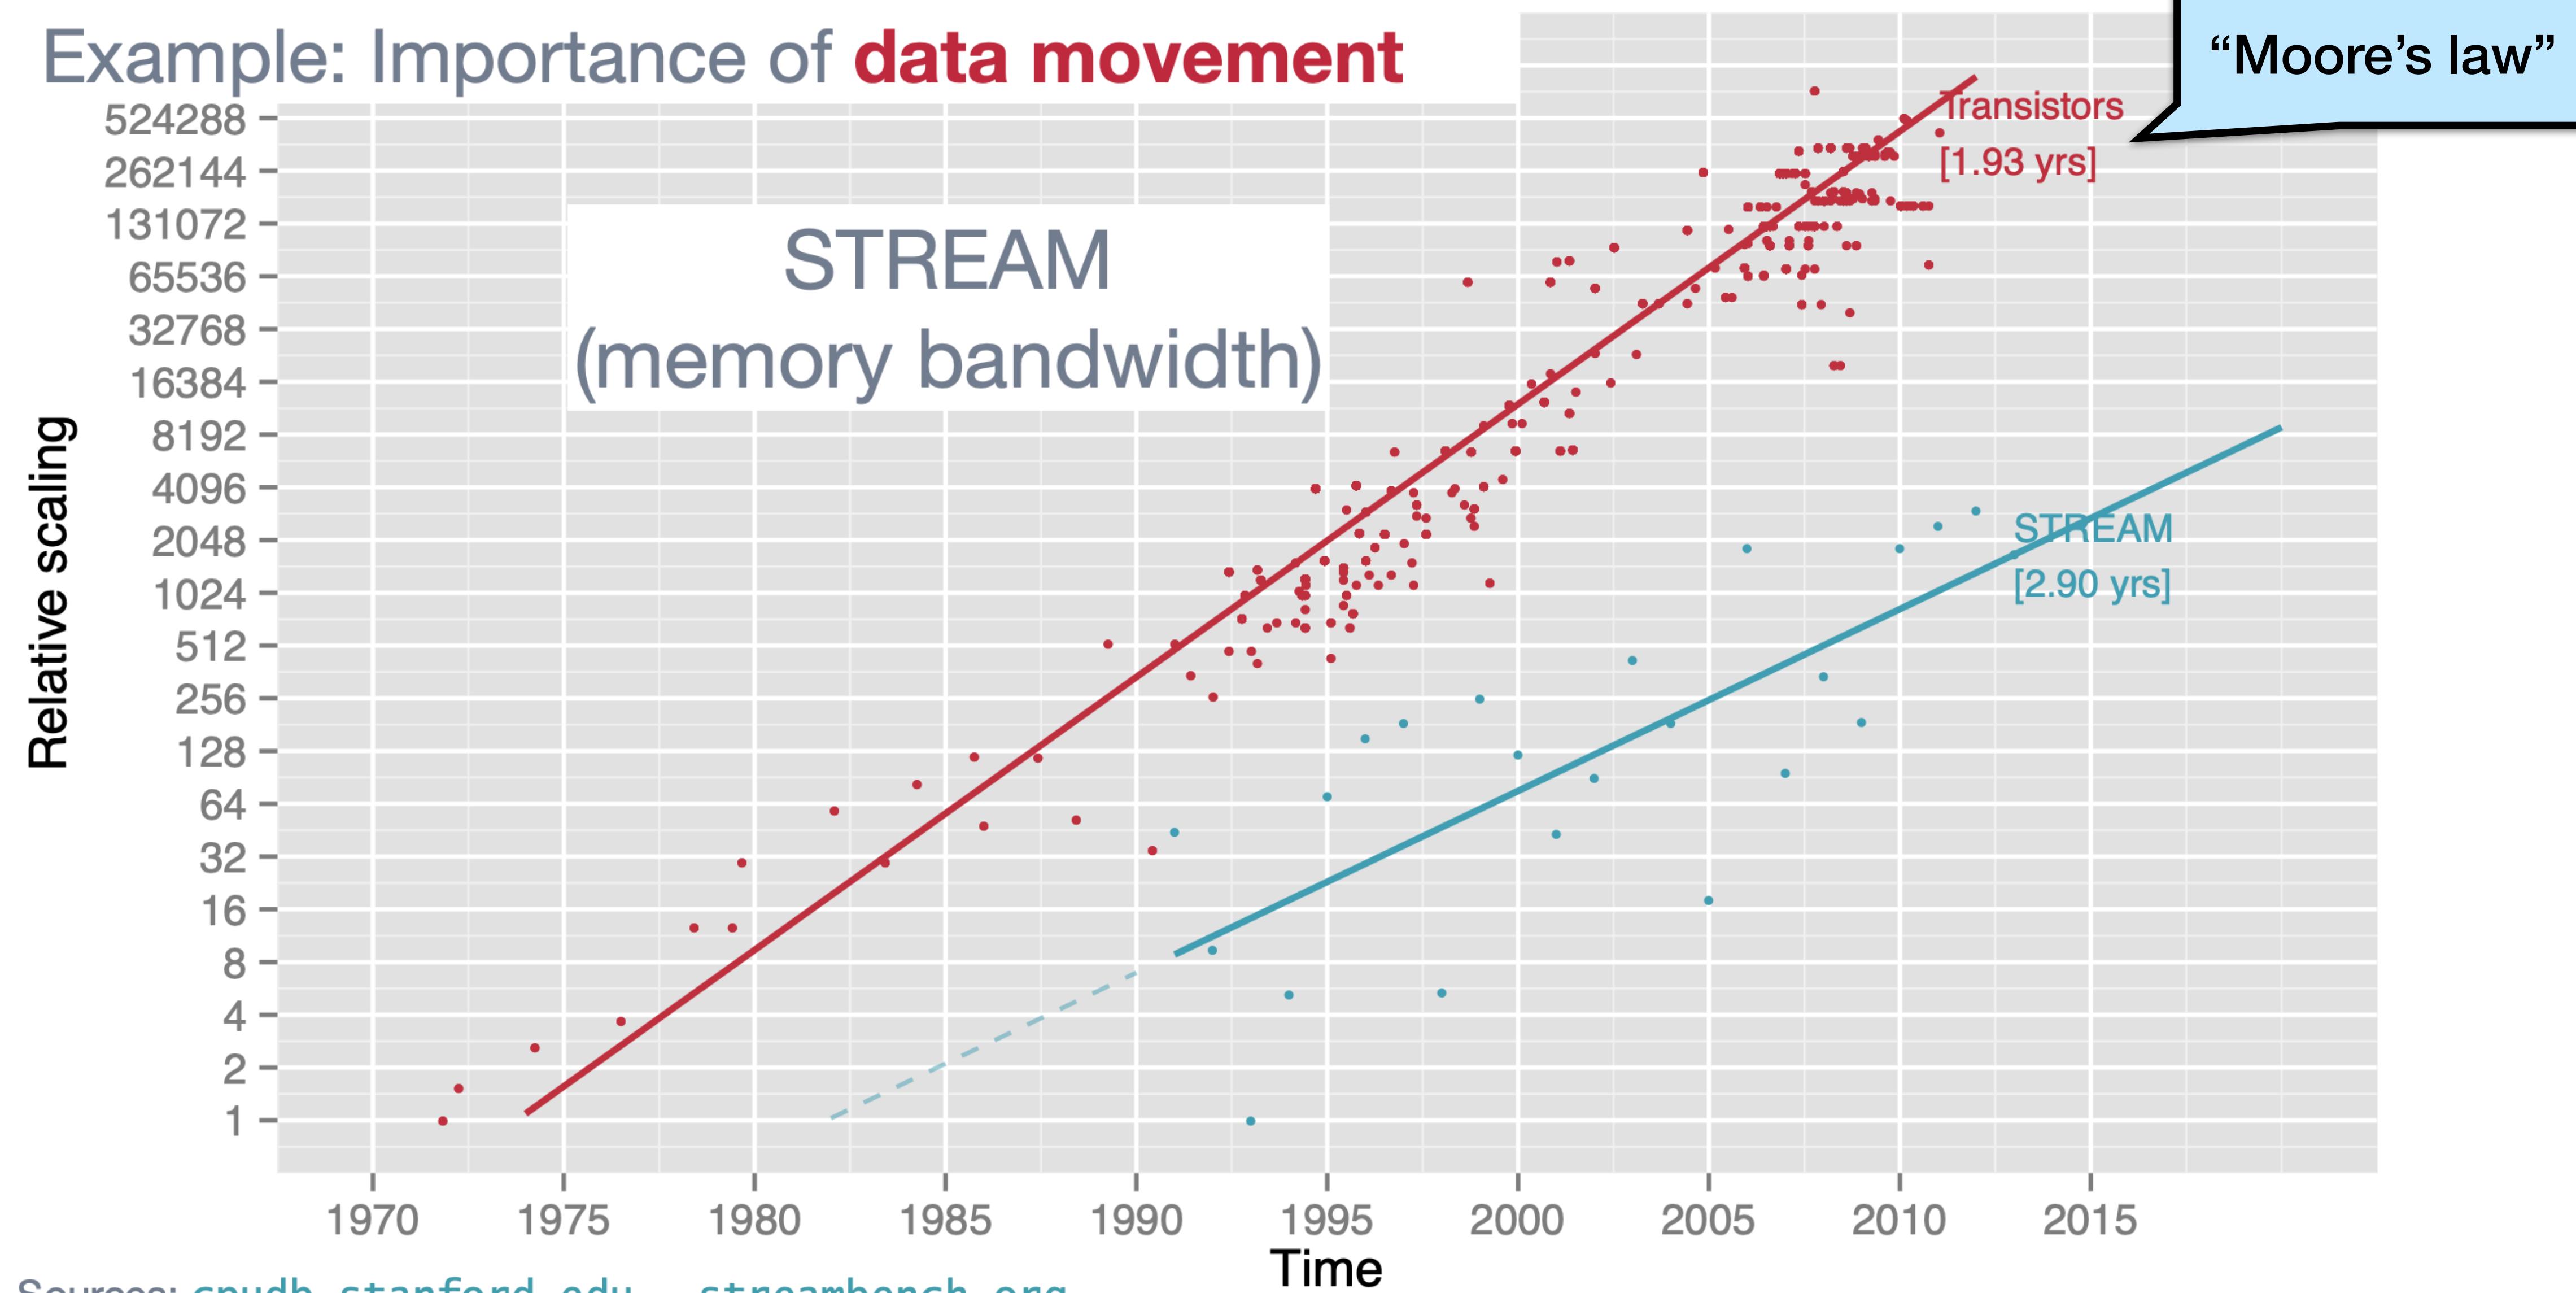


Parallelization alone is not enough due to the **processor-memory gap**, or the exponentially increasing gap between CPU and memory speeds [Hennessy & Patterson 2011].

The memory wall, or the **gap between memory latency and latency of floating-point operations**, is increasing exponentially over time.



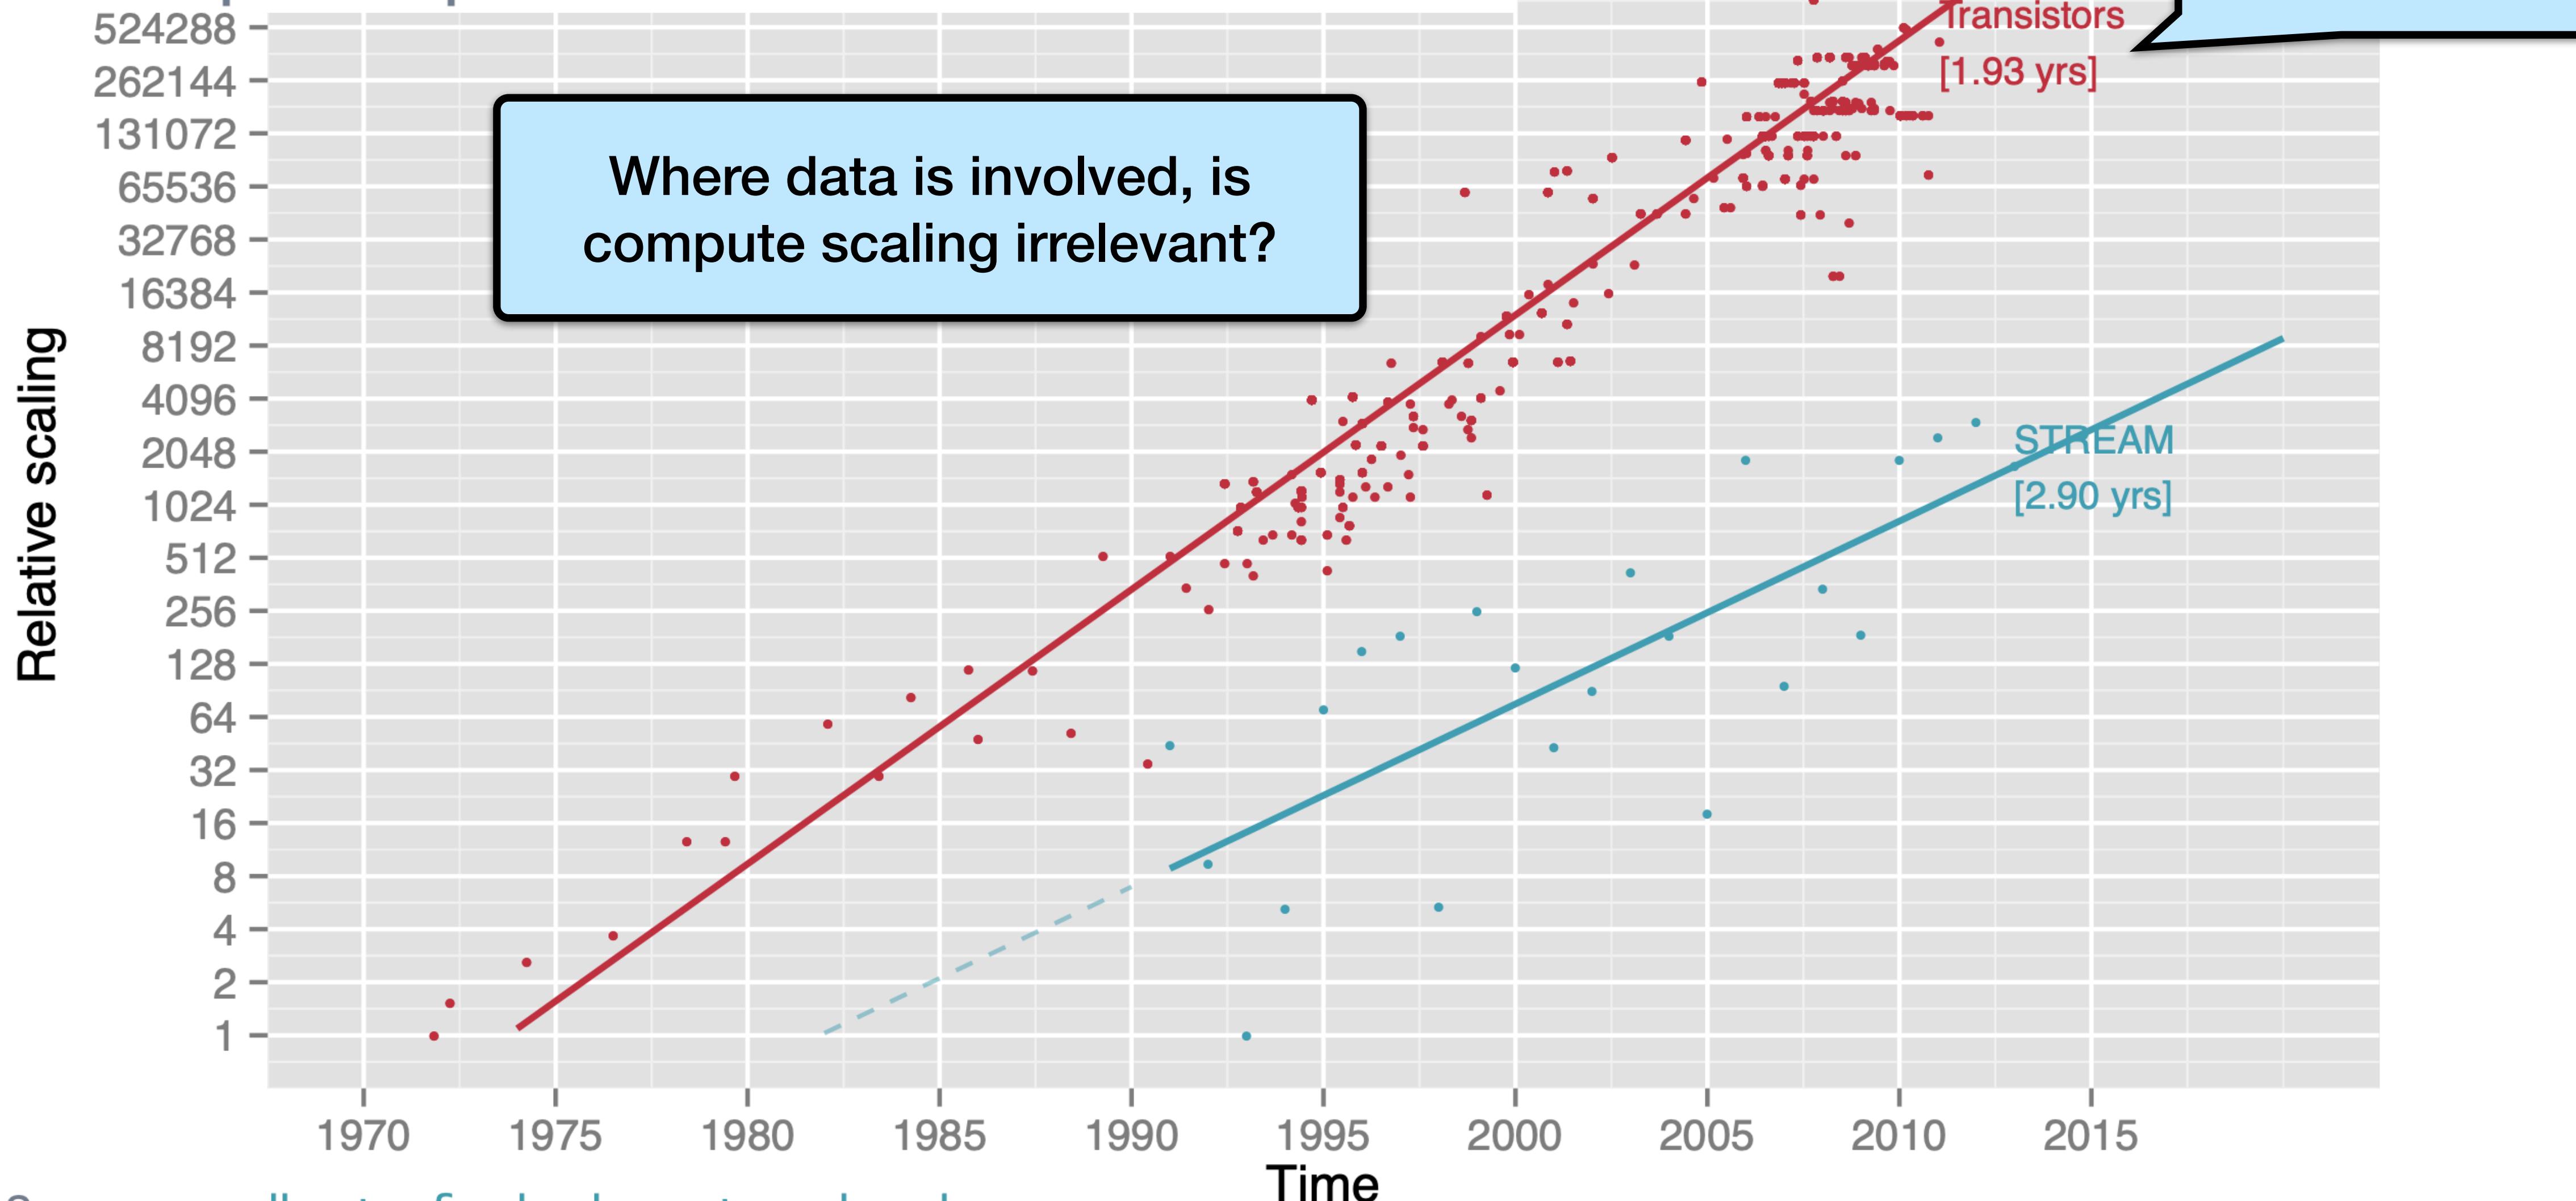
Example: Importance of **data movement**



Sources: cpudb.stanford.edu – streambench.org



Example: Importance of **data movement**



Motivating trends

How?

✓ **Parallelism,**

✓ **data locality,**

and **specialization.**

Driven by?

Power,

memory,

and **physics.**

Not covered in detail in this course

Example: Specialization (in hardware)

Understanding Sources of Inefficiency in General-Purpose Chips

Rehan Hameed¹, Wajahat Qadeer¹, Megan Wachs¹, Omid Azizi¹, Alex Solomatnikov²,
Benjamin C. Lee¹, Stephen Richardson¹, Christos Kozyrakis¹ and Mark Horowitz¹

¹Dept. of Electrical Engineering
Stanford University, Stanford, CA
{rhameed, wqadeer, wachs, oazizi,
bcclee, stevenr, kozyraki, horowitz}@stanford.edu

²Hicamp Systems,
Menlo Park, CA
solomatnikov@gmail.com

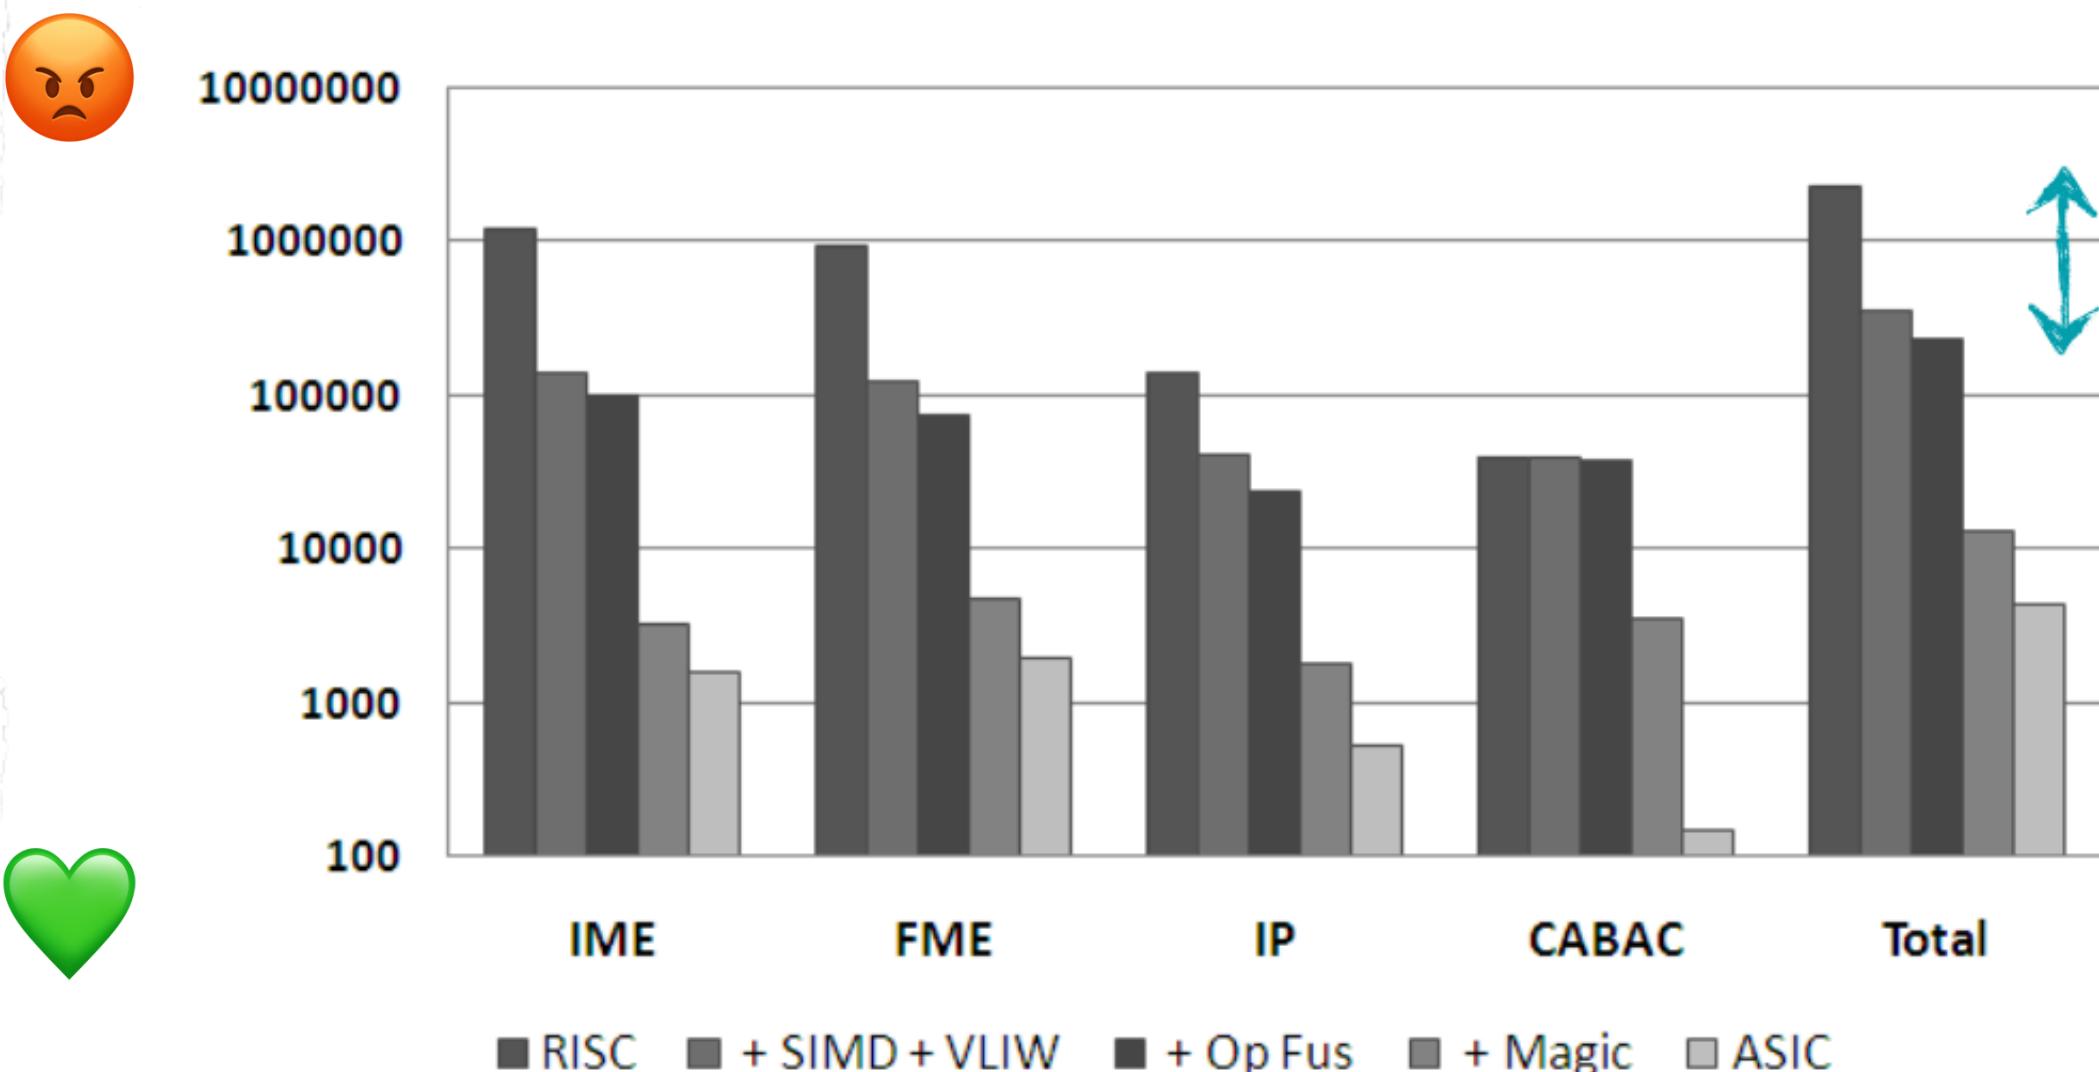


Figure 2. Each set of bar graphs represents energy consumption (μJ) at each stage of optimization for IME, FME, IP and CABAC respectively. Each optimization builds on the ones in the previous stage with the first bar in each set representing RISC energy dissipation followed by generic optimizations such as SIMD and VLIW, operation fusion and ending with “magic” instructions

~ 10x in energy (left) & time (right)
from “leaner” processor design

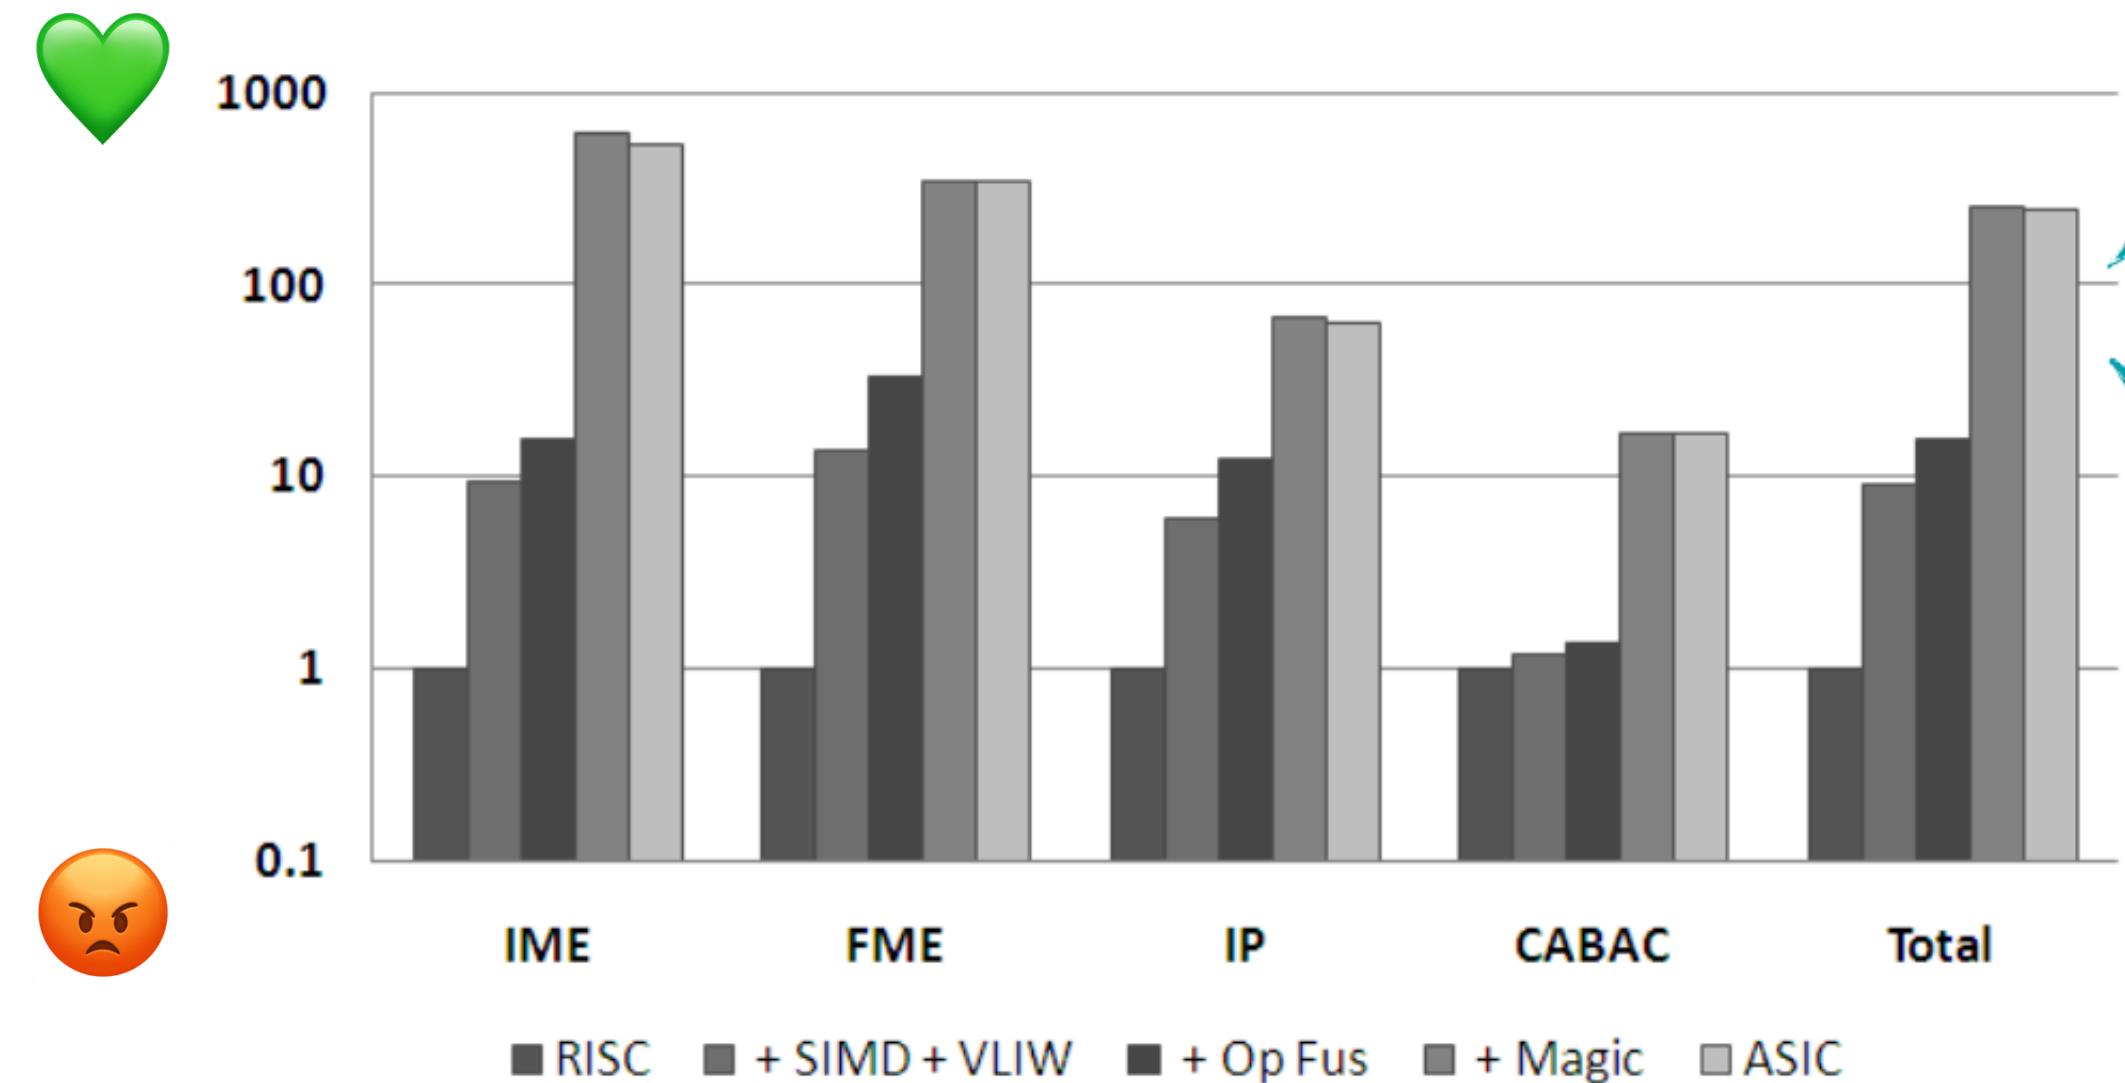


Figure 3. Each set of bar graphs represents speedup at each stage of optimization. Each optimization builds on those of the previous stage with the first bar in each set representing RISC speedup, followed by generic optimizations such as SIMD and VLIW, then operation fusion and finally “magic” instructions occurring complex instruction subgraphs. Operation fusion is

Example: Specialization (in hardware)

Understanding Sources of Inefficiency in General-Purpose Chips

Rehan Hameed¹, Wajahat Qadeer¹, Megan Wachs¹, Omid Azizi¹, Alex Solomatnikov², Benjamin C. Lee¹, Stephen Richardson¹, Christos Kozyrakis¹ and Mark Horowitz¹

¹Dept. of Electrical Engineering
Stanford University, Stanford, CA
{rhameed, wqadeer, wachs, oazizi,
bcclee, steve, kozyraki, horowitz}@stanford.edu

²Hicamp Systems,
Menlo Park, CA
solomatnikov@gmail.com

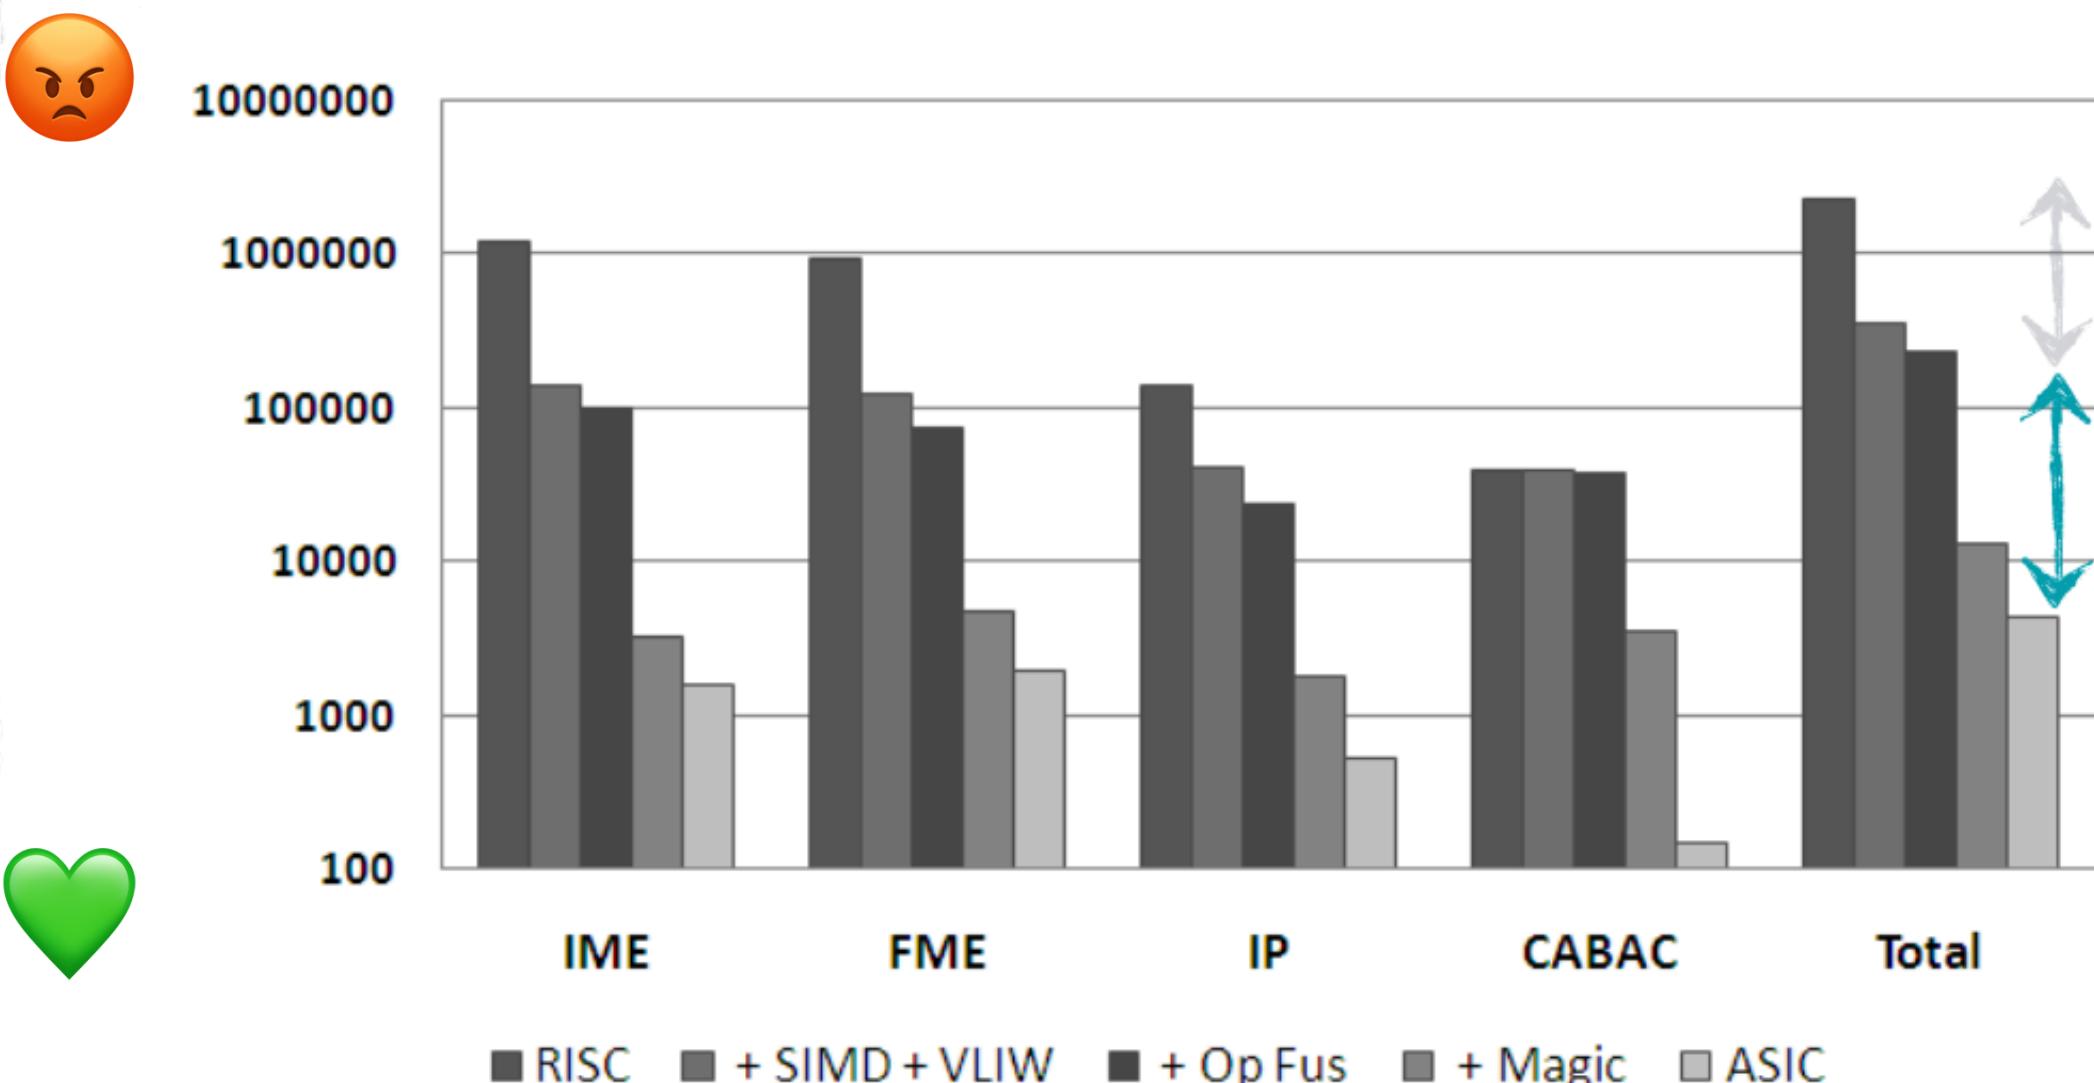


Figure 2. Each set of bar graphs represents energy consumption (μJ) at each stage of optimization for IME, FME, IP and CABAC respectively. Each optimization builds on the ones in the previous stage with the first bar in each set representing RISC energy dissipation followed by generic optimizations such as SIMD and VLIW, operation fusion and ending with “magic” instructions

~ 10x in energy (left) & time (right)
from “leaner” processor design

~ 10x+ in energy (left) & time (right)
from aggressive customization

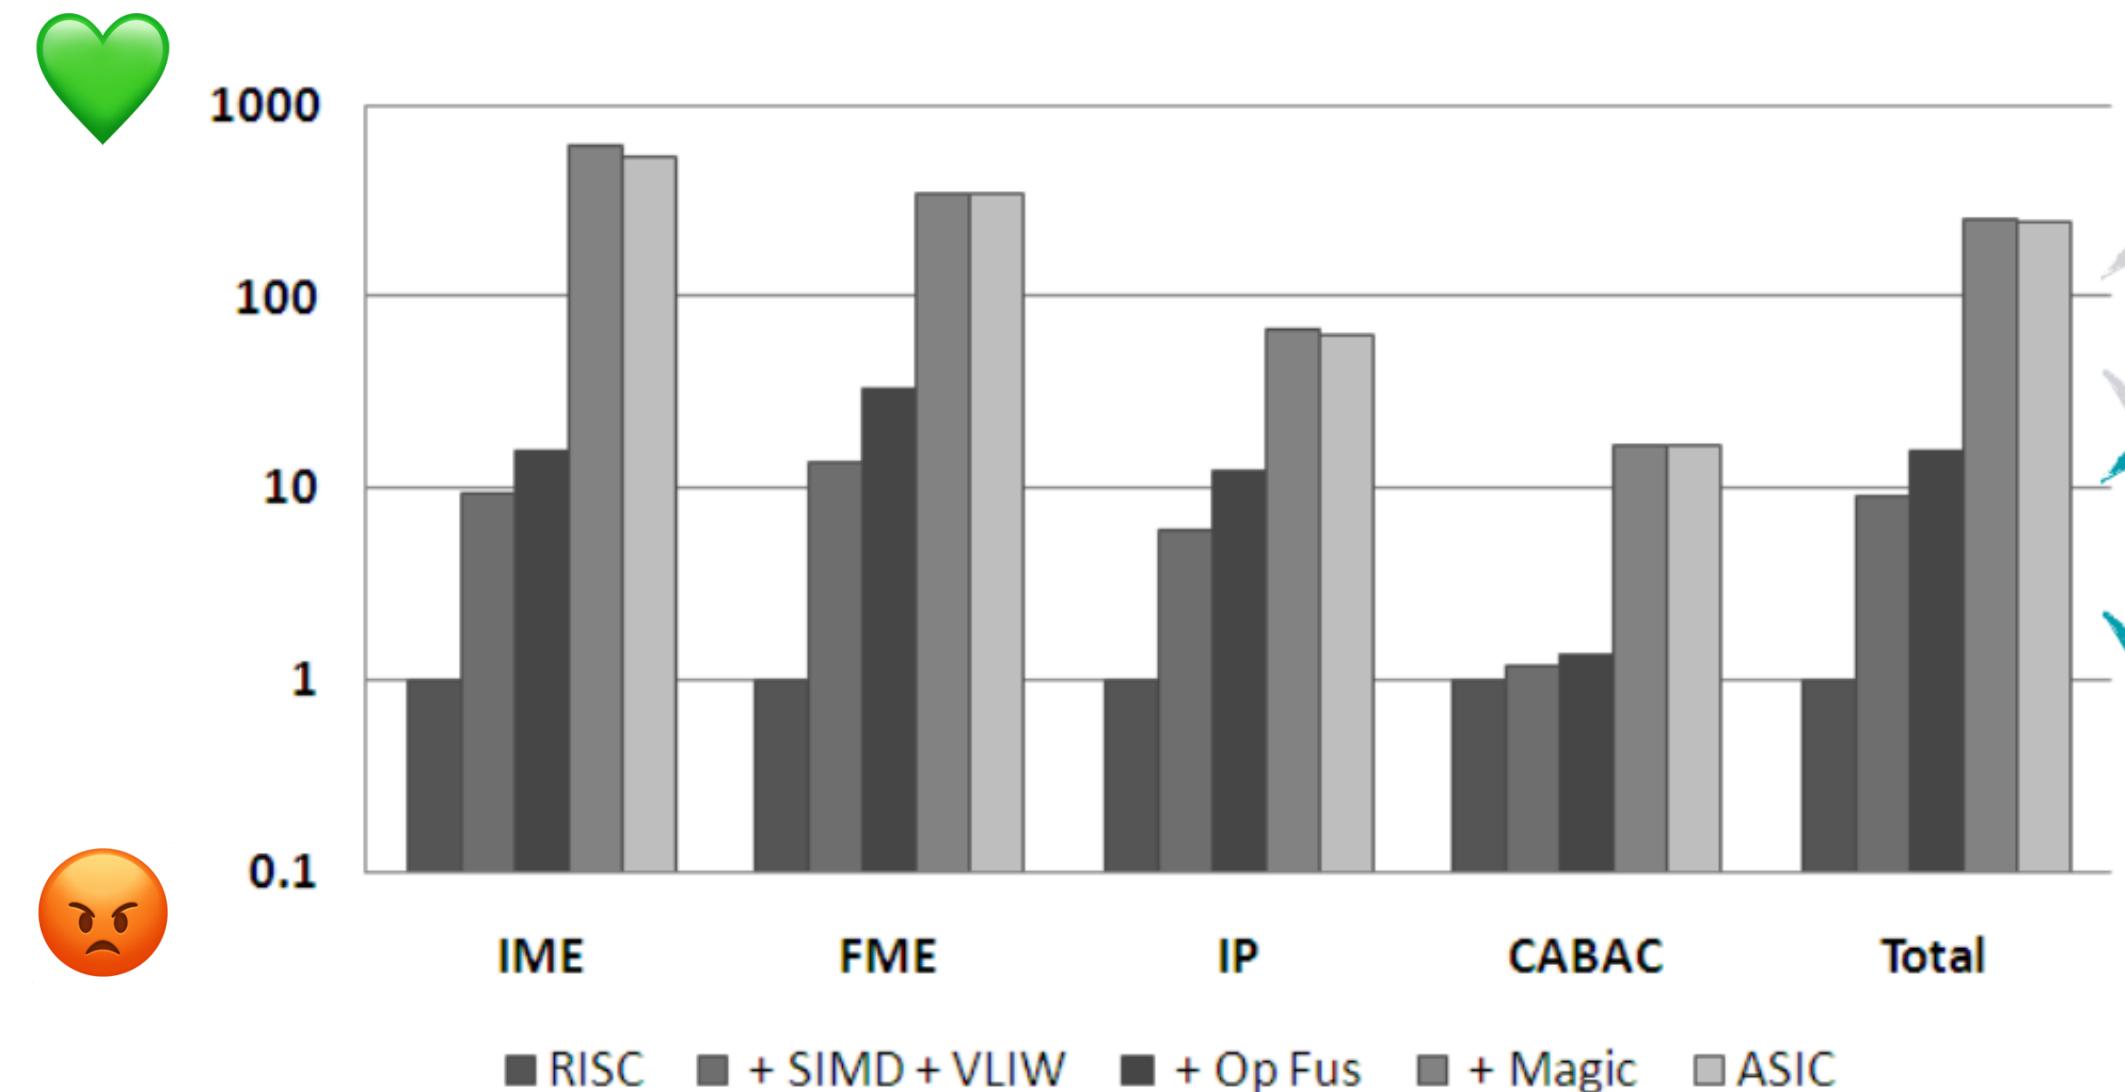
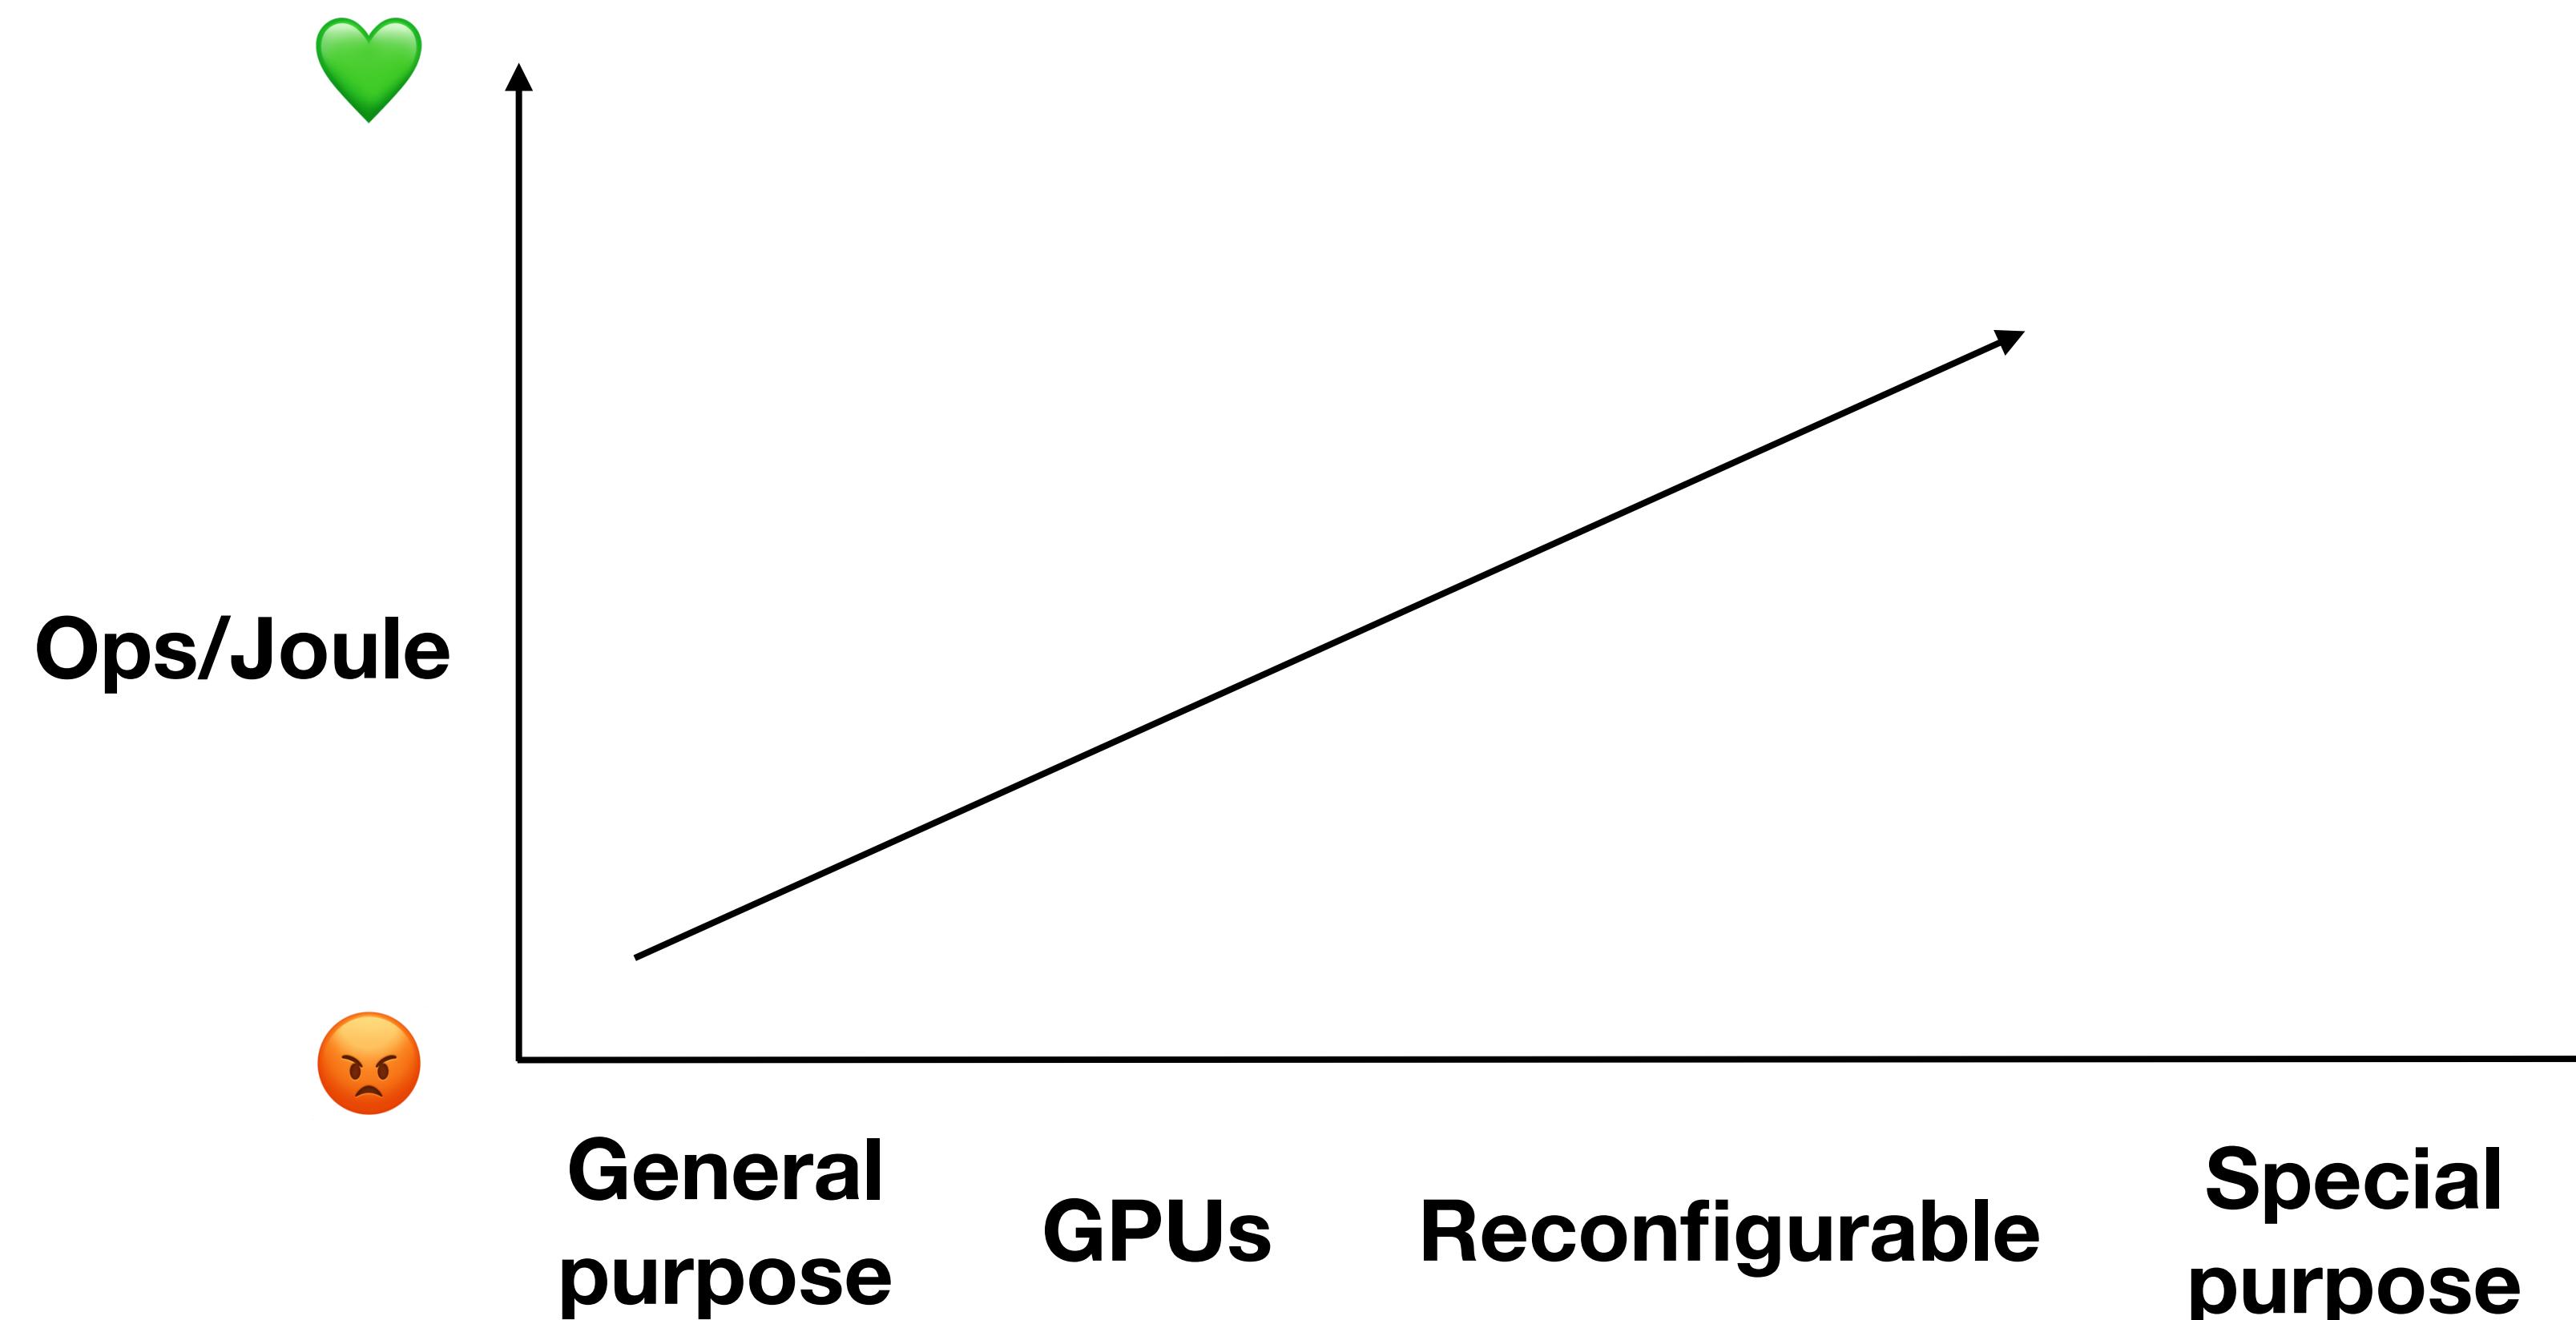
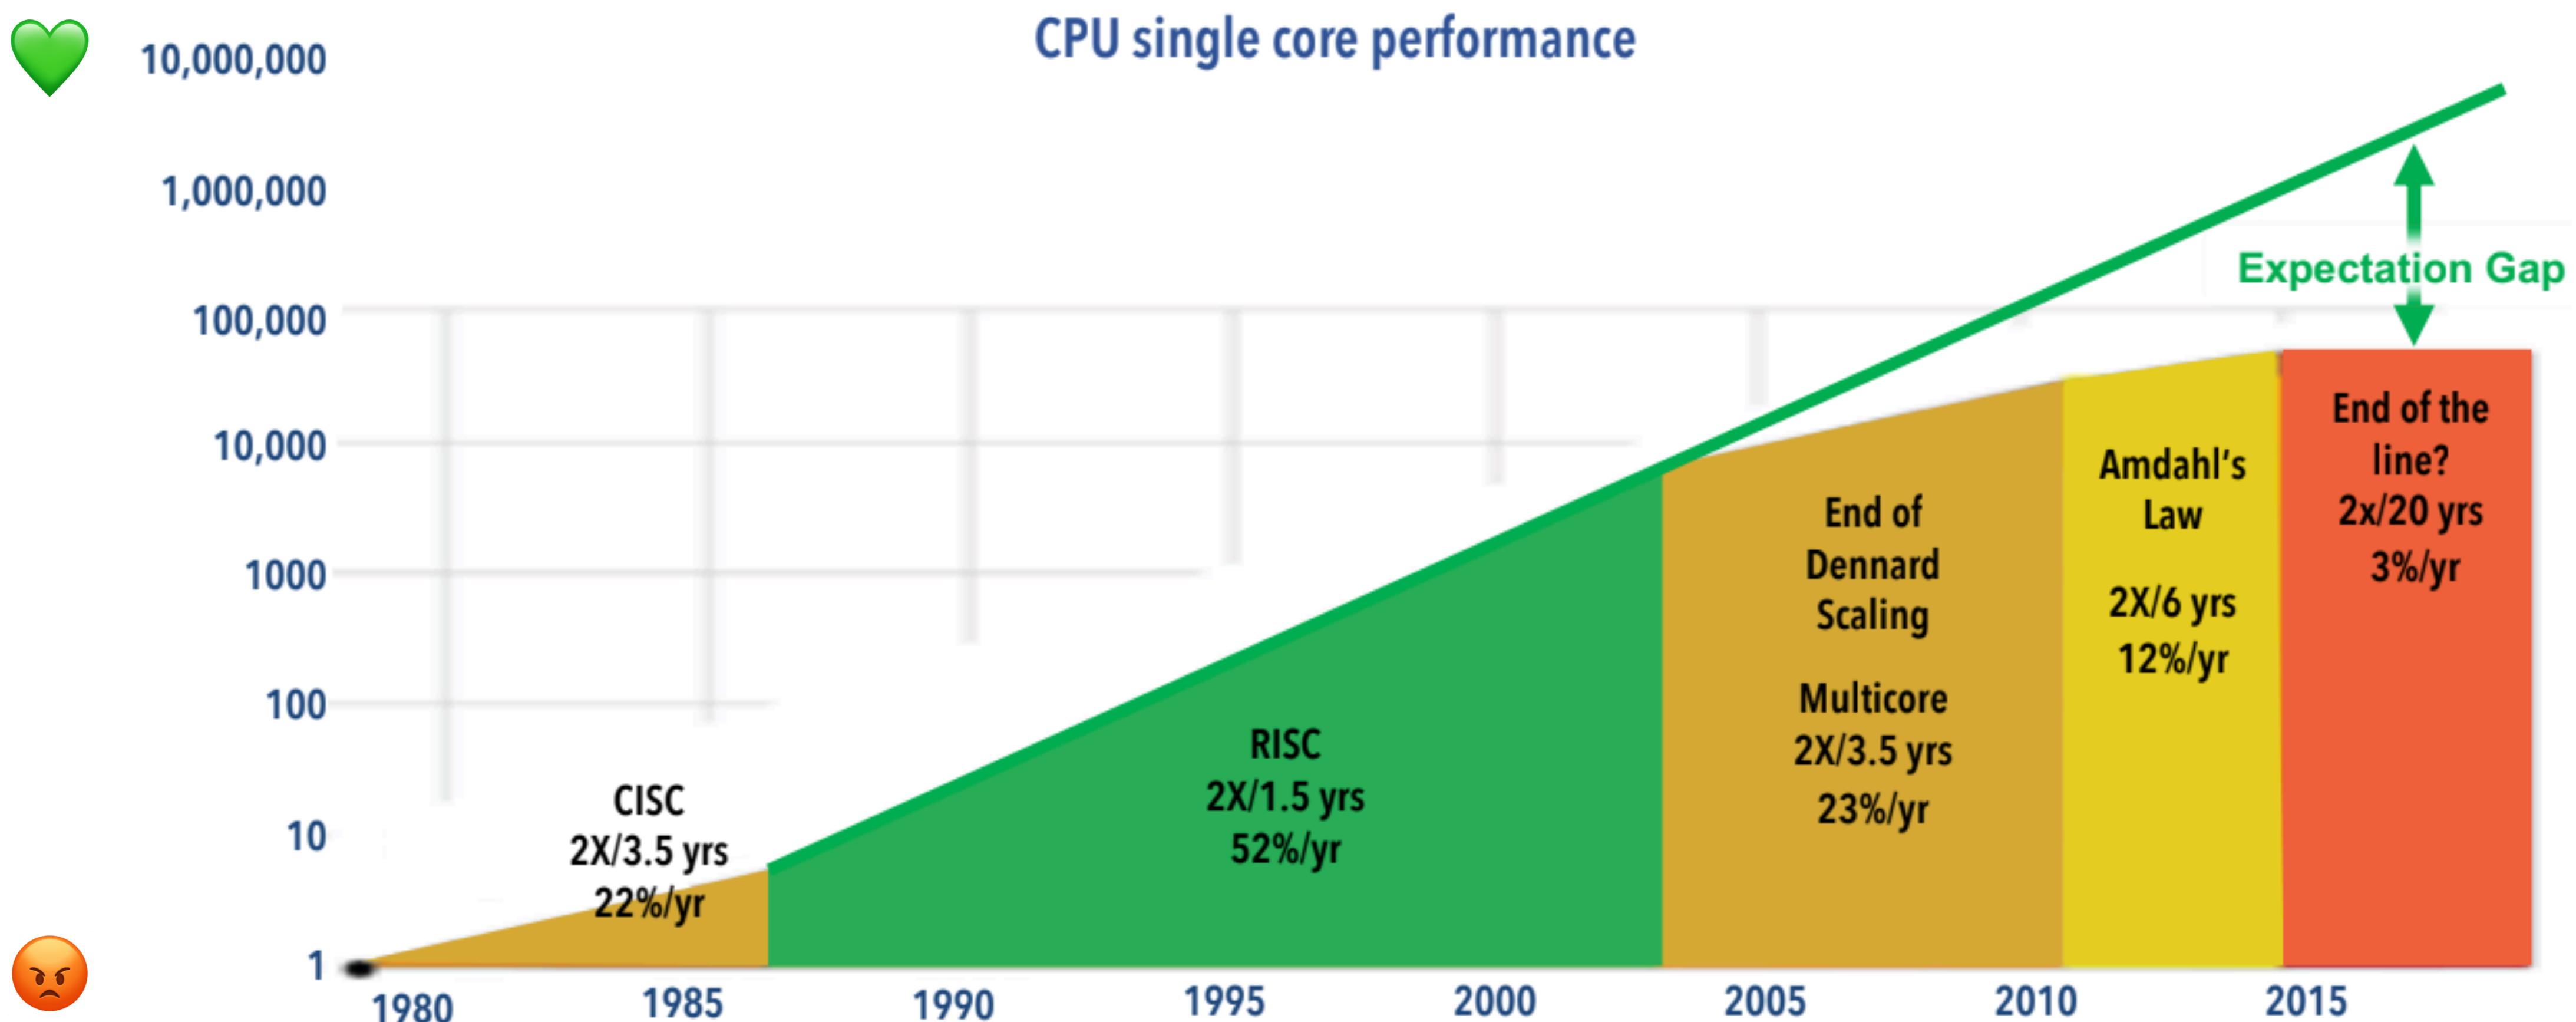


Figure 3. Each set of bar graphs represents speedup at each stage of optimization. Each optimization builds on those of the previous stage with the first bar in each set representing RISC speedup, followed by generic optimizations such as SIMD and VLIW, then operation fusion and finally “magic” instructions occurring complex instruction subgraphs. Operation fusion is

Specialization: End game for Moore's Law

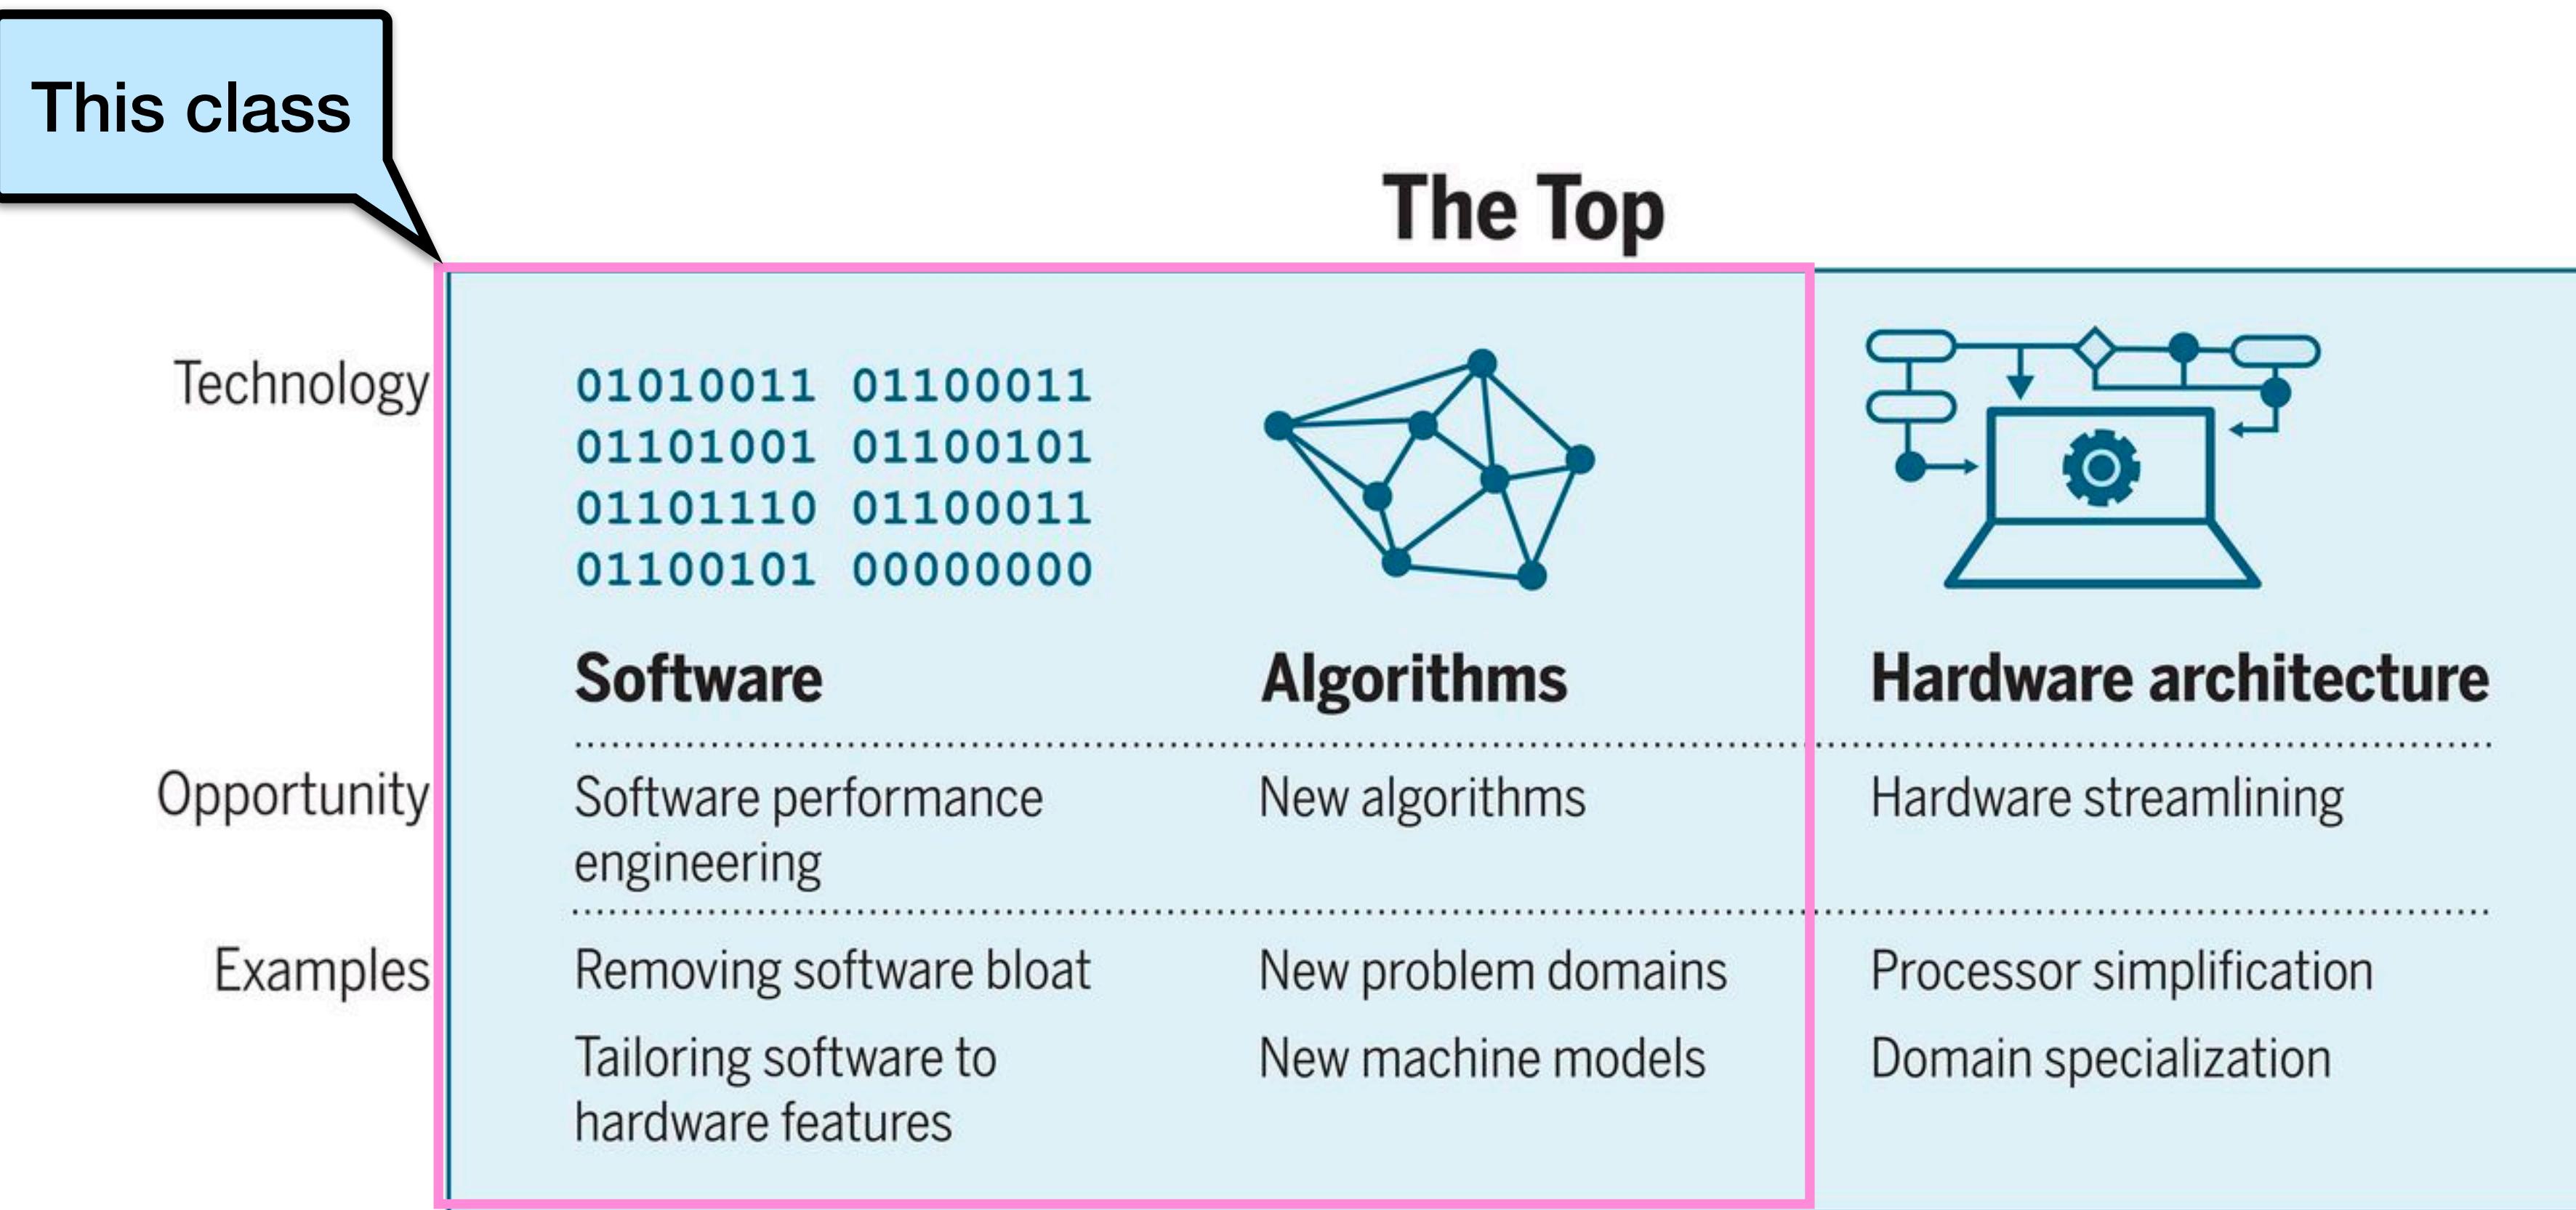


Traditional scaling is coming to an end



Hennessey/Patterson Turning Award Presentation: Golden Age of Computer Architecture Research

Performance gains after Moore's law ends



The Bottom

for example, semiconductor technology