

# CSE 6220/CX 4220

## Introduction to HPC

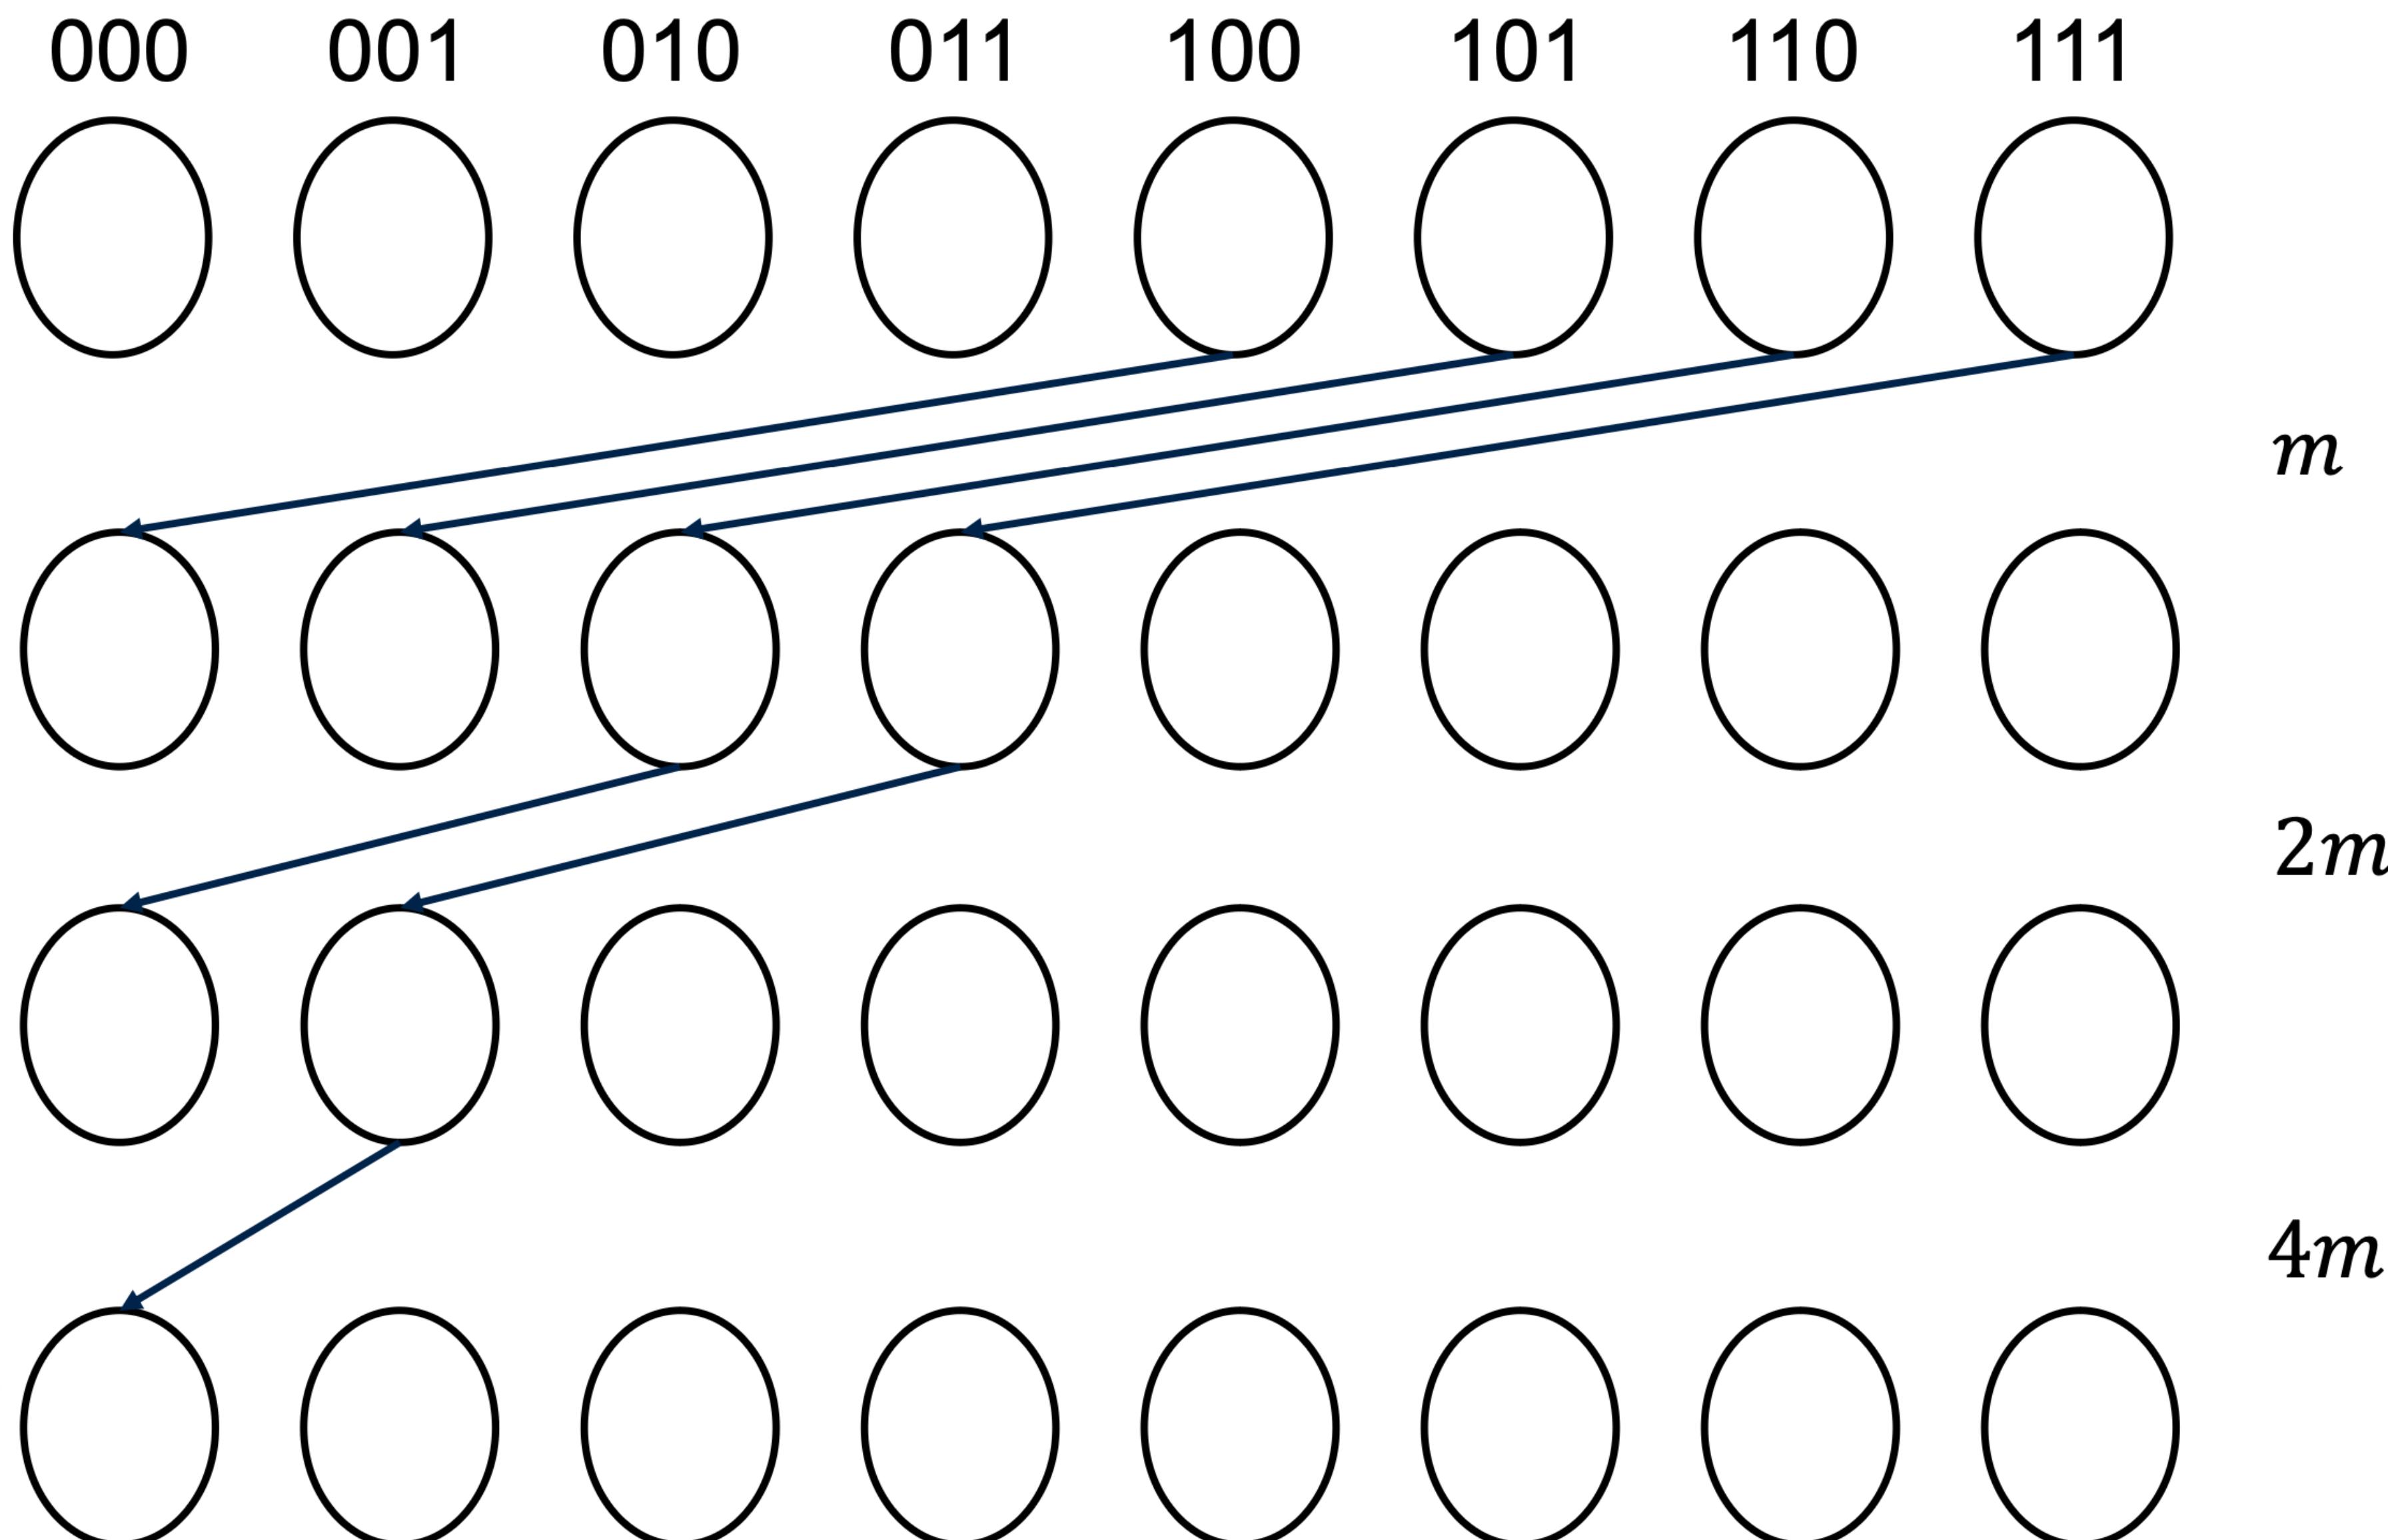
# Lecture 8: Applications of Prefix Sums

Helen Xu  
[hxu615@gatech.edu](mailto:hxu615@gatech.edu)



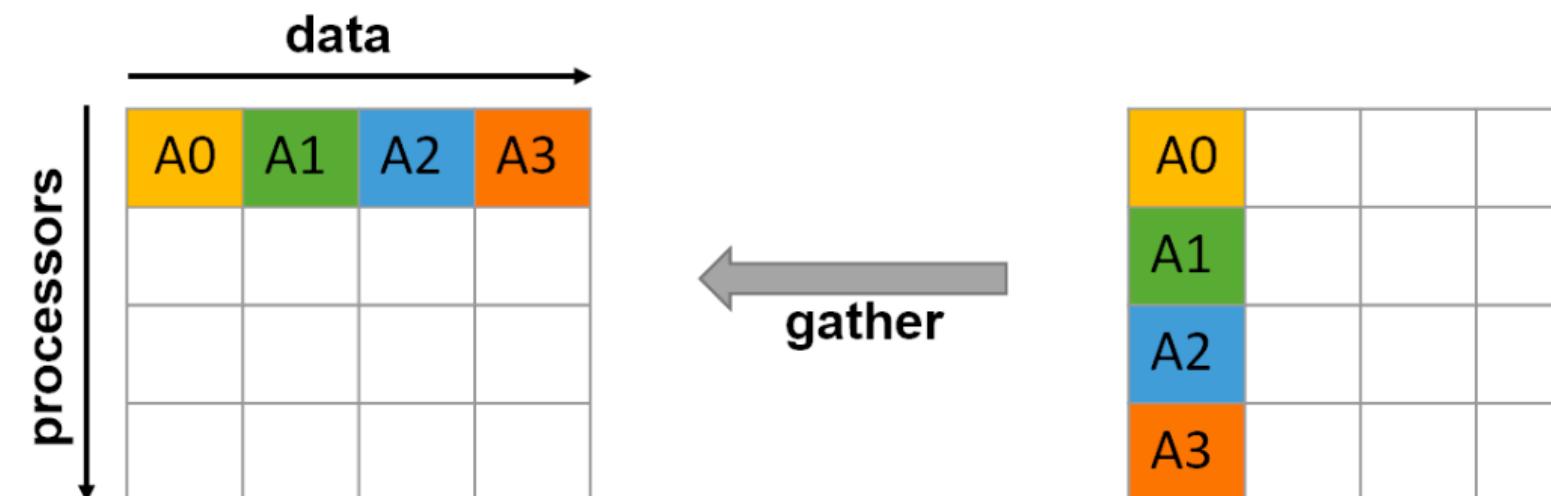
Georgia Tech College of Computing  
School of Computational  
Science and Engineering

# Super fast sidebar into Gather (for PA1)



# Super fast sidebar into Gather (for PA1)

Gather data from all processes on one process(=root)



```
int MPI_Gather(void* sbuf, int scount, MPI_Datatype stype,
               void* rbuf, int rcount, MPI_Datatype rtype,
               int root, MPI_Comm comm);
```

```
#include <mpi.h>
#include <iostream >
#include <vector >
int main(int argc, char* argv[]) {
    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    std::vector<int> buf;
    if (rank == 0) buf.resize(size, -1);
    MPI_Gather(&rank, 1, MPI_INT,
               &buf[0], 1, MPI_INT, 0, MPI_COMM_WORLD);
    if (rank == 0) std::cout << buf.back() << std::endl;
    return MPI_Finalize();
}
```

At this point, what is the content of the vector buf?

# Recap of Shared-Memory Parallel Prefix

# Work-span examples

```
for(int i = 0; i < n; i++) {  
    parallel_for(int j = 0; j < N; j++ {  
        // O(1)  
    }  
}
```

What is the work and span?

# Work-span examples

```
for(int i = 0; i < n; i++) {  
    parallel_for(int j = 0; j < N; j++ {  
        // O(1)  
    }  
}
```

What is the work and span?

Work =  $\Theta(n^2)$

Span =  $\Theta(n \lg n)$

# Work-span examples

```
for(int i = 0; i < n; i++) {  
    parallel_for(int j = 0; j < N; j++ {  
        parallel_for(int k = 0; k < N; k++ {  
            // O(1)  
        }  
    }  
}
```

What is the work and span?

# Work-span examples

```
for(int i = 0; i < n; i++) {  
    parallel_for(int j = 0; j < N; j++ {  
        parallel_for(int k = 0; k < N; k++ {  
            // O(1)  
        }  
    }  
}
```

What is the work and span?

$$\text{Work} = \Theta(n^3)$$

$$\text{Span} = \Theta(n \cdot (\lg n + \lg n)) = \Theta(n \lg n)$$

# Work-span examples

```
for(int i = 0; i < n; i++) {  
    parallel_for(int j = 0; j < N; j++ {  
        for(int k = 0; k < N; k++ {  
            // O(1)  
        }  
    }  
}
```

What is the work and span?

# Work-span examples

```
for(int i = 0; i < n; i++) {  
    parallel_for(int j = 0; j < N; j++ {  
        for(int k = 0; k < N; k++ {  
            // O(1)  
        }  
    }  
}
```

What is the work and span?

$$\text{Work} = \Theta(n^3)$$

$$\text{Span} = \Theta(n(\lg n + n)) = \Theta(n^2)$$

# Hillis-Steele Prefix Sum

for i = 0 up to  $\log(n)$ :

**parallel\_for** j = 0 up to n-1:

if  $j < 2^i$ :

$$x_j^{i+1} \leftarrow x_j^i$$

else:

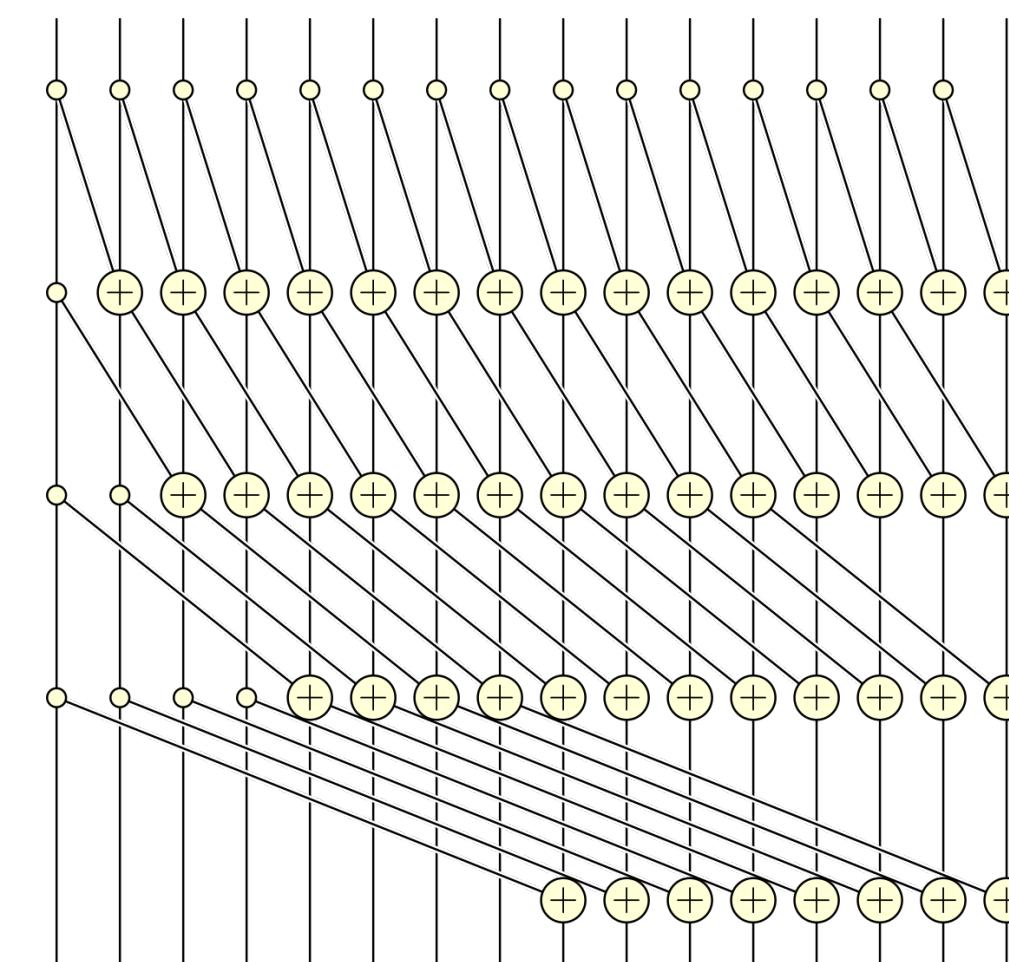
$$x_j^{i+1} \leftarrow x_j^i + x_{j-2^i}^i$$

What is the work and span?

Work =  $O(n \lg n)$

Span =  $O(\lg^2 n)$

Better, but still not  
work efficient



# Work-Efficient Parallel Prefix Pseudocode

Upsweep:

for  $d$  from 0 to  $\lg(n) - 1$ :

parallel\_for  $i$  from 0 to  $n - 1$ ,  $i += 2^{d+1}$ :

$A[i + 2^{d+1} - 1] \leftarrow A[i + 2^d - 1] + A[i + 2^{d+1} - 1]$

Downsweep:

for  $d$  from  $\lg(n) - 1$  to 0:

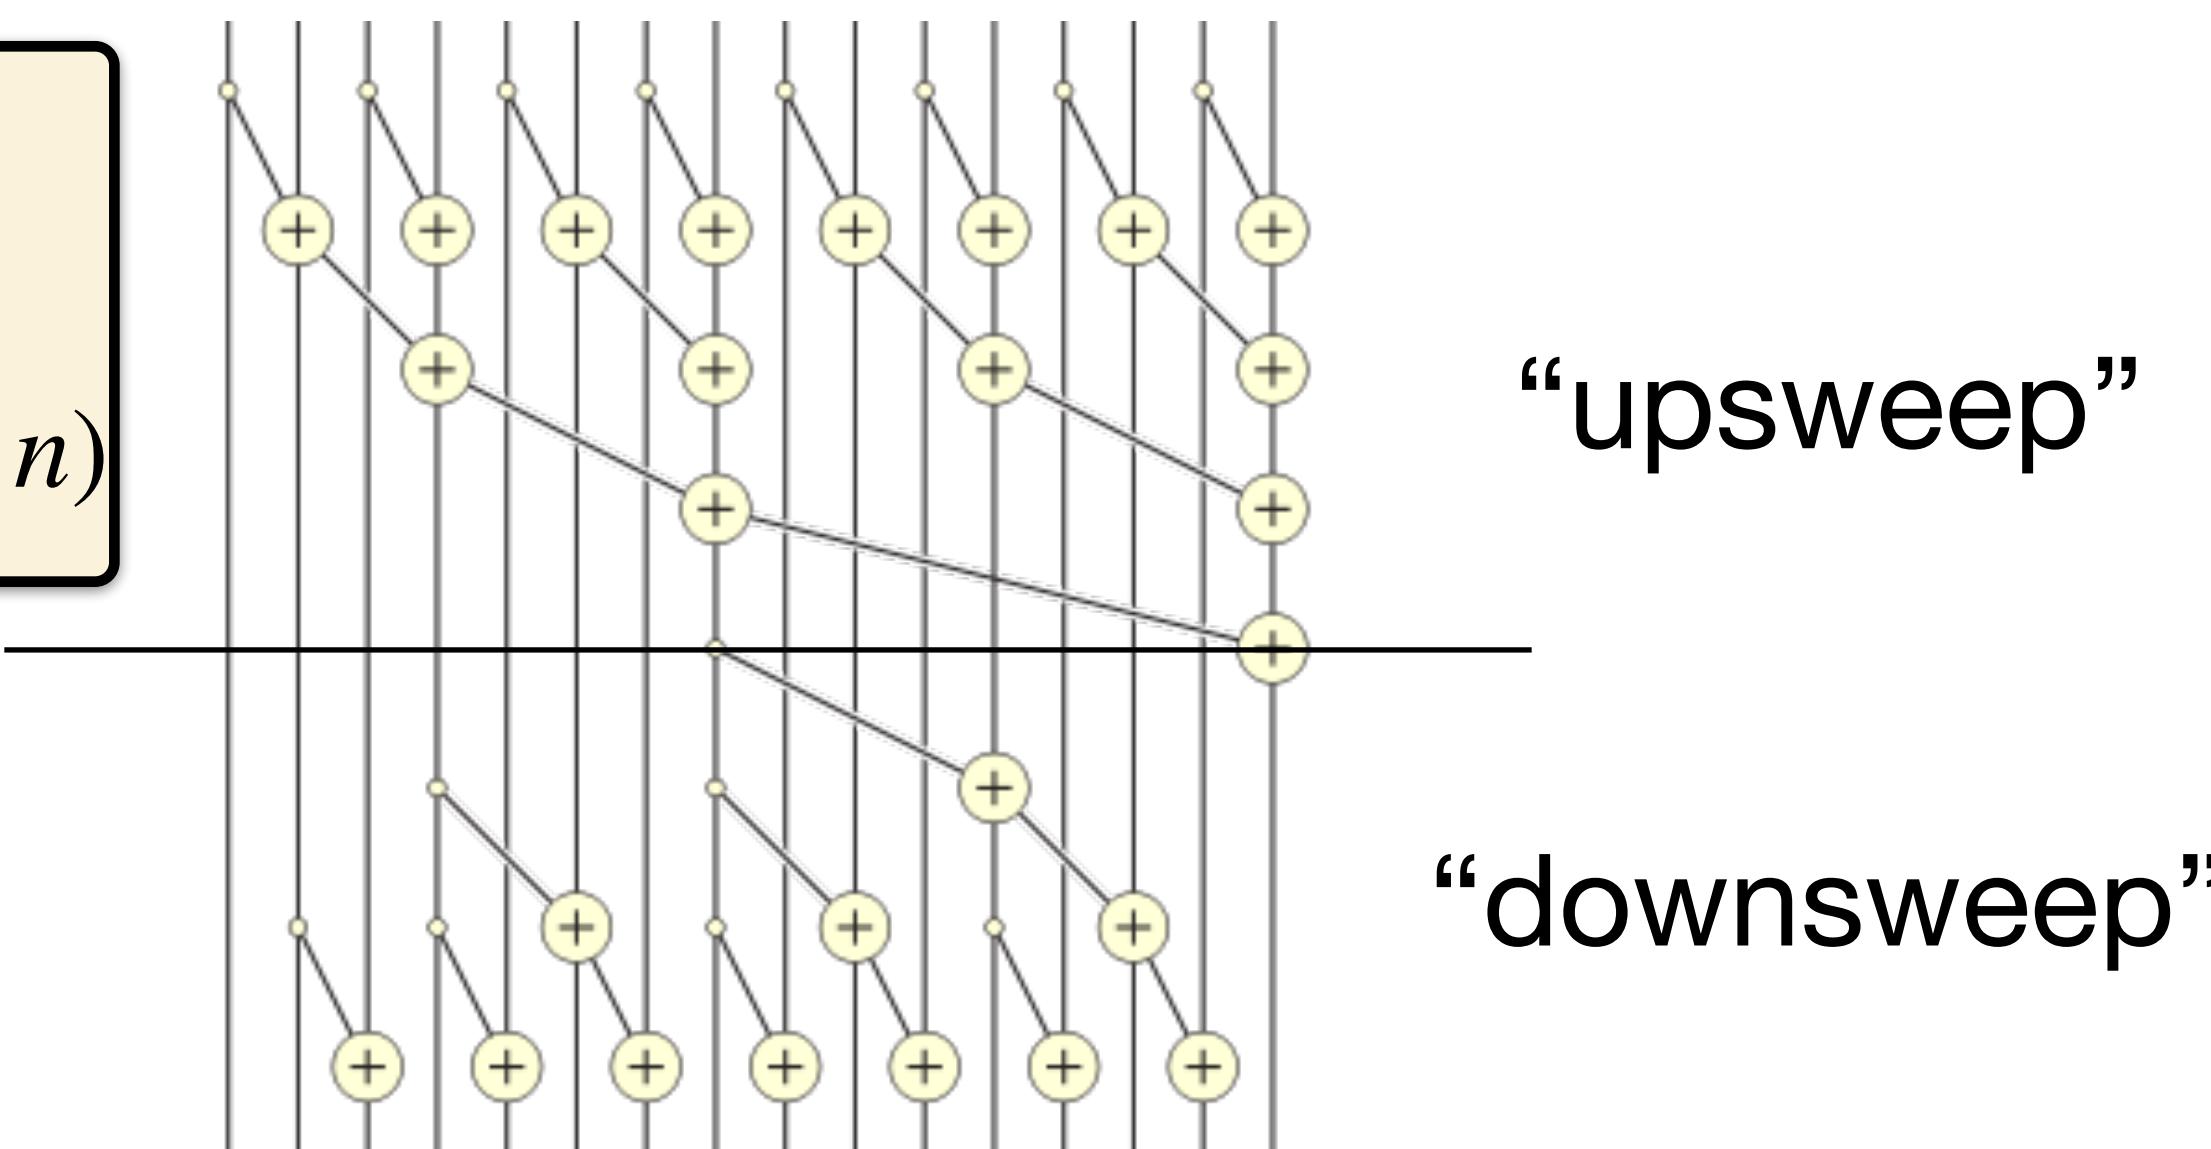
parallel\_for  $i$  from  $2^d - 1$  to  $n - 1 - 2^d$ ,  $i += 2^{d+1}$ :

if  $i - 2^d \geq 0$ :

$A[i] = A[i] + A[i - 2^d]$

Span in looping formulation  
follows the expression:

$$\lg(n/2) + \lg(n/4) + \dots + 1 = \Theta(\lg^2 n)$$



# **More Applications of Parallel Prefix**

# Application: Radix Sort (Serial)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

 $b_0 = 0$  $b_0 = 1$ 

Sort on least significant bit ( $b_0$  in  $b_2b_1b_0$ )  
 $\text{XX0} < \text{XX1}$  (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

 $b_1 = 0$  $b_1 = 1$ 

Stably sort entire array on next bit  
 $\text{X0X} < \text{X1X}$

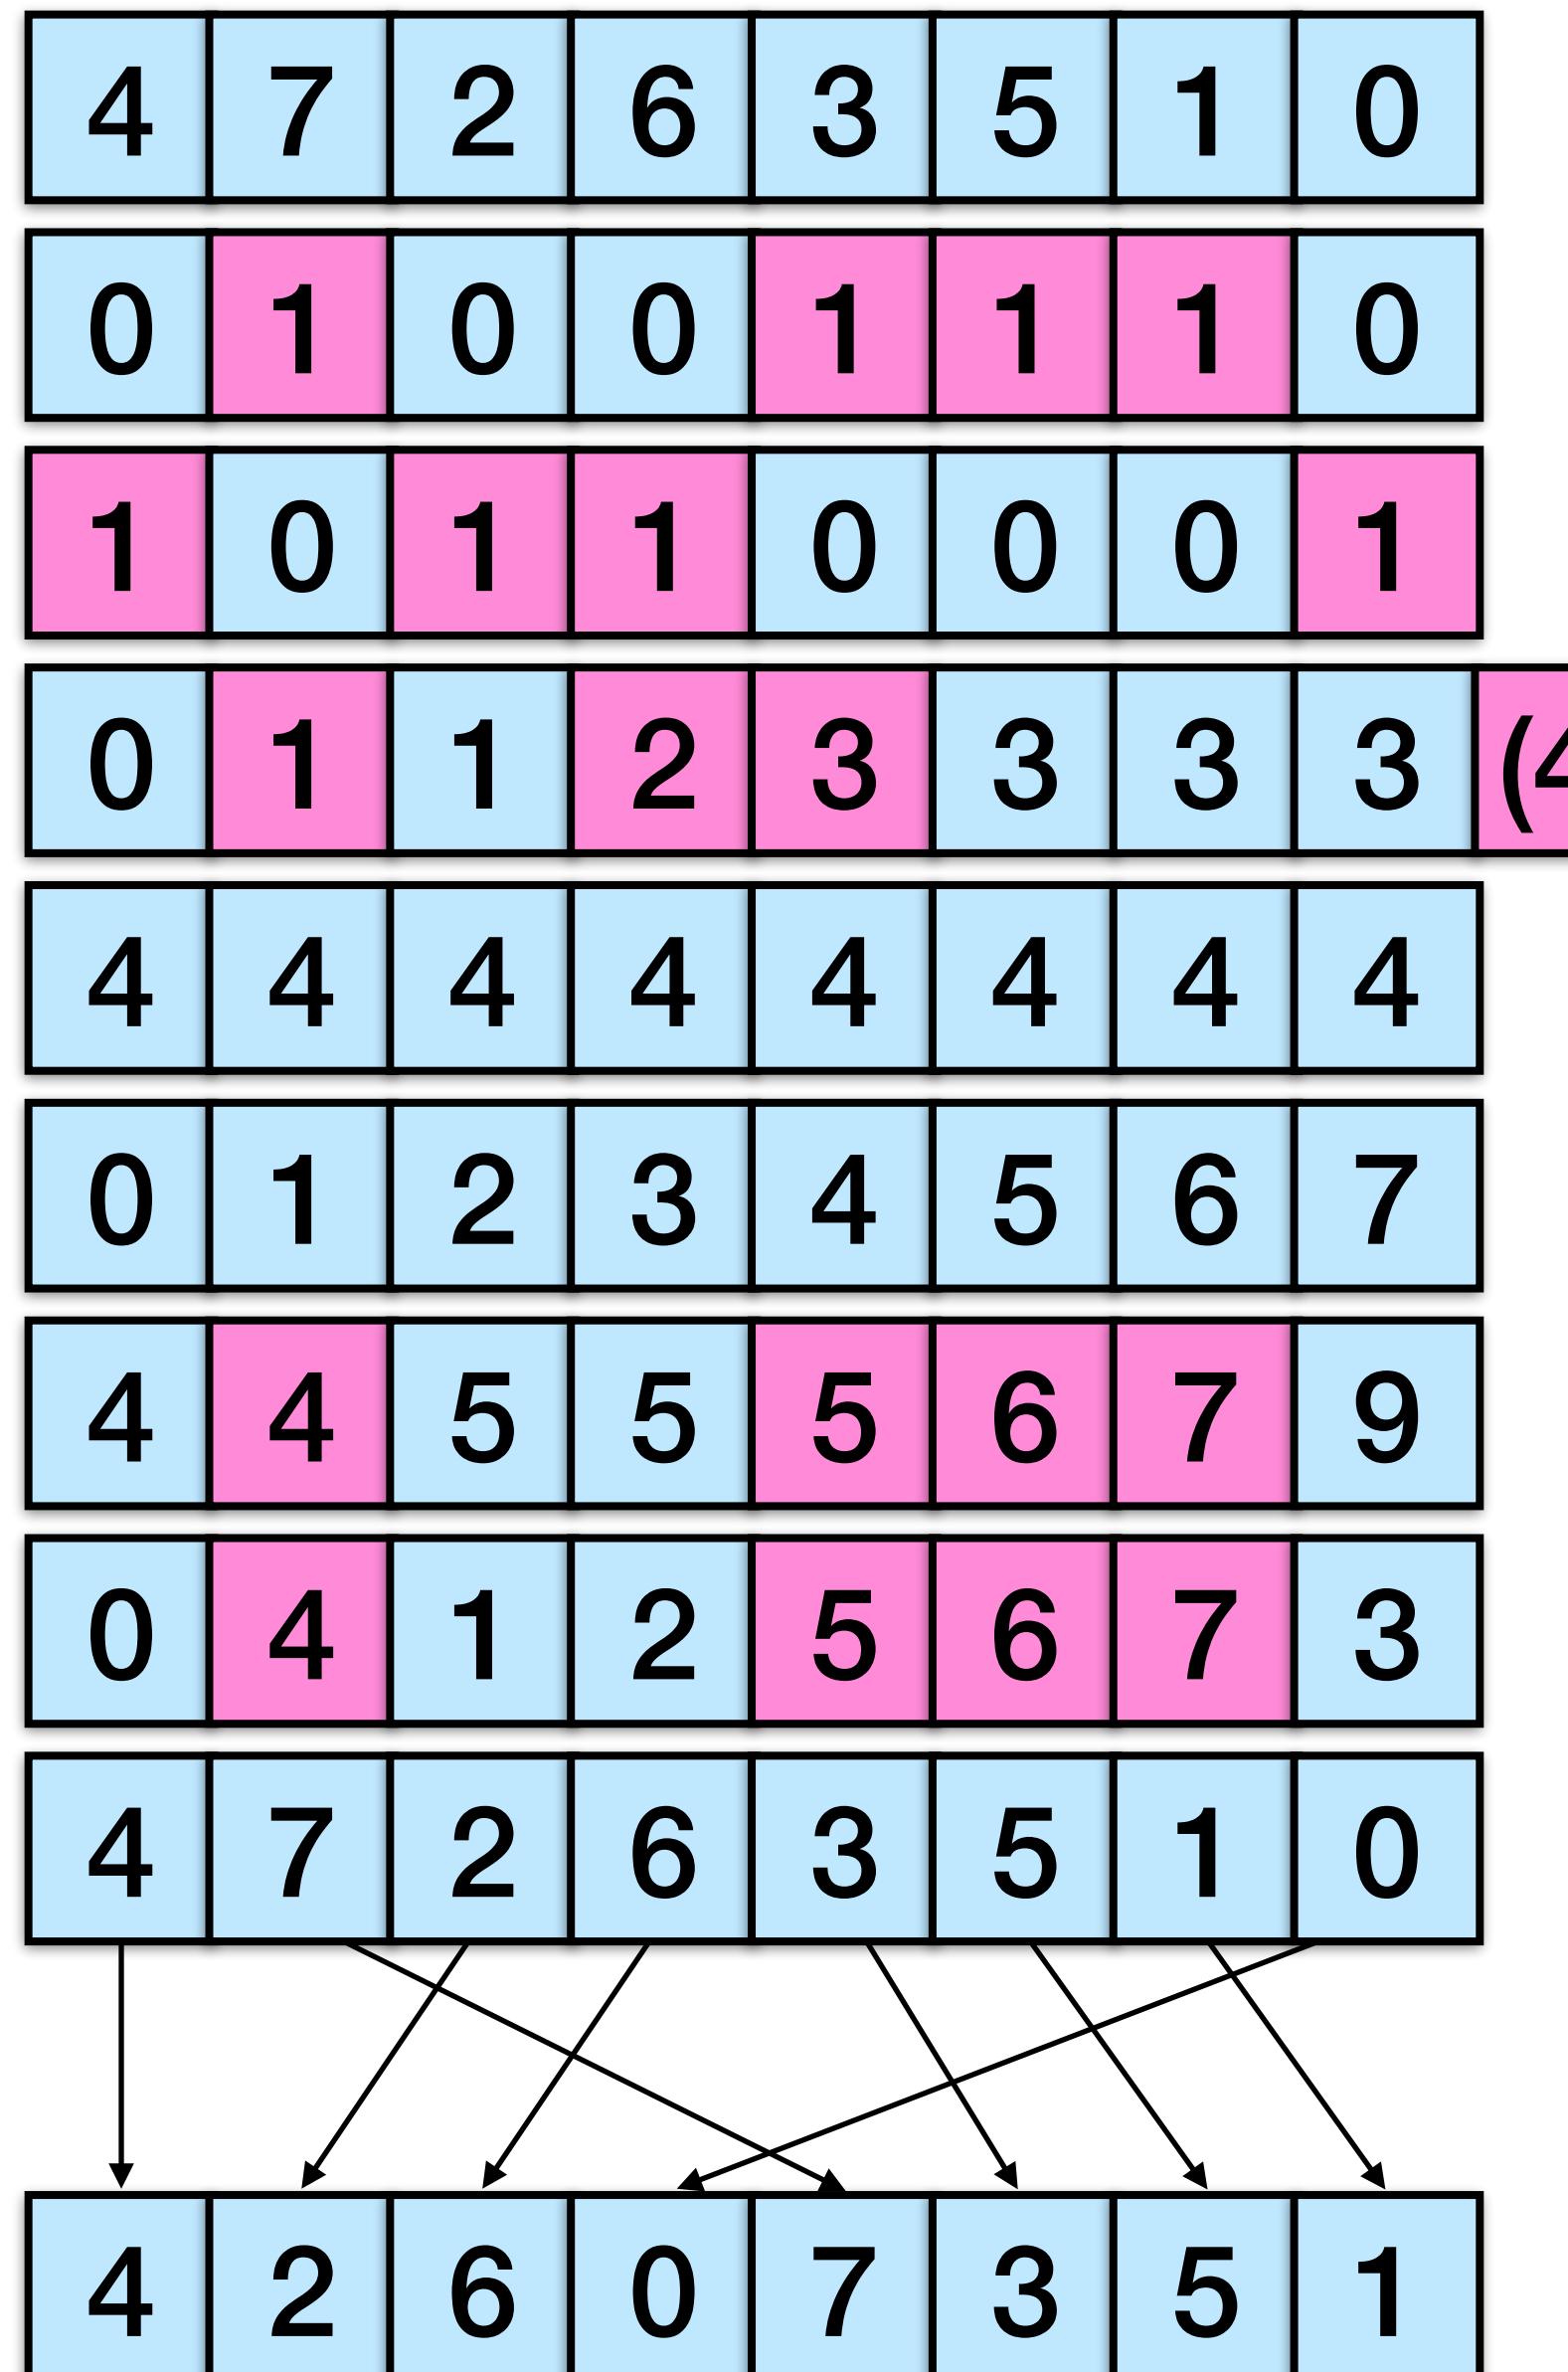
4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

 $b_2 = 0$  $b_2 = 1$ 

Stably sort entire array on next bit  
 $0\text{XX} < 1\text{XX}$

# Application: Data-Parallel Radix Sort



Input

Odds = last bit of each element

Evens = complement of odds

Even\_positions = exclusive scan of evens

totalEvens = broadcast last element

index = constant array of 0 .. n

odd\_positions = #evens + idx - even\_pos

pos = get positions using masked assignment

Scatter input according to pos

(repeat with next bit until you are out of bits)

# Analyzing Data-Parallel Radix Sort

**ALGORITHM:** RADIX-SORT( $A, b$ )

```
1  for  $i$  from 0 to  $b - 1$ 
2    FLAGS :=  $\{(a >> i) \bmod 2 : a \in A\}$ 
3    NOTFLAGS :=  $\{1 - b : b \in \text{FLAGS}\}$ 
4     $R_0 := \text{SCAN}(\text{NOTFLAGS})$ 
5     $s_0 := \text{SUM}(\text{NOTFLAGS})$ 
6     $R_1 := \text{SCAN}(\text{FLAGS})$ 
7     $R := \{\text{if } \text{FLAGS}[j] = 0 \text{ then } R_0[j] \text{ else } R_1[j] + s_0 : j \in [0..|A|]\}$ 
8     $A := A \leftarrow \{(R[j], A[j]) : j \in [0..|A|]\}$ 
9  return  $A$ 
```

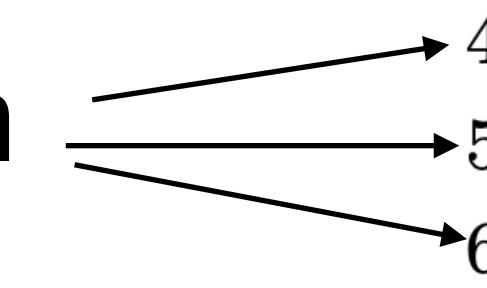
What is the work and span?

# Analyzing Data-Parallel Radix Sort

**ALGORITHM:** RADIX-SORT( $A, b$ )

- 1 **for**  $i$  **from** 0 **to**  $b - 1$
- 2     FLAGS :=  $\{(a >> i) \bmod 2 : a \in A\}$
- 3     NOTFLAGS :=  $\{1 - b : b \in \text{FLAGS}\}$
- 4      $R_0 := \text{SCAN}(\text{NOTFLAGS})$
- 5      $s_0 := \text{SUM}(\text{NOTFLAGS})$
- 6      $R_1 := \text{SCAN}(\text{FLAGS})$
- 7      $R := \{\text{if } \text{FLAGS}[j] = 0 \text{ then } R_0[j] \text{ else } R_1[j] + s_0 : j \in [0..|A|]\}$
- 8      $A := A \leftarrow \{(R[j], A[j]) : j \in [0..|A|]\}$
- 9 **return**  $A$

$O(\log n)$  span



What is the work and span?

Work: There are  $b$  iterations of the for loop, and iteration takes  $O(n)$  work, so the total work is  $O(bn)$ .

Span: There are  $b$  iterations of the for loop, and iteration has  $O(\log n)$  span, so the total span is  $O(b \log n)$ .

# Application: Adding n-bit Integers

Problem: Computing sum of two n-bit binary numbers,  $a$  and  $b$ .

$$a = a_{n-1}a_{n-2}\dots a_0 \text{ and } b = b_{n-1}b_{n-2}\dots b_0$$

$$s = a + b = s_ns_{n-1}\dots s_0 \text{ (using carry bit array } c = c_{n-1}, \dots, c_0, c_{-1})$$

```
c[-1] = 0 // rightmost carry bit
for i = 0 to n - 1: //compute right to left
    s[i] = (a[i] xor b[i]) xor c[i-1] // one or three 1s
    c[i] = ((a[i] xor b[i]) and c[i-1]) or (a[i] and b[i]) // next carry bit
```

Example:

$$a = 22$$

$$b = 29$$

	a	1 0 1 1 0	$s[0]$ depends on these
	b	1 1 1 0 1	
	c	1 1 1 0 0 0	
	s	1 1 0 0 1 1	

Goal: Compute all  $c_i$  in  $O(\log n)$  span via parallel prefix

# Application: Adding n-bit Integers

```
c[-1] = 0 // rightmost carry bit  
for i = 0 to n - 1: //compute right to left  
    c[i] = ((a[i] xor b[i]) and c[i-1]) or (a[i] and b[i]) // next carry bit
```

Idea: Split carry bit into two cases that indicate information about the carry-out ( $c_i$ ) regardless of carry-in ( $c_{i-1}$ ):

- Generate ( $g_i$ ): This column will generate a carry-out **whether or not** the carry-in is 1.

$$g_i = a_i \& \& b_i$$

- Propagate ( $p_i$ ): This column will propagate a carry-in **if there is one** to the carry-out.

$$p_i = a_i || b_i$$

can be computed in parallel

$$c_i = g_i + p_i c_{i-1}$$

# Carry Lookahead Logic

Idea: Define each carry-in in terms of  $p_i$ ,  $g_i$  and the initial carry in  $c_{i-1}$  and not in terms of carry chain (i.e., unwind the recursion):

- $c_0 = g_0 + p_0 c_{-1}$
- $c_1 = g_1 + p_1 c_0 = g_1 + p_1 g_0 + p_1 p_0 c_{-1}$
- ...

Can be expressed with 2-by-2 boolean matrix multiplication:

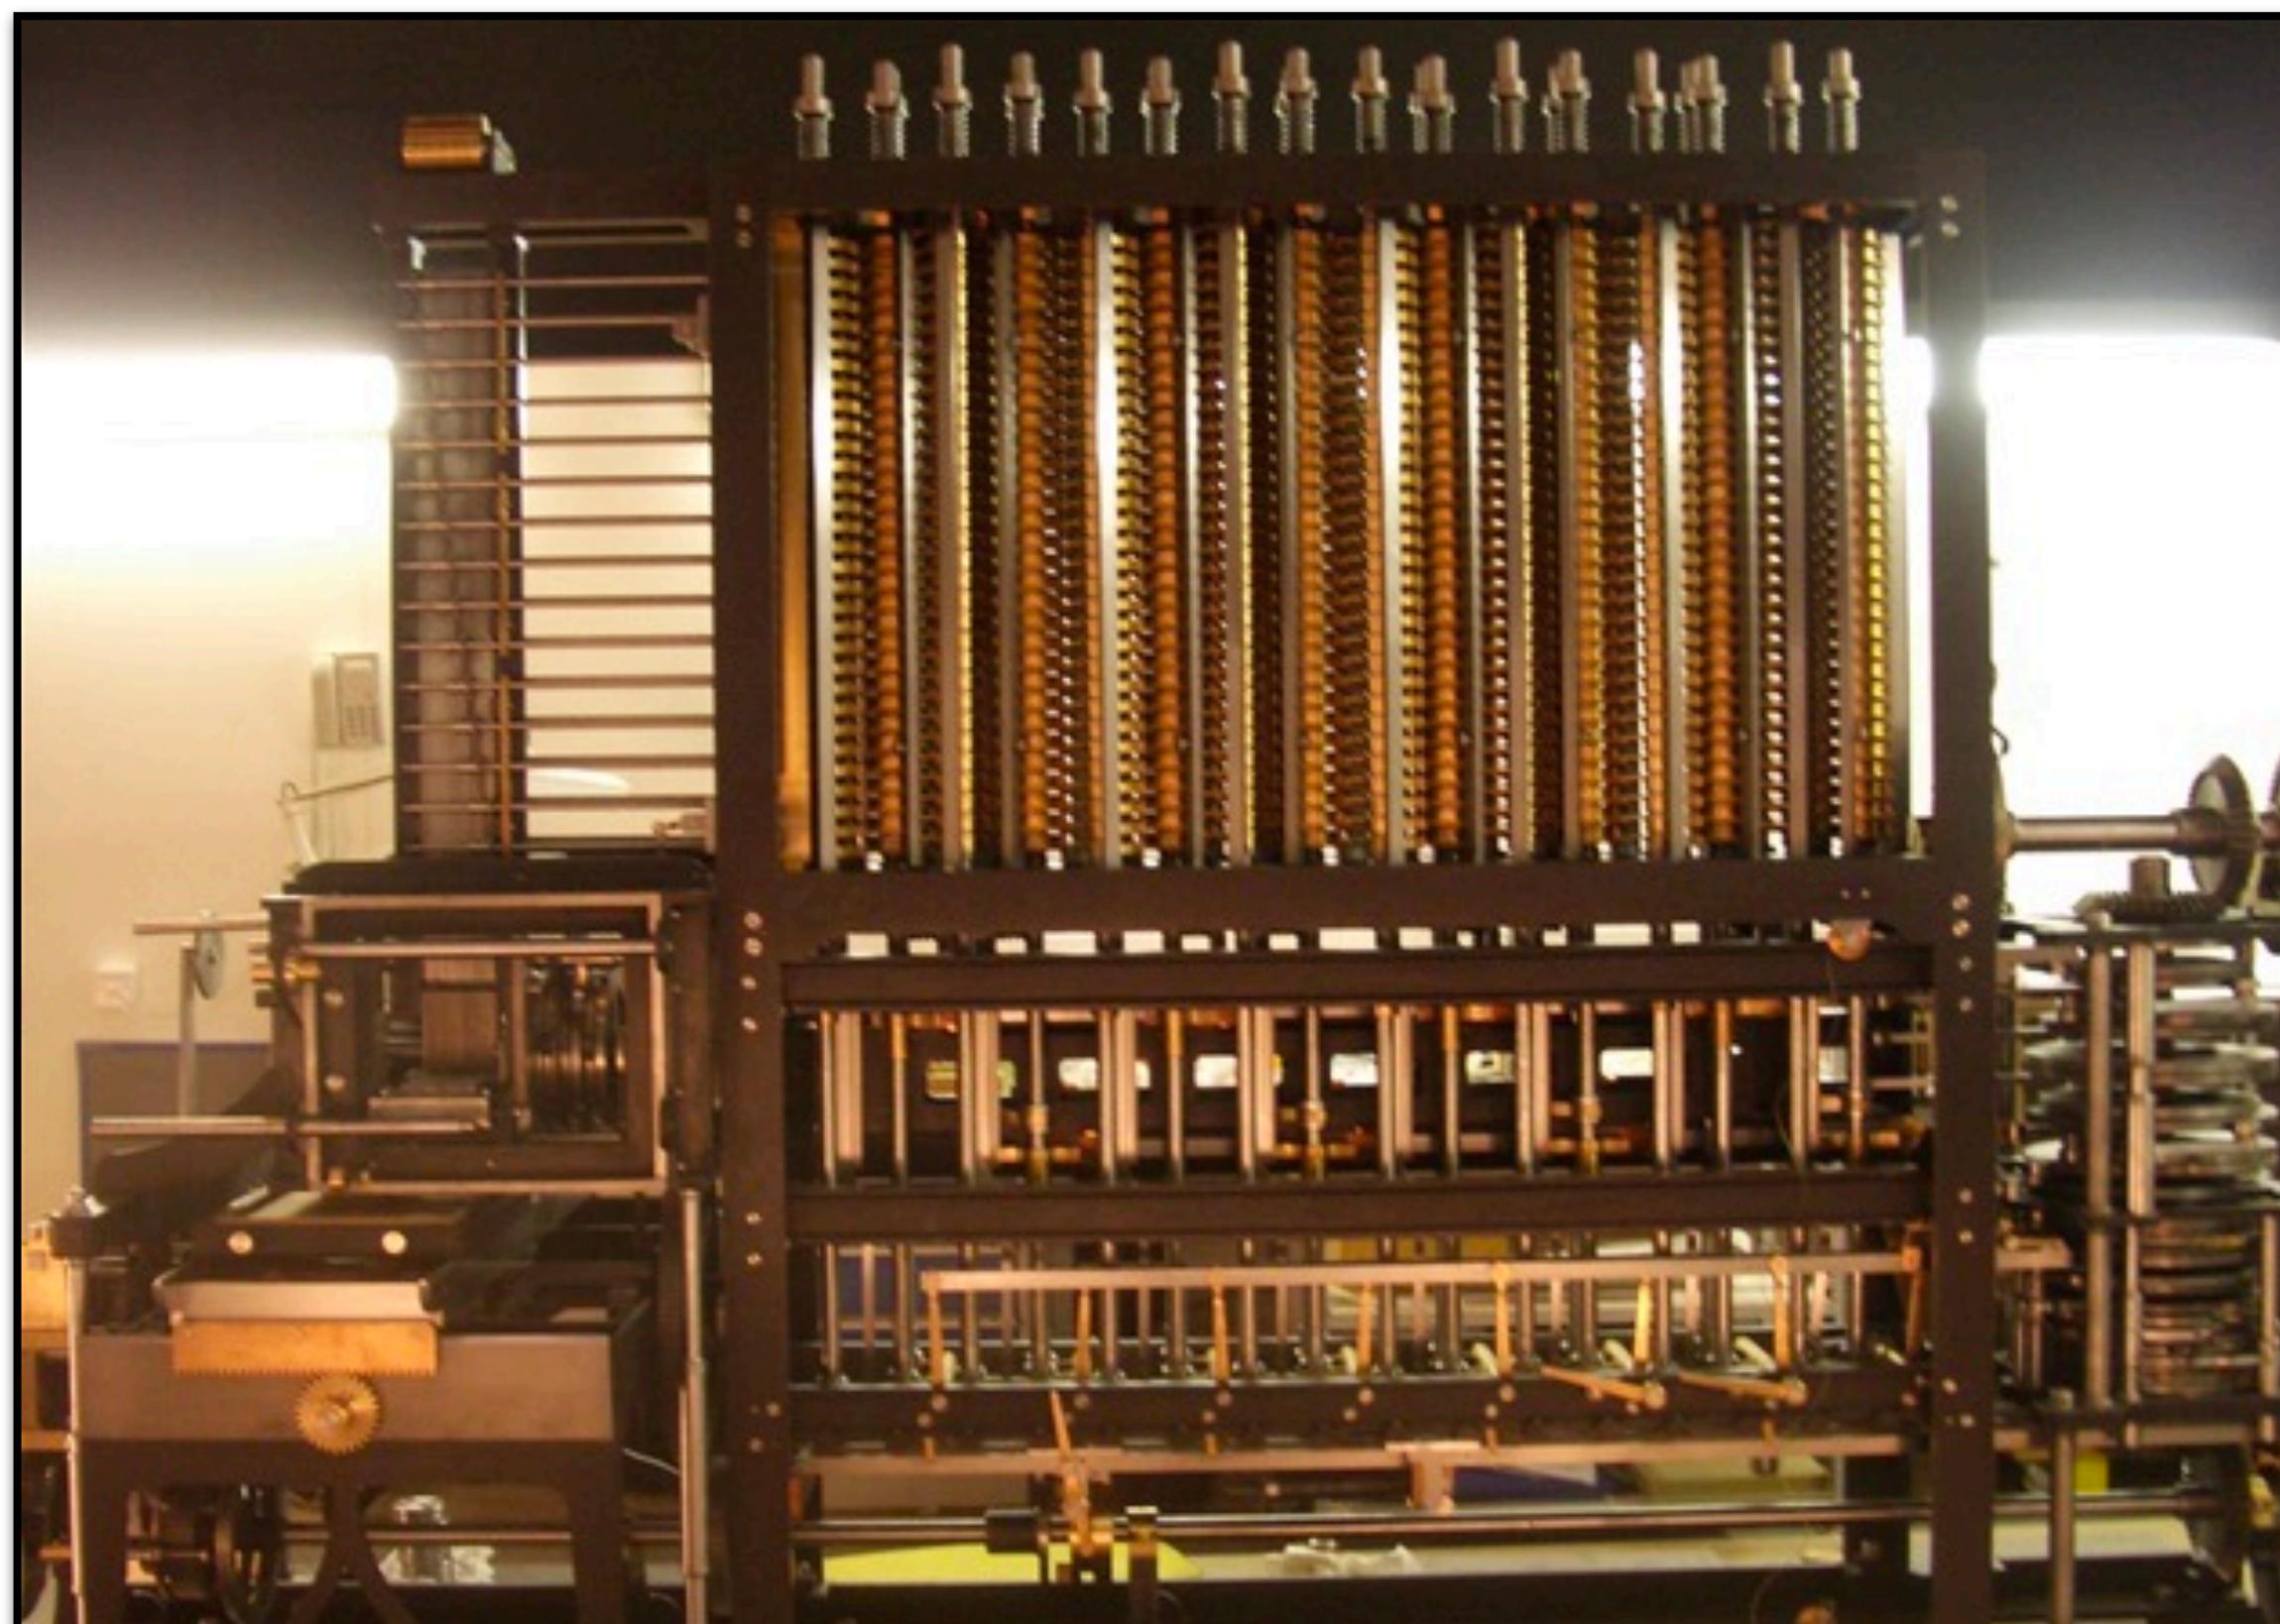
$$M[i] = \begin{pmatrix} p[i] & g[i] \\ 0 & 1 \end{pmatrix}$$

$$\begin{pmatrix} c[i] \\ 1 \end{pmatrix} = M[i] \times M[i-1] \times \dots \times M[0] \times \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

**Carry-lookahead addition** is used in all computers

# This idea is used in all hardware

The design goes back to Babbage in the 1800s:



# Application: Lexical Analysis

Lexical analysis **divides a long string of characters into tokens** - often the first thing a compiler does when processing a program.

Suppose we have a regular language - we can represent it with a **finite-state automaton** that begins in a certain state and makes transitions between states based on the characters read.

```
if x <= n then print ( "x = " , x ) ;
```

TABLE I. A Finite-State Automaton for Recognizing Tokens

Old State	Character Read													New line
	•	A	B	...	Y	Z	+	-	*	<	>	=	"	Space
N	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
A	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
Z	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
*	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
<	A	A	...	A	A	*	*	*	<	<	=	Q	N	N
=	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
Q	S	S	...	S	S	S	S	S	S	S	S	E	S	S
S	S	S	...	S	S	S	S	S	S	S	S	E	S	S
E	E	E	...	E	E	*	*	*	<	<	*	S	N	N

Seems to depend on the previous state, which depends on characters read up to some point

Goal: Perform lexical analysis in parallel

# Application: Lexical Analysis

Idea: replace every character in the string with the array representation of its **state-to-state function** (column).

Then perform a parallel-prefix operation with  $\oplus$  as the **array composition**. Each character becomes an array representing the **state-to-state function** for that prefix.

Use the **initial state** to index into the arrays.

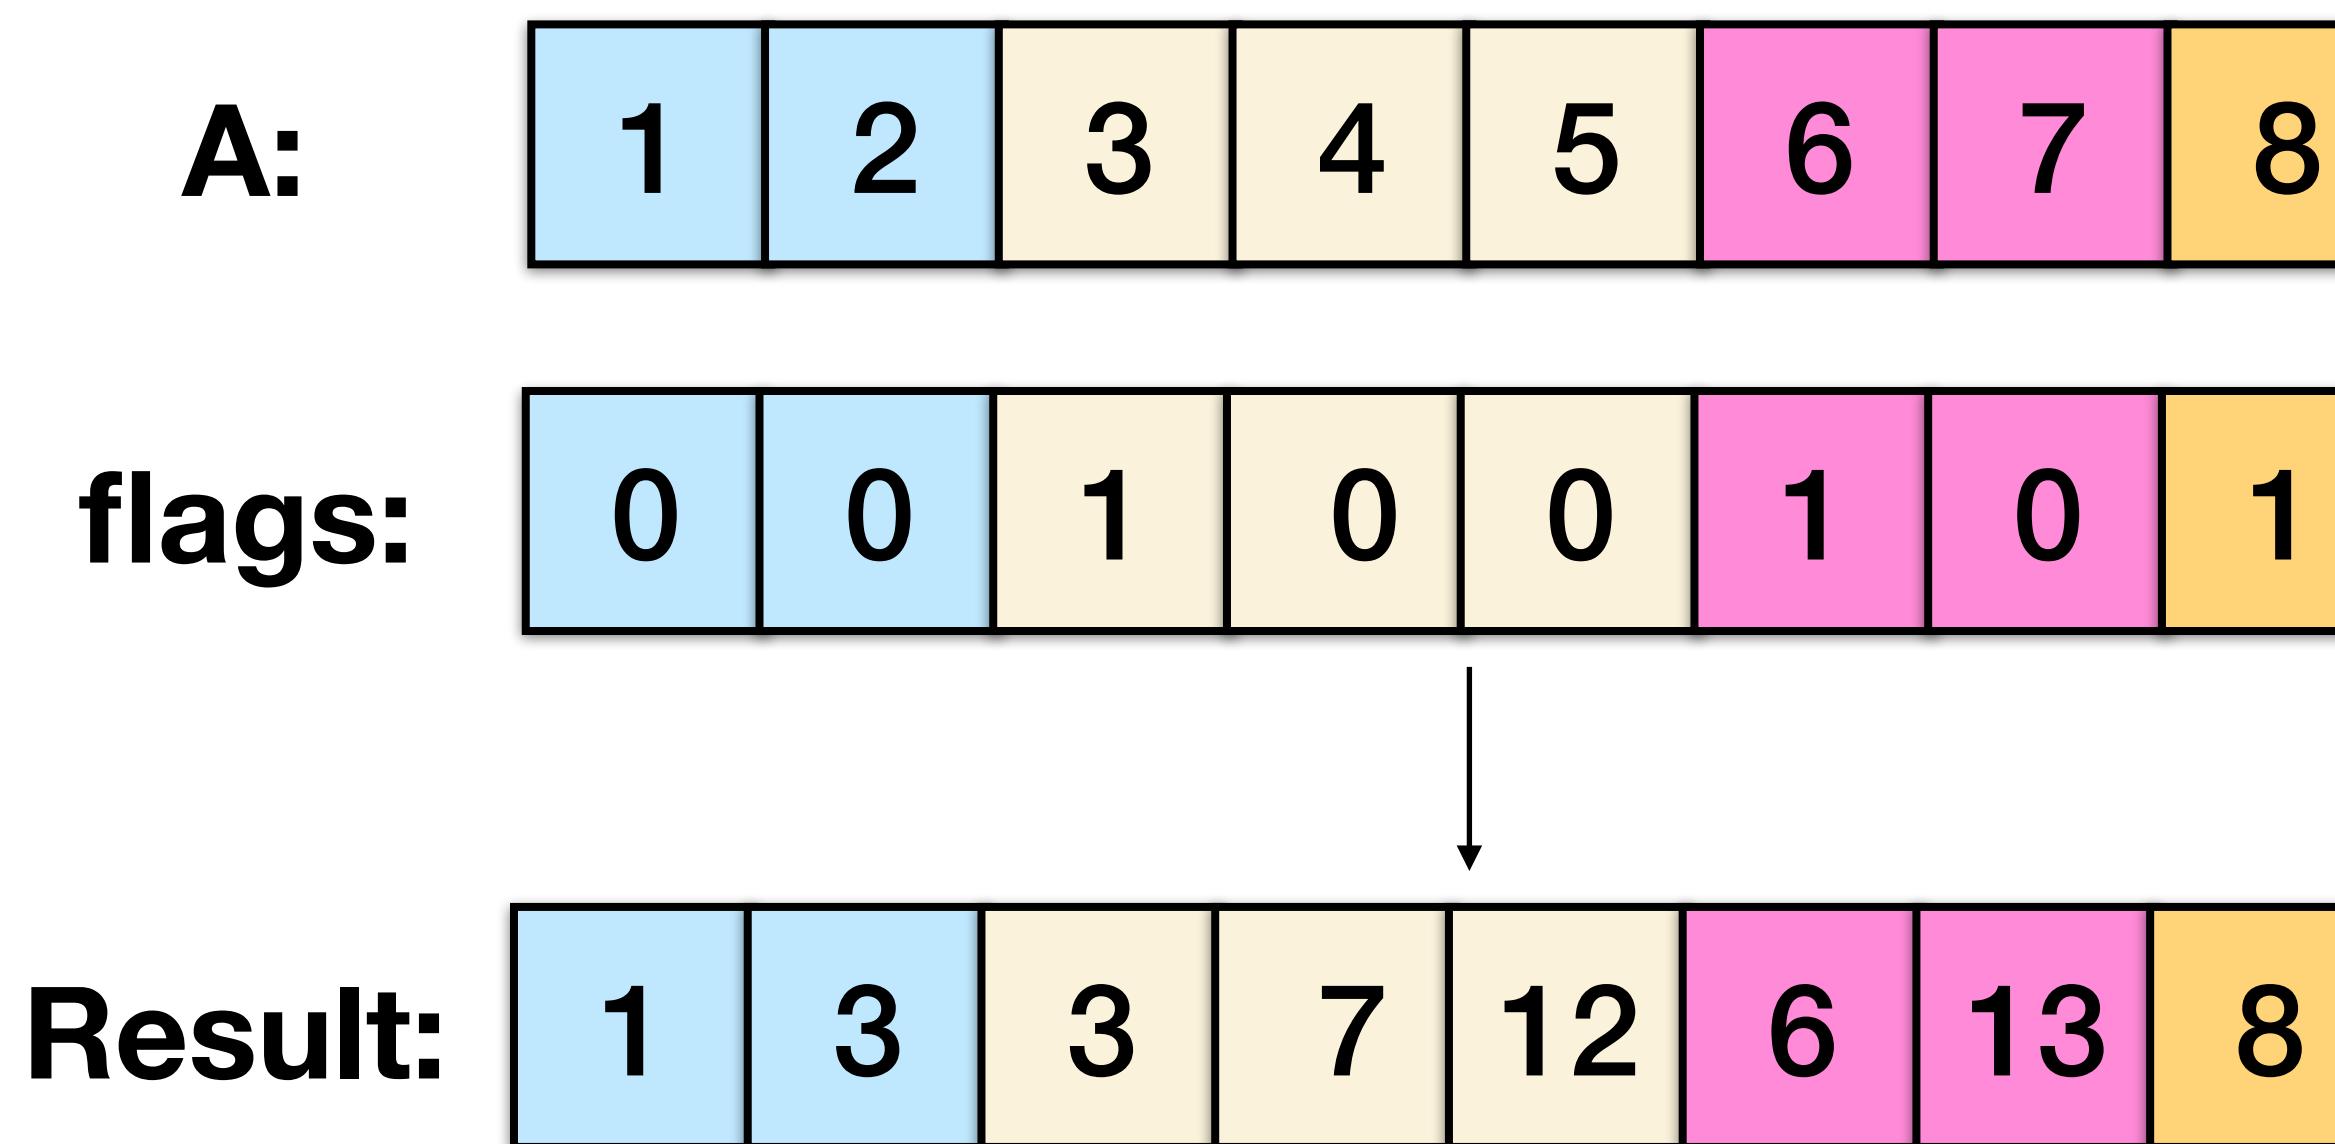
TABLE I. A Finite-State Automaton for Recognizing Tokens

Old State	Character Read													Space	New line
	•	A	B	...	Y	Z	+	-	*	<	>	=	"		
N	A	A	...	A	A	*	*	*	<	<	*	Q	N	N	
A	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N	
Z	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N	
*	A	A	...	A	A	*	*	*	<	<	*	Q	N	N	
<	A	A	...	A	A	*	*	*	<	<	=	Q	N	N	
=	A	A	...	A	A	*	*	*	<	<	*	Q	N	N	
Q	S	S	...	S	S	S	S	S	S	S	E	S	S	S	
S	S	S	...	S	S	S	S	S	S	S	E	S	S	S	
E	E	E	...	E	E	*	*	*	<	<	*	S	N	N	

Example of state-to-state function for the space character

# Application: Segmented Scans

Inputs: value array, flag array, associative operator  $\oplus$



Can be used to parallelize sparse-matrix vector multiply (SpMV)

# Application: Image Processing

- A **summed-area table** is an algorithm/data structure for quickly generating the sum of values in some rectangular subset of a grid.
- Often used in image processing [Crow, 84].
- Computes **prefix sums in both dimensions**, and then **inclusion-exclusion on the corners** to compute the sum within any rectangular area.

$$I(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y')$$

$$A = (x_0, y_0), B = (x_1, y_0), C = (x_0, y_1), D = (x_1, y_1)$$

$$\sum_{x_0 < x \leq x_1, y_0 < y \leq y_1} i(x, y) = I(D) + I(A) - I(B) - I(C)$$

1.

31	2	4	33	5	36
12	26	9	10	29	25
13	17	21	22	20	18
24	23	15	16	14	19
30	8	28	27	11	7
1	35	34	3	32	6

2.

31	33	37	70	75	111
43	71	84	127	161	222
56	101	135	200	254	333
80	148	197	278	346	444
110	186	263	371	450	555
111	222	333	444	555	666

$$15 + 16 + 14 + 28 + 27 + 11 = \\ 101 + 450 - 254 - 186 = 111$$

[https://en.wikipedia.org/wiki/Summed-area\\_table](https://en.wikipedia.org/wiki/Summed-area_table)

# Application: Image Processing

- A **summed-area table** is an algorithm/data structure for quickly generating the sum of values in some rectangular subset of a grid.
- Often used in image processing [Crow, 84].
- Computes **prefix sums in both dimensions**, and then **inclusion-exclusion on the corners** to compute the sum within any rectangular area.

$$I(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y')$$

$$A = (x_0, y_0), B = (x_1, y_0), C = (x_0, y_1), D = (x_1, y_1)$$

$$\sum_{x_0 < x \leq x_1, y_0 < y \leq y_1} i(x, y) = I(D) + I(A) - I(B) - I(C)$$

Requires inverse

1.

31	2	4	33	5	36
12	26	9	10	29	25
13	17	21	22	20	18
24	23	15	16	14	19
30	8	28	27	11	7
1	35	34	3	32	6

2.

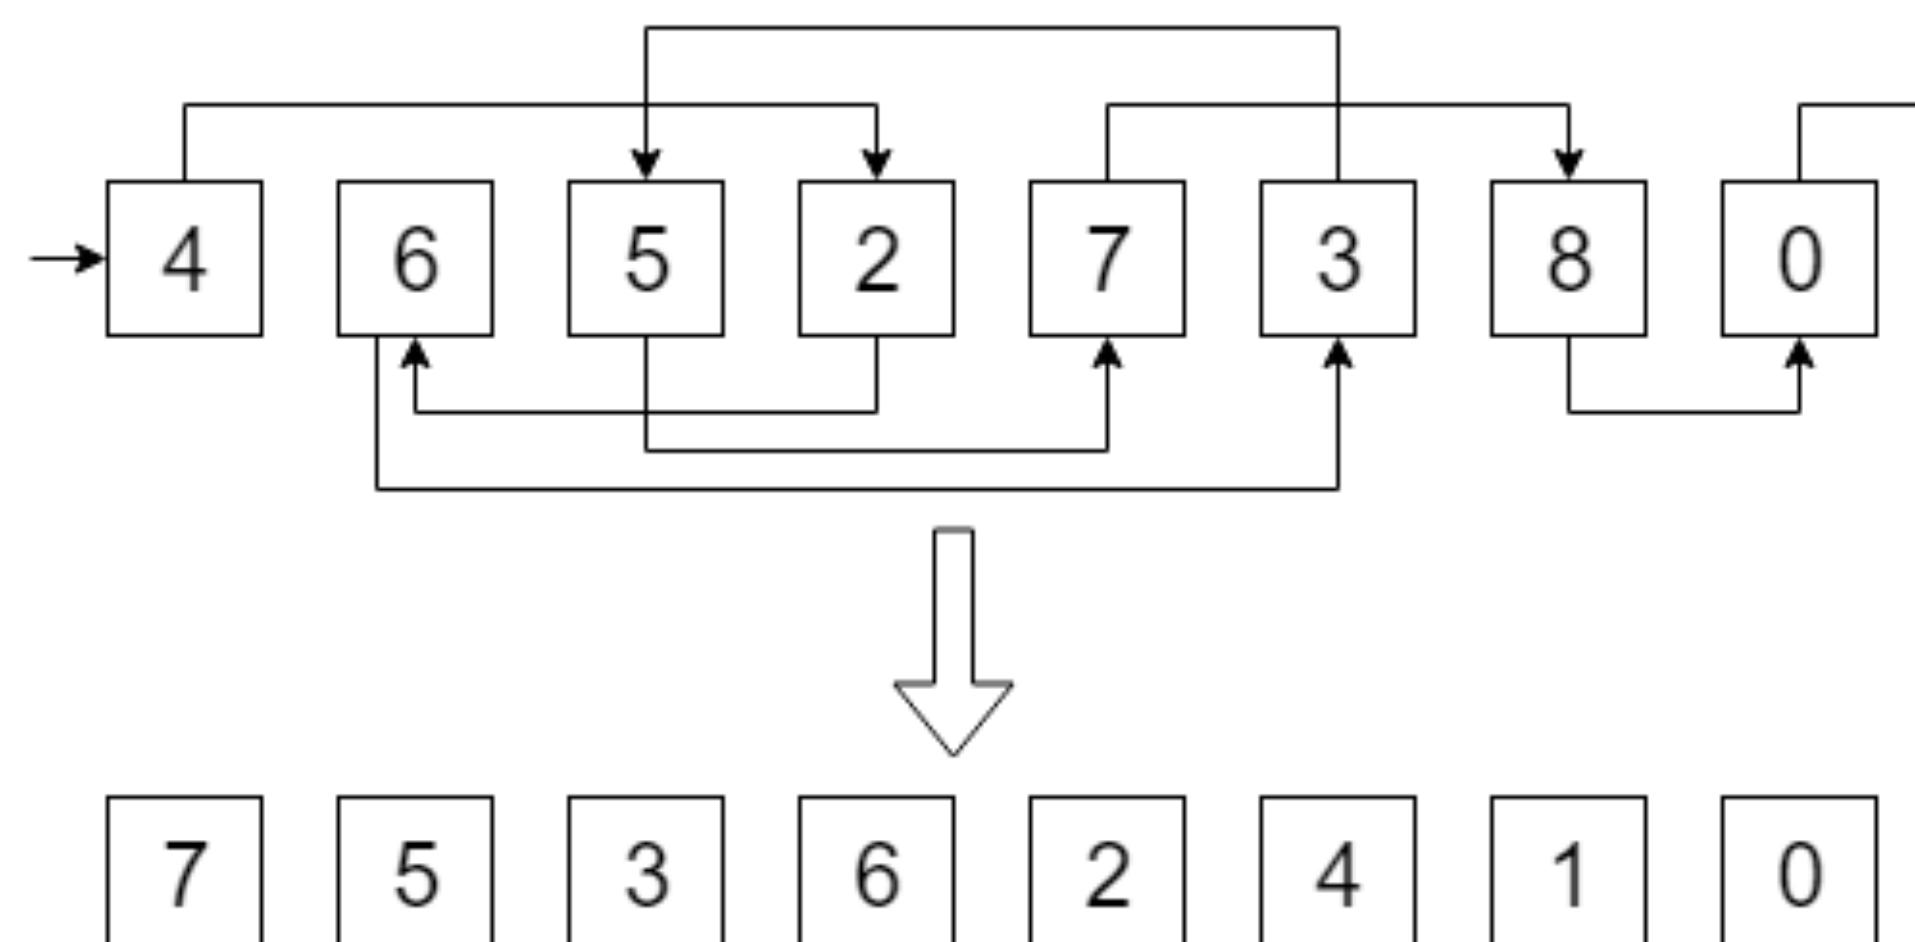
31	33	37	70	75	111
43	71	84	127	161	222
56	101	135	200	254	333
80	148	197	278	346	444
110	186	263	371	450	555
111	222	333	444	555	666

$$15 + 16 + 14 + 28 + 27 + 11 = \\ 101 + 450 - 254 - 186 = 111$$

[https://en.wikipedia.org/wiki/Summed-area\\_table](https://en.wikipedia.org/wiki/Summed-area_table)

# Application: List Ranking

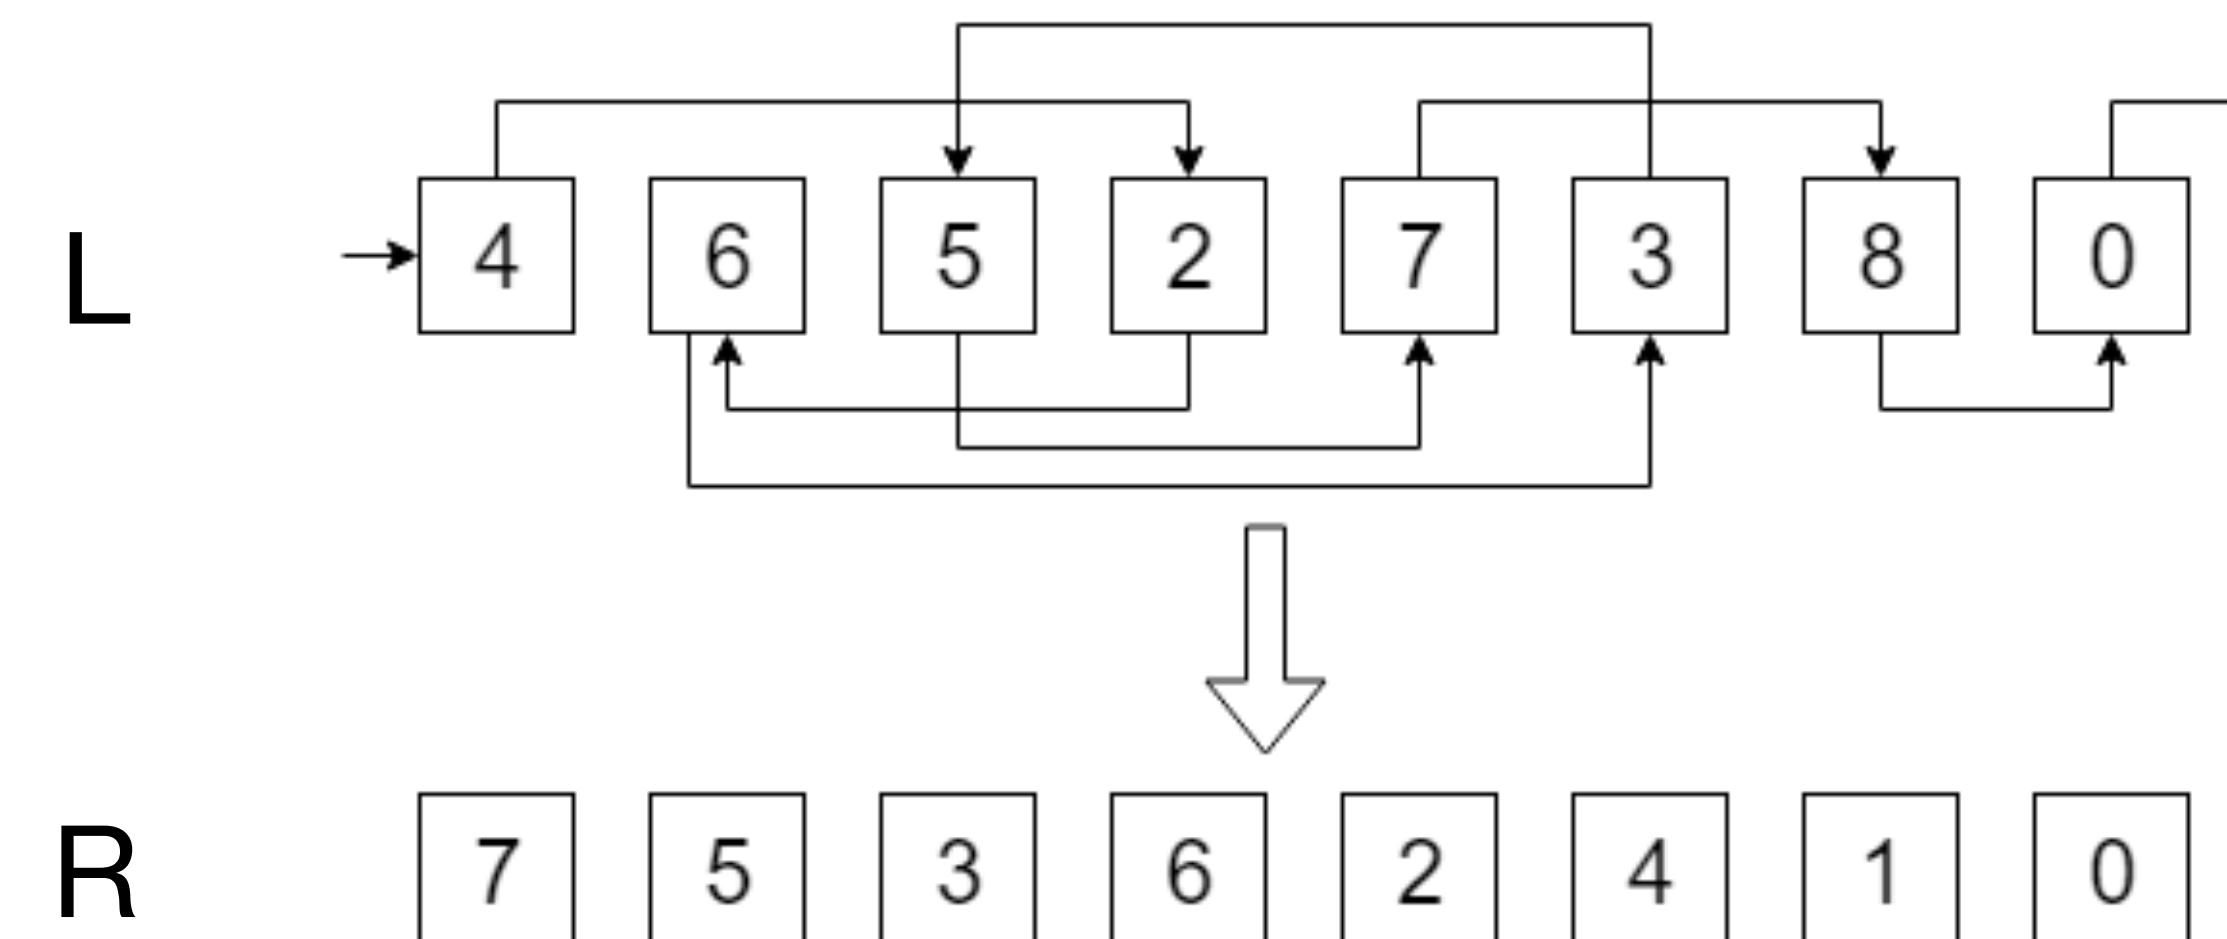
- The list ranking problem is to determine the **distance from every node in a singly-linked list to the end**.
- In serial, the problem can be solved by traversing the list in  $O(n)$  time.
- List ranking underlies many applications like prefix sum over linked lists, and tree algorithms.



# Application: List Ranking

Input: A linked list  $L$  of  $n$  nodes with an array  $S$  specifying their order. That is, every **entry in the list  $L$  is a pointer to the index of its successor** (or 0, if it is the end node).

Output: A list  $R$  of length  $n$  containing the **rank** (distance from the end of the list) of each element.



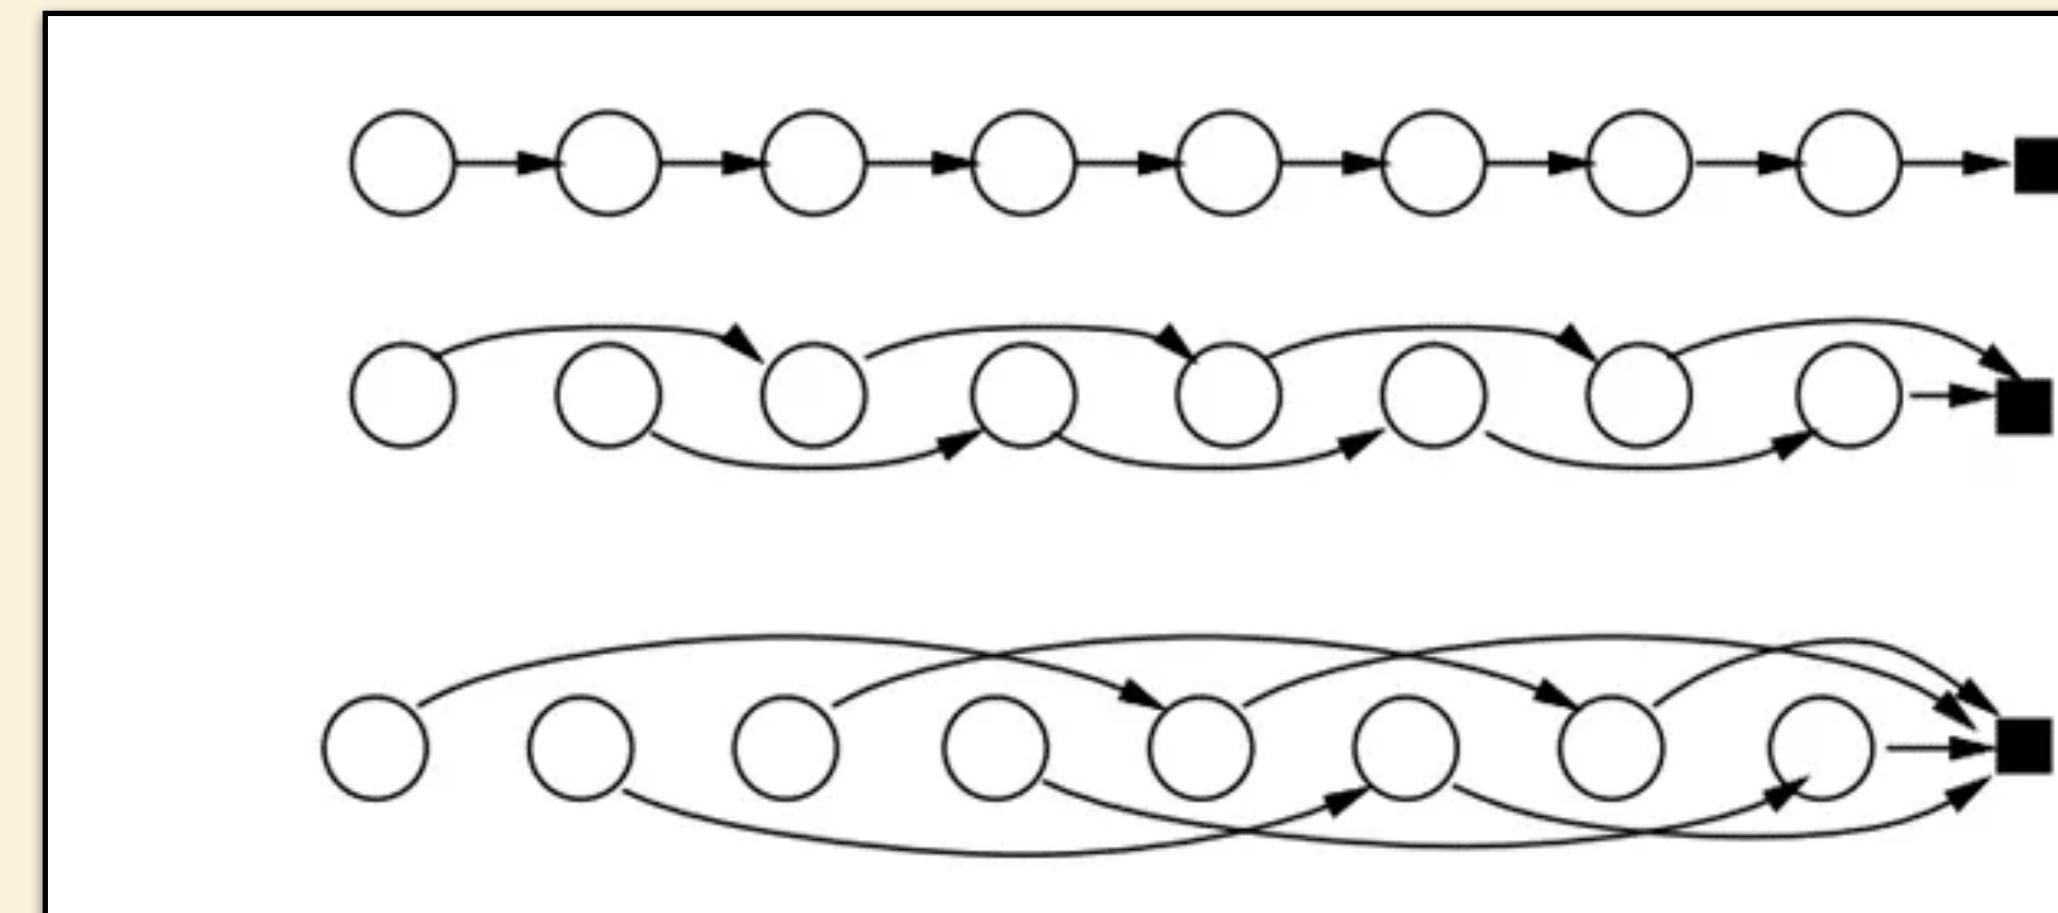
# Wyllie's Algorithm for List Ranking

$L$  = list of input successor pointers in an array  
 $R$  = output list of ranks (distance from end)

```

parallel_for i = 0 to n-1:
    if L[i] != 0: R[i] = 1
    else: R[i] = 0

for j = 0 to ceil(lg n) - 1:
    temp, temp2;
    parallel_for i = 0 to n-1:
        temp[i] = R[L[i]];
        temp2[i] = L[L[i]];
    parallel_for i = 0 to n-1:
        R[i] = R[i] + temp[i];
        L[i] = temp2[i];
    
```



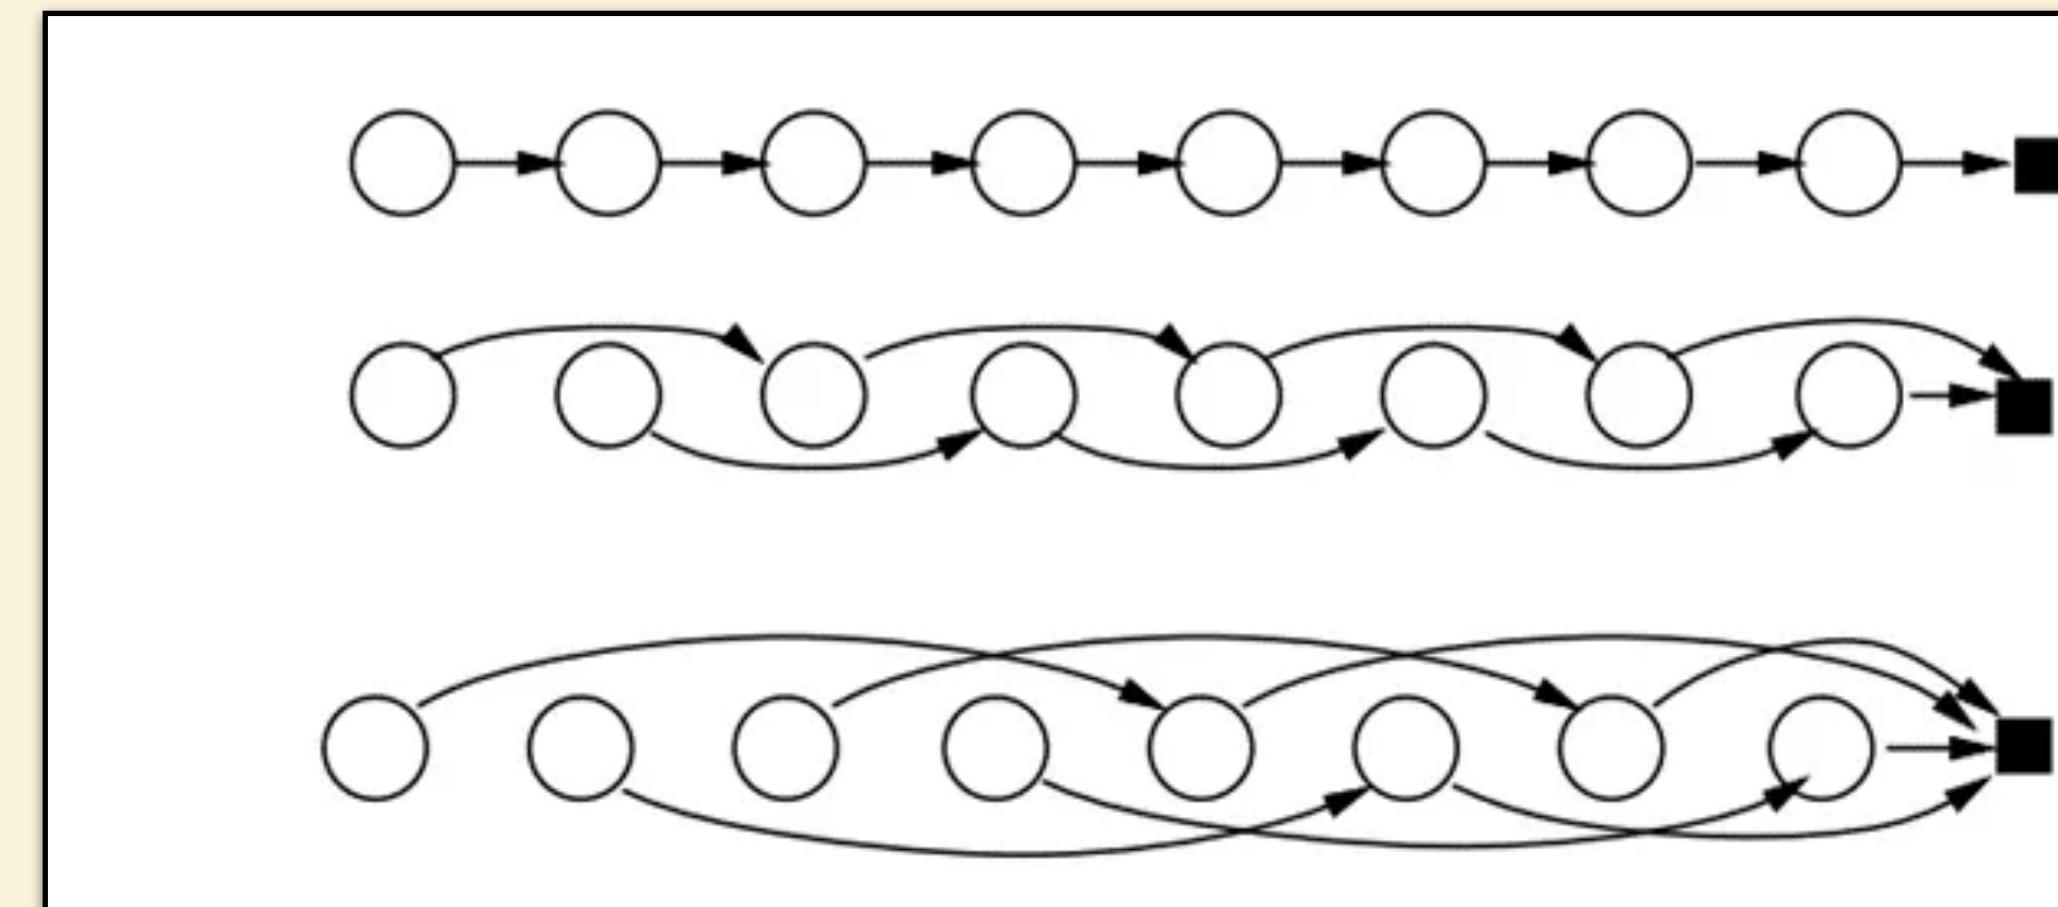
# Wyllie's Algorithm for List Ranking

$L$  = list of input successor pointers in an array  
 $R$  = output list of ranks (distance from end)

```

parallel_for i = 0 to n-1:
    if L[i] != 0: R[i] = 1
    else: R[i] = 0

for j = 0 to ceil(lg n) - 1:
    temp, temp2;
    parallel_for i = 0 to n-1:
        temp[i] = R[L[i]];
        temp2[i] = L[L[i]];
    parallel_for i = 0 to n-1:
        R[i] = R[i] + temp[i];
        L[i] = temp2[i];
    
```



What is the work and the span?

# Work-Span Analysis

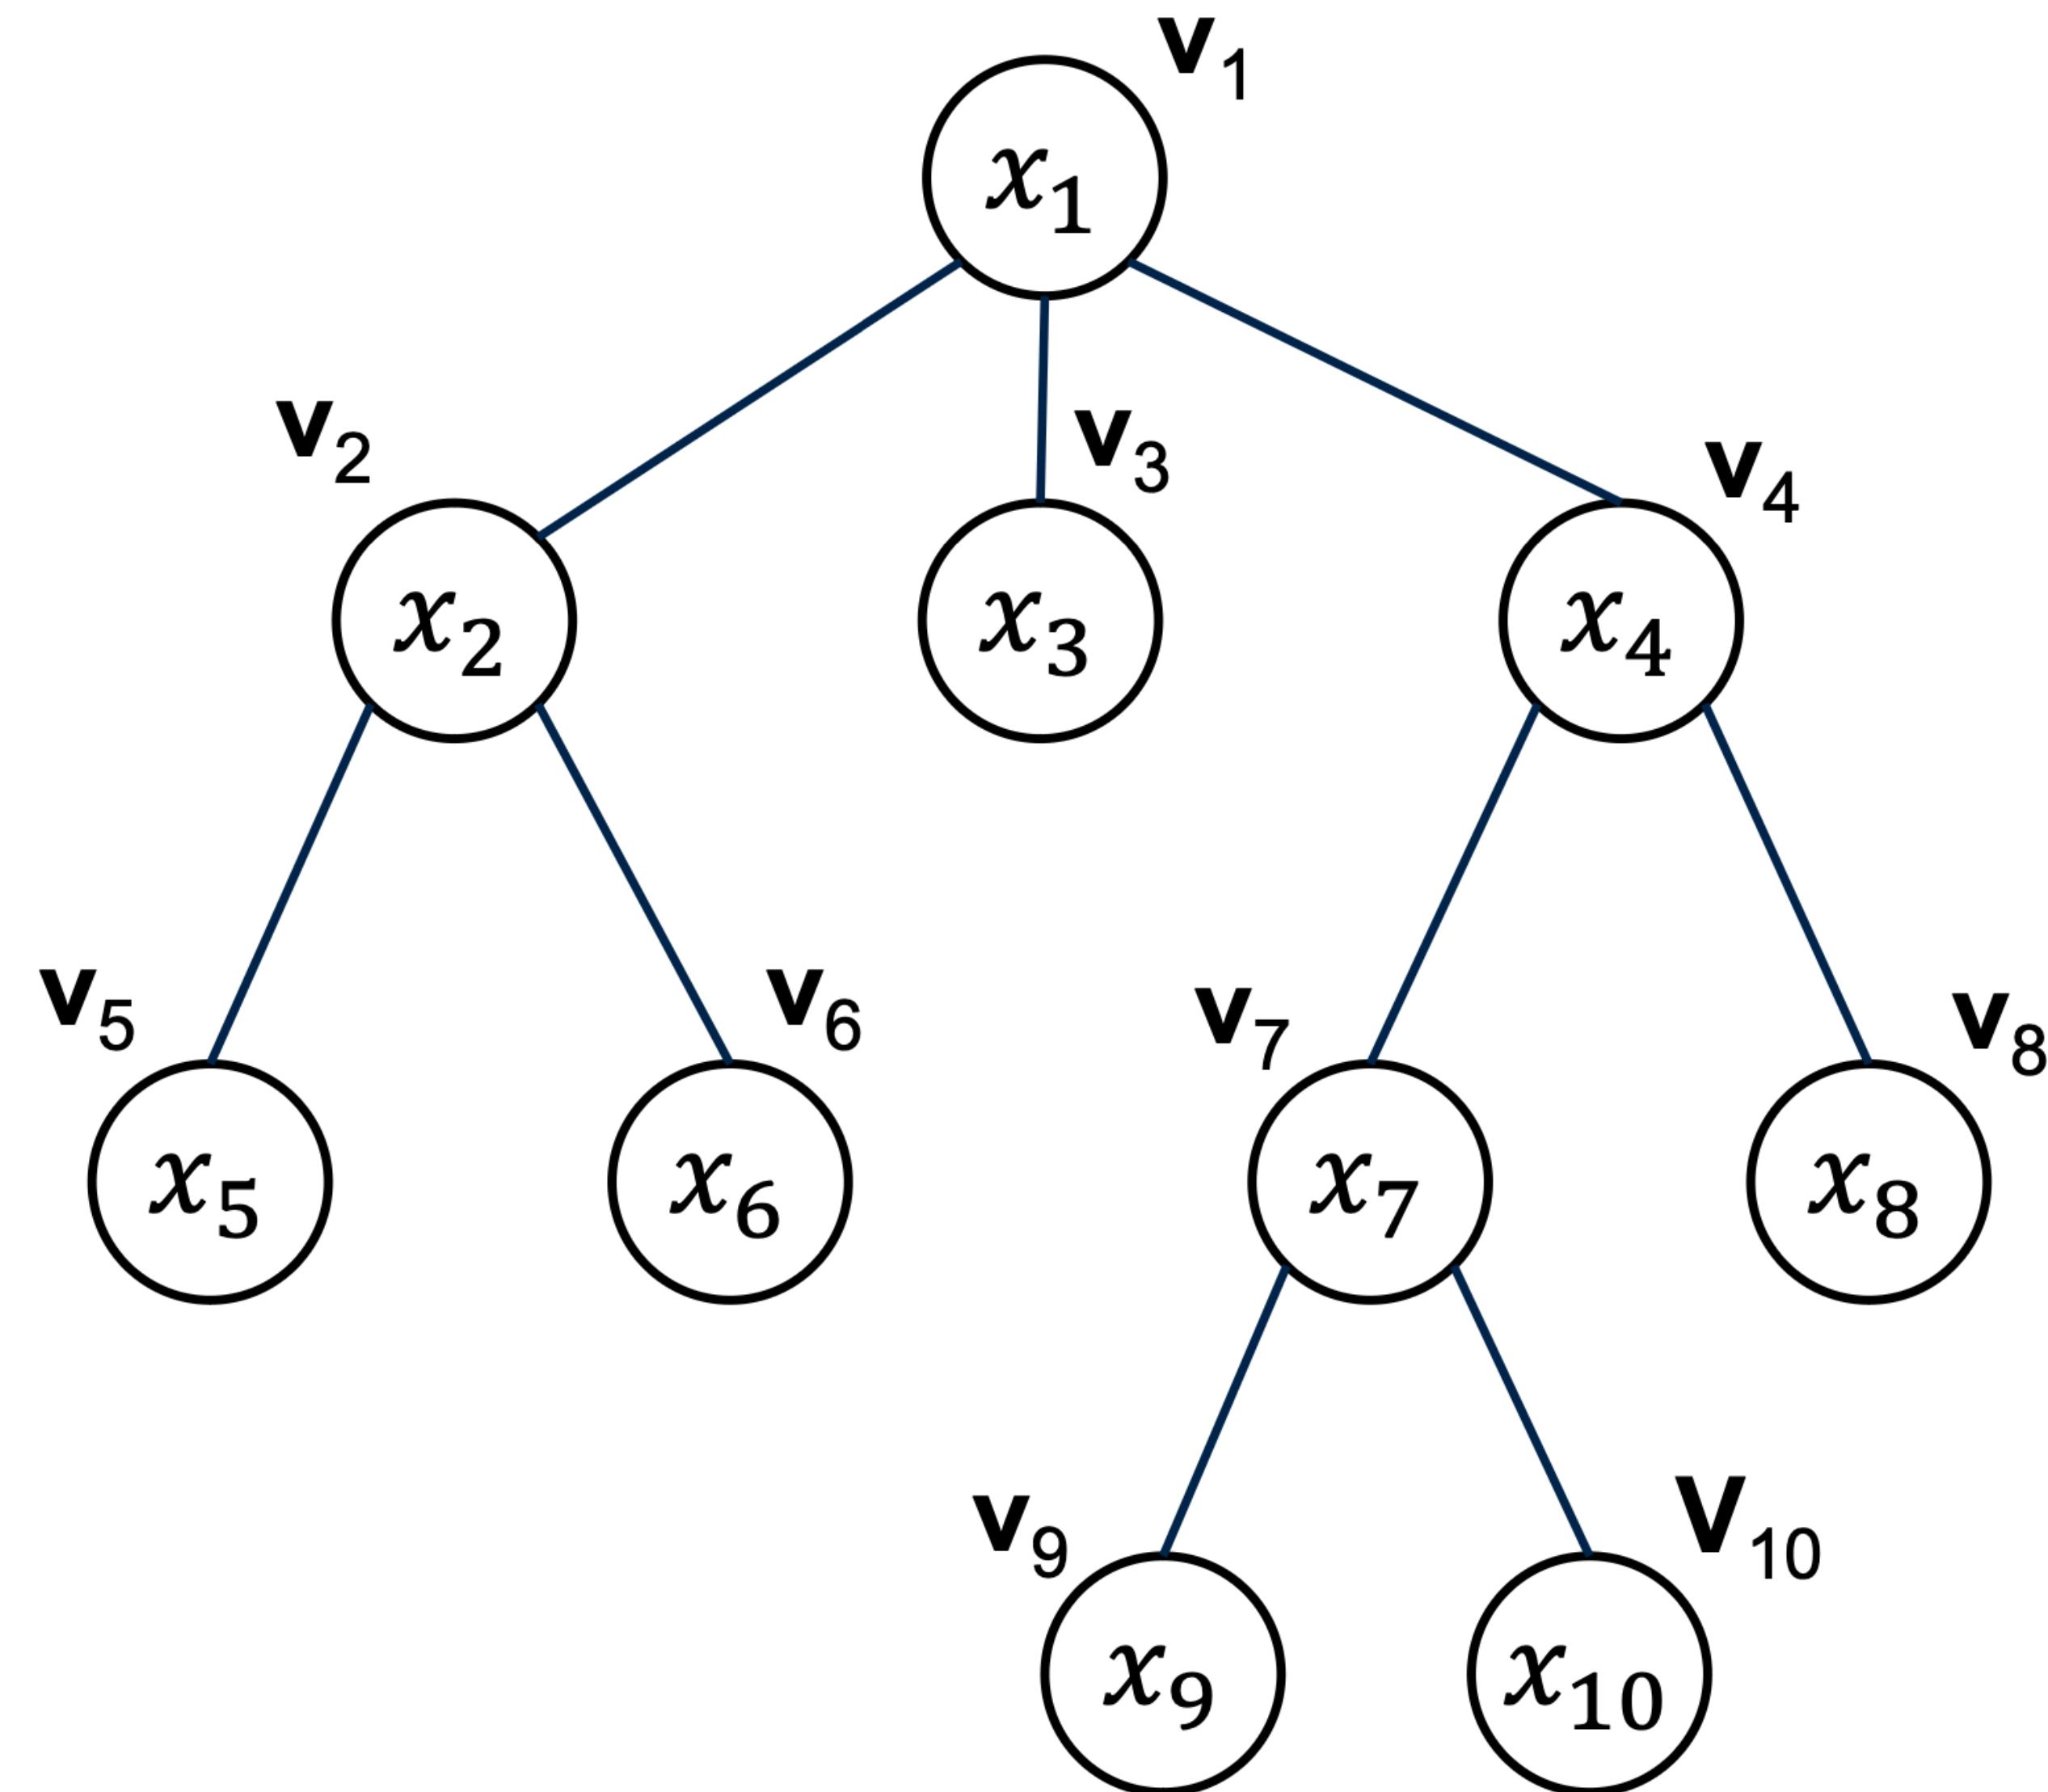
L = list of input successor pointers in an array  
R = output list of ranks (distance from end)

```
parallel_for i = 0 to n-1:  
    if L[i] != 0: R[i] = 1  
    else: R[i] = 0  
  
for j = 0 to ceil(lg n) - 1:  
    temp, temp2;  
    parallel_for i = 0 to n-1:  
        temp[i] = R[L[i]];  
        temp2[i] = L[L[i]];  
    parallel_for i = 0 to n-1:  
        R[i] = R[i] + temp[i];  
        L[i] = temp2[i];
```

Not work efficient, since  
serial algorithm is  $O(n)$

Work =  $O(n \log n)$   
Span =  $O(\log^2 n)$

# Application: Tree Accumulation



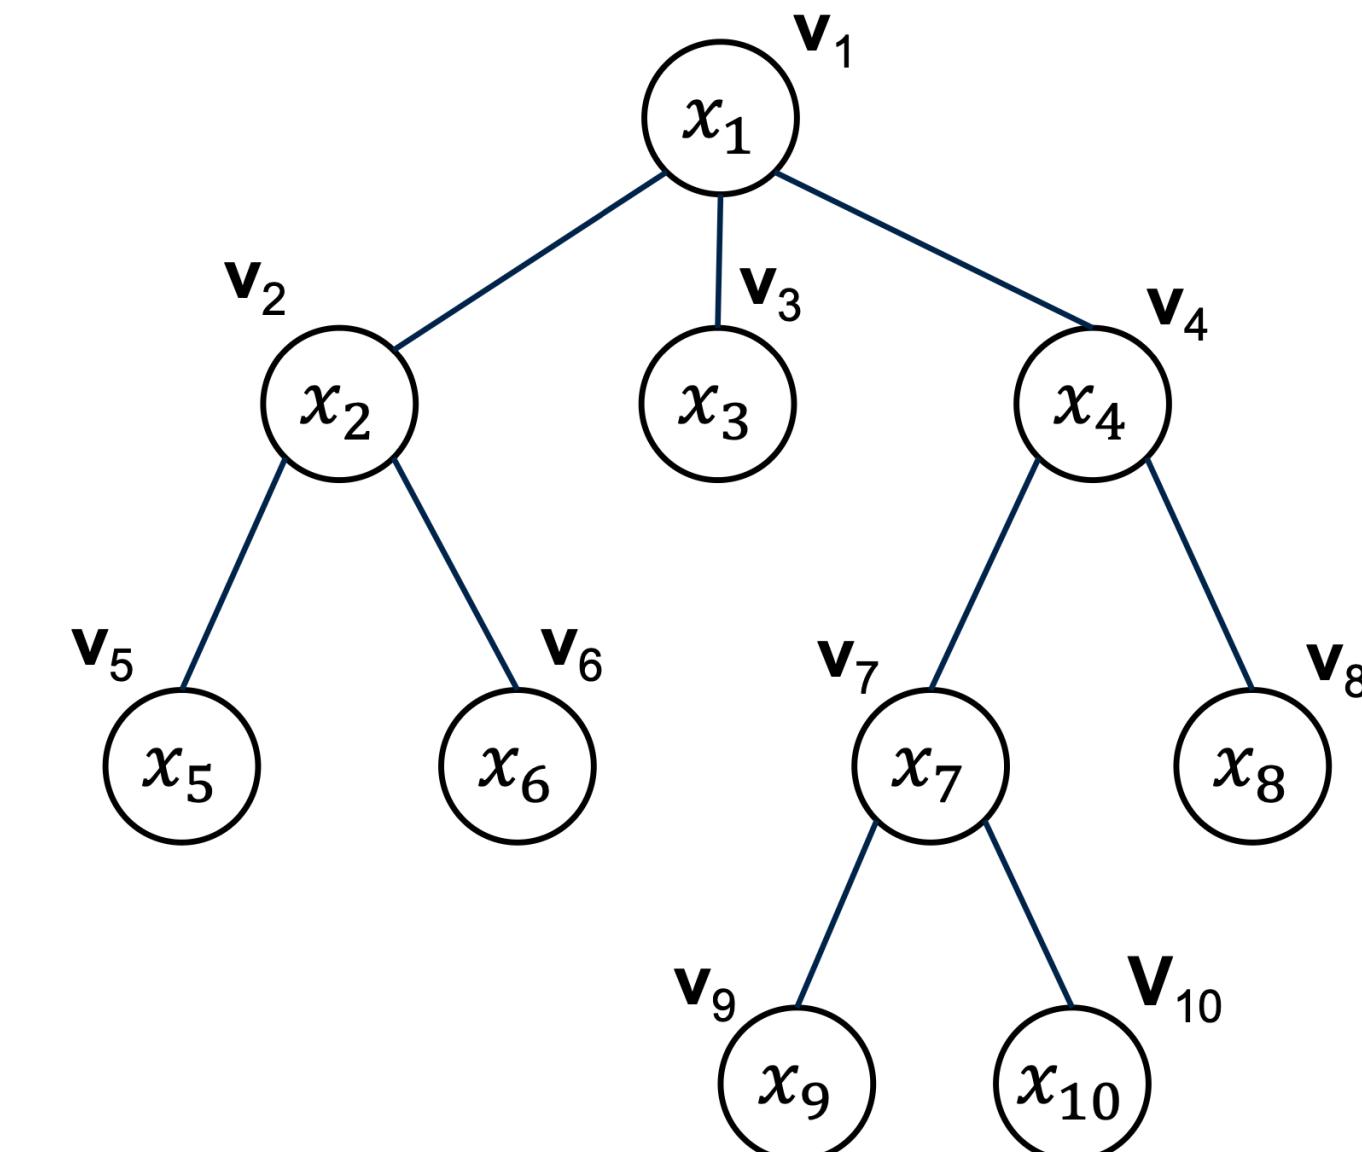
# Application: Tree Accumulation

**Upward Accumulation (UA):** For a node  $v$ , sum all numbers in the subtree

- Example:  $UA(v_4) = x_4 + x_7 + x_8 + x_9 + x_{10}$
- We want to compute it **for all nodes**

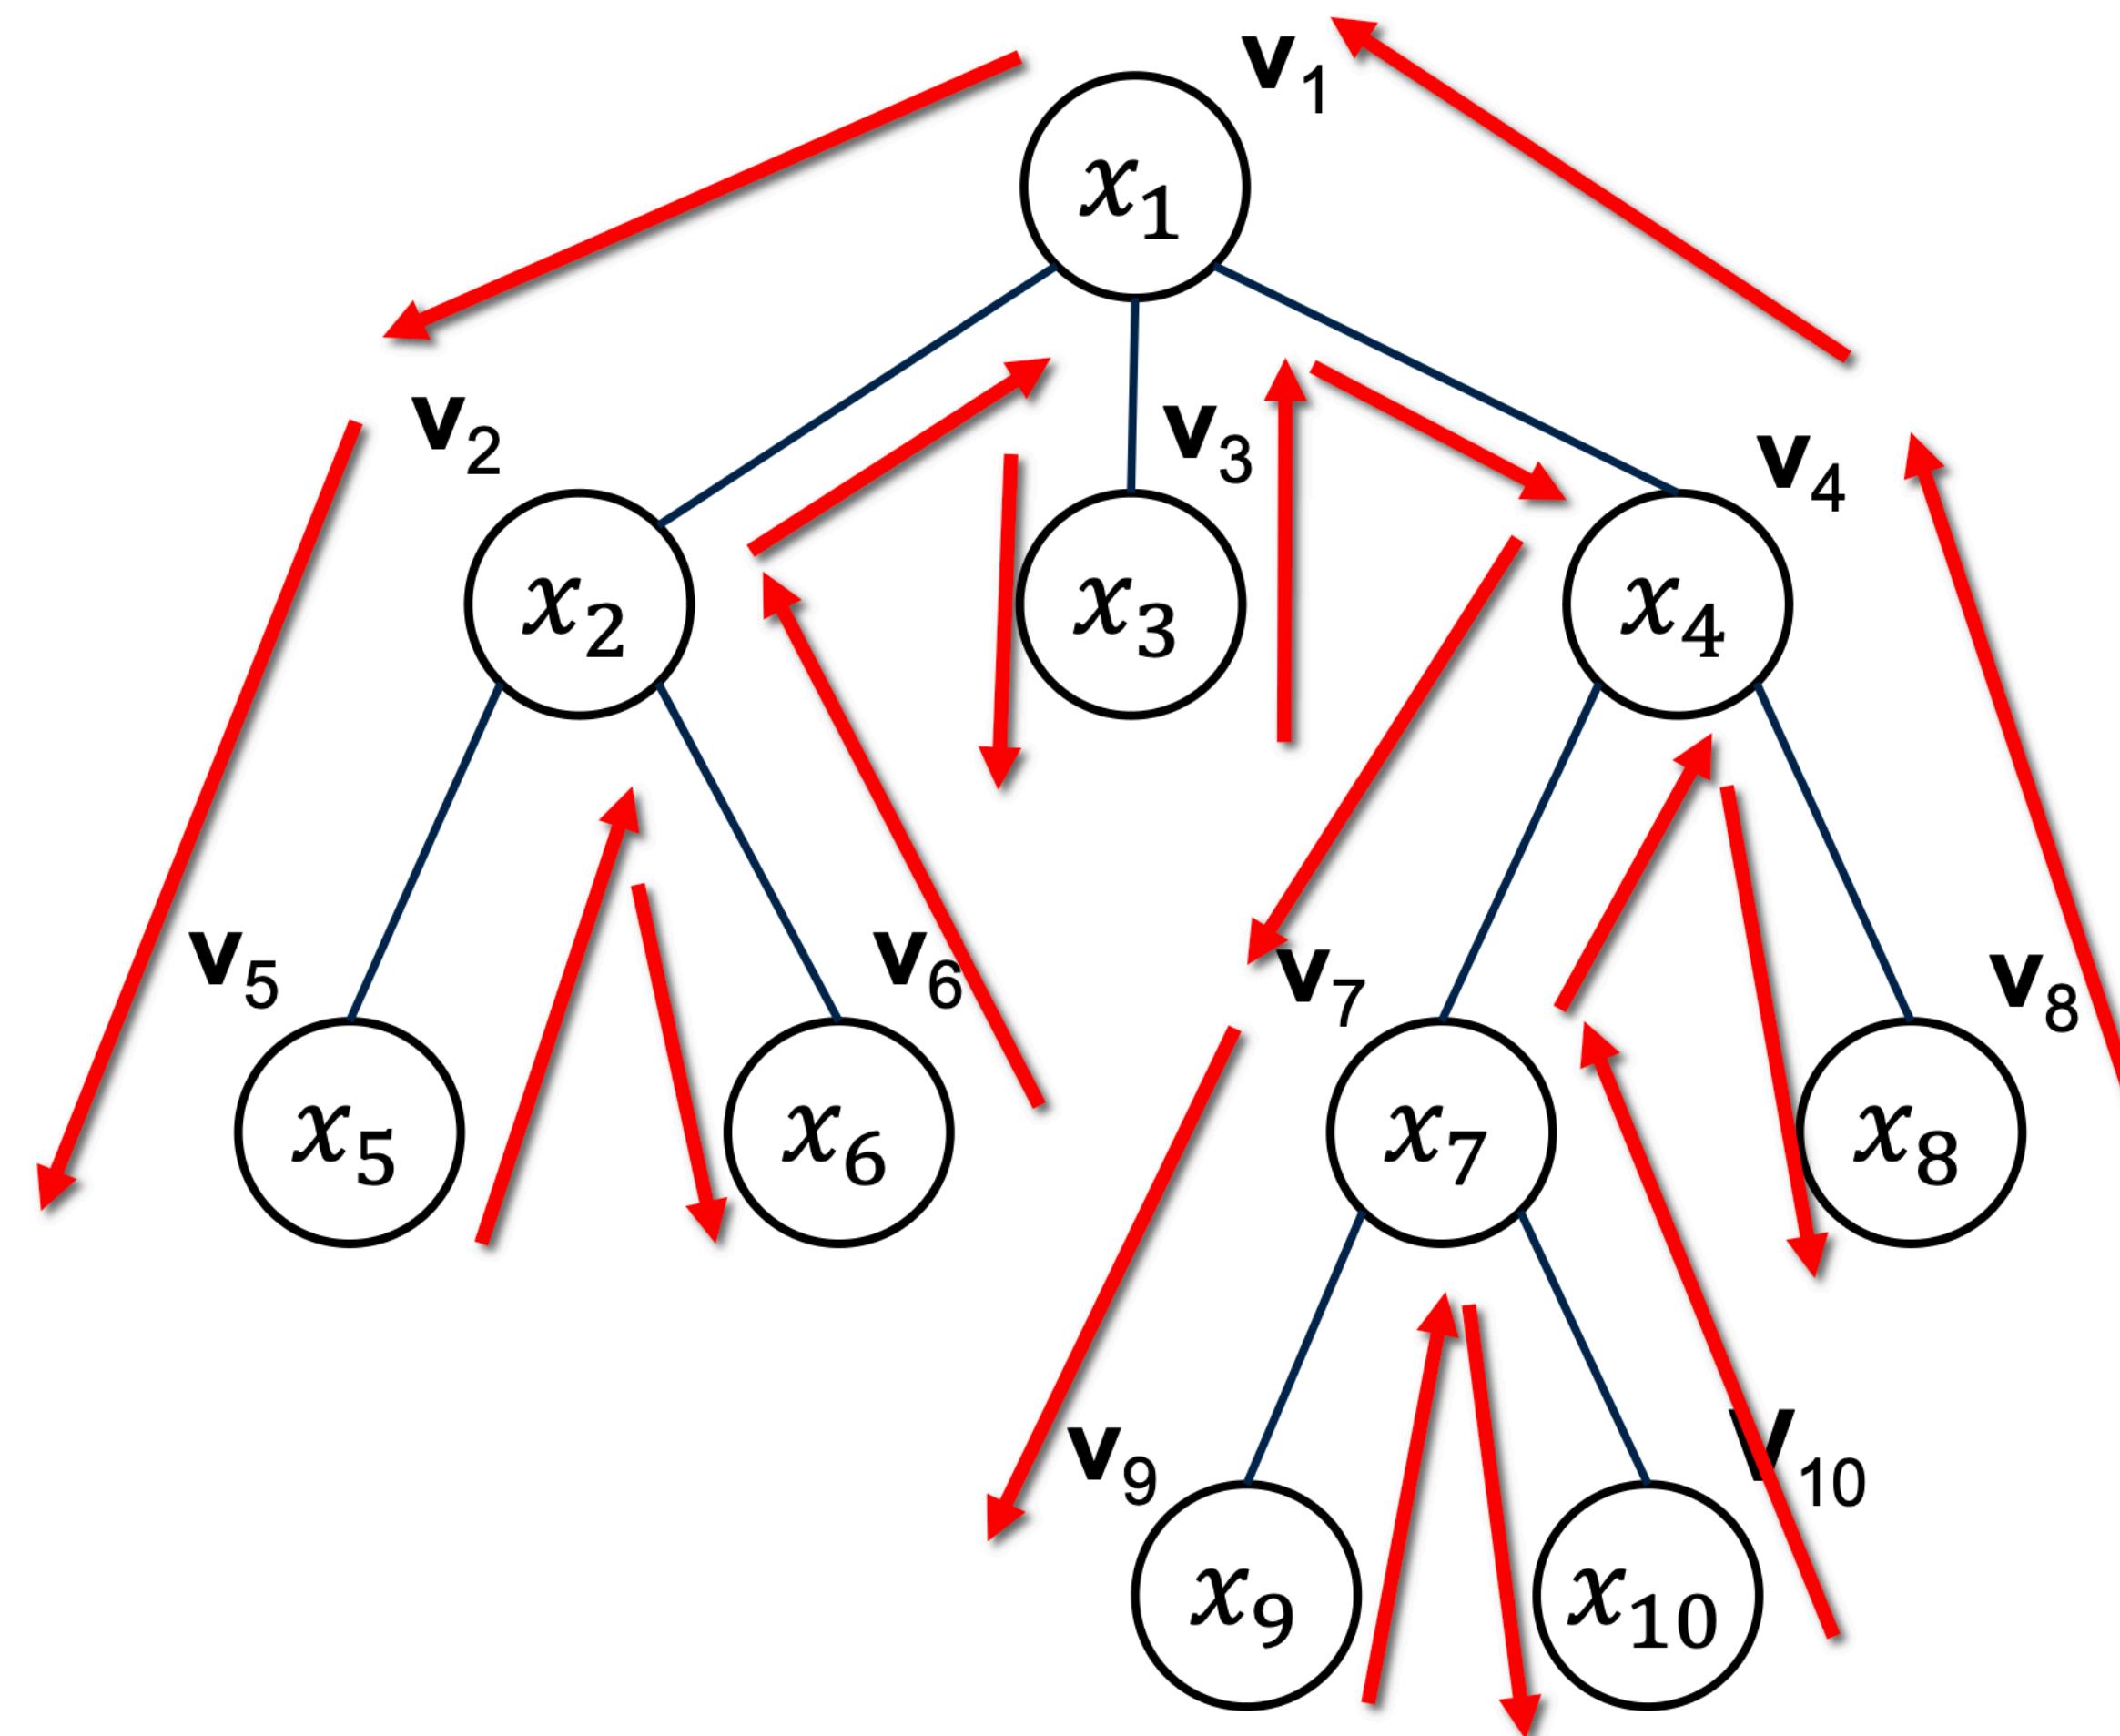
**Downward Accumulation (DA):** For each node  $v$ , sum all numbers in ancestors of  $v$

- Example:  $DA(v_7) = x_7 + x_4 + x_1$



Serial computation time is  $\Theta(n)$  for both UA and DA

# Euler Tour Example

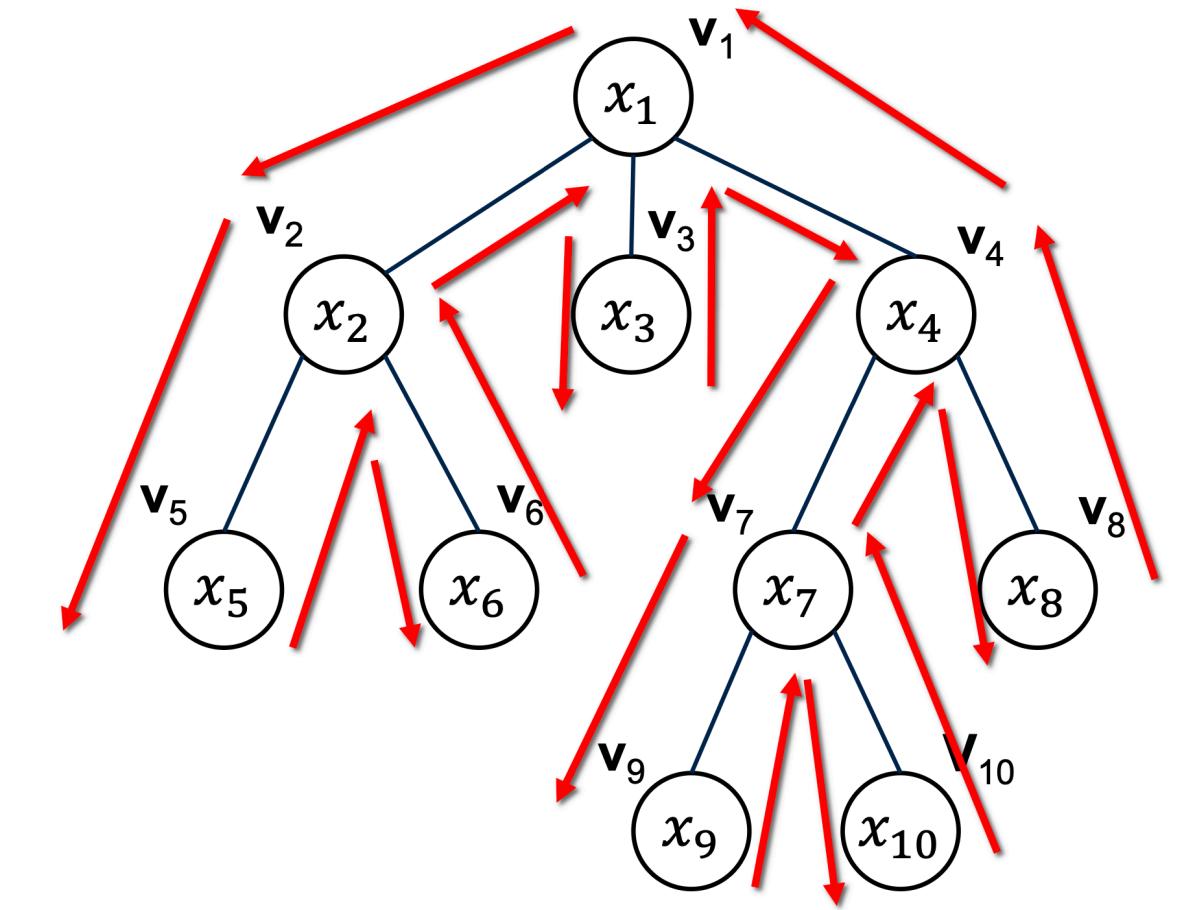


Euler Tour:  $v_1 \ v_2 \ v_5 \ v_2 \ v_6 \ v_2 \ v_1 \ v_3 \ v_1 \ v_4 \ v_7 \ v_9 \ v_7 \ v_{10} \ v_7 \ v_4 \ v_8 \ v_4 \ v_1$

# Tree Accumulation: UA

Sum in subtree

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- $|ET| = 1 + 2(n - 1) = 2n - 1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Otherwise put 0



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_4$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	0	0	$x_3$	0	$x_4$	$x_7$	$x_9$	0	$x_{10}$	0	0	$x_8$	0	0

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $i_l$ : index of last occurrence of  $v$
- $UA(v) = PS[i_l] - PS[i_f - 1]$

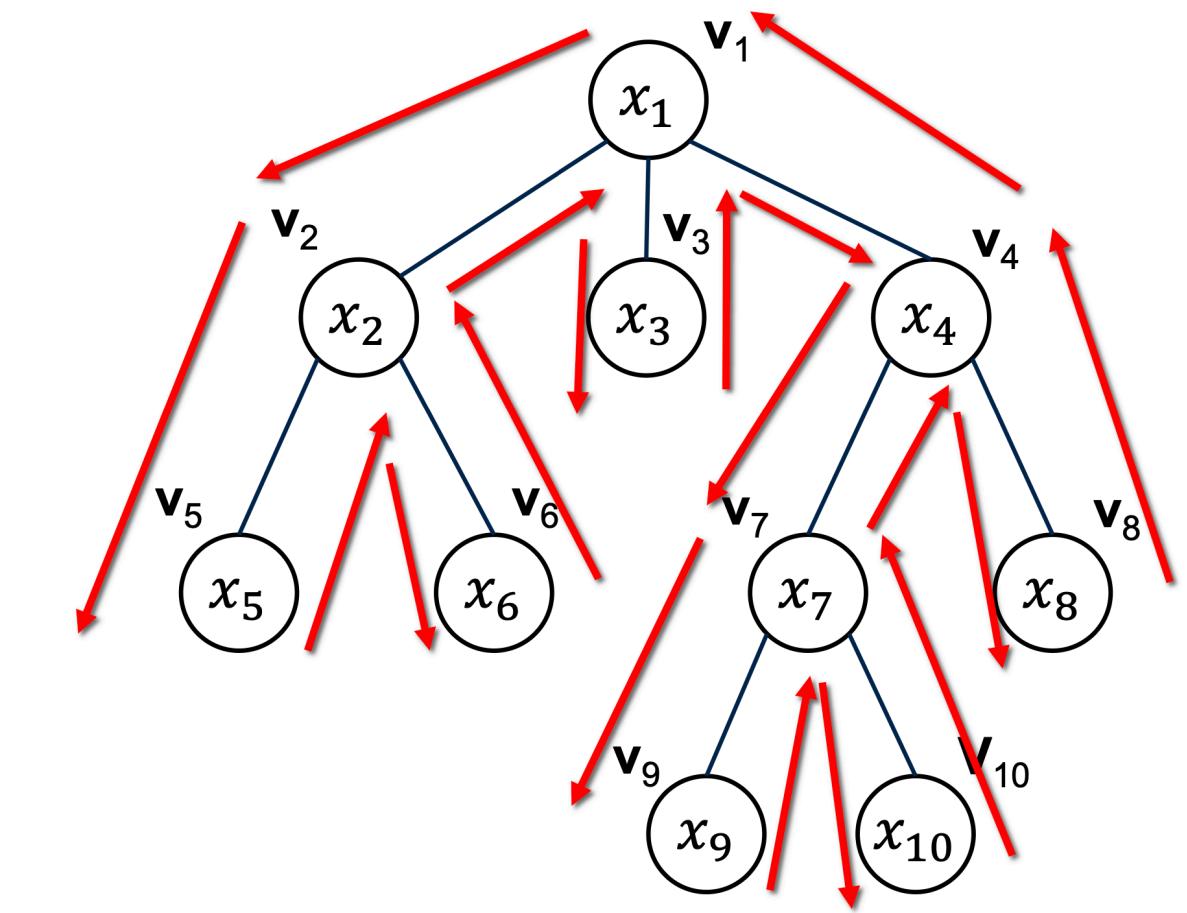
# Tree Accumulation: UA

Sum in subtree

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- $|ET| = 1 + 2(n - 1) = 2n - 1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Otherwise put 0

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_4$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	0	0	$x_3$	0	$x_4$	$x_7$	$x_9$	$x_{10}$	0	0	$x_8$	0	0	

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $i_l$ : index of last occurrence of  $v$
- $UA(v) = PS[i_l] - PS[i_f - 1]$



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$

How do we tell if an entry is the first occurrence? We need to find it in parallel and each one in  $O(1)$

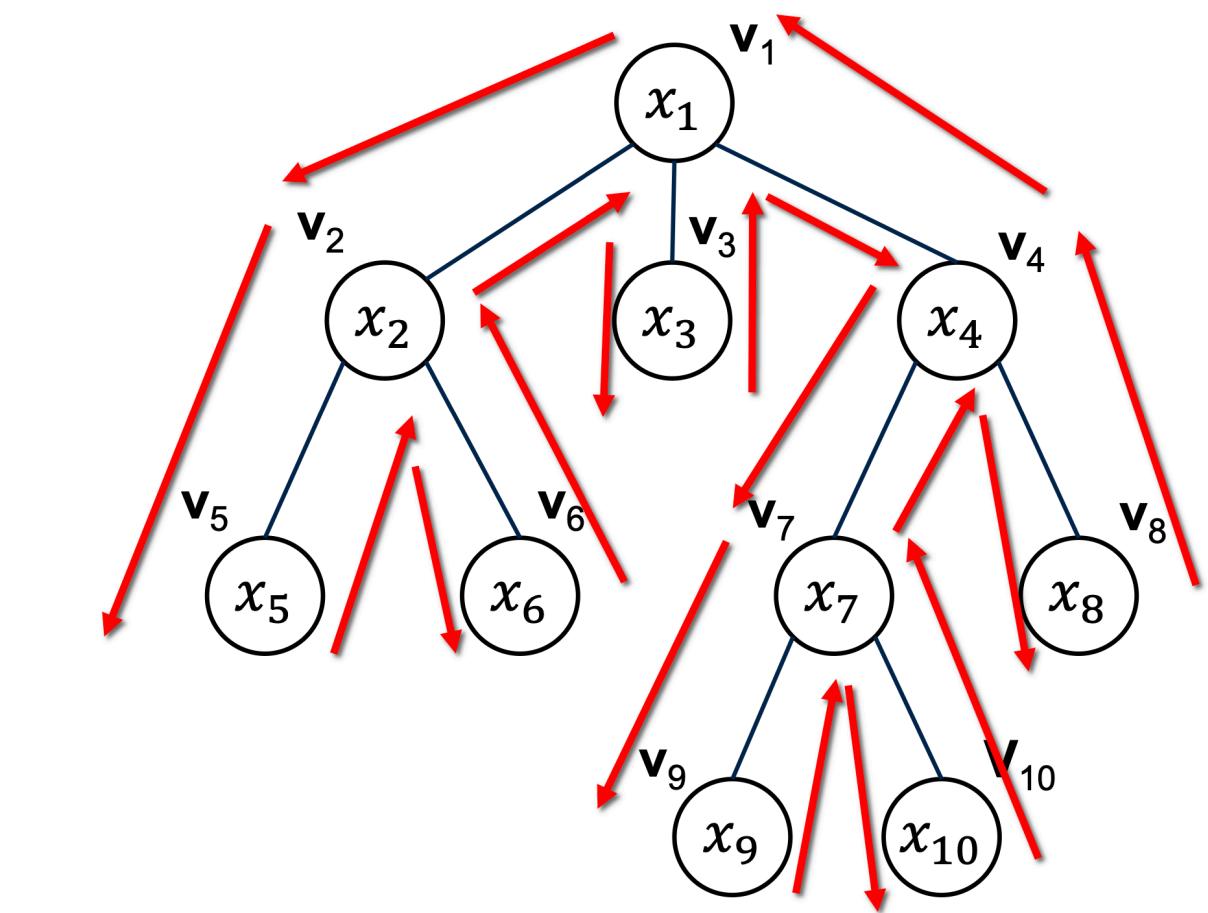
# Tree Accumulation: UA

Sum in subtree

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- $|ET| = 1 + 2(n - 1) = 2n - 1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Otherwise put 0

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_4$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	0	0	$x_3$	0	$x_4$	$x_7$	$x_9$	$x_{10}$	0	0	$x_8$	0	0	

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $i_l$ : index of last occurrence of  $v$
- $UA(v) = PS[i_l] - PS[i_f - 1]$



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$

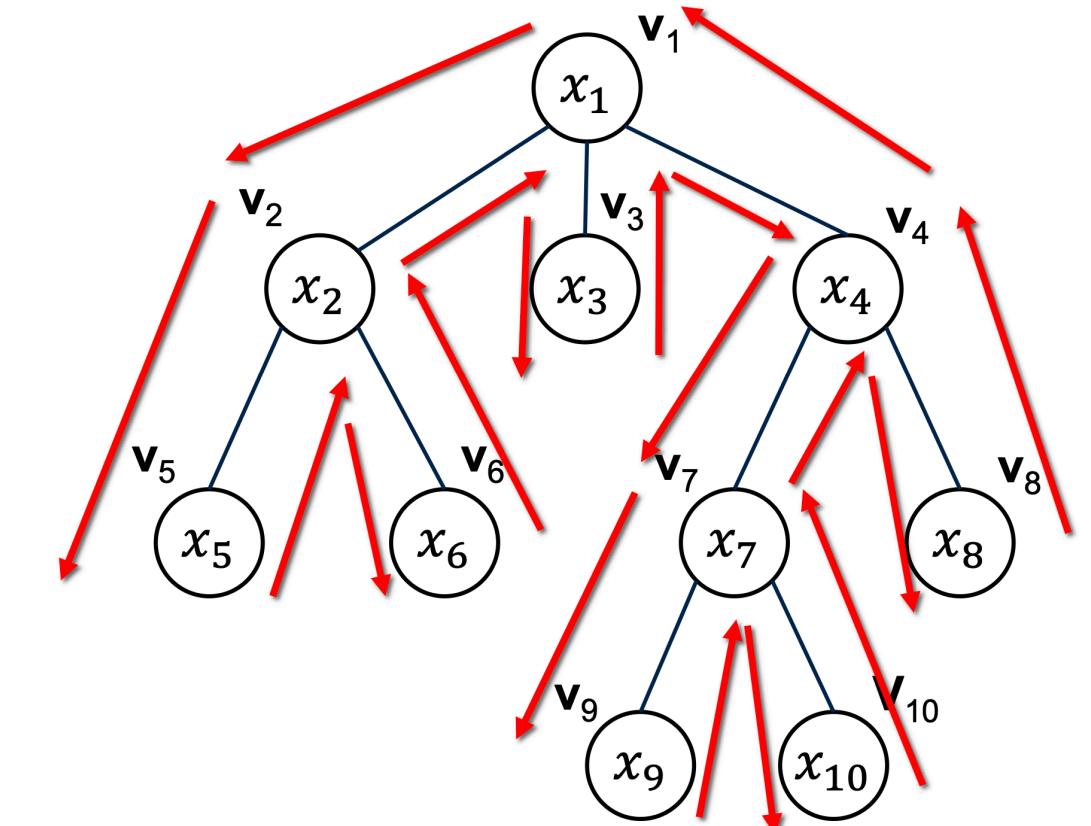
How do we tell if an entry is the first occurrence? We need to find it in parallel and each one in  $O(1)$

Look to your left - if it is the parent, you are the first occurrence

Sum up self +  
ancestors

# Tree Accumulation: DA

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Last occurrence of  $v_i$  put  $-x_i$
  - Otherwise put 0



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_4$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	$-x_2$	0	$x_3$	0	$x_4$	$x_9$	0	$x_{10}$	$-x_7$	0	$x_8$	$-x_4$	$-x_1$	

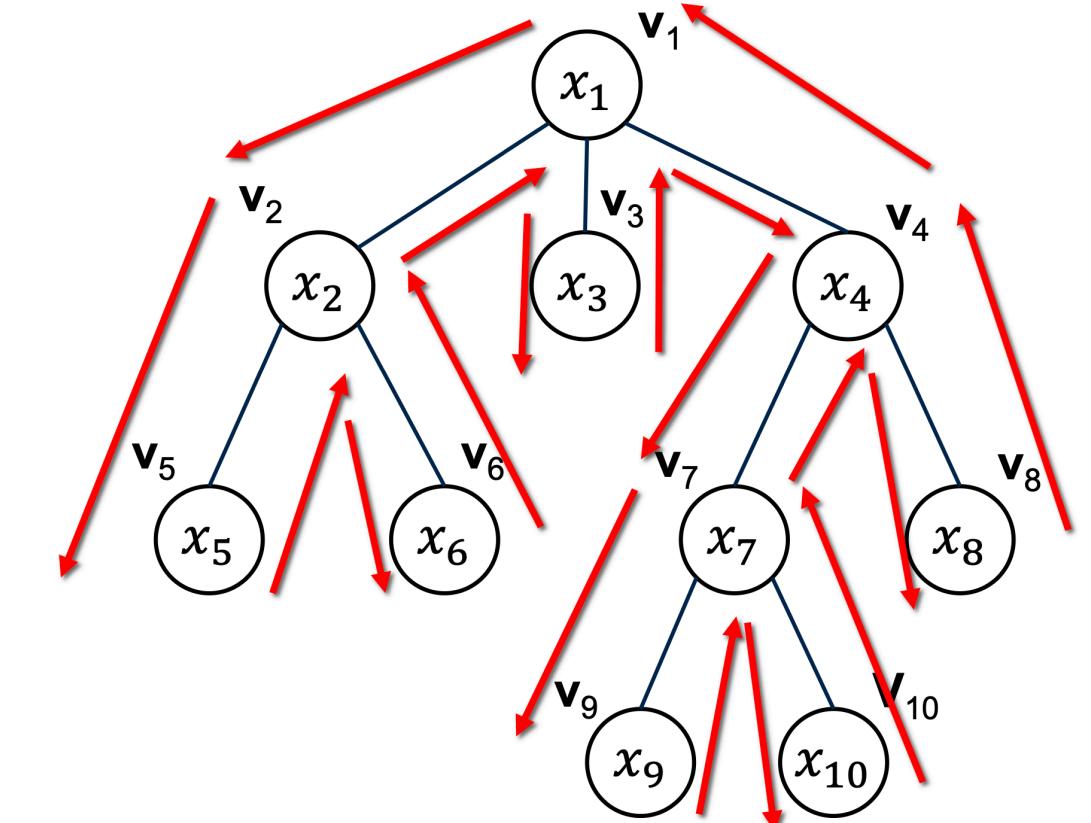
- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $DA(v) = PS[i_f]$

How do we know which is the last occurrence? We need to find it in parallel and each one in  $O(1)$

Sum up self +  
ancestors

# Tree Accumulation: DA

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Last occurrence of  $v_i$  put  $-x_i$
  - Otherwise put 0



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	$-x_2$	0	$x_3$	0	$x_4$	$x_9$	0	$x_{10}$	$-x_7$	0	$x_8$	$-x_4$	$-x_1$

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $DA(v) = PS[i_f]$

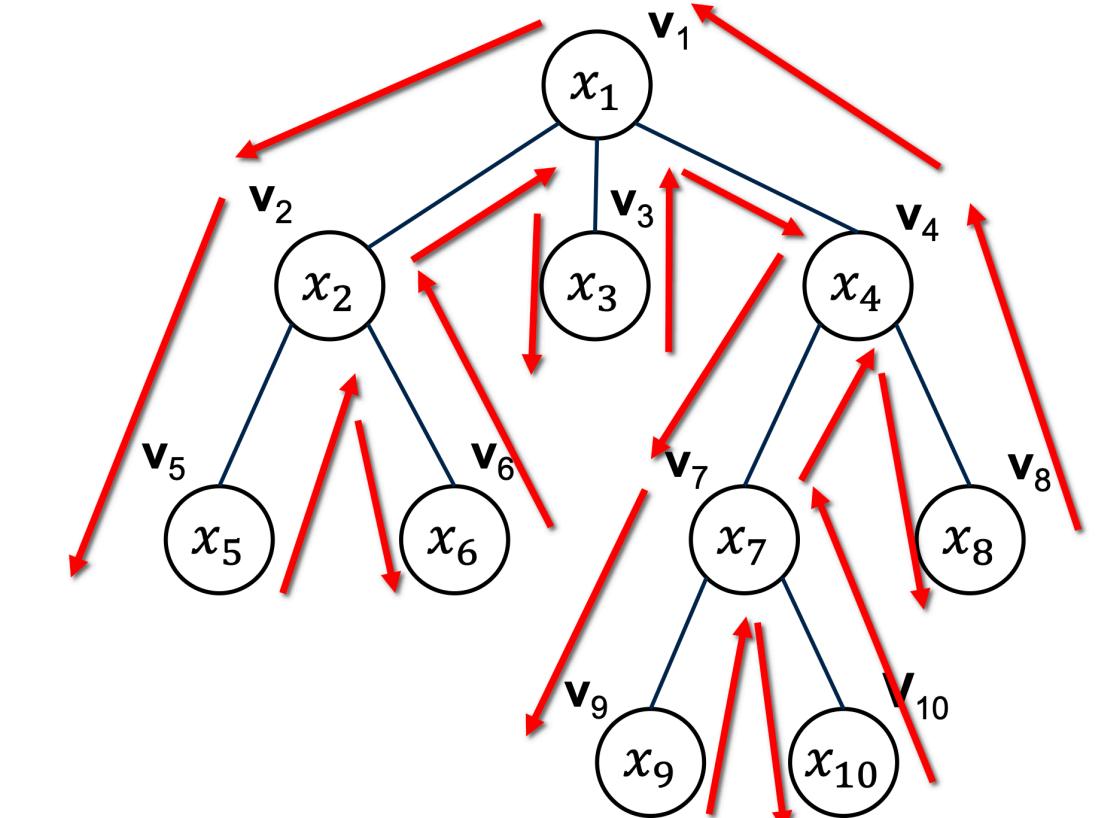
How do we know which is the last occurrence? We need to find it in parallel and each one in  $O(1)$

Look to your right - if it is the parent, you are the last occurrence

Sum up self +  
ancestors

# Tree Accumulation: DA

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Last occurrence of  $v_i$  put  $-x_i$
  - Otherwise put 0



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_{10} v_7 v_4 v_8 v_4 v_1$

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	$-x_2$	0	$x_3$	0	$x_4$	$x_9$	0	$x_{10}$	$-x_7$	0	$x_8$	$-x_4$	$-x_1$

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $DA(v) = PS[i_f]$

Is this correct? Hint:  
Look at the leaves

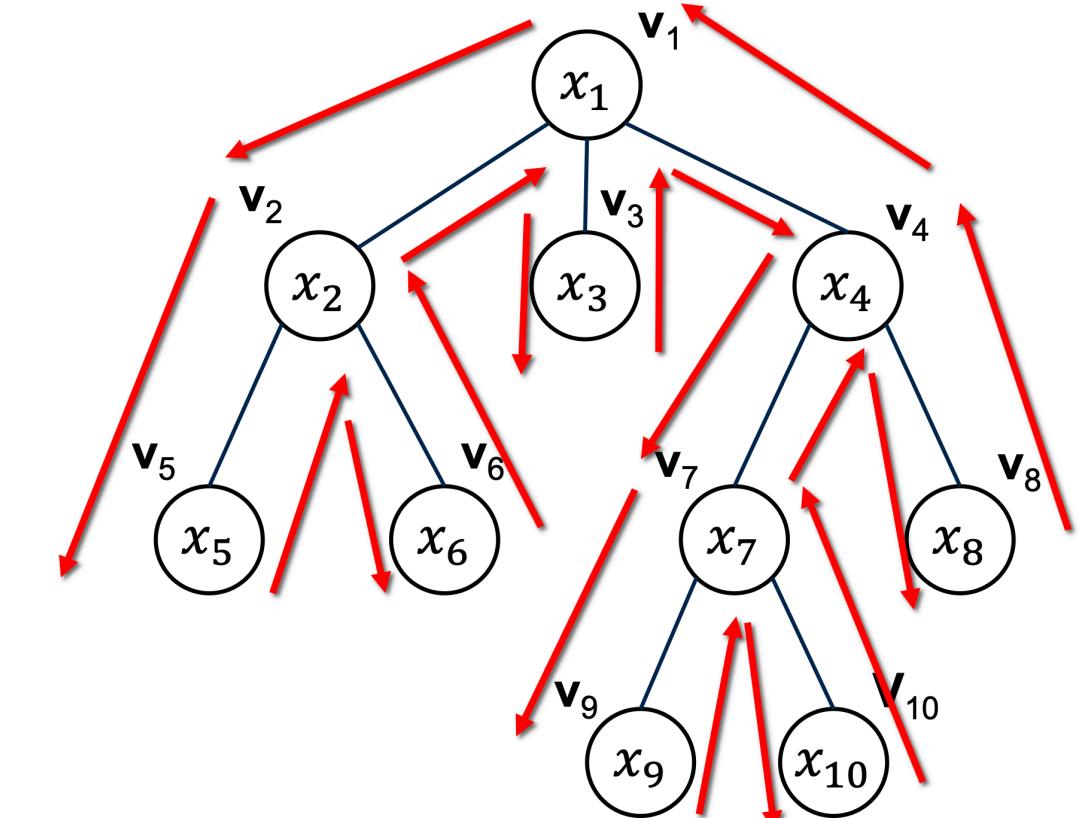
How do we know which is the last occurrence? We need to find it in parallel and each one in  $O(1)$

Look to your right - if it is the parent,  
you are the last occurrence

Sum up self +  
ancestors

# Tree Accumulation: DA

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- Assume ET is block partitioned among processors.
  - First occurrence of  $v_i$  put  $x_i$
  - Last occurrence of  $v_i$  put  $-x_i$
  - Otherwise put 0



Euler Tour:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_{10} v_7 v_4 v_8 v_4 v_1$

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	$x_5$	0	$x_6$	$-x_2$	0	$x_3$	0	$x_4$	$x_9$	0	$x_{10}$	$-x_7$	0	$x_8$	$-x_4$	$-x_1$

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- $DA(v) = PS[i_f]$

Is this correct? Hint:  
Check the leaves

How do we know which is the last occurrence? We need to find it in parallel and each one in  $O(1)$

Look to your right - if it is the parent,  
you are the last occurrence

Sum up self +  
ancestors

# Tree Accumulation: DA

- ET:  $v_1 v_2 v_5 v_2 v_6 v_2 v_1 v_3 v_1 v_4 v_7 v_9 v_7 v_{10} v_7 v_4 v_8 v_4 v_1$
- Assume ET is block partitioned among processors.
  - Internal nodes:
    - First occurrence of  $v_i$  put  $x_i$
    - Last occurrence of  $v_i$  put  $-x_i$
    - Otherwise put 0
  - Leaf nodes: put 0

$v_1$	$v_2$	$v_5$	$v_2$	$v_6$	$v_2$	$v_1$	$v_3$	$v_1$	$v_4$	$v_7$	$v_9$	$v_7$	$v_{10}$	$v_7$	$v_4$	$v_8$	$v_4$	$v_1$
$x_1$	$x_2$	0	0	0	$-x_2$	0	0	0	$x_4$	$x_7$	0	0	0	$-x_7$	0	0	$-x_4$	$-x_1$

- PS: resulting prefix sums array
- $i_f$ : index of first occurrence of  $v$
- If  $v_i$  is an internal node
  - $DA(v_i) = PS[i_f]$
- If  $v_i$  is a leaf node
  - $DA(v_i) = PS[i_f] + x_i$

Over all possible  
alignments

# Sequence Alignment

- An important problem in computational biology.
- DNA sequences: Strings over {A,C,G,T}.
- Goal: To find out how “evolutionarily close” the sequences are. This is tested by how “well” the sequences align.
- Alignment: Arranging characters of each sequence into columns to expose similarity.
- Gaps (-) may be inserted to denote insertions/deletions.
- Optimal Alignment: Minimum number of substitutions, insertions, and deletions needed to turn one sequence into another.

# Example Alignment and Score Computation

- Alignment has a score that shows quality.
- Every column of an alignment is a match, mismatch or a gap.
- Matches are preferred and hence have a positive score, others have a negative score.
- Example: If match = 1, mismatch = 0 and gap = -1, then for the following alignment

A	T	G	A	-	C	C
A	-	G	A	A	T	C
1	-1	1	1	-1	0	1

$$\text{Score}(\text{ATGACC}, \text{ AGAATC}) = 2$$

# Problem Definition

## Input:

- Two sequences  $A = a_1, a_2, \dots, a_m$  and  $B = b_1, b_2, \dots, b_n$
- Scores for match/mismatch ( $f(a, b)$ ) and gap ( $g$ )

## Algorithm:

Dynamic programming solution:

- $T$  = Table of size  $(m + 1) \times (n + 1)$
- $T[i, j]$  = best score between  $a_1, a_2, \dots, a_i$  and  $b_1, b_2, \dots, b_j$

$$\bullet T[i, j] = \max \begin{cases} T[i - 1, j - 1] + f(a_i, b_j) \\ T[i - 1, j] - g \\ T[i, j - 1] - g \end{cases}$$

match/mismatch

gap in B

gap in A

- Sequential time =  $O(mn)$

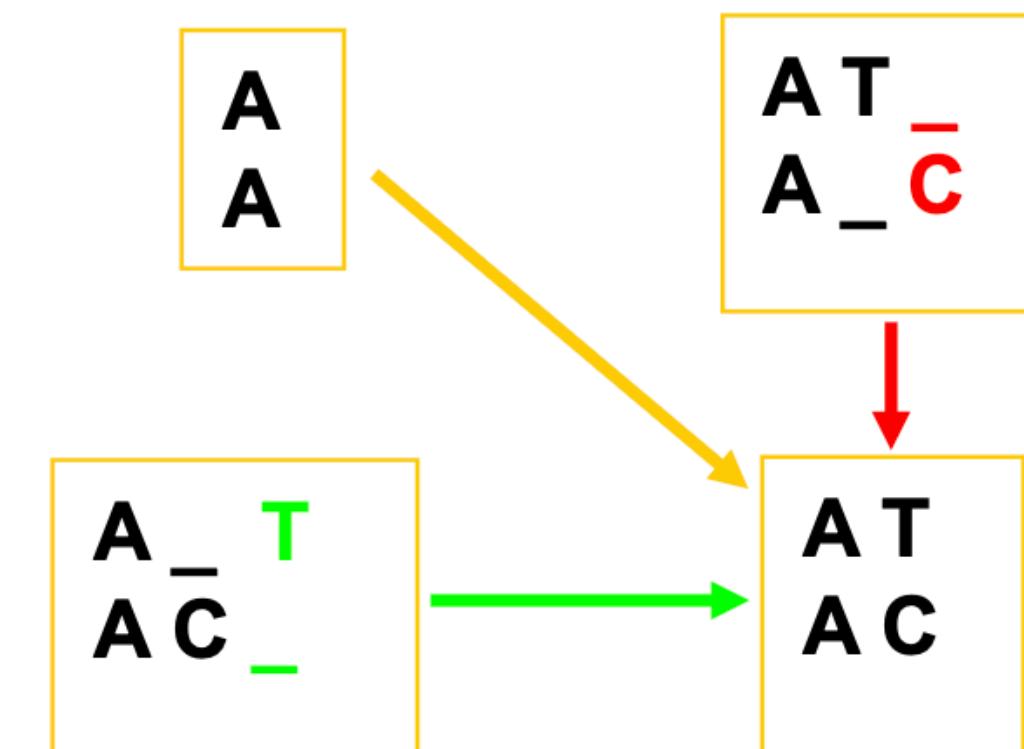
# Dynamic Programming Matrix

$$T[i, j] = \max \left\{ \begin{array}{l} T[i - 1, j - 1] + f(a_i, b_j) \\ T[i - 1, j] - g \\ T[i, j - 1] - g \end{array} \right.$$

Match = 2

Substitution = 0

Gap = -1



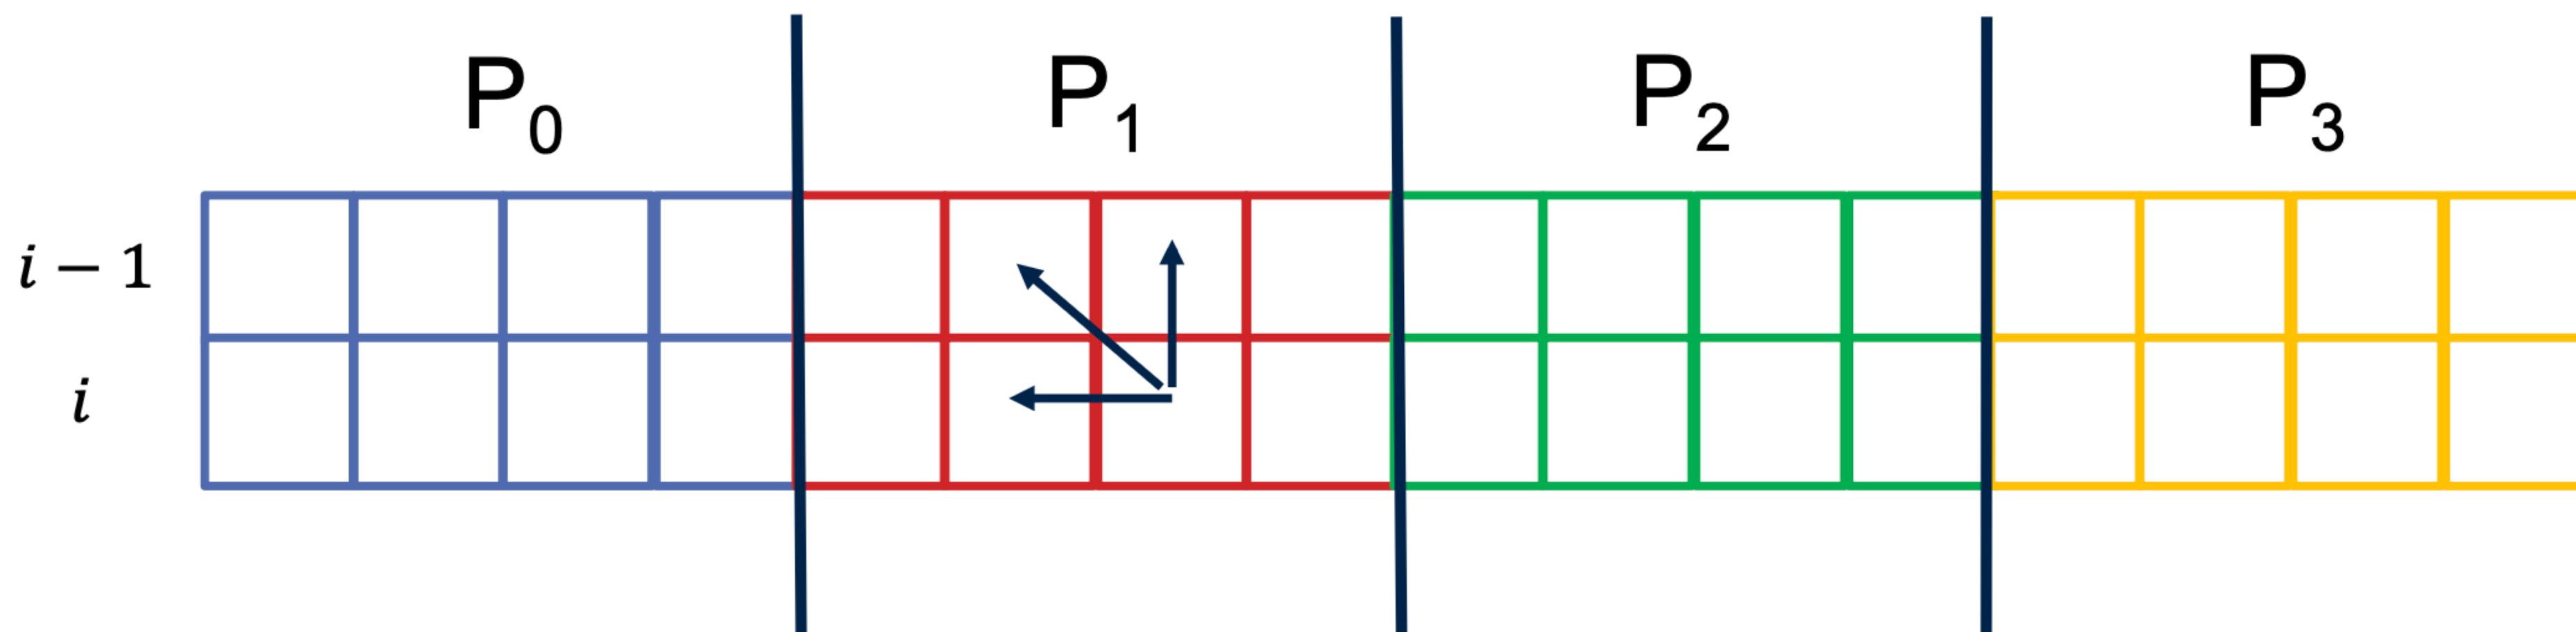
Initialize all gaps

	0	1	2	3	4	5	6
0		Gap	A	T	G	C	T
1	Gap	0	-1	-2	-3	-4	-5
2	A	-1	2 (2) ↓ 1 (0)	1 (0)	0	0	0
3	C	-2	1 (0) ↓ 2 (0)	2 (0)	1	2	1
4	C	-3	0	1 (0) ↓ 3 (2)	2	3	2
5	G	-4	0	0	3 (2) ↓ 5 (2)	2	3
6	C	-5	0	0	2	5 (2) ↓ 7 (2)	4
7	T	-6	0	2	1	4	7 (2)

Max score

# Solution Using Parallel Prefix

- We compute each row of T in parallel:



$$\bullet w[j] = \max \begin{cases} T[i-1, j-1] + f(a_i, b_j) \\ T[i-1, j] - g \end{cases}$$

Can be computed in parallel upfront with previous row

$$\bullet T[i, j] = \max \begin{cases} w[j] \\ T[i, j-1] - g \end{cases}$$

max is a binary associative operator, but  $-g$  poses issue for immediate use of parallel prefix

# Solution Using Parallel Prefix

j is a global index

$$\bullet x[j] = T[i, j] + j \cdot g$$

$$\bullet x[j] = \max \begin{cases} w[j] + jg \\ T[i, j - 1] - g + jg \end{cases}$$

$$\bullet x[j] = \max \begin{cases} w[j] + jg \\ T[i, j - 1] + (j - 1)g \end{cases}$$

$$\bullet x[j] = \max \begin{cases} w[j] + jg \\ x[j - 1] \end{cases}$$

compute  $x[j]$  using parallel prefix, then get  $T[i, j]$

# Parallel Algorithm: Compute row $i$ from $i-1$

1. Use right shift to send last element of  $(i - 1)^{\text{th}}$  row on each processor.
2. Create vector A

$$A[j] = jg + \max \begin{cases} T[i-1, j-1] + f(a_i, b_j) \\ T[i-1, j] - g \end{cases} \quad O\left(\frac{n}{p}\right)$$

3. Compute parallel prefix on A using max as operator and store results in  $x$
4. Compute  $T[i, j]$ 's using

$$T[i, j] = x[j] - jg \quad O\left(\frac{n}{p} + (\tau + \mu) \log p\right) \quad O\left(\frac{n}{p}\right)$$

Computation Time =  $O\left(\frac{mn}{p} + m \log p\right)$

Communication Time =  $O((\tau + \mu)m \log p)$

Efficient when  $p = O\left(\frac{n}{\log n}\right)$

Needs to be done  
for each row

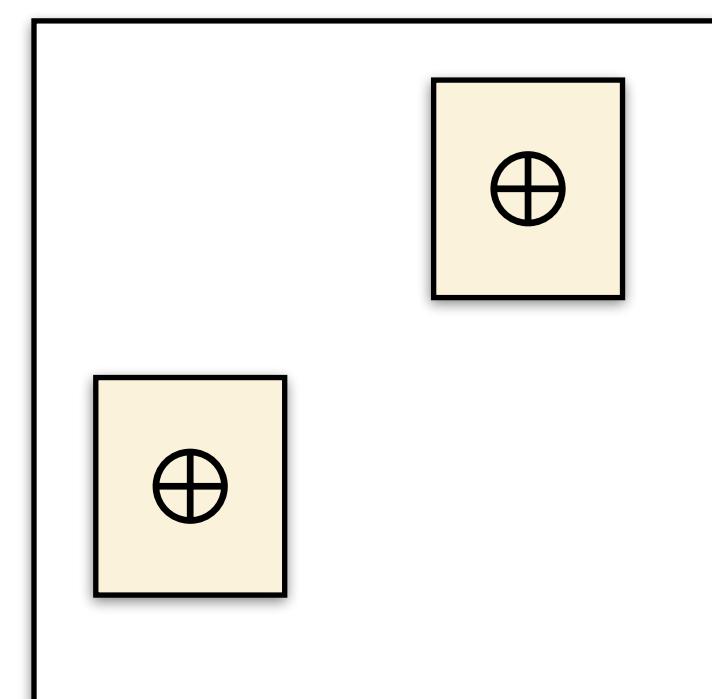
# Application: Tensor Region Sums

Several scientific computing applications involve reducing **many (potentially overlapping) regions** of a tensor to a single value for each region, using a binary associative operator  $\oplus$ .

**Inclusion:** The **summed-area table** (SAT) method preprocesses an image to answer queries about the sum in rectangular subregions of a tensor [C84].

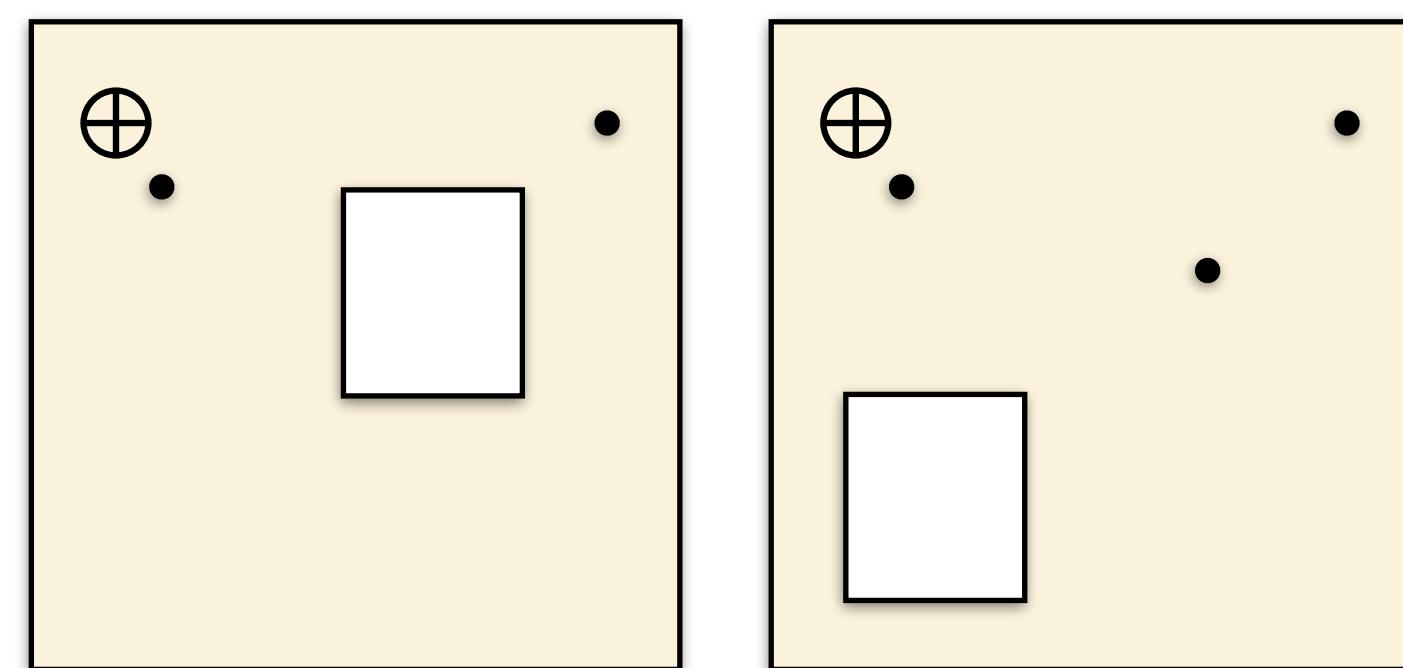
**Exclusion:** The essence of the **fast multipole method** (FMM) is a reduction of a subregion's elements, excluding elements too close [BG97, CRW93, D00].

Summed-area Table



computes reductions for all N points

Fast Multipole Method



computes reductions for all excluded regions

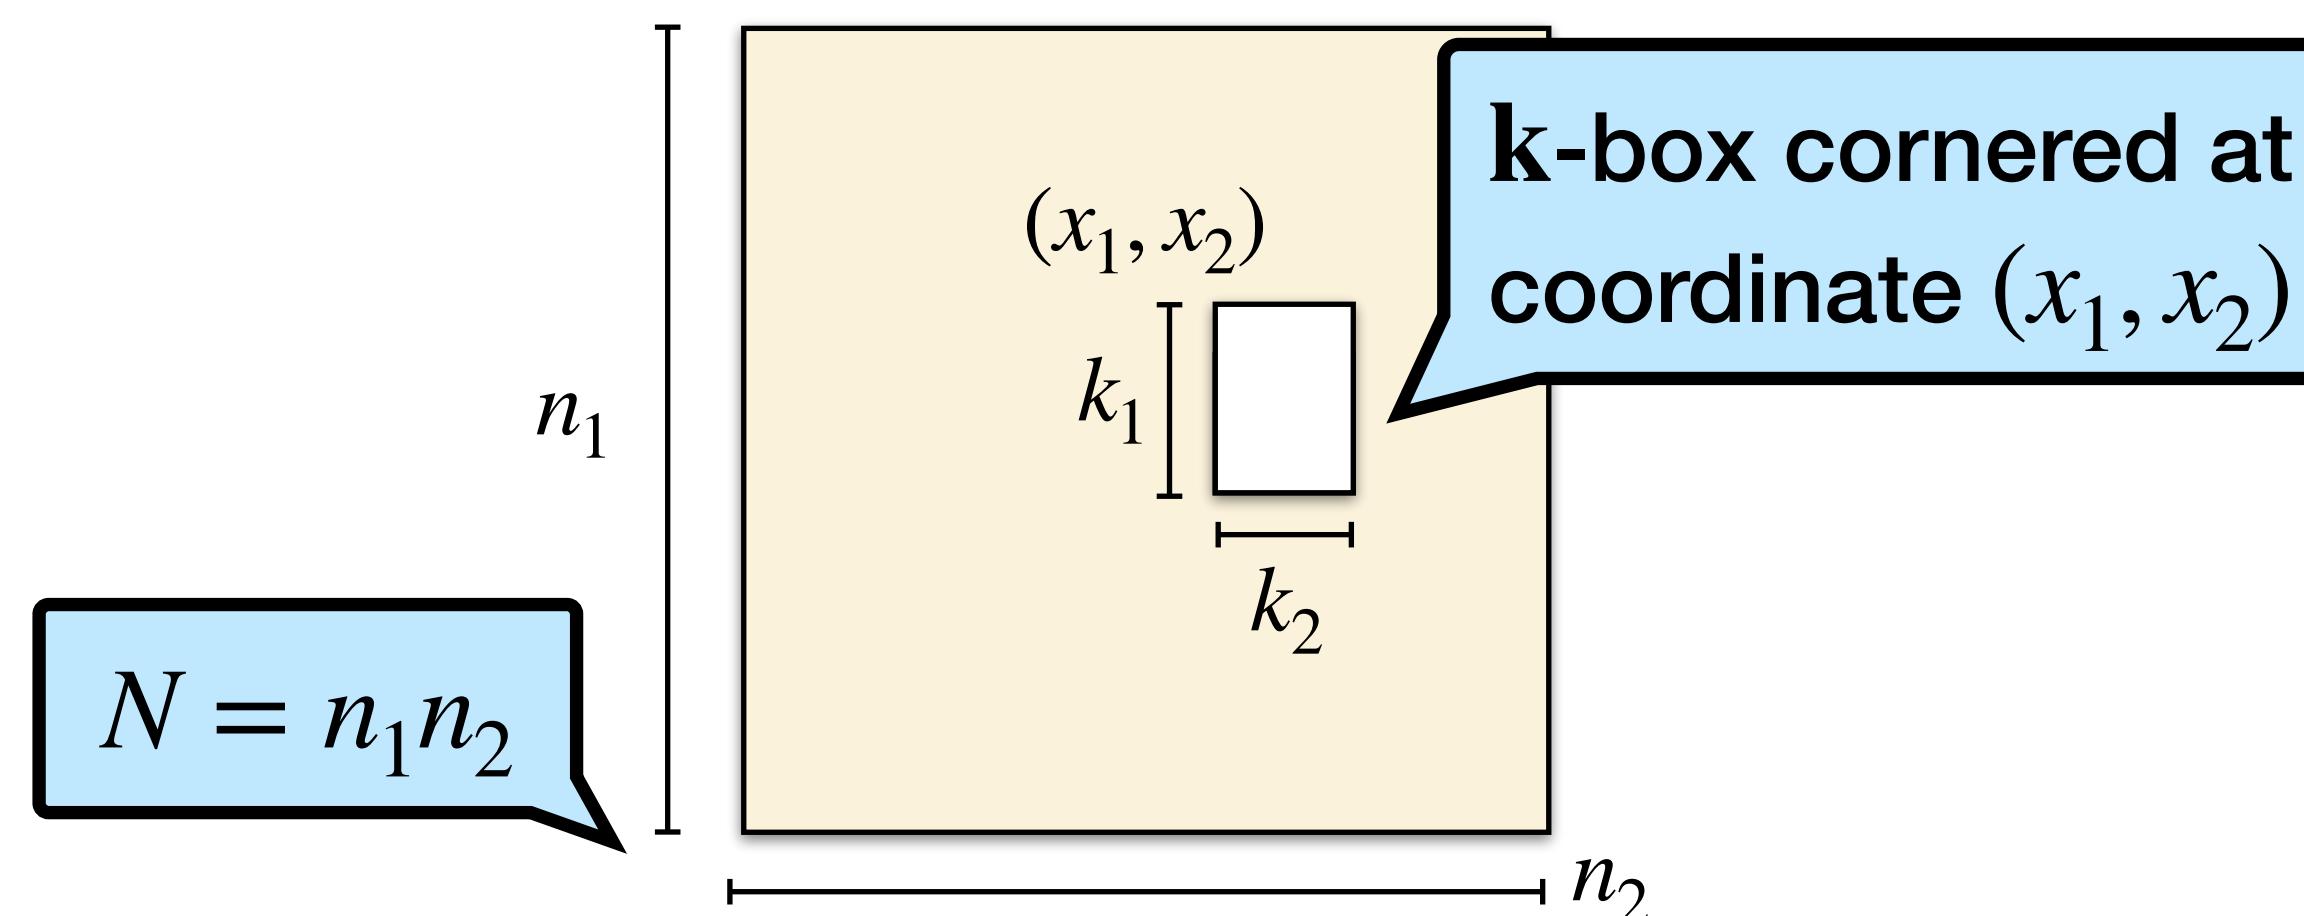
Binary associative operator

# Excluded-sums problem

The excluded-sums problem [DDELP05] underlies applications that require **reducing regions of a tensor** to a single value using  $\oplus$ .

In 2D, it takes as input an  $n_1 \times n_2$  matrix  $A$  and “box size”  $\mathbf{k} = (k_1, k_2)$  where  $k_1 \leq n_1, k_2 \leq n_2$ .

The problem involves reducing the excluded region outside of **every  $\mathbf{k}$ -box** in the matrix.



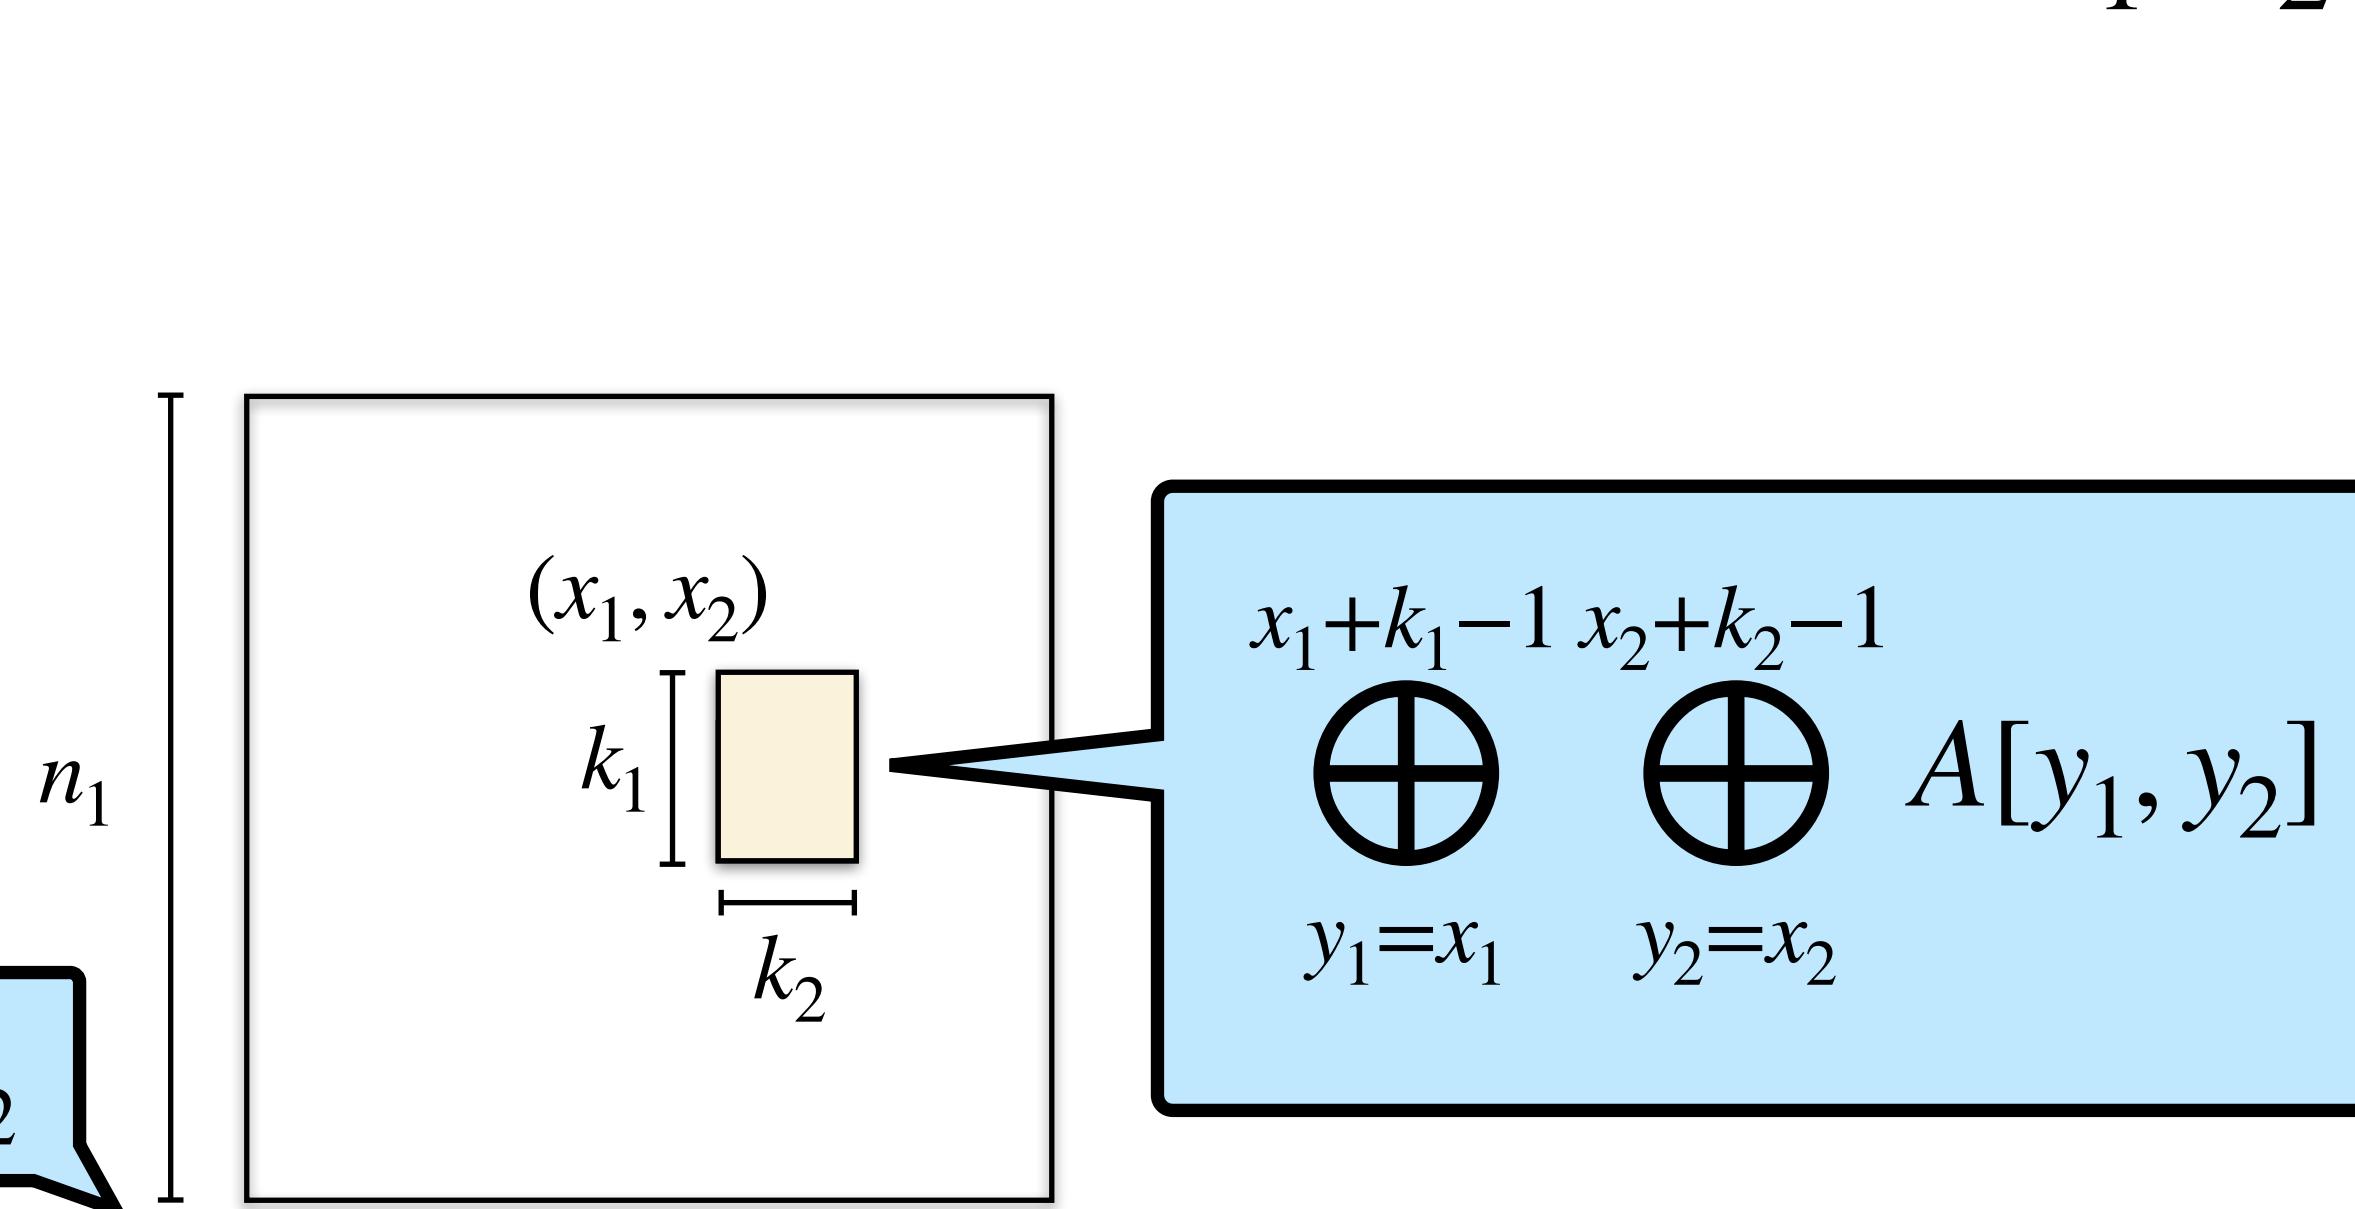
# Included-sums Problem

The included-sums problem takes the same input as the excluded-sums problem.

In 2D, the included sum at coordinate  $(x_1, x_2)$  involves reducing (accumulating with  $\oplus$ ) all elements **in the k-box** cornered at  $(x_1, x_2)$ .

Can be computed straightforwardly with four nested loops in  $\Theta(n_1 n_2 k_1 k_2)$  time.

$$N = n_1 n_2$$

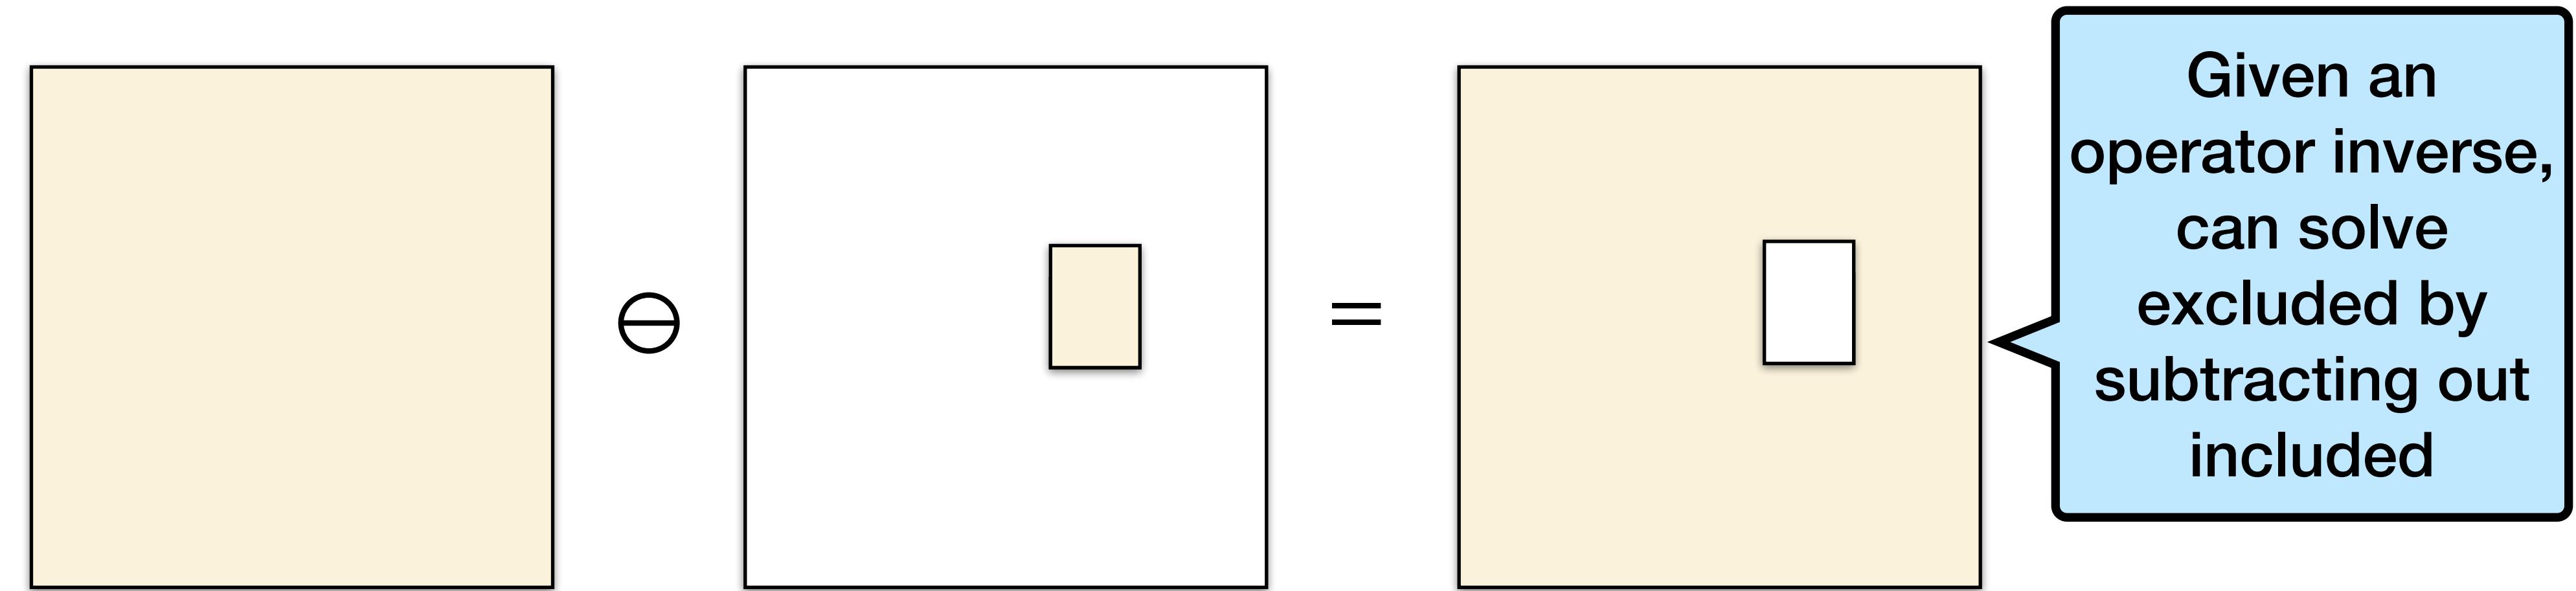


# Inclusion and Exclusion Example

Input					Inclusion					Exclusion				
1	3	6	2	5	16	19	10	10	7	75	72	81	81	84
3	9	1	1	2	18	16	10	8	4	73	75	81	83	87
5	1	5 3		2	13	11	10	14	11	78	80	81	77	80
4	3	2 0		9	15	8	10	24	17	76	83	81	67	74
6	2	1	7	8	8	3	8	15	8	83	78	83	76	83

We present an example with addition for ease of understanding, but in general an algorithm for these problems should work with general operators.

# Included and Excluded Sums With and Without Operator Inverse



This approach fails for **operators without inverse** such as max, or the FMM's functions, which may exhibit singularities [DDELP05].

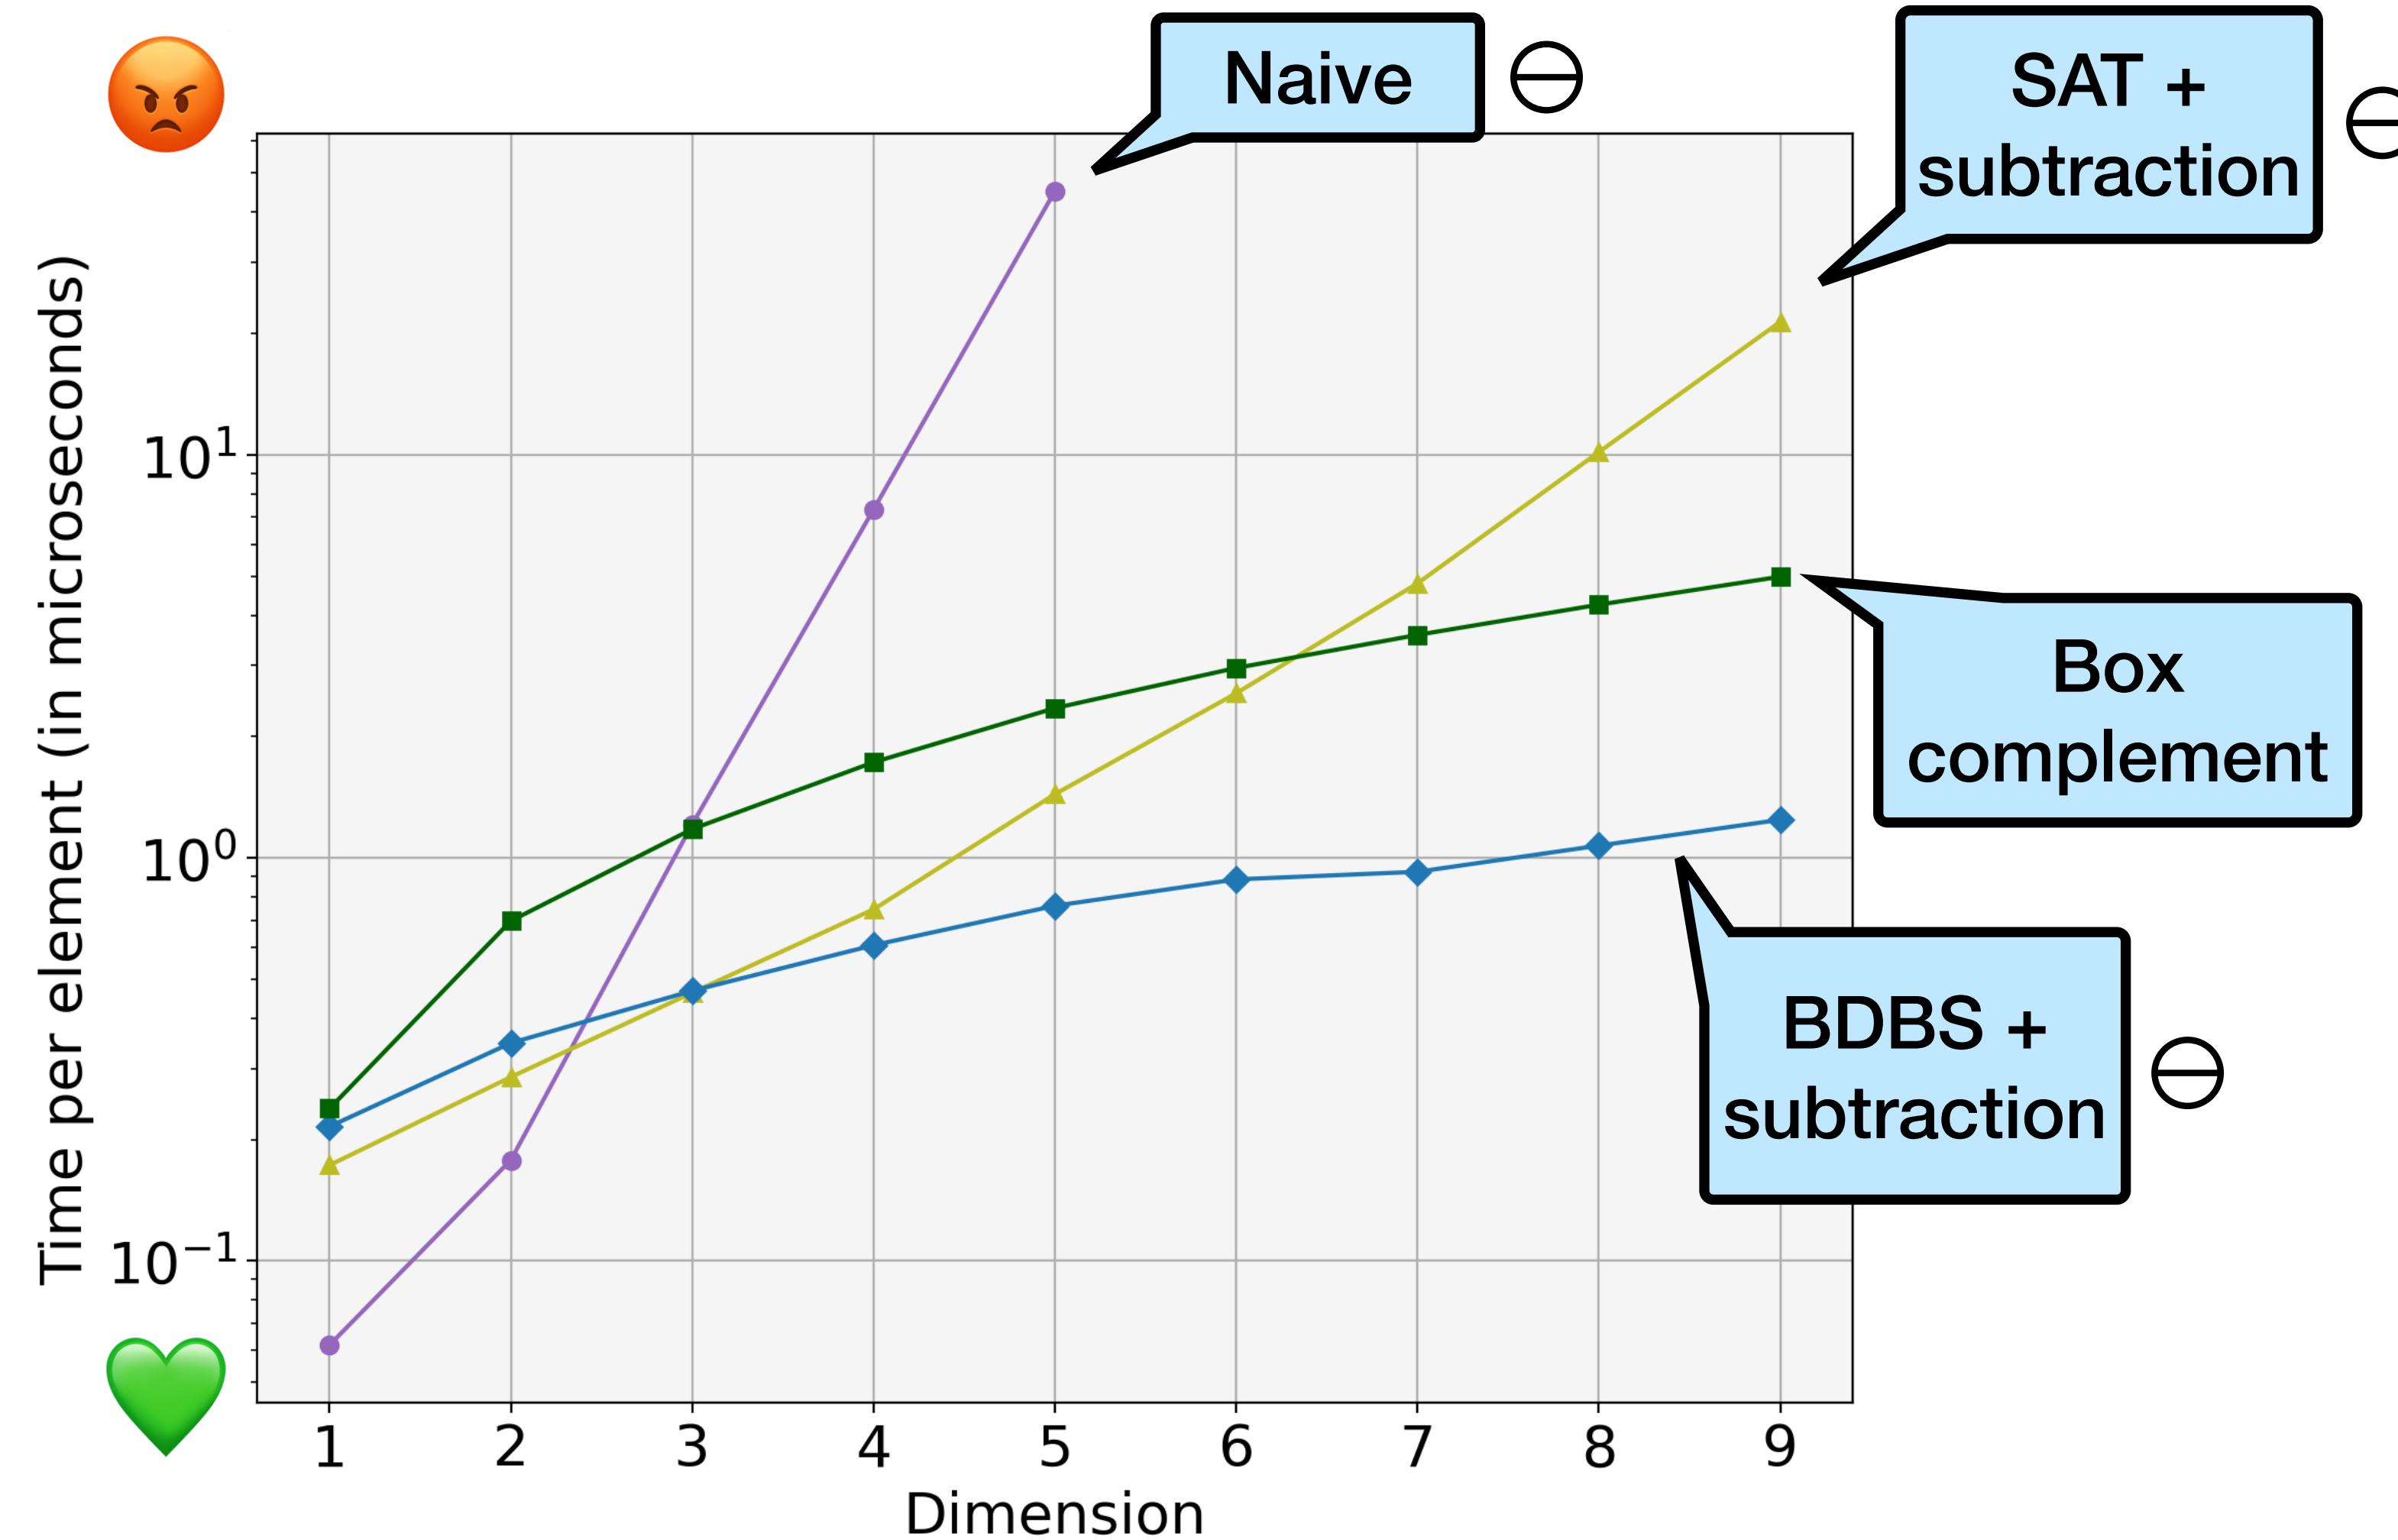
We refine the included- and excluded-sums problems into **weak** and **strong** versions. The weak version requires an operator inverse, while the strong version does not.

# Algorithmic Bounds

Given a  $d$ -dimensional tensor with  $N$  elements:

Algorithm	Problem	Weak/Strong	Time	Space
Summed-area table	Included	Weak 😐	$\Theta(2^d N)$ 😐	$\Theta(N)$ ❤️
Corners( $c$ )	Excluded	Strong ❤️	$\Omega(2^d N)$ 😐	$\Theta(cN)$ 😠
Bidirectional box-sum (BDBS)	Included	Strong ❤️	$\Theta(dN)$ ❤️	$\Theta(N)$ ❤️
Box complement	Excluded	Strong ❤️	$\Theta(dN)$ ❤️	$\Theta(N)$ ❤️

# Weak and Strong Excluded Sums in Higher Dimensions



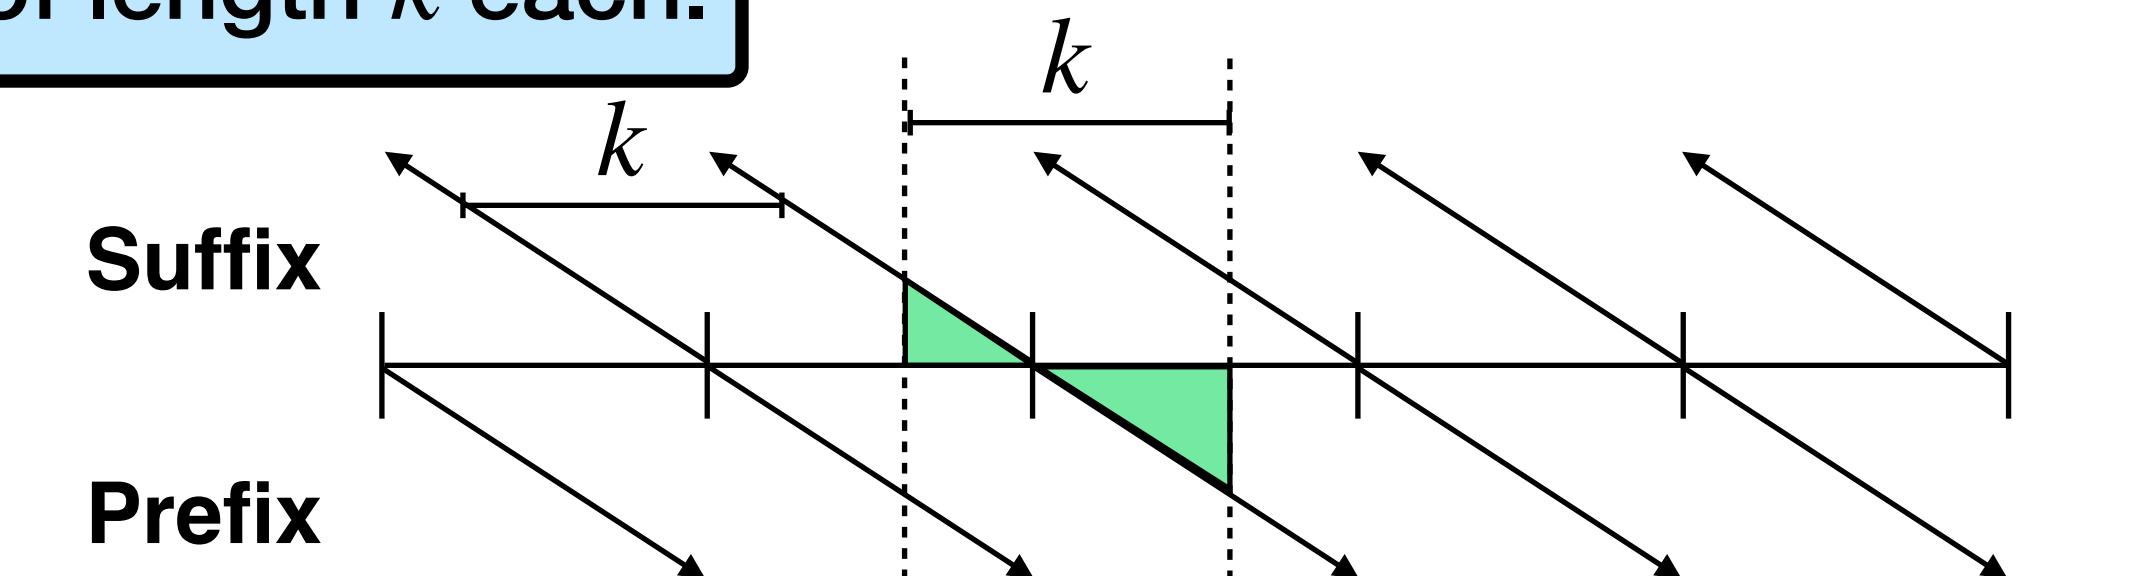
# Bidirectional Box-sum Algorithm for Strong Included Sums

We will start with the bidirectional box-sum algorithm (BDBS) in one dimension then show how to extend the technique to higher dimensions.

Given a list  $A$  of length  $N$  and a (scalar) box size  $k$ , output a list  $A'$  of included sums.

Position	1	2	3	4	5	6	7	8
Input $\rightarrow A$	2	5	3	1	6	3	9	0
Prefix $\rightarrow A_p$	2	7	10	11	6	9	18	18
Suffix $\rightarrow A_s$	11	9	4	1	18	12	9	0
Target $\rightarrow A'$	11	15	13	19	18	12	9	0

Compute intermediate prefix and suffix arrays with  $N/k$  prefixes and suffixes of length  $k$  each.

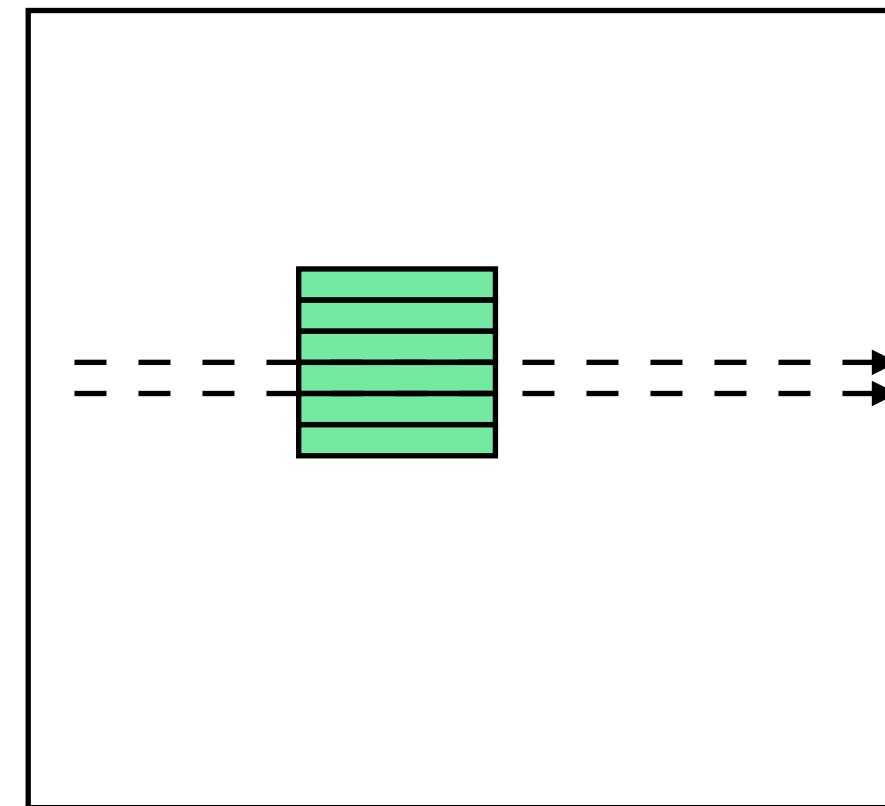


$\Theta(N)$  time, space in 1D

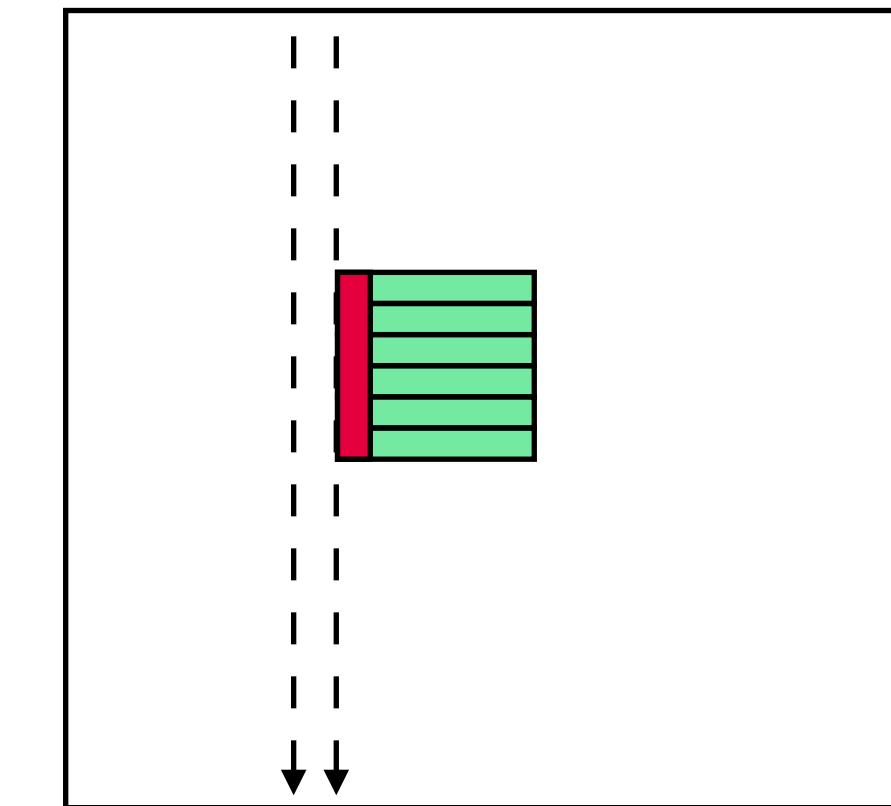
# Multidimensional Bidirectional Box Sum

The BDBS technique extends into arbitrary dimensions by performing the prefixes and suffixes along each dimension in turn.

**Bidirectional box sum  
along first dimension**



**Bidirectional box sum along  
second dimension on result**

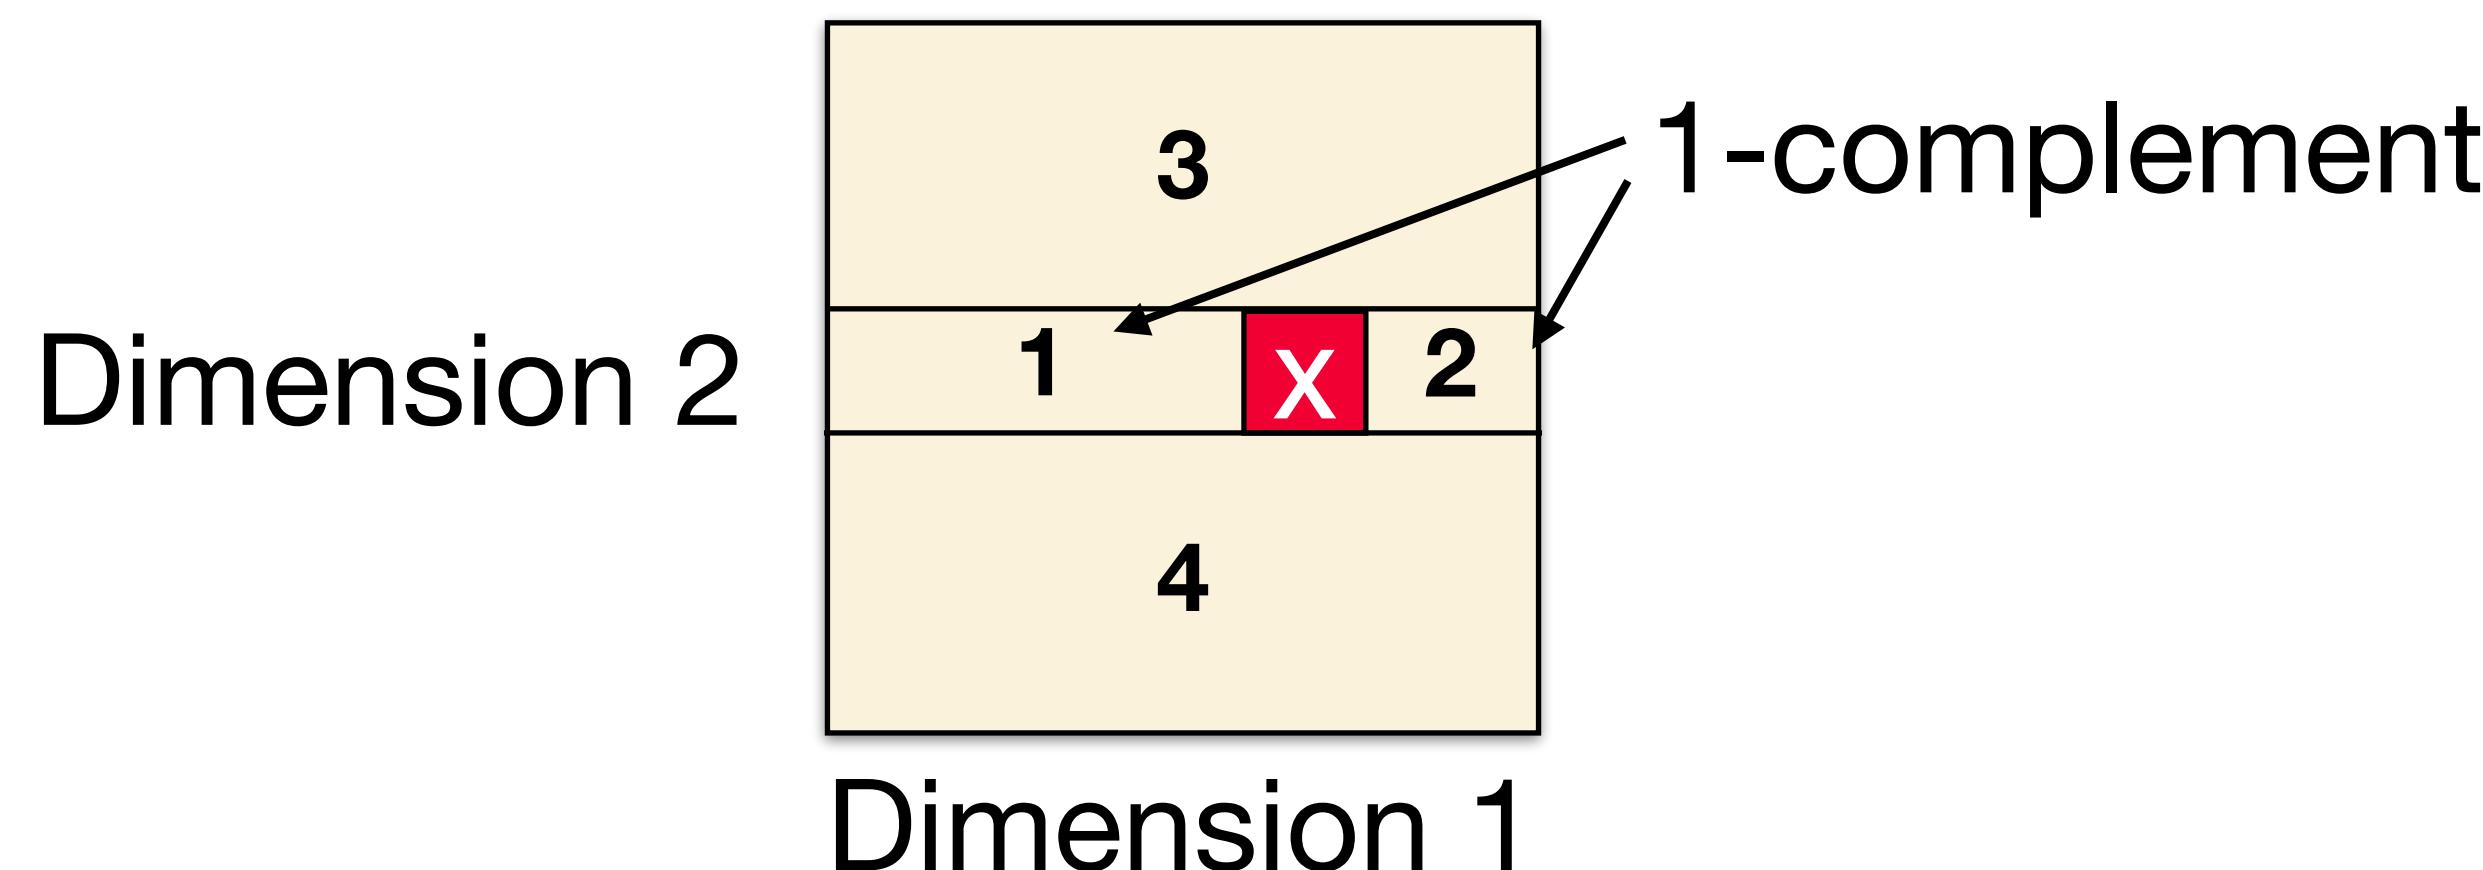


Given a  $d$ -dimensional tensor with  $N$  elements, BDBS solves the strong included-sums problem in  $\Theta(dN)$  time and  $\Theta(N)$  space.

# Formulating the Excluded Sum as the Box Complement

Given a  $d$ -dimensional tensor and a “box size”, we will first sketch how to decompose the excluded region for each point into  $2d$  disjoint regions.

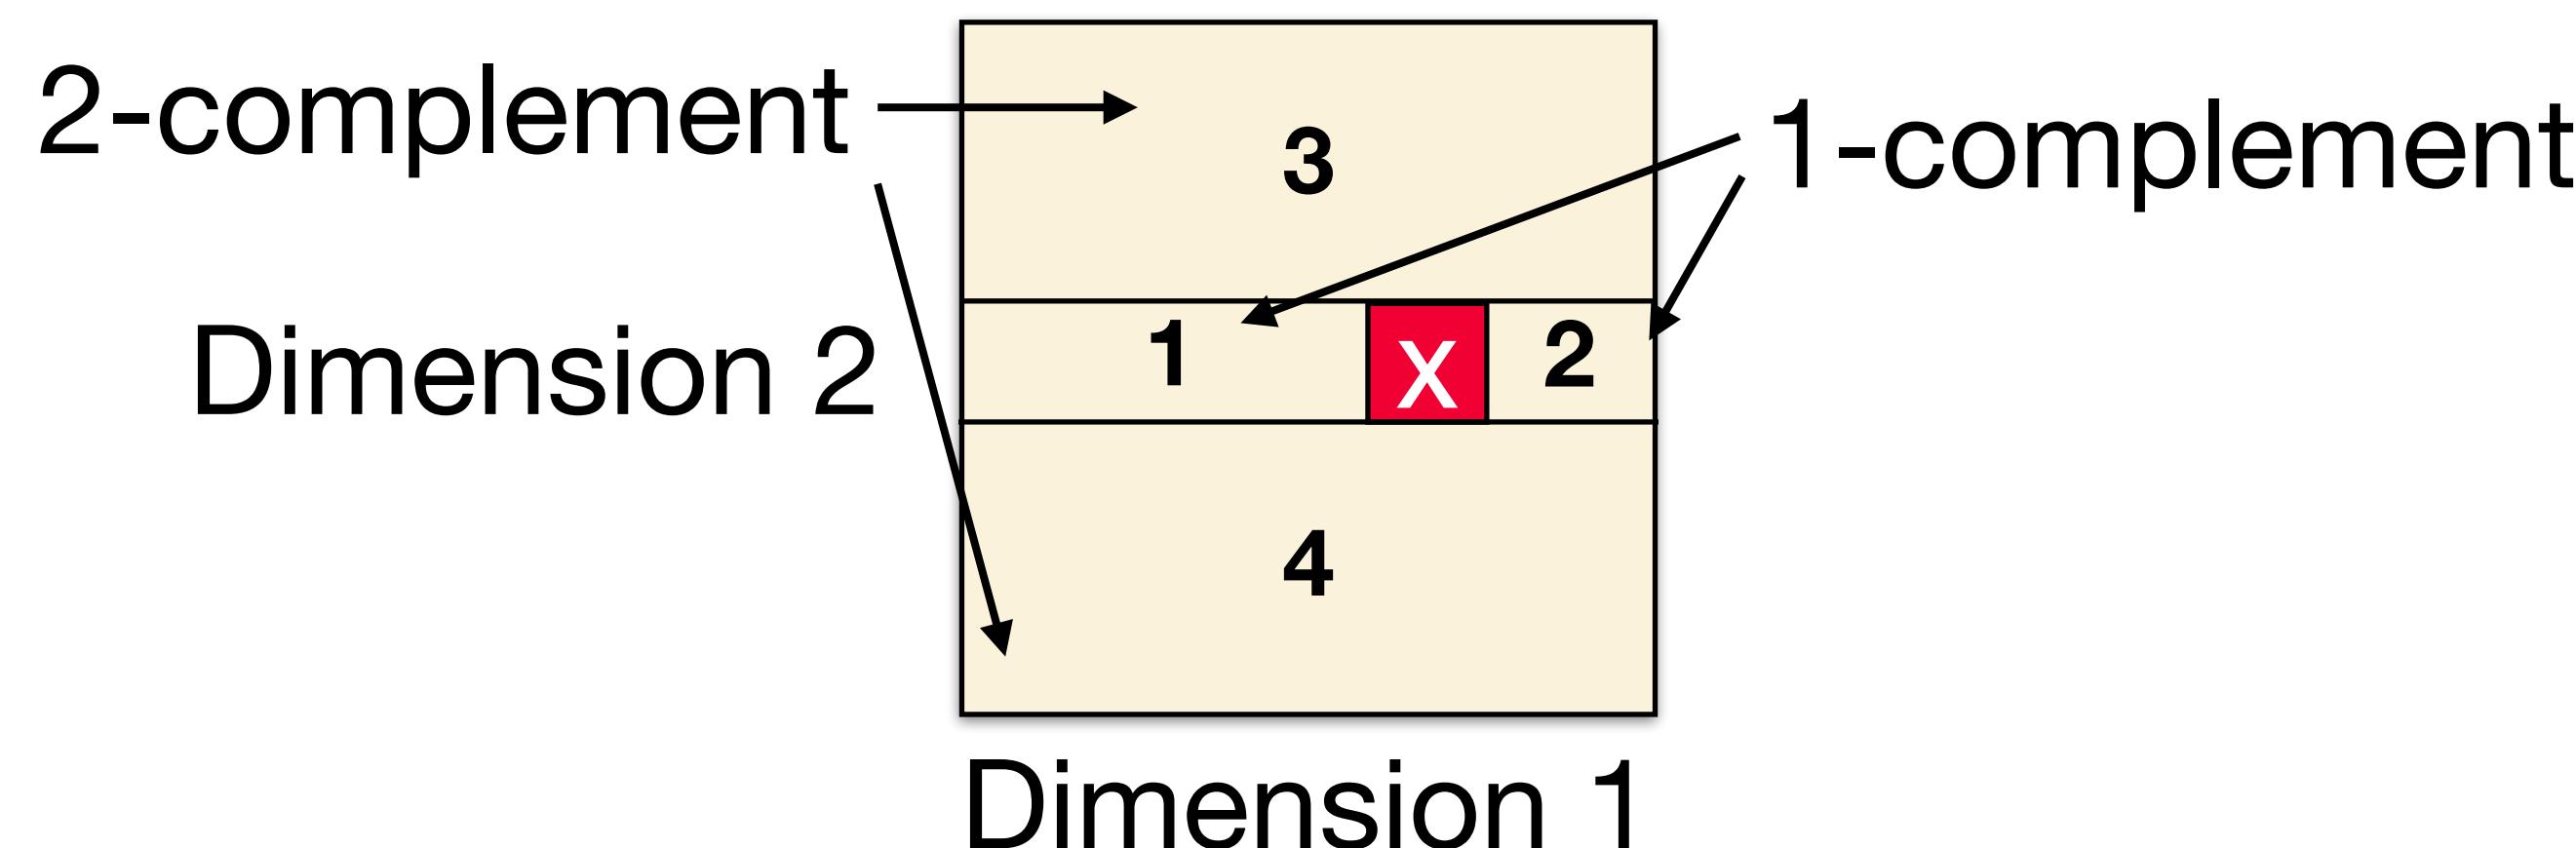
At a high level, the “ $i$ -complement” of a box such that there is some coordinate in dimension  $j \in [1, i]$  that is “out of range” in dimension  $j$ , and the coordinates are “in range” for all dimensions  $m \in [i + 1, d]$ .



# Formulating the Excluded Sum as the Box Complement

Given a  $d$ -dimensional tensor and a “box size”, we will first sketch how to decompose the excluded region for each point into  $2d$  disjoint regions.

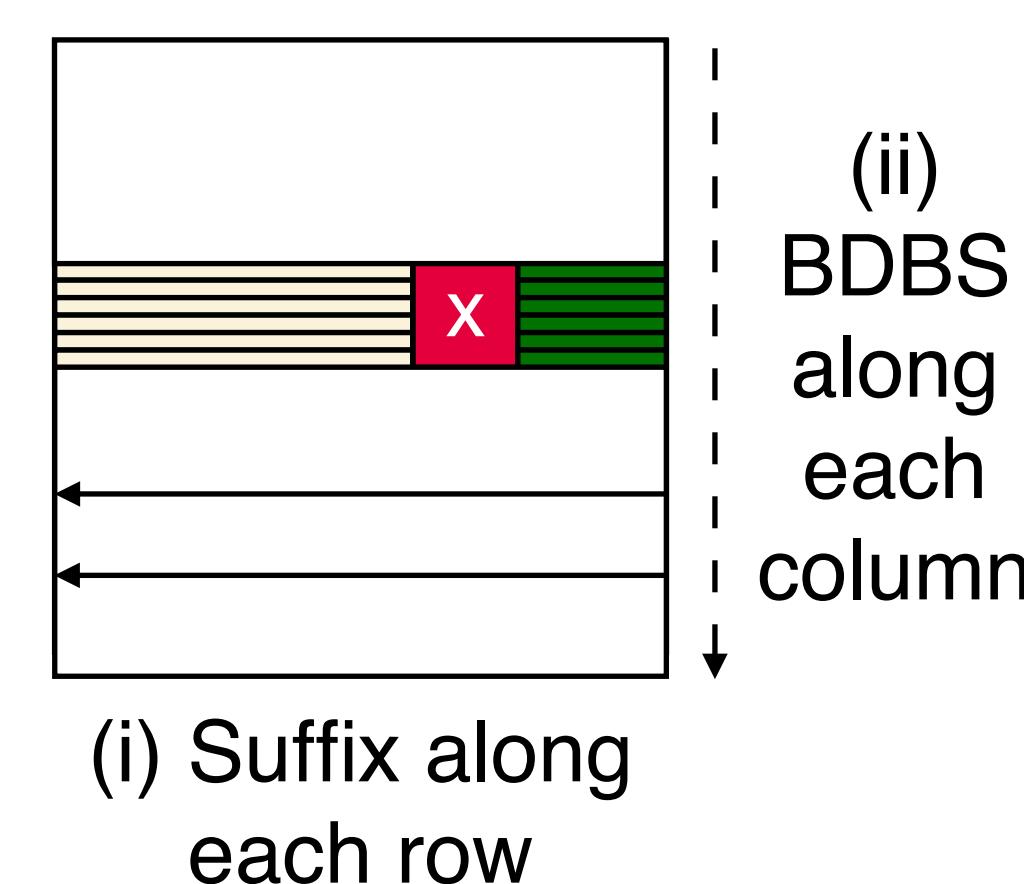
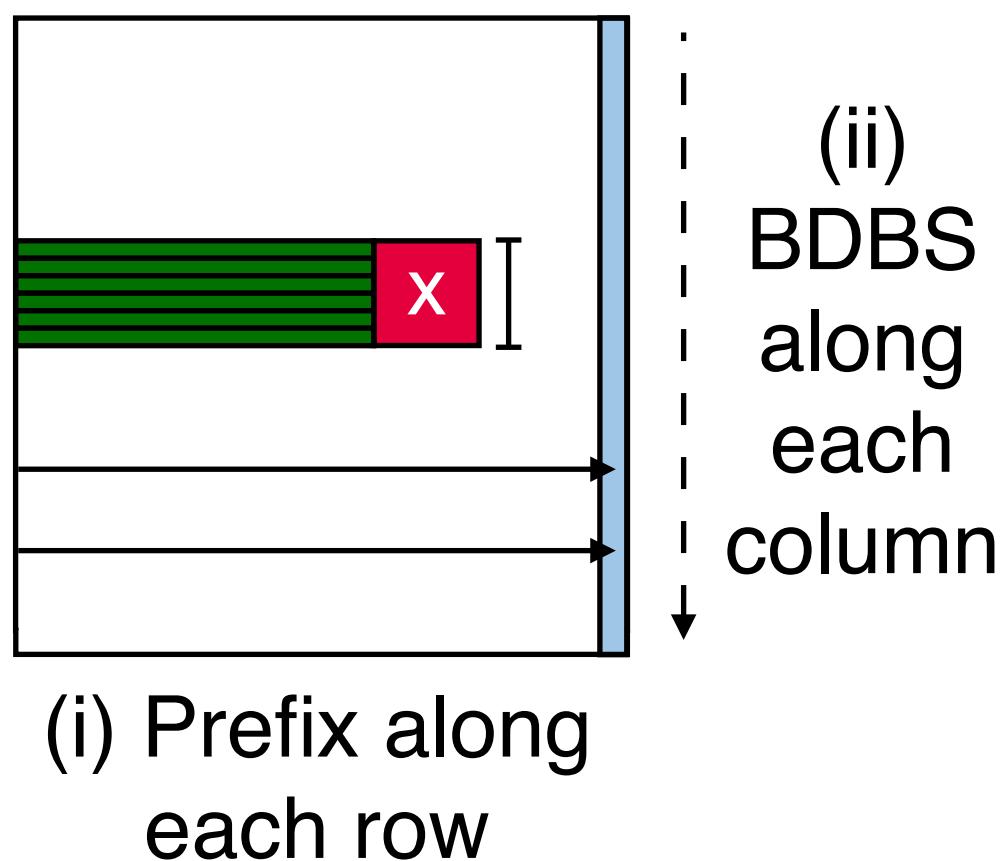
At a high level, the “ $i$ -complement” of a box such that there is some coordinate in dimension  $j \in [1, i]$  that is “out of range” in dimension  $j$ , and the coordinates are “in range” for all dimensions  $m \in [i + 1, d]$ .



# Box-Complement Algorithm for Strong Excluded Sums

The box-complement algorithm uses **dimension reduction** to compute the “ $i$ -complement” for all  $i = 1, \dots, d$ .

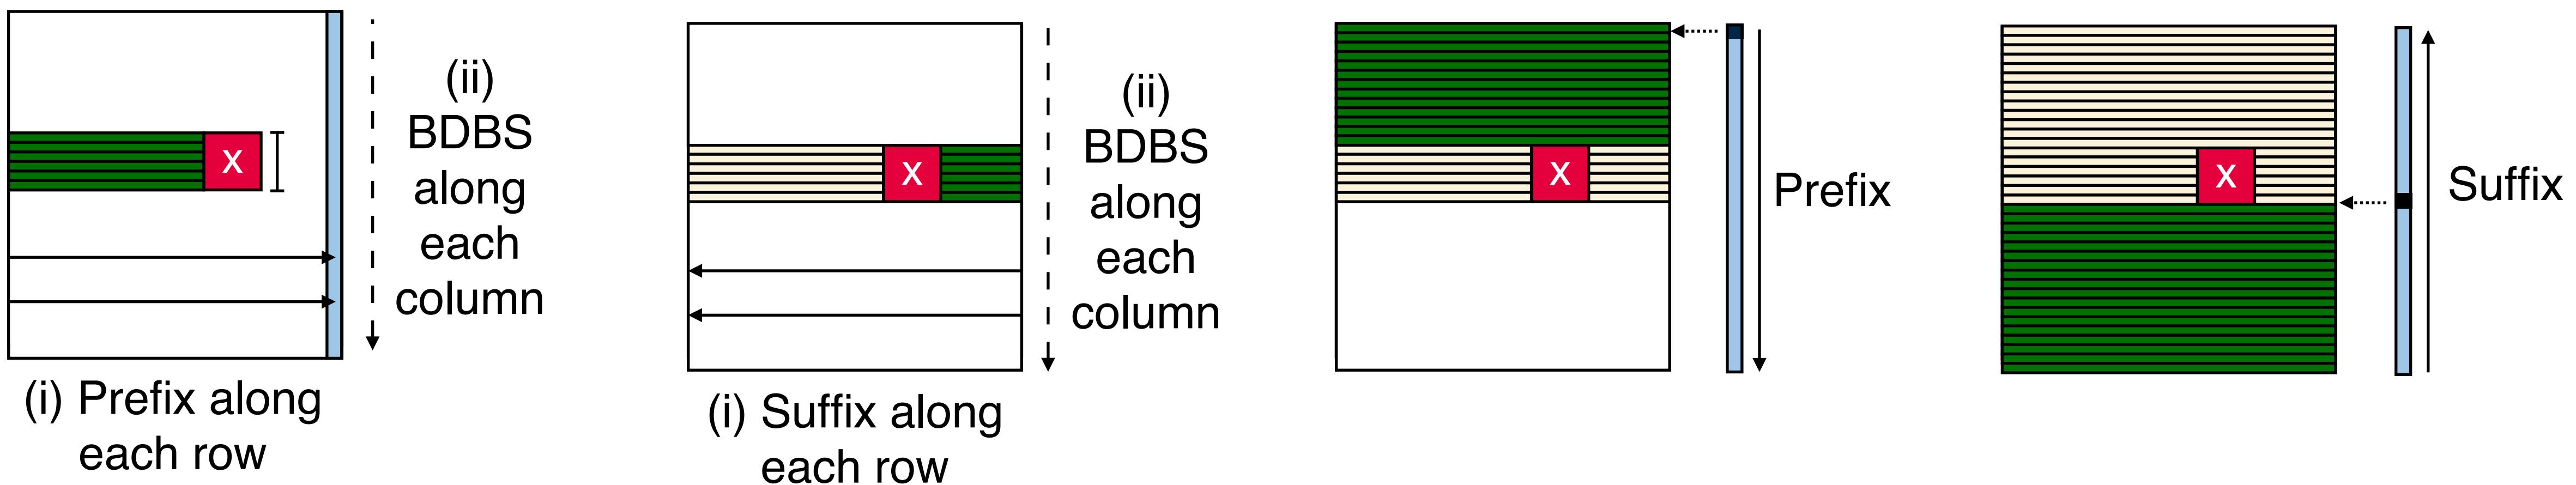
The **BDBS algorithm** for included sums is a major subroutine in the box-complement algorithm to sum up elements “in the range.”



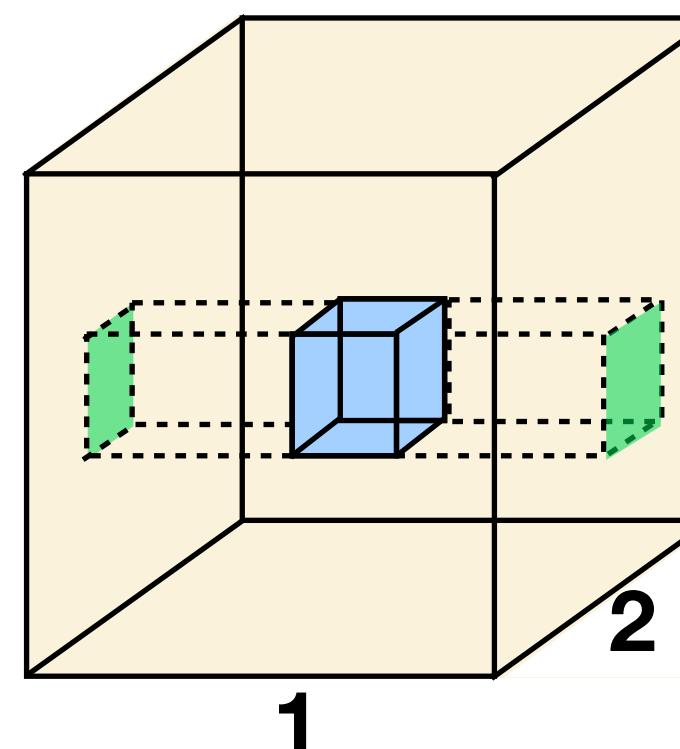
# Box-Complement Algorithm for Strong Excluded Sums

The box-complement algorithm uses **dimension reduction** to compute the “ $i$ -complement” for all  $i = 1, \dots, d$ .

The **BDBS algorithm** for included sums is a major subroutine in the box-complement algorithm to sum up elements “in the range.”

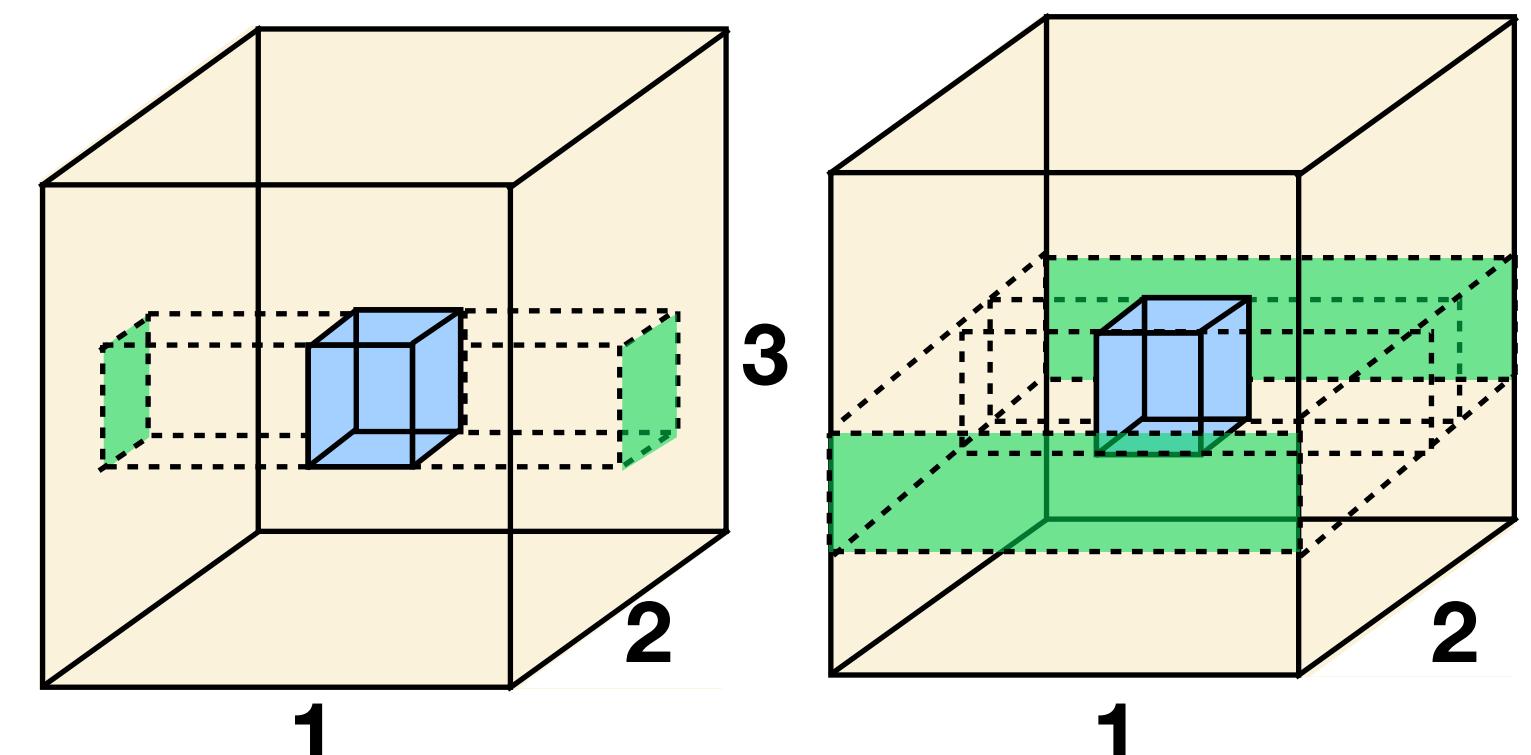


# Extending the Box-complement Algorithm to Higher Dimensions



**Full Prefix / Suffix  
Dimensions:** 1  
**Included Sum  
Dimensions:** 2, 3

# Extending the Box-complement Algorithm to Higher Dimensions

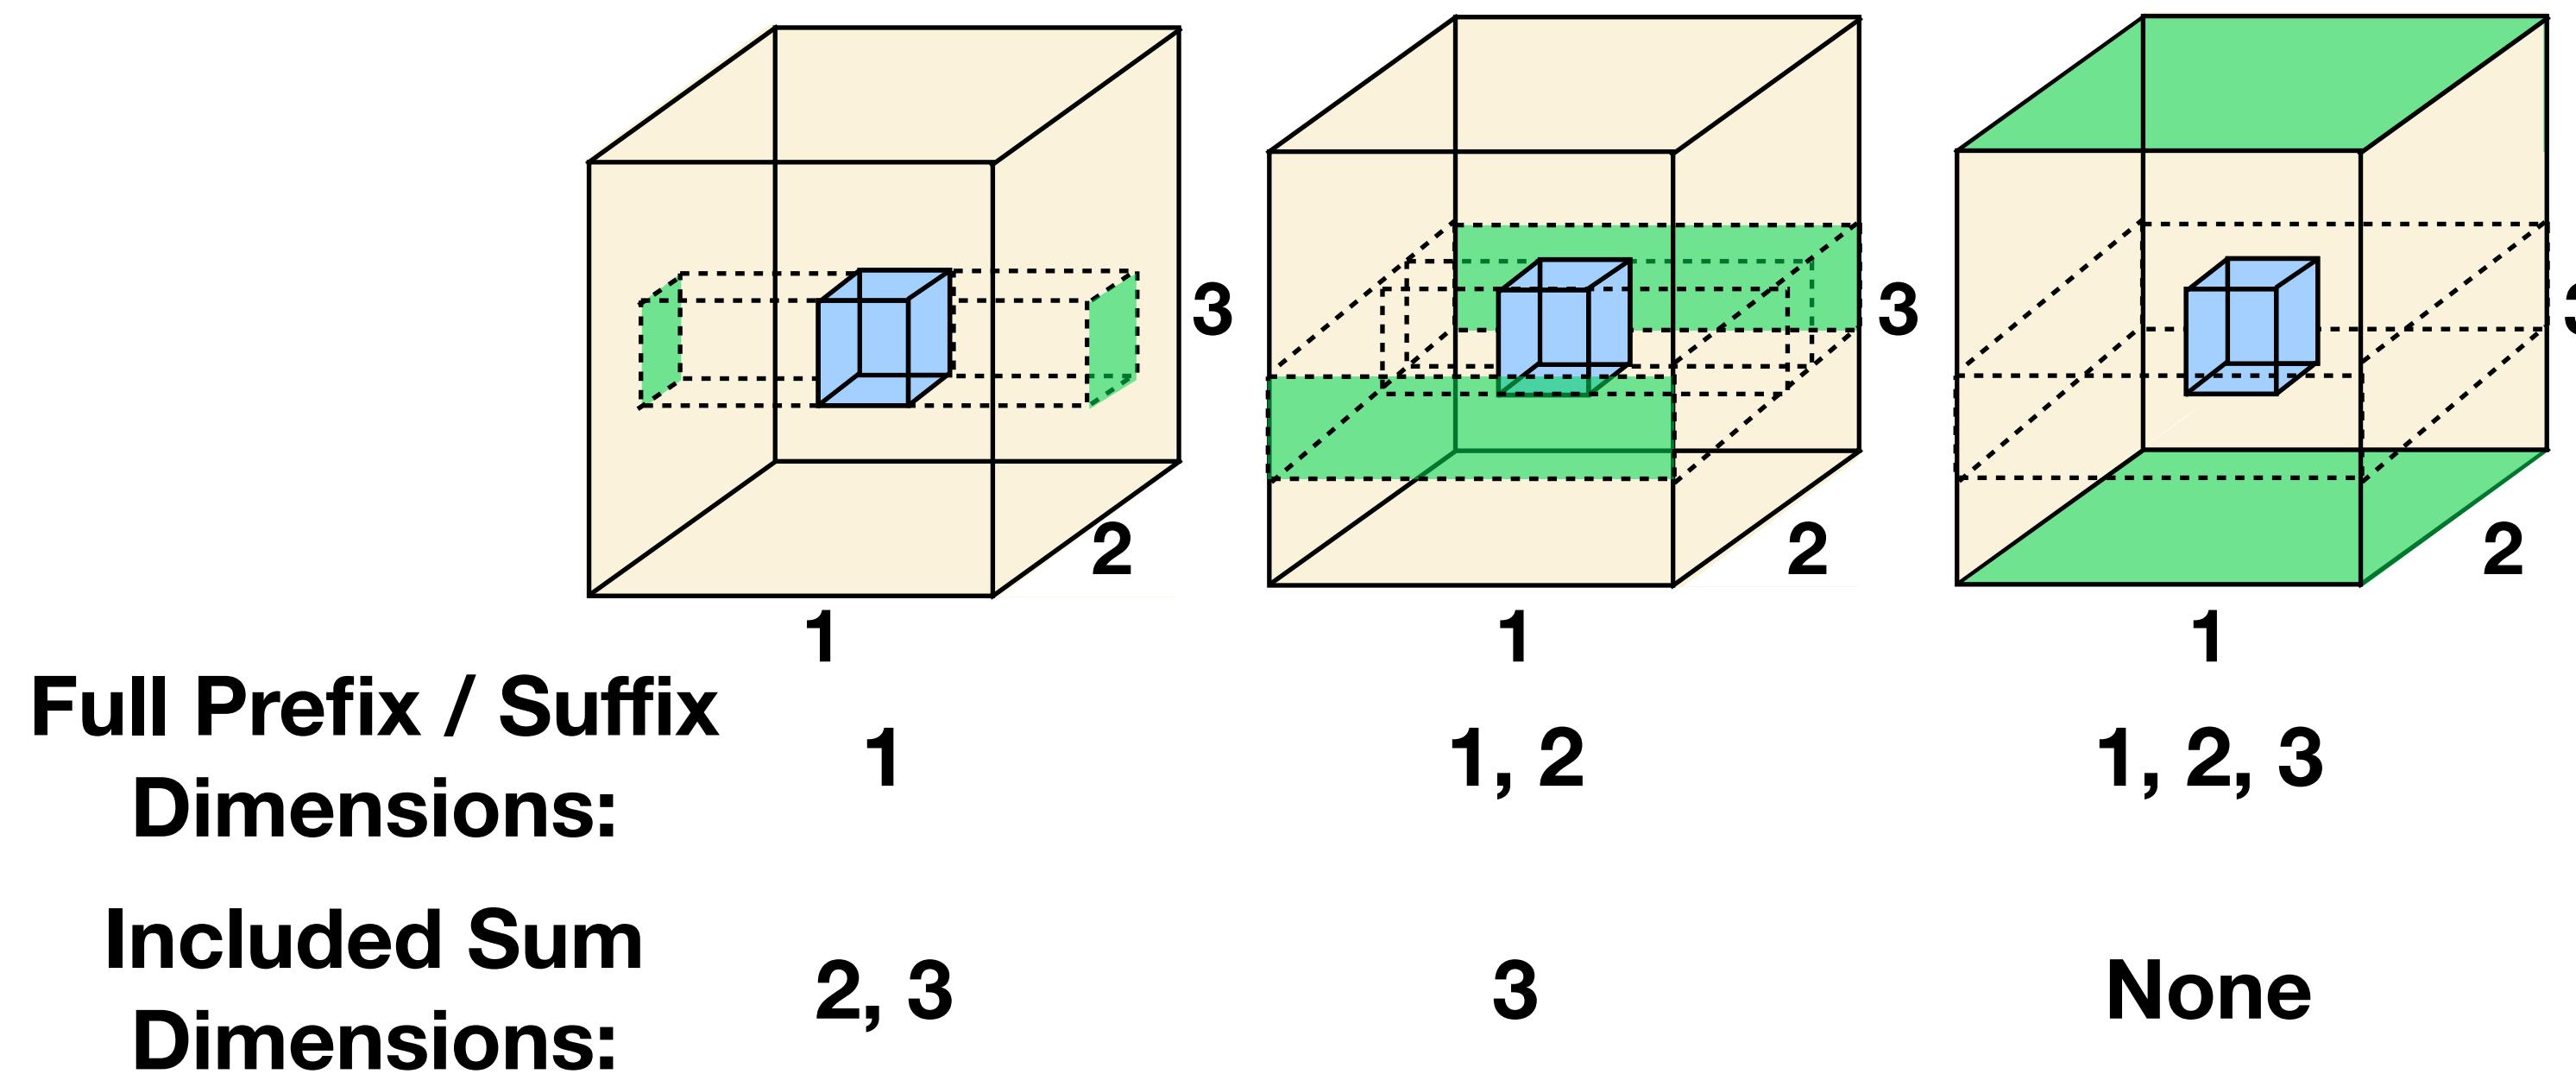


**Full Prefix / Suffix  
Dimensions:**

**Included Sum  
Dimensions:**      2, 3

3

# Extending the Box-complement Algorithm to Higher Dimensions



Each dimension-reduction step takes  $\Theta(N)$  time and reuses the same temporaries, for a total of  $\Theta(dN)$  time and  $\Theta(N)$  space.