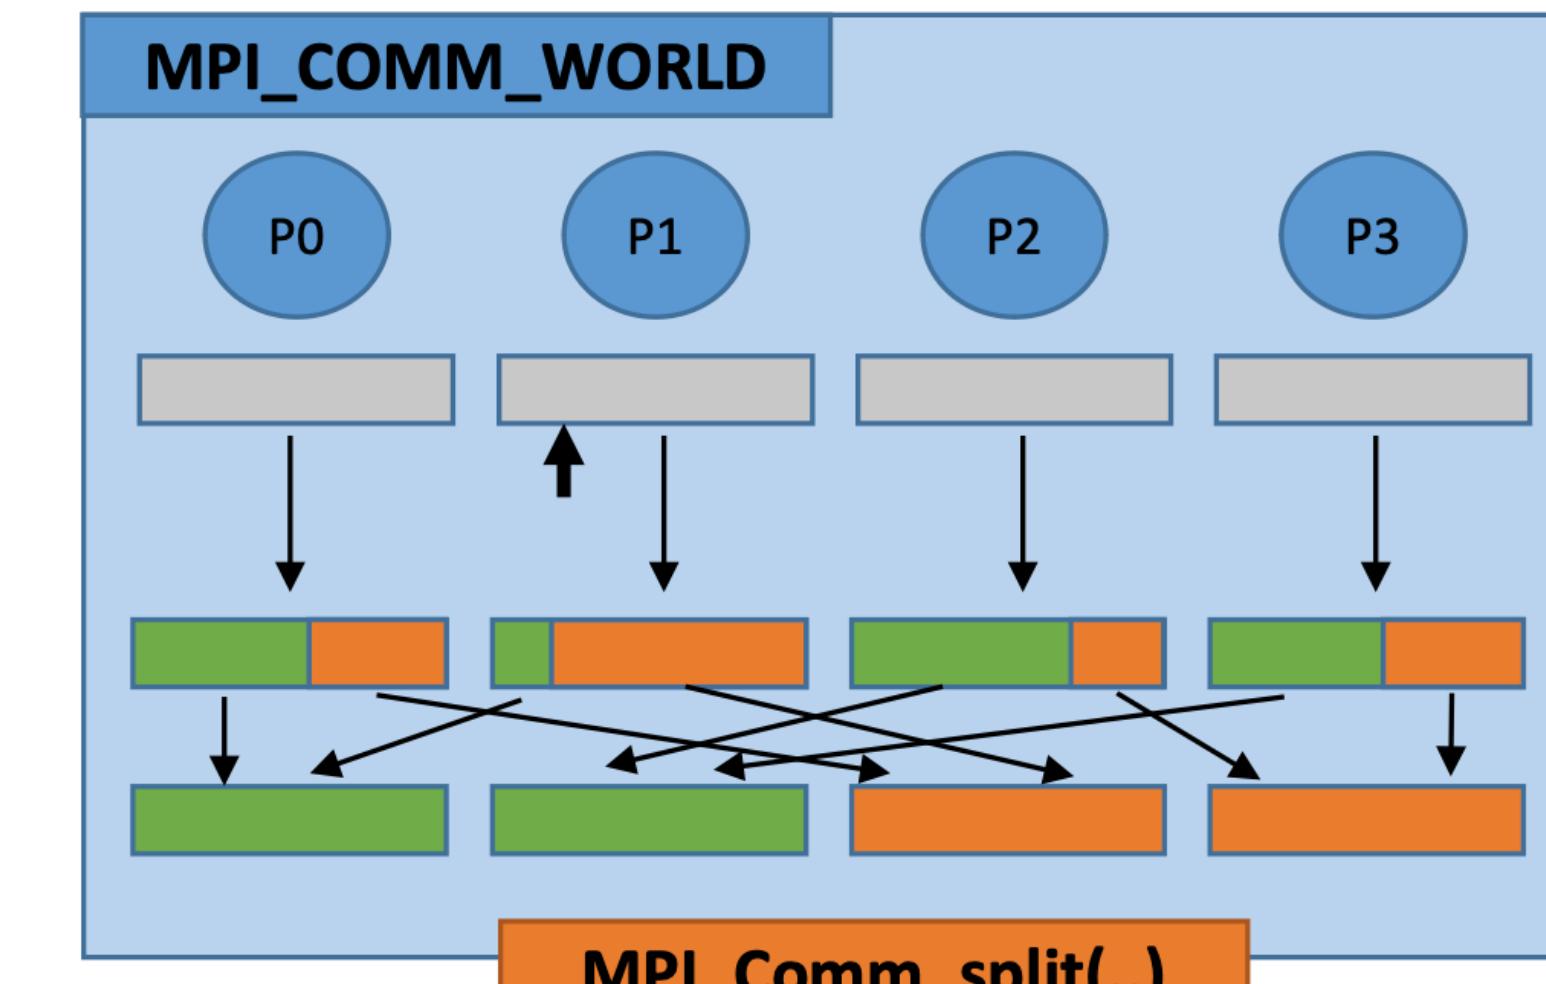


# Update: Parallel Quicksort

## Parallel Quicksort

1. Pick pivot and broadcast  $O(\log p(\tau + \mu))$
2. Partition locally  $O\left(\frac{n}{p}\right)$
3. Re-arrange partitions  $O\left(\tau + \mu \frac{n}{p}\right)$
4. Split processors and recurse
5. If  $p = 1$ , locally sort

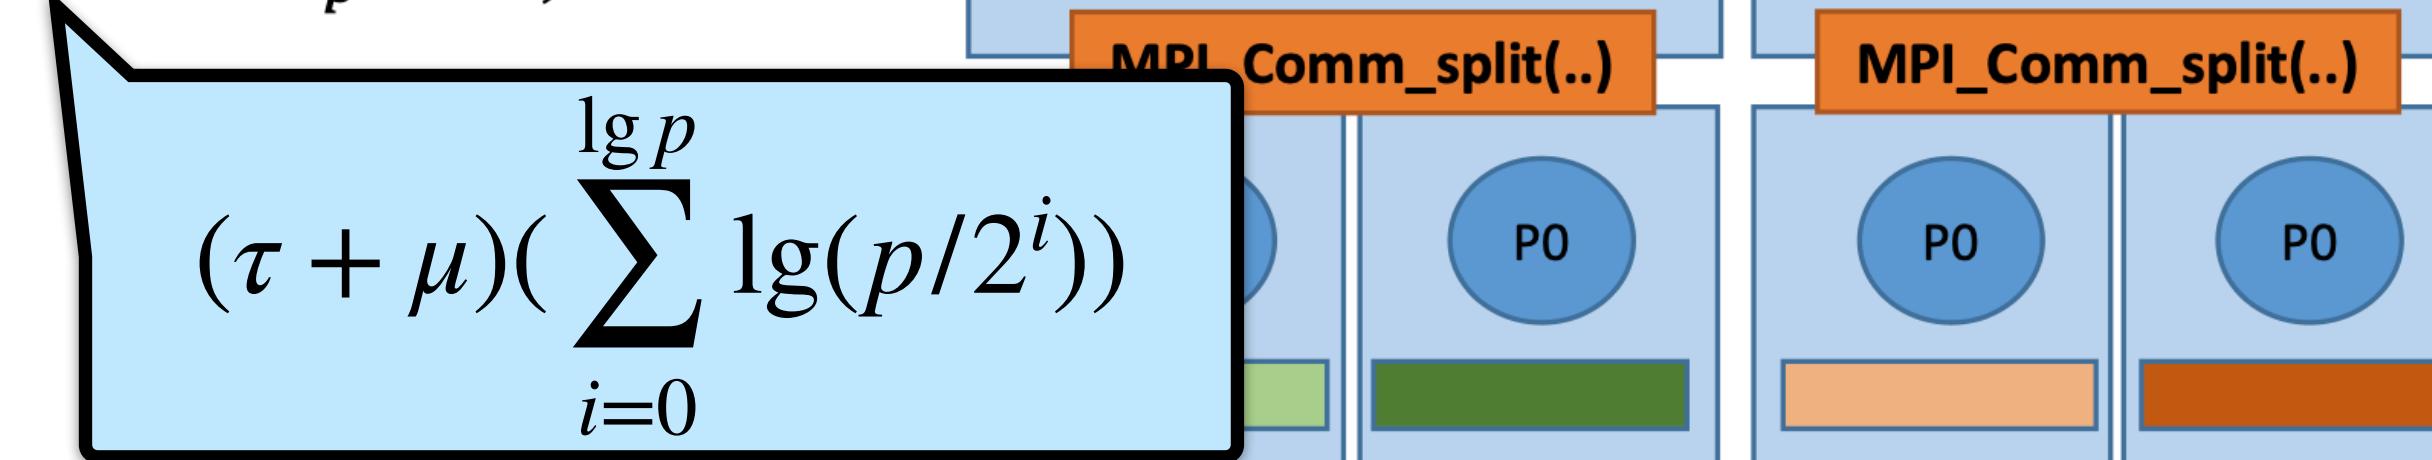


Best/expected case (assuming no imbalance):

- $O(\log p)$  iterations before local sort
- Computation:  $O\left(\frac{n}{p} \log p + \frac{n}{p} \log \frac{n}{p}\right) = O\left(\frac{n \log n}{p}\right)$
- Communication:  $O\left(\tau \lg^2 p + \mu \frac{n}{p} \log p\right)$

Worst case:

- Much worse!



$$(\tau + \mu) \left( \sum_{i=0}^{\lg p} \lg(p/2^i) \right)$$

# CSE 6220/CX 4220

## Introduction to HPC

# Lecture 14: Dense Matrix Algorithms

Helen Xu  
[hxu615@gatech.edu](mailto:hxu615@gatech.edu)

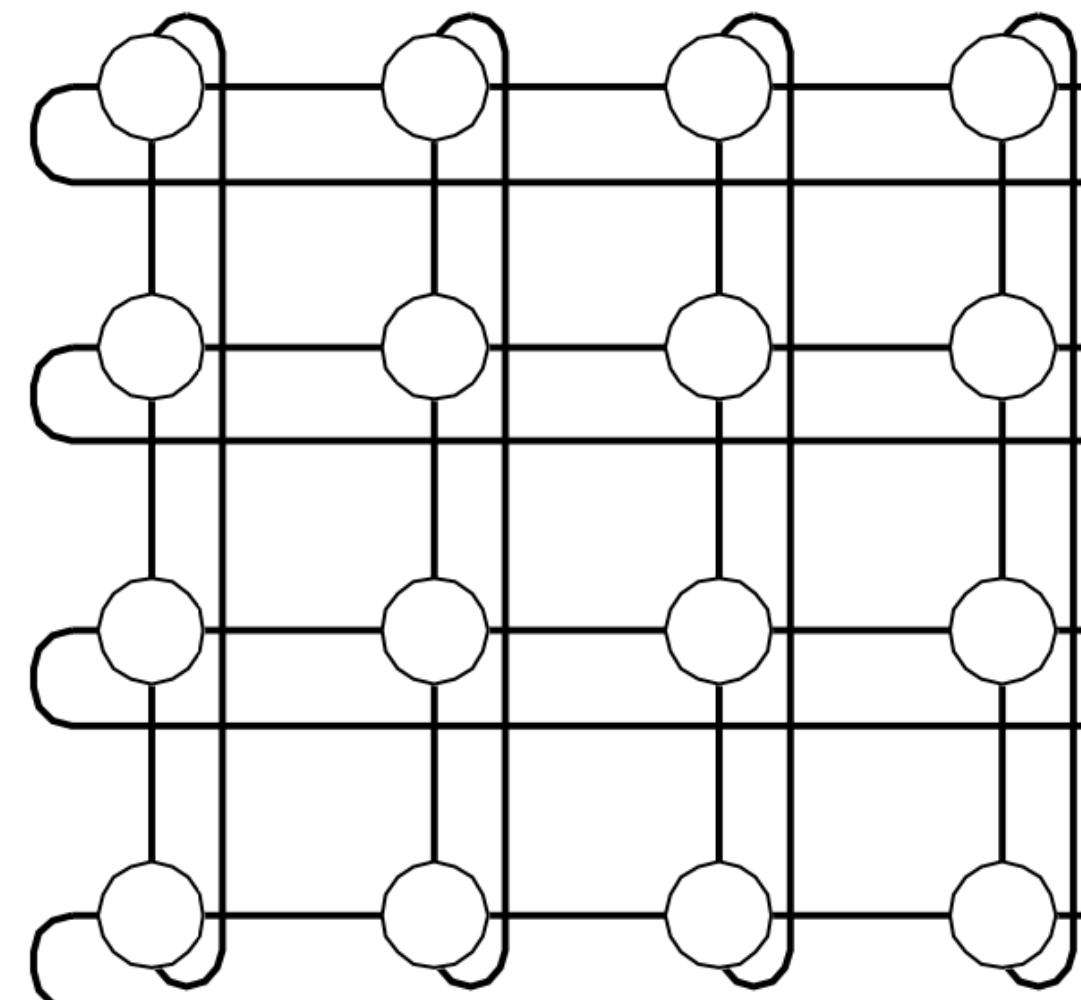


# Recap: Communication Primitives

- Broadcast
  - Reduce
  - AllReduce
  - Scan
  - Gather
  - AllGather
  - Scatter
  - All\_to\_All
    - Arbitrary Permutations:  $O(\tau \cdot p + \mu \cdot m \cdot p)$
    - Hypercubic Permutations:  $O(\tau \cdot \log p + \mu \cdot m \cdot p \cdot \log p)$
- 
- $$\Theta(\tau \log p + \mu m \log p)$$
- $$\Theta(\tau \log p + \mu m p)$$

# Matrix Algorithms

- Regular structure
- Well-suited to **static partitioning** of computation
- Computation typically partitioned based on the data
  - input, output, intermediate
- Typical data partitioning strategies for matrix algorithms
  - 1D and 2D block, cyclic, block-cyclic
- Easier to think logically about the algorithms on a 2D mesh of processors:



# **Matrix-Vector Multiplication**

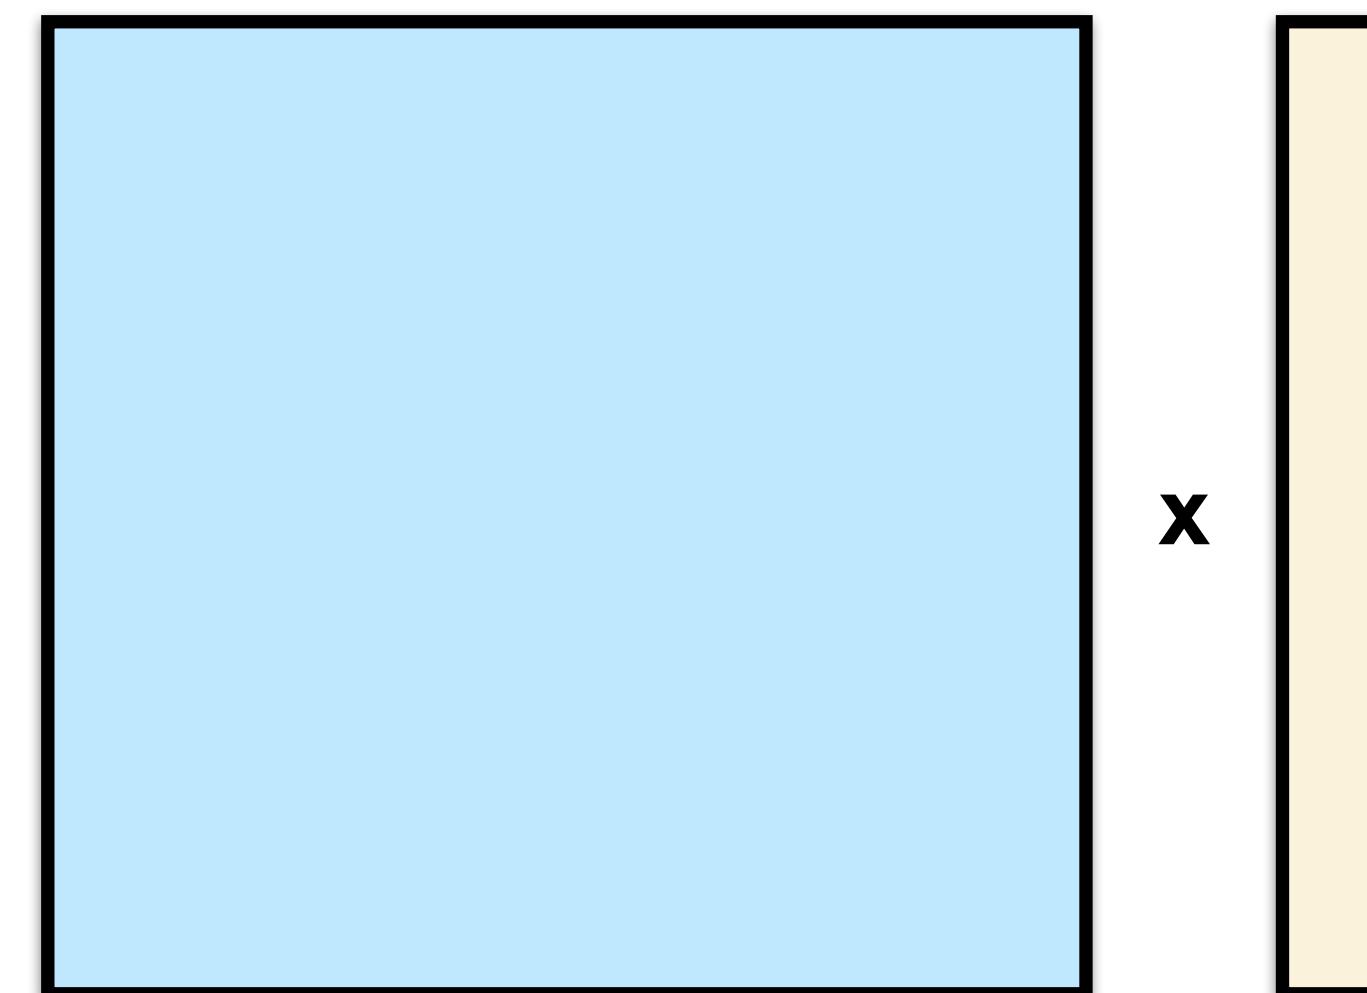
# Matrix-Vector Multiplication

Problem:

- multiply a dense  $n \times n$  matrix  $A$  with an  $n \times 1$  vector  $x$
- yield the  $n \times 1$  result vector  $y$

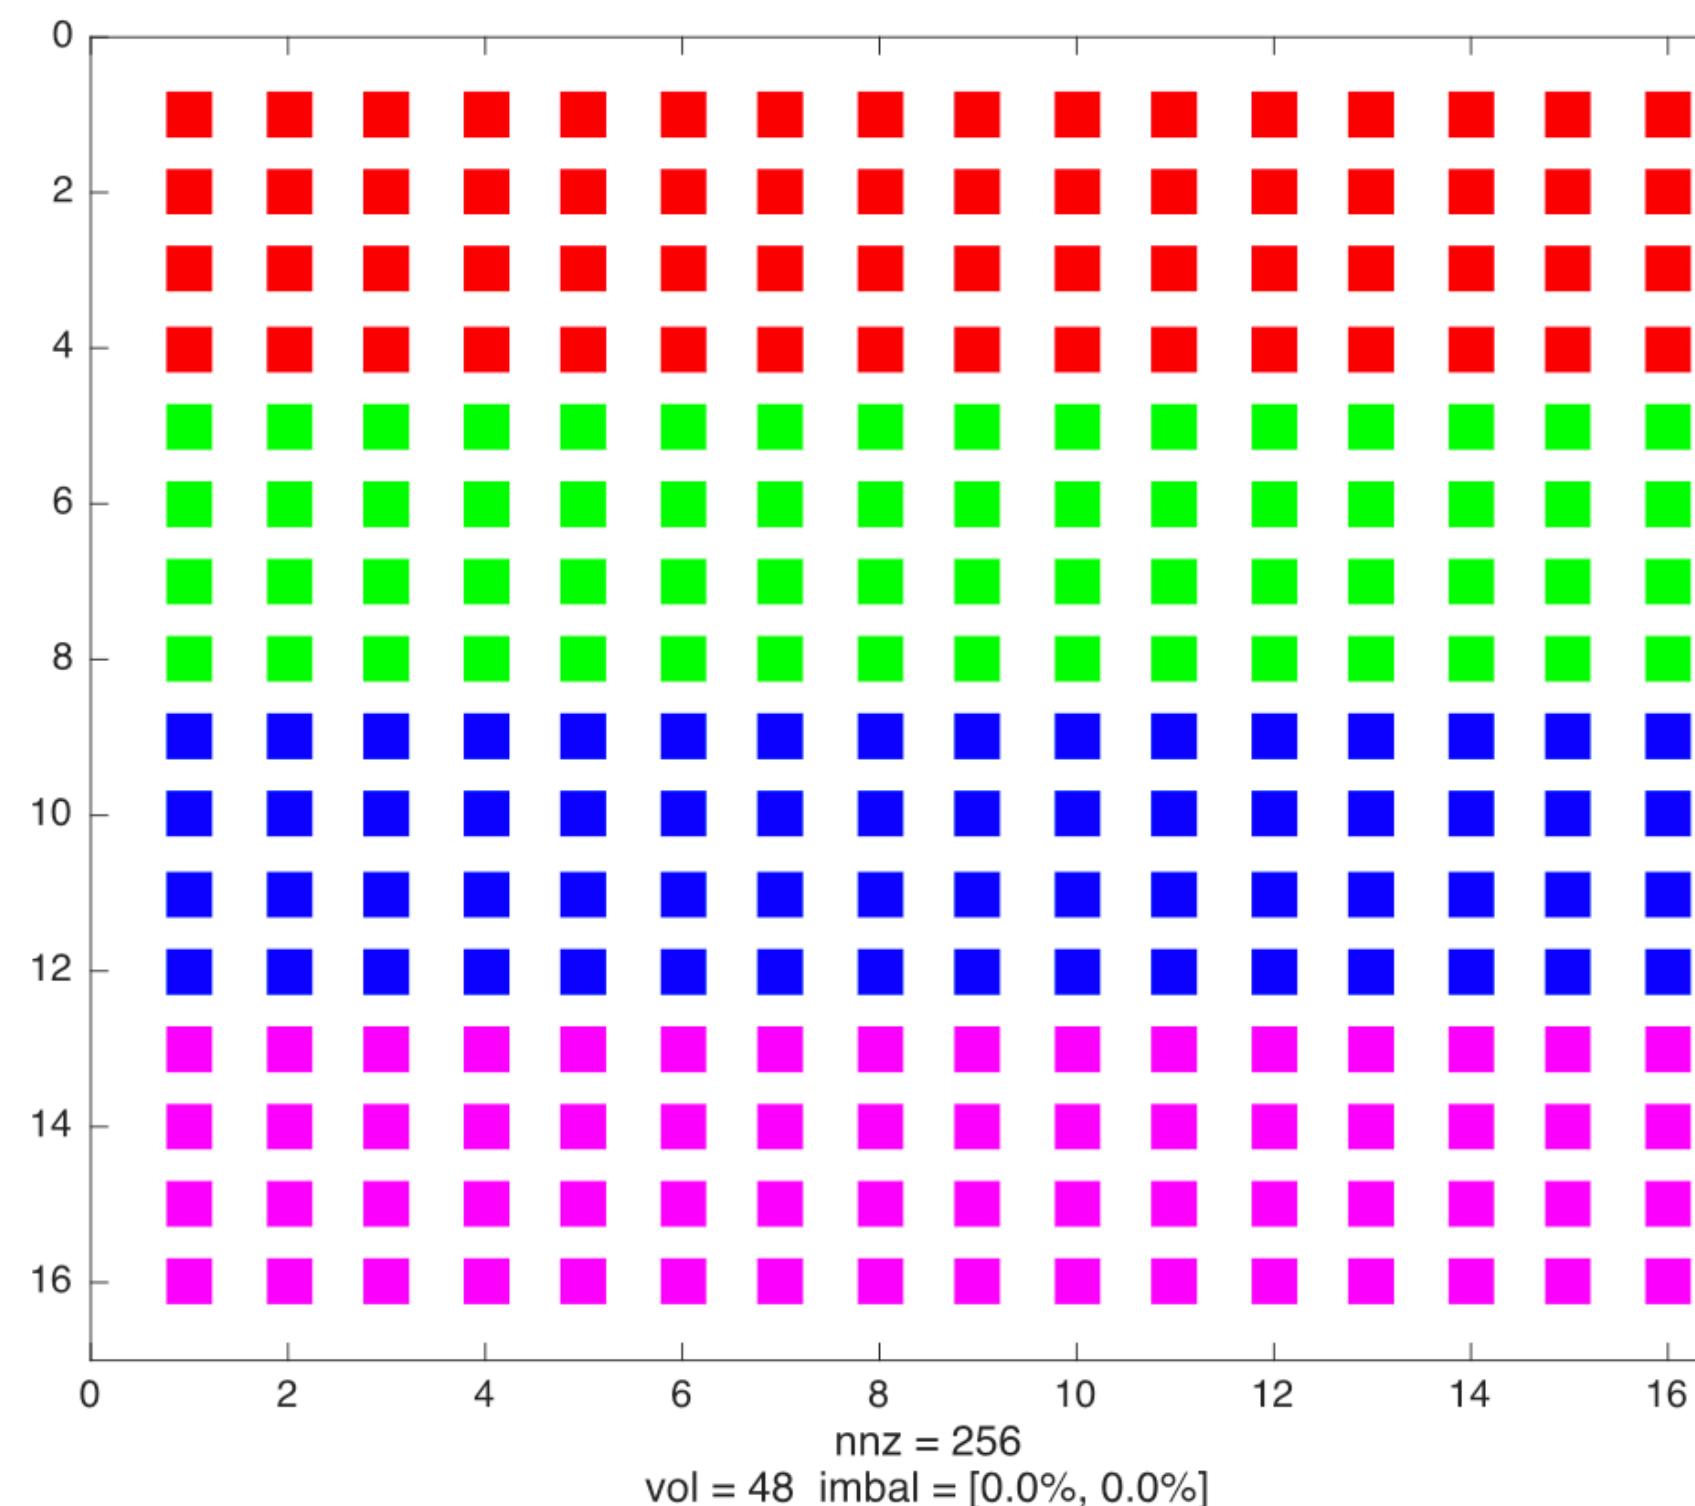
What is the serial time complexity?

- Serial algorithm requires  $n^2$  multiplications and additions



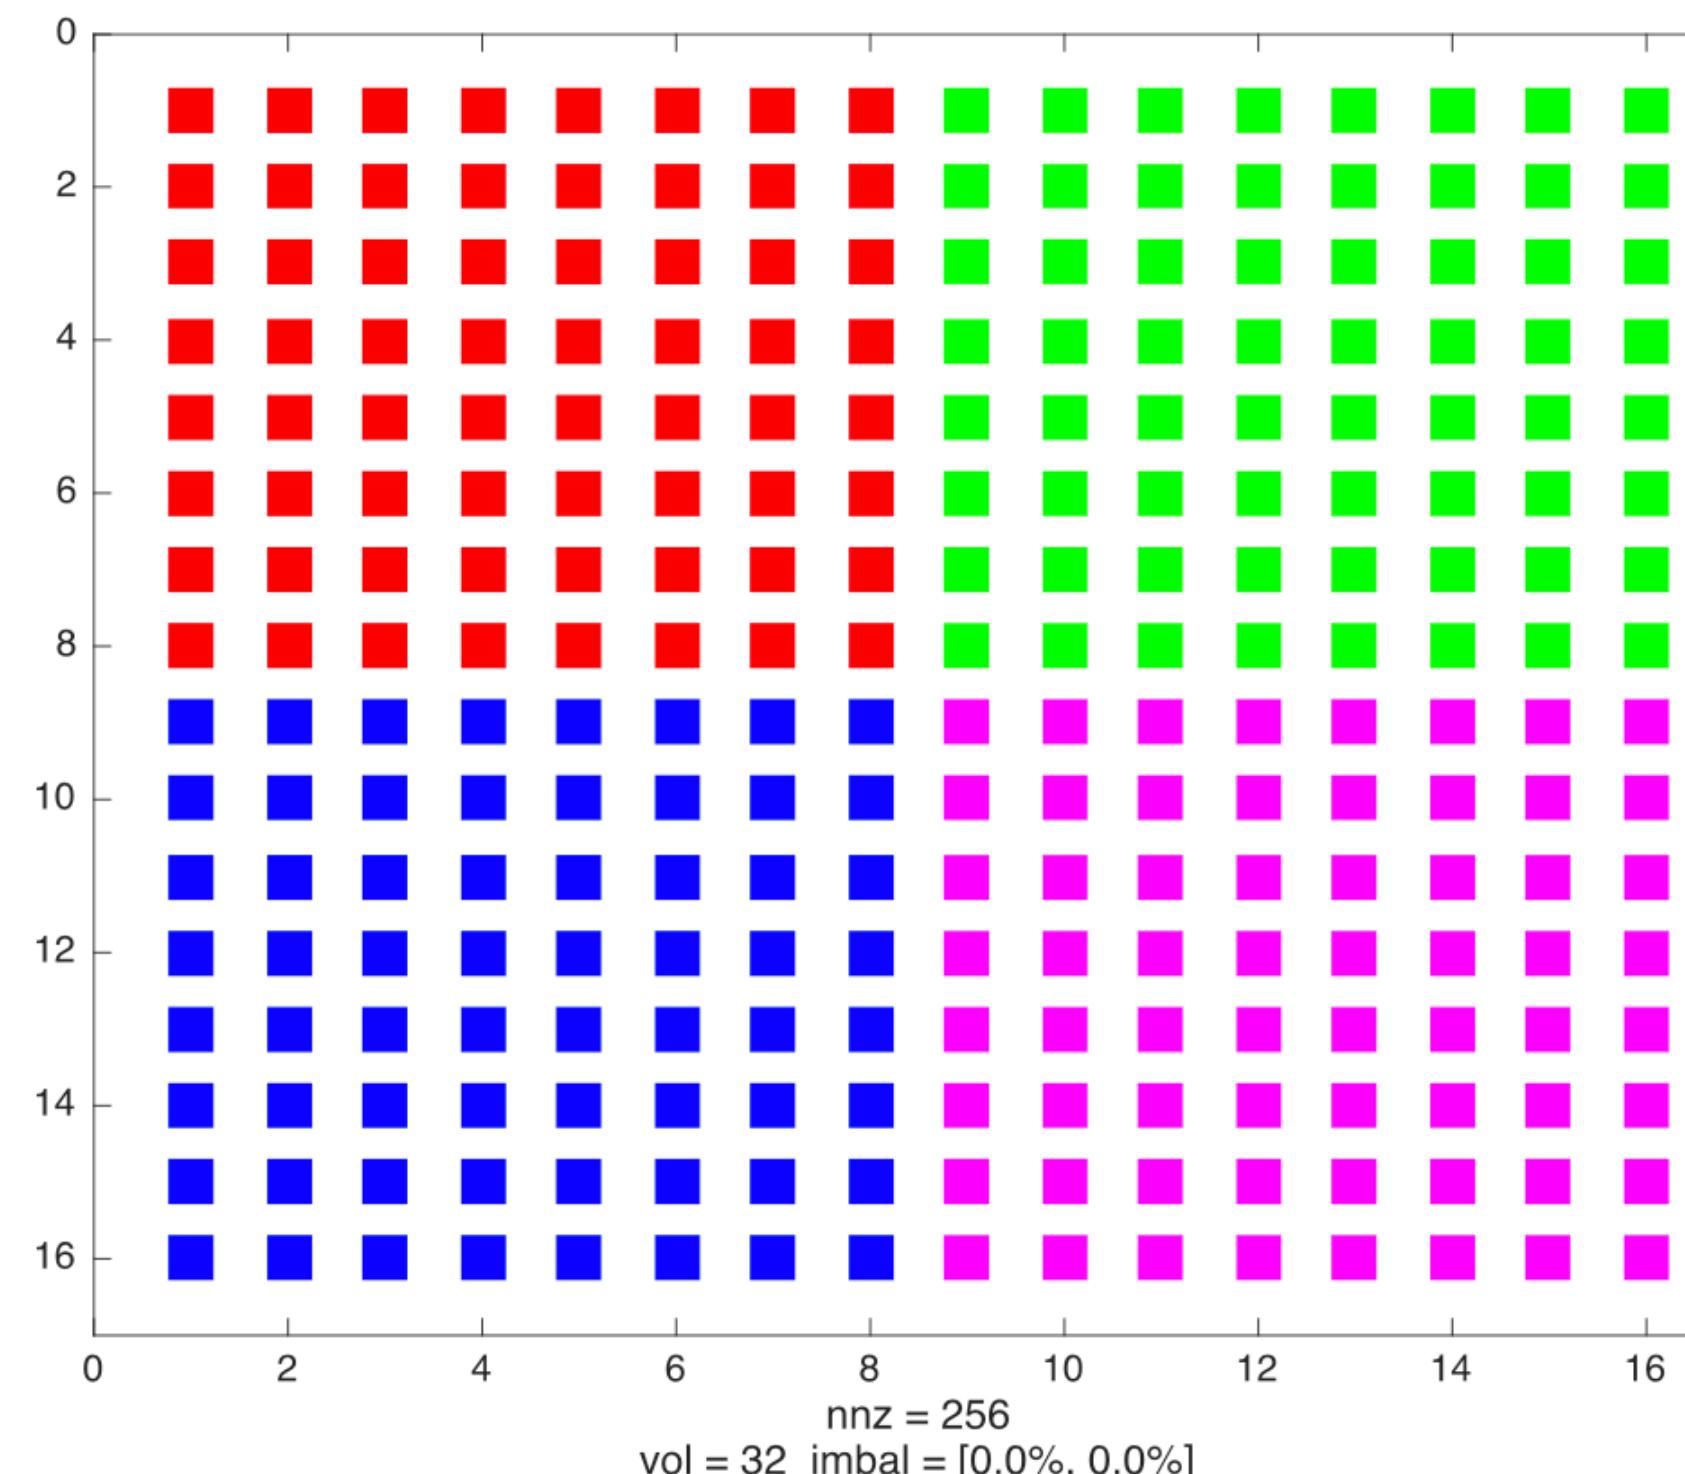
# 1D Partitioning for Matrix-Vector Multiply

- 1D partitioning splits up the matrix into **block rows** with at least one row on each processor.
- Given  $n$  rows, 1D partitioning can use at most  $n$  processors.



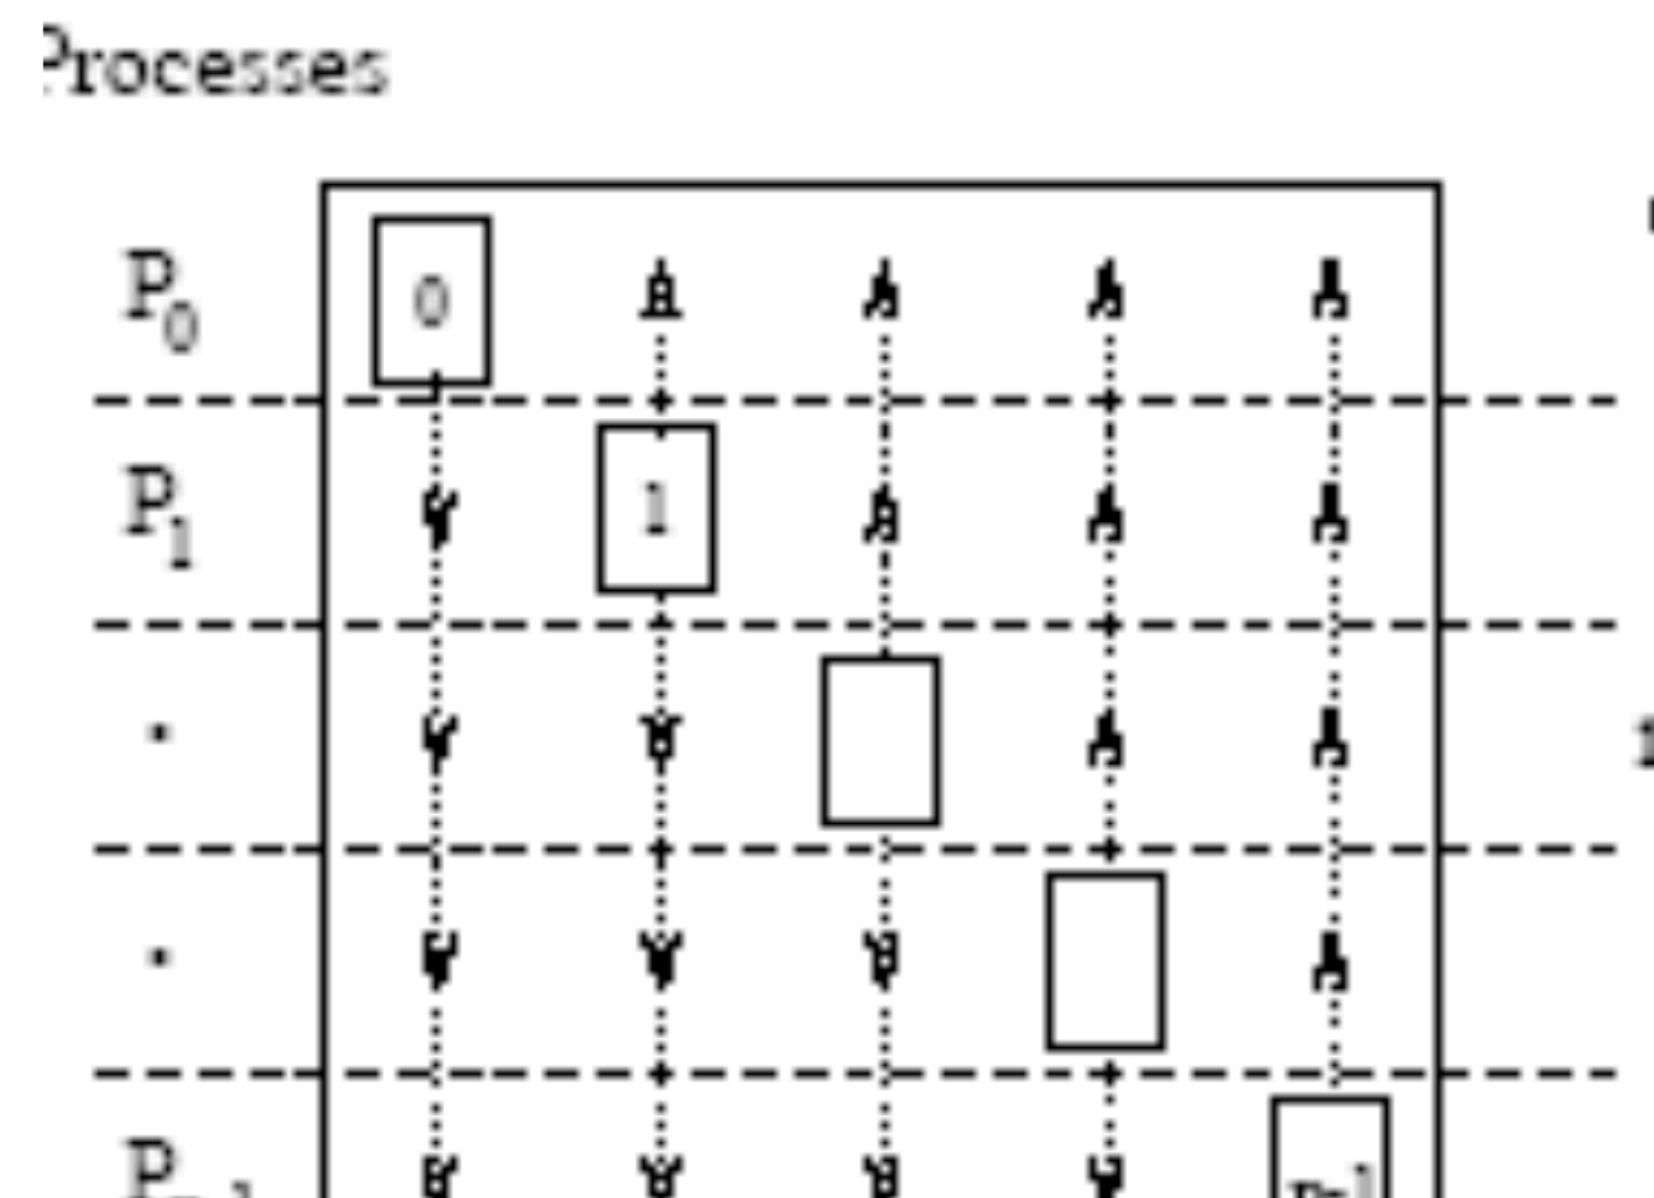
# 2D Partitioning for Matrix-Vector Multiply

- 2D partitioning splits up the matrix into **square sub-matrices** - at least 1 element per processor
- Theoretically, 2D partitioning can use up to  $n^2$  processors (equal to number of elements)



# Analysis of 1D Partitioning

- First, consider simple case:  $p = n$
- $n \times n$  matrix is partitioned among  $p$  processors
  - *Each processor stores a complete row of the matrix*
- $n \times 1$  vector  $x$  is also partitioned in a conformal manner
  - Each process owns one element of  $x$



# Analysis of 1D Partitioning

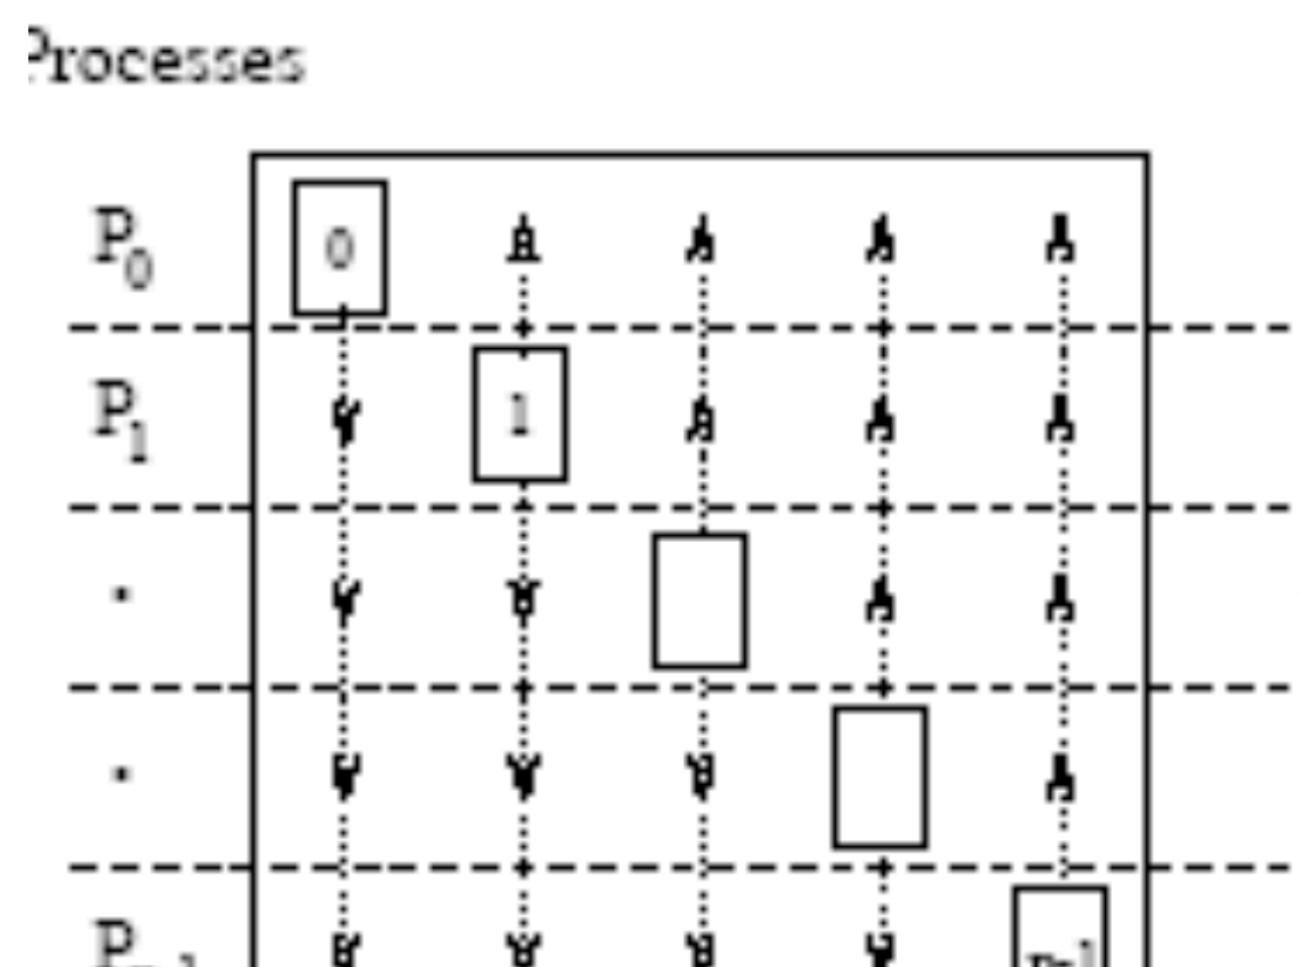
- Step 1: Use AllGather to distribute all of  $x$  to each processor

$$O(\tau \log n + \mu \cdot n)$$

$\mu n > n$ , so time is basically  
the cost of allgather

- Step 2: Processor  $P_i$  computes  $y[i] = \sum_{j=0}^{n-1} A[i, j] \cdot x[j]$

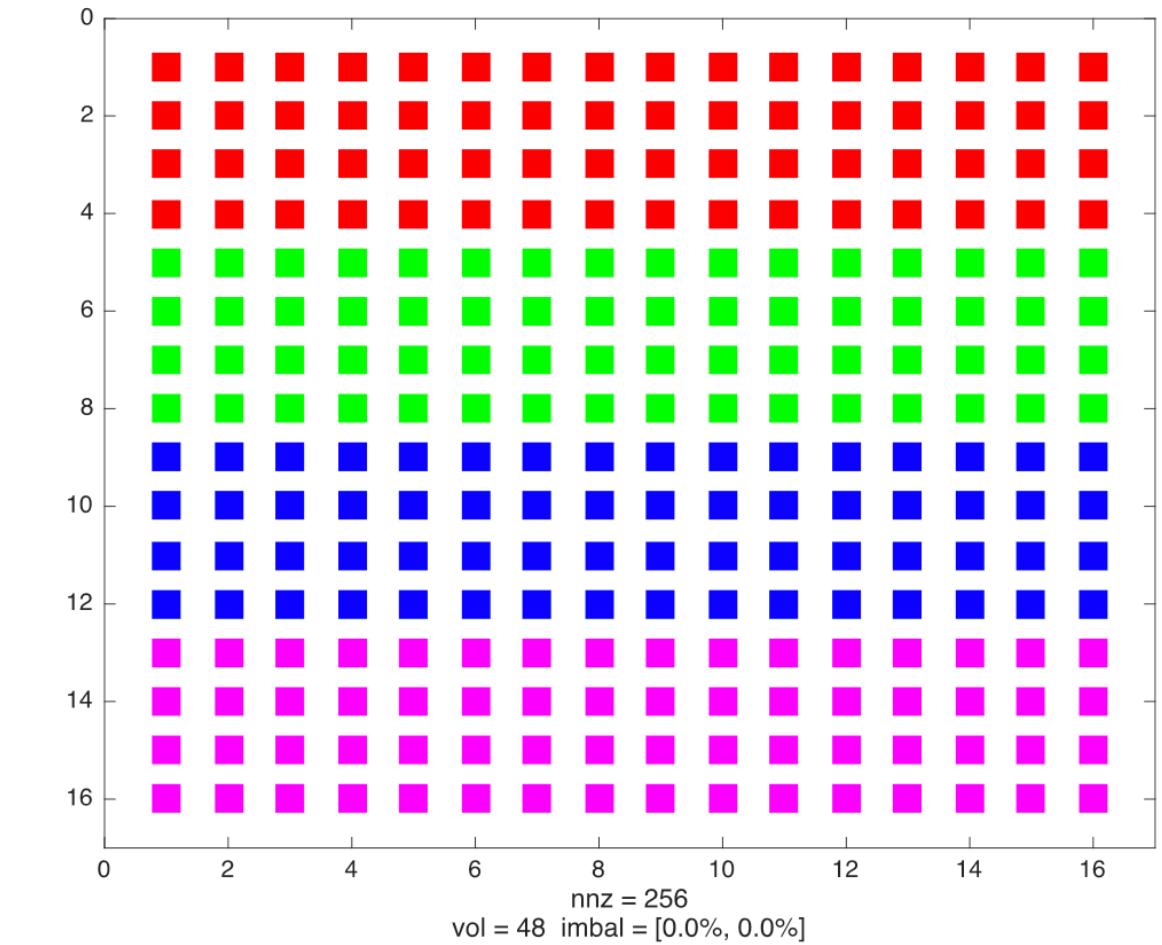
$$O(n)$$



# Analysis of 1D Partitioning

For  $n > p$

- Initially, each processor stores
  - $n/p$  complete rows of the matrix  $A$
  - $n/p$  elements of the vector  $x$
- Distribute all of  $x$  vector to each processor
  - AllGather among  $p$  processors, messages of size  $n/p$
- On each processor:  $n/p$  local dot products on vectors (rows) of length  $n$

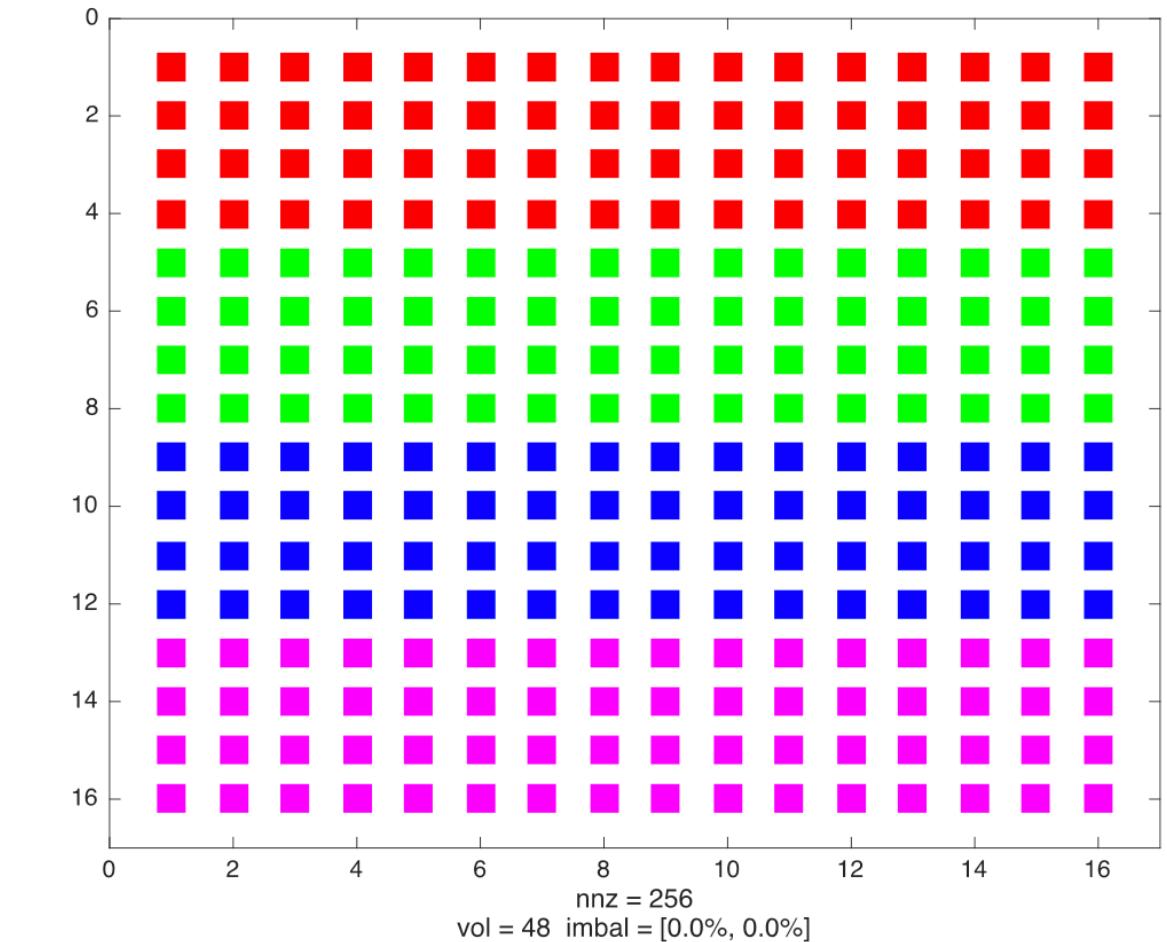


- Step 1: Comm. Time = **What is the communication time?**
- Step 2: Computation time = **What is the computation time?**
- Efficient up to  $O(n)$  processors, the maximum the alg. can use.

# Analysis of 1D Partitioning

For  $n > p$

- Initially, each processor stores
  - $n/p$  complete rows of the matrix  $A$
  - $n/p$  elements of the vector  $x$
- Distribute all of  $x$  vector to each processor
  - AllGather among  $p$  processors, messages of size  $n/p$
- On each processor:  $n/p$  local dot products on vectors (rows) of length  $n$

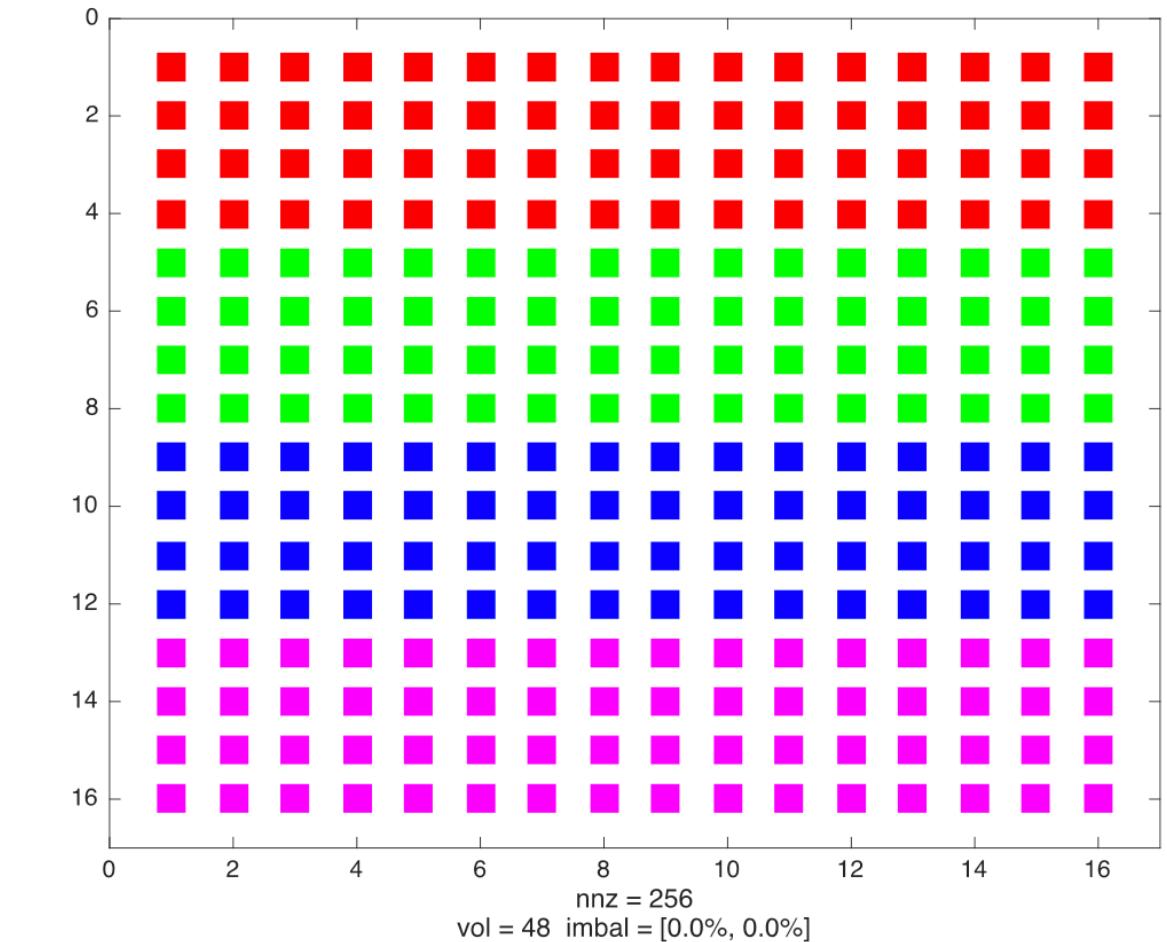


- Step 1: Comm. Time =  $O\left(\tau \log p + \mu \cdot \frac{n}{p} \cdot p\right)$  =  $O(\tau \log p + \mu \cdot n)$
- Step 2: Computation time = What is the computation time?
- Efficient up to  $O(n)$  processors, the maximum the alg. can use.

# Analysis of 1D Partitioning

For  $n > p$

- Initially, each processor stores
  - $n/p$  complete rows of the matrix  $A$
  - $n/p$  elements of the vector  $x$
- Distribute all of  $x$  vector to each processor
  - AllGather among  $p$  processors, messages of size  $n/p$
- On each processor:  $n/p$  local dot products on vectors (rows) of length  $n$



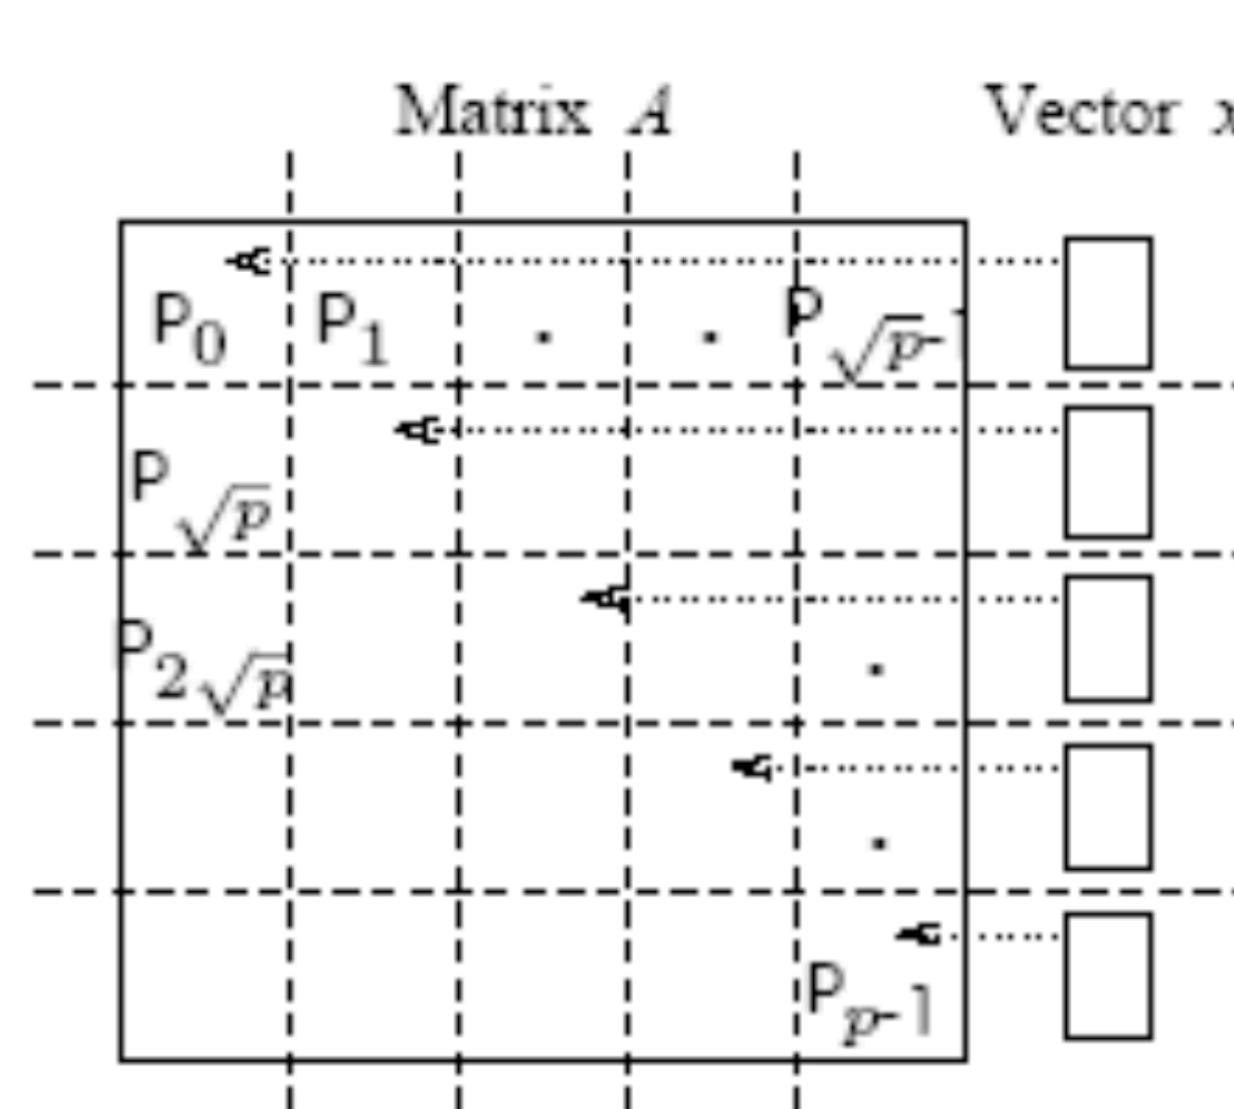
- Step 1: Comm. Time =  $O\left(\tau \log p + \mu \cdot \frac{n}{p} \cdot p\right) = O(\tau \log p + \mu \cdot n)$
- Step 2: Computation time =  $O\left(\frac{n^2}{p}\right)$
- Efficient up to  $O(n)$  processors, the n

Can only use up to  $n/\mu$  processors because computation time needs to dominate for efficiency

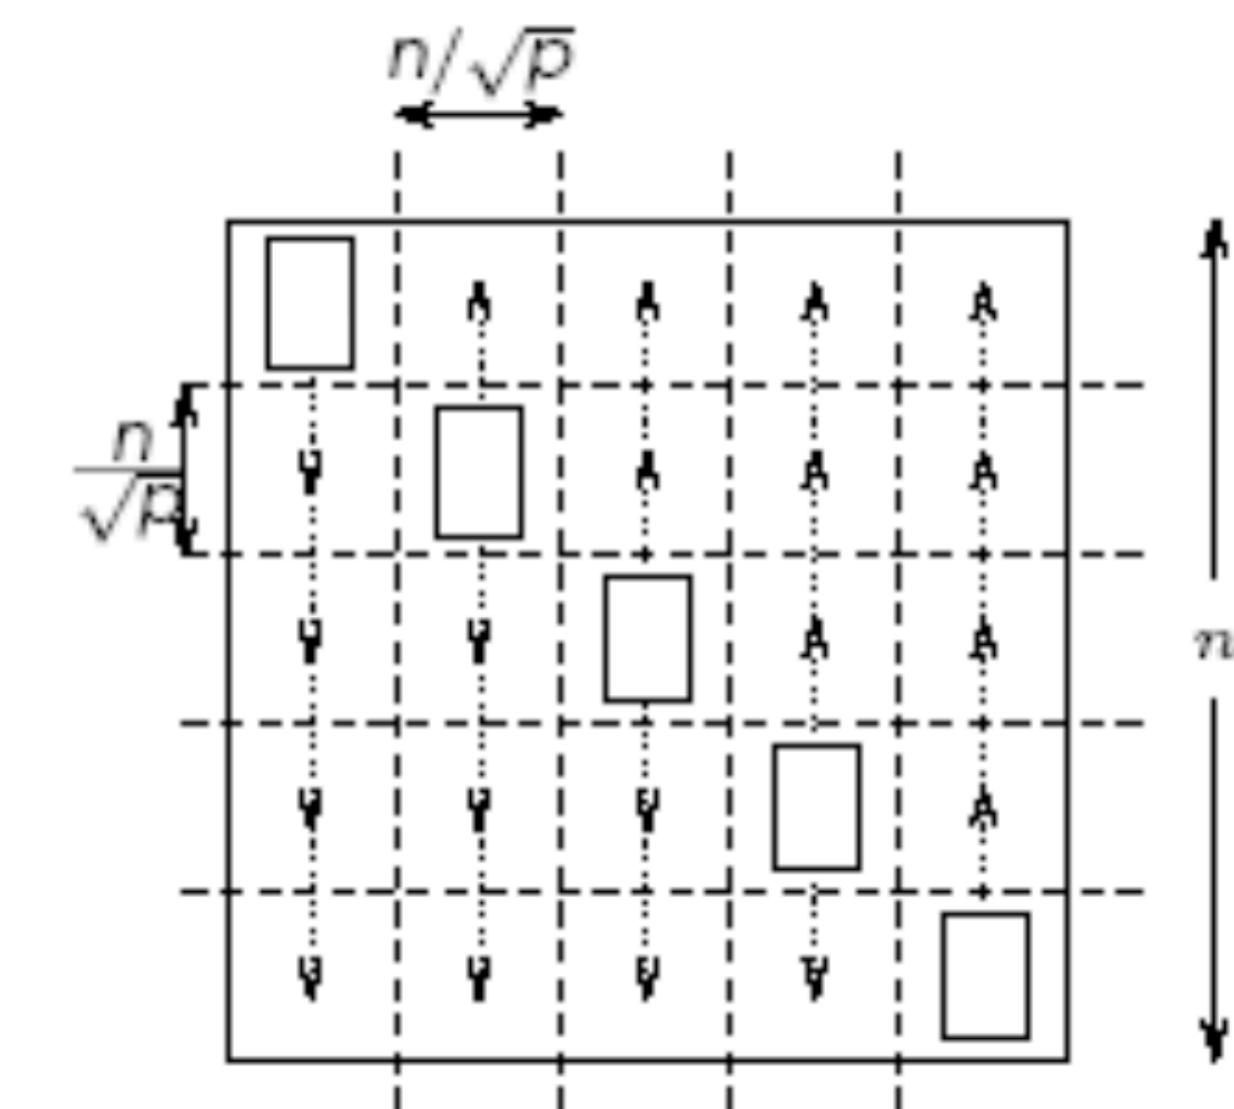
# Analysis of 2D Partitioning

- Consider :  $p = n^2$
- $n \times n$  matrix is partitioned among  $p = n^2$  processors (virtual grid topology)
  - *Each processor owns a single element*
- $n \times 1$  vector  $x$  is distributed in the last column of  $n$  processors

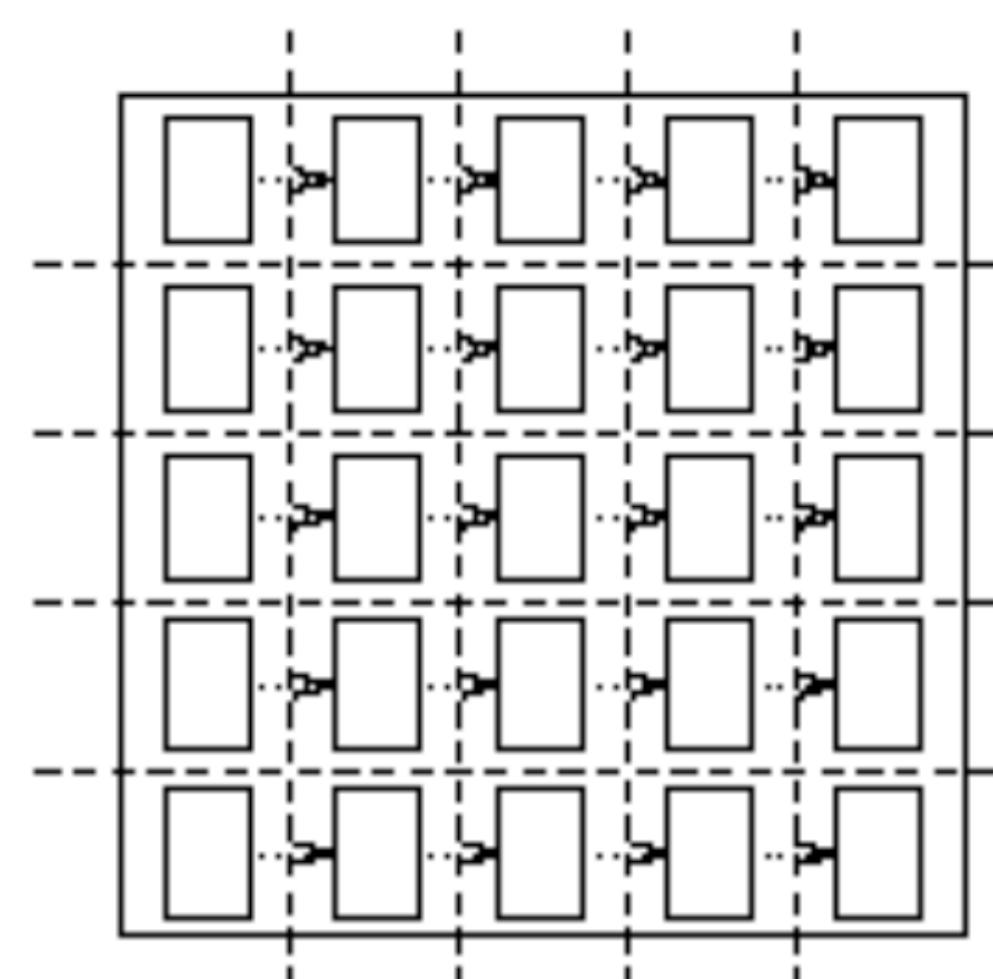
What is a good algorithm given this distribution?



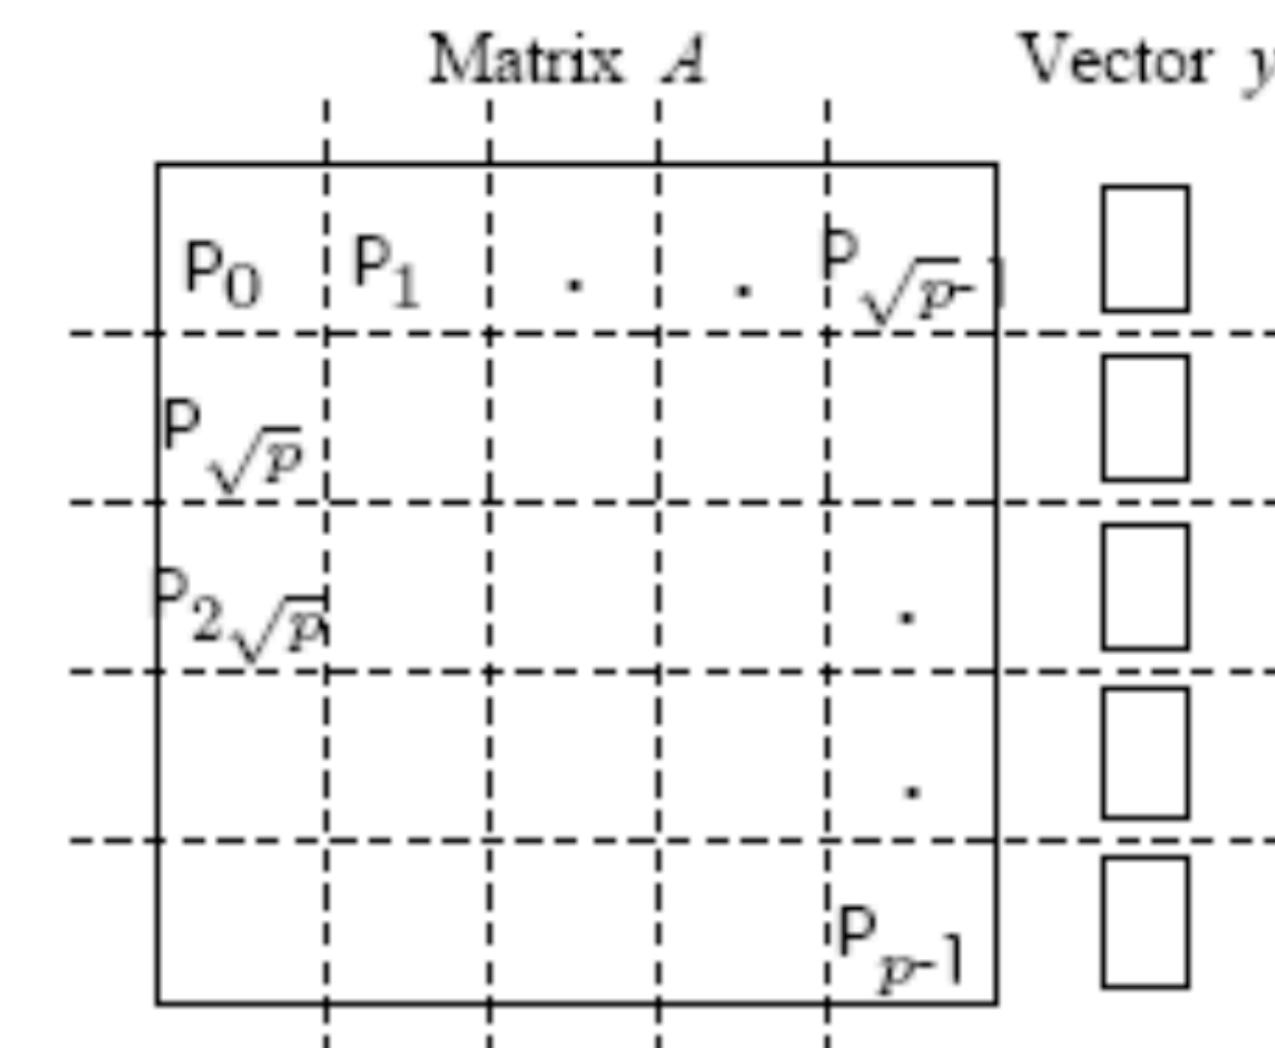
(a) Initial data distribution and communication steps to align the vector along the diagonal



(b) One-to-all broadcast of portions of the vector along process columns



(c) All-to-one reduction of partial results



(d) Final distribution of the result vector

# Analysis of 2D Partitioning

$$p = n^2$$

- Algorithm

- align vector along the main diagonal: one-to-one communication  $(\tau + \mu) \lg n$
- Broadcast vector element to  $n$  processors in column: one-to-all broadcast  $(\tau + \mu) \lg n$
- Local multiplication 1
- Sum partial  $y$  values in each row: all-to-one reduction  $(\tau + \mu) \lg n + \lg n$

# Analysis of 2D Partitioning

$$p = n^2$$

- Algorithm

- align vector along the main diagonal: one-to-one communication  $(\tau + \mu) \lg n$
- Broadcast vector element to  $n$  processors in column: one-to-all broadcast  $(\tau + \mu) \lg n$
- Local multiplication 1
- Sum partial  $y$  values in each row: all-to-one reduction  $(\tau + \mu) \lg n + \lg n$

- Communication time =  $O(\tau \log n + \mu \cdot \log n)$

- Computation time =  $O(\log n)$

From reduction

Is the algorithm efficient?

# Analysis of 2D Partitioning

$$p = n^2$$

- Algorithm

- align vector along the main diagonal: one-to-one communication  $(\tau + \mu) \lg n$
- Broadcast vector element to  $n$  processors in column: one-to-all broadcast  $(\tau + \mu) \lg n$
- Local multiplication 1
- Sum partial  $y$  values in each row: all-to-one reduction  $(\tau + \mu) \lg n + \lg n$

- Communication time =  $O(\tau \log n + \mu \cdot \log n)$

- Computation time =  $O(\log n)$

From reduction

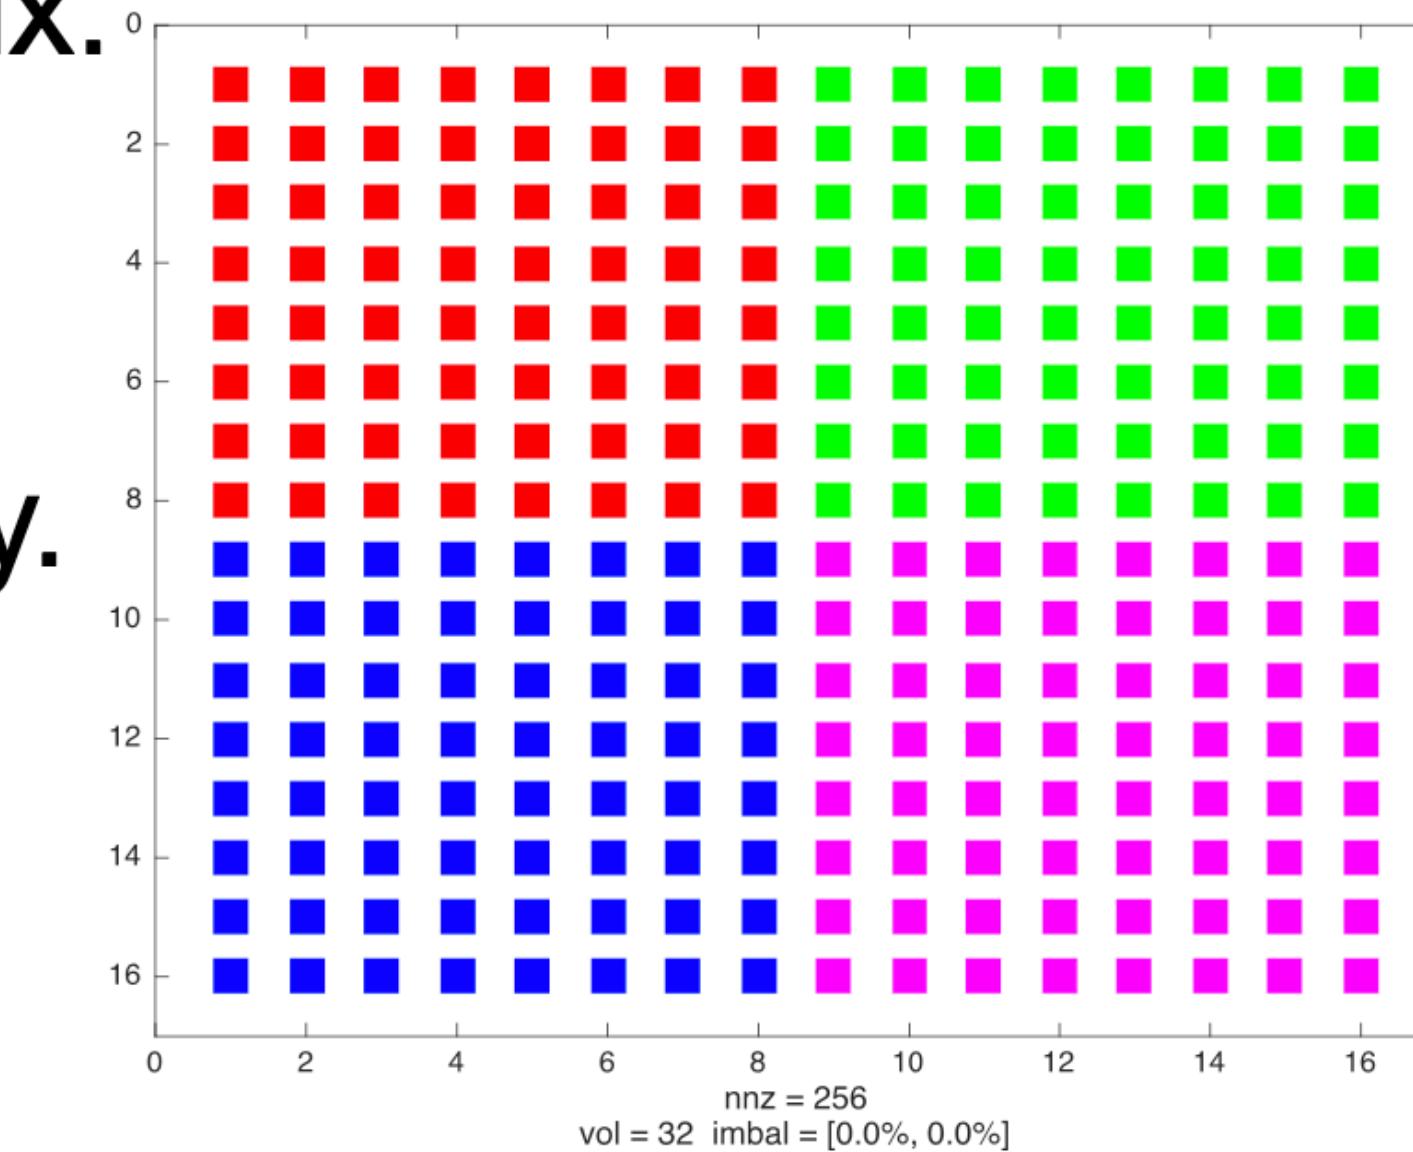
- Is the algorithm efficient?

- No – we need  $O(1)$  run-time!

# 2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- When using fewer than  $n^2$  processors, each process owns an  $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$  block of the matrix.
- The vector is distributed in portions of  $\frac{n}{\sqrt{p}}$  elements in the last processor-column only.

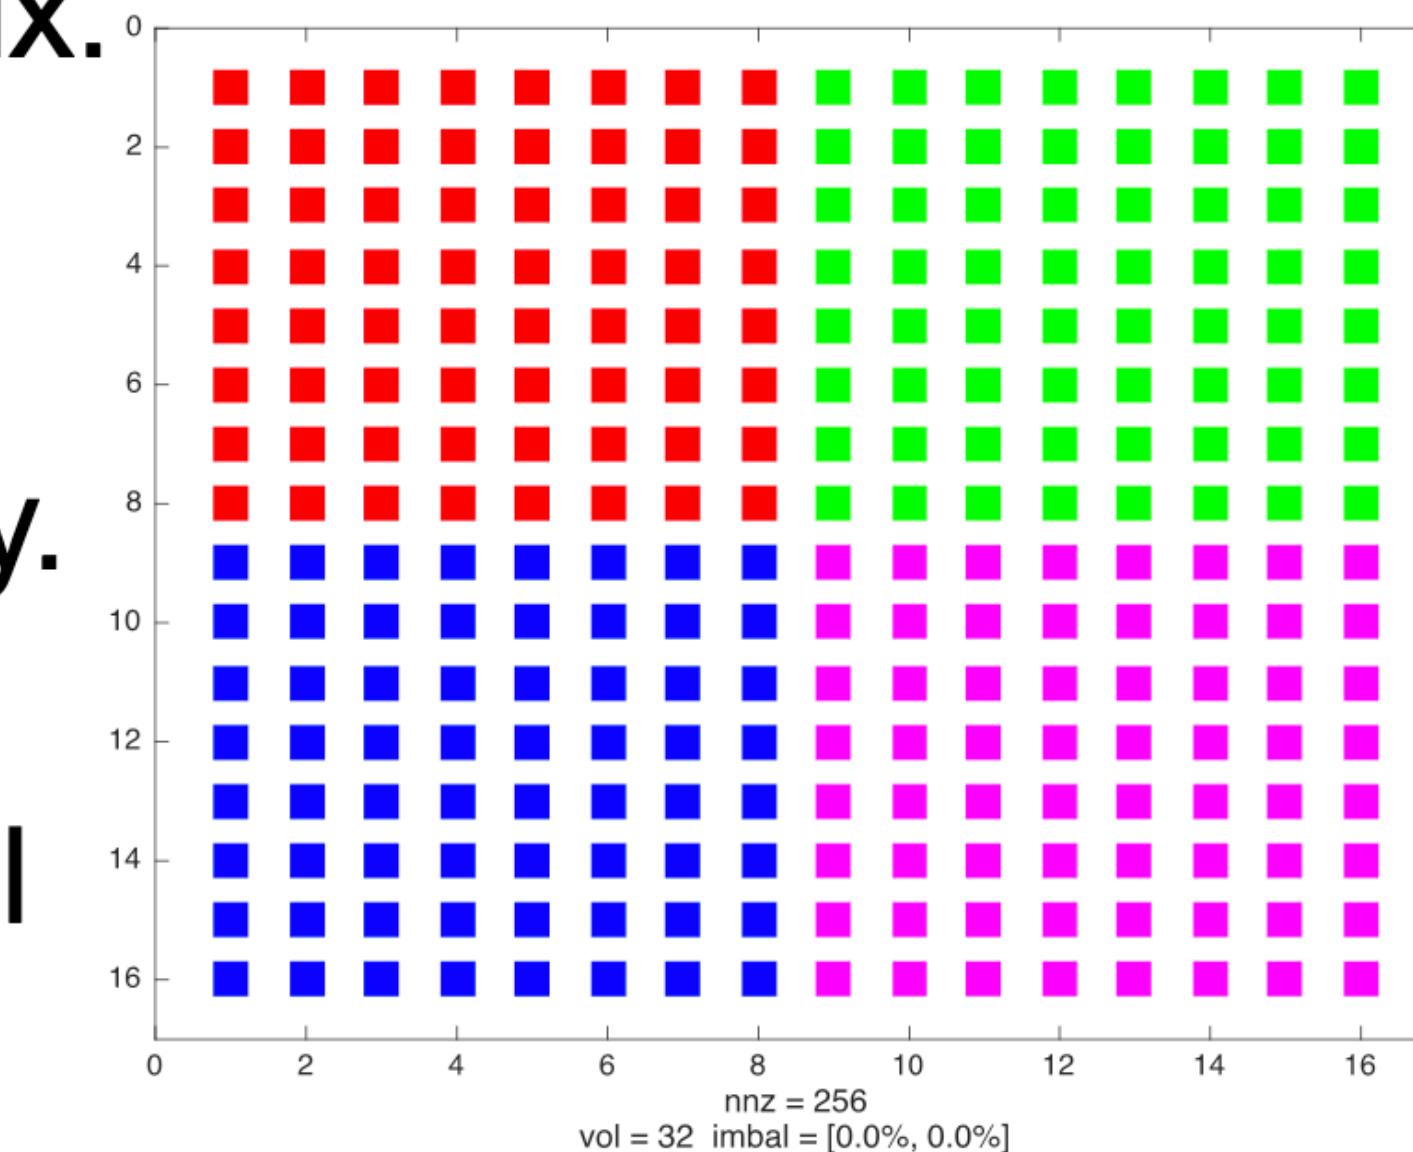


# 2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- When using fewer than  $n^2$  processors, each process owns an  $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$  block of the matrix.
- The vector is distributed in portions of  $\frac{n}{\sqrt{p}}$  elements in the last processor-column only.
- In this case, the message sizes for the alignment, broadcast, and reduction are all  $\frac{n}{\sqrt{p}}$ .

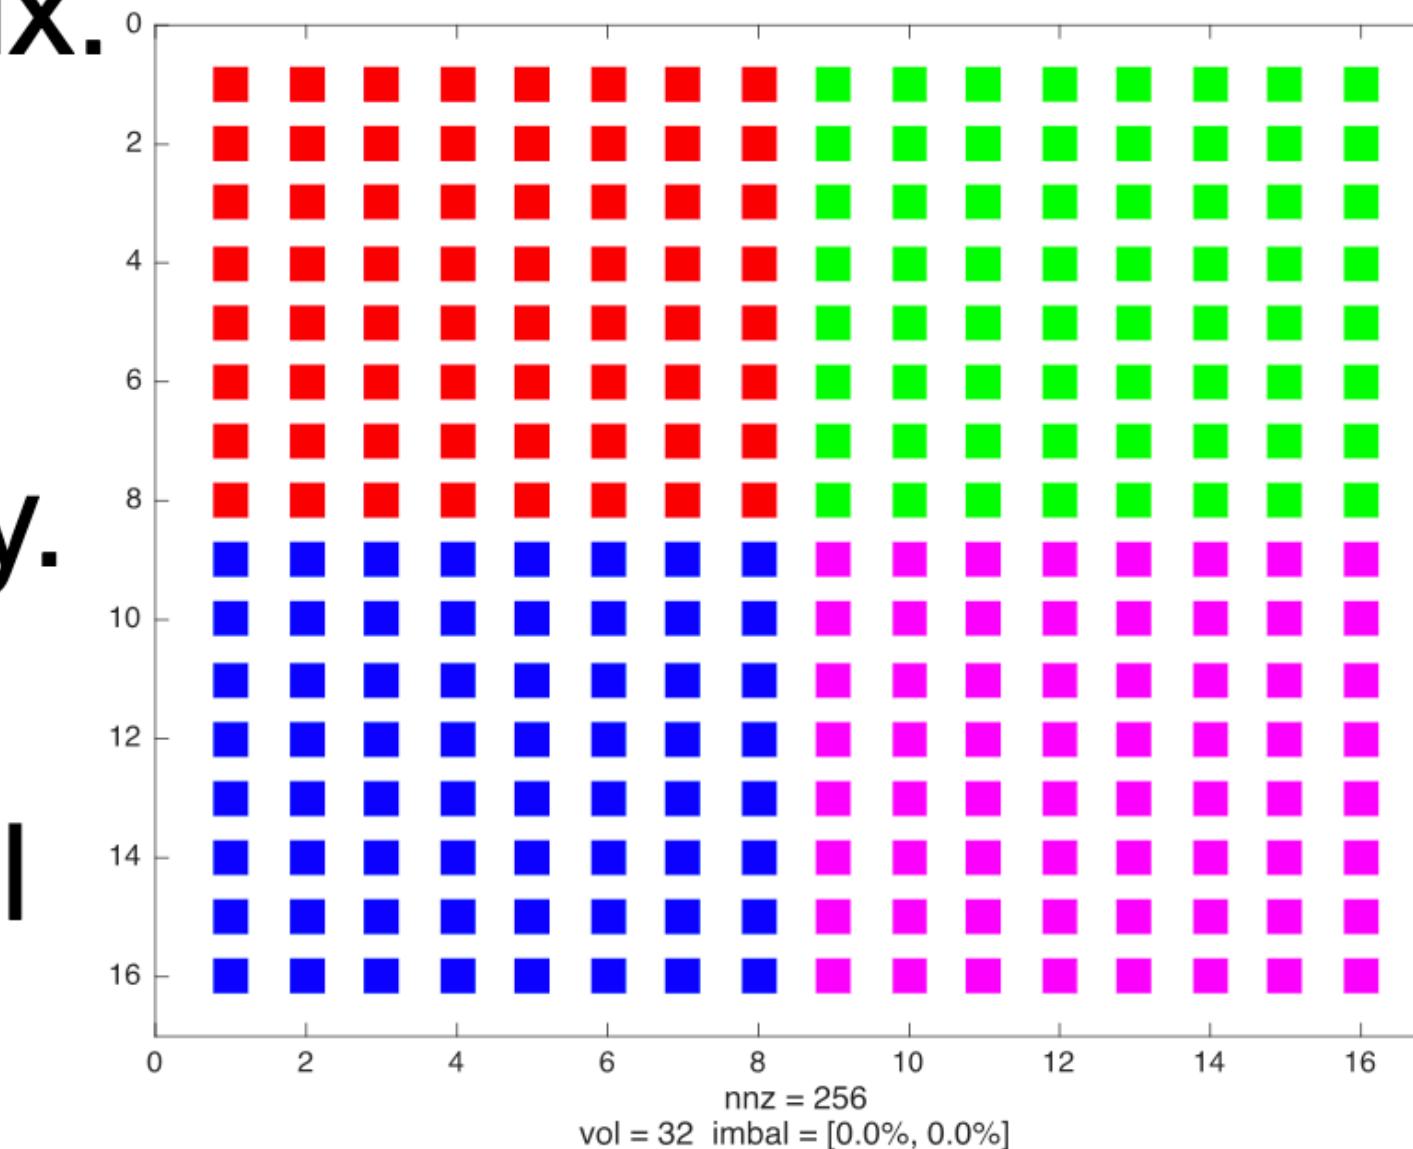
Part of the vector



# 2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- When using fewer than  $n^2$  processors, each process owns an  $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$  block of the matrix.
- The vector is distributed in portions of  $\frac{n}{\sqrt{p}}$  elements in the last processor-column only.
- In this case, the message sizes for the alignment, broadcast, and reduction are all  $\frac{n}{\sqrt{p}}$ .
- The computation is a product of an  $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$  submatrix with a vector of length  $\frac{n}{\sqrt{p}}$ .



Run serial algorithm on the submatrix on each processor

# 2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm
  - align vector along the main diagonal:  $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
  - Broadcast vector elements to processors in column:  
 $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
  - Local multiplication:  $O\left(\frac{n^2}{p}\right)$
  - Sum partial y values in each row:  $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$

# 2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm
  - align vector along the main diagonal:  $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
  - Broadcast vector elements to processors in column:  $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
  - Local multiplication:  $O\left(\frac{n^2}{p}\right)$
  - Sum partial y values in each row:  $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
  - Communication time =  $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
  - Computation time =  $O\left(\frac{n^2}{p}\right)$
  - Efficient up to  $O\left(\frac{n^2}{\log^2 n}\right)$  processors.

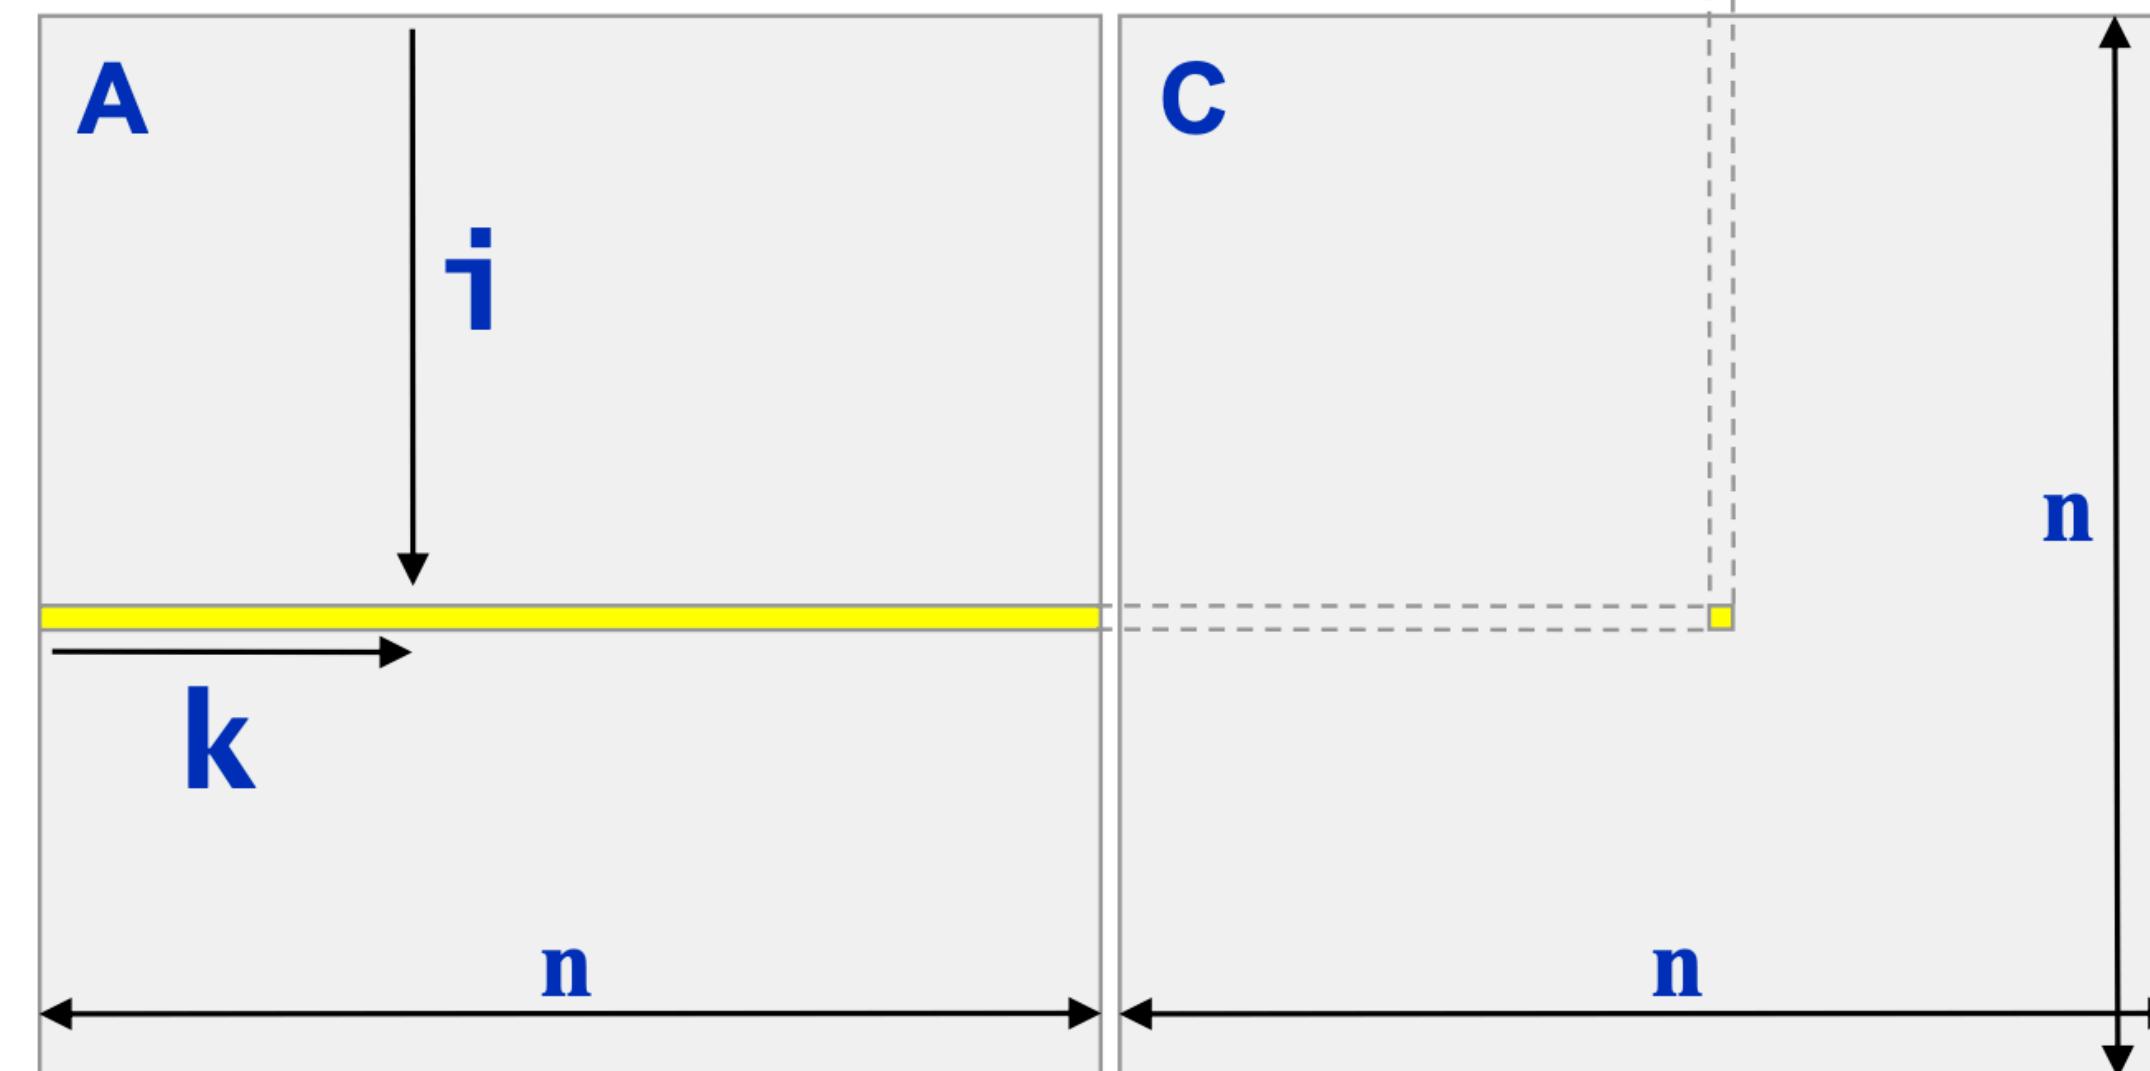
More scalable than 1D  
on the hypercube

# **Matrix-Matrix Multiplication**

# Matrix-Matrix Multiplication

- Consider the problem of multiplying two  $n \times n$  dense, square matrices  $A$  and  $B$  to yield the product matrix  $C = A \times B$ .
- The serial complexity is  $O(n^3)$ .

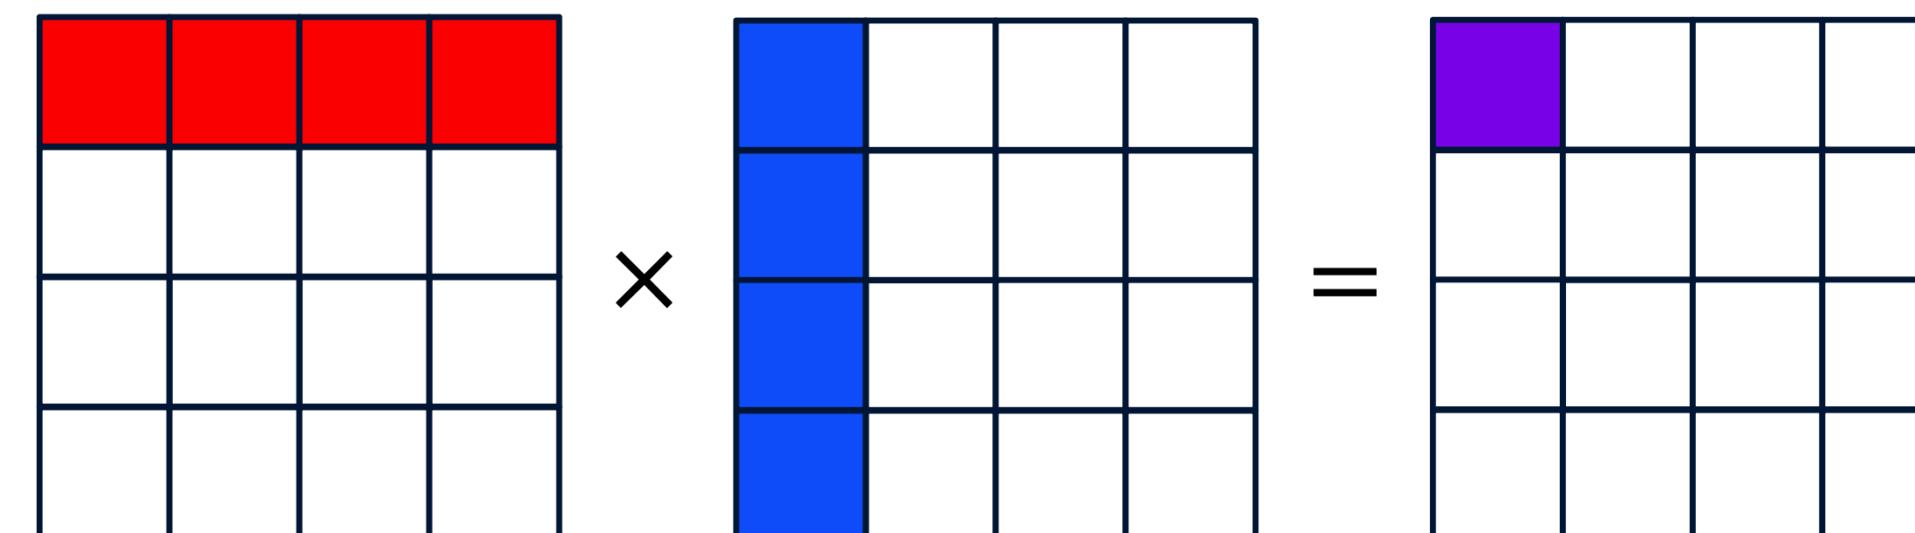
```
for (i = 0; i < n; ++i)
    for (j = 0; j < n; ++j) {
        double sum = 0;
        for (k = 0; k < n; ++k)
            sum += A[i][k] * B[k][j];
        C[i][j] = sum;
    }
```



# Initial Algorithm for Matrix-Matrix Multiplication

Assume we have  $n^2$  processors, where every processor  $(i, j)$  has elements  $A[i, j]$  and  $B[i, j]$  in the initial distribution.

For processor  $(i, j)$  to compute  $C[i, j]$ , it needs the  $i$ -th row of  $A$  and  $j$ -th col of  $B$ .



Idea:

- Allgather  $i$ -th row of  $A$  onto each processor for each row  $i$
- Allgather  $j$ -th col of  $B$  onto each processor for each col  $j$
- Perform dot product of row and col to get  $C[i, j]$

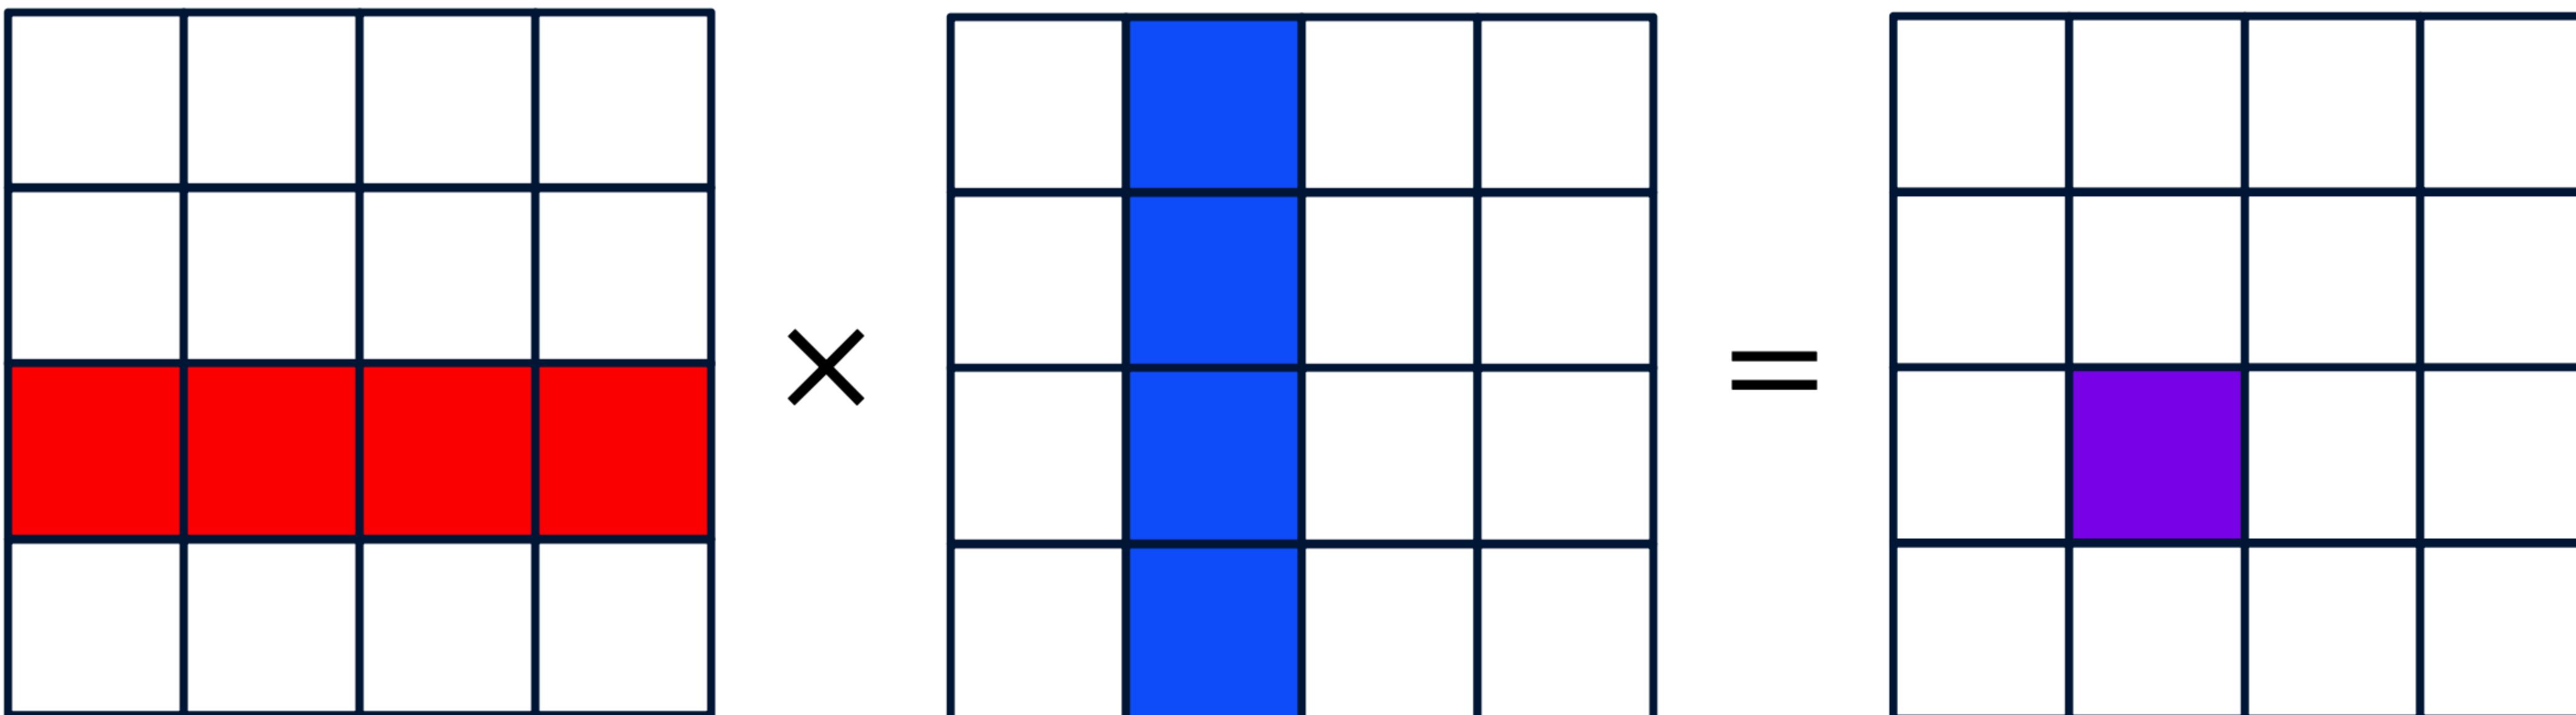
# Data Requirements for Matrix-Matrix Multiply

Consider data needed for output matrix block shown in purple

$$\begin{array}{|c|c|c|c|} \hline \textcolor{red}{\square} & \textcolor{red}{\square} & \textcolor{red}{\square} & \textcolor{red}{\square} \\ \hline \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \end{array} \times \begin{array}{|c|c|c|c|} \hline \textcolor{blue}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{blue}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{blue}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{blue}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline \textcolor{purple}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} & \textcolor{white}{\square} \\ \hline \end{array}$$

# Data Requirements for Matrix-Matrix Multiply

Consider data needed for output matrix block shown in purple



# Matrix-Matrix Multiplication

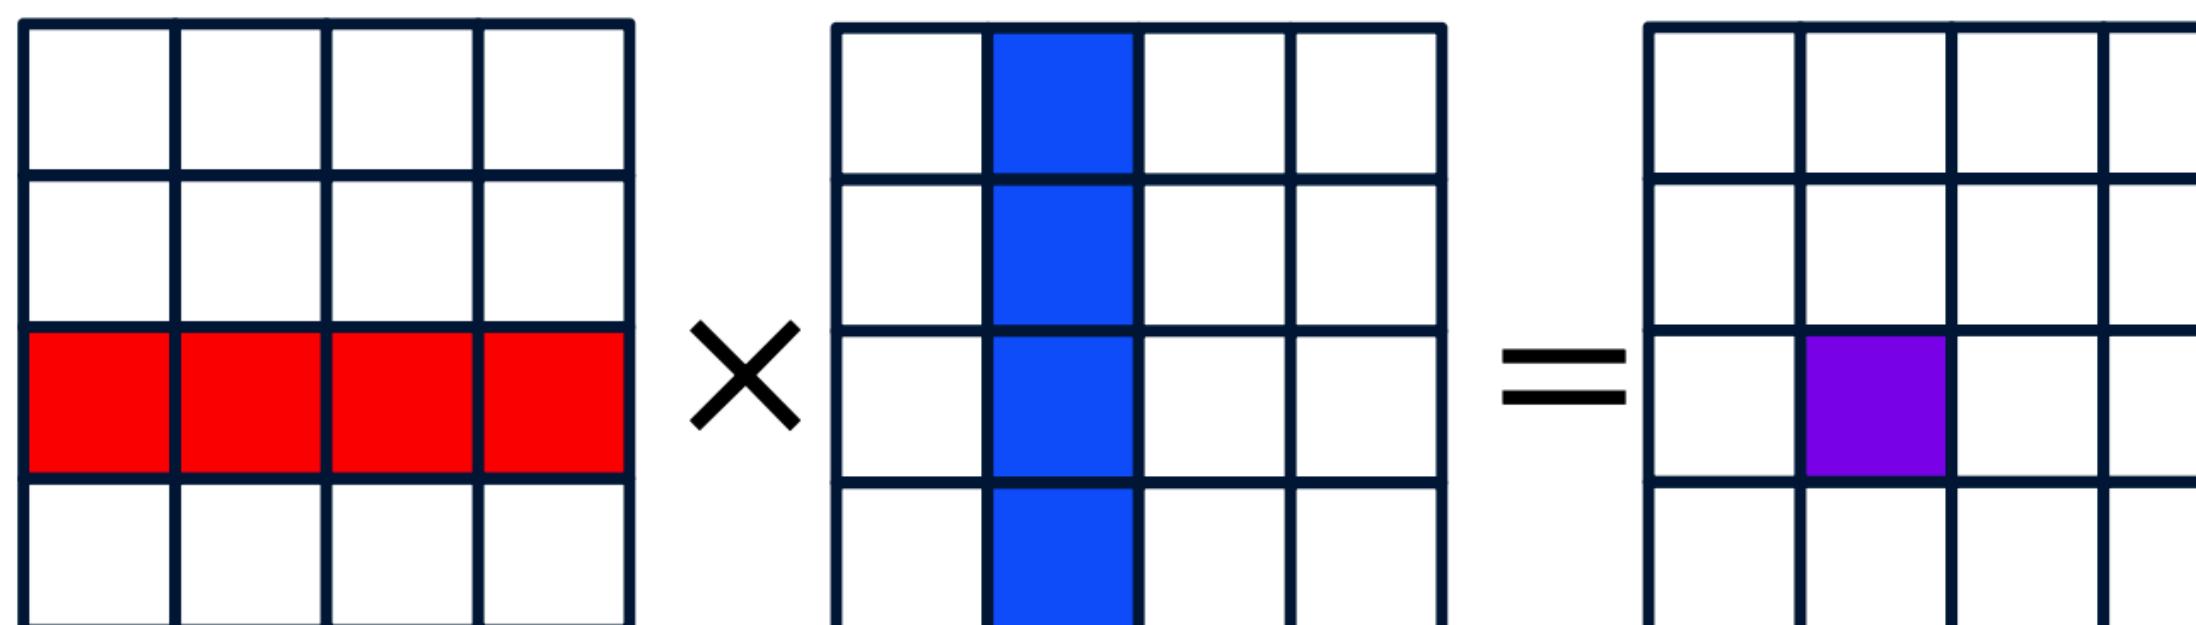
- Consider the problem of multiplying two  $n \times n$  dense, square matrices  $A$  and  $B$  to yield the product matrix  $C = A \times B$ .
- The serial complexity is  $O(n^3)$ .
  - We do not consider better serial algorithms (Strassen's method)  
Note: these can still be used as serial kernels in the parallel algorithms.

# Matrix-Matrix Multiplication

- Consider the problem of multiplying two  $n \times n$  dense, square matrices  $A$  and  $B$  to yield the product matrix  $C = A \times B$ .
- The serial complexity is  $O(n^3)$ .
  - We do not consider better serial algorithms (Strassen's method)  
Note: these can still be used as serial kernels in the parallel algorithms.
- A useful concept in this case is called *block* operations
  - an  $n \times n$  matrix  $A$  can be regarded as a  $q \times q$  array of blocks
  - $A_{i,j}$  ( $0 \leq i, j < q$ ) such that each block is an  $(n/q) \times (n/q)$  submatrix.
- In this view, we perform  $q^3$  matrix multiplications, each involving  $(n/q) \times (n/q)$  matrices.

# Blocked Matrix-Matrix Multiply

- Consider two  $n \times n$  matrices  $A$  and  $B$  partitioned into  $p$  blocks  $A_{i,j}$  and  $B_{i,j}$  ( $0 \leq i, j < \sqrt{p}$ ) of size  $(n/\sqrt{p}) \times (n/\sqrt{p})$  each.
- Process  $P_{i,j}$  initially stores  $A_{i,j}$  and  $B_{i,j}$  and computes block  $C_{i,j}$  of the result matrix.
- Computing submatrix  $C_{i,j}$  requires all submatrices  $A_{i,k}$  and  $B_{k,j}$  for  $0 \leq k < \sqrt{p}$ .
  - AllGather blocks of  $A$  along rows and  $B$  along columns.
  - Perform local submatrix multiplication.



# Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:

- $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$

# Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
  - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:
  - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$

# Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
  - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:
  - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
- Total Time:  $O\left(\frac{n^3}{p} + \tau \log p + \mu \frac{n^2}{\sqrt{p}}\right)$
- Efficient for  $p = O(n^2)$

# Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
  - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:
  - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
- Total Time:  $O\left(\frac{n^3}{p} + \tau \log p + \mu \frac{n^2}{\sqrt{p}}\right)$
- Efficient for  $p = O(n^2)$

Any issues with this algorithm?

# Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
  - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:
  - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
- Total Time:  $O\left(\frac{n^3}{p} + \tau \log p + \mu \frac{n^2}{\sqrt{p}}\right)$
- Efficient for  $p = O(n^2)$
- Major drawback of the algorithm is that it is not memory optimal.

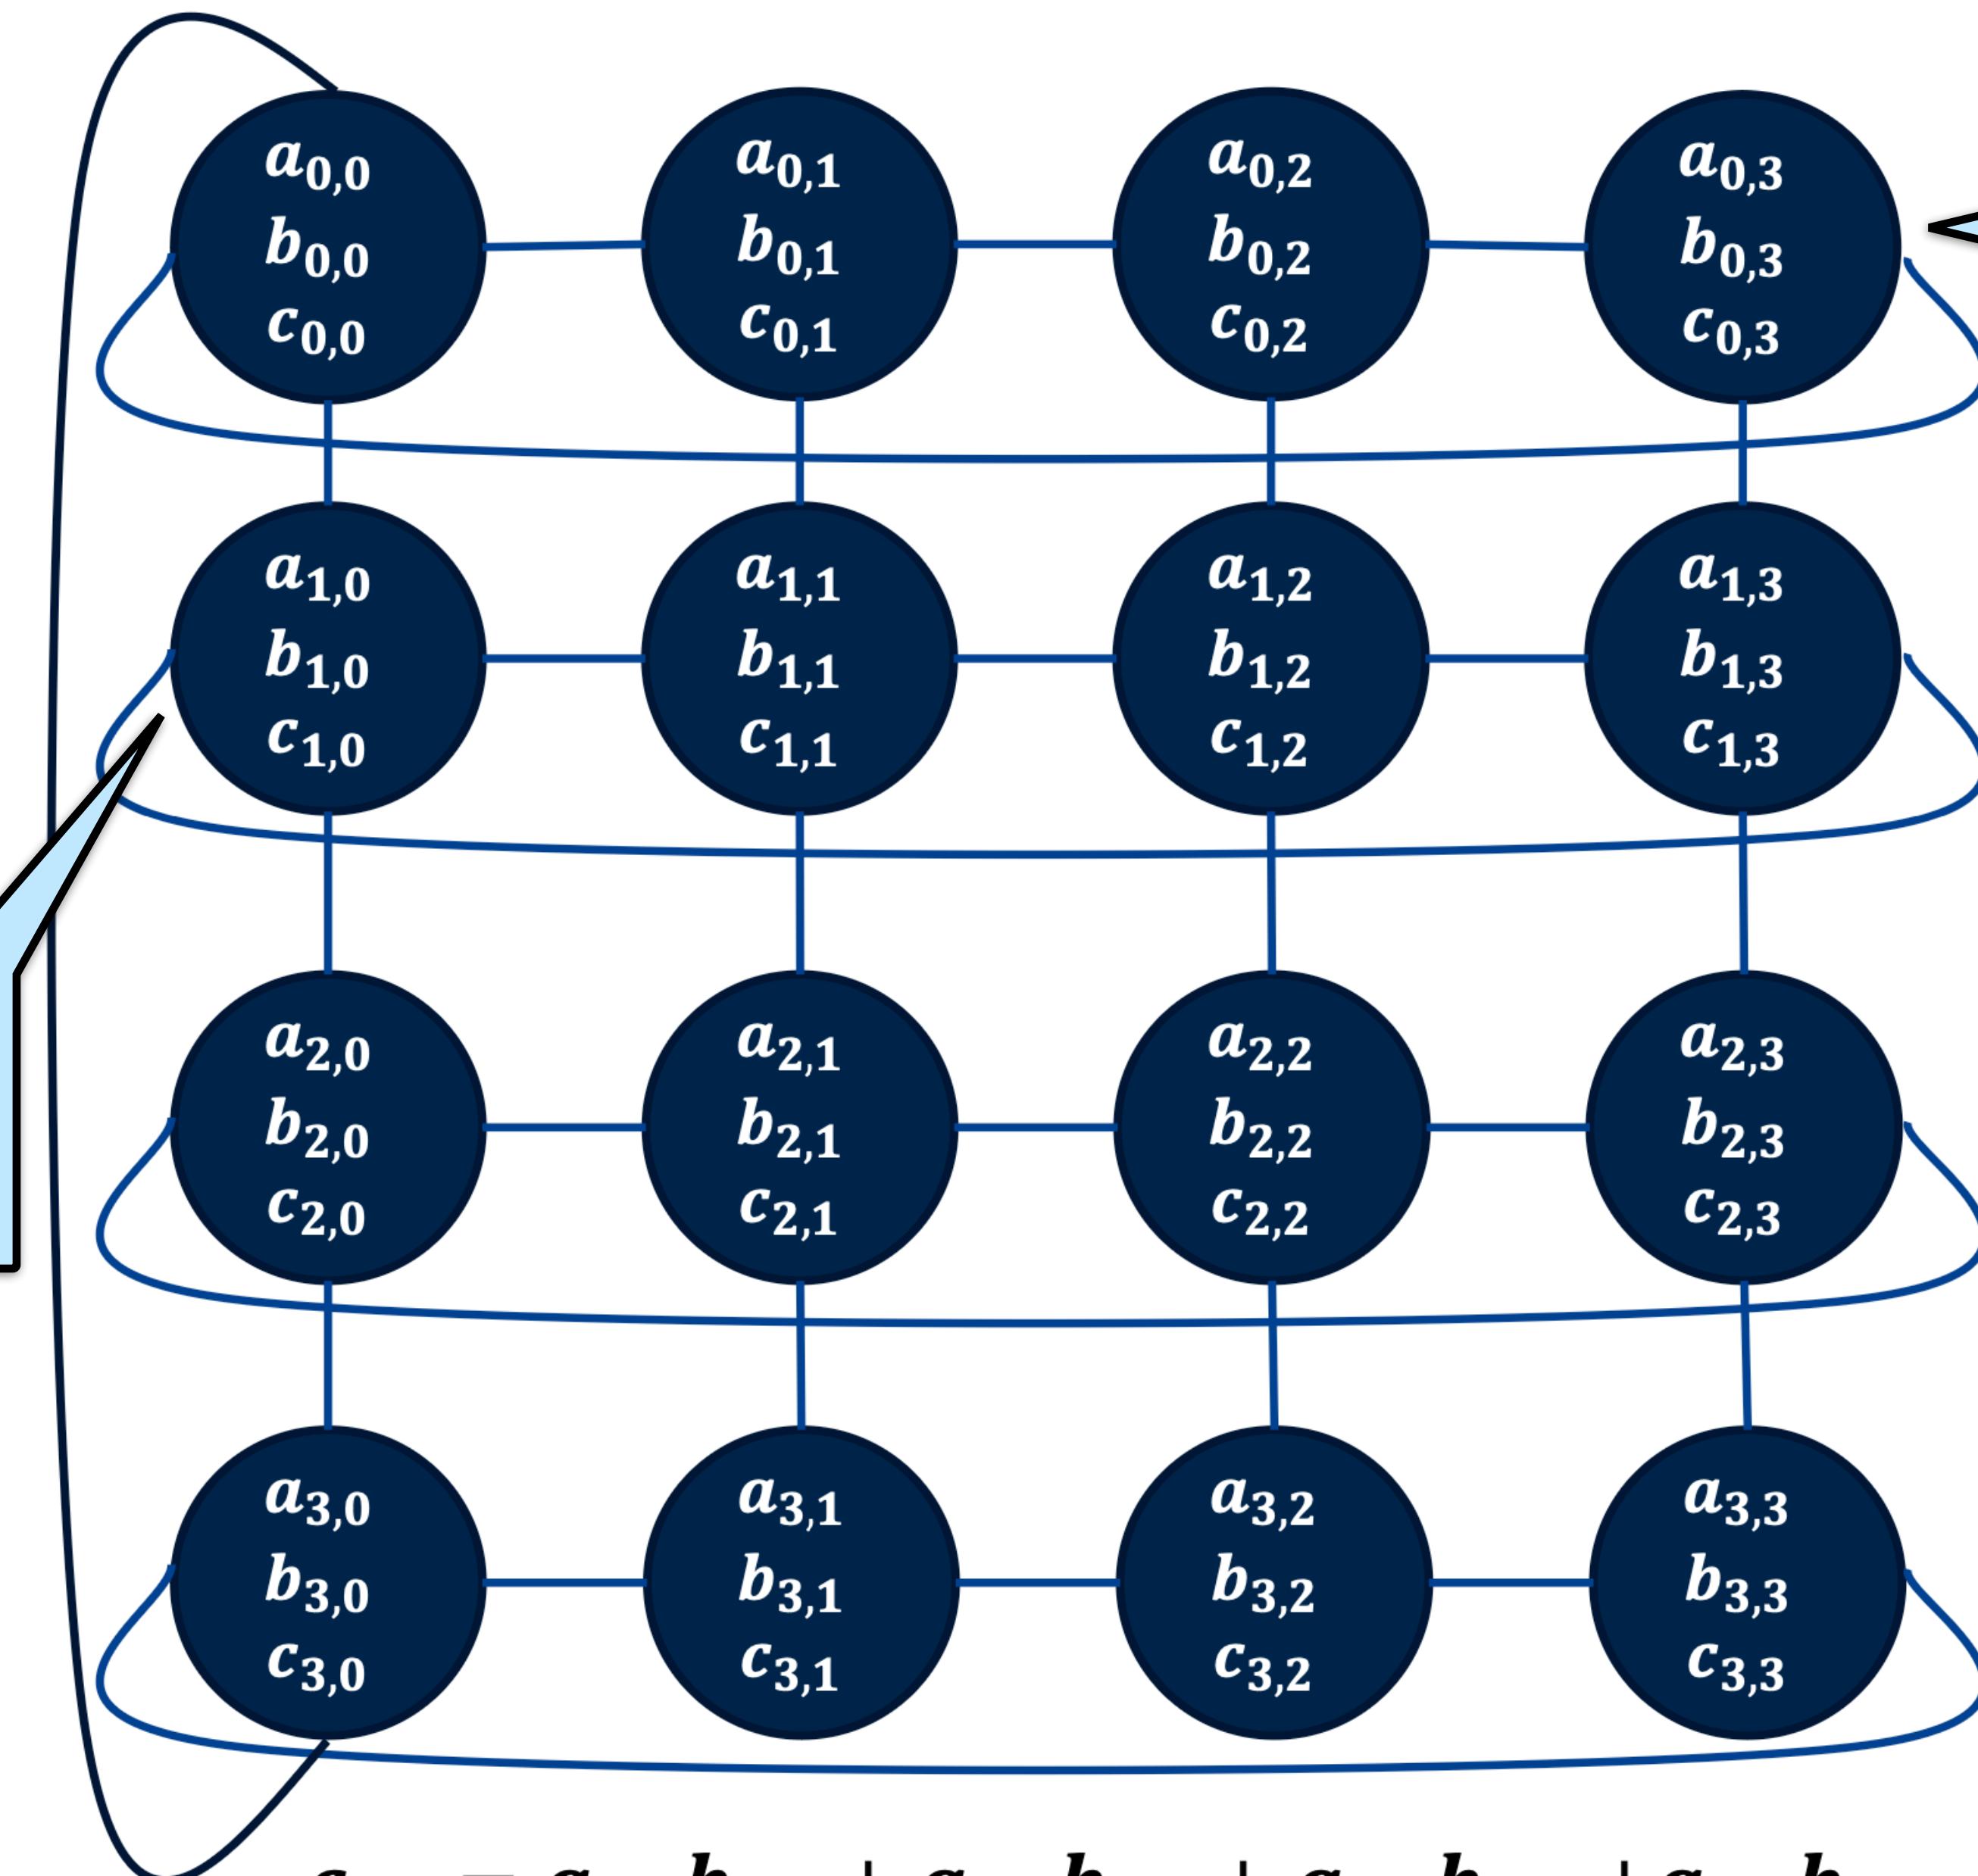
Extra  $\sqrt{p}$  factor of memory consumption!

Ideal: Memory is serial\_memory / p

# **START PART 2 HERE**

Restriction - one instance of each element in A, B across all processors

Send b elements to the north one hop at a time



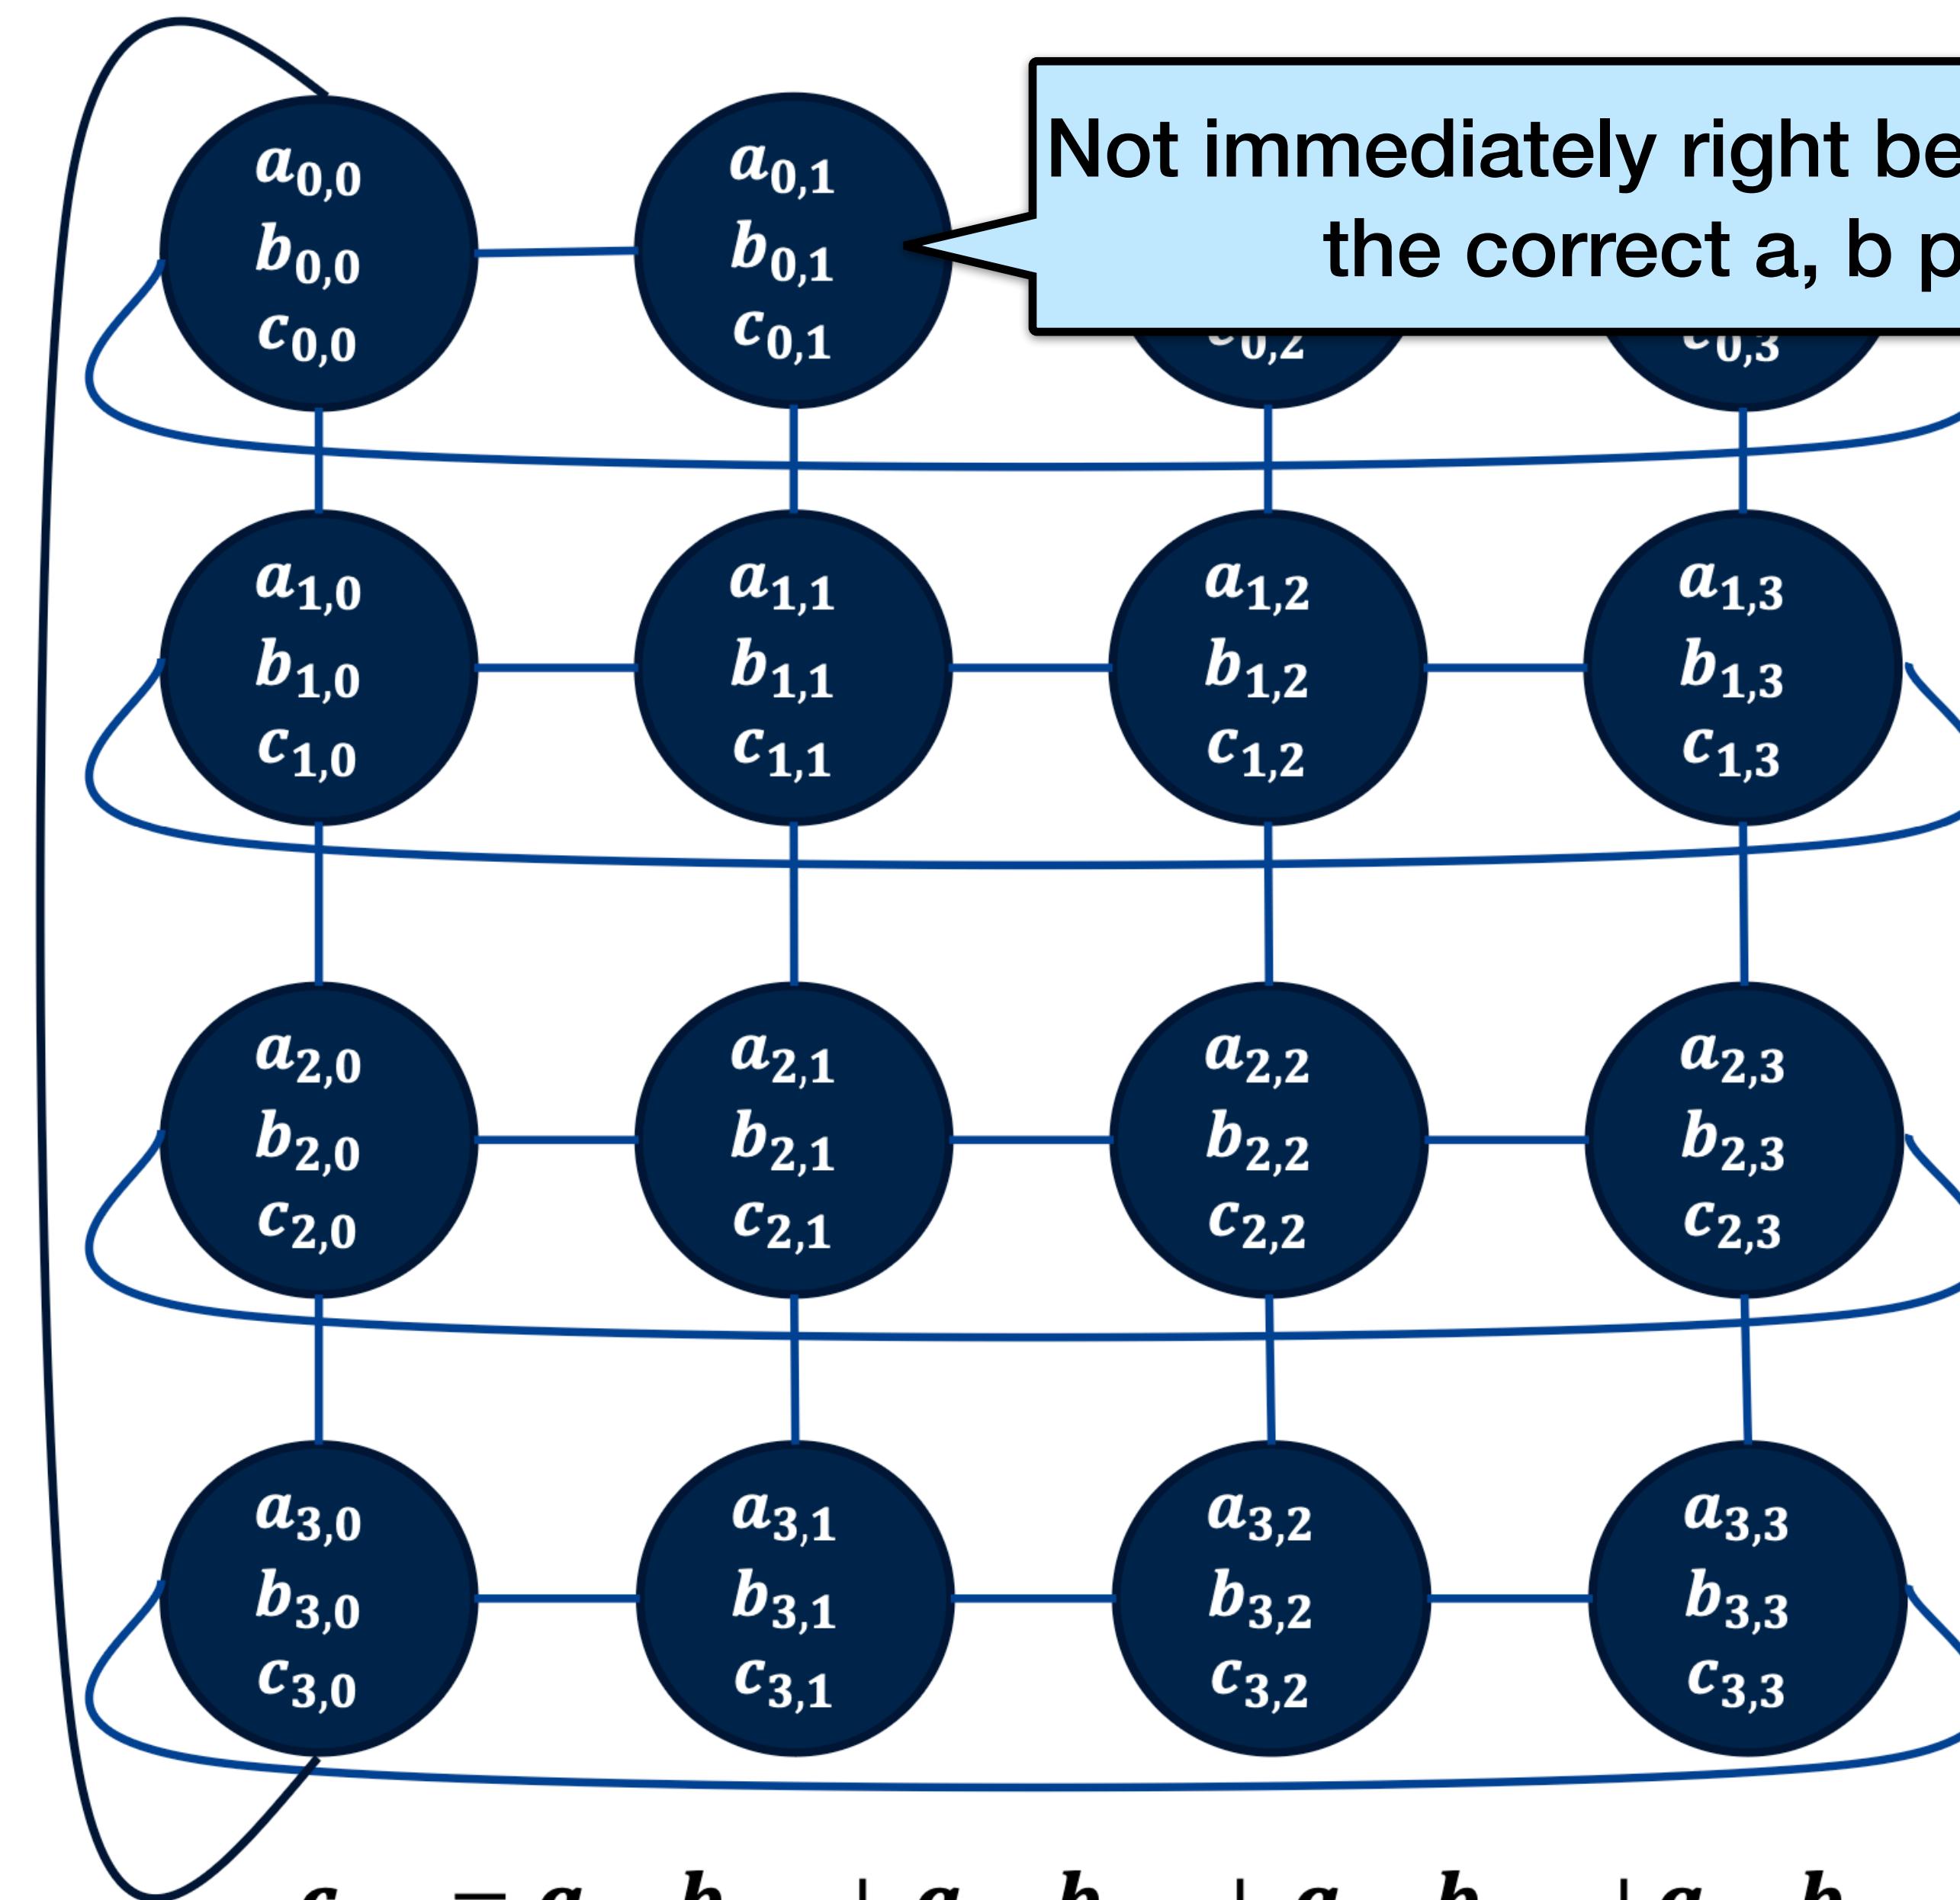
Send a elements to the left one hop at a time

$$c_{0,0} = a_{0,0}b_{0,0} + a_{0,1}b_{1,0} + a_{0,2}b_{2,0} + a_{0,3}b_{3,0}$$

$$c_{0,1} = \boxed{a_{0,0}}b_{0,1} + \boxed{a_{0,1}}b_{1,1} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

Goal - all processors are computing in each round

We need the corresponding a and b on each processor in each round



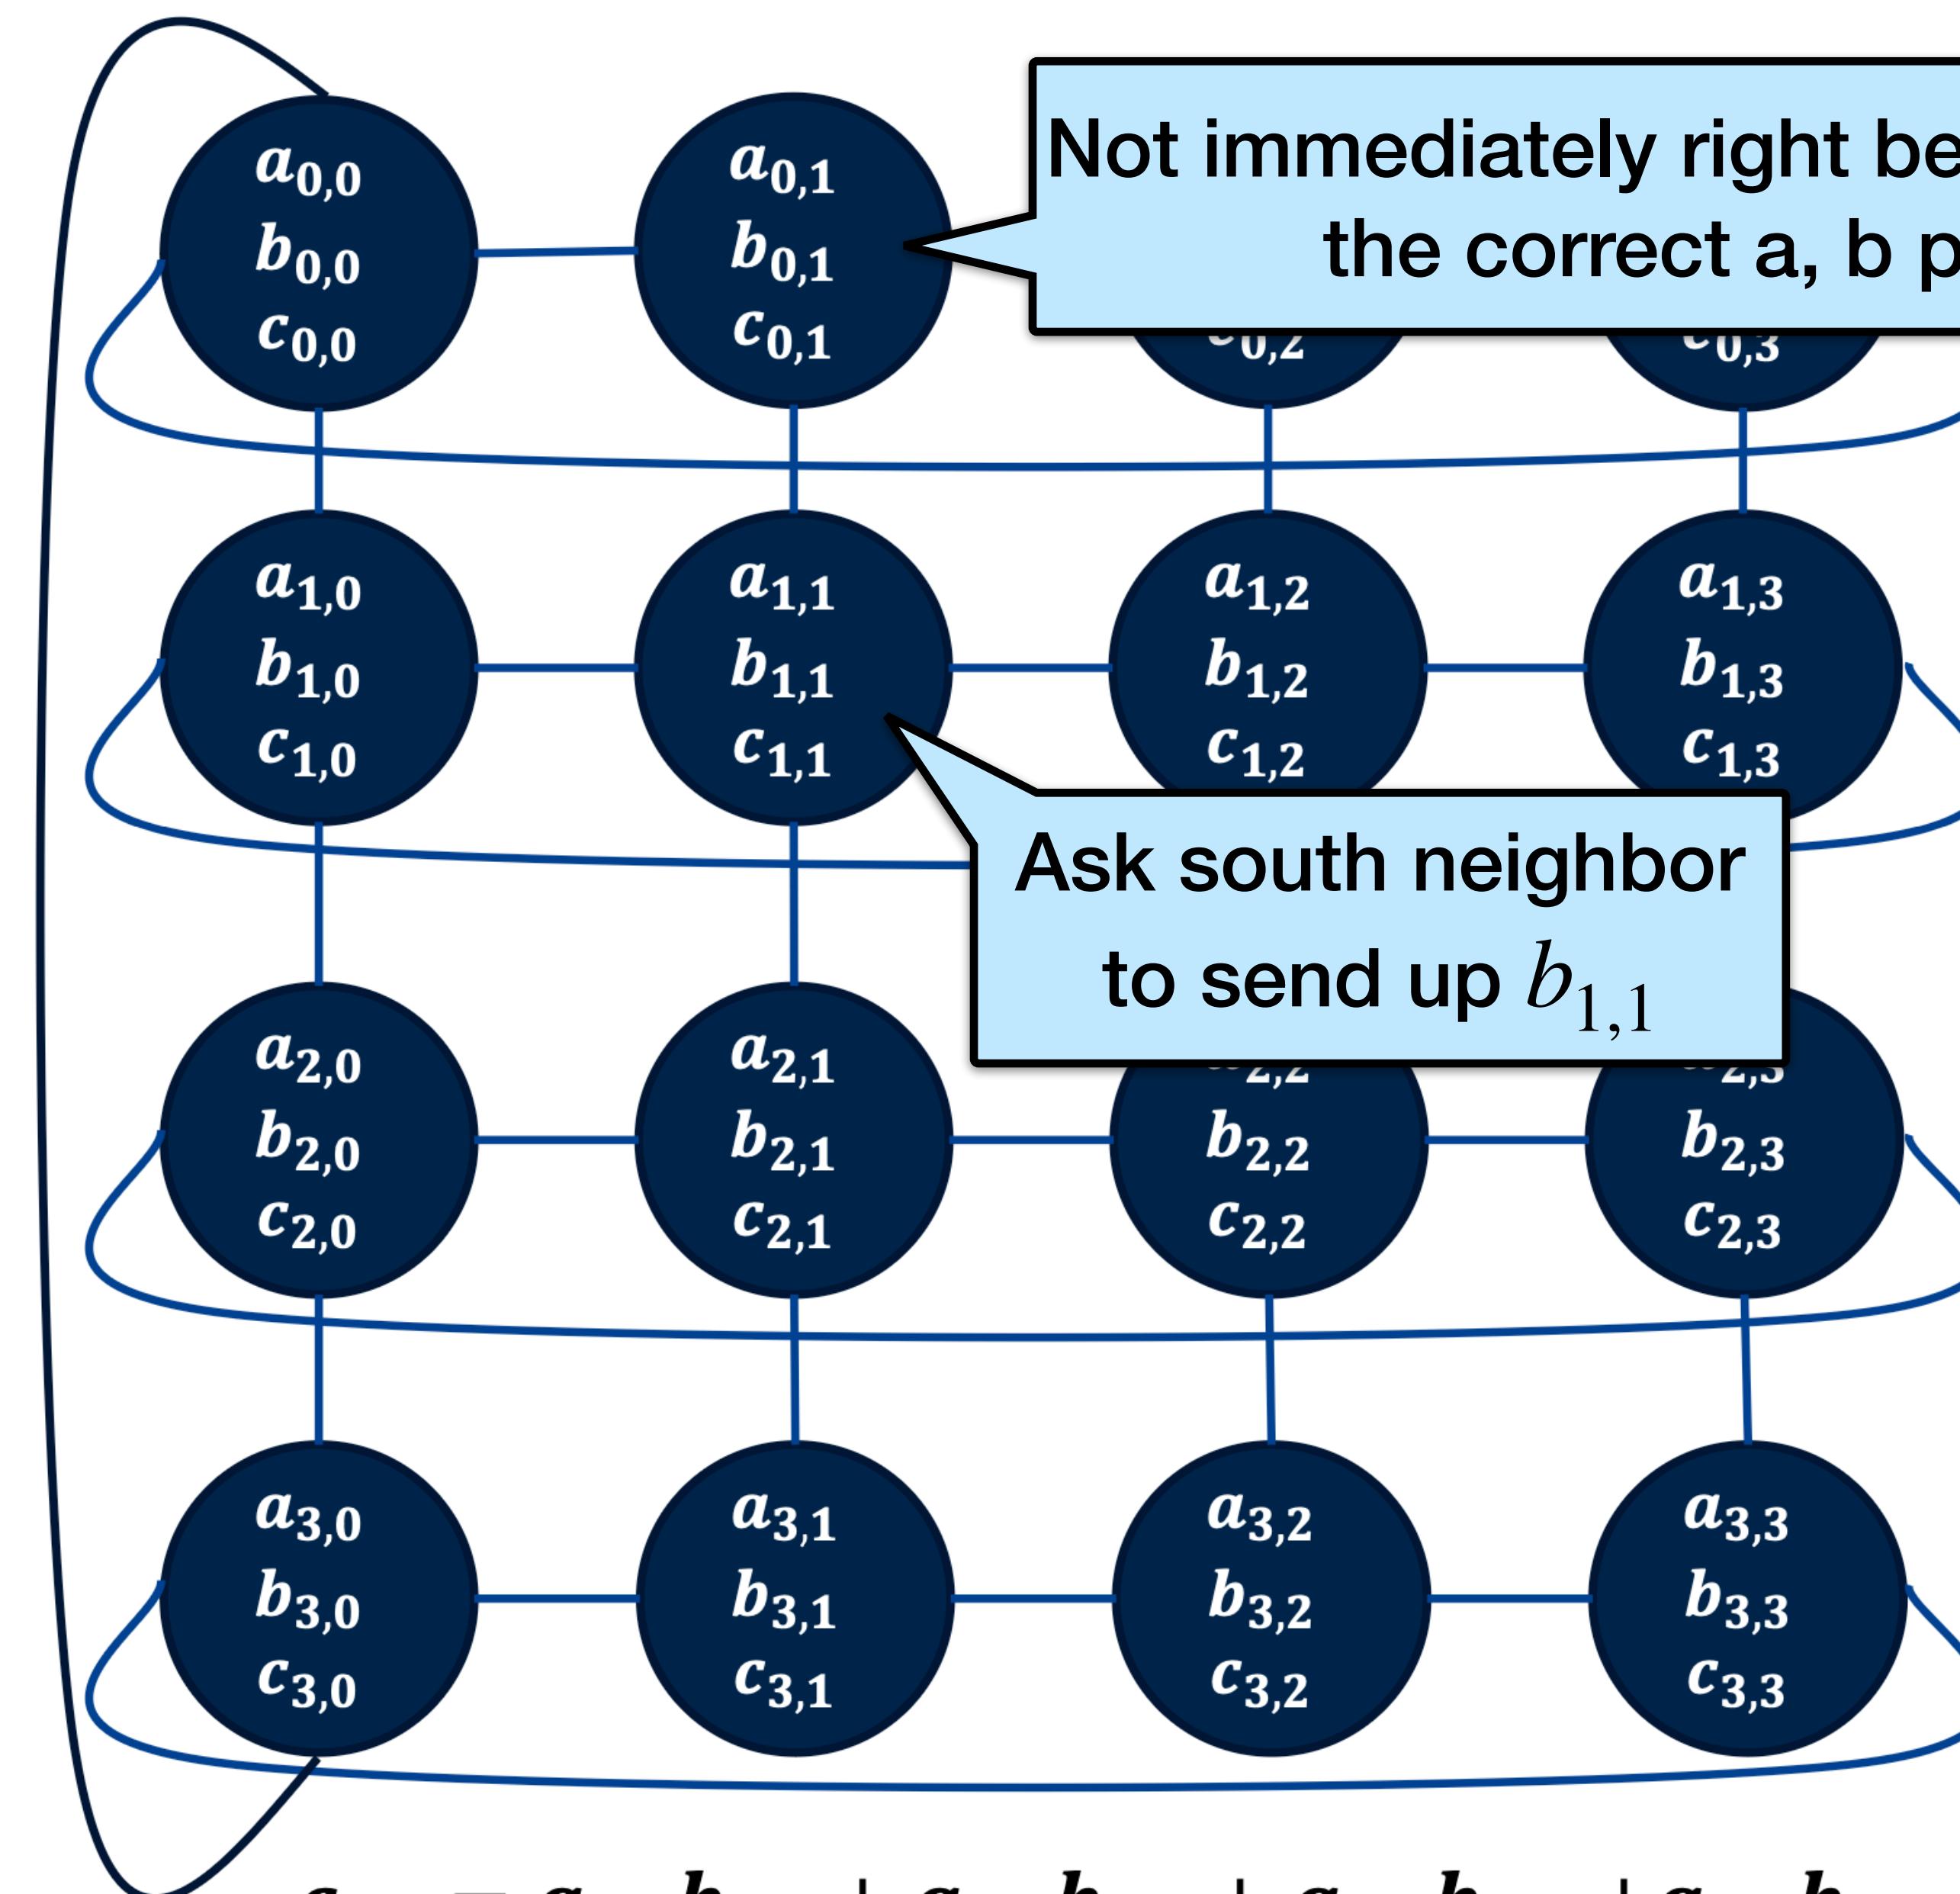
Not immediately right because not the correct a, b pair

$$c_{0,0} = a_{0,0}b_{0,0} + a_{0,1}b_{1,0} + a_{0,2}b_{2,0} + a_{0,3}b_{3,0}$$
$$c_{0,1} = \boxed{a_{0,0}}b_{0,1} + \boxed{a_{0,1}}b_{1,1} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

Can compute the terms in any order

Goal - all processors are computing in each round

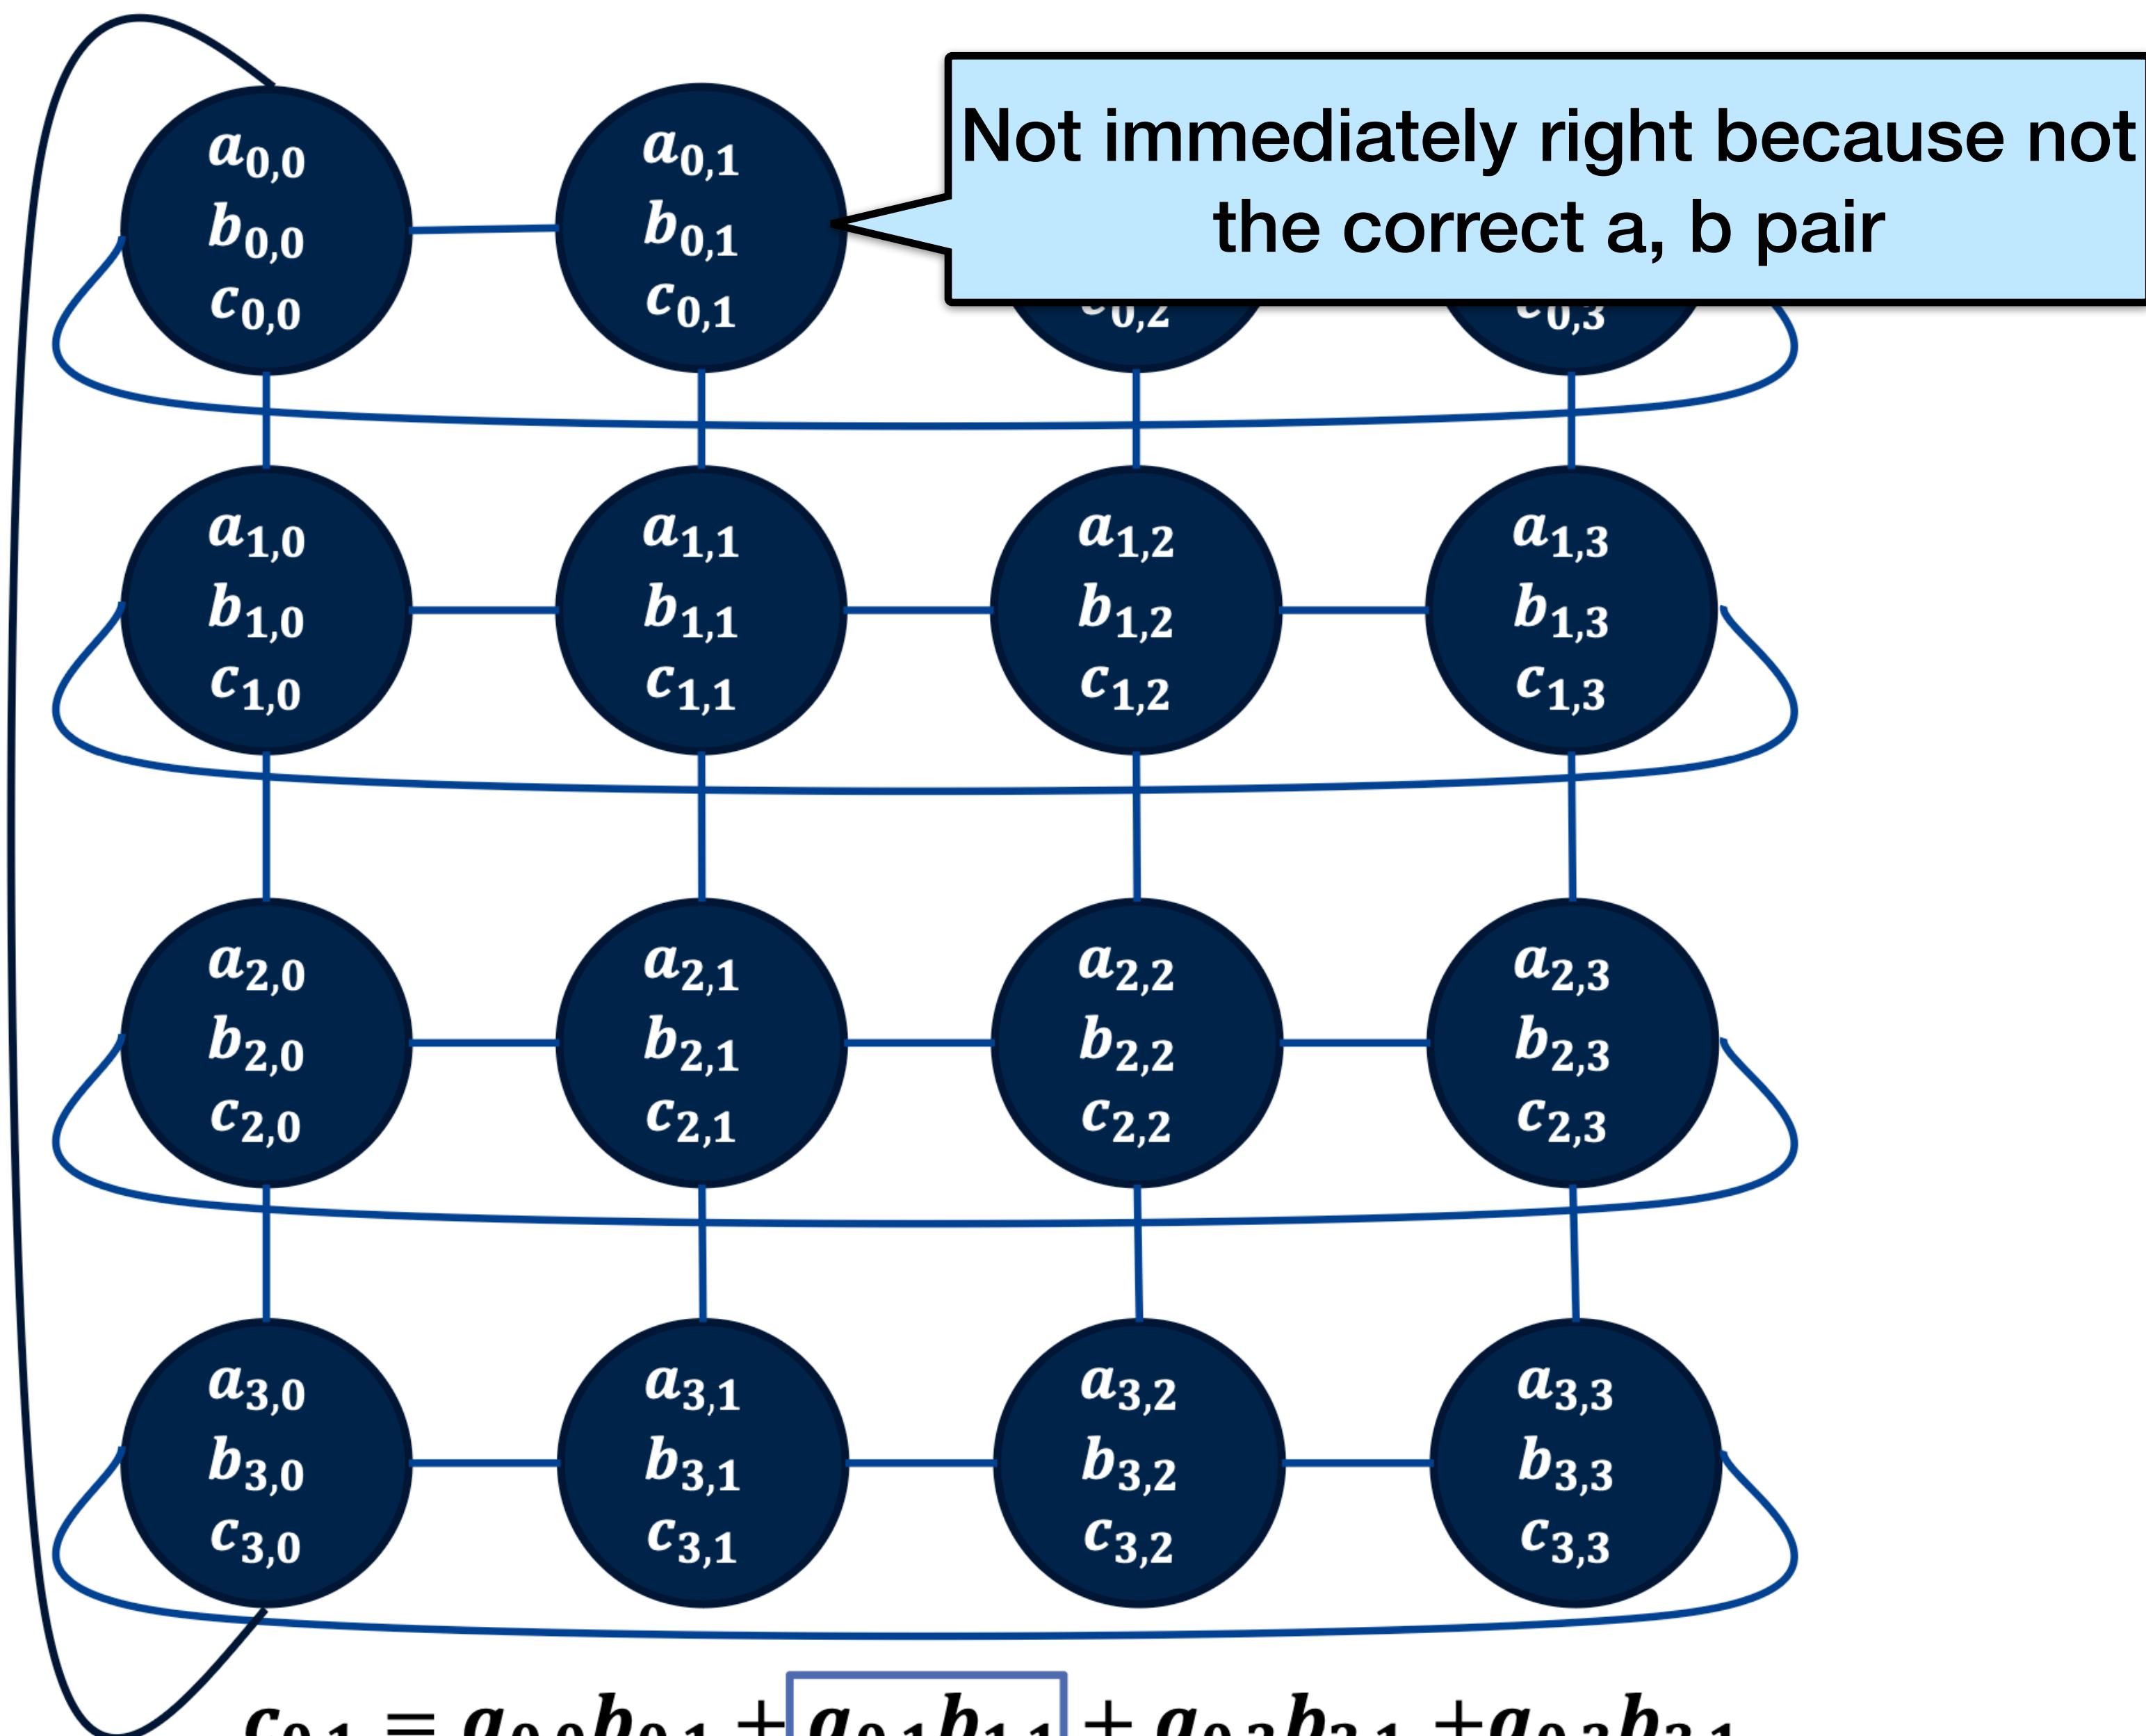
We need the corresponding a and b on each processor in each round



$$c_{0,0} = a_{0,0}b_{0,0} + a_{0,1}b_{1,0} + a_{0,2}b_{2,0} + a_{0,3}b_{3,0}$$

$$c_{0,1} = \boxed{a_{0,0}}b_{0,1} + \boxed{a_{0,1}}b_{1,1} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

Can compute the terms in any order

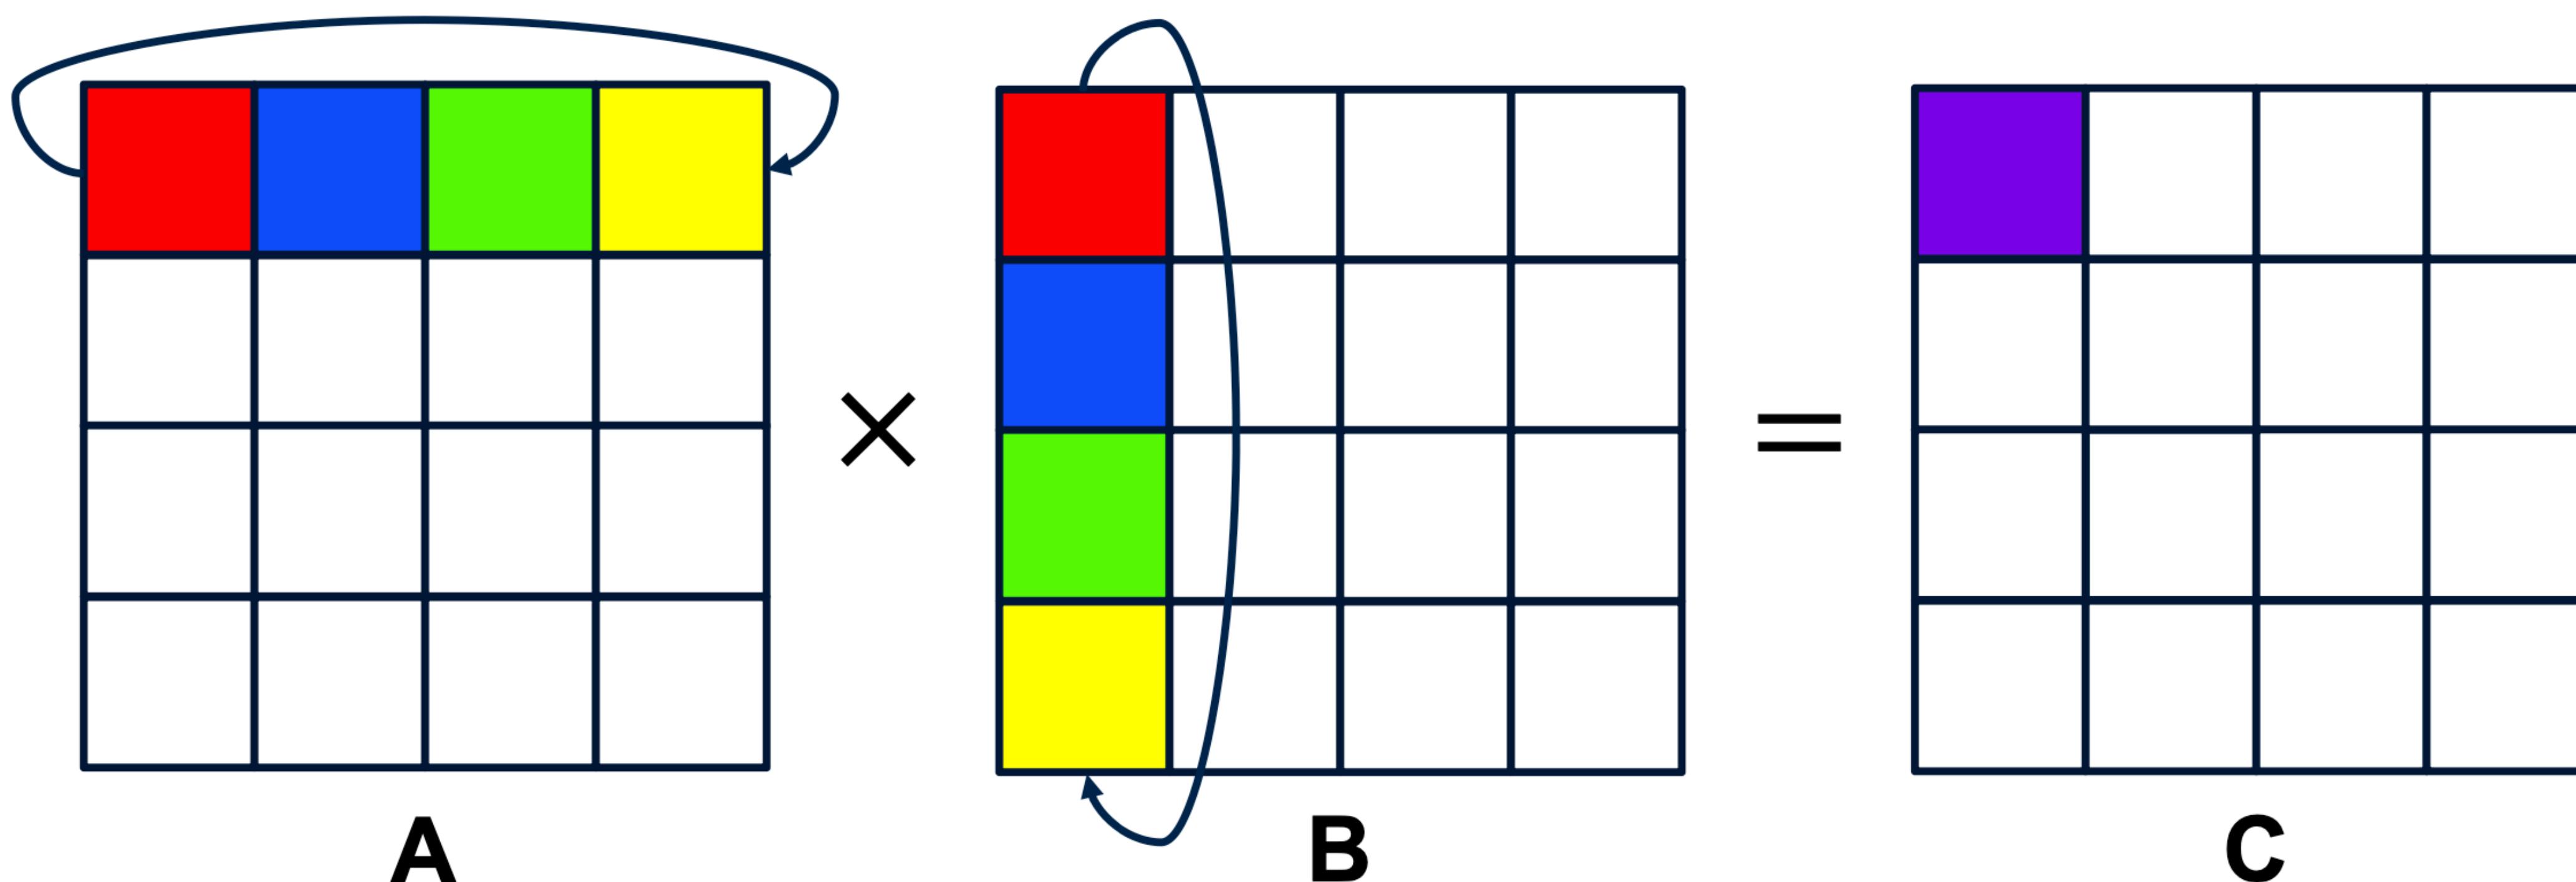


$$c_{0,1} = a_{0,0}b_{0,1} + \boxed{a_{0,1}b_{1,1}} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

$$c_{0,2} = a_{0,0}b_{0,2} + a_{0,1}b_{1,2} + \boxed{a_{0,2}b_{2,2}} + a_{0,3}b_{3,2}$$

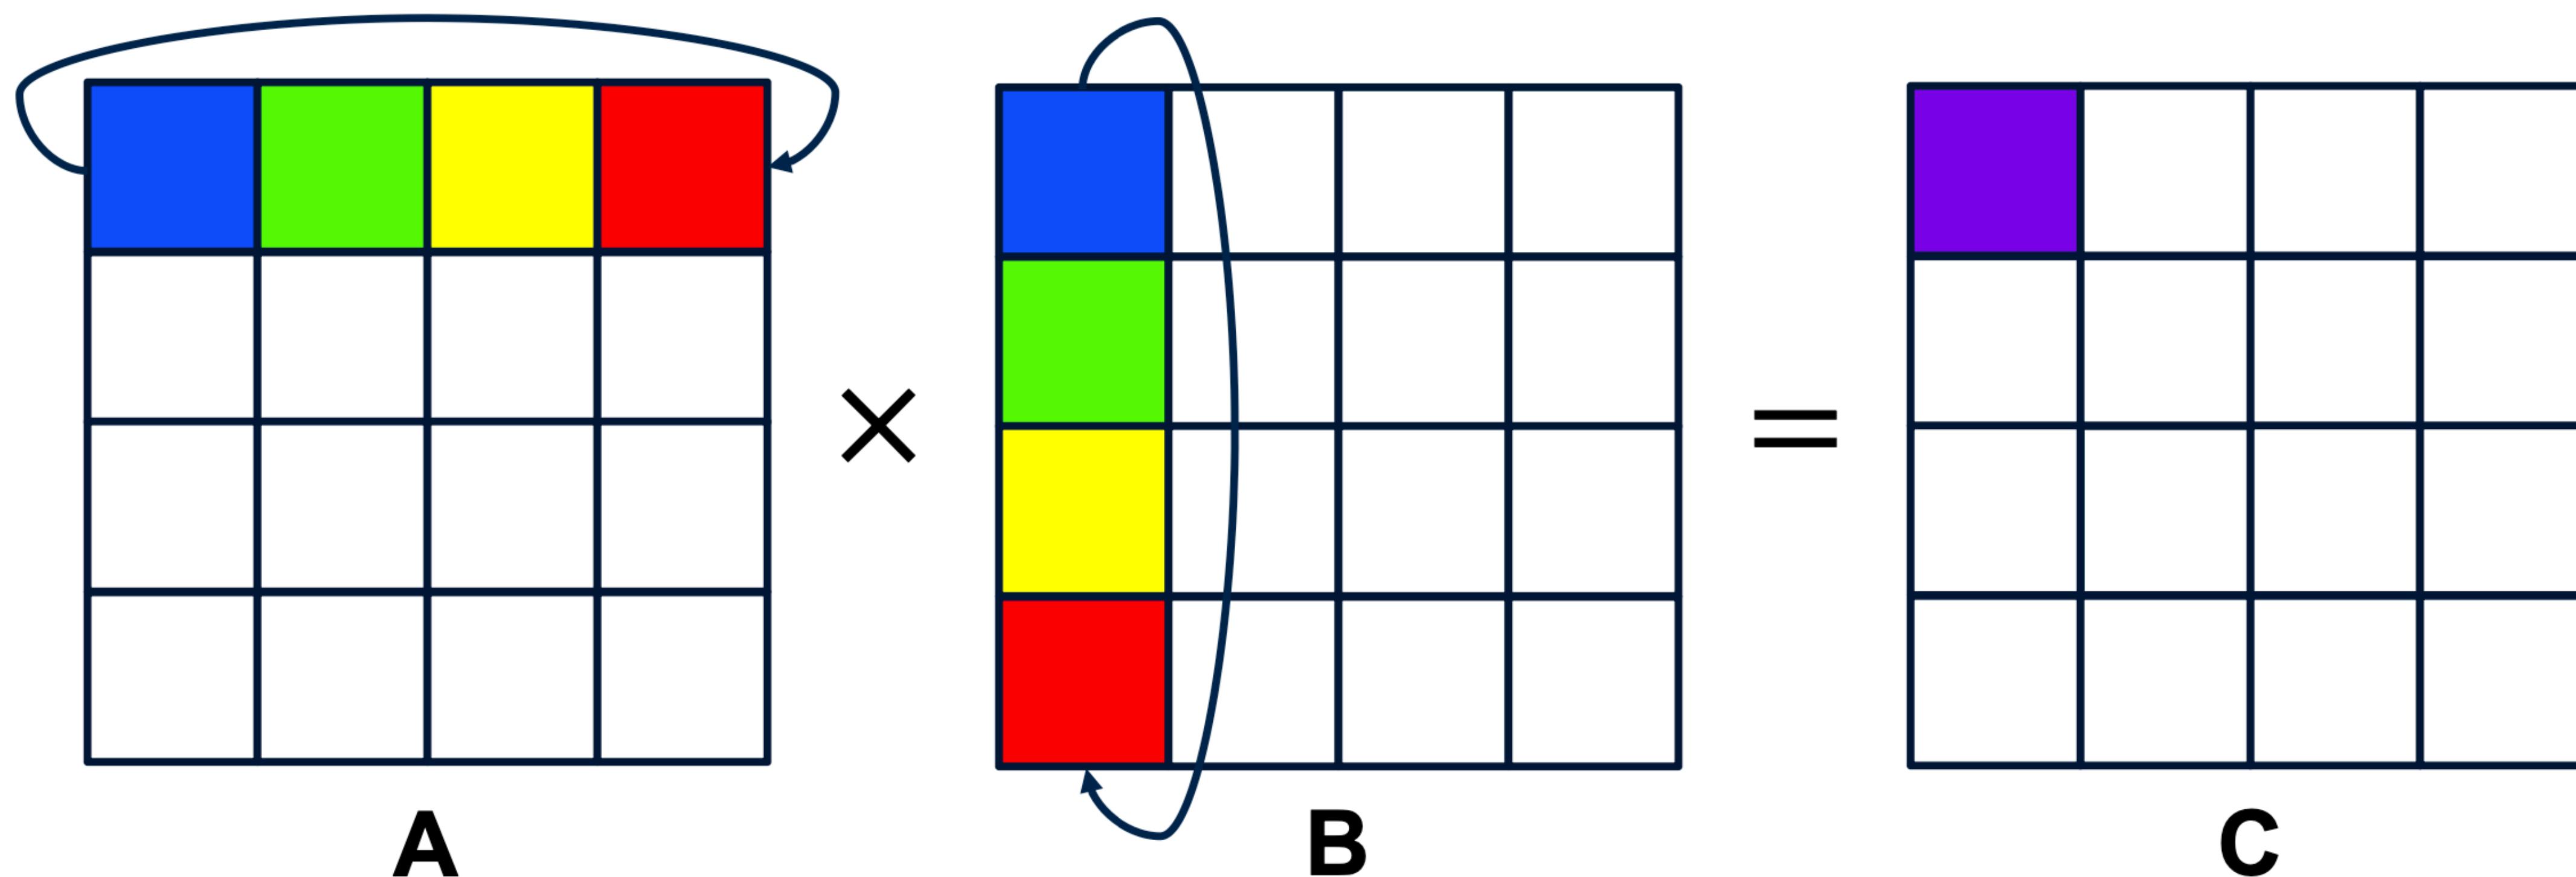
# Cannon's Algorithm: Rotations

Consider processor  $P_{0,0}$ : Step 1 - Multiply the **Red** blocks



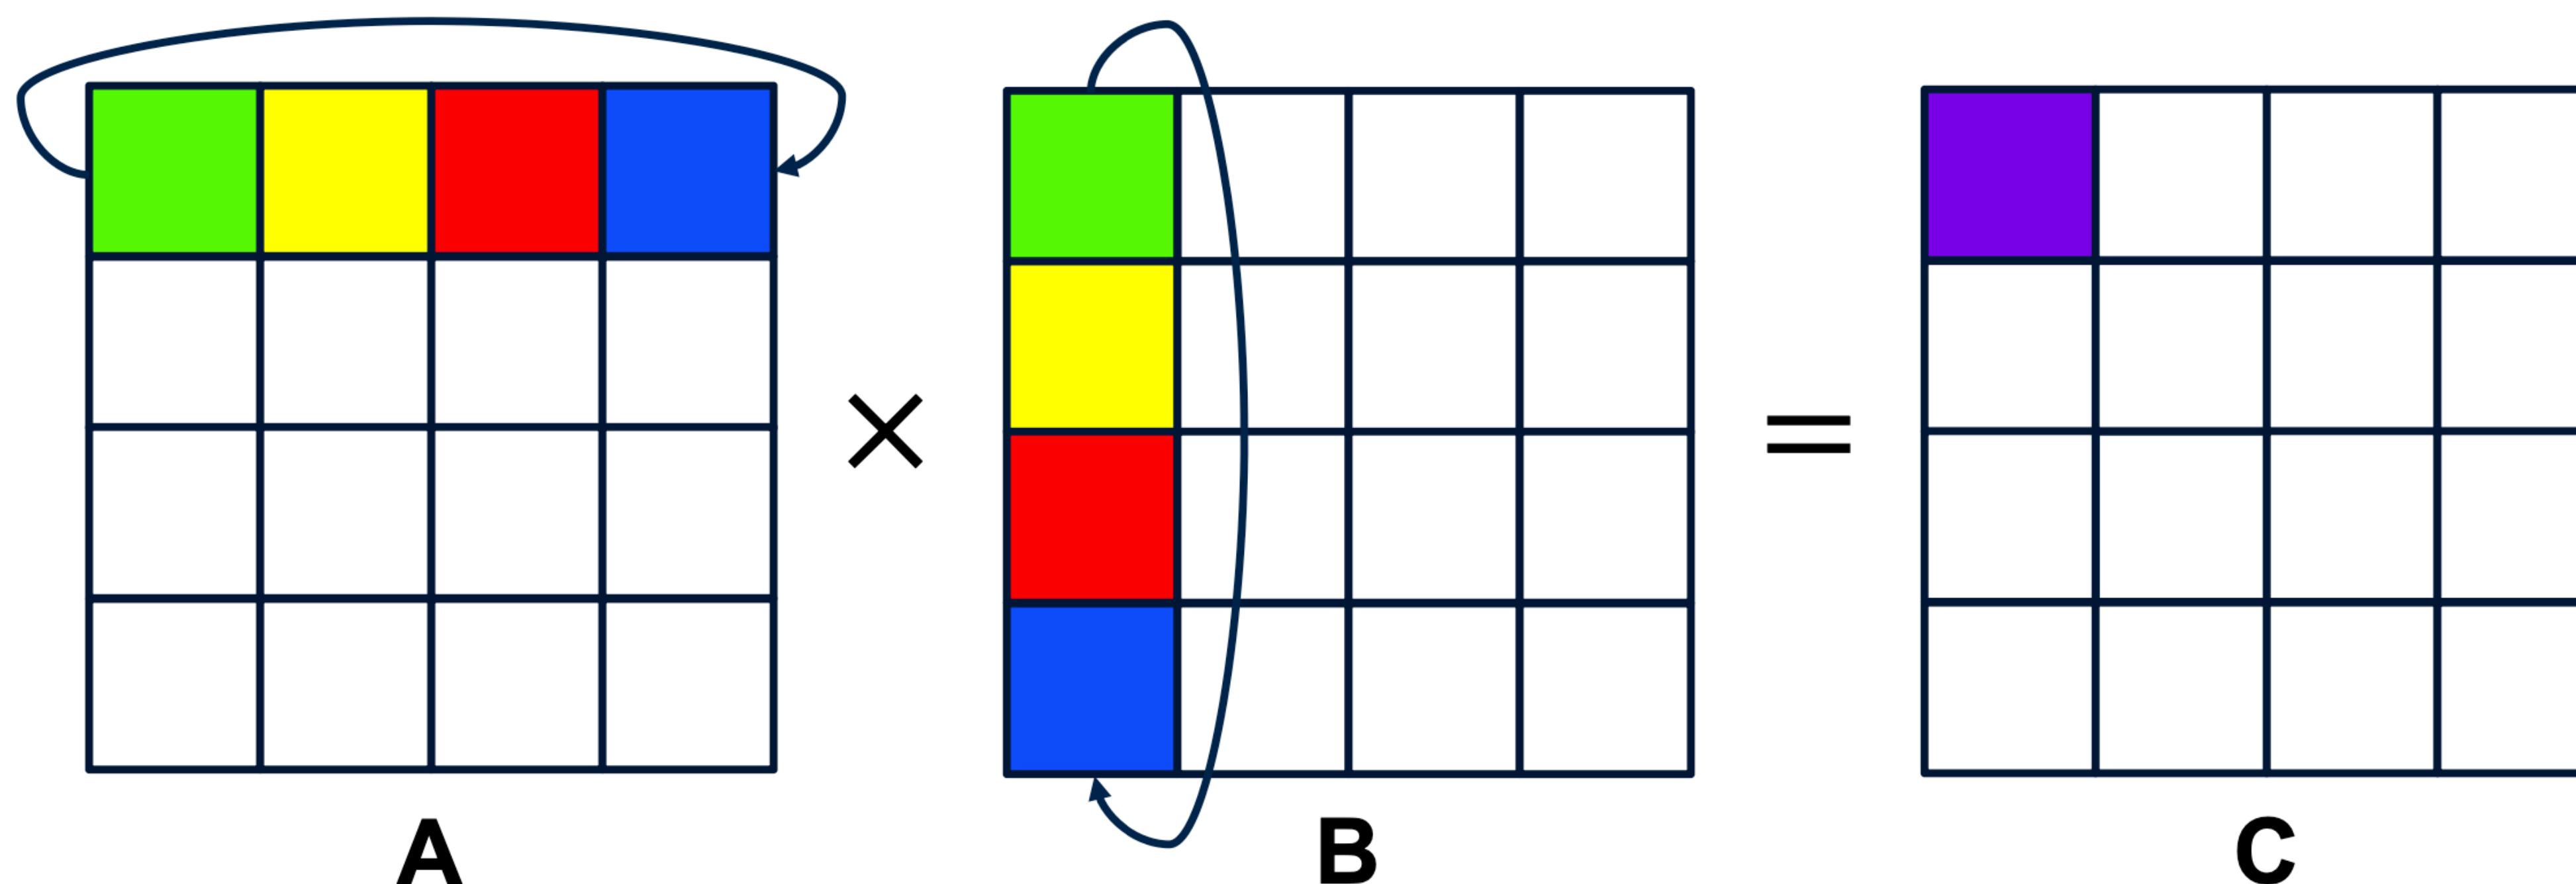
# Cannon's Algorithm: Rotations

$P_{0,0}$ : Step 2 – Left shift A, Up shift B, Multiply the **Blue** blocks



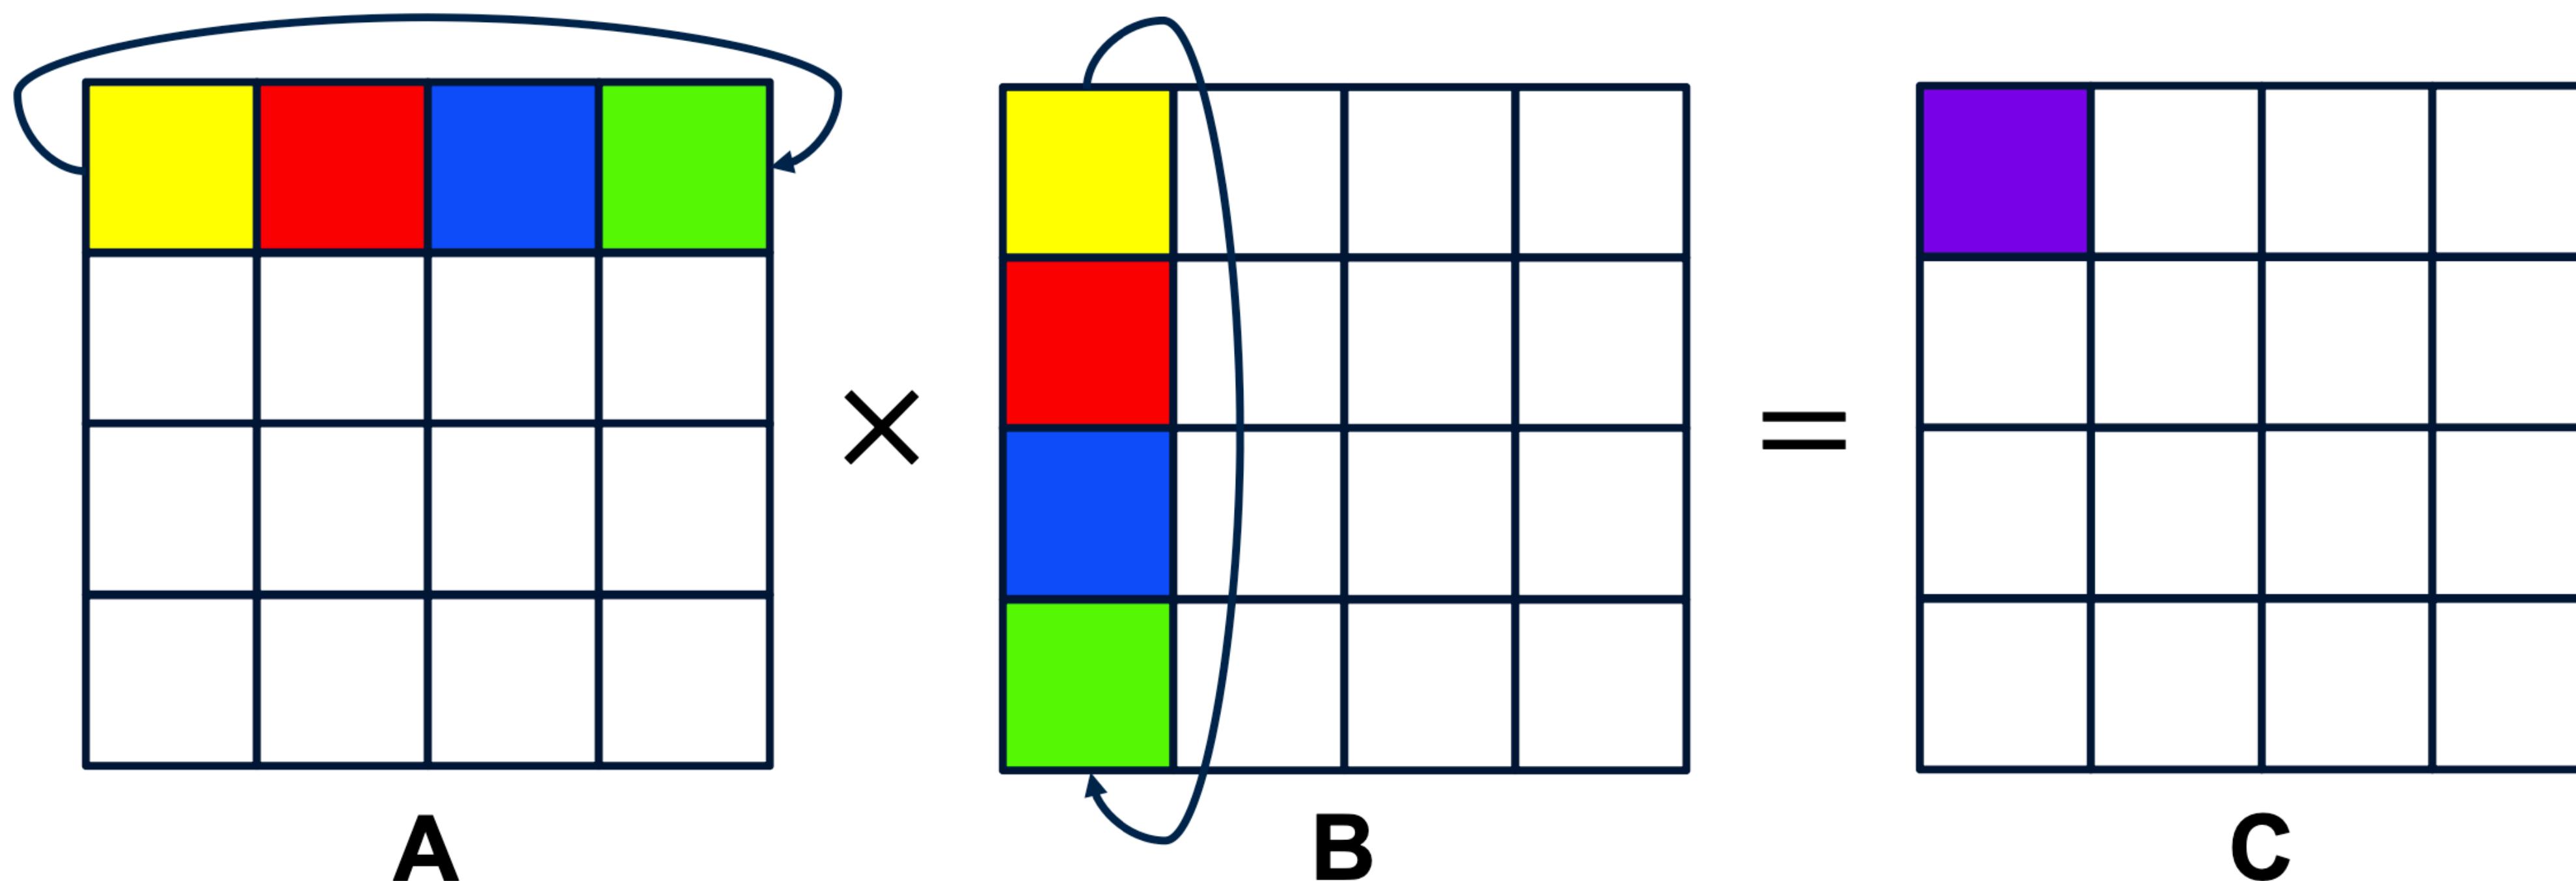
# Cannon's Algorithm: Rotations

$P_{0,0}$ : Step 3 – Left shift A, Up shift B, Multiply the **Green** blocks



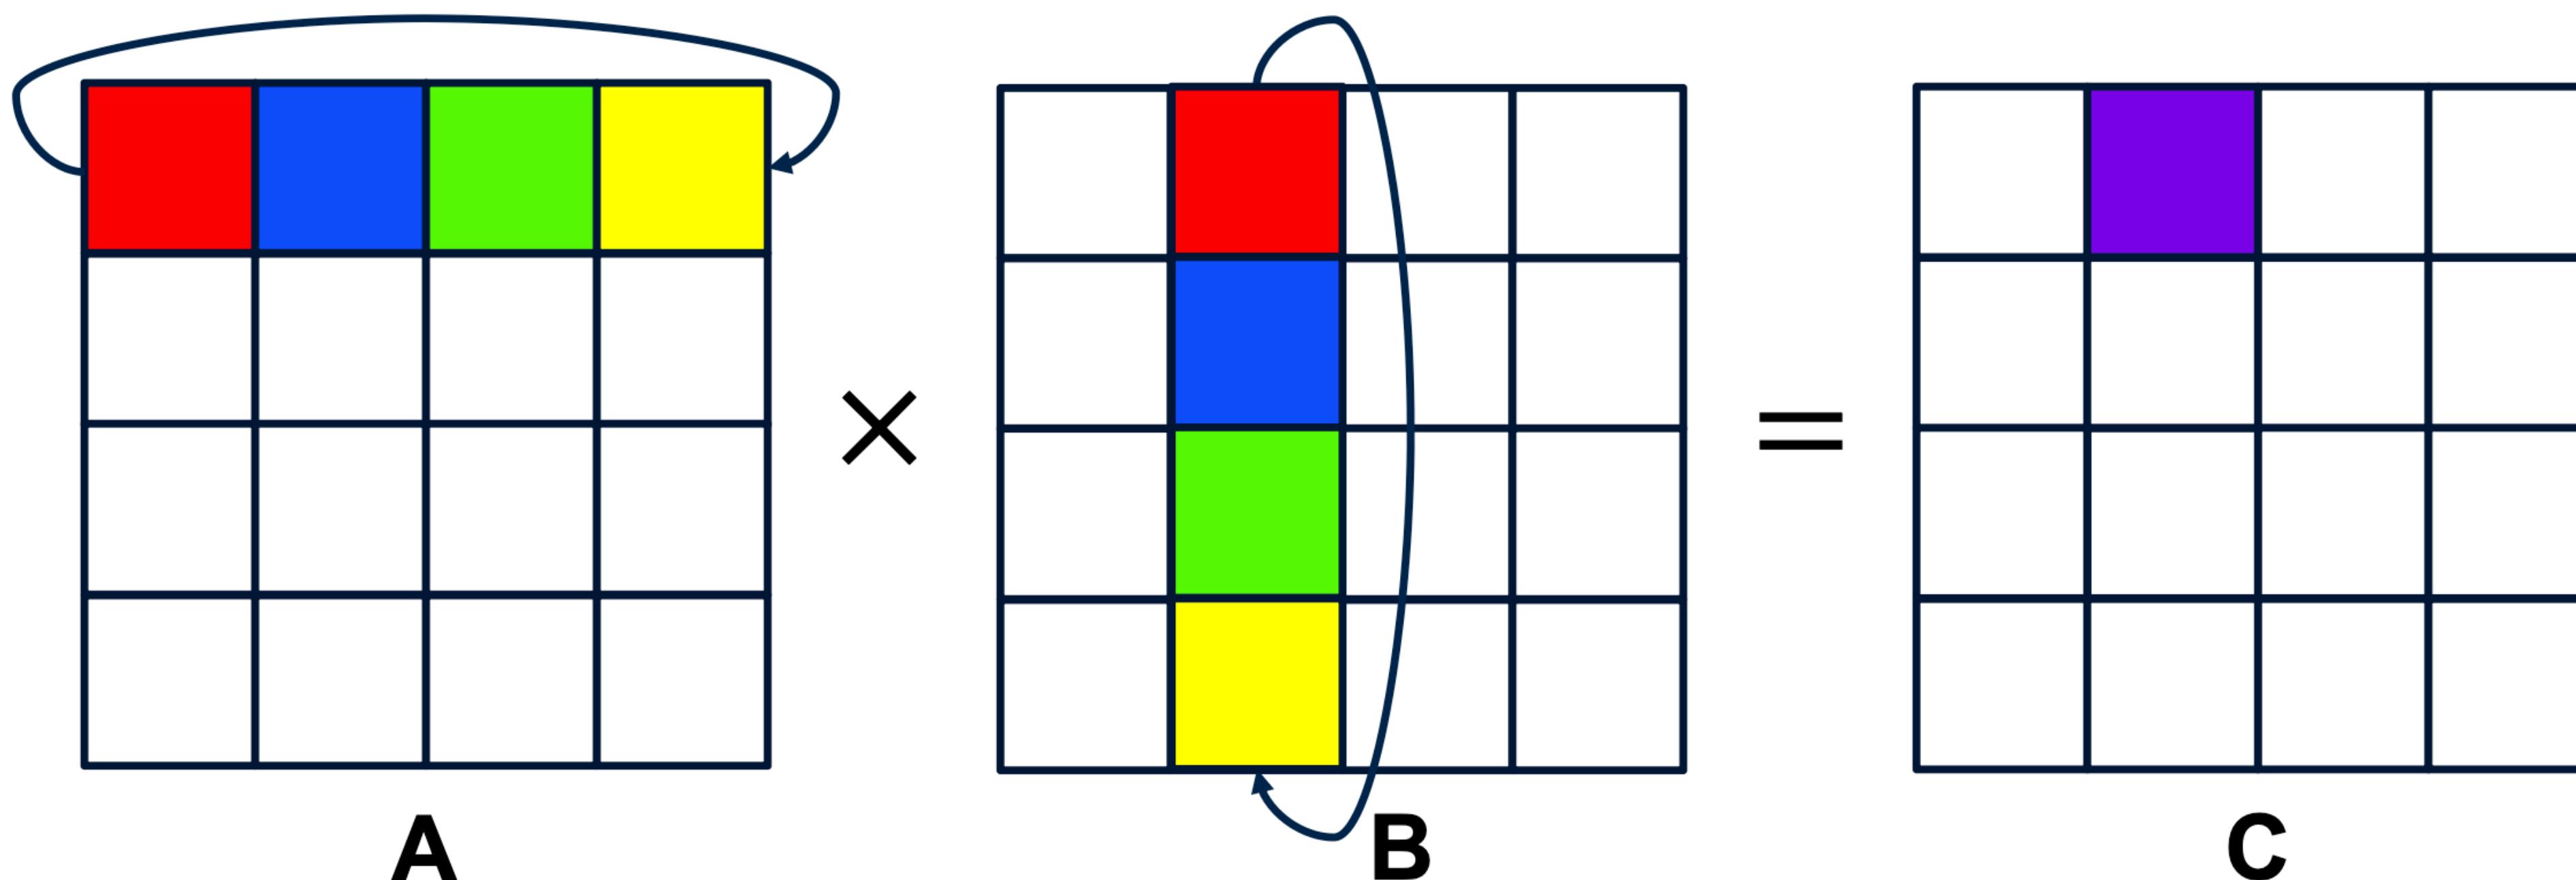
# Cannon's Algorithm: Rotations

$P_{0,0}$ : Step 4 – Left shift A, Up shift B, Multiply the Yellow blocks



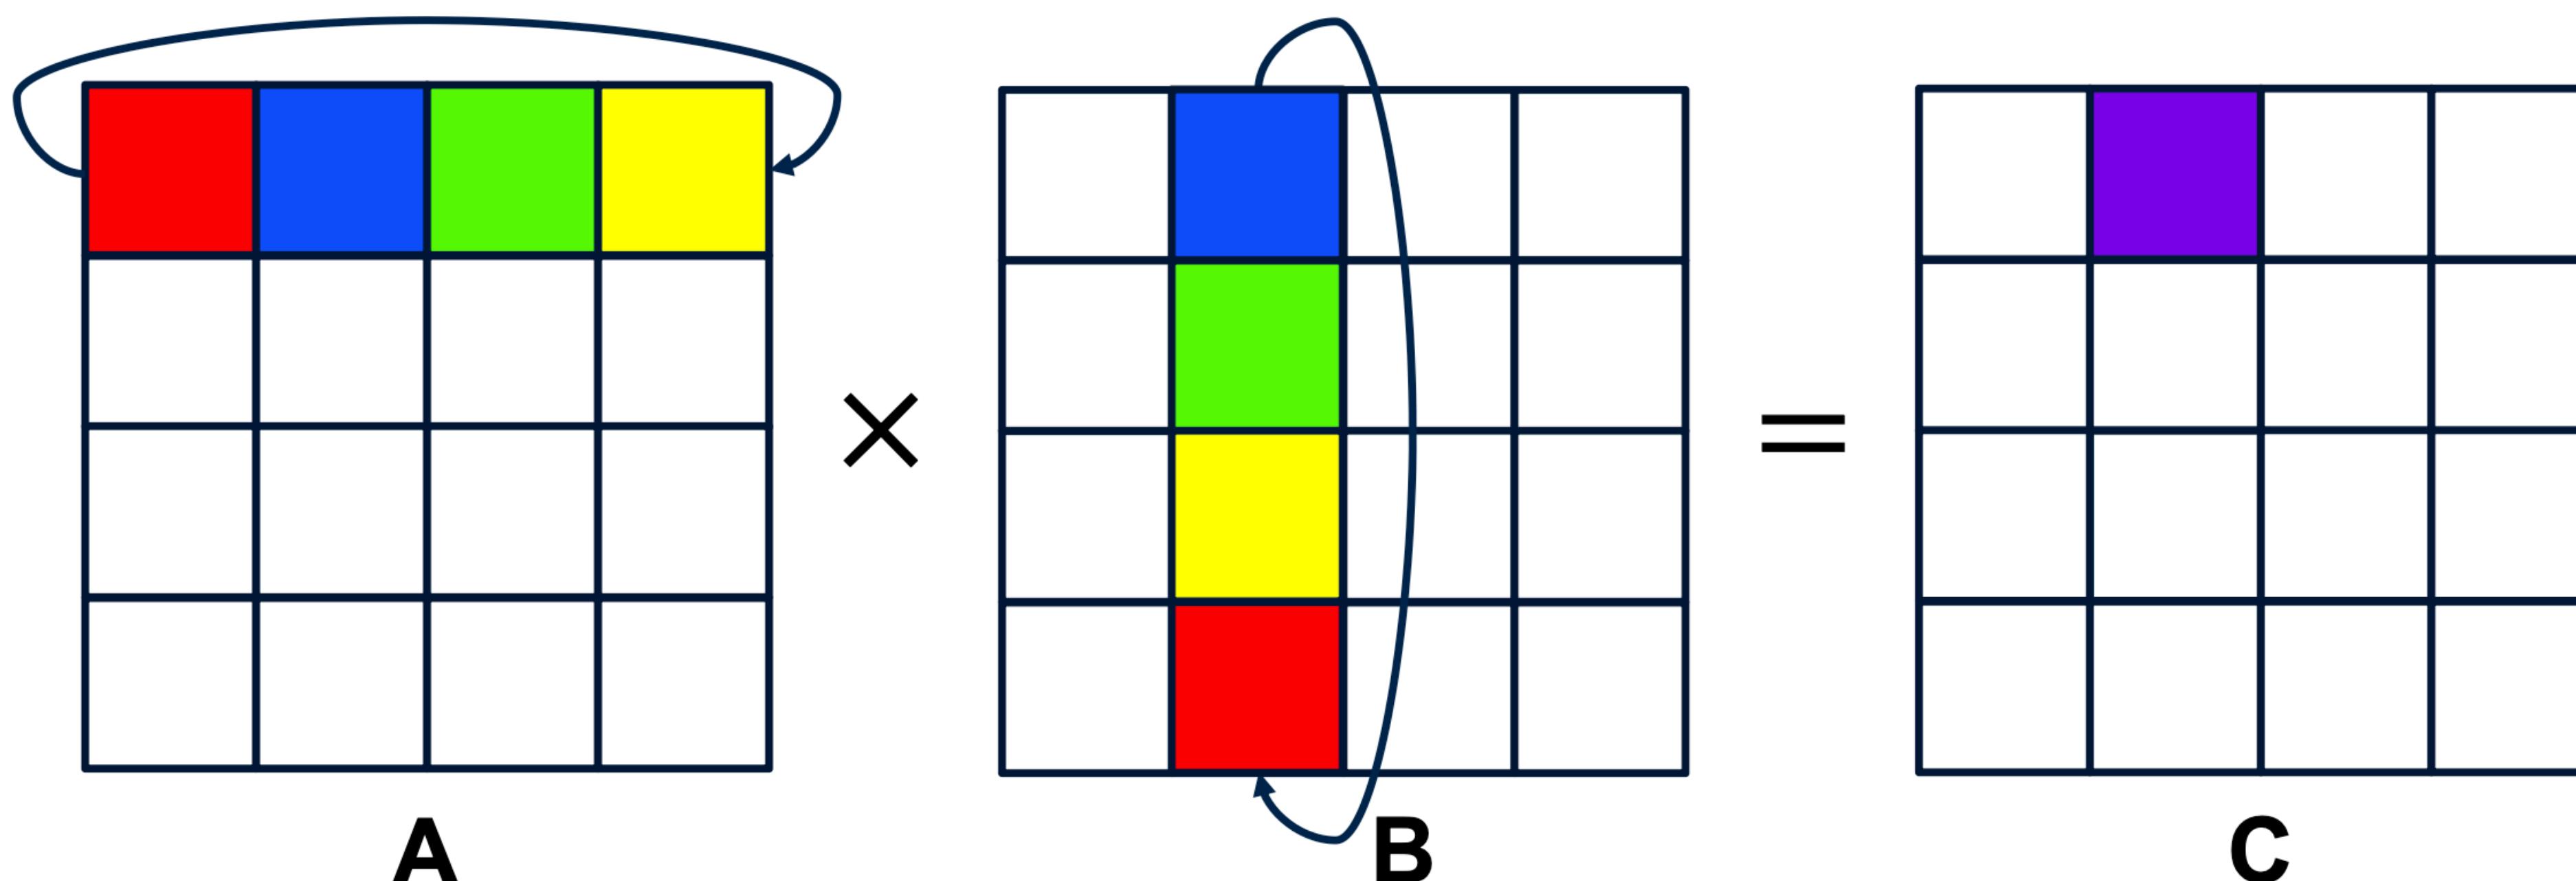
# Cannon's Algorithm: Initialization

Consider processor  $P_{0,1}$



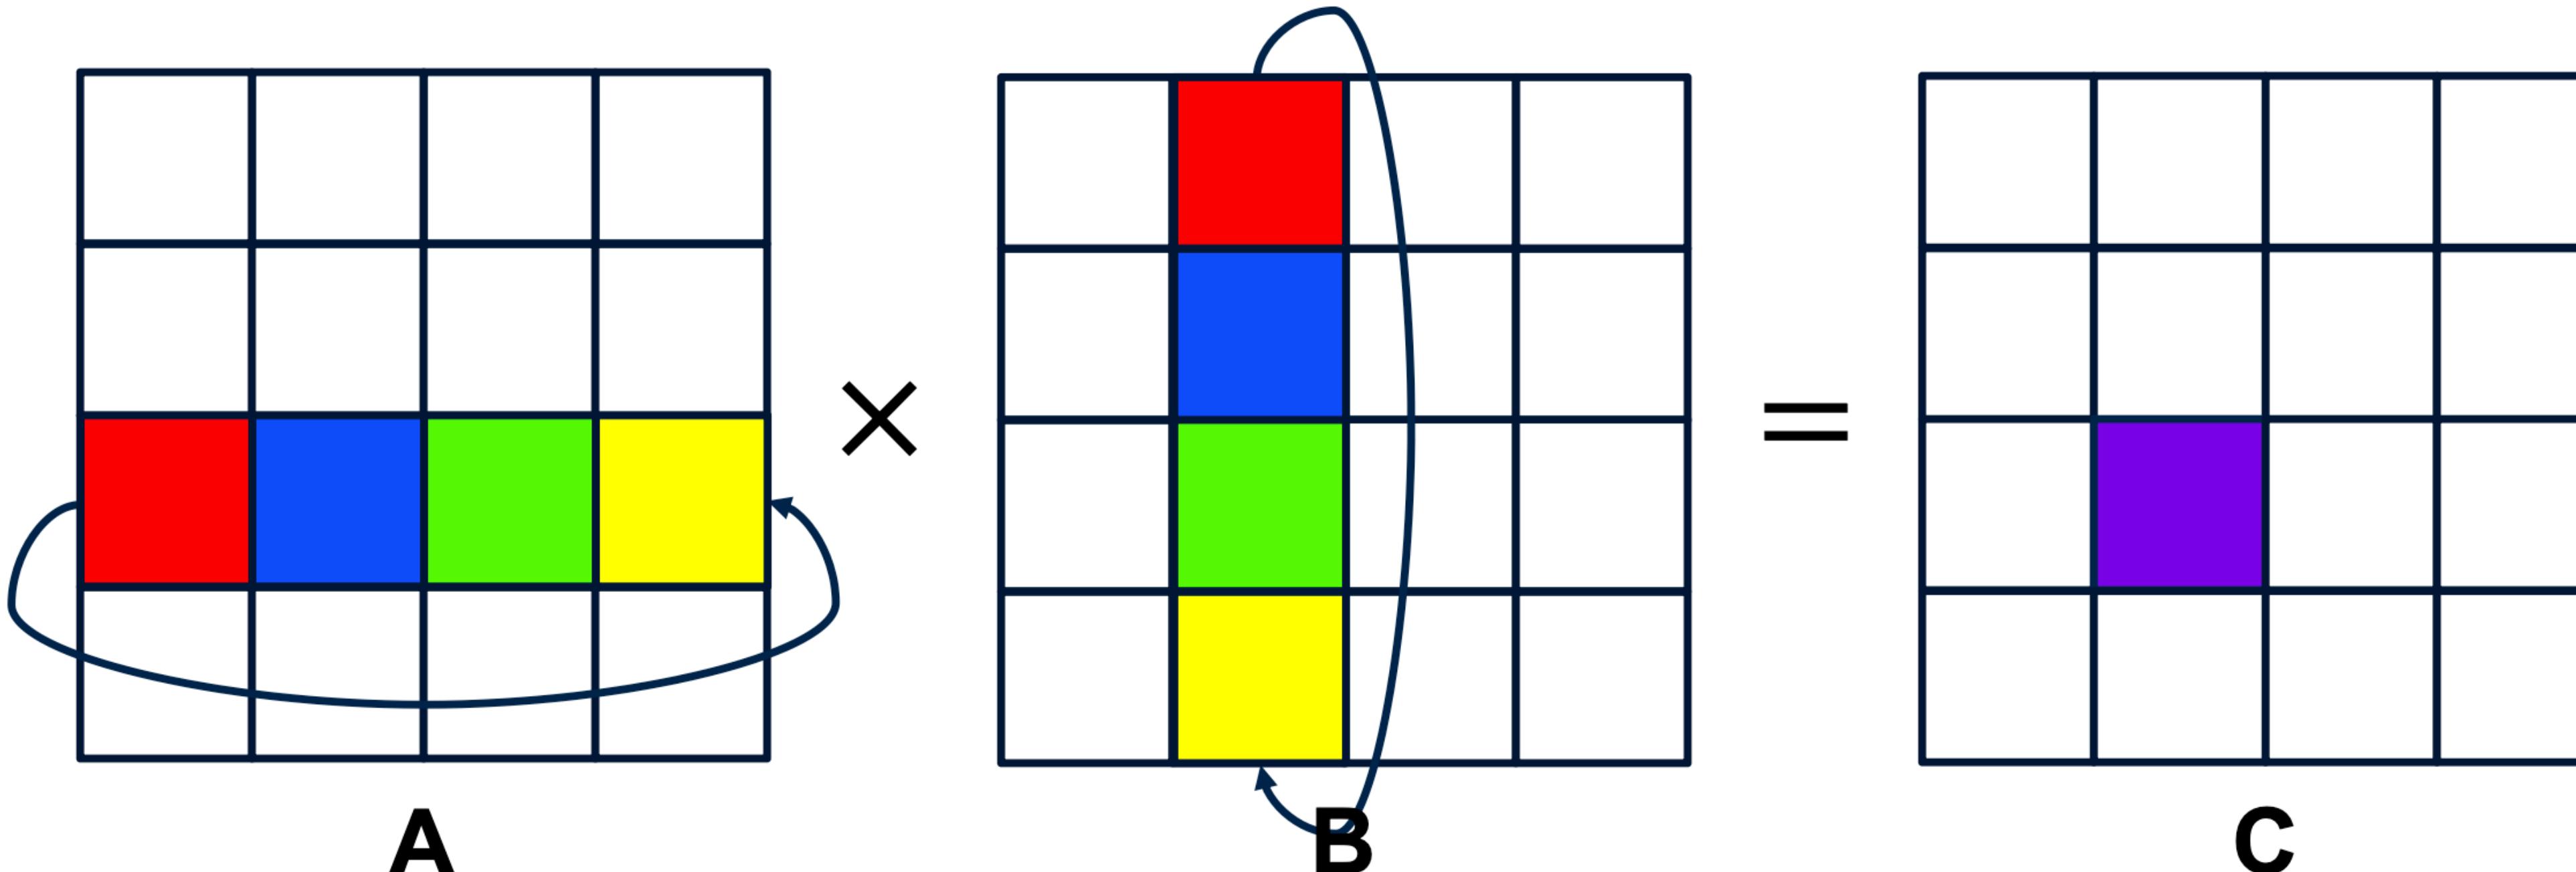
# Cannon's Algorithm: Initialization

Consider processor  $P_{0,1}$ : Align blocks by shifting B up



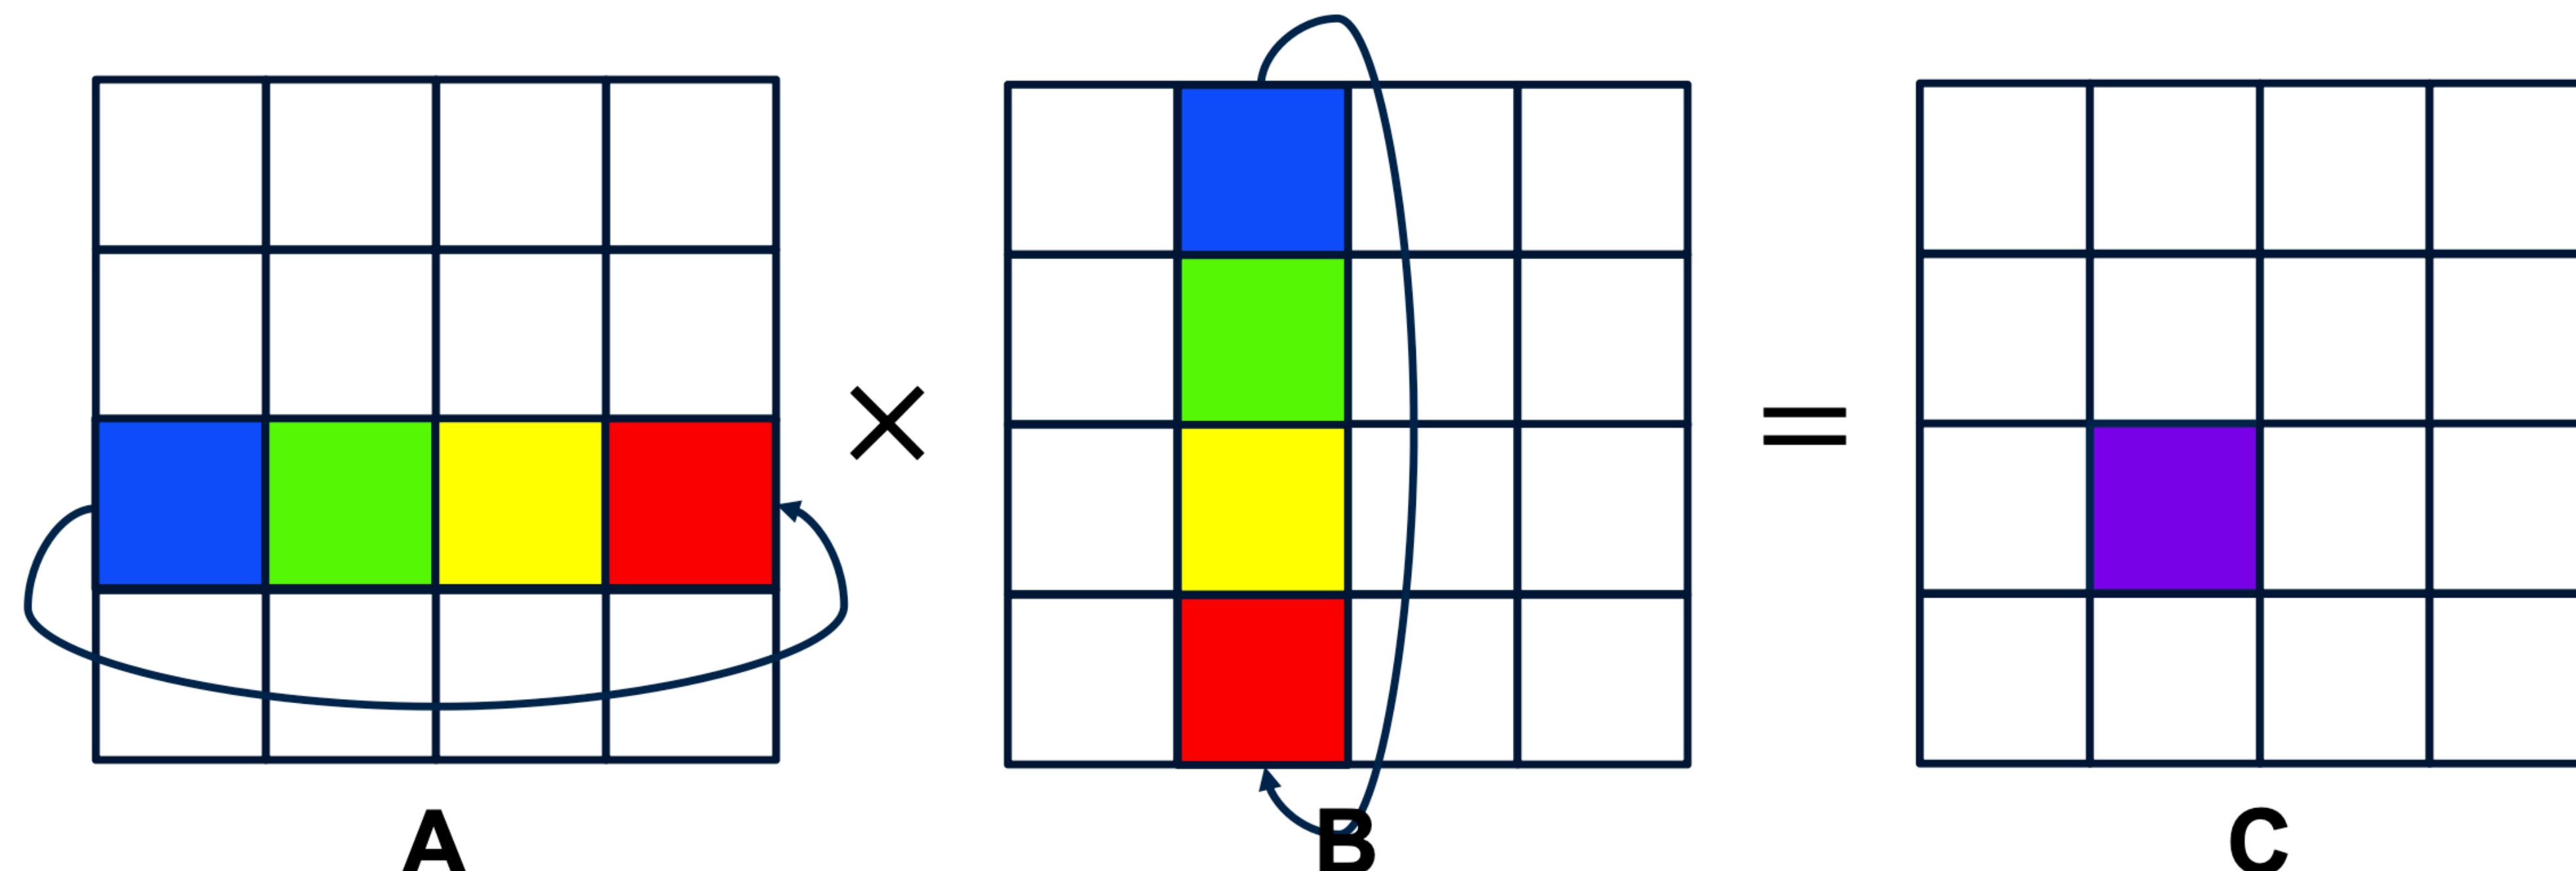
# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- We need to align the blocks!
  - In row 2 shift  $A_{ij}$  left by 2
  - In column 1 shift  $B_{ij}$  up by 1



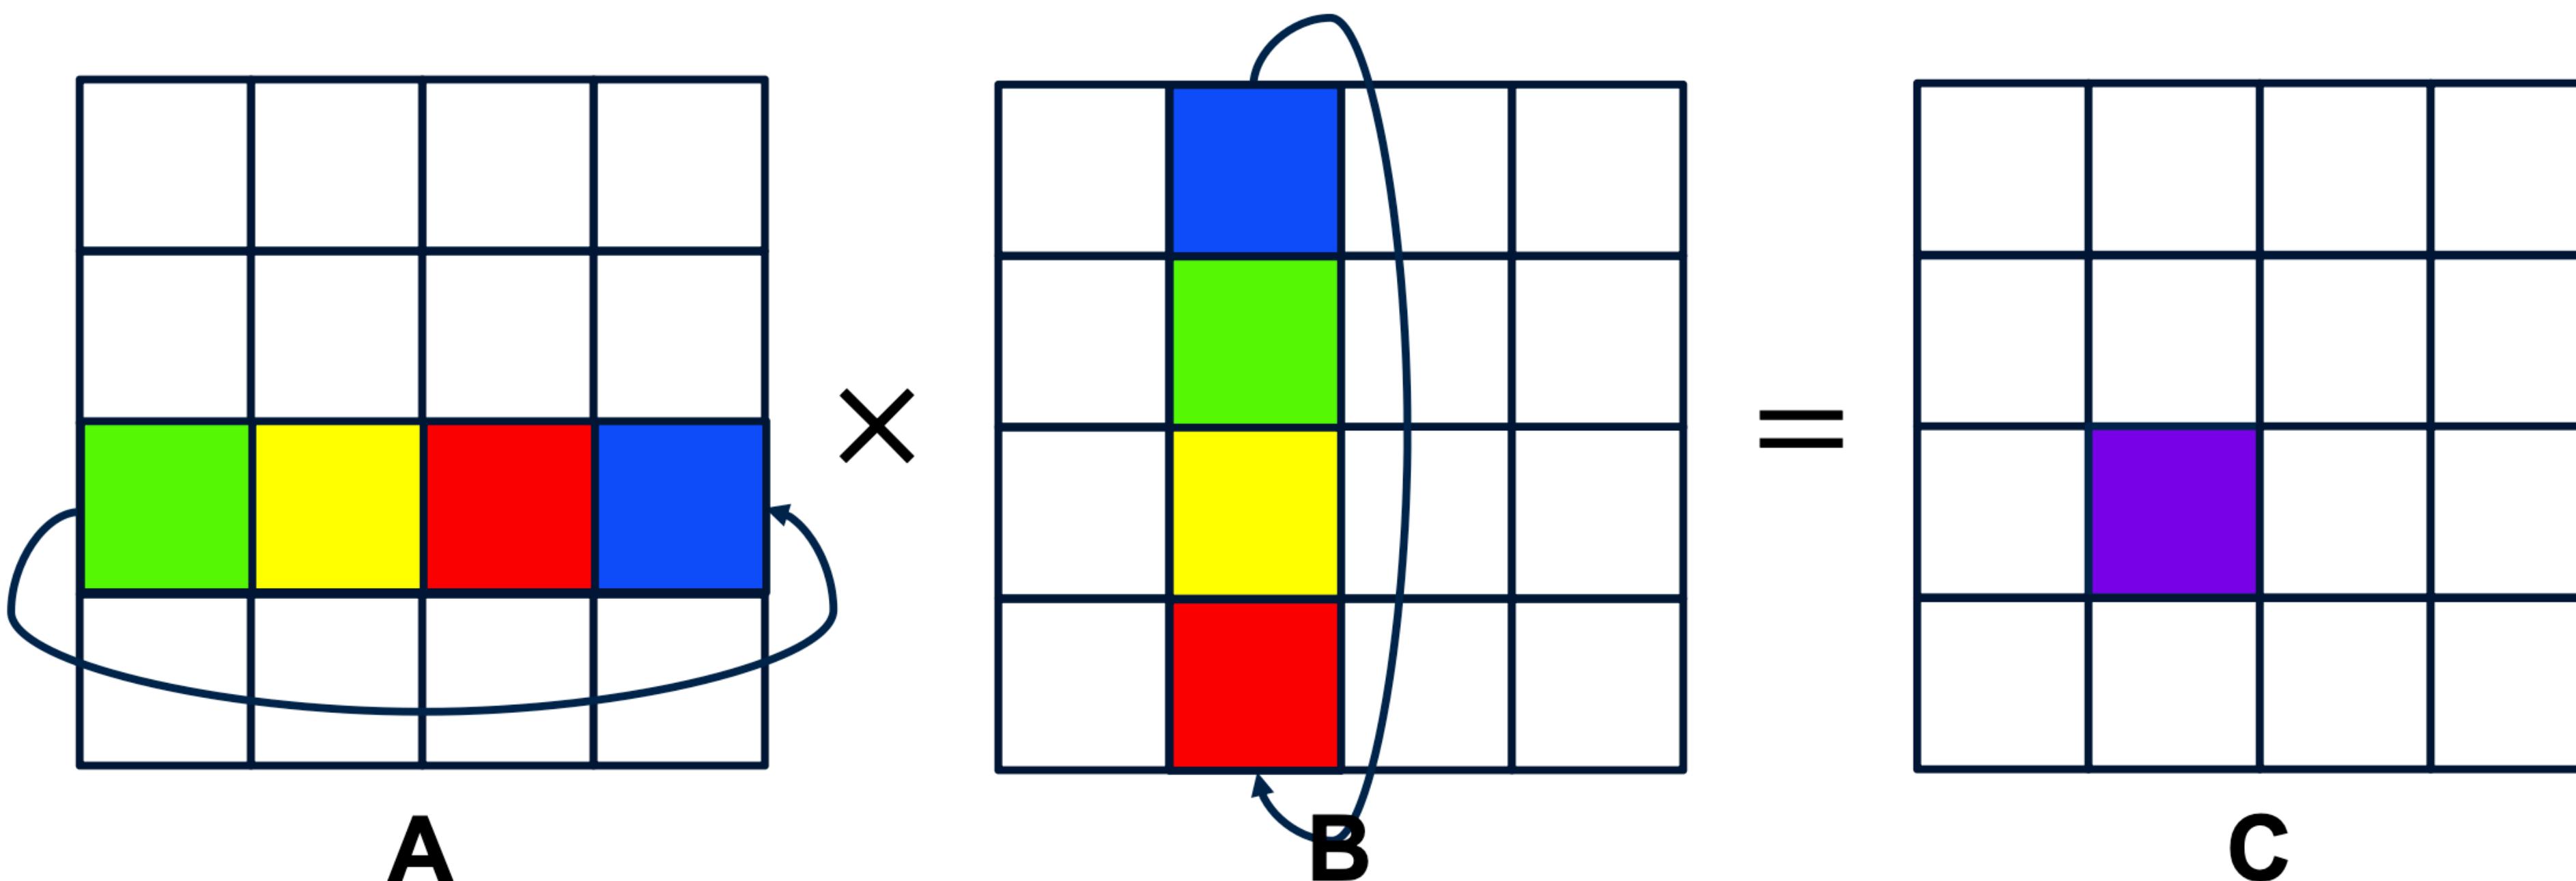
# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- Alignment Step 1



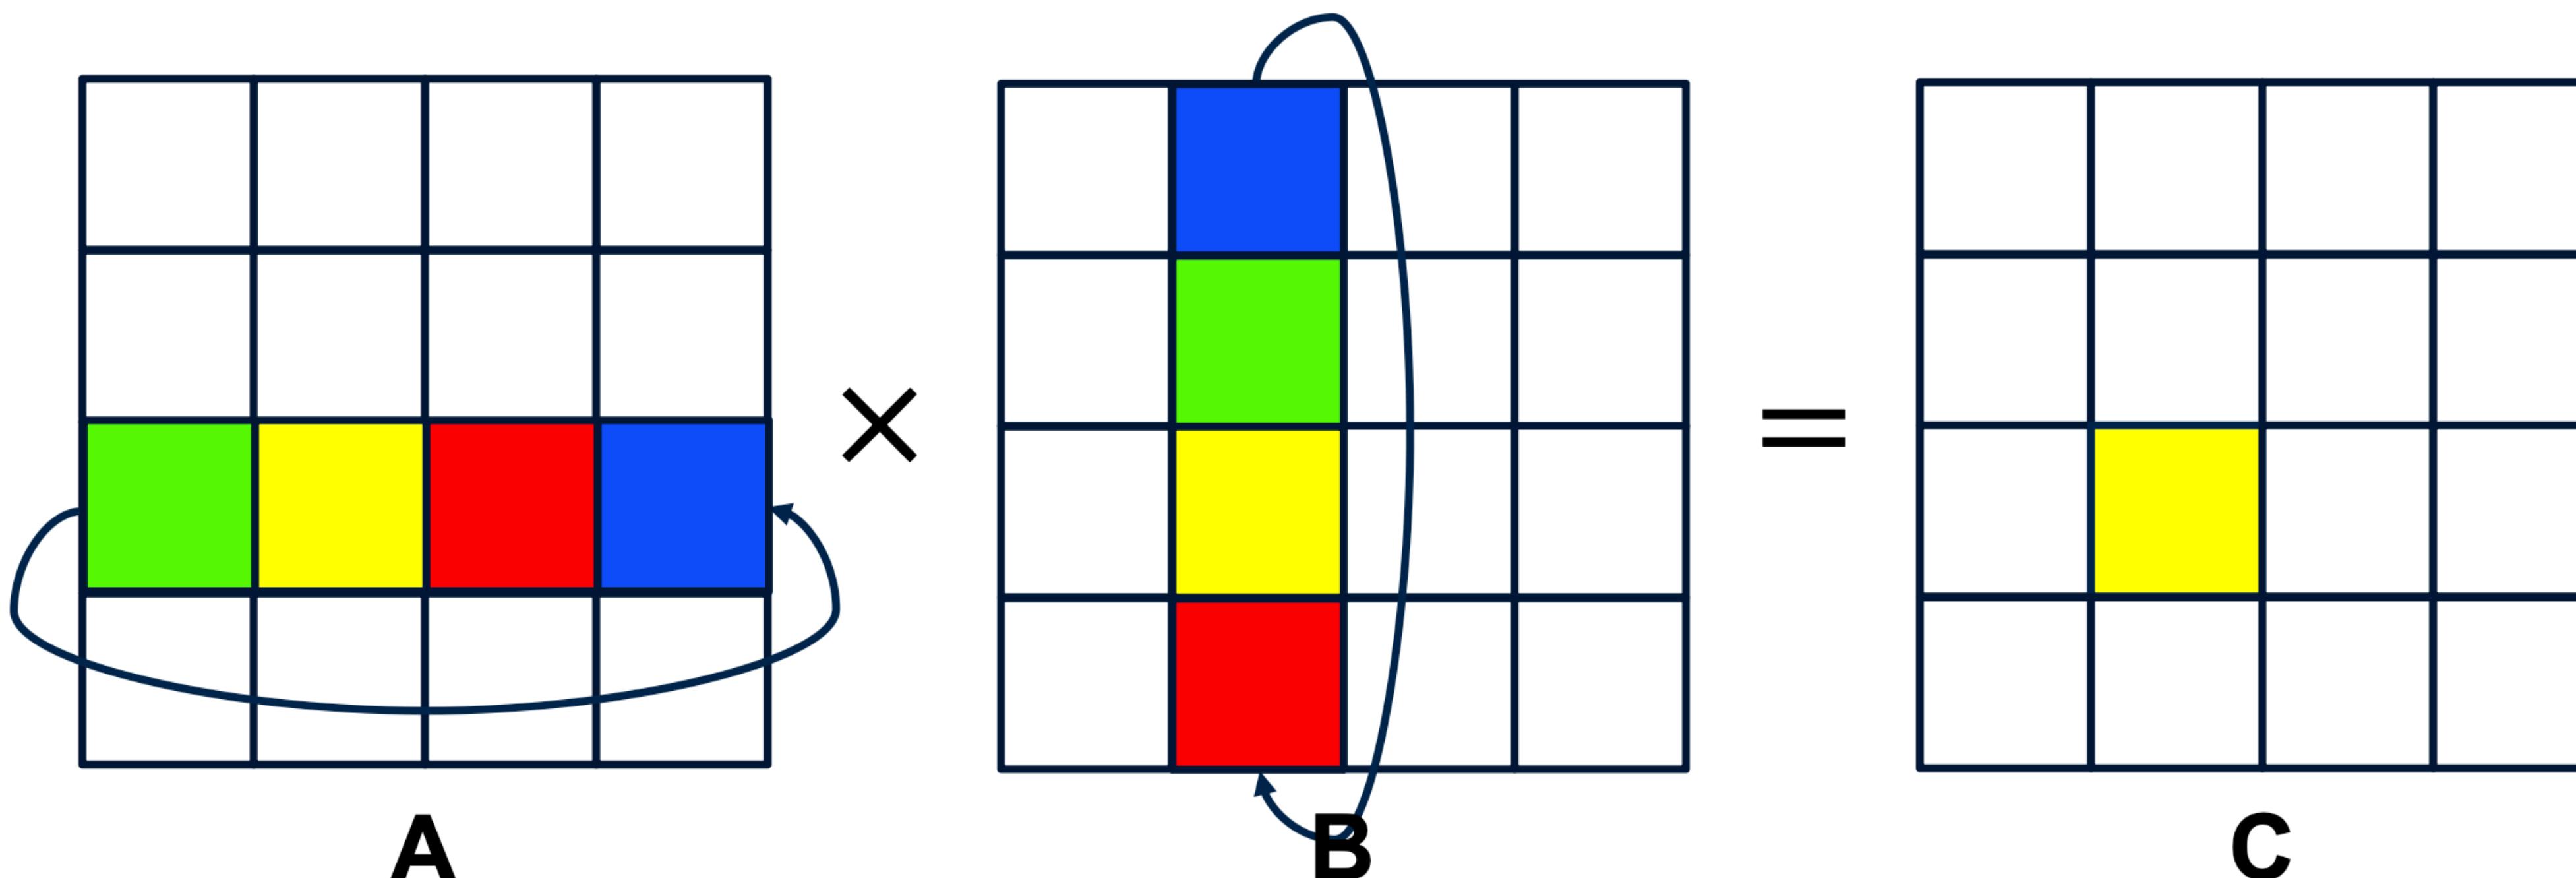
# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- Alignment Step 2



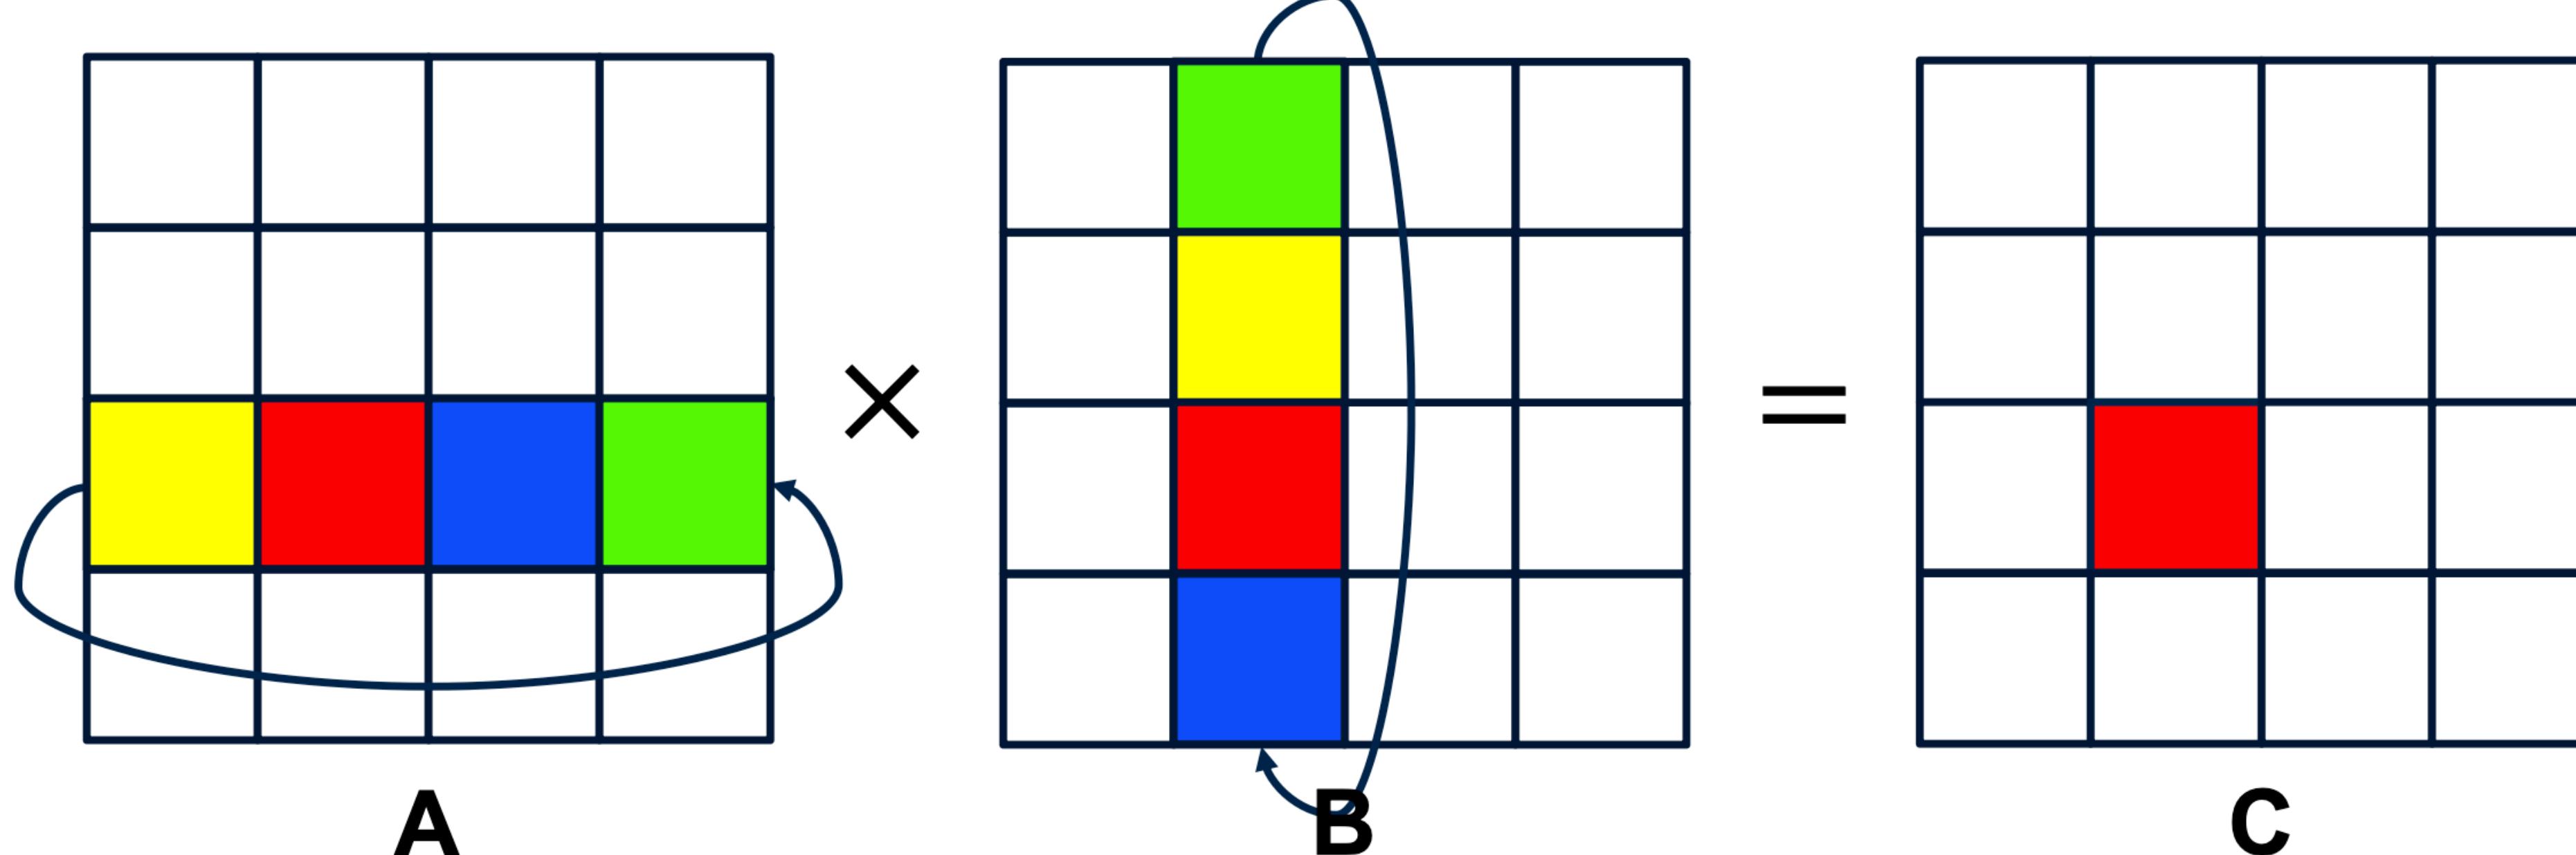
# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- Multiplication Step 1



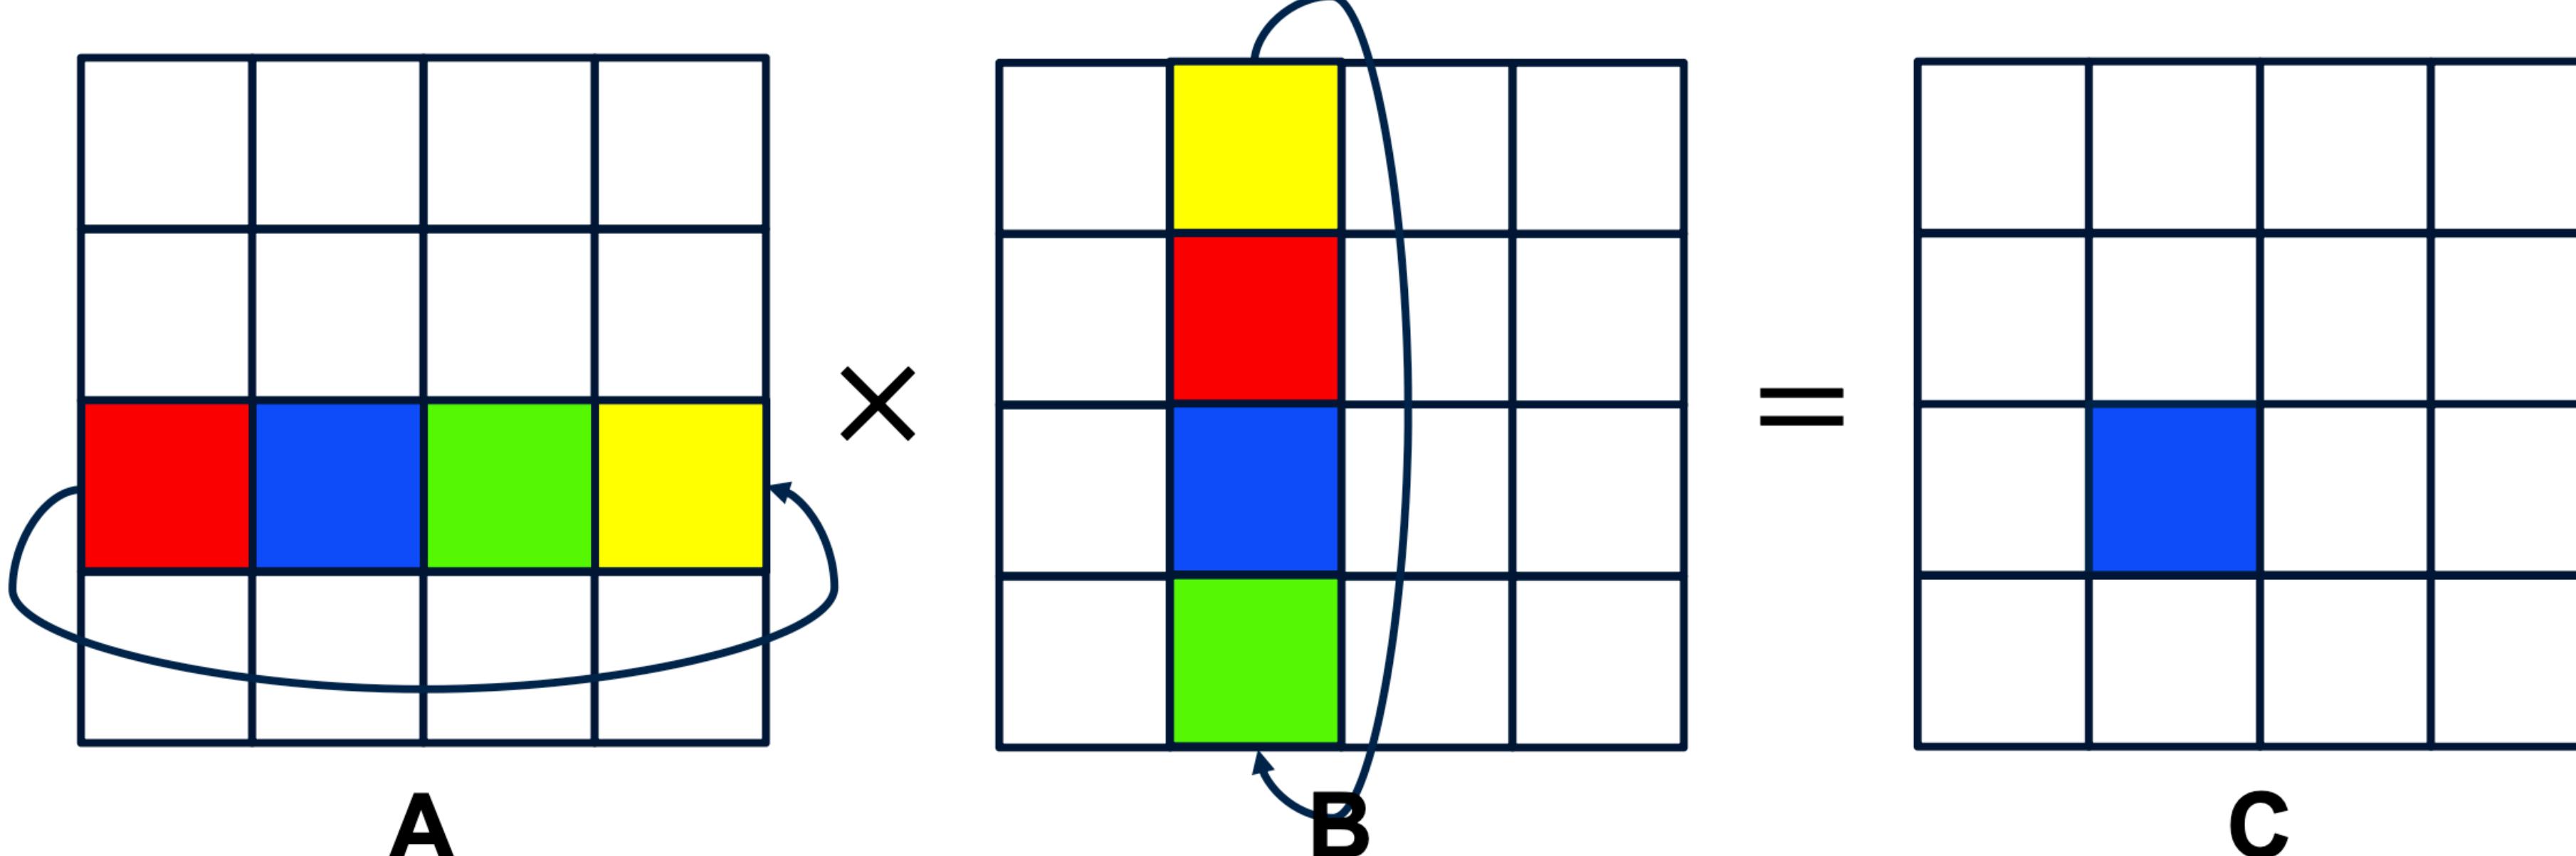
# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- Multiplication Step 2



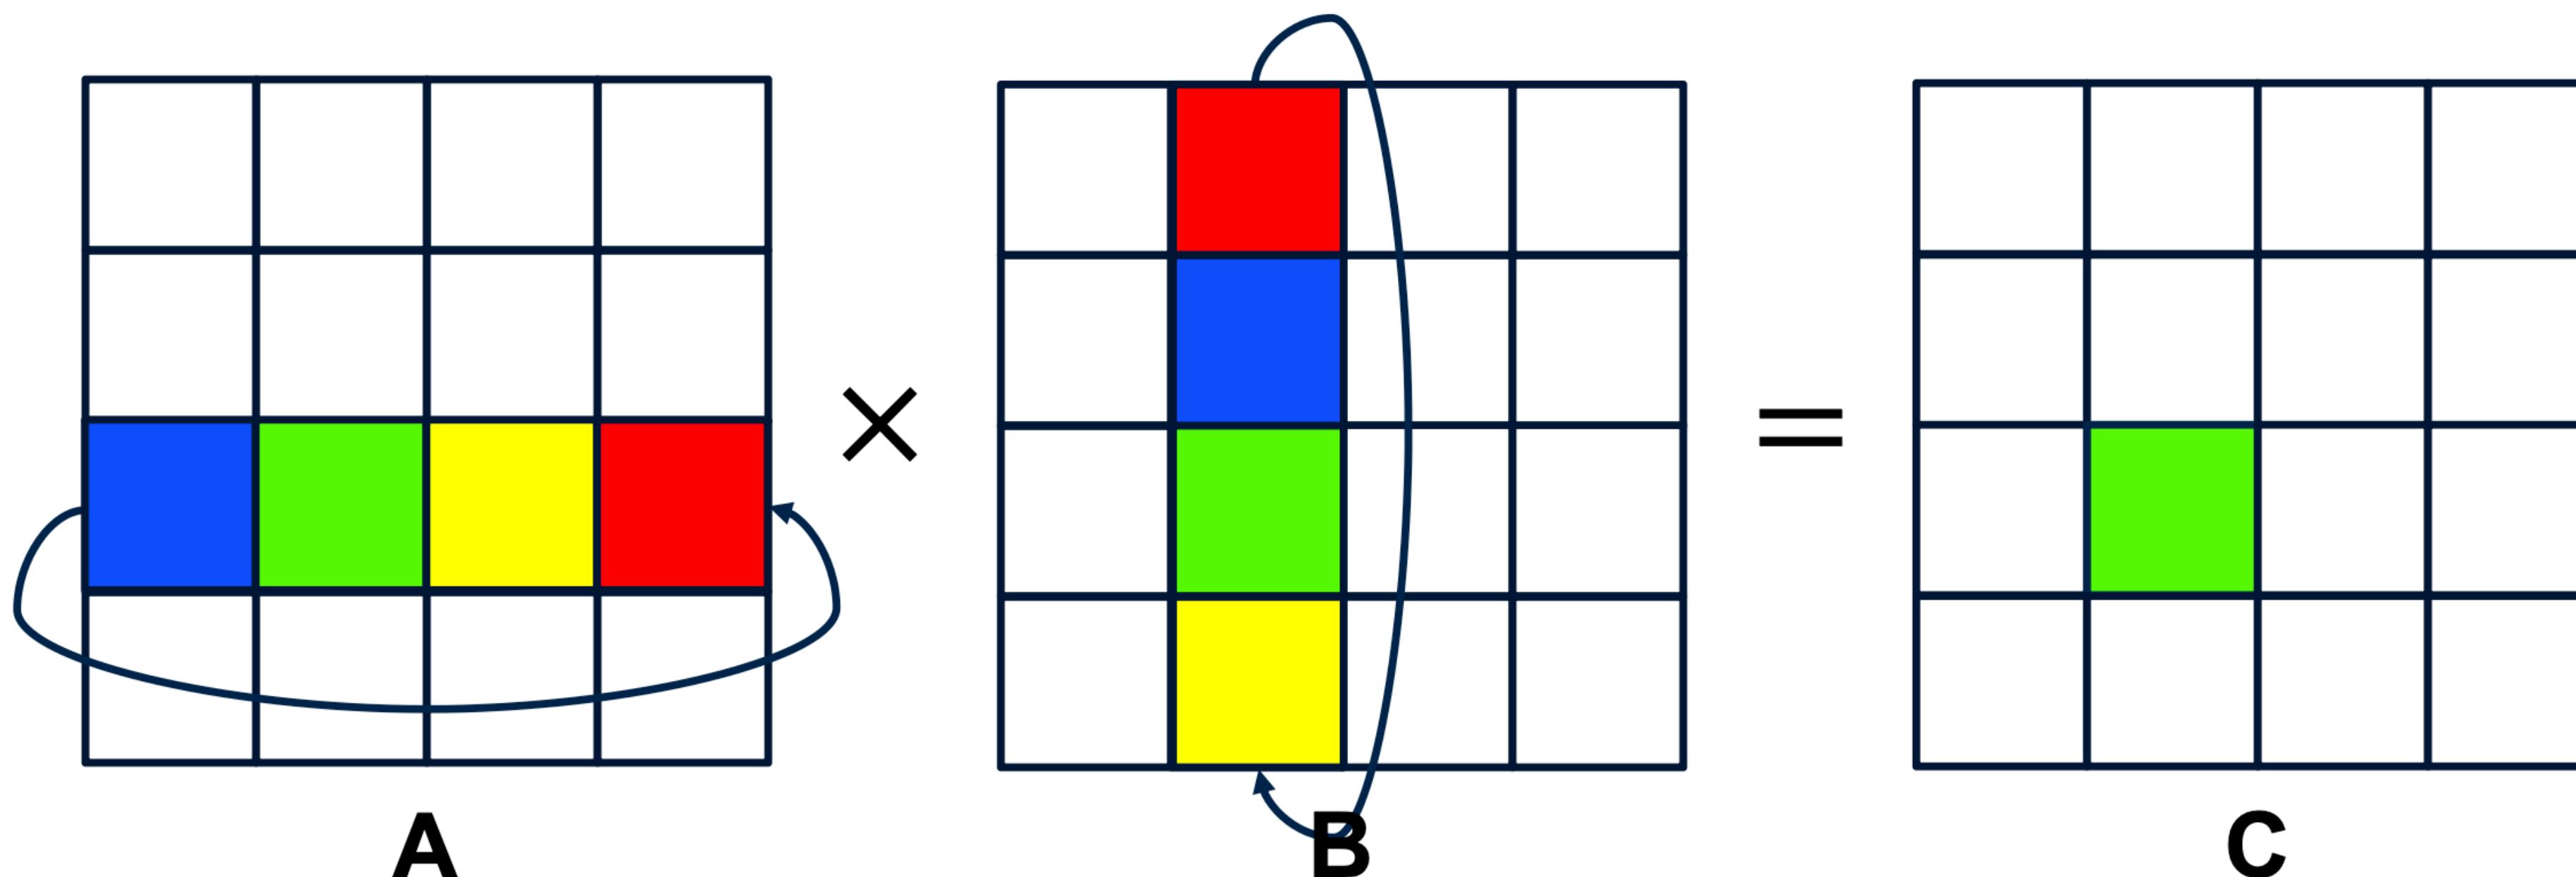
# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- Multiplication Step 3



# Cannon's Algorithm

- Consider processor  $P_{2,1}$
- Multiplication Step 4



# Cannon's Algorithm: Initial Alignment

Consider processor  $P_{i,j}$

- Aligning blocks:
  - Shift  $A_{i,j}$  left by  $i$ 
    - $A_{i,(j+i)}$
  - Shift  $B_{i,j}$  up by  $j$ 
    - $B_{(i+j),j}$
  - $P_{i,j}$  will have  $A_{i,(j+i)} \bmod \sqrt{p}$  and  $B_{(i+j),j} \bmod \sqrt{p}, j$

Line up inner dimension

# Analysis of Cannon's Algorithm

- In the alignment step,
  - maximum distance over which a block shifts is  $\sqrt{p} - 1$ ,
  - two shift operations require a total of  $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$  time.

# Analysis of Cannon's Algorithm

- In the alignment step,
  - maximum distance over which a block shifts is  $\sqrt{p} - 1$ ,
  - two shift operations require a total of  $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$  time.
- Compute-and-shift phase has  $\sqrt{p}$  steps
  - Computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:  $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
  - Communication requires  $\sqrt{p}$  shifts  $O(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p})$

# Analysis of Cannon's Algorithm

- In the alignment step,
  - maximum distance over which a block shifts is  $\sqrt{p} - 1$ ,
  - two shift operations require a total of  $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$  time.
- Compute-and-shift phase has  $\sqrt{p}$  steps
  - Computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:  $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
  - Communication requires  $\sqrt{p}$  shifts  $O(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p})$
- Total Time:  $O\left(\frac{n^3}{p} + \tau\sqrt{p} + \mu\frac{n^2}{\sqrt{p}}\right)$

# Analysis of Cannon's Algorithm

- In the alignment step,
  - maximum distance over which a block shifts is  $\sqrt{p} - 1$ ,
  - two shift operations require a total of  $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$  time.
- Compute-and-shift phase has  $\sqrt{p}$  steps
  - Computation requires  $\sqrt{p}$  multiplications of  $(n/\sqrt{p}) \times (n/\sqrt{p})$  sized submatrices:  $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
  - Communication requires  $\sqrt{p}$  shifts  $O(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p})$
- Total Time:  $O\left(\frac{n^3}{p} + \tau\sqrt{p} + \mu\frac{n^2}{\sqrt{p}}\right)$
- **More messages, but Cannon's Algorithm is memory optimal.**

```
1 MatrixMatrixMultiply(int n, double *a, double *b, double *c,
2                         MPI_Comm comm)
3 {
4     int i;
5     int nlocal;
6     int npes, dims[2], periods[2];
7     int myrank, my2drank, mycoords[2];
8     int uprank, downrank, leftrank, rightrank, coords[2];
9     int shiftsource, shiftdest;
10    MPI_Status status;
11    MPI_Comm comm_2d;
12
13    /* Get the communicator related information */
14    MPI_Comm_size(comm, &npes);
15    MPI_Comm_rank(comm, &myrank);
16
17    /* Set up the Cartesian topology */
18    dims[0] = dims[1] = sqrt(npes);
19
20    /* Set the periods for wraparound connections */
21    periods[0] = periods[1] = 1;
22
23    /* Create the Cartesian topology, with rank reordering */
24    MPI_Cart_create(comm, 2, dims, periods, 1, &comm_2d);
25
26    /* Get the rank and coordinates with respect to the new topology */
27    MPI_Comm_rank(comm_2d, &my2drank);
28    MPI_Cart_coords(comm_2d, my2drank, 2, mycoords);
29
30    /* Compute ranks of the up and left shifts */
31    MPI_Cart_shift(comm_2d, 0, -1, &rightrank, &leftrank);
32    MPI_Cart_shift(comm_2d, 1, -1, &downrank, &uprank);
33
34    /* Determine the dimension of the local matrix block */
35    nlocal = n/dims[0];
36
37    /* Perform the initial matrix alignment. First for A and then for B */
38    MPI_Cart_shift(comm_2d, 0, -mycoords[0], &shiftsource, &shiftdest);
39    MPI_Sendrecv_replace(a, nlocal*nlocal, MPI_DOUBLE, shiftdest,
40                          1, shiftsource, 1, comm_2d, &status);
```

```
41
42     MPI_Cart_shift(comm_2d, 1, -mycoords[1], &shiftsource, &shiftdest);
43     MPI_Sendrecv_replace(b, nlocal*nlocal, MPI_DOUBLE,
44                           shiftdest, 1, shiftsource, 1, comm_2d, &status);
45
46 /* Get into the main computation loop */
47 for (i=0; i<dims[0]; i++) {
48     MatrixMultiply(nlocal, a, b, c); /* c = c + a*b */
49
50 /* Shift matrix a left by one */
51 MPI_Sendrecv_replace(a, nlocal*nlocal, MPI_DOUBLE,
52                      leftrank, 1, rightrank, 1, comm_2d, &status);
53
54 /* Shift matrix b up by one */
55 MPI_Sendrecv_replace(b, nlocal*nlocal, MPI_DOUBLE,
56                      uprank, 1, downrank, 1, comm_2d, &status);
57 }
58
59 /* Restore the original distribution of a and b */
60 MPI_Cart_shift(comm_2d, 0, +mycoords[0], &shiftsource, &shiftdest);
61 MPI_Sendrecv_replace(a, nlocal*nlocal, MPI_DOUBLE,
62                      shiftdest, 1, shiftsource, 1, comm_2d, &status);
63
64 MPI_Cart_shift(comm_2d, 1, +mycoords[1], &shiftsource, &shiftdest);
65 MPI_Sendrecv_replace(b, nlocal*nlocal, MPI_DOUBLE,
66                      shiftdest, 1, shiftsource, 1, comm_2d, &status);
67
68 MPI_Comm_free(&comm_2d); /* Free up communicator */
69 }
70
71 /* This function performs a serial matrix-matrix multiplication c = a*b */
72 MatrixMultiply(int n, double *a, double *b, double *c)
73 {
74     int i, j, k;
75
76     for (i=0; i<n; i++)
77         for (j=0; j<n; j++)
78             for (k=0; k<n; k++)
79                 c[i*n+j] += a[i*n+k]*b[k*n+j];
80 }
```

# Outer Product Matrix Multiplication