

# CSE 6220/CX 4220

## Introduction to HPC

# Lecture 19: Programming Models / Architectures and Interconnect Networks

**Helen Xu**

**[hxu615@gatech.edu](mailto:hxu615@gatech.edu)**




Georgia Tech College of Computing  
School of Computational  
Science and Engineering

Slide credits to Prof. Srinivas Aluru

# Parallel Models / Architectures

First half - revisiting shared-memory models for idealized parallel processors


The first model was the **Parallel Random Access Model** (PRAM) - before real parallel computers were built



# Architecture of an Ideal Parallel Computer

A natural extension of the Random Access Machine (RAM) serial architecture is the **Parallel Random Access Machine**, or PRAM.

- PRAMs consist of **unlimited processors** and a **global memory of unbounded size** that is uniformly accessible to all processors.
- Processors share a common clock but may execute different instructions in each cycle.
- Each processor can simultaneously undertake  $O(1)$  work in  $O(1)$  time.
- Each processor can access any memory location in  $O(1)$  time.



# PRAM Model

Depending on how simultaneous memory accesses are handled, PRAMs can be divided into four subclasses.

- Exclusive-read, exclusive-write (EREW) PRAM.
- Concurrent-read, exclusive-write (CREW) PRAM.
- Concurrent-read, concurrent-write (CRCW) PRAM.
- Exclusive-read, concurrent-write (ERCW) PRAM.


Fourth one is for completeness but not used much.

The first three are the most popular - listed in **increasing levels of power**

# Concurrent Writes

What does **concurrent write** mean, anyway?

- Common: write only if all values are identical.
- Priority: follow a predetermined priority order (e.g., min, max, etc.).
- Sum: Write the sum of all data items.
- Arbitrary: write the data from a randomly selected processor.



# Example: Sum $p$ numbers on PRAM

What is the time complexity of **adding**  $p$  numbers if we have:

Concurrent-read, concurrent-write (CRCW) PRAM: Sum model

- All processors write to the sum variable concurrently
- $O(1)$

Exclusive-read, exclusive-write (EREW) PRAM

- Similar to the tree-like algorithms we have done so far
- $O(\log p)$

# Problem: Compute logical OR on PRAM

What is the time complexity of **computing logical OR** of  $p$  Booleans if we have:

Concurrent-read, concurrent-write (CRCW) PRAM: Common model (only write identical values)

- Start by writing 0 in the result, then a processor writes 1 iff it has 1
- $O(1)$

Exclusive-read, exclusive-write (EREW) PRAM

- Similar to the tree-like algorithms we have done so far
- $O(\log p)$



# Complexity Theory for Parallel Algorithms


- **P-completeness:** Given any polynomial number of processors - can it be solved in polylog time?
- Every algorithm that is **NP-hard serially is also NP-hard in parallel** - by reducing the parallel algorithm to a serial algorithm
- Given a polynomial number of processors, time would only go up polynomially to run it on one processor.
- There are some problems that are polynomially solvable serially but **hard to parallelize** - e.g., some forms of depth-first search



# Physical Complexity of PRAM Computer

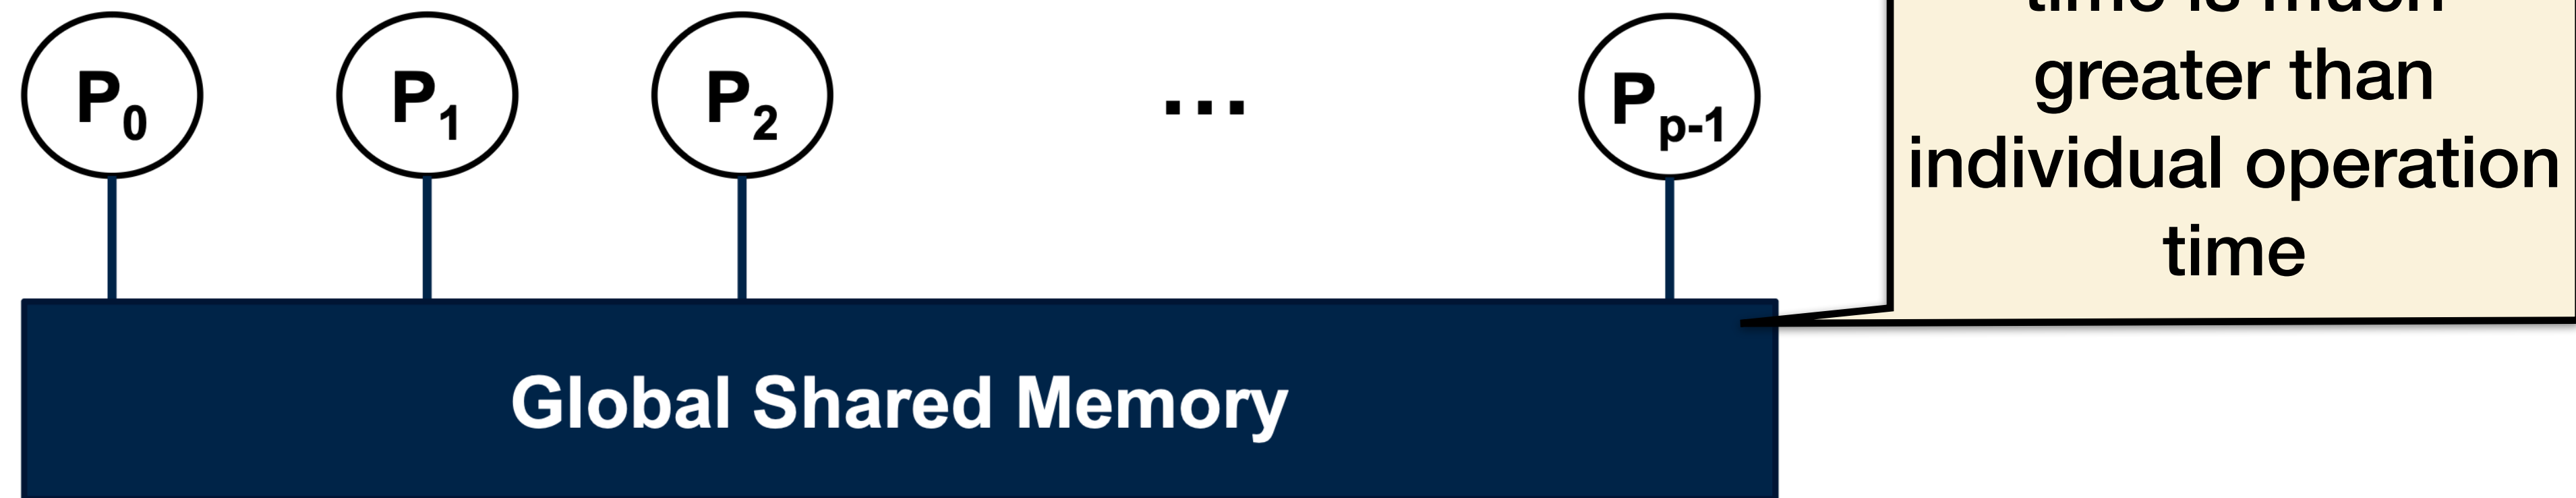
Processors and memories are **connected via switches**.

- Since these switches must operate in  $O(1)$  time at the level of words, for a system of  $p$  processors and  $m$  words, the switch complexity is  $O(mp)$ .
- Clearly, for meaningful values of  $p$  and  $m$ , **a true PRAM is not realizable**.



# PRAM Limitations

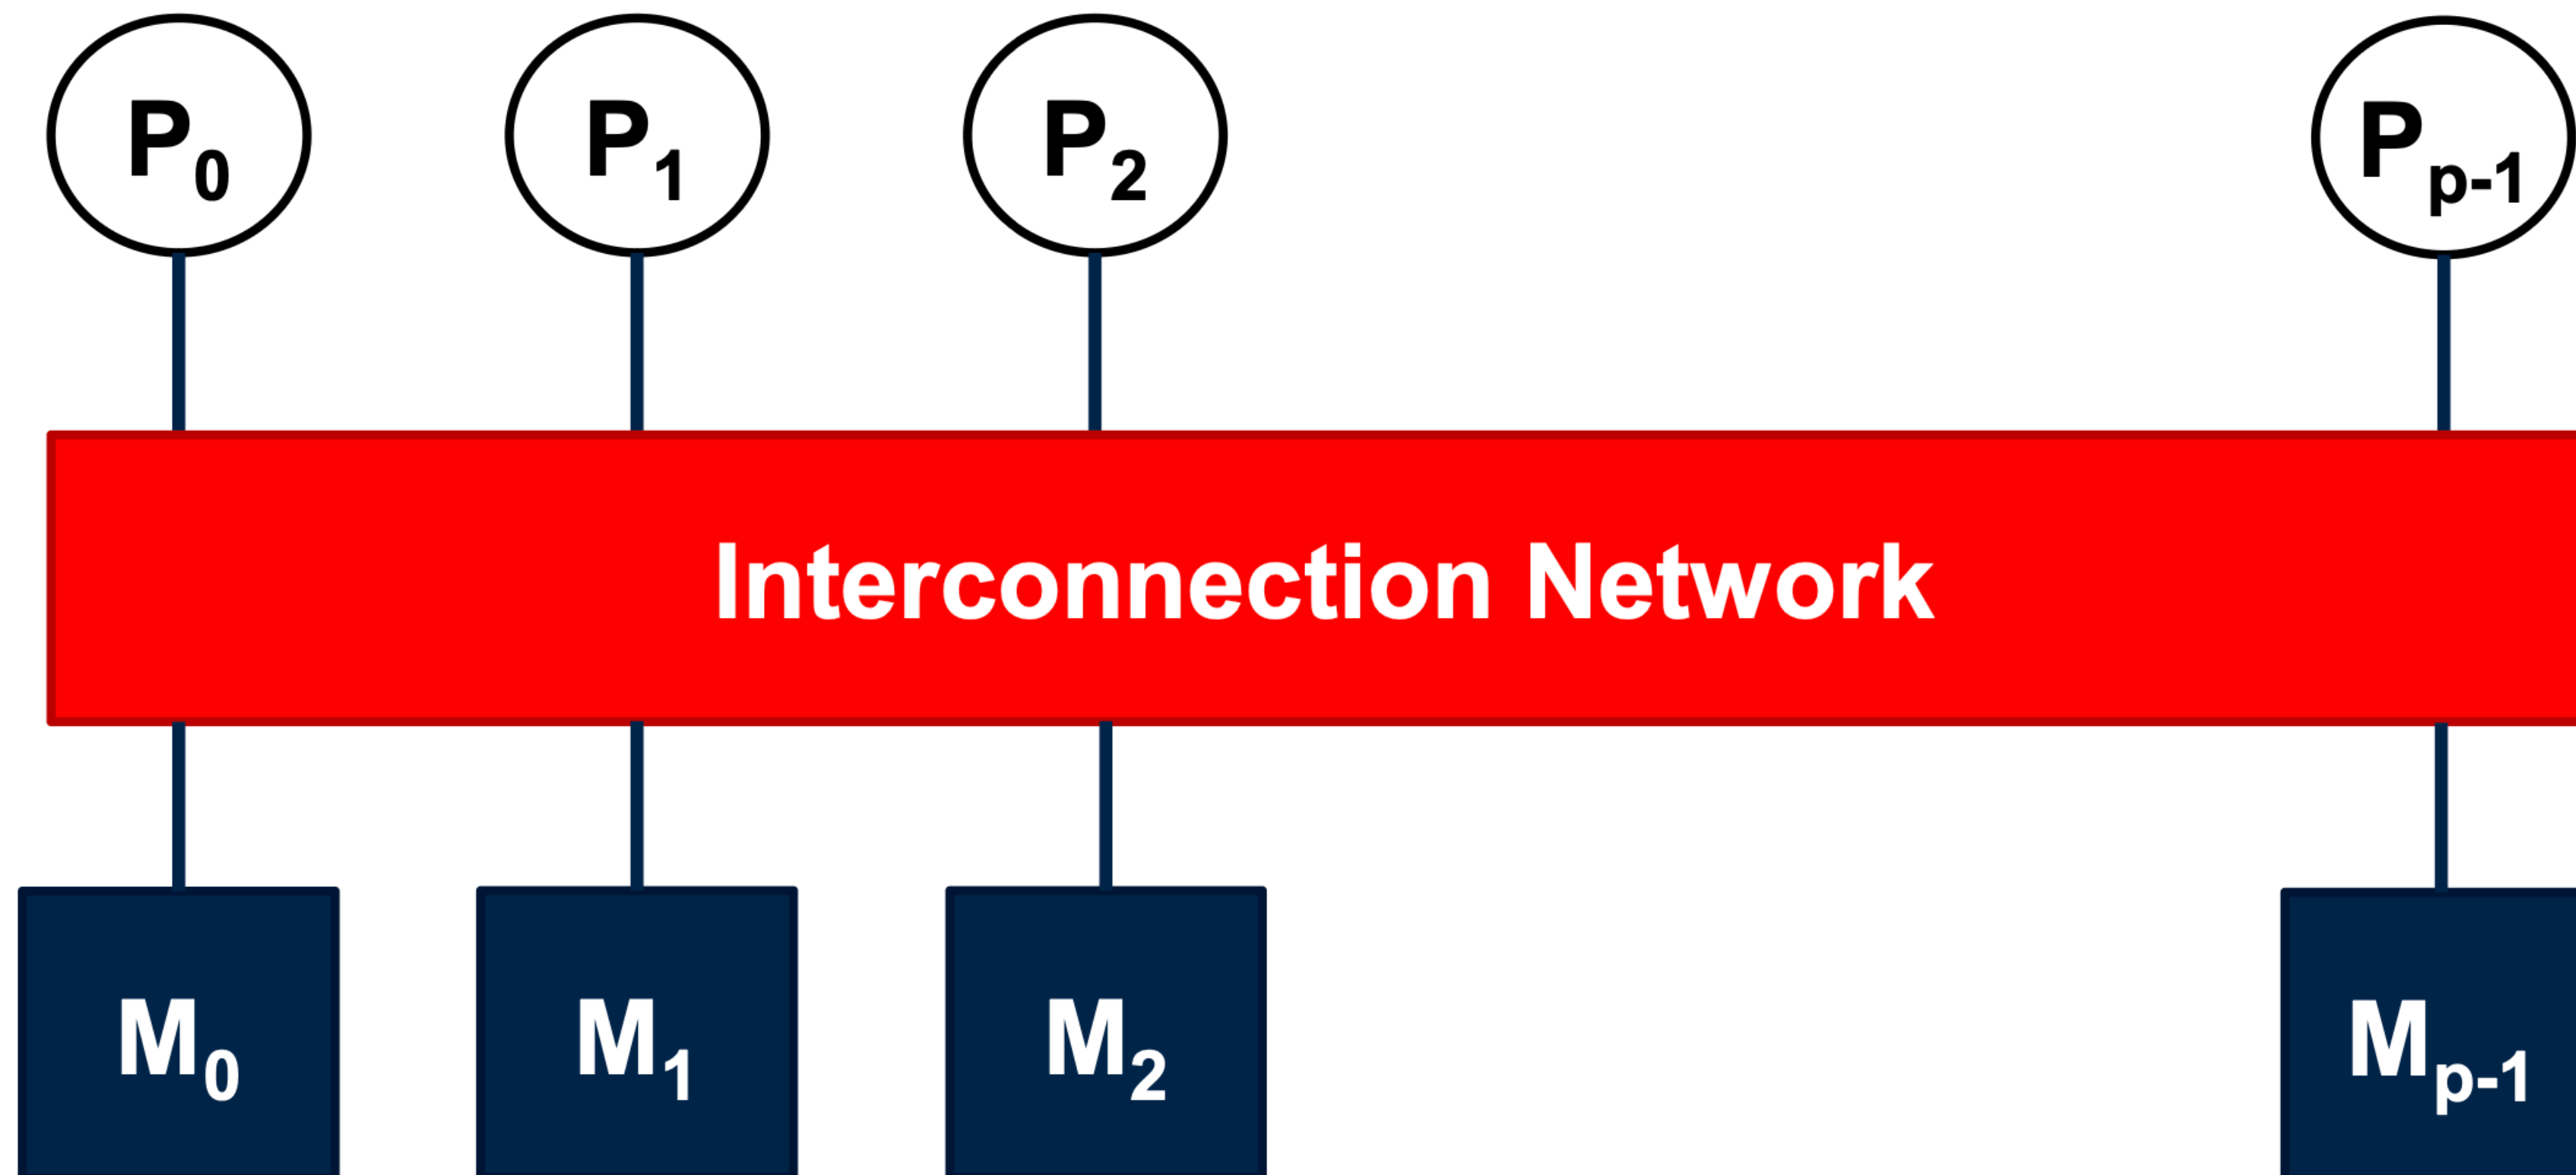
Many of the original PRAM algorithms didn't work well in practice because they didn't account for **communication**.



The PRAM model is still useful as a **first pass at designing** a parallel algorithm, since other models are more complex.

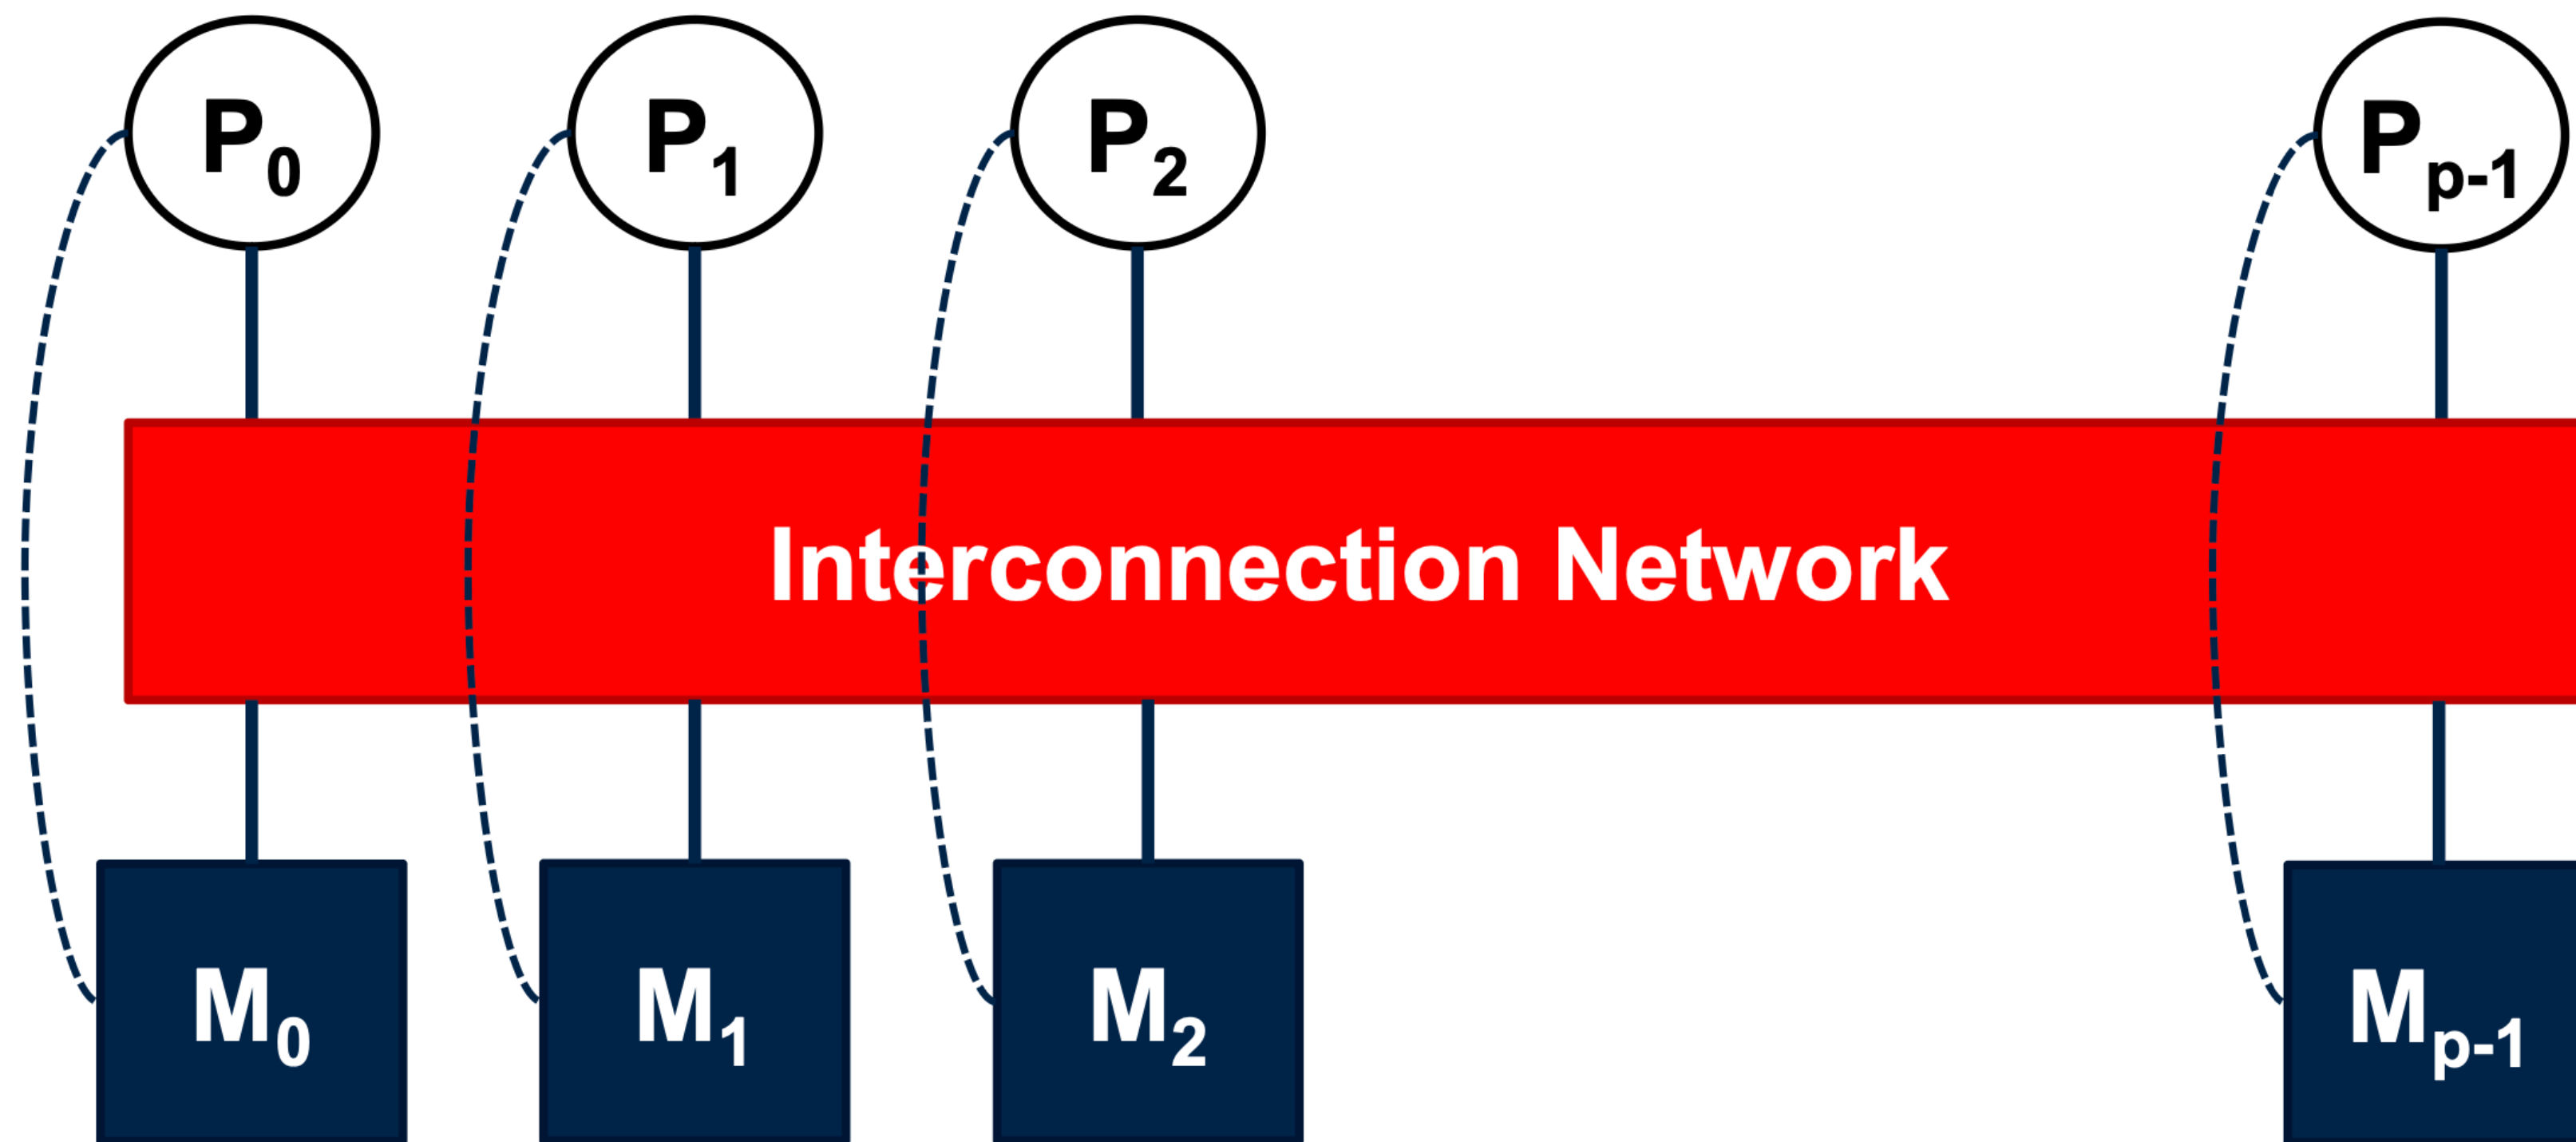
# Networked Shared Memory Model

- **Interconnection network** between processors and memory modules (might be larger than one element)



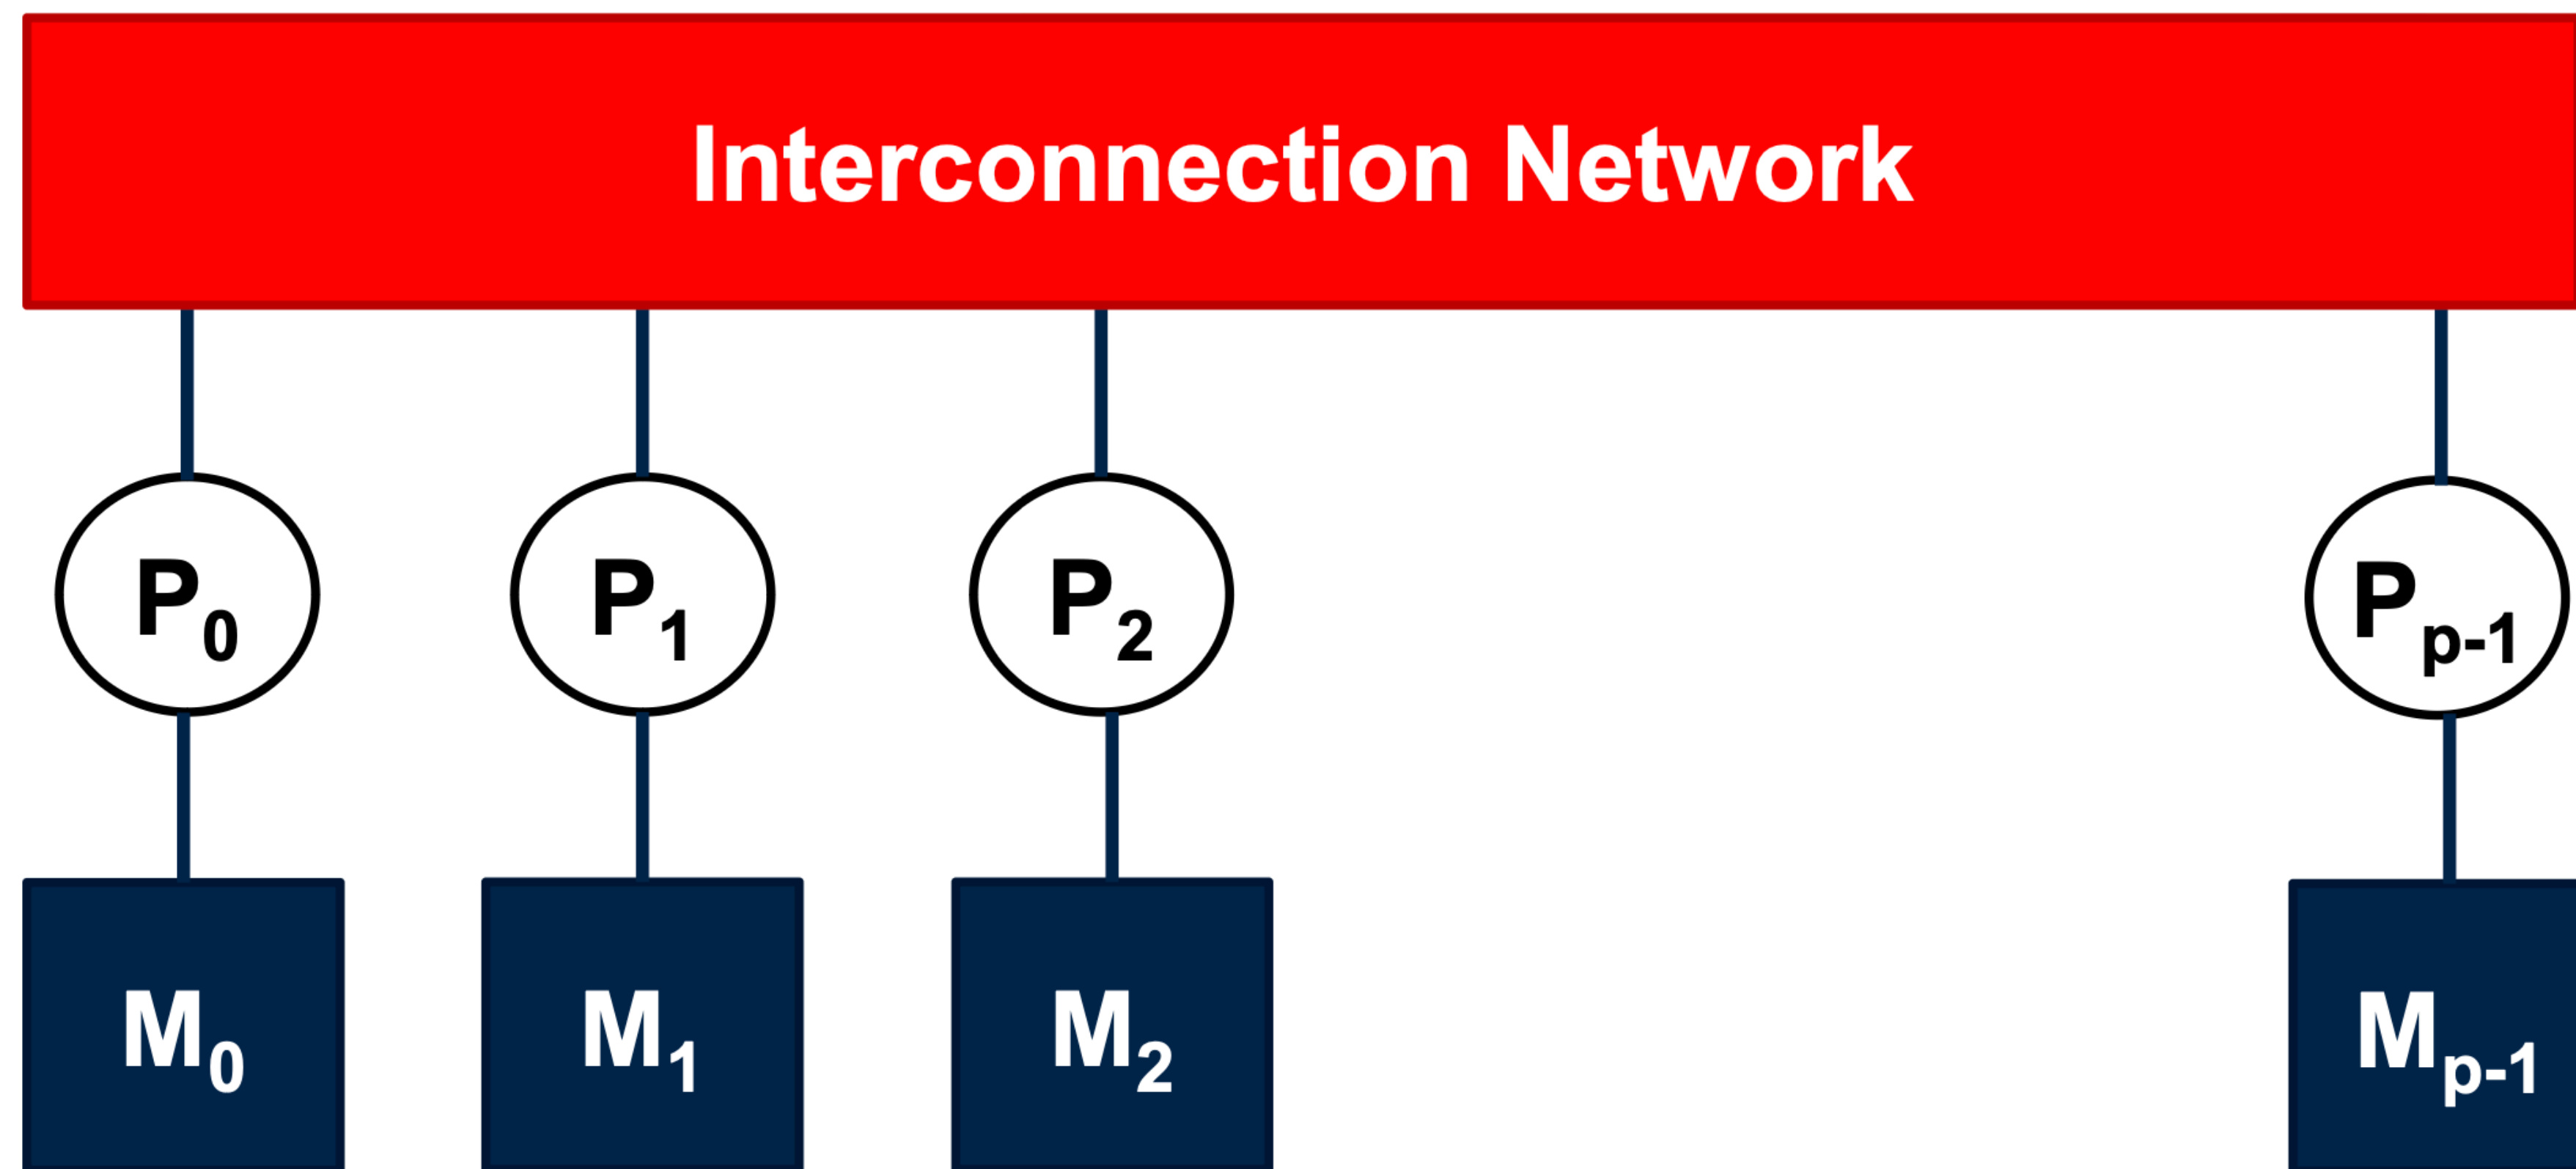
# Networked Shared Memory Model

One type of non-uniform memory access (NUMA) - **private link between each processor and its related memory module** (bypassing the interconnect).



# Networked Distributed-Memory Model

Another way to view the previous architecture is the distributed-memory model we have been working with so far.



# Distinguishing between shared and distributed memory

There are multiple levels at which we can specify whether we are using shared or distributed memory:

- Algorithm (depends on the model)
- Programming language / abstraction (e.g., MPI)
- Actual hardware/architecture

# Abstractions

Assume  $p$  processors with a memory of  $m$  per processor

Total Memory =  $mp = M$

	Shared	Distributed
Algorithm	$O(1)$ time to access memory	$O(1)$ time to access to only local memory
Programming	$\langle 0 \dots M - 1 \rangle$	$\langle 0 \dots p - 1, 0 \dots m - 1 \rangle$

How to support a shared-memory programming abstraction on top of distributed memory?

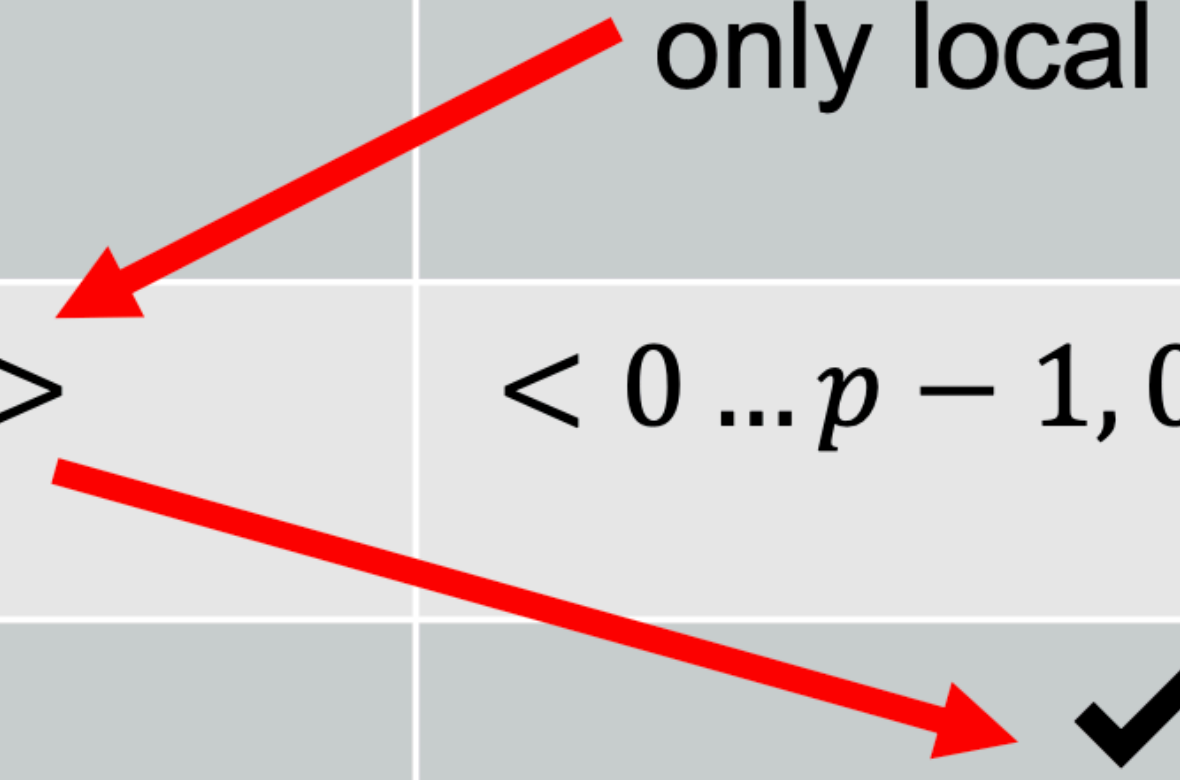


# Abstractions

How to support shared memory programming abstraction on top of a distributed memory?

$$\langle i \rangle \rightarrow \langle i / m, i \% m \rangle$$

	Shared	Distributed
Algorithm	$O(1)$ time to access memory	$O(1)$ time to access to only local memory
Programming	$\langle 0 \dots M - 1 \rangle$	$\langle 0 \dots p - 1, 0 \dots M - 1 \rangle$
Architecture	✗	✓

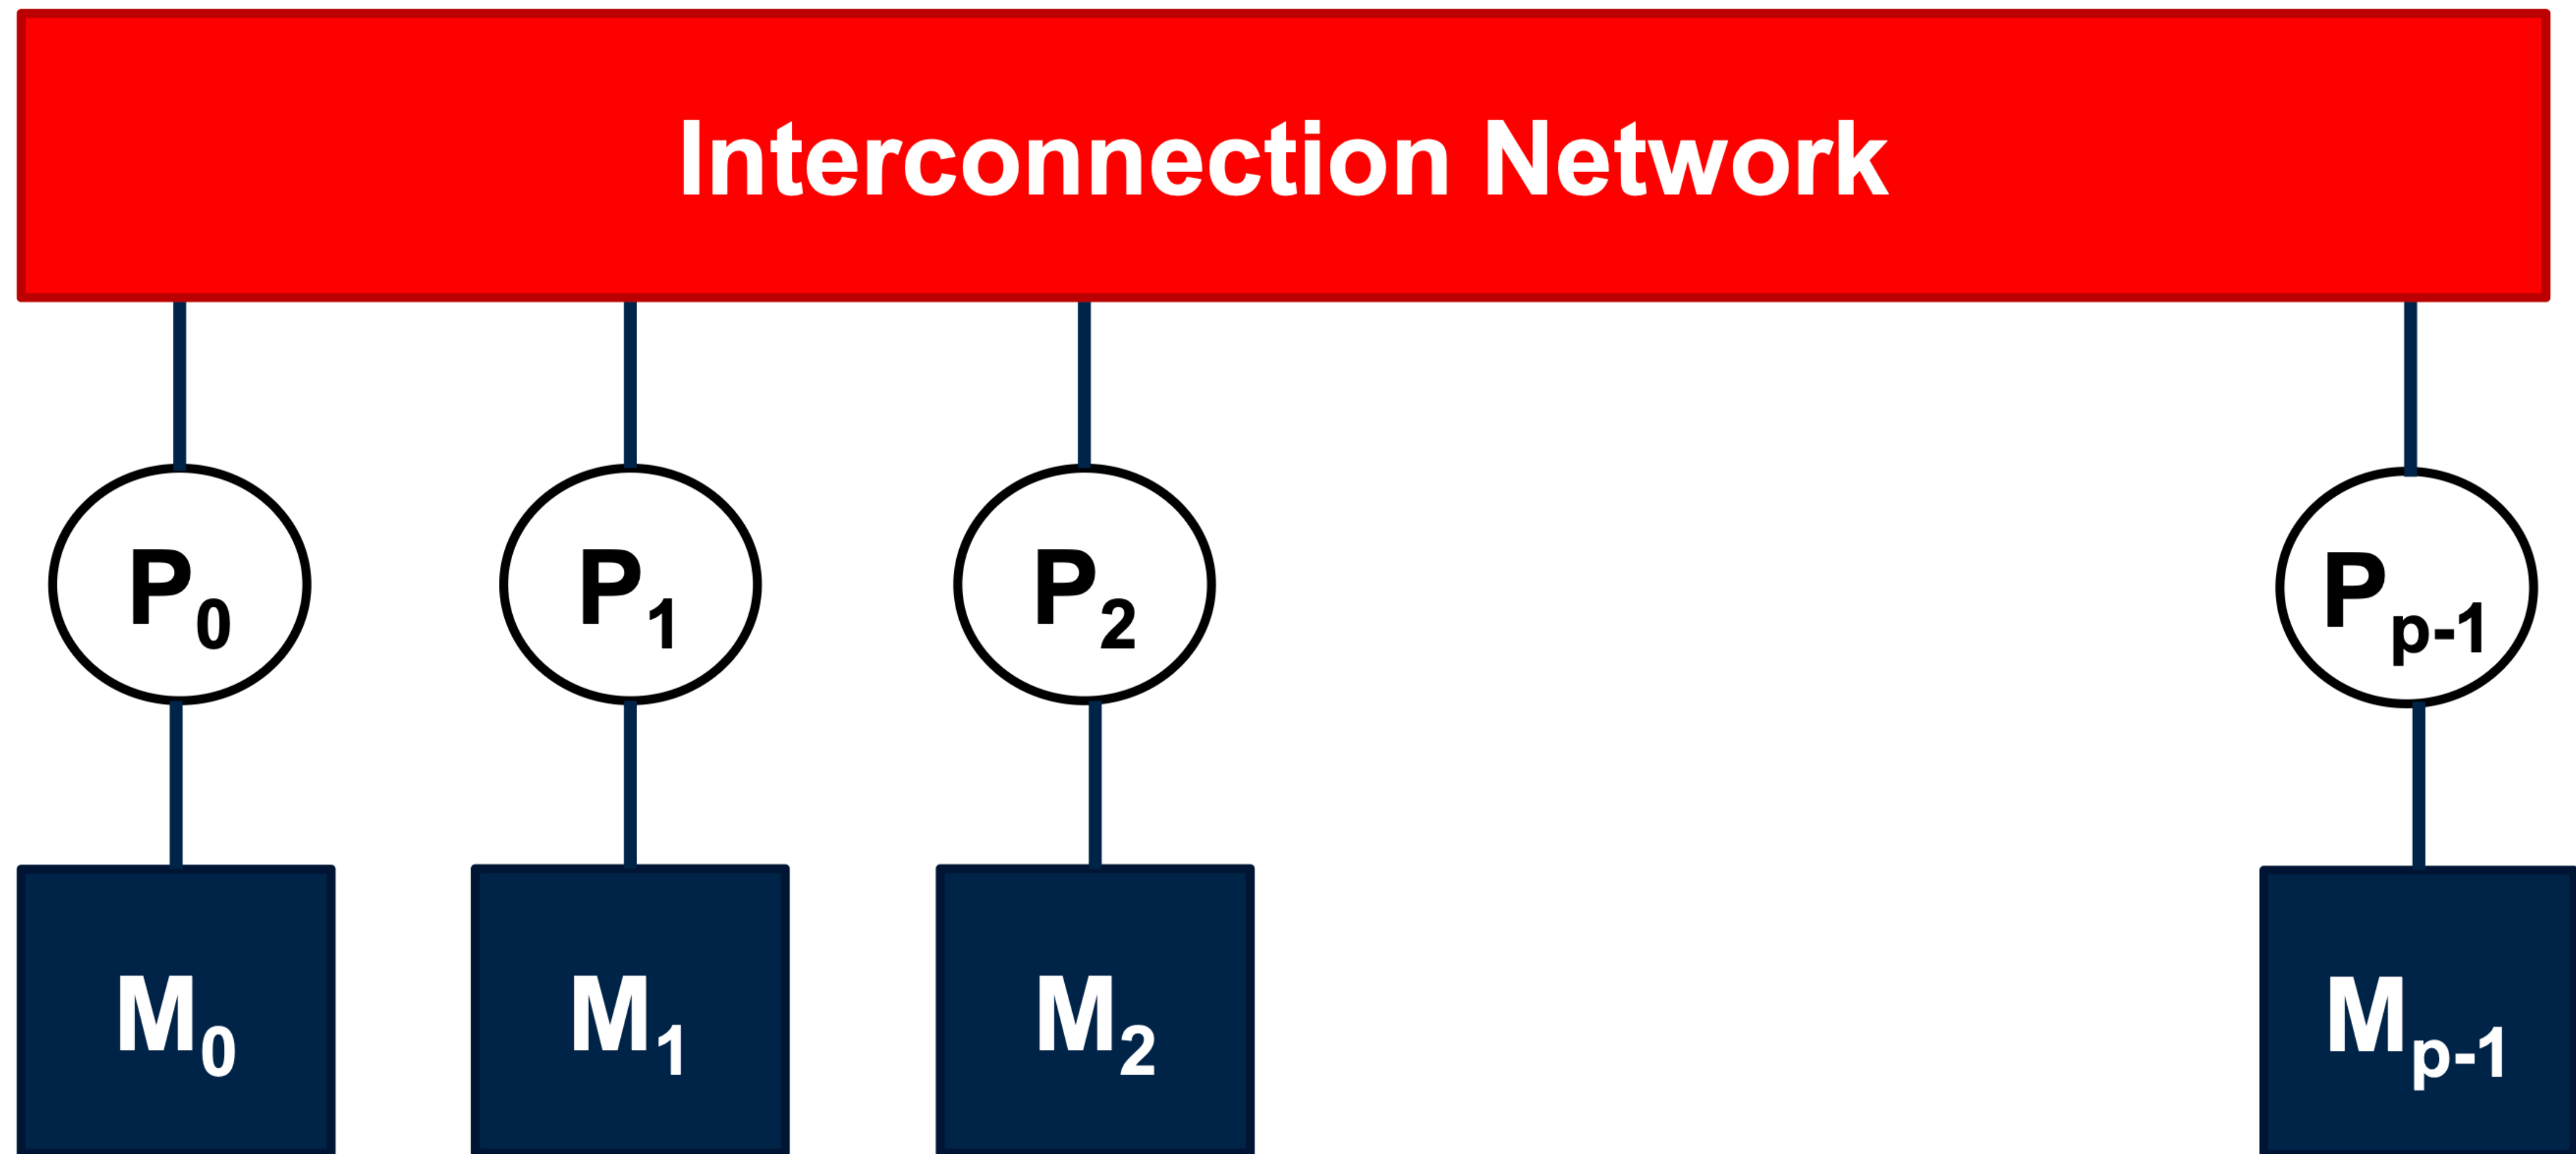


# Example

Consider the following simple problem - every processor tries to read the same global memory location (e.g.,  $a$ ) and set their own local variable (e.g.,  $x$ ).

- In CRCW/CREW, it would take  $O(1)$  time.
- In practice, it would actually take something like  $O(p)$  time since processors have to queue up at memory locations.
- To resolve the issue - declare array of size  $p$ , one processor reads  $a$ , writes it to another location, then 2 processors read those and write  $a$  to 2 more locations, etc. Then all  $p$  processors set  $x = a$  from different slots -  $O(\log p)$

# What should the interconnect network be?



# Interconnection Networks for Parallel Computers

Interconnection networks carry data between processors and to memory.

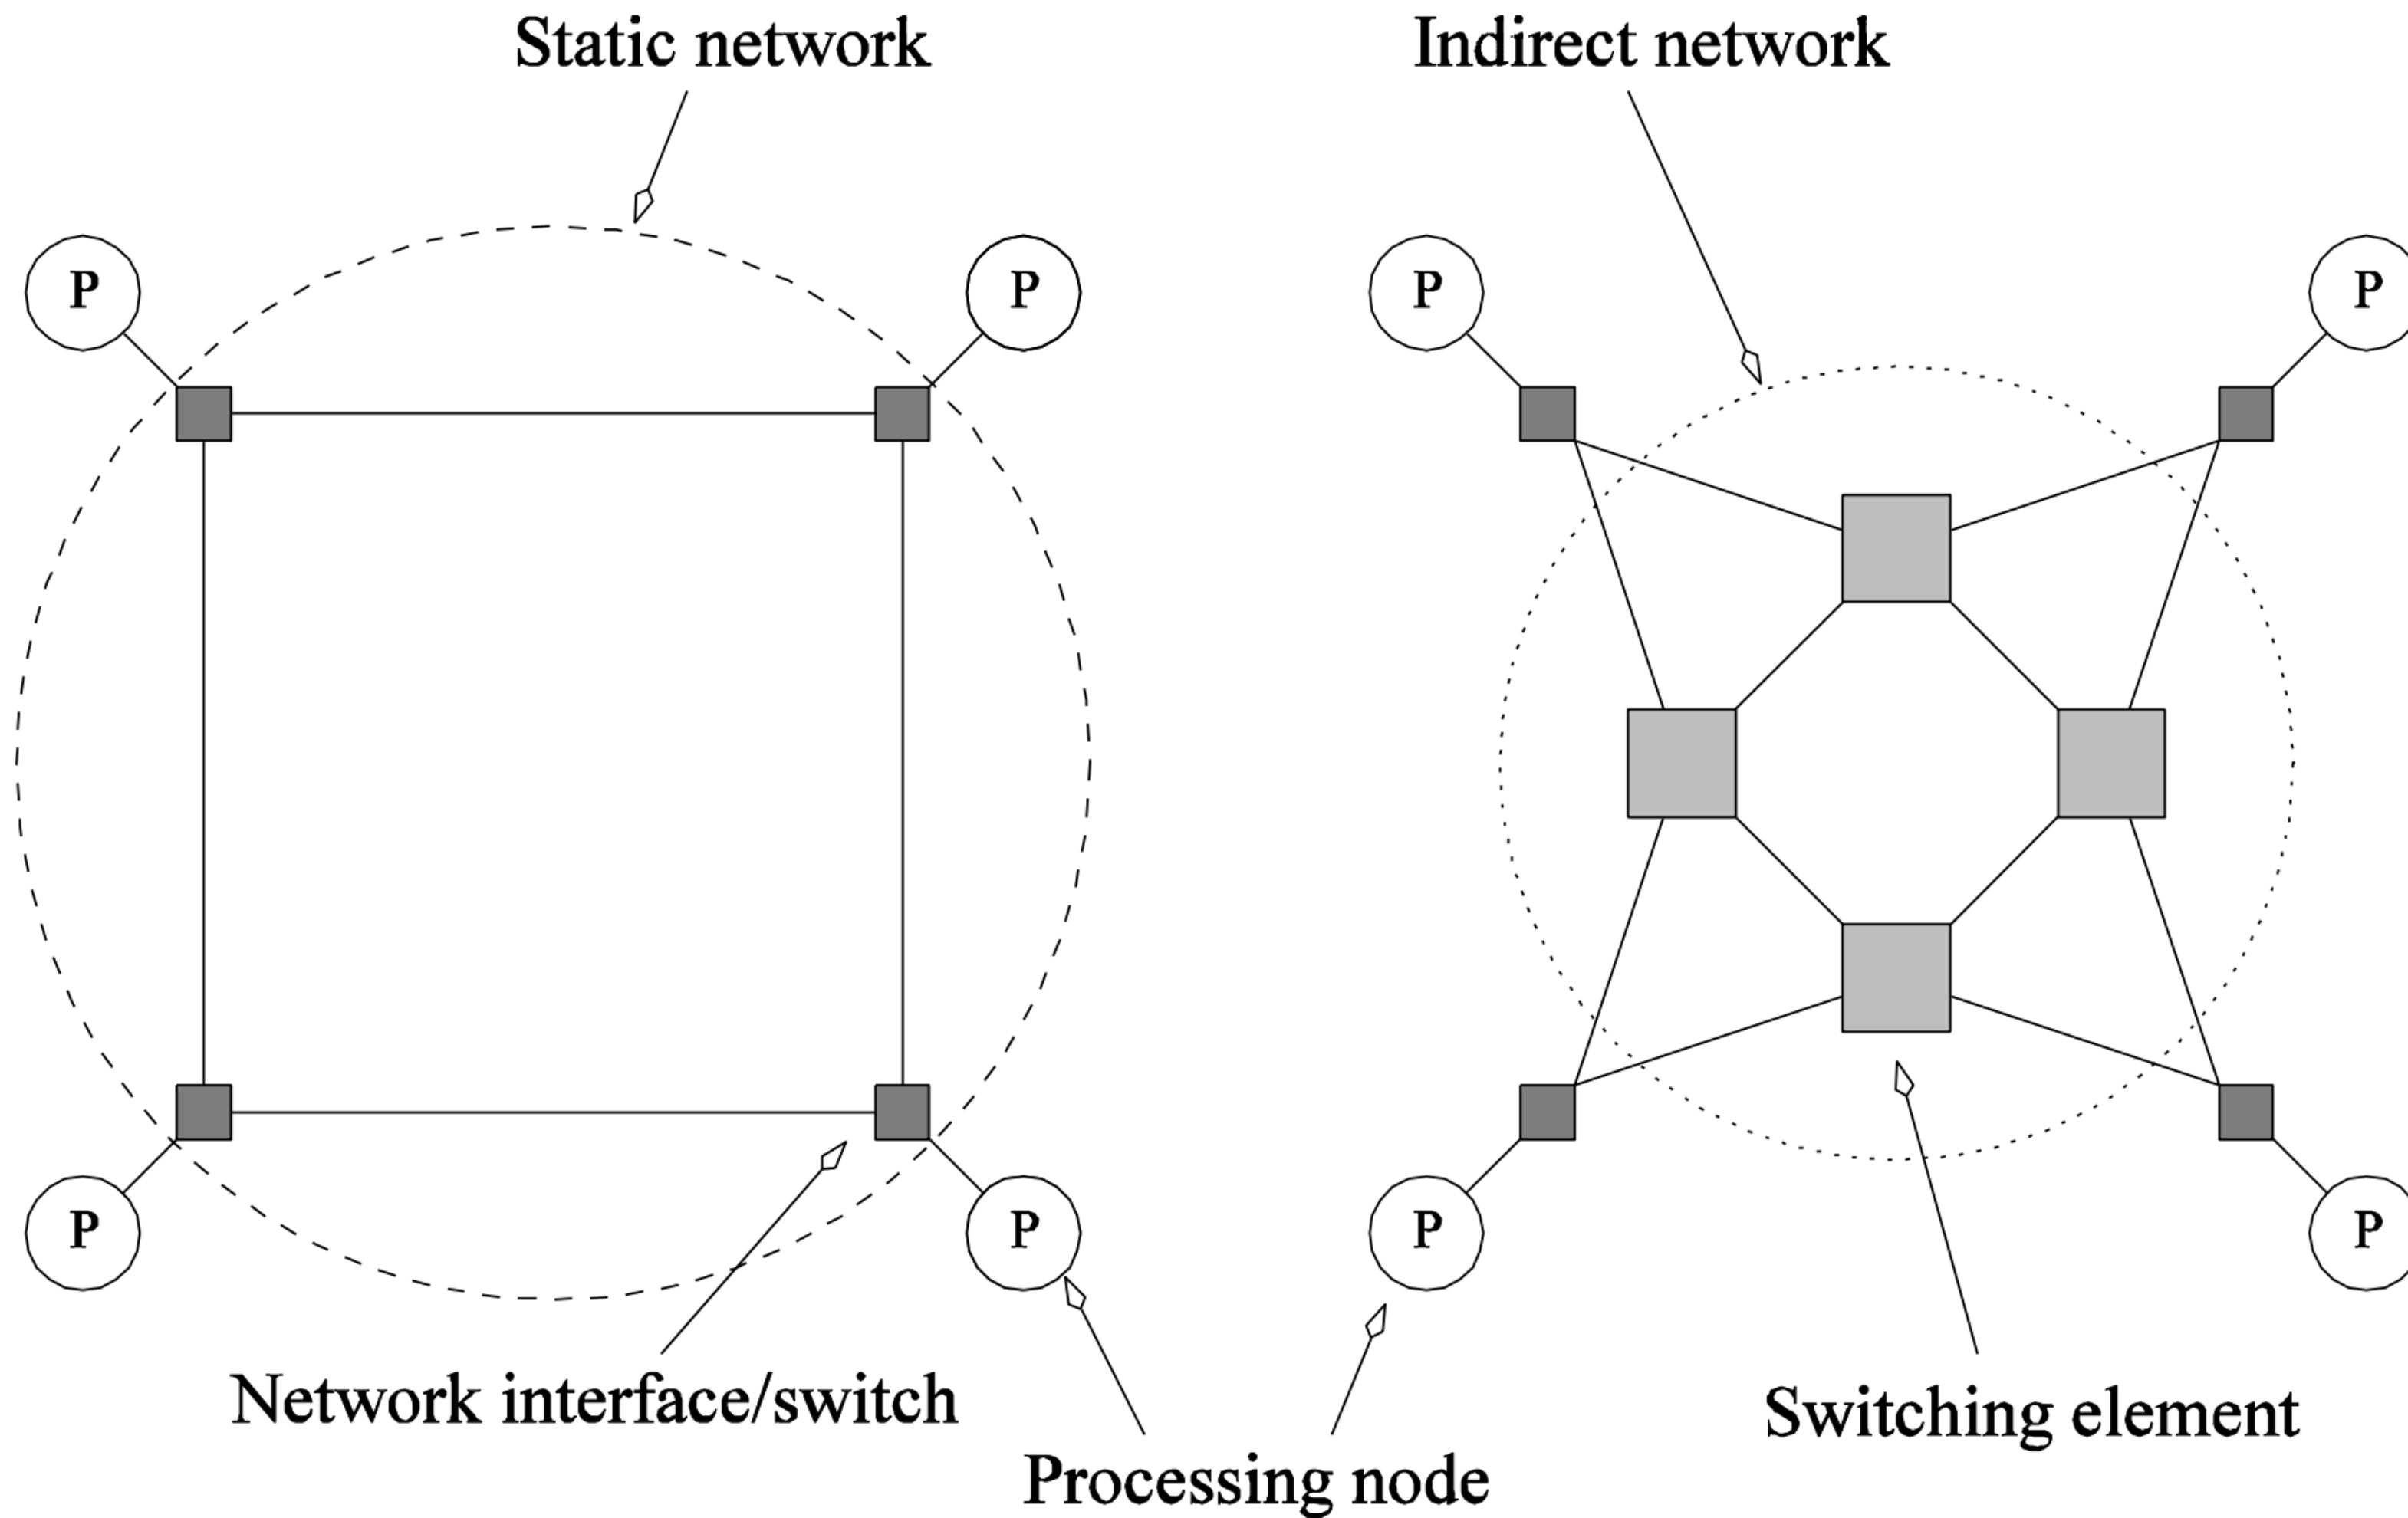
Interconnects are made of switches and links (wires, fiber).

Interconnects are classified as **static** or **dynamic**.

Static networks consist of **point-to-point communication links** among processing nodes and are also referred to as **direct** networks.

Dynamic networks are built using **switches and communication links**.  
Dynamic networks are also referred to as **indirect networks**.

# Static and Dynamic Interconnection Networks



# Evaluating Static Networks

**Distance:** *between a pair of processors is the number of links on a shortest path connecting the two.*

**Diameter:** The largest of distances between pairs of processors.

**Ideal Diameter:**  $\Theta(1)$

**Bisection Width (BW):** The minimum number of wires you must cut to divide the network into two subnetwork with  $p/2$  processors each.

**Ideal Bisection Width:**  $\Theta(p)$

# Evaluating Static Networks

**Cost:** The number of links or switches (whichever is asymptotically higher) is a meaningful measure of the cost. However, a number of other factors, such as the ability to layout the network, the length of wires, etc., also factor in to the cost.

***Ideal Cost:  $O(p)$***




# Network Topologies: Completely Connected Network

- Each processor is connected to every other processor.

- Diameter = 1


- $BW = \frac{p}{2} \times \frac{p}{2} = \frac{p^2}{4} = \Theta(p^2)$



- Number of Links per node:  $p - 1$
- Cost : The number of links in the network scales as  $O(p^2)$ .
- While the performance scales very well, the hardware complexity is not realizable for large values of  $p$ .

# Network Topologies: Star Connected Network


- Every node is connected only to a common node at the center.



- Diameter = 2
- Distance between any pair of nodes is  $O(1)$ .
- However, the central node becomes a bottleneck.
- BW = 1
- Number of Links per node: maximum of  $p - 1$
- Cost :  $p - 1$

# Network Topologies: Tree Properties


- The distance between any two nodes is no more than  $2 \log p$ .
- Links higher up the tree potentially carry more traffic than those at the lower levels.
- Trees can be laid out in 2D with no wire crossings. This is an attractive property of trees.



# Fat Trees


**Fat trees** [Leiserson 85] double links (and bandwidth) at each level (starting at leaves and going up) to handle more traffic on the way up.

Example supercomputers include the Connection Machine CM-5, Summit and Sierra (until recently, the two fastest), etc.




# Network Topologies - Linear Arrays / Rings

In a linear array, each node has two neighbors, one to its left and one to its right. If the nodes at either end are connected, we refer to it as a 1-D torus or a ring.



(a)




(b)

**Linear arrays: (a) with no wraparound links; (b) with wraparound link.**




# Network Topologies: 2D and 3D Meshes


- A generalization to 2 dimensions has nodes with 4 neighbors
- A further generalization to  $d$  dimensions has nodes with  $2d$  neighbors.



(a)



(b)



(c)

Two and three dimensional meshes:

- (a) 2-D mesh with no wraparound;  
(b) 2-D mesh with wraparound link (2-D torus); and  
(c) a 3-D mesh with no wraparound.

$$d \text{ dimension } k\text{-ary mesh } p = k^d$$

# Evaluating Static Networks

Down is good

Up is good

Down is good

Down is good

Network	Diameter	Bisection Width	Max. Links per Node	Cost (total no. of links)
Completely-connected	1	$p^2/4$	$p - 1$	$p(p - 1)/2$
Star	2	1	$p - 1$	$p - 1$
Linear array	$p - 1$	1	2	$p - 1$
Complete binary tree	$2 \log((p + 1)/2)$	1	3	$p - 1$
2-D mesh, no wraparound	$2(\sqrt{p} - 1)$	$\sqrt{p}$	4	$2(p - \sqrt{p})$
2-D wraparound mesh (torus)	$2 \lfloor \sqrt{p}/2 \rfloor$	$2\sqrt{p}$	4	$2p$
Wraparound $k$ -ary $d$ -dimensional mesh ( $p = k^d$ )	$d \lfloor k/2 \rfloor$	$2k^{d-1}$	$2d$	$dp$
	$d \lfloor p^{1/d}/2 \rfloor$	$2p^{1-\frac{1}{d}}$		



# How To Maximize Mesh Dimensionality?

$p = 32$ . How many dimensions can we make?

- 8 X 4 2-D mesh
- 4 X 4 X 2 3-D mesh
- 4 X 2 X 2 X 2 4-D mesh
- 2 X 2 X 2 X 2 5-D mesh
- Can't go any further! Also, torus just duplicates mesh links.
- To maximize dimensionality, choose variable number of dimensions based on  $p$ . For  $p = 2^d$ , choose  $d = \log p$  as the number of dimensions.

# Hypercube: Mesh of Maximum Dimensionality

$$d = \log p ; p = 2^d$$

- Diameter =  $d(p^{1/d} - 1) = d(2 - 1) = d = \log p$
- Bisection Width =  $p^{1-1/d} = \frac{p}{p^{1/d}} = p/2$

How do we address processors in hypercube and find neighbors?

- It's just a mesh (2 processors in each dimension).
- $d$ -dimensional mesh has ranks described by  $d$  coordinates.
- To find neighbors along a dimension, hold all other dimensions fixed, and increment/decrement chosen dimension by 1.

# Hypercube: Processor Representation and Neighbor Finding

$d$ -dimensions. Coordinate along each dimension is 0 or 1.

Example: (1,0,0,1,0,1) in 6-dimensional hypercube.

Neighbor along a dimension:

- If coordinate is 0, neighbor is 1 (other neighbor doesn't exist)
- If coordinate is 1, neighbor is 0 (other neighbor doesn't exist)
- Every processor is a corner processor!

Since each coordinate is 0 or 1, just write them consecutively: (1,0,0,1,0,1) → 100101

Neighbor along a dimension: invert the bit

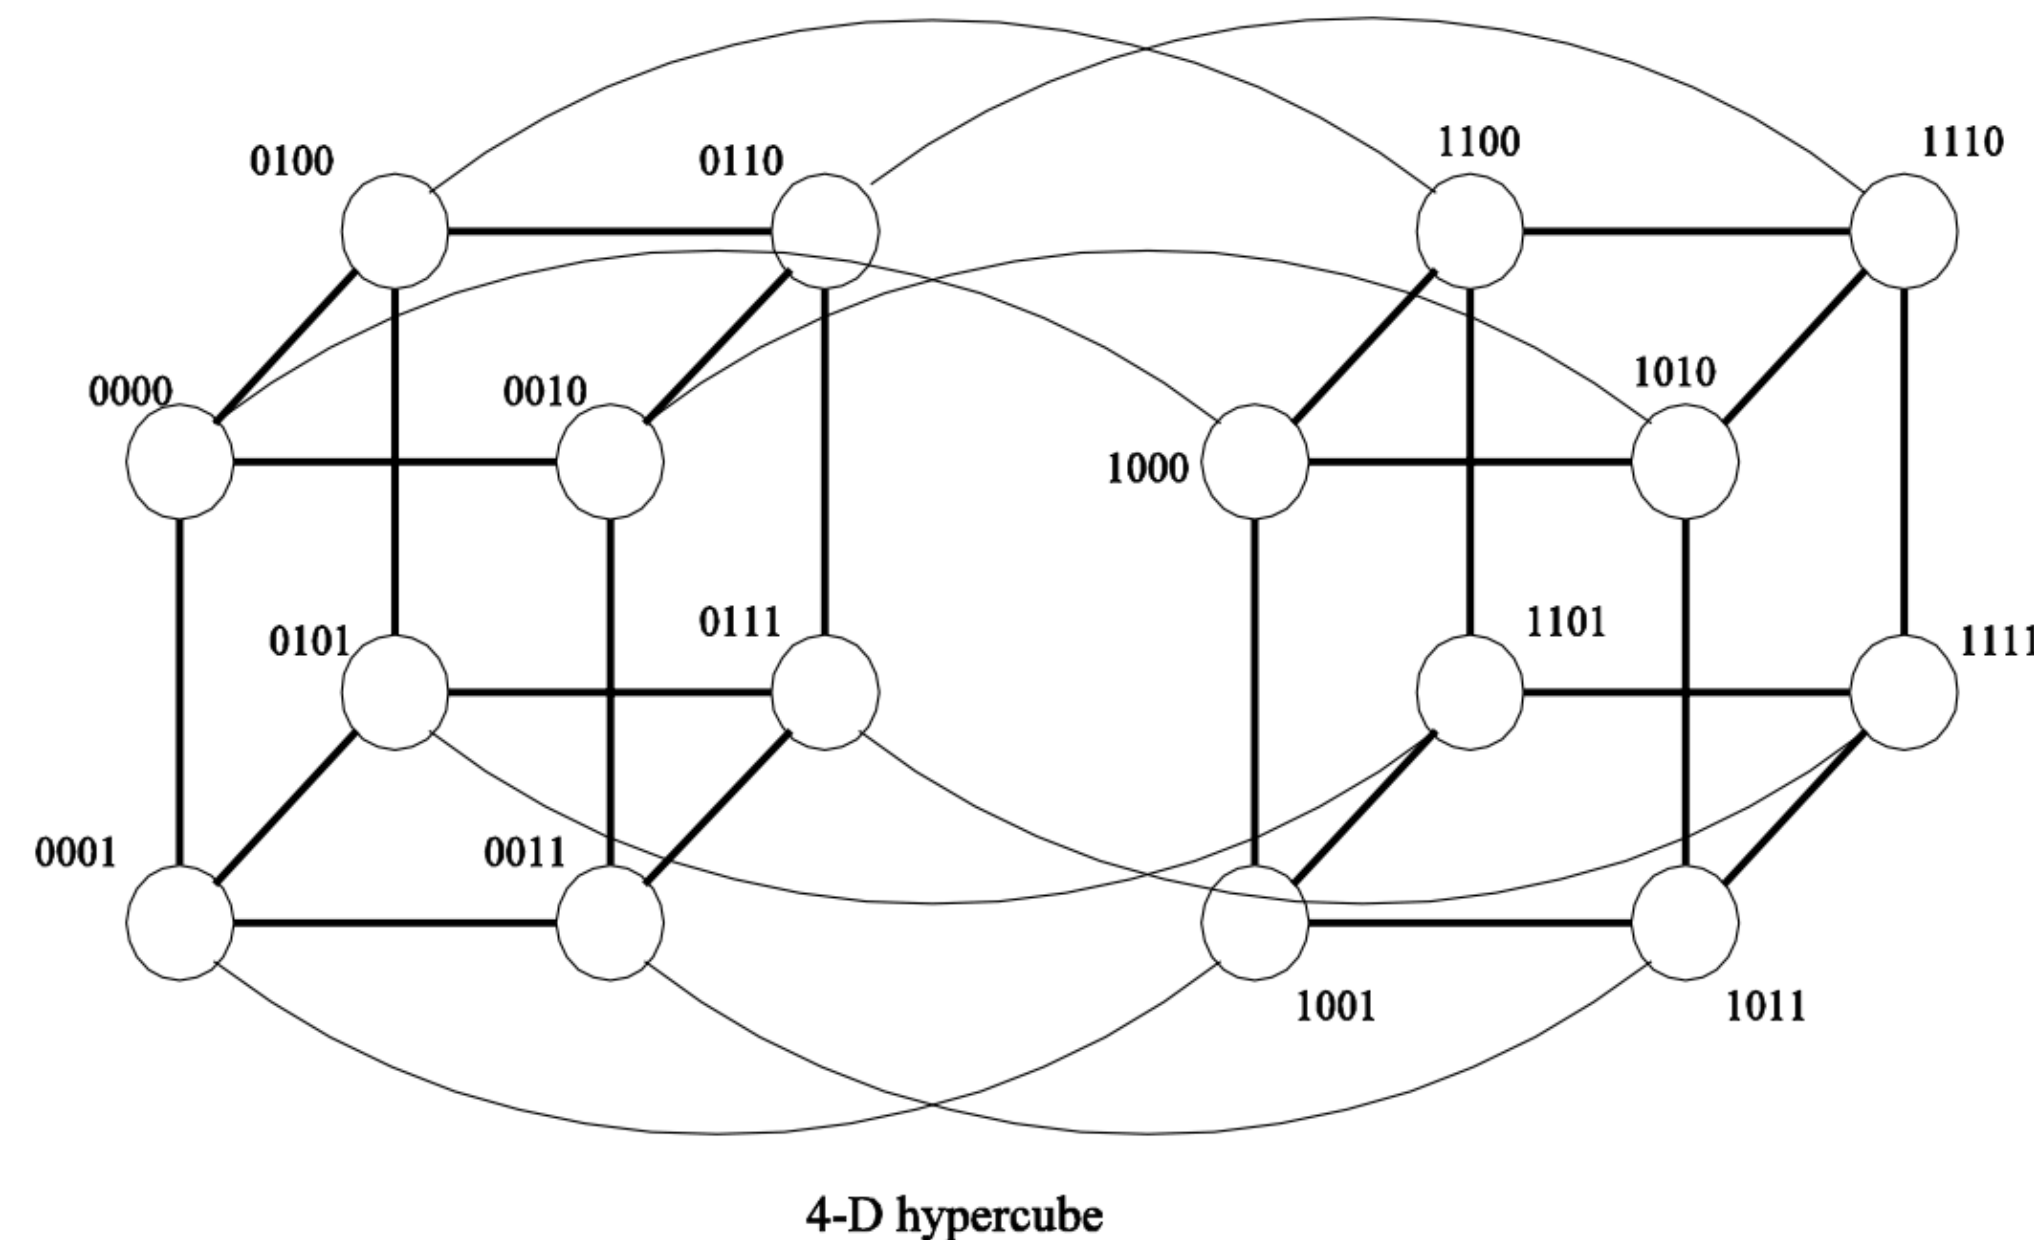
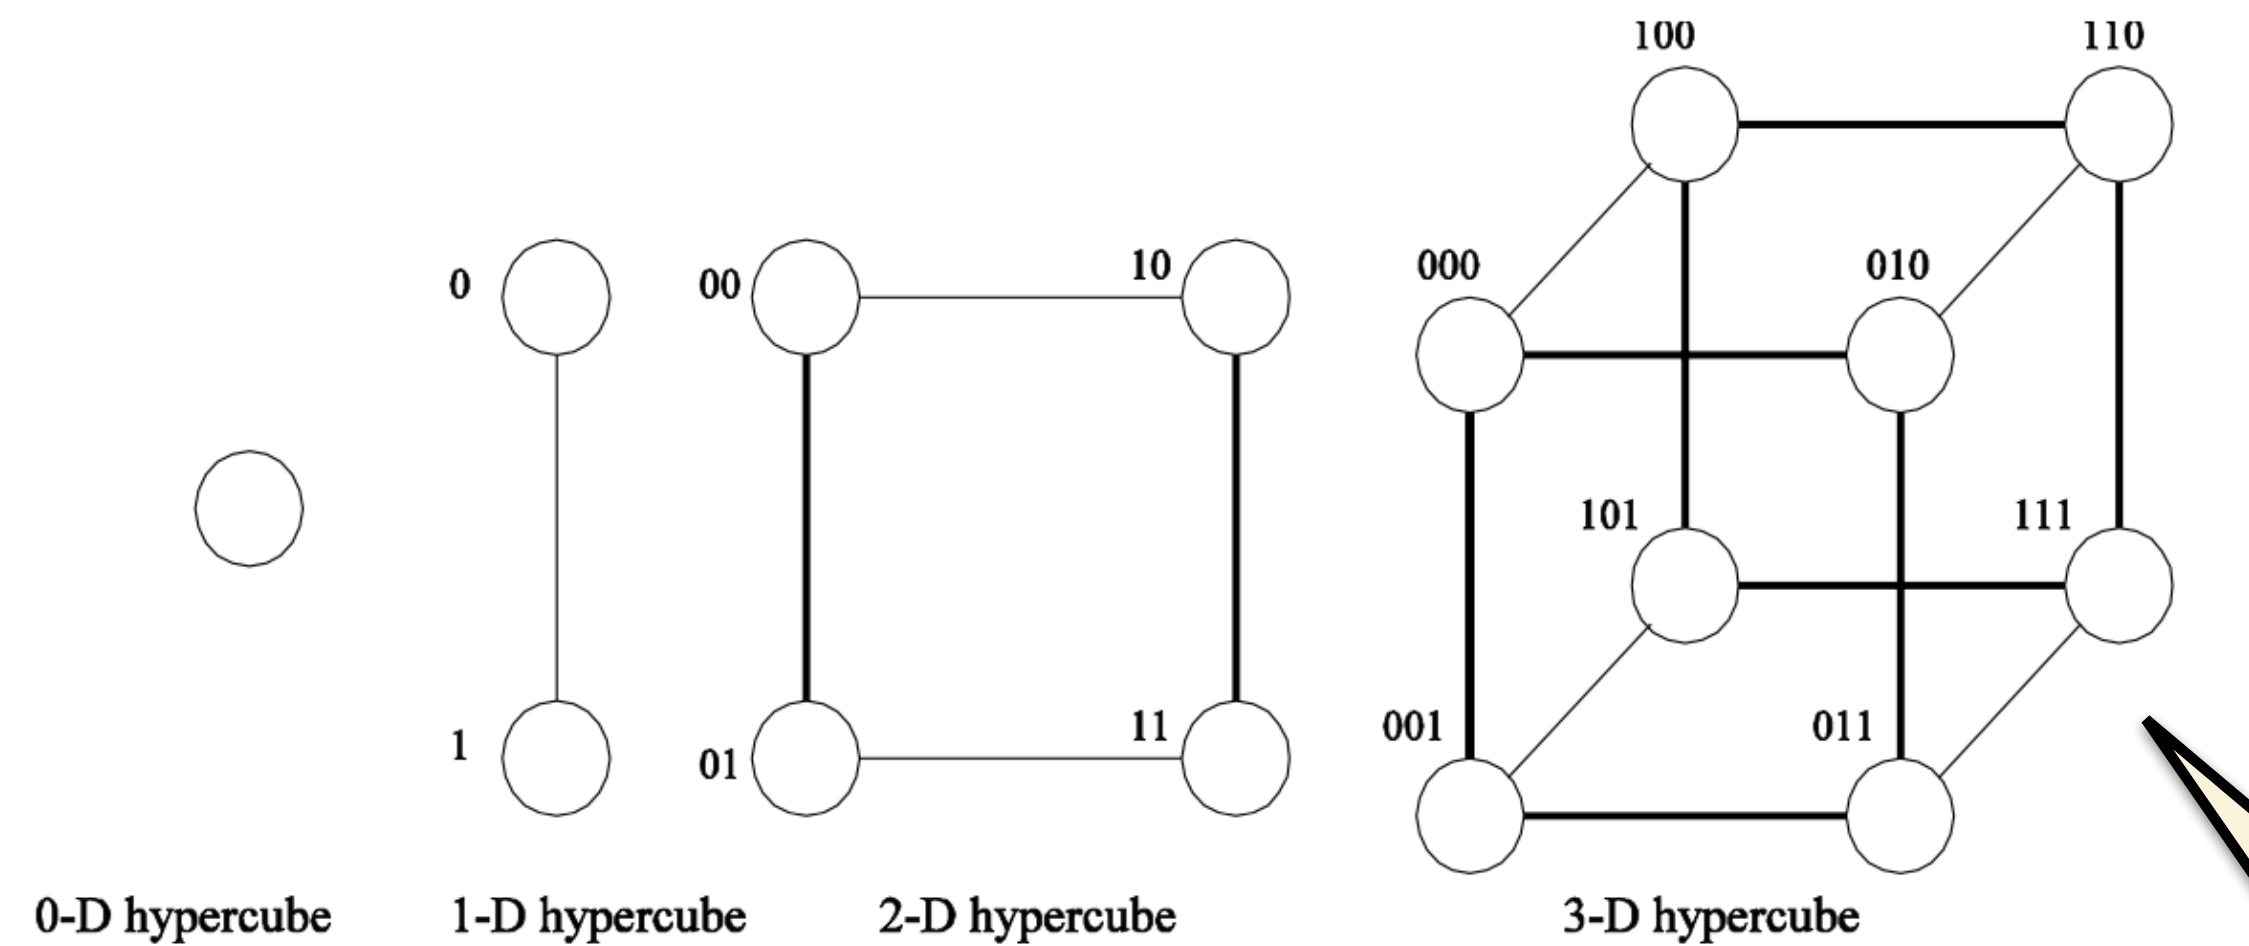
# Network Topologies: Properties of Hypercubes

The distance between any two nodes is at most  $\log p$  (diameter).

Each node has  $\log p$  neighbors.

The distance between two nodes is given by the number of bit positions at which the two nodes differ.

# Alternative View of Hypercubes



Construction of  
hypercubes from  
hypercubes of  
lower dimension

# Difficulty of Constructing Hypercubes

How do we **reuse nodes in smaller dimensional hypercubes** to build larger dimension hypercubes?


- In practice: **set maximum number of dimensions** allowed.

Hypercubes **can't be embedded in physical three-dimensional space** without crisscrossing links and non-uniform length cables.

- Solution: spread the links over **multiple stages**



# Static Butterfly Interconnection Network







# Dynamic Multistage Interconnection Networks

Basic switching element is a 2X2 crossbar

- A crossbar can connect any input to any output
- The switches operate in two modes – crossover or passthrough.



(a)




(b)

**Two switching configurations of the  $2 \times 2$  switch:**  
**(a) Pass-through; (b) Cross-over.**

# Perfect Shuffle

Goal: Route arbitrary  
hypercubic permutations  
efficiently




L	R	L	R	L	R
0	4	0	2	0	1
1	5	4	6	2	3
2	6	1	3	4	5
3	7	5	7	6	7

Communication  
pattern is the same  
across stages

take 1 from the left  
and 1 from the  
right in stages

# Omega Network


- An omega network has  $p/2 \times \log p$  switching nodes, and the cost of such a network grows as  $O(p \log p)$ .



# Routing in an Omega Network


- Destination Tag Algorithm (DTA) – consider bits of destination processor from most significant to least significant
  - Bit is 0: connect to upper output
  - Bit is 1: connect to lower output

**Example:  
010 to 101**




# DTA based on Perfect Shuffle

Perfect shuffle brings together pairs of elements with flipped MSB



L	R	L	R	L	R
0	4	0	2	0	1
1	5	4	6	2	3
2	6	1	3	4	5
3	7	5	7	6	7


# Not Every Permutation Can be Supported



**An example of congestion in omega network: one of the messages  $010 \rightarrow 111$  and  $110 \rightarrow 100$  is blocked at link AB.**



# Omega Network - Hypercubic Permutations



**Omega multistage interconnection network can route any hypercubic permutation without congestion.**