

# CSE 6220/CX 4220

## Introduction to HPC

# Lecture 20: Sparse Matrix-Vector Multiplication (SpMV)

Helen Xu  
[hxu615@gatech.edu](mailto:hxu615@gatech.edu)



# 1990 Nobel Prize in Economics



Our models strained the computer capabilities of the day [1950s]. I observed that **most of the coefficients in our matrices were zero**; i.e. , the nonzeros were “sparse” in the matrix

- Harry Markowitz

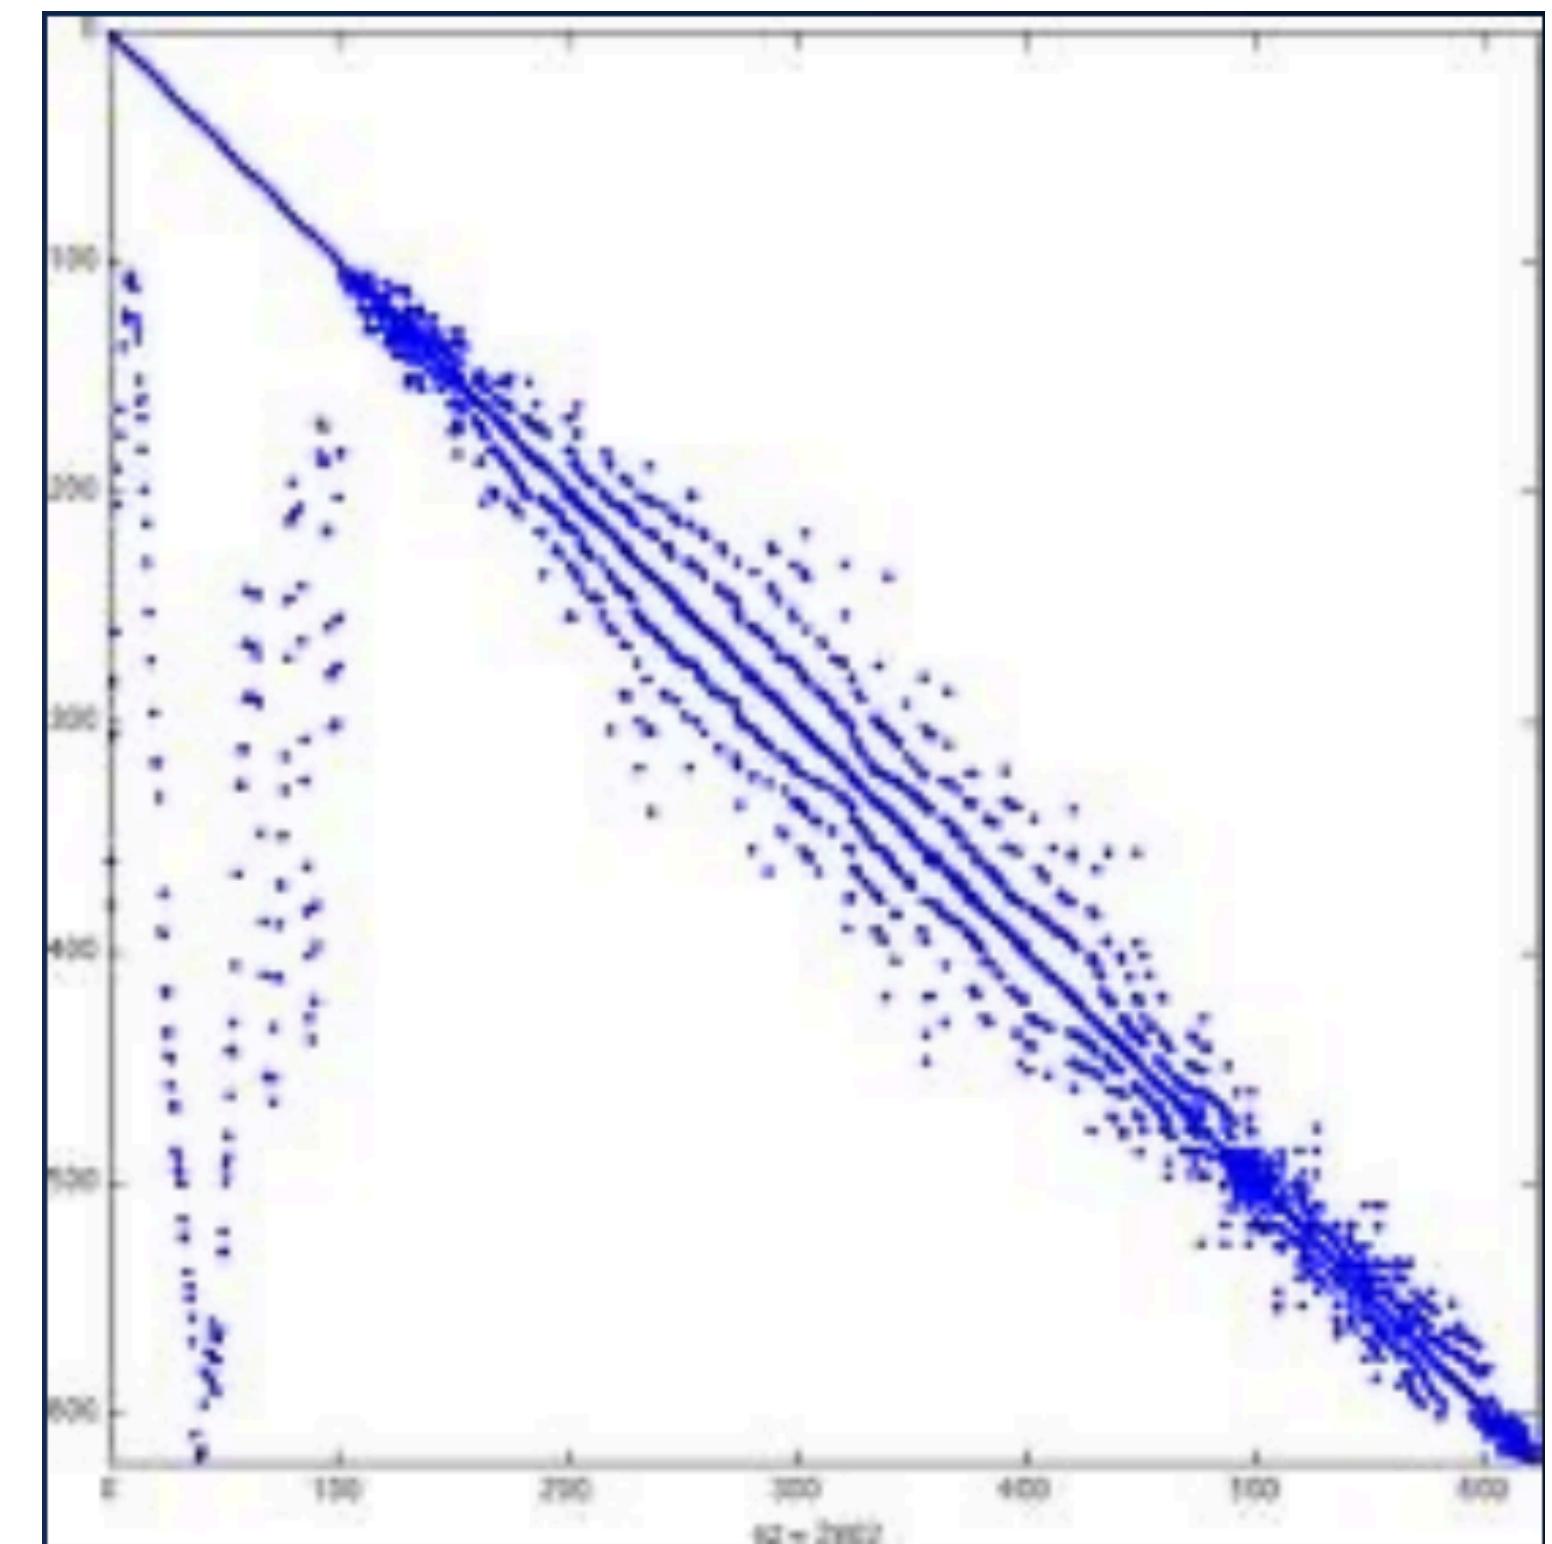
# What is a sparse matrix?

A matrix with **primarily zeros** (>> 90%).

Representing them using dense data structures  
**wastes memory and computation.**

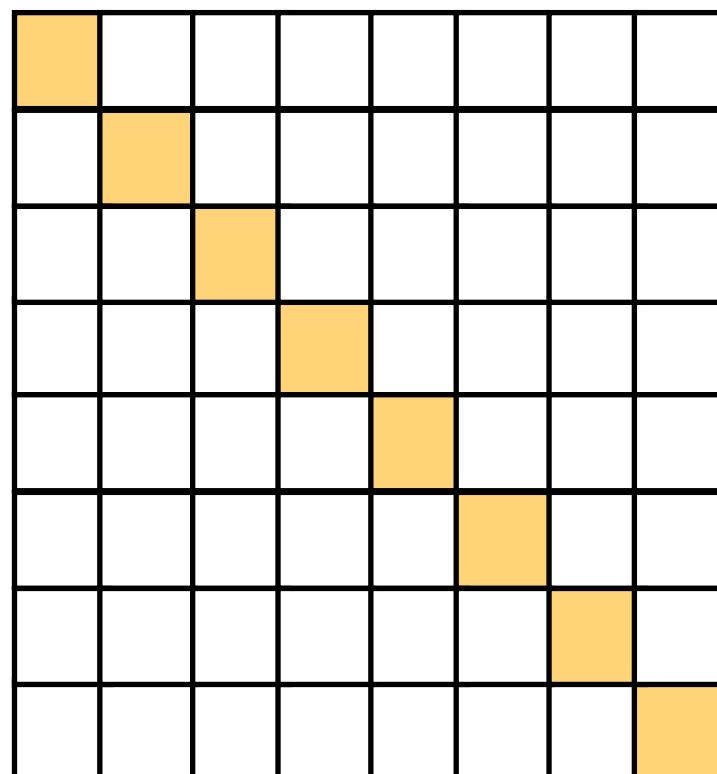
Sparse matrices arise in many applications

- Simulating climate
- Analyzing images (photos, MRIs,...)
- Web page ranking for search
- Graphs, including Graph Neural Nets

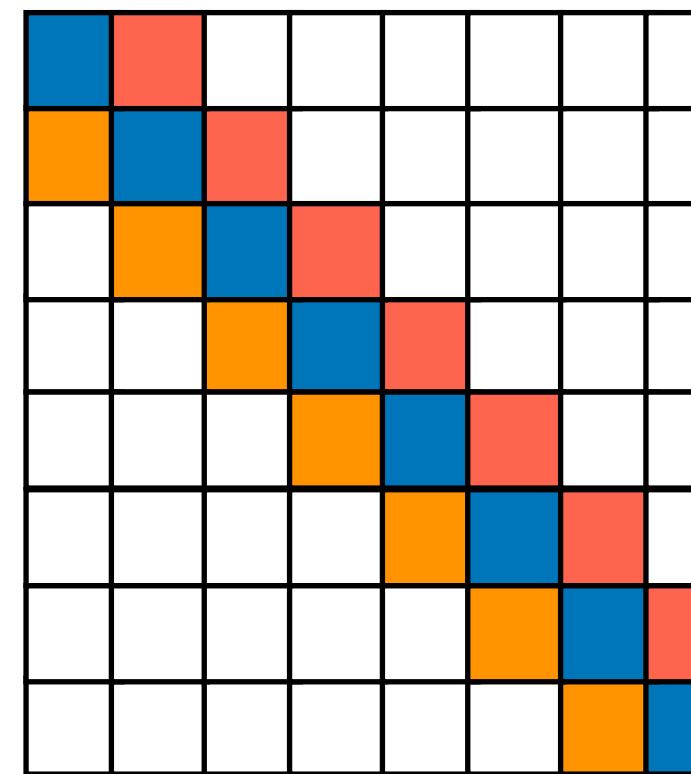


<https://sparse.tamu.edu/>

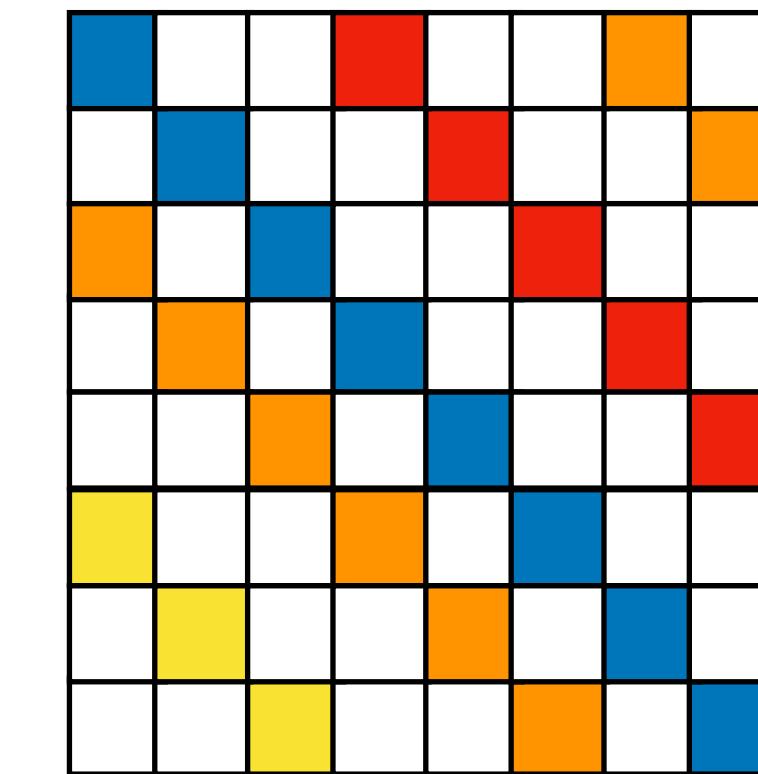
# Examples of sparse matrices



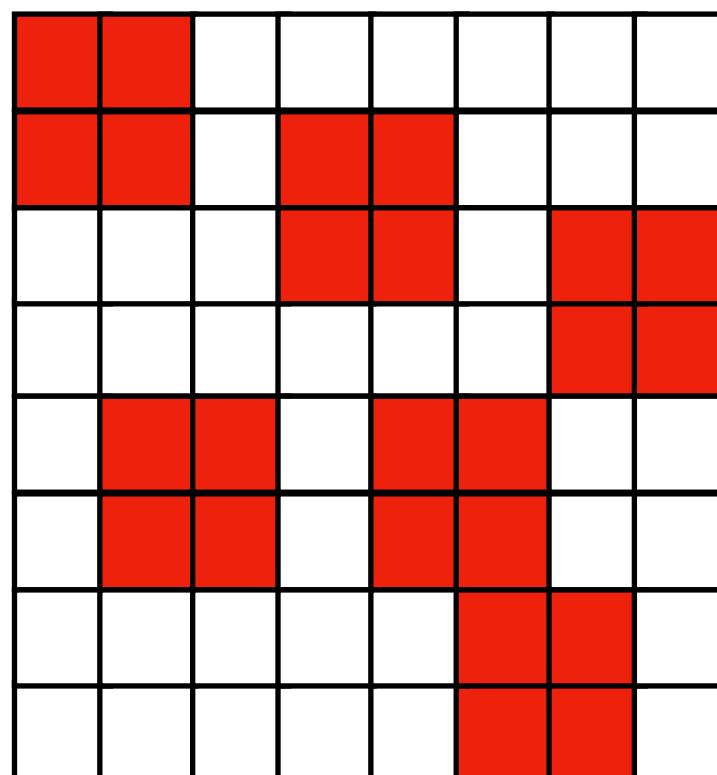
**Diagonal**



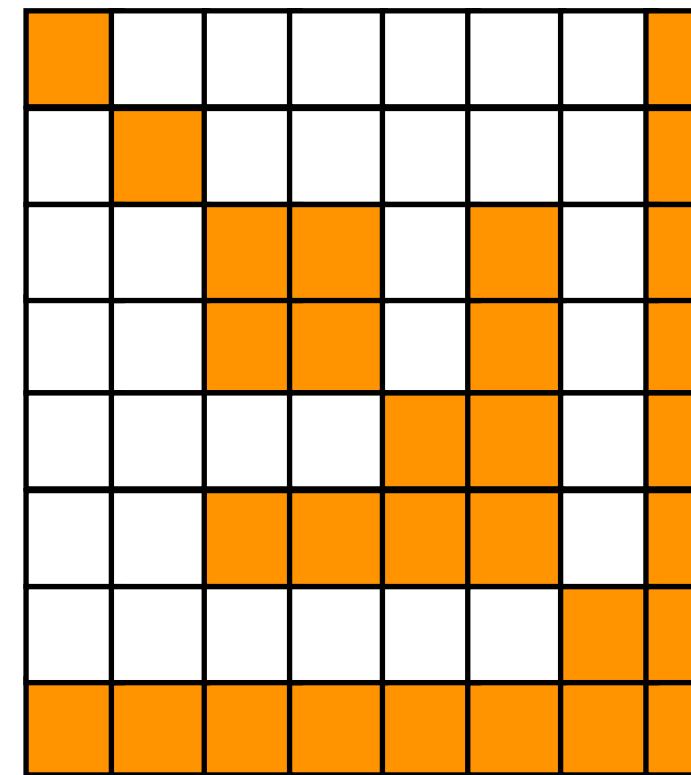
**Tridiagonal**



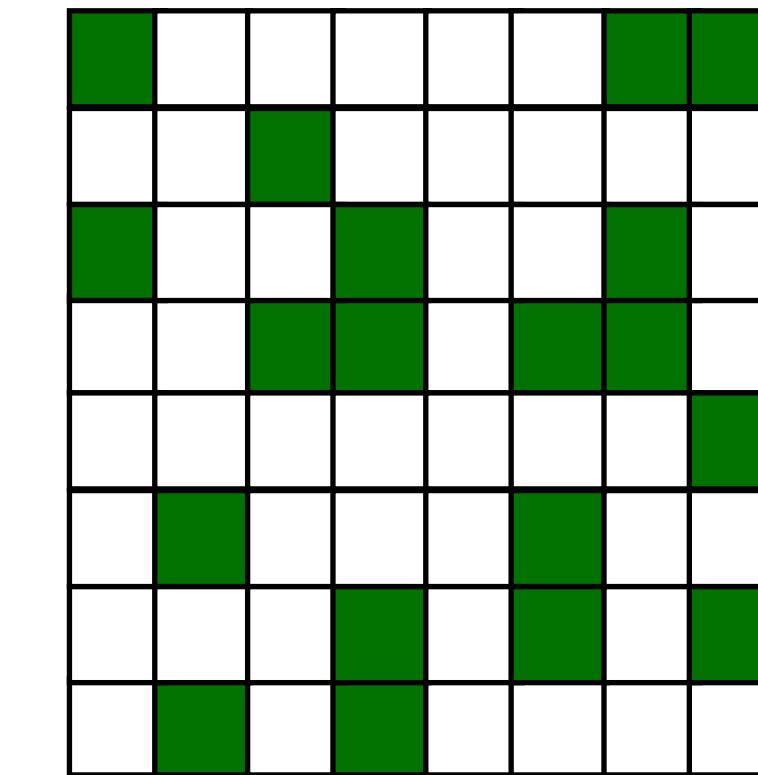
**“Generalized diagonal”**



**Block matrix  
(2x2 dense blocks)**



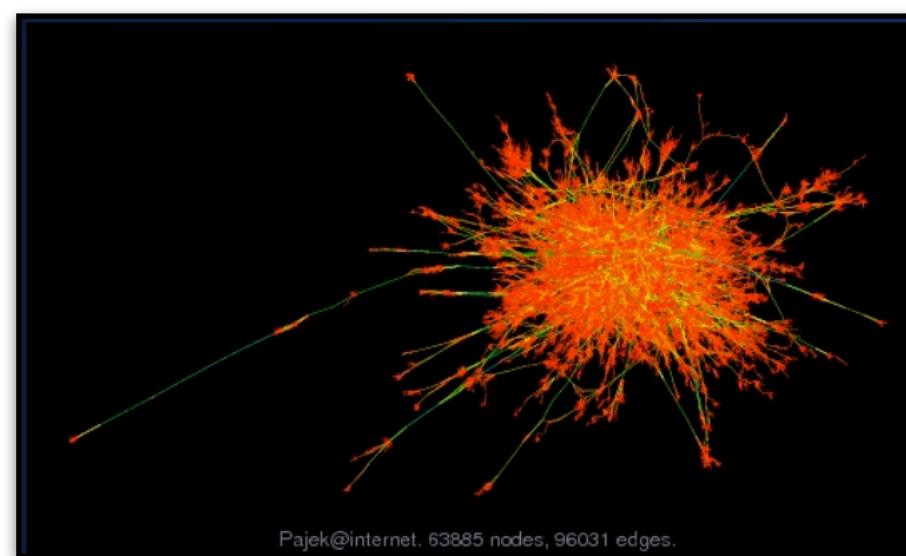
**Symmetric**



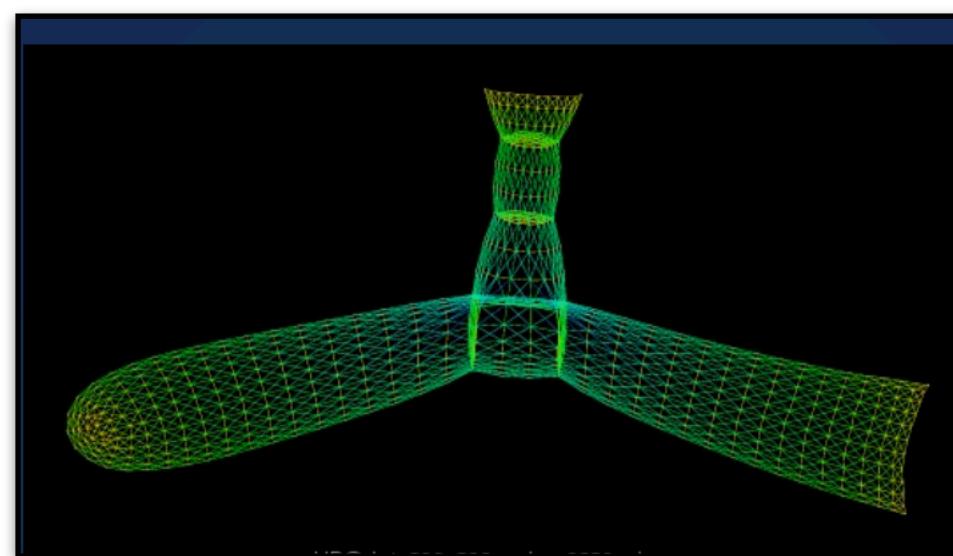
**Irregular**

# Sparse matrices are everywhere

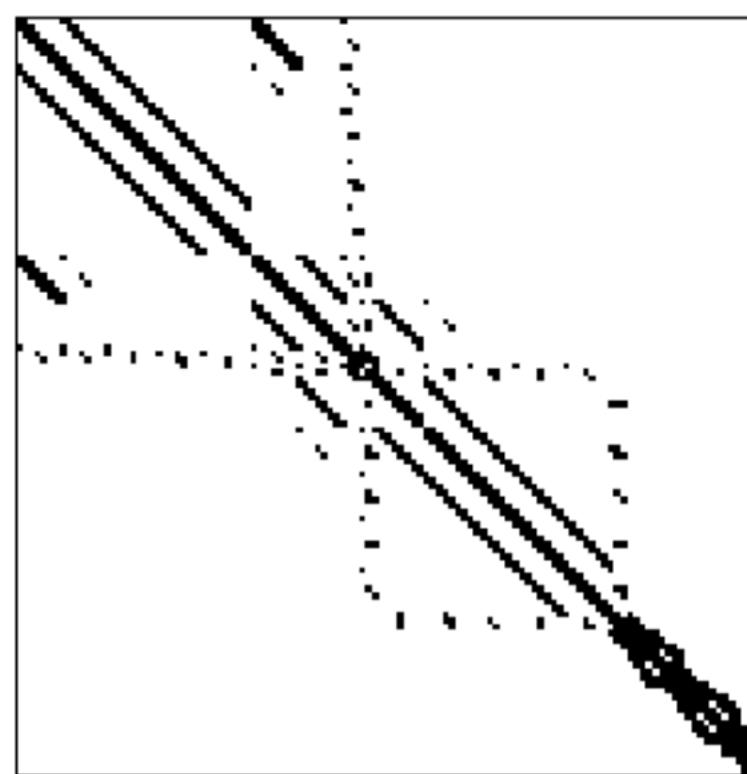
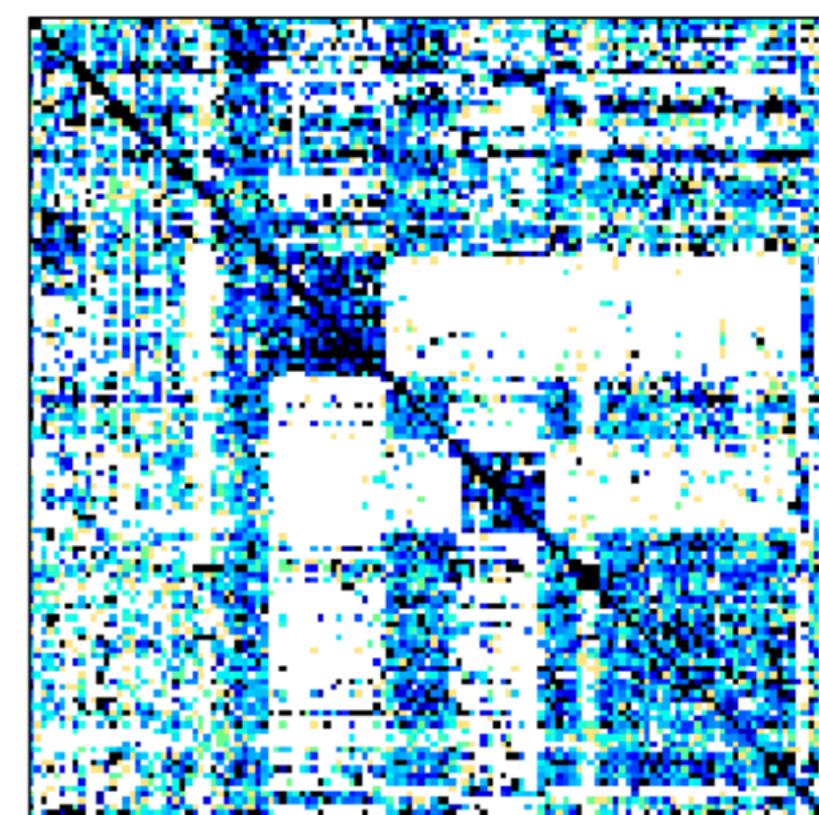
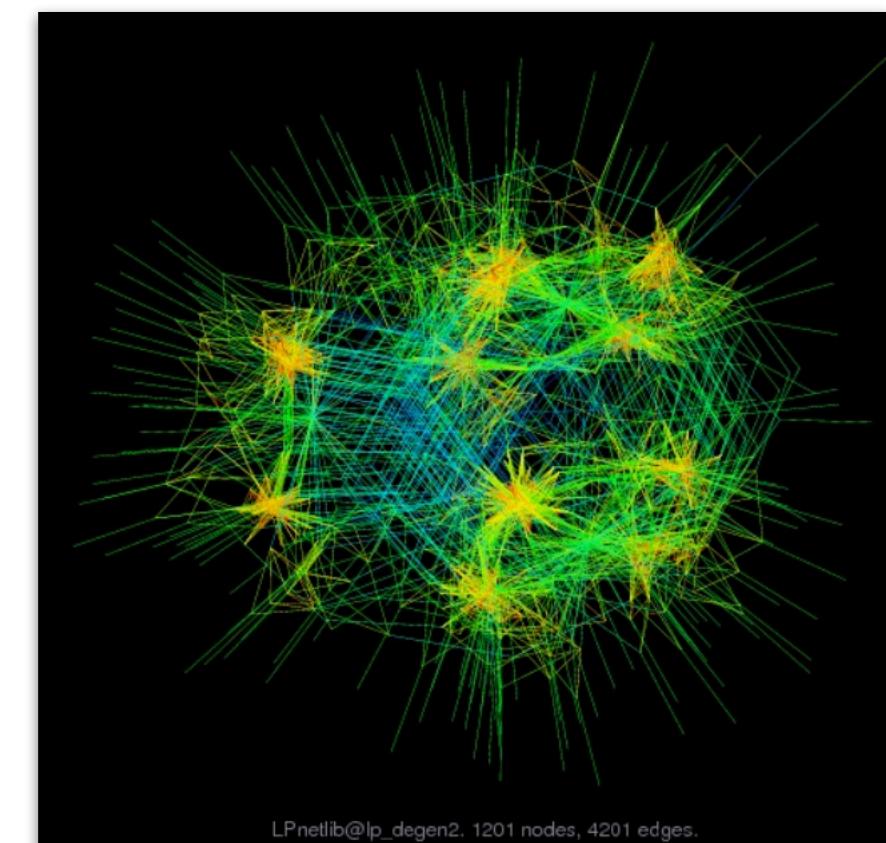
Internet connectivity



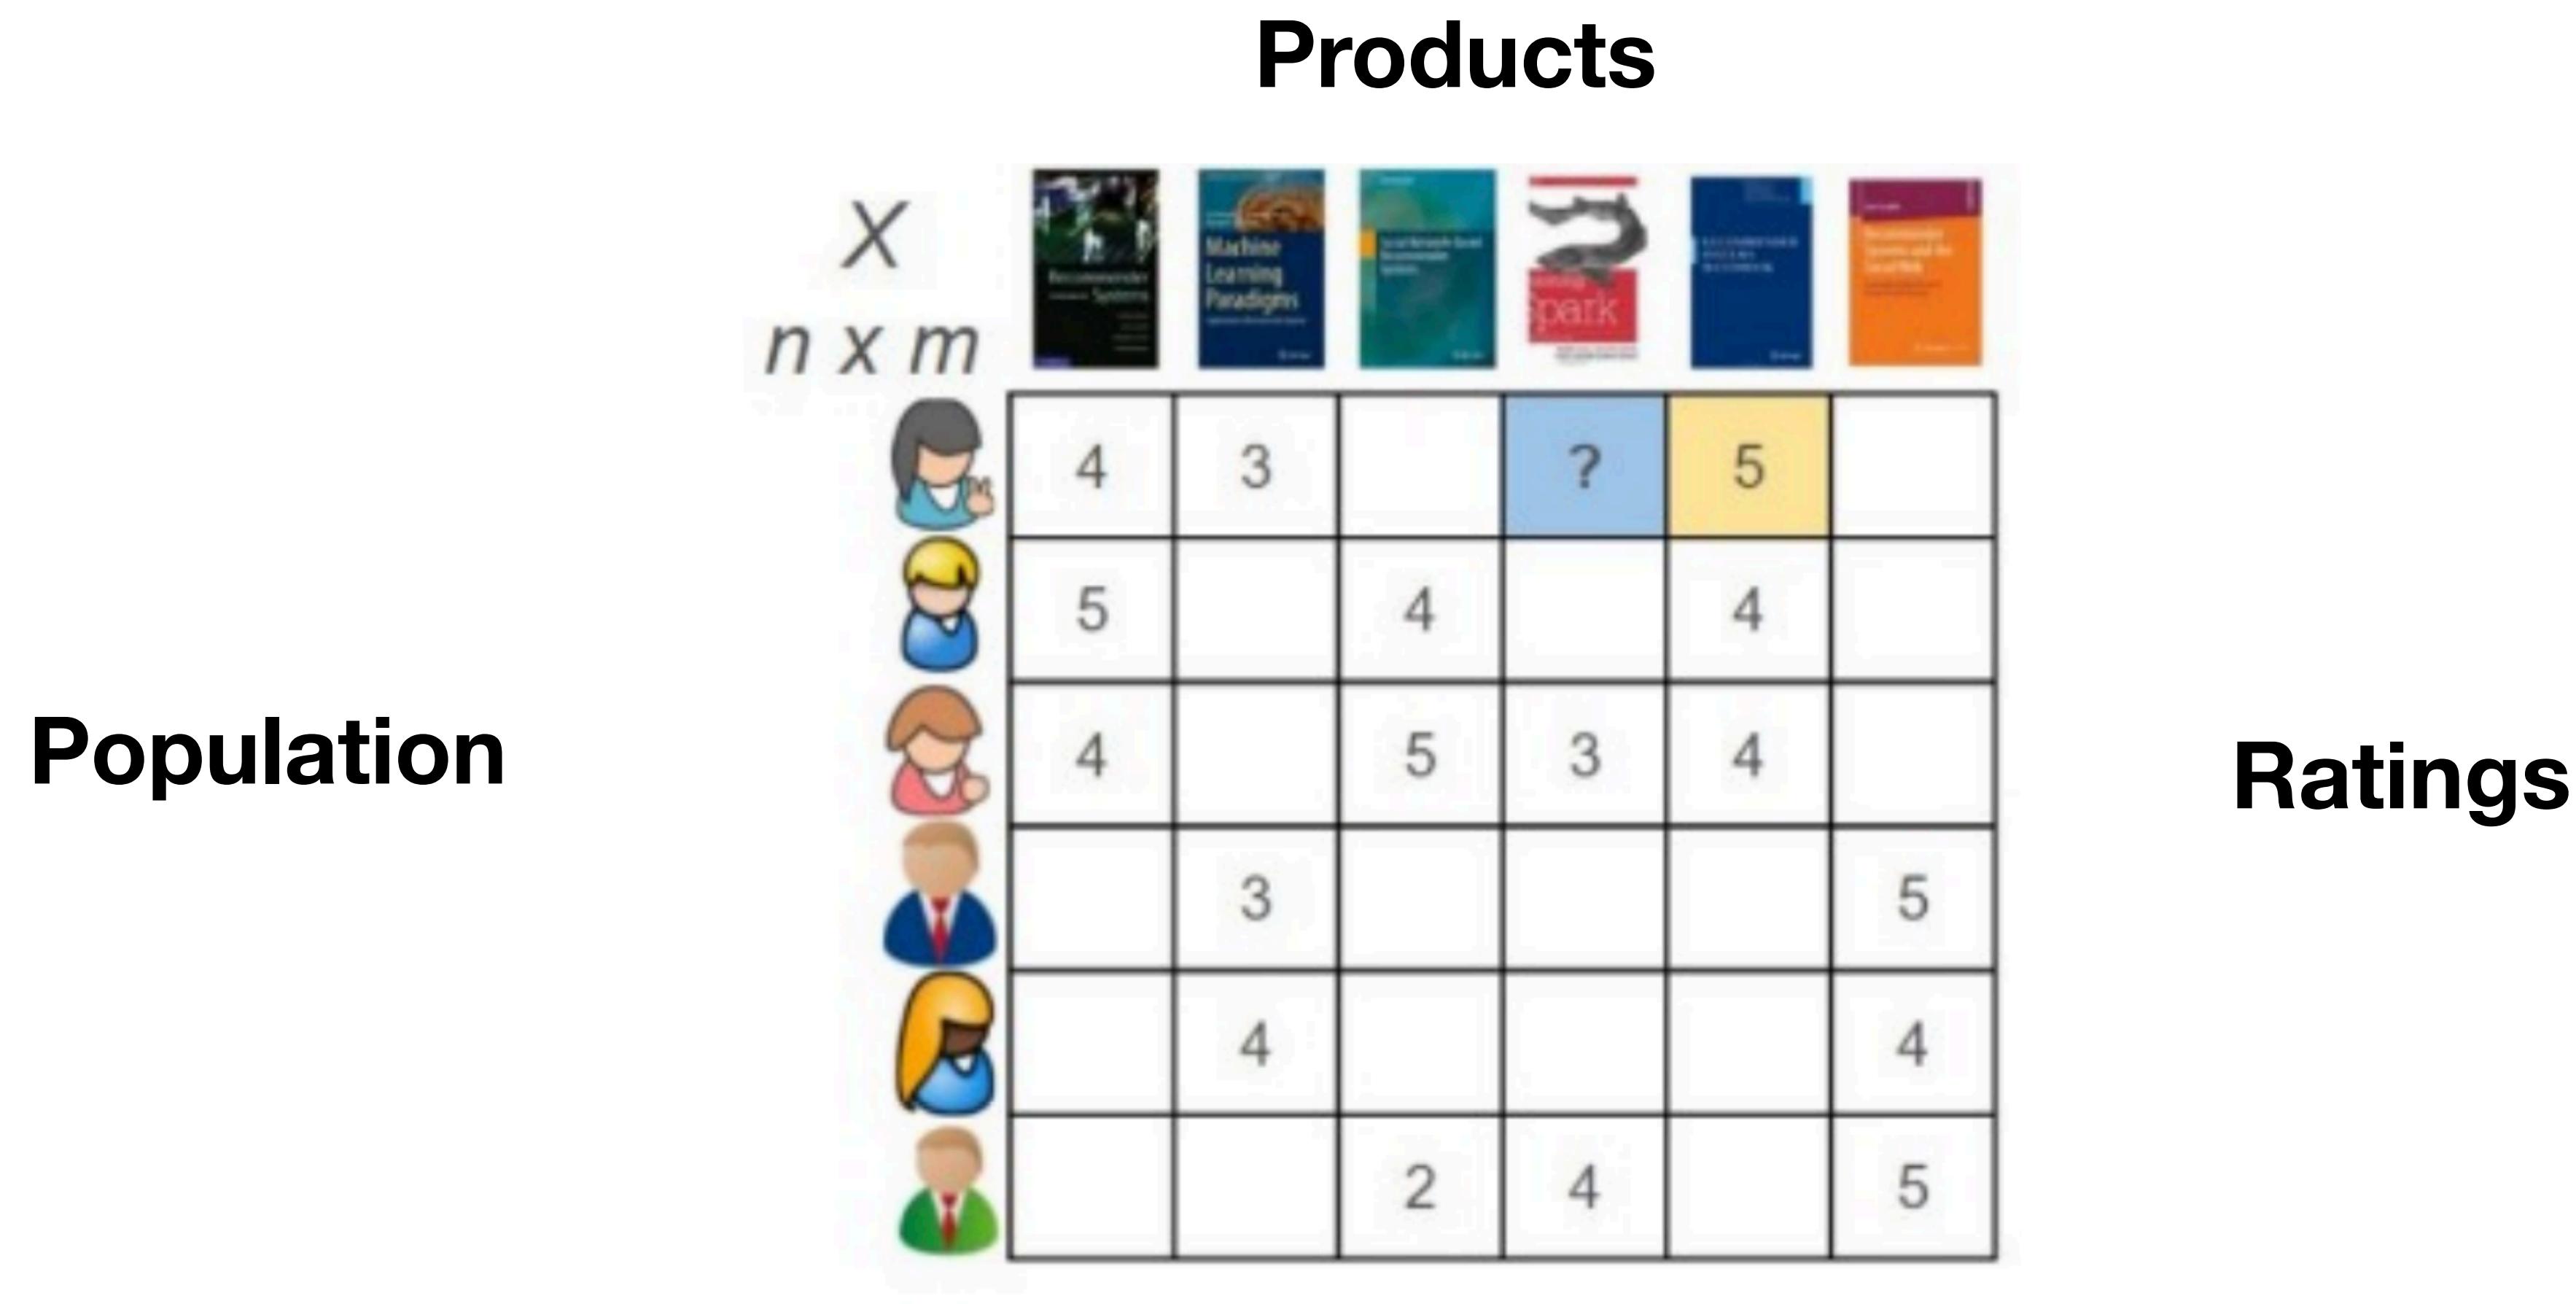
Structural design



Linear programming



# Recommendation matrix



# Applications involving sparse matrices

## Graphs

- Google's PageRank (originally an eigen-problem on the web adjacency matrix)
- Transportation network analysis

## Text analysis

- Latent Semantic Indexing finds topics in a document corpus by doing a singular value decomposition of a bag-of-words matrix

## Scientific and engineering

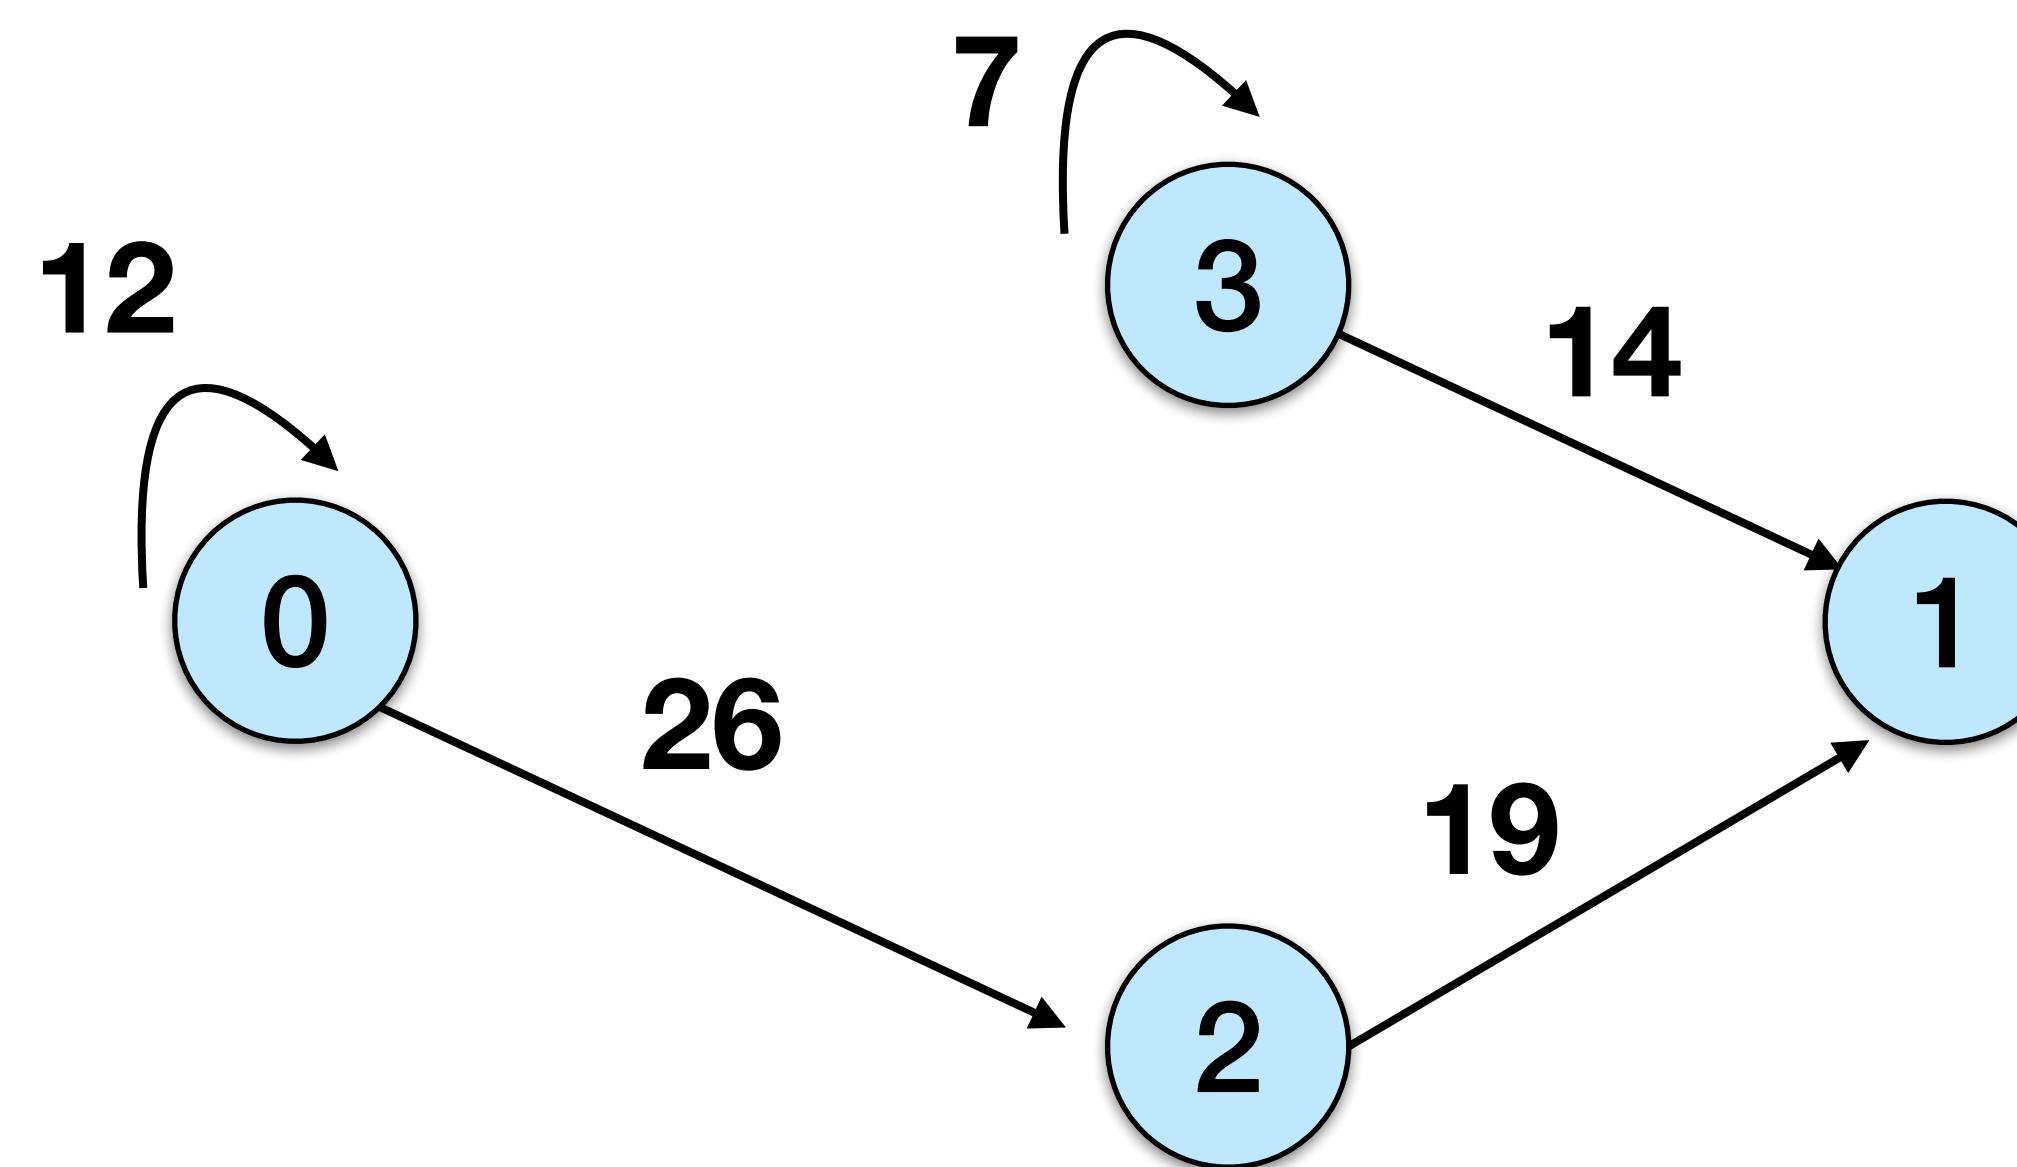
- Solving differential equations (climate modeling, etc.)
- Optimization problems

And many more...

# **Sparse Matrix Formats**

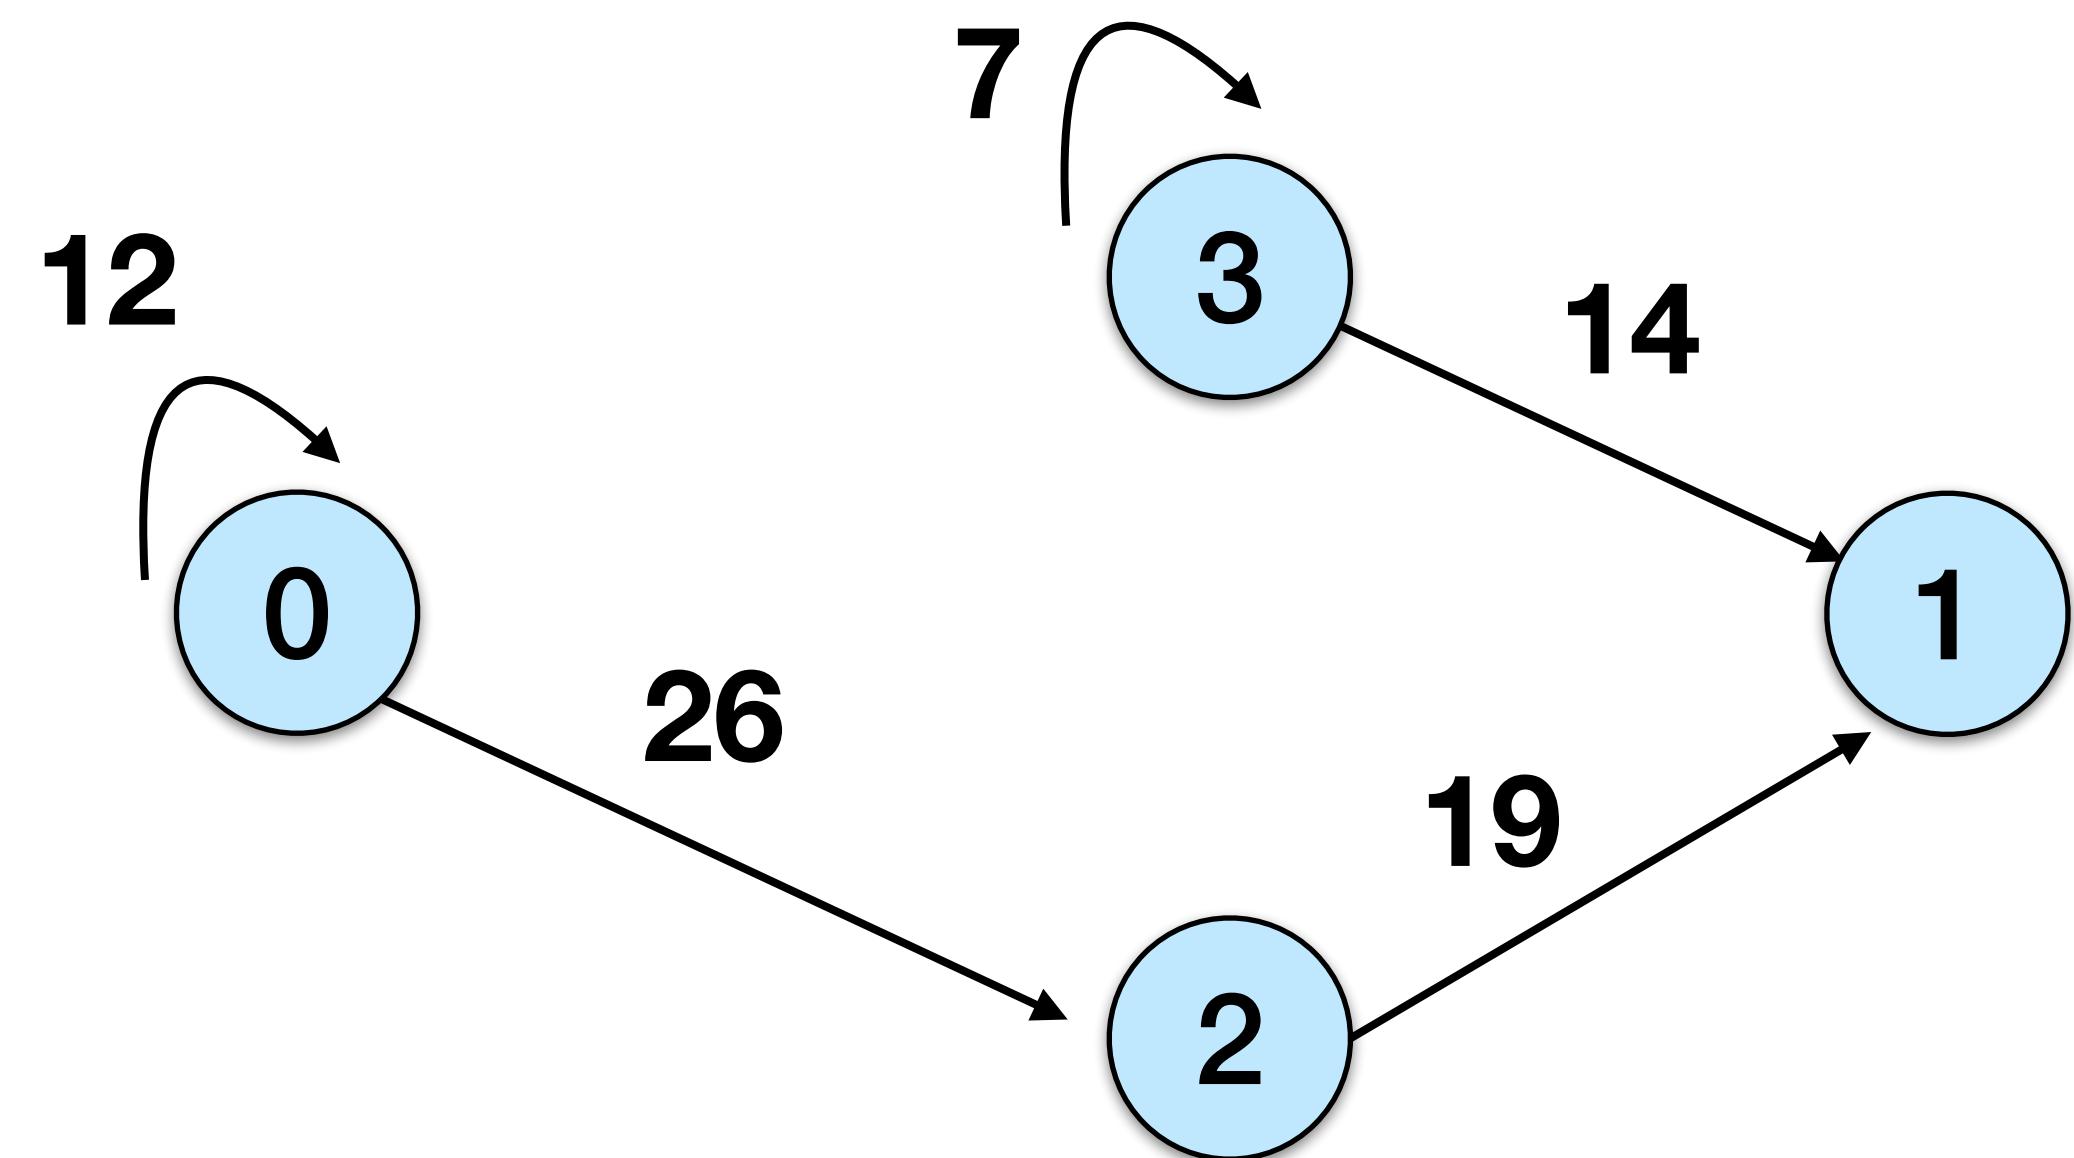
# Adjacency matrix representation

Any sparse matrix representation can be used for sparse graphs, and vice versa.



0	1	2	3
0	12		26
1			
2	19		
3	14		7

# Coordinate representation (COO)

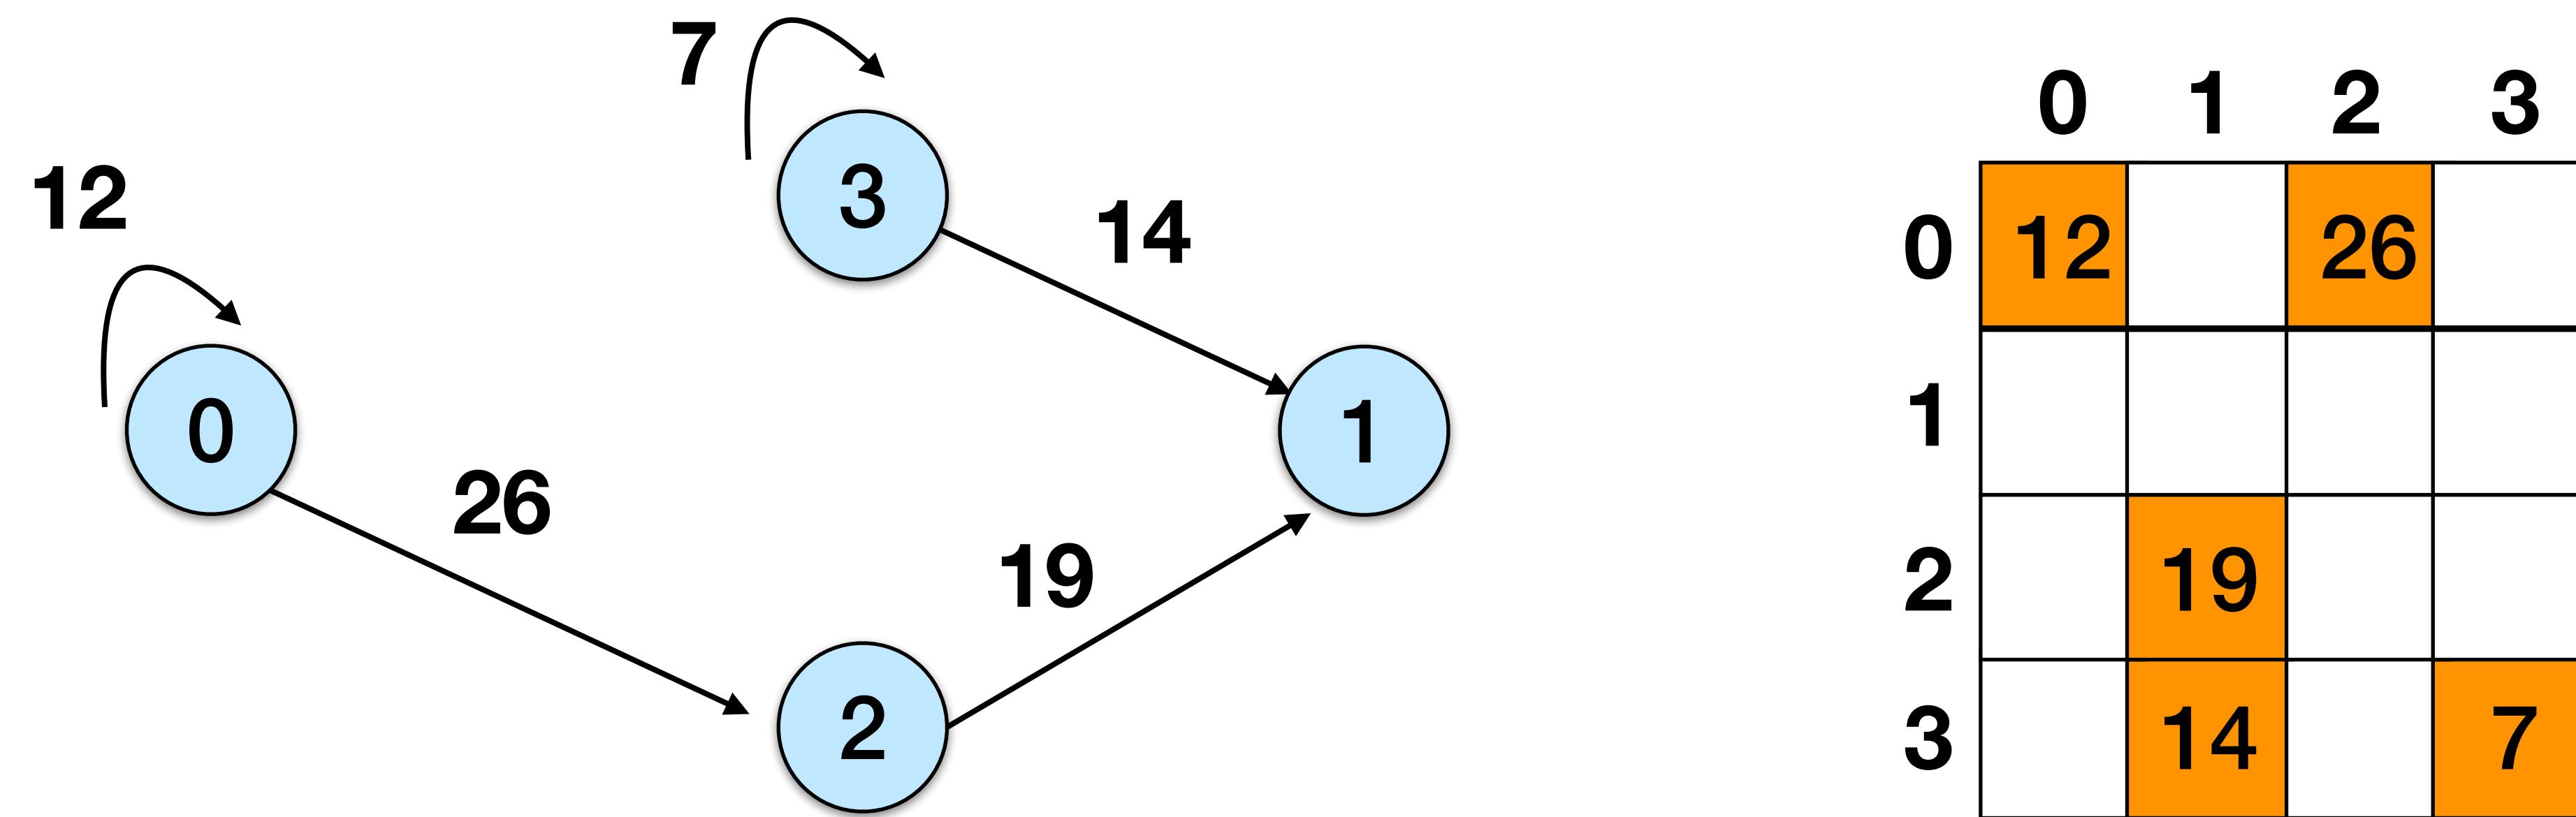


	0	1	2	3
0	12		26	
1				
2		19		
3			14	7

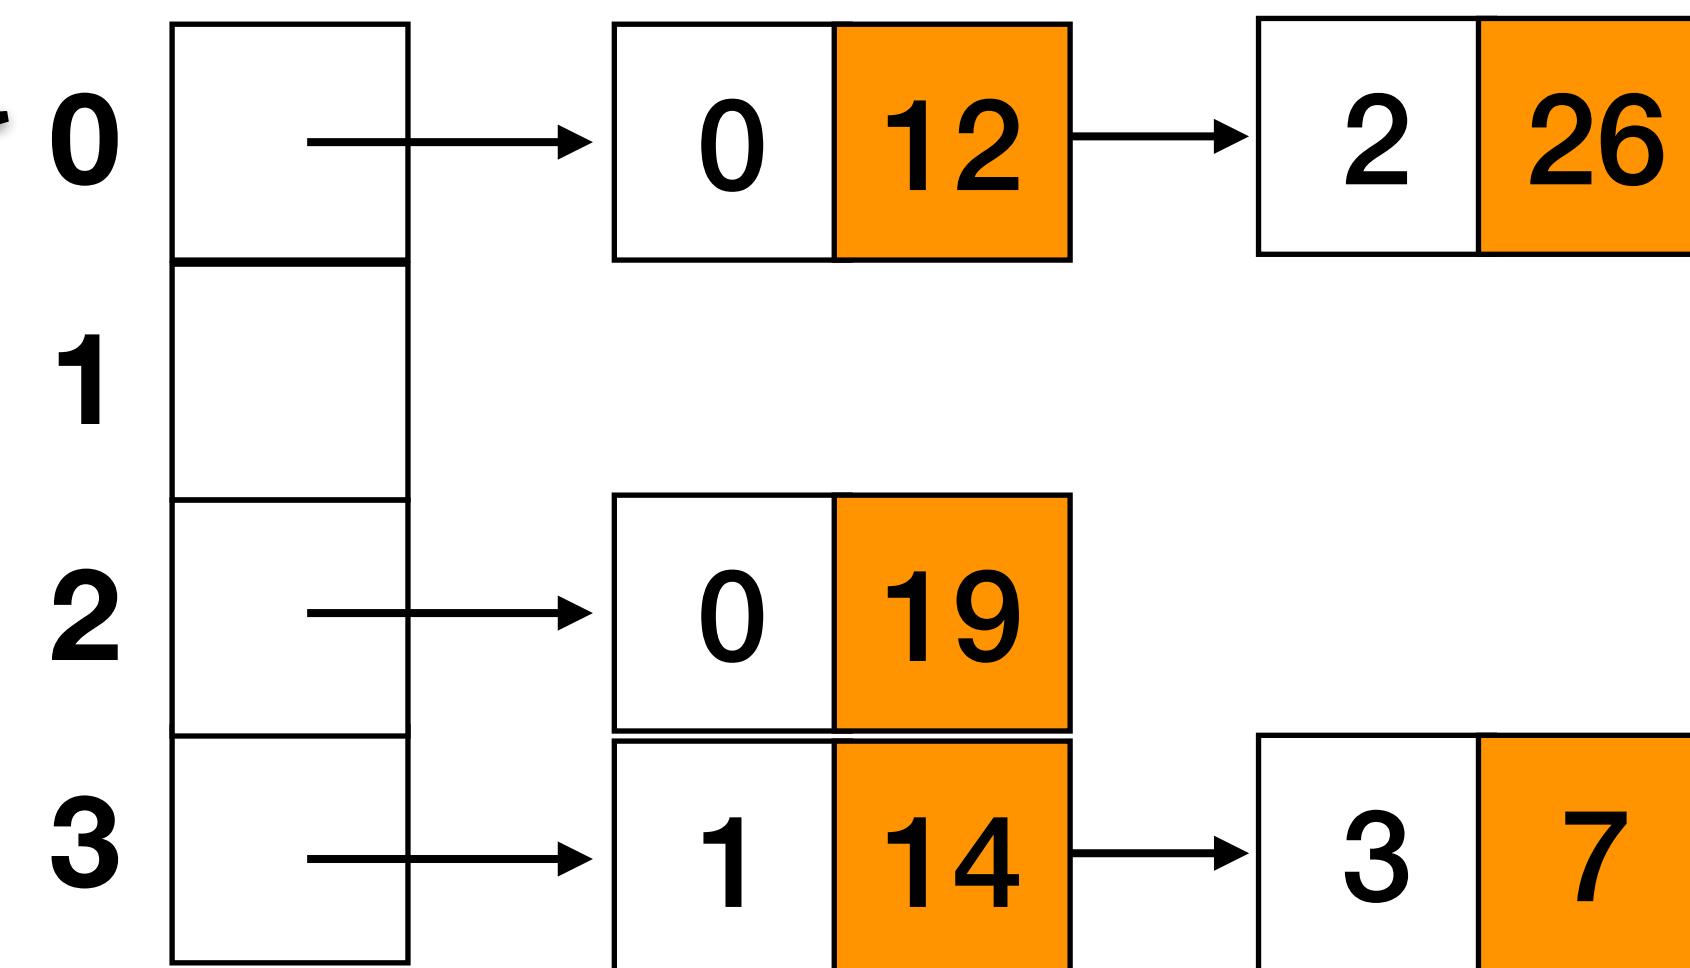
(0, 0, 12)  
(0, 2, 26)  
(2, 1, 19)  
(3, 1, 14)  
(3, 3, 7)

row + column index +  
weight per nonzero  
(easy to build / modify)

# Adjacency lists

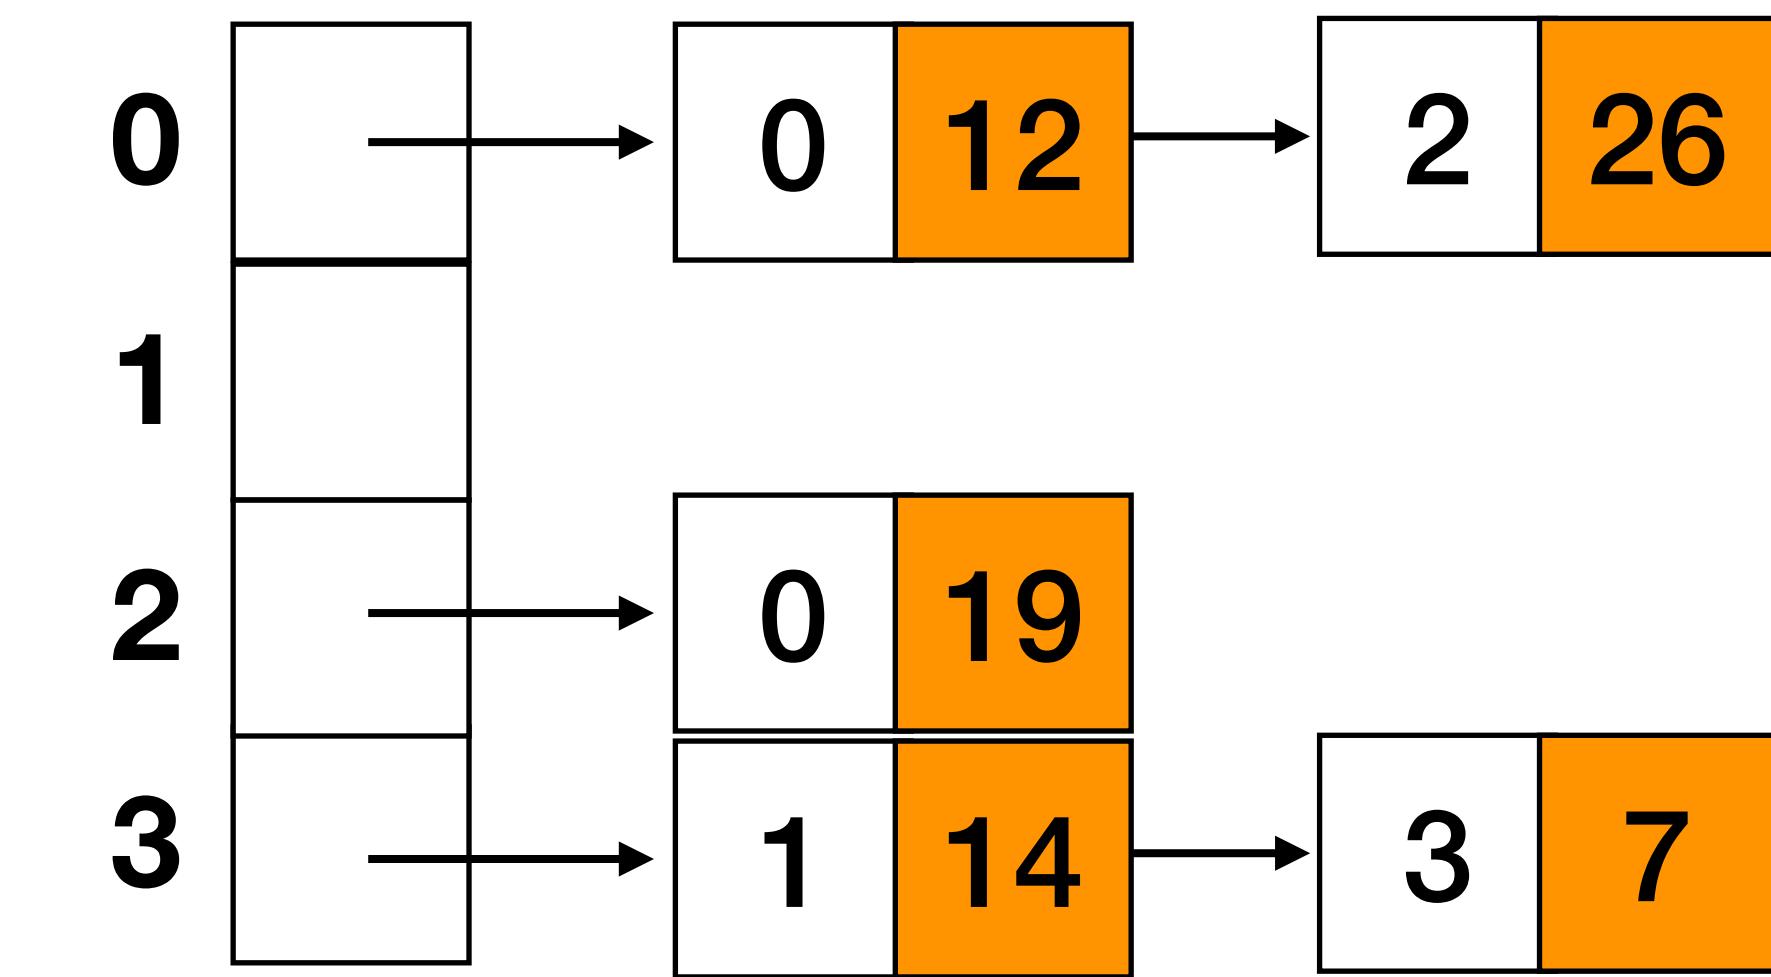
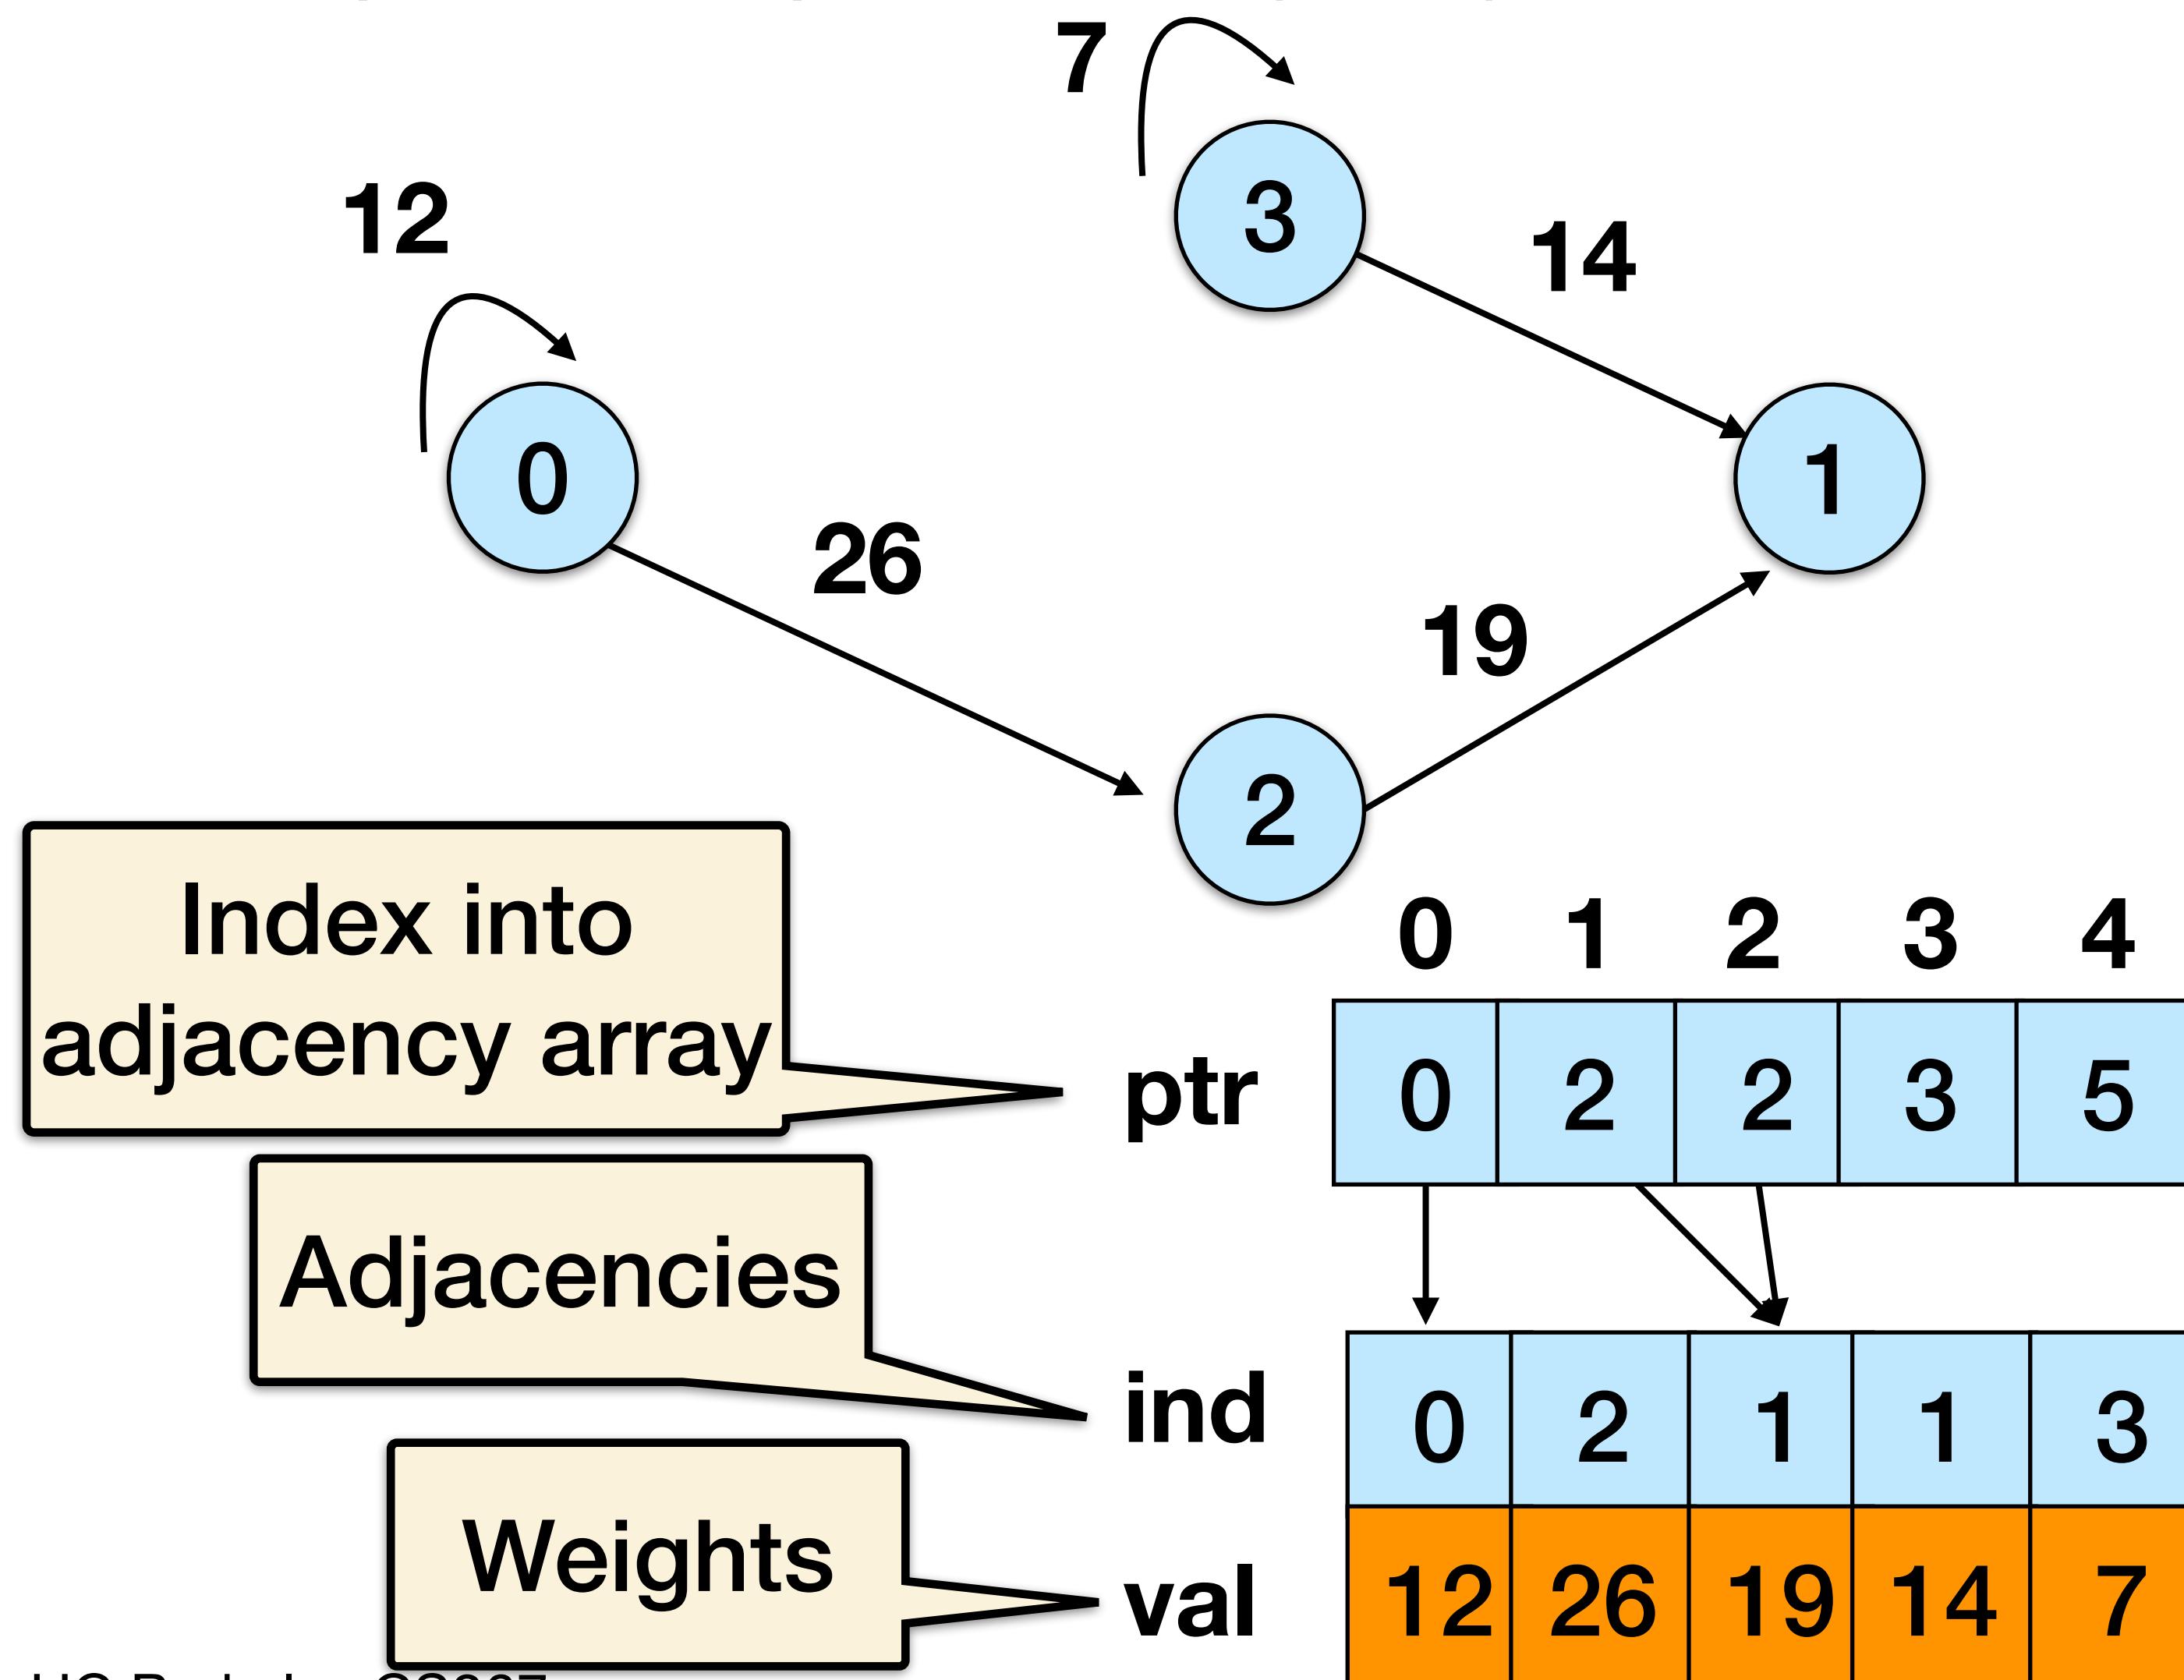


Source stored  
implicitly



# Compressed Sparse Row (CSR)

Compressed sparse row (CSR) = **cache-efficient adjacency lists**



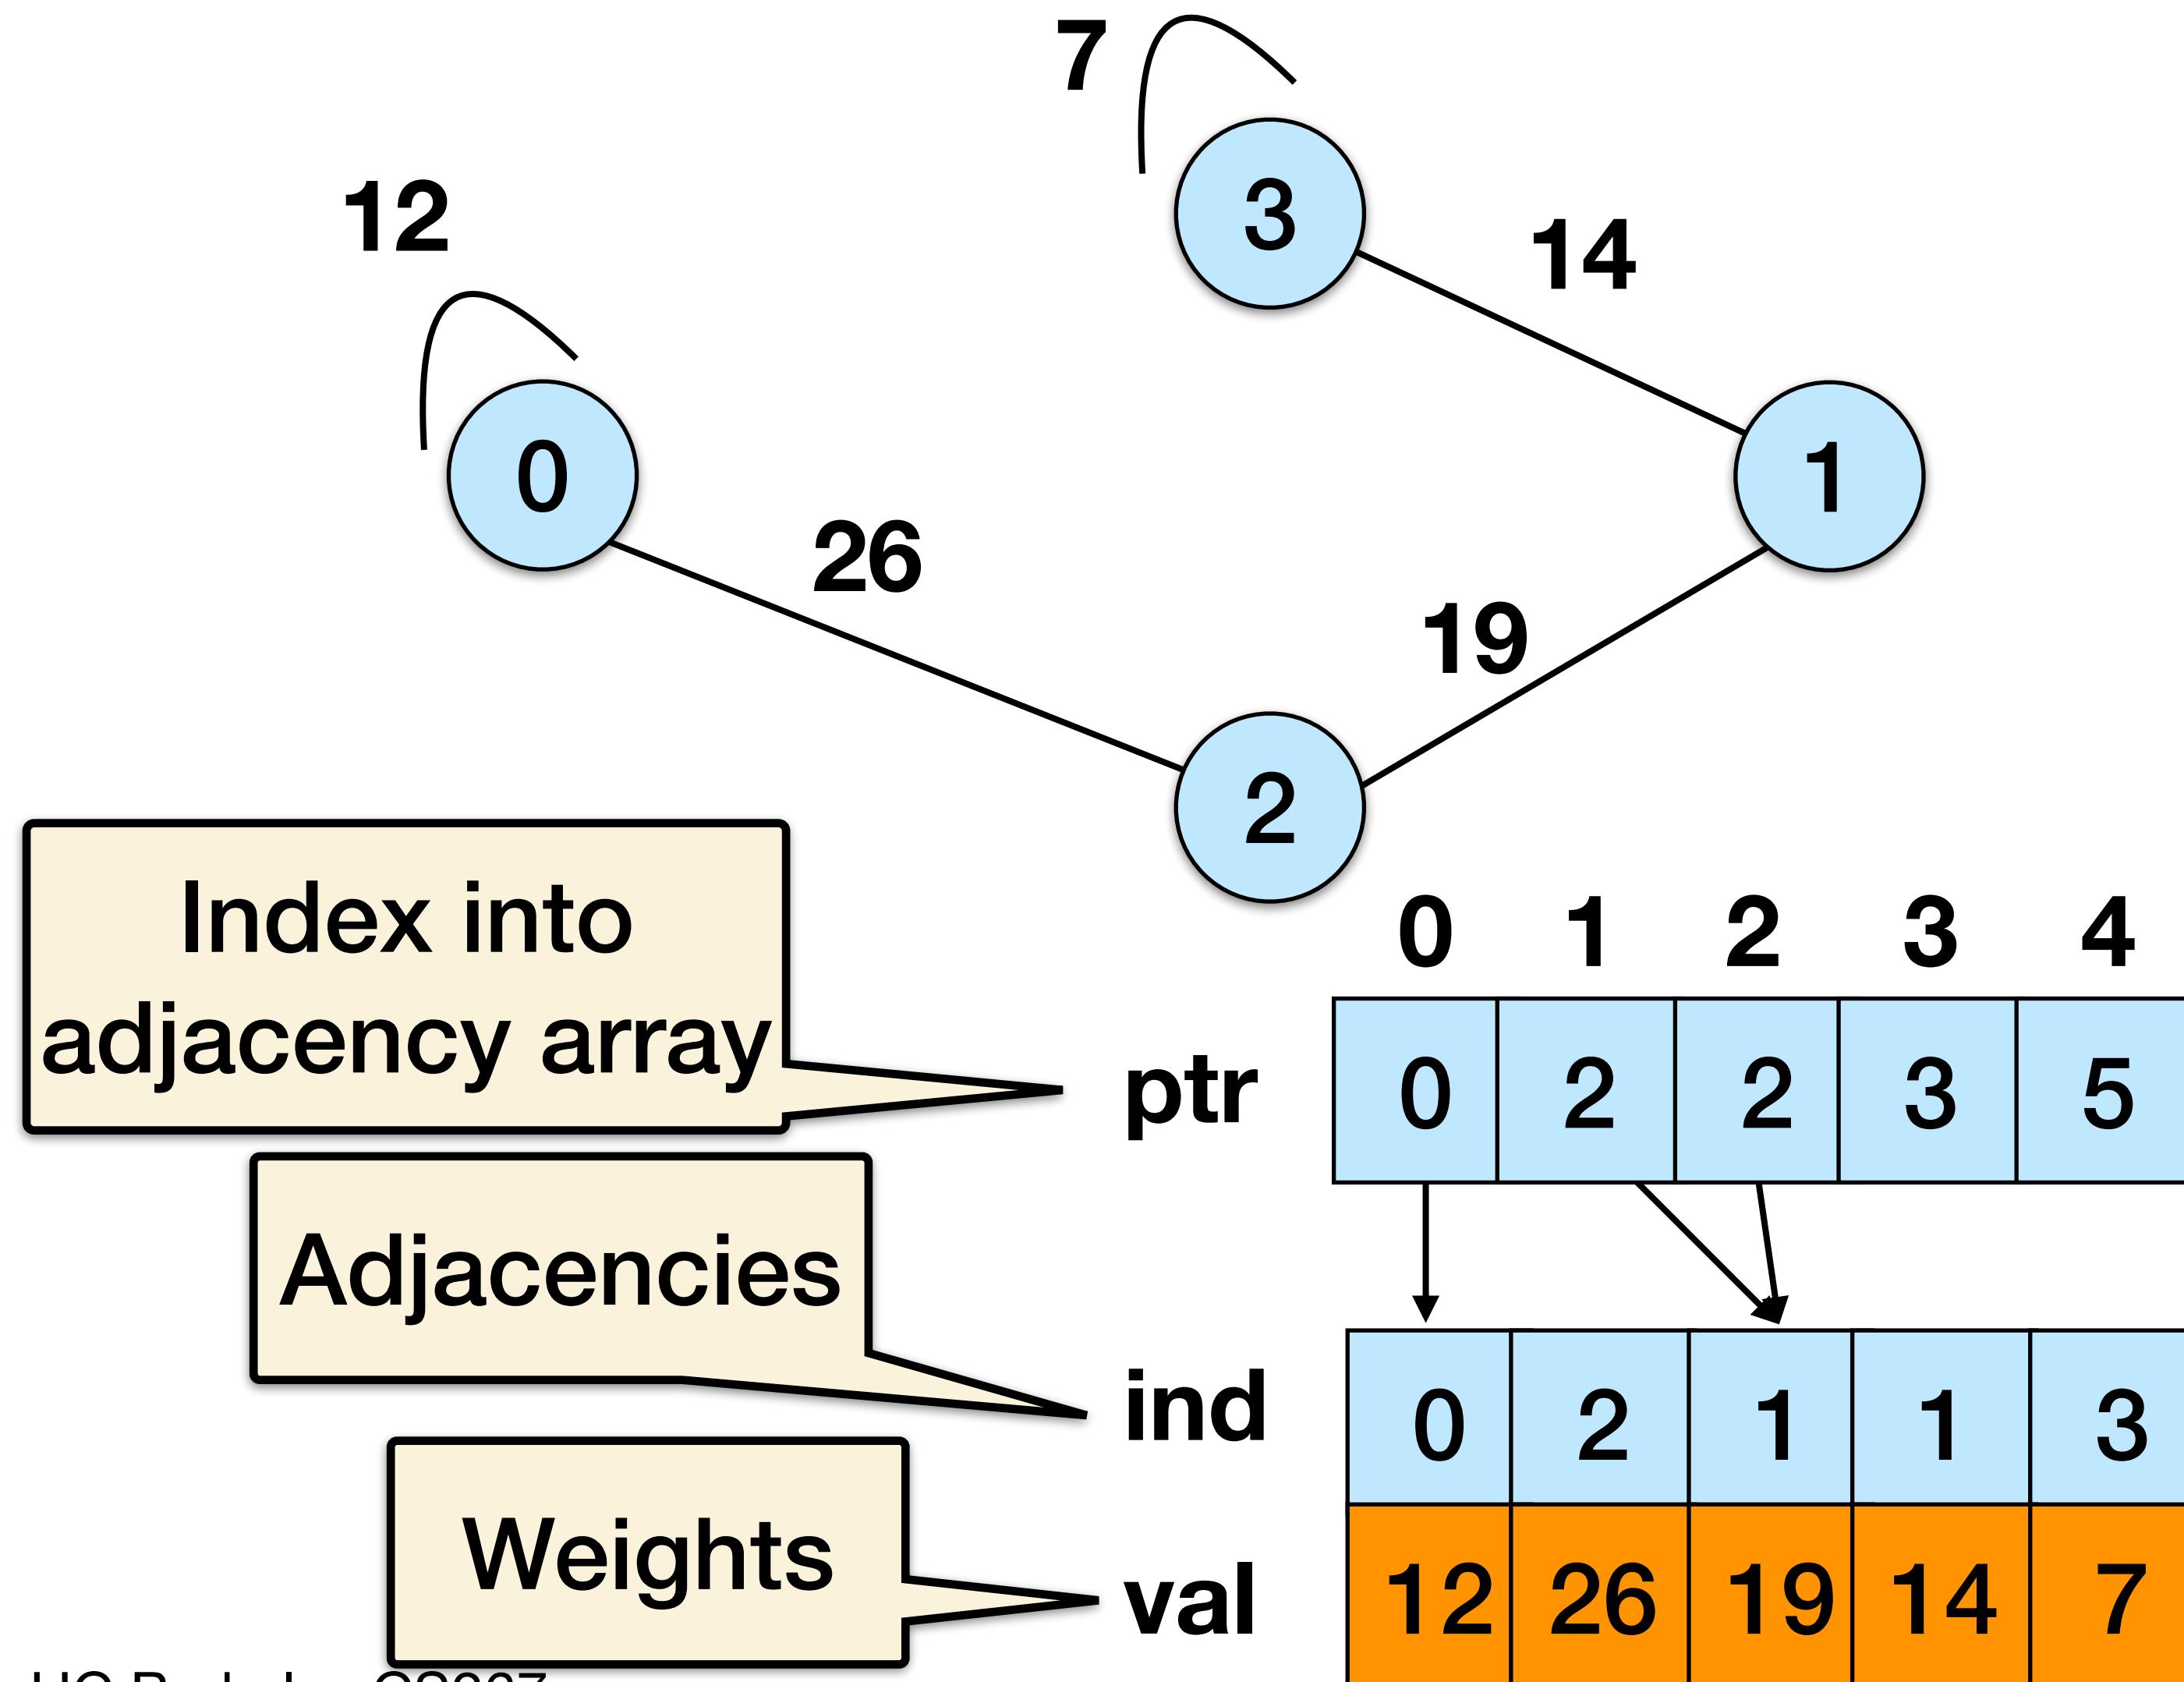
(row starts in CSR)

(column ids in CSR)

(numerical values in CSR)

# Directed vs Undirected

An undirected graph corresponds to a symmetric matrix - only need to store **half**



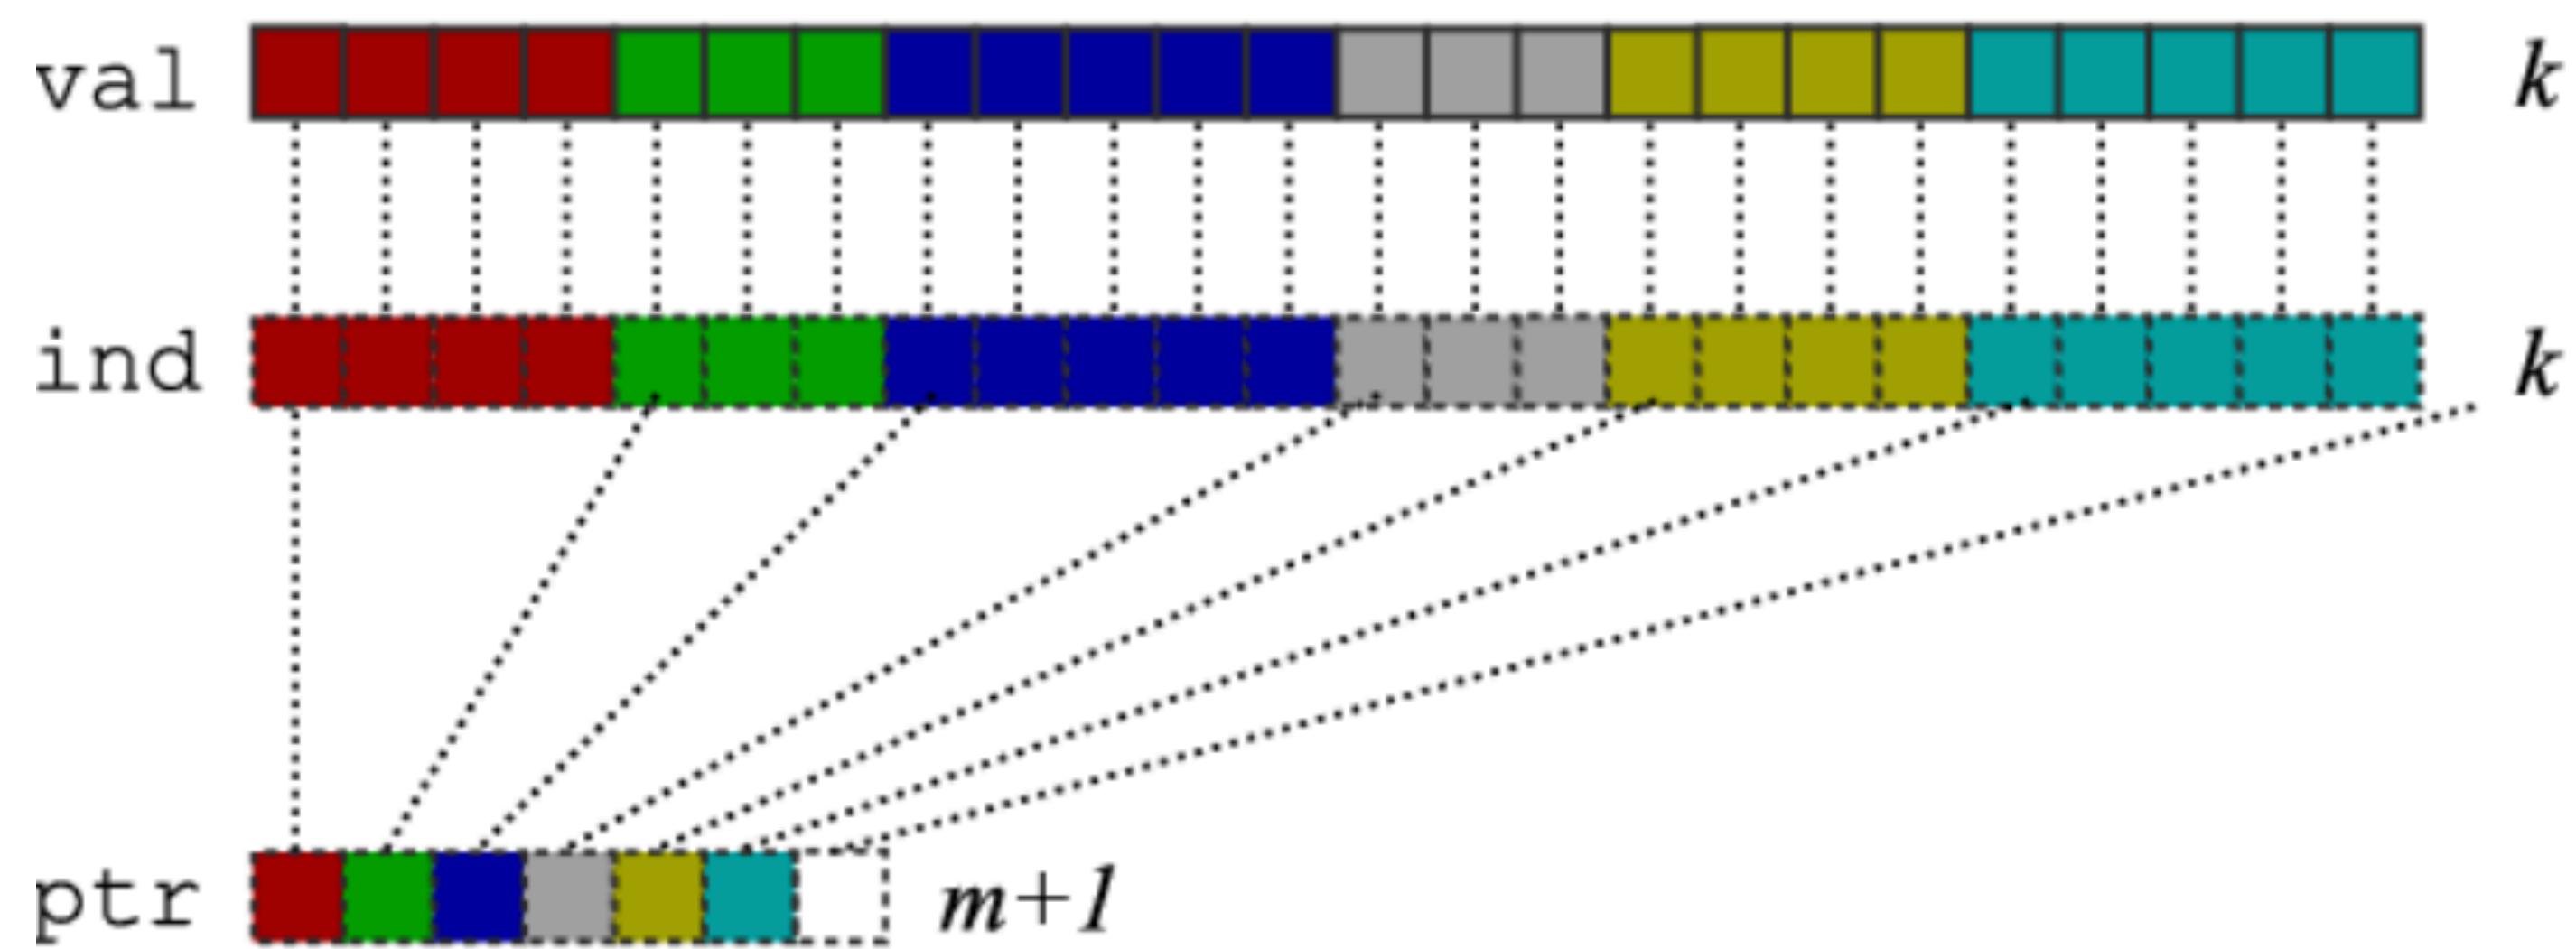
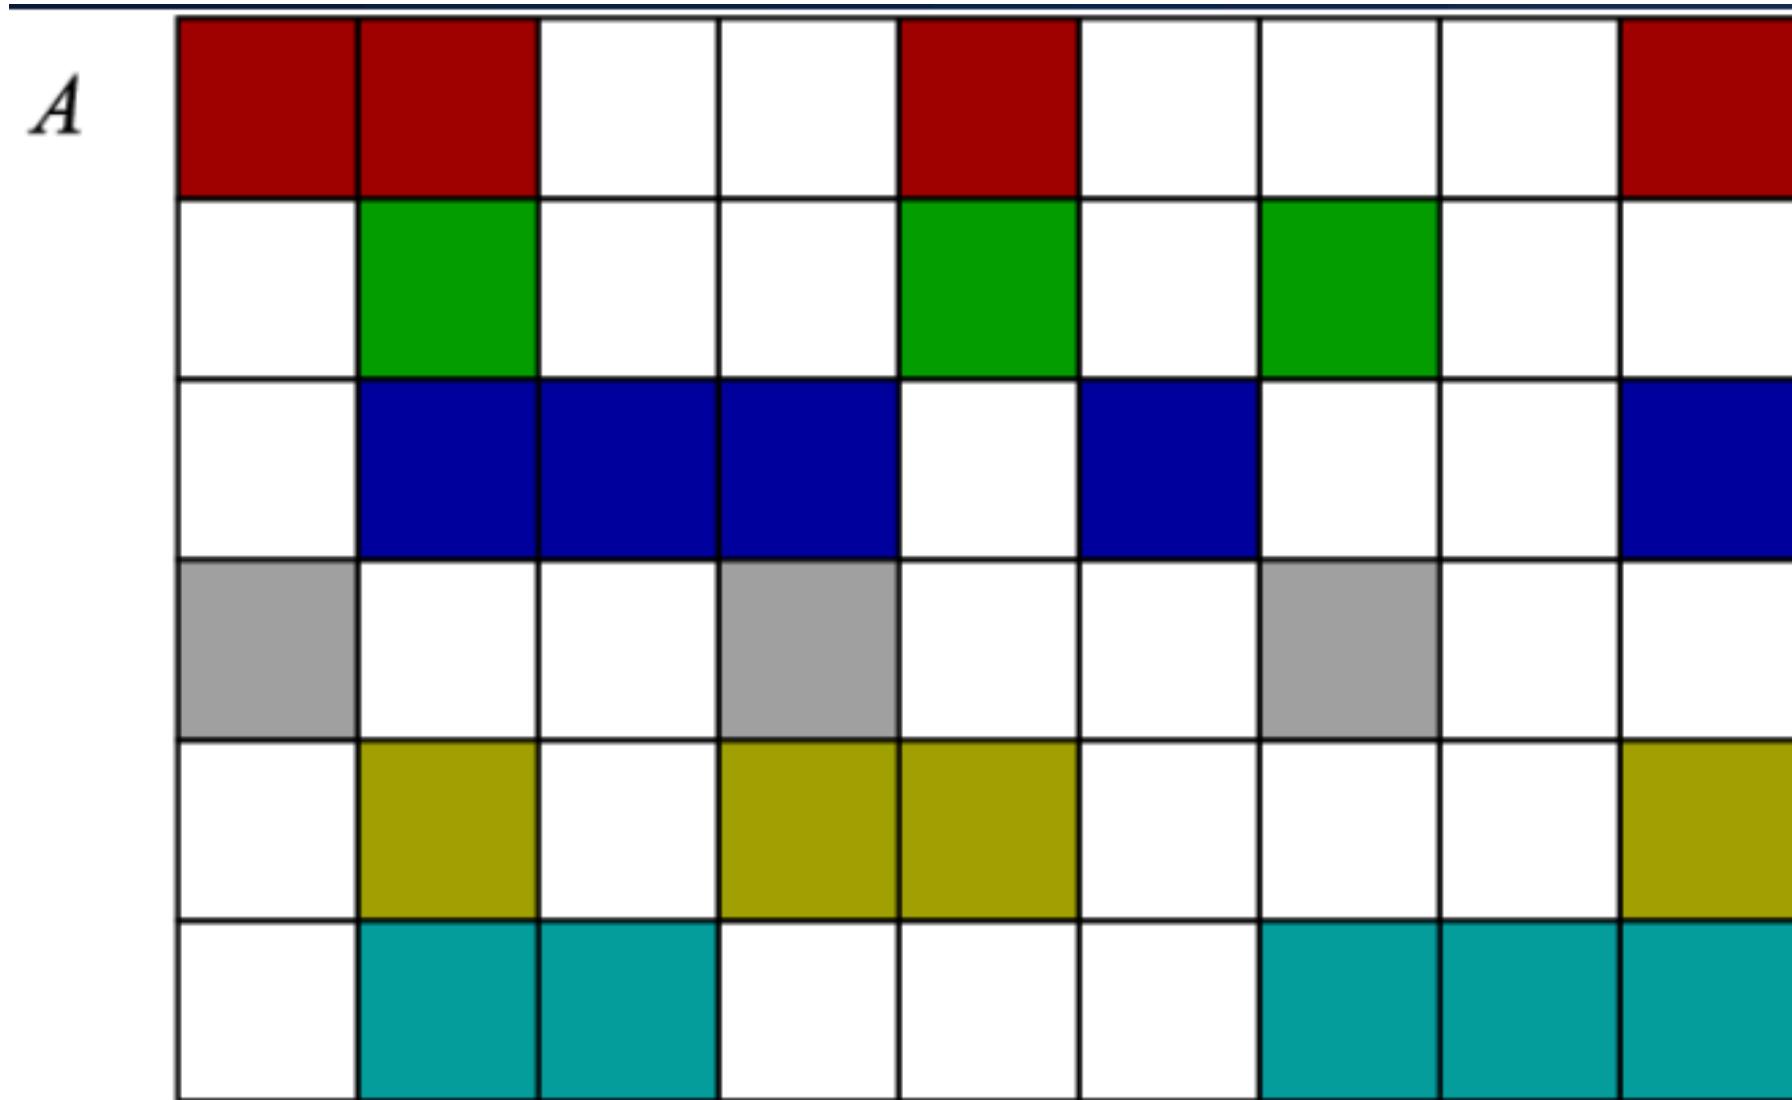
	0	1	2	3
0	12		26	
1		19	14	
2	26	19		
3		14		7

(row starts in CSR)

(column ids in CSR)

(numerical values in CSR)

# Compressed Sparse Row (CSR) Storage



CSR has:

- Size **nnz** = number of nonzeros
- Array of the nonzero **values** (**val**) of size nnz
- Array of the column **indices** (**ind**) for each value of size nnz
- Array of row start **pointers** (**ptr**) of size n = number of rows

# Other Storage Formats

A 6x10 grid of colored squares. The colors are arranged in a repeating pattern: Red (top row), White, Green, Blue, Grey, and Yellow. Each color appears exactly once in every row and column. The grid is bounded by black lines.

# Compressed Sparse Row (CSR) is the most common and our focus today

# **Others include**

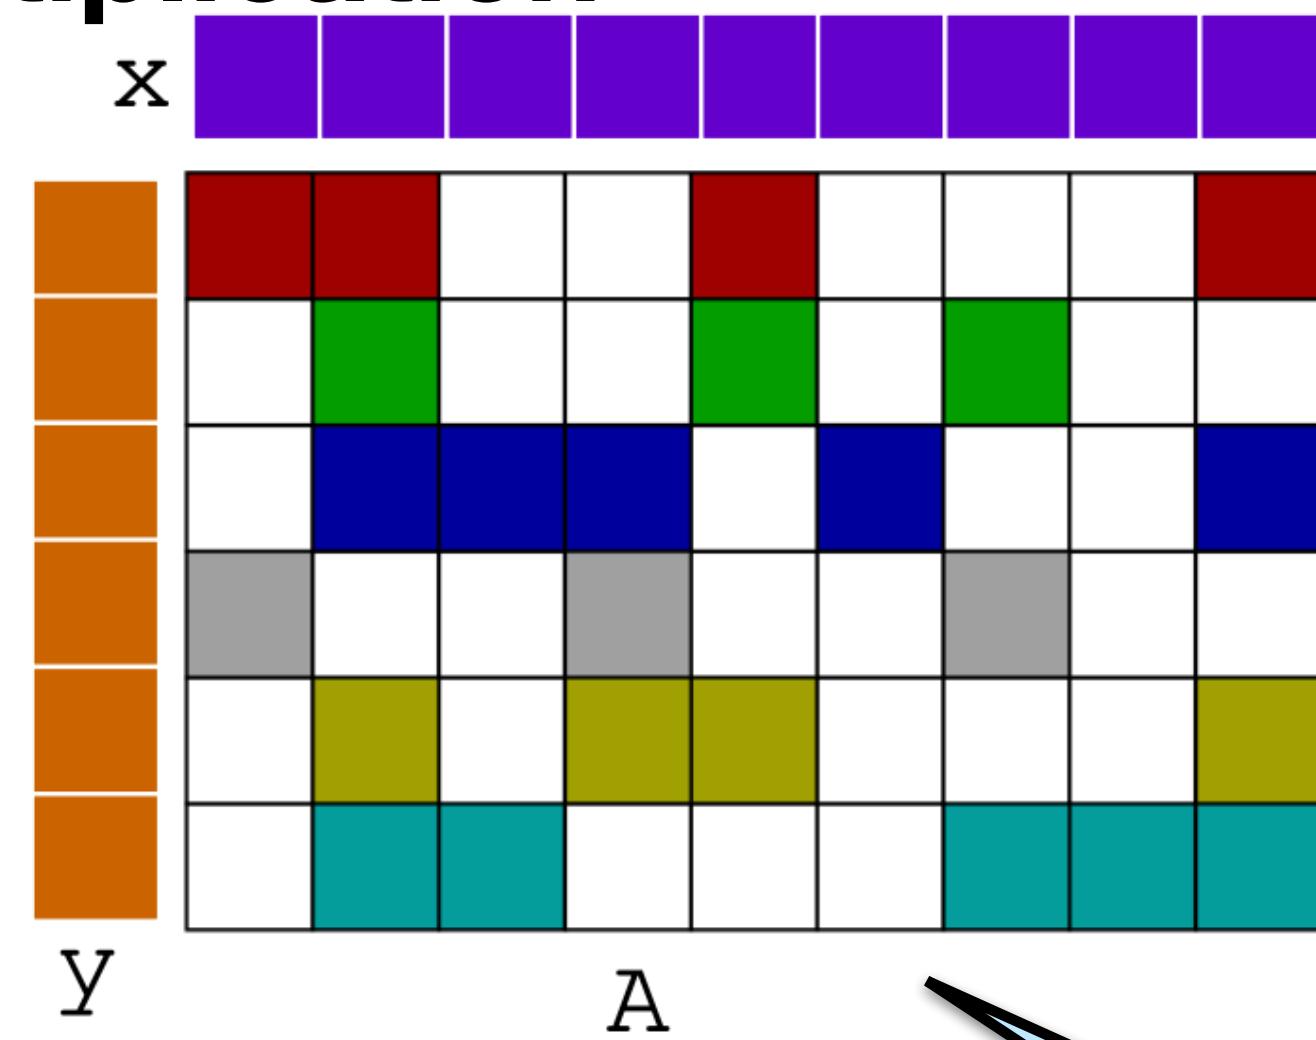
- Compressed Sparse Column (CSC)
  - Diagonal (DIAG): store main diagonal as 1D array; or diagonal bands as 2D (padded)
  - Symmetric: store only  $\frac{1}{2}$  the array (indexing more complicated)
  - **Blocked**: store each block contiguously
    - **Register blocked**: blocks are small and dense, avoid indexes within blocks
    - **Cache blocked**: blocks are large and themselves sparse

...and many other specialized formats!

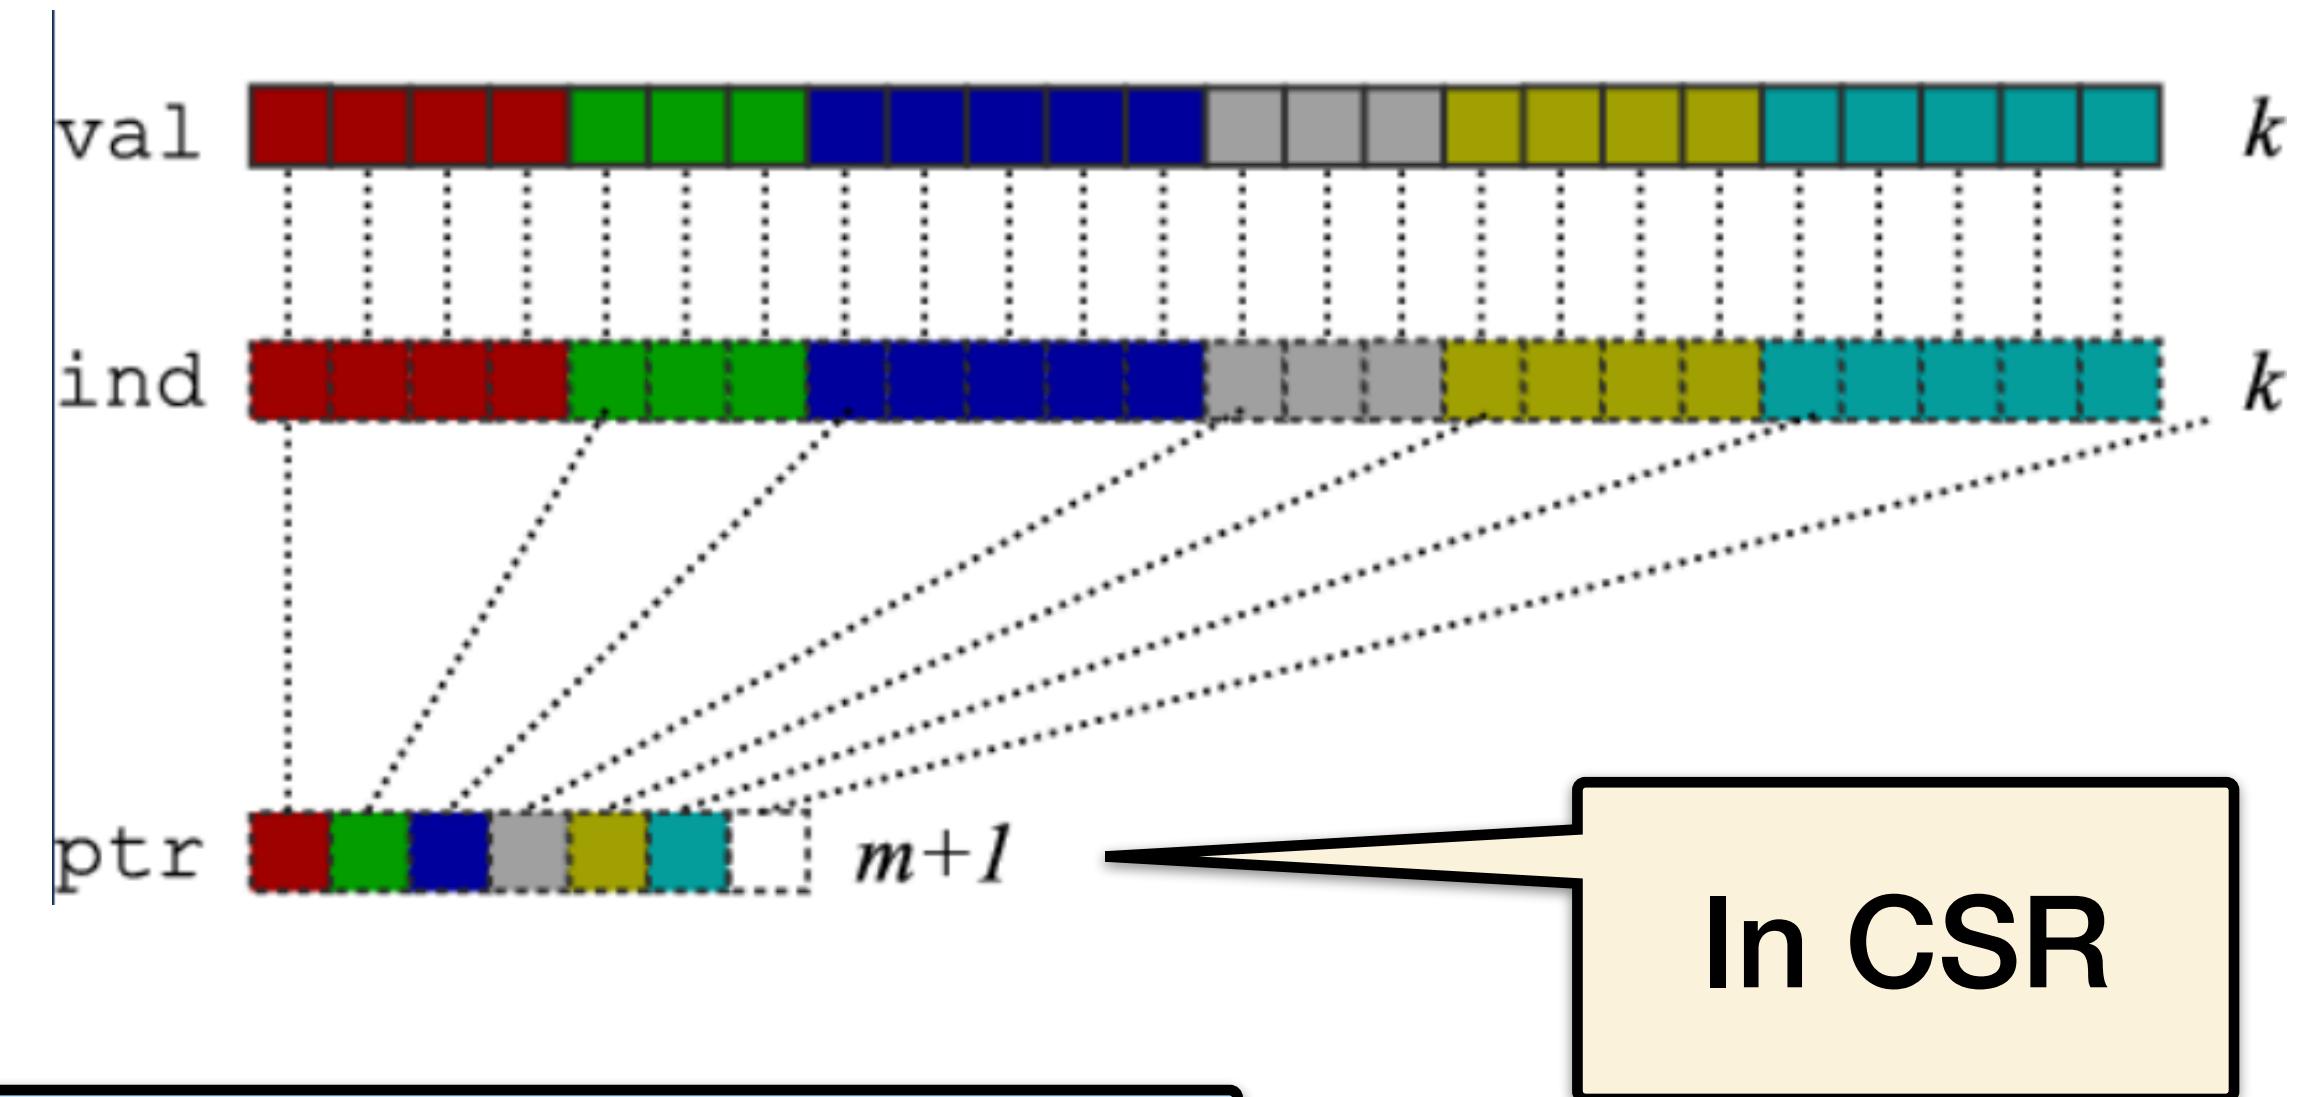
# **Serial SpMV**

# Serial SpMV in CSR

## SpMV: Sparse Matrix-Vector Multiplication



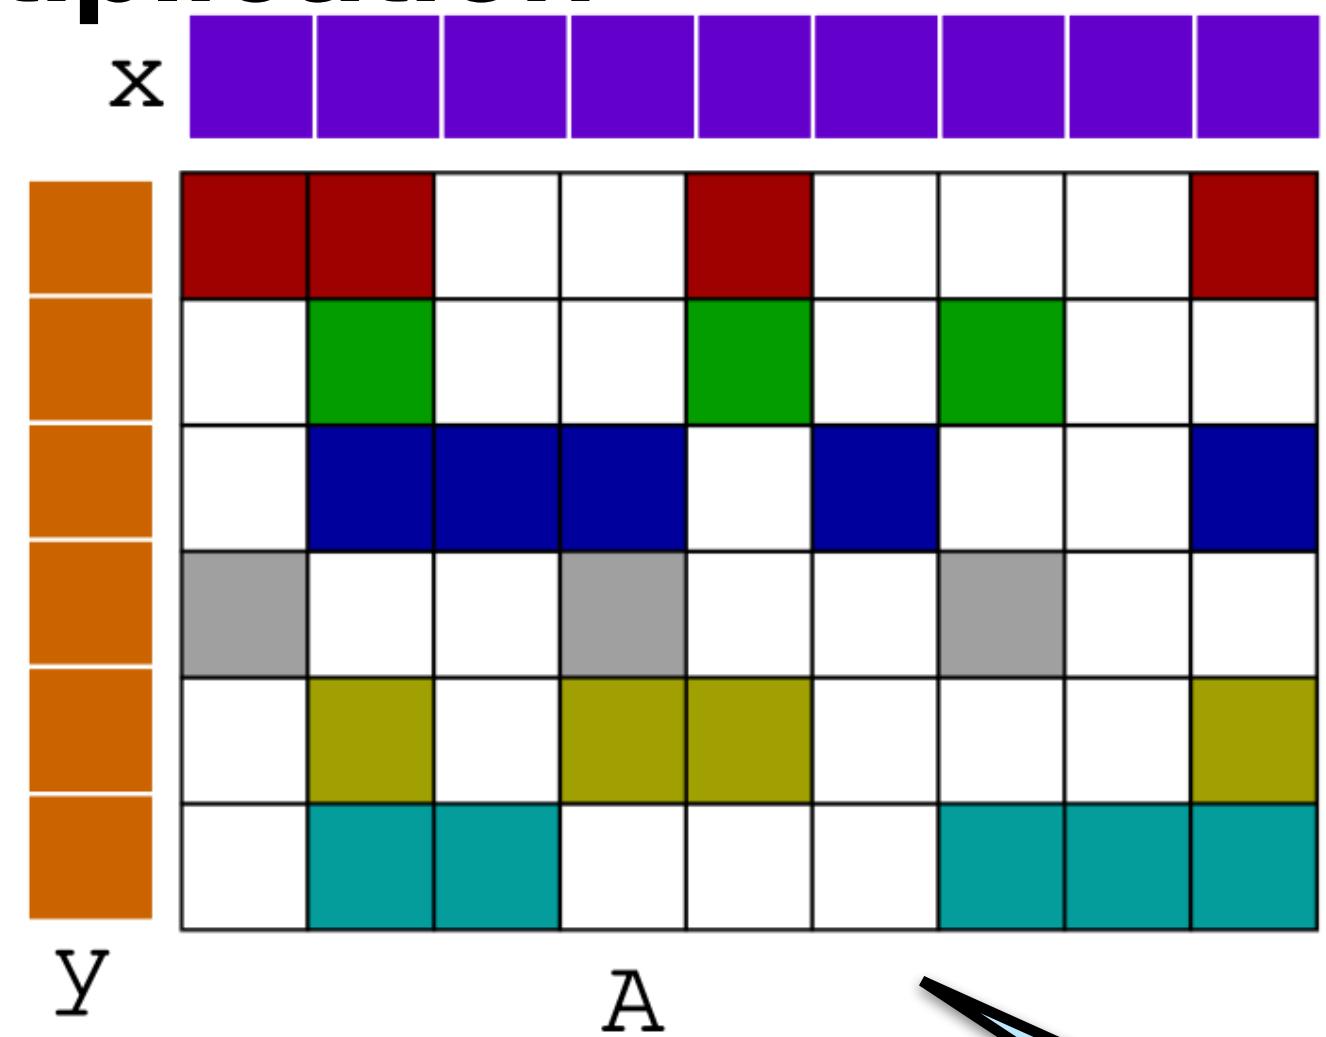
## Representation of A



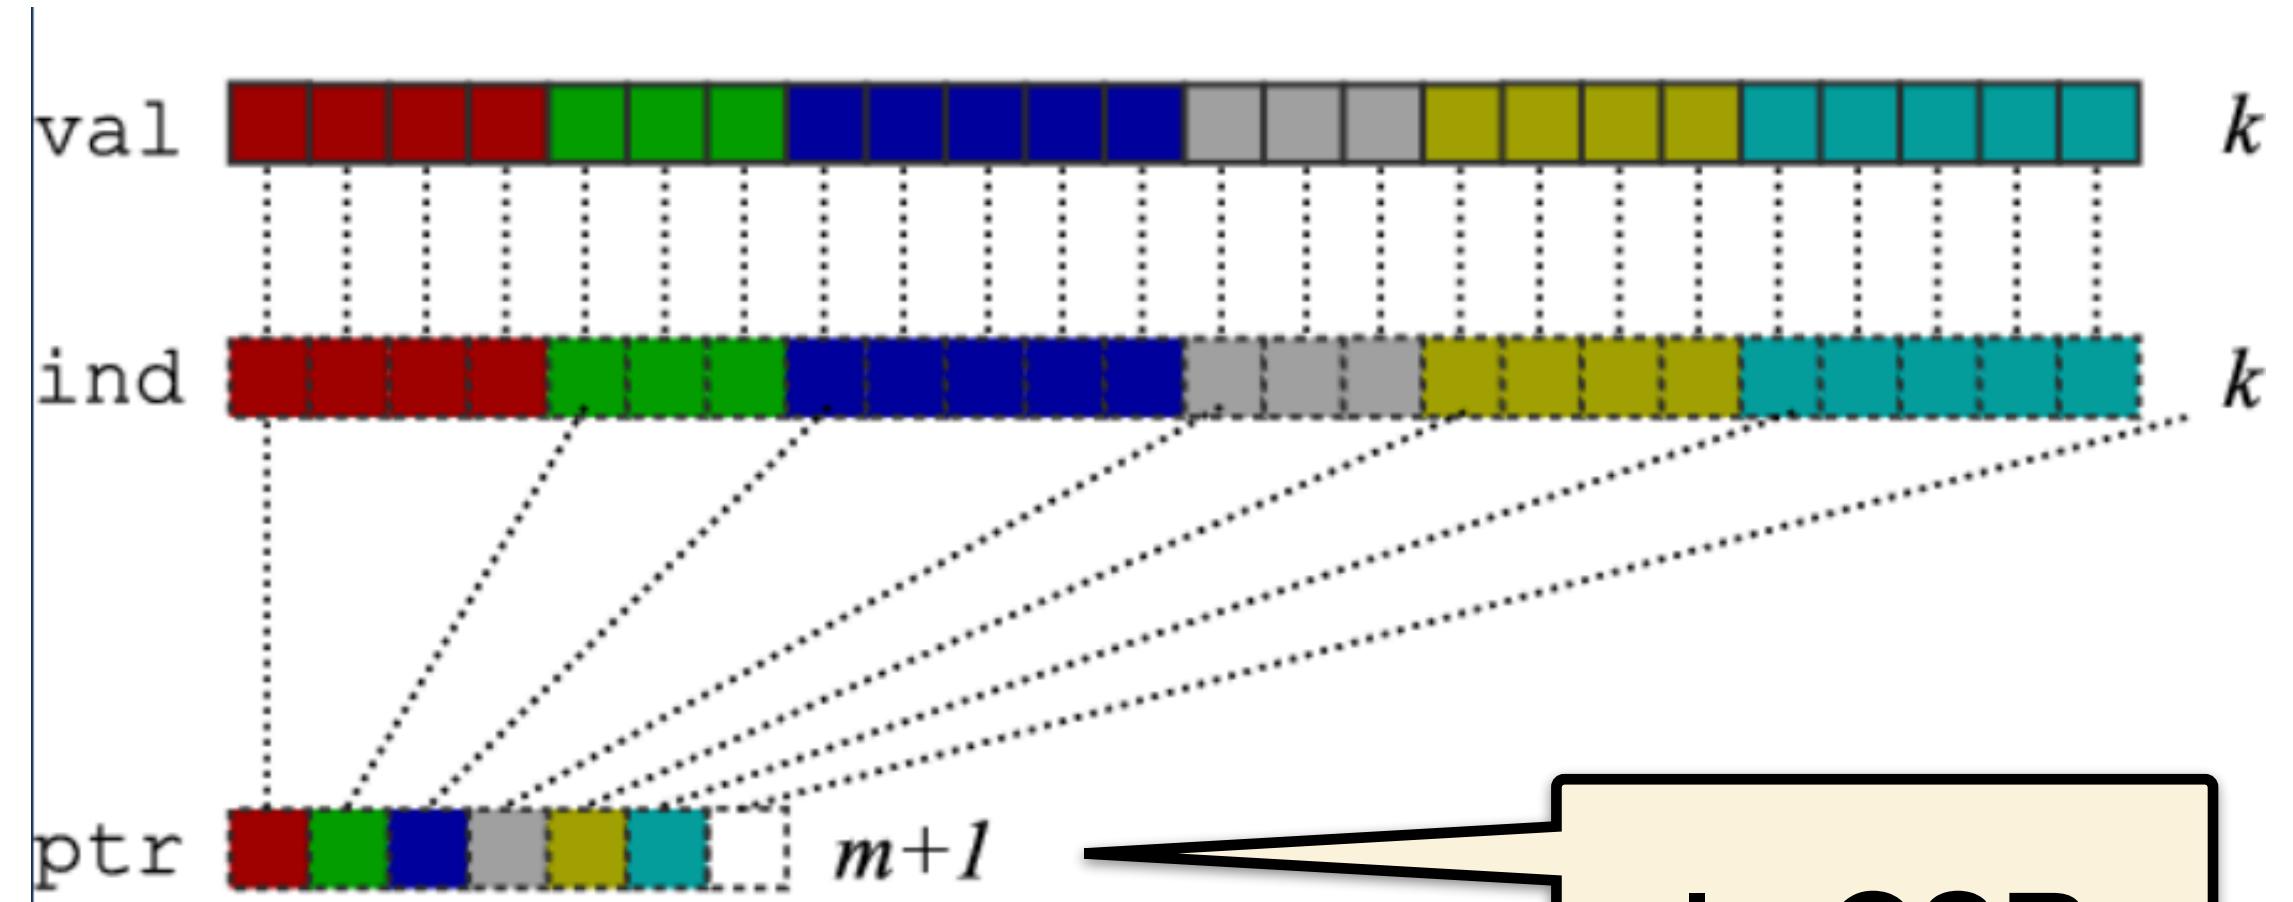
Compute  $y(i) = y(i) + A(i, j) * x(j)$

# Serial SpMV in CSR

## SpMV: Sparse Matrix-Vector Multiplication



## Representation of A



Compute  $y(i) = y(i) + A(i, j) * x(j)$

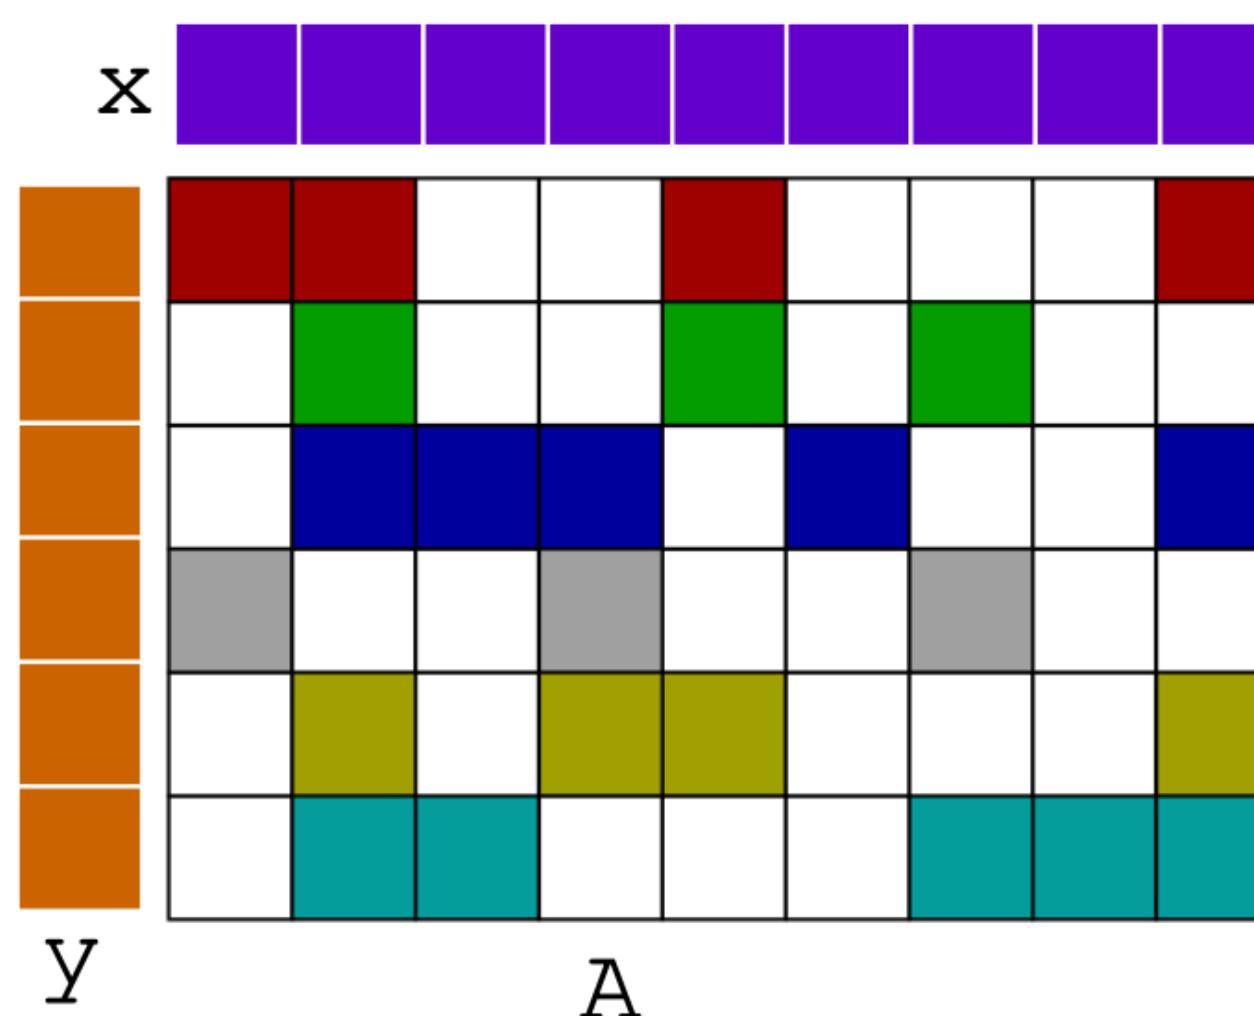
```
for each row i:  
    for k = ptr[i] to ptr[i+1] do  
        y[i] += val[k] * x[ind[k]]
```

- No reuse in A
- Maximum reuse in y as written
- Reuse in x?

# Parallel SpMV

# Parallel SpMV in CSR

```
parallel_for (int i = 0; i < m; i++) {  
    for (int j = ptr[i]; j < ptr[i+1]; j++) {  
        tmp = ind[j];  
        y[i] += val[j] * x[tmp]  
    }  
}
```



# Segmented Suffix Scan

0 1 2 3 4 5 6

x = 

1	2	1	2	1	2	1
---	---	---	---	---	---	---

ptr = 

0	2	5	7	9
---	---	---	---	---

A  
ind = 

1	2	0	3	4	1	2	5	6	0	1	4
---	---	---	---	---	---	---	---	---	---	---	---

val = 

1	1	1	2	2	2	1	2	2	1	3	1
---	---	---	---	---	---	---	---	---	---	---	---

flag = 

1	0	1	0	0	1	0	1	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---

from x

x(ind) = 

2	1	1	2	1	2	1	2	1	1	2	1
---	---	---	---	---	---	---	---	---	---	---	---

prod =

val \* x(ind)

prod = 

2	1	1	4	2	4	1	4	2	1	6	1
---	---	---	---	---	---	---	---	---	---	---	---

# Segmented Suffix Scan

0 1 2 3 4 5 6

x = 

1	2	1	2	1	2	1
---	---	---	---	---	---	---

ptr = 

0	2	5	7	9
---	---	---	---	---

A  
ind = 

1	2	0	3	4	1	2	5	6	0	1	4
---	---	---	---	---	---	---	---	---	---	---	---

val = 

1	1	1	2	2	2	1	2	2	1	3	1
---	---	---	---	---	---	---	---	---	---	---	---

flag = 

1	0	1	0	0	1	0	1	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---

prod = 

2	1	1	4	2	4	1	4	2	1	6	1
---	---	---	---	---	---	---	---	---	---	---	---

sums = 

2	3	1	5	7	4	5	4	6	1	7	8
---	---	---	---	---	---	---	---	---	---	---	---

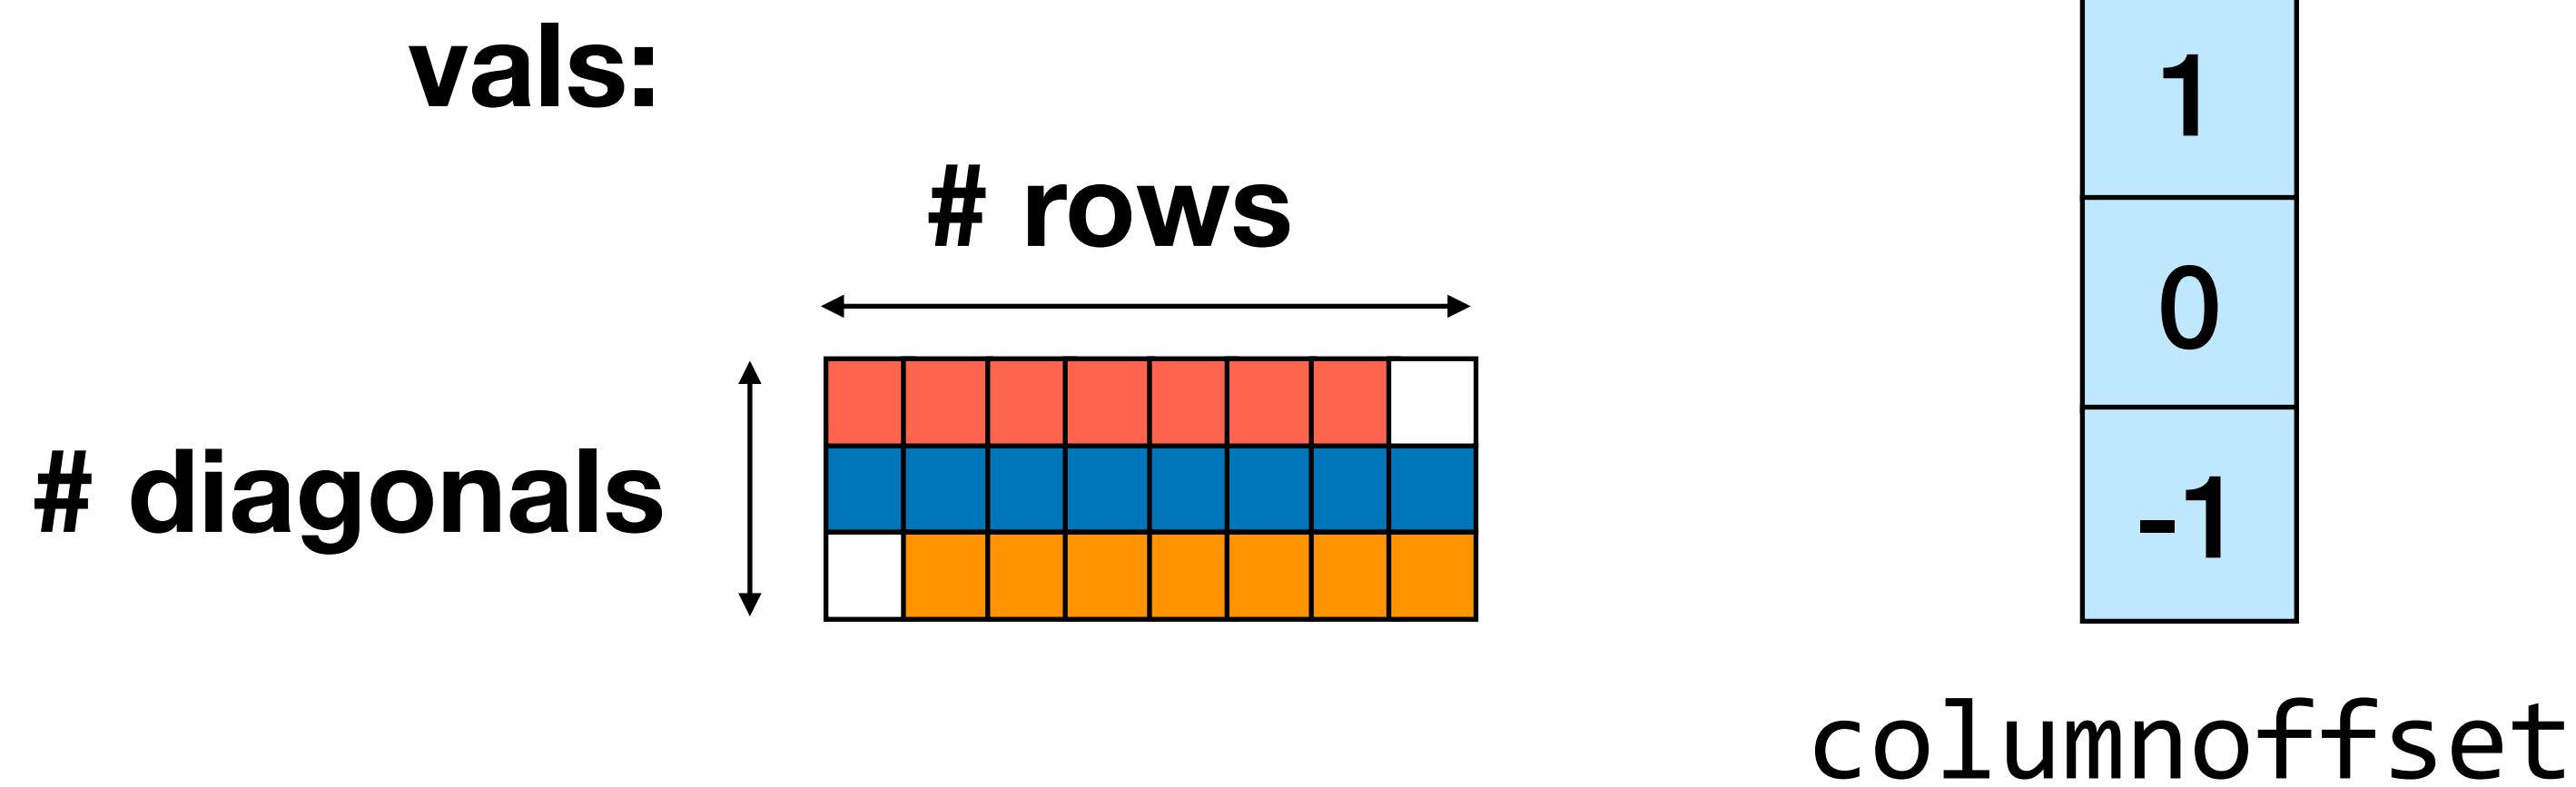
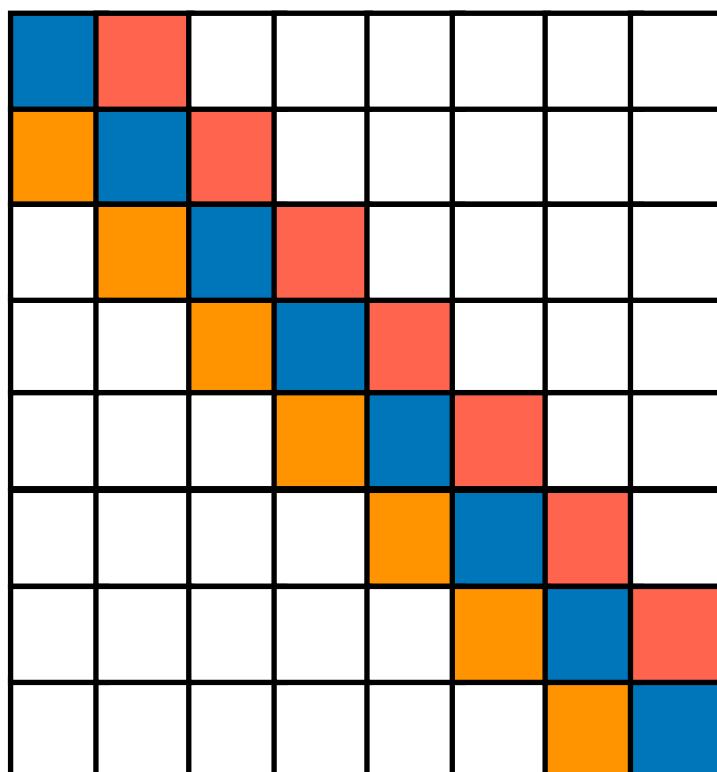
flag = zeros  
+ ones(ptr)

prod =  
val \* x(ind)

segmented scan  
on flag, prod

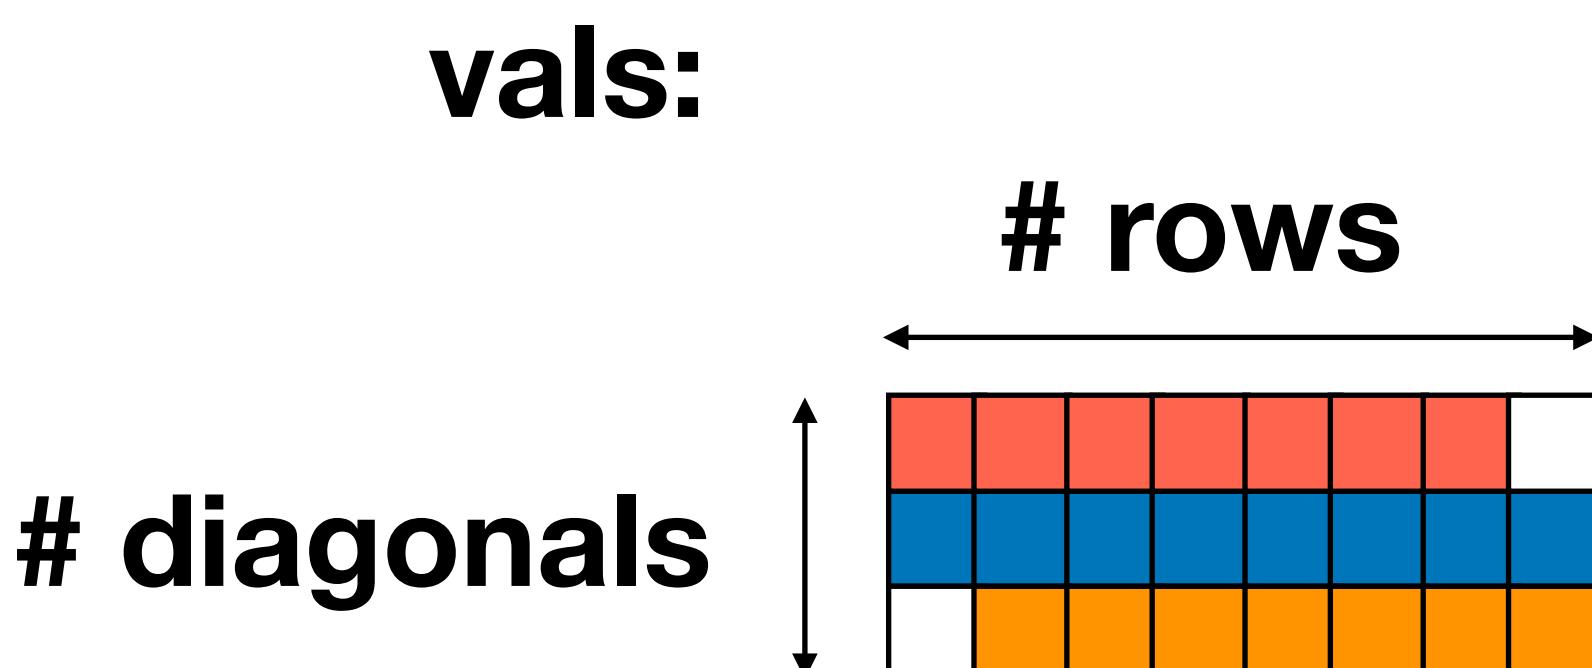
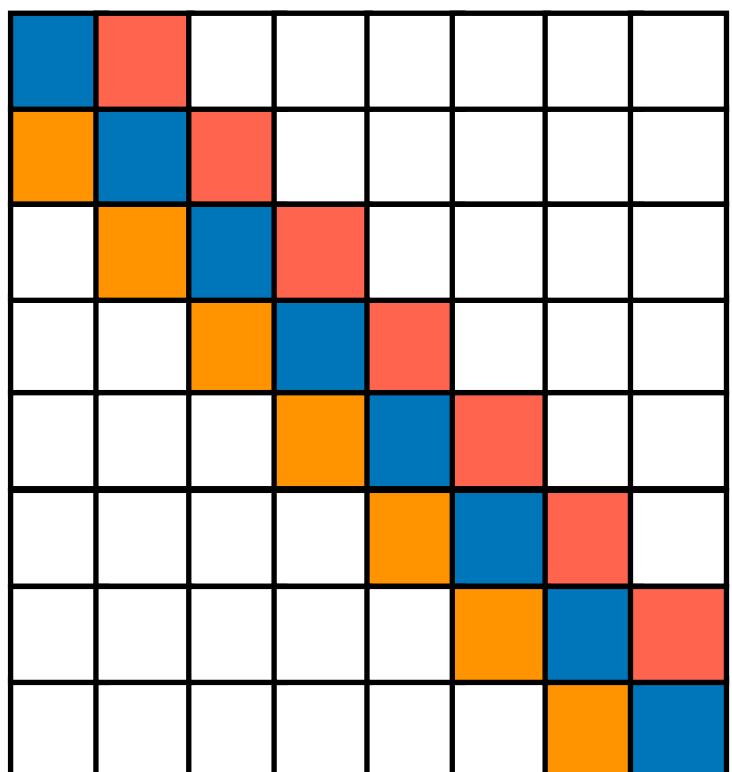
y = end of  
segments

# Parallel SpMV in Diagonal Format



```
for each diagonal k do
    parallel_for each row i do
        column = i + columnoffset[k]
        if (column >= 0 && column < n)
            y[i] = y[i] + val[k][column] * x[column]
```

# Parallel SpMV in Diagonal Format



columnoffset

```
for each diagonal k do
    parallel_for each row i do
        column = i + columnoffset[k]
        if (column >= 0 && column < n)
            y[i] = y[i] + val[k][column] * x[column]
```

**Diagonal is also popular for GPUs/vectors**

# Distributed SpMV

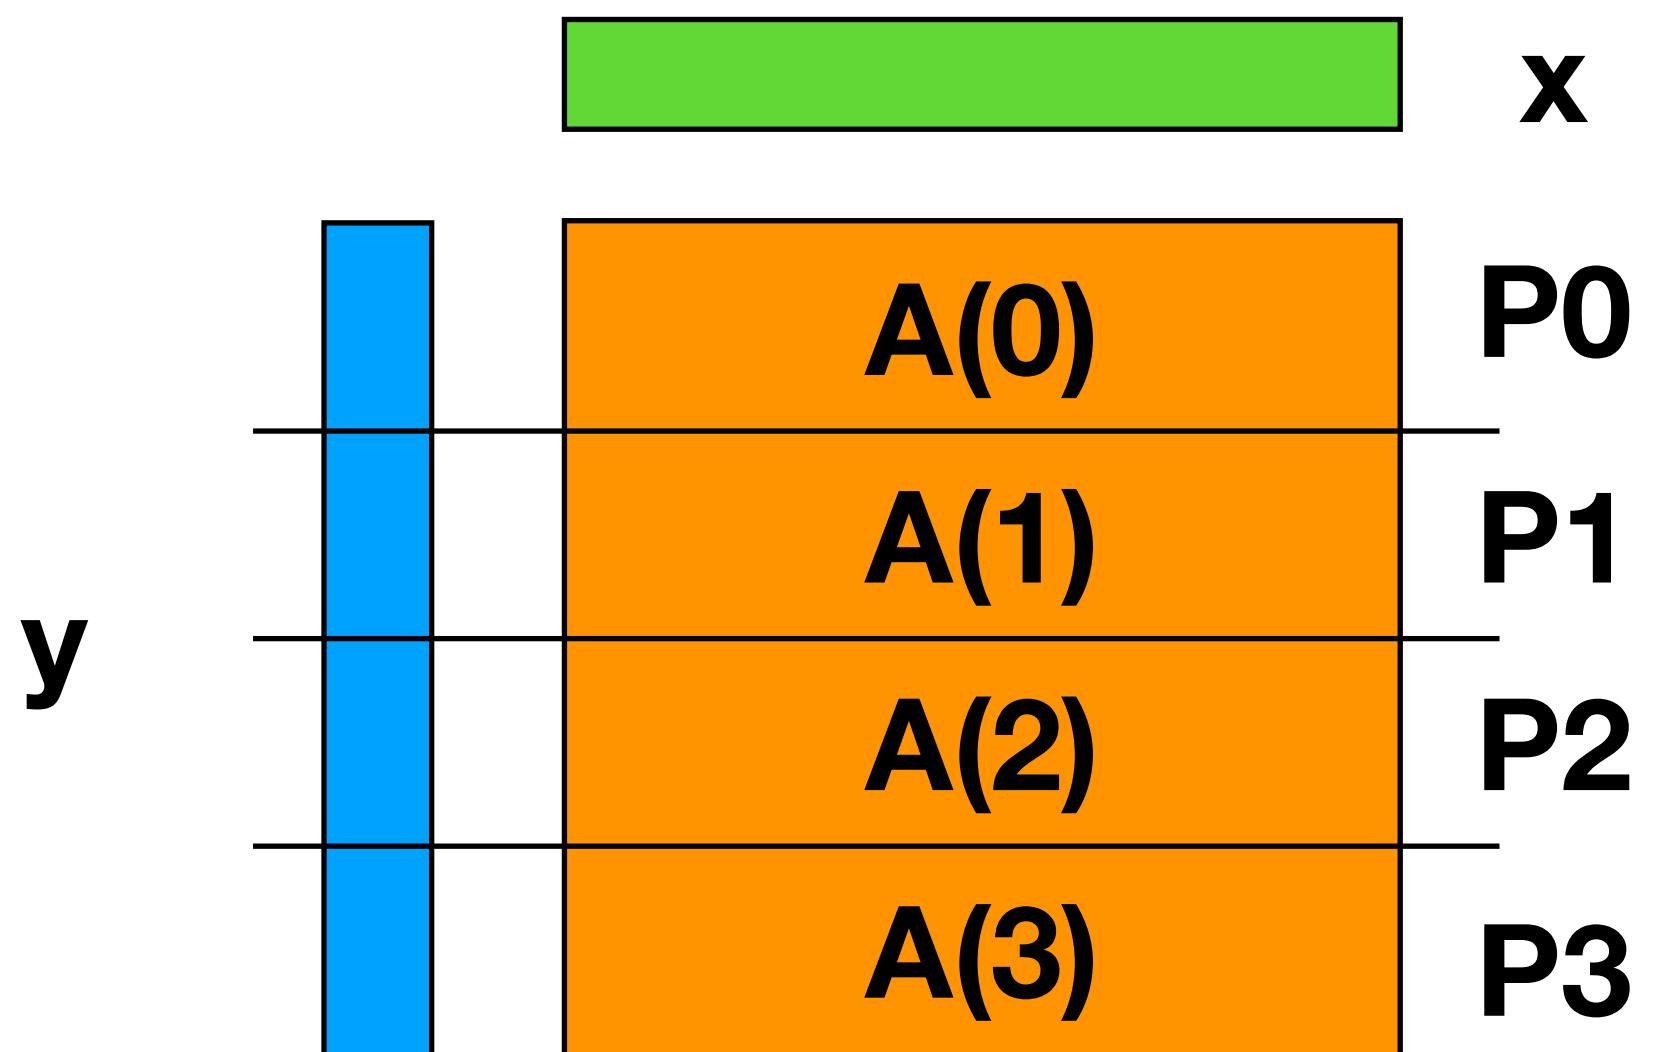
# Distributed Dense Matrix-Vector Product (Row Major)

warmup on  
dense case

- Compute  $y = y + A^*x$ , where  $A$  is a dense matrix
- Layout: **1D row blocked**
- Algorithm:

Allgather  $x$

```
for all local i
    for all j
        compute  $y[i] += A[i, j]*x[j]$  locally
```

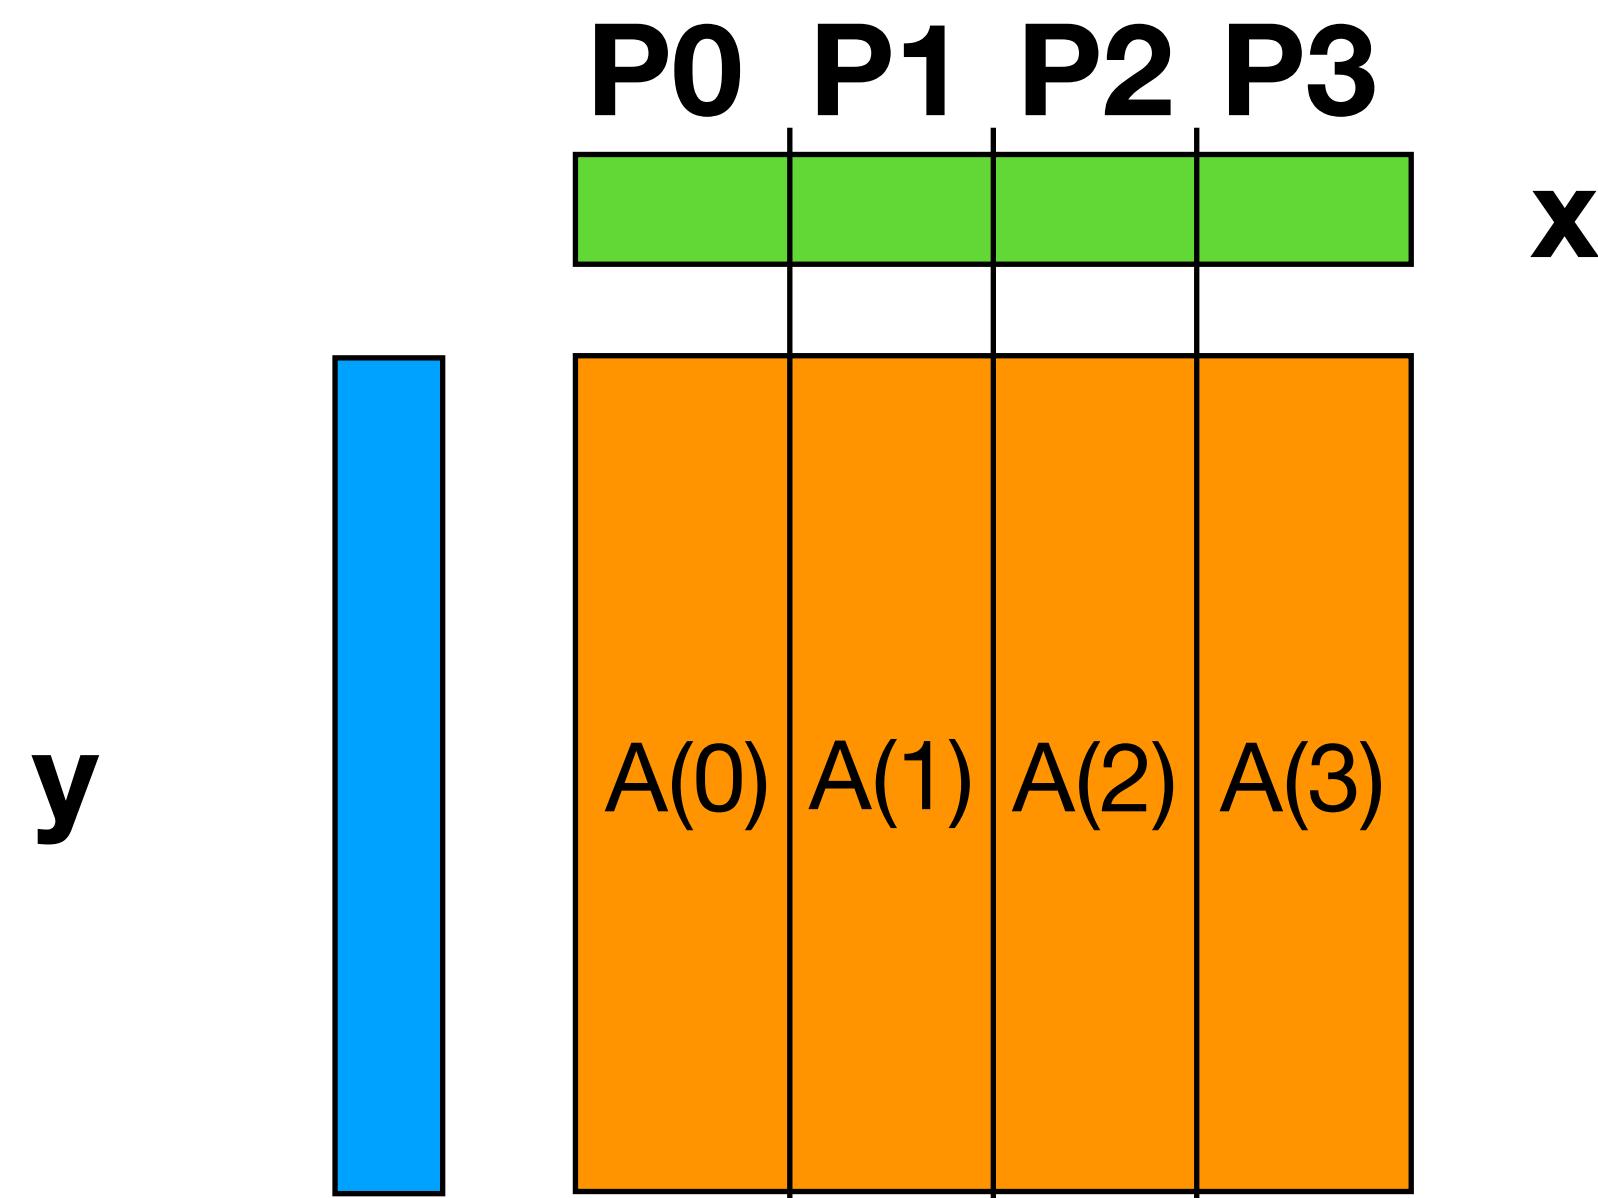


Broadcast + local  
dot product

# Distributed Dense Matrix-Vector Product (Column Major)

- Compute  $y = y + A^*x$ , where  $A$  is a dense matrix
- Layout: **1D column blocked**
- Algorithm:

```
foreach processor p
    make a temp vector of size n
    for all local j
        for all i
            compute temp[i] += A[i, j]*x[j] locally
    y[i] += SumReduce(temp[i])
```



Local daxpy and parallel reduction

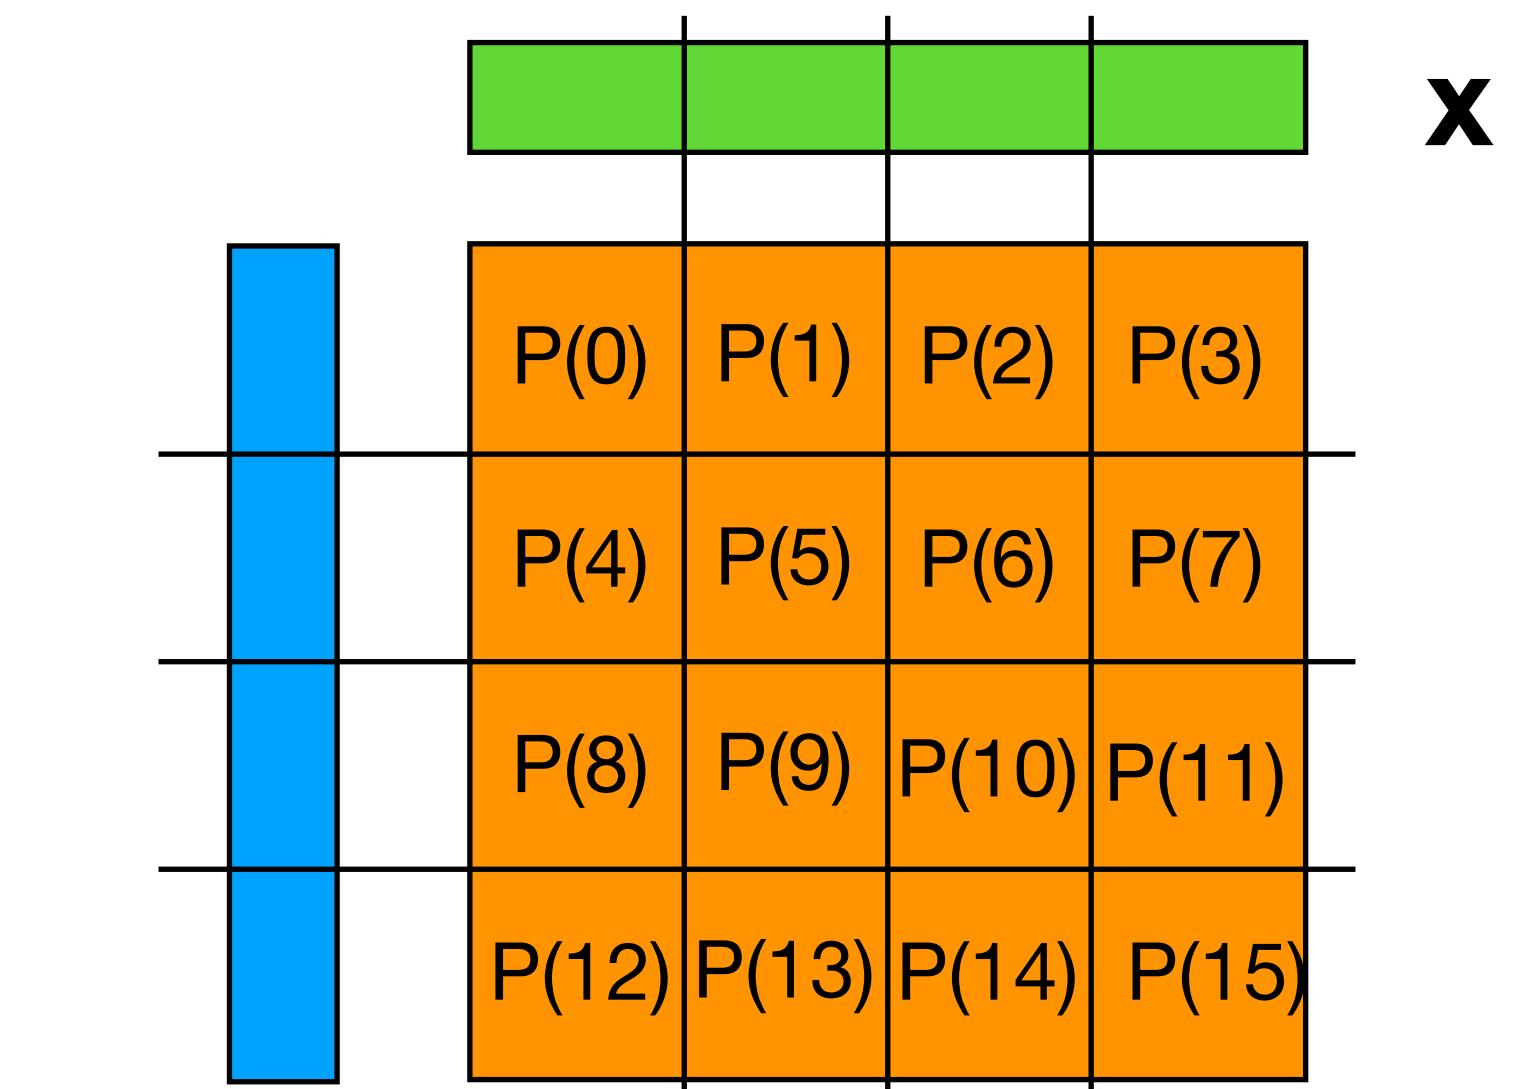
Reduce across all rows

# Distributed Dense Matrix-Vector Product (2D Blocked)

A 2D blocked layout uses a broadcast of  $x$  and reduction into  $y$ .

Both on a subset of processors

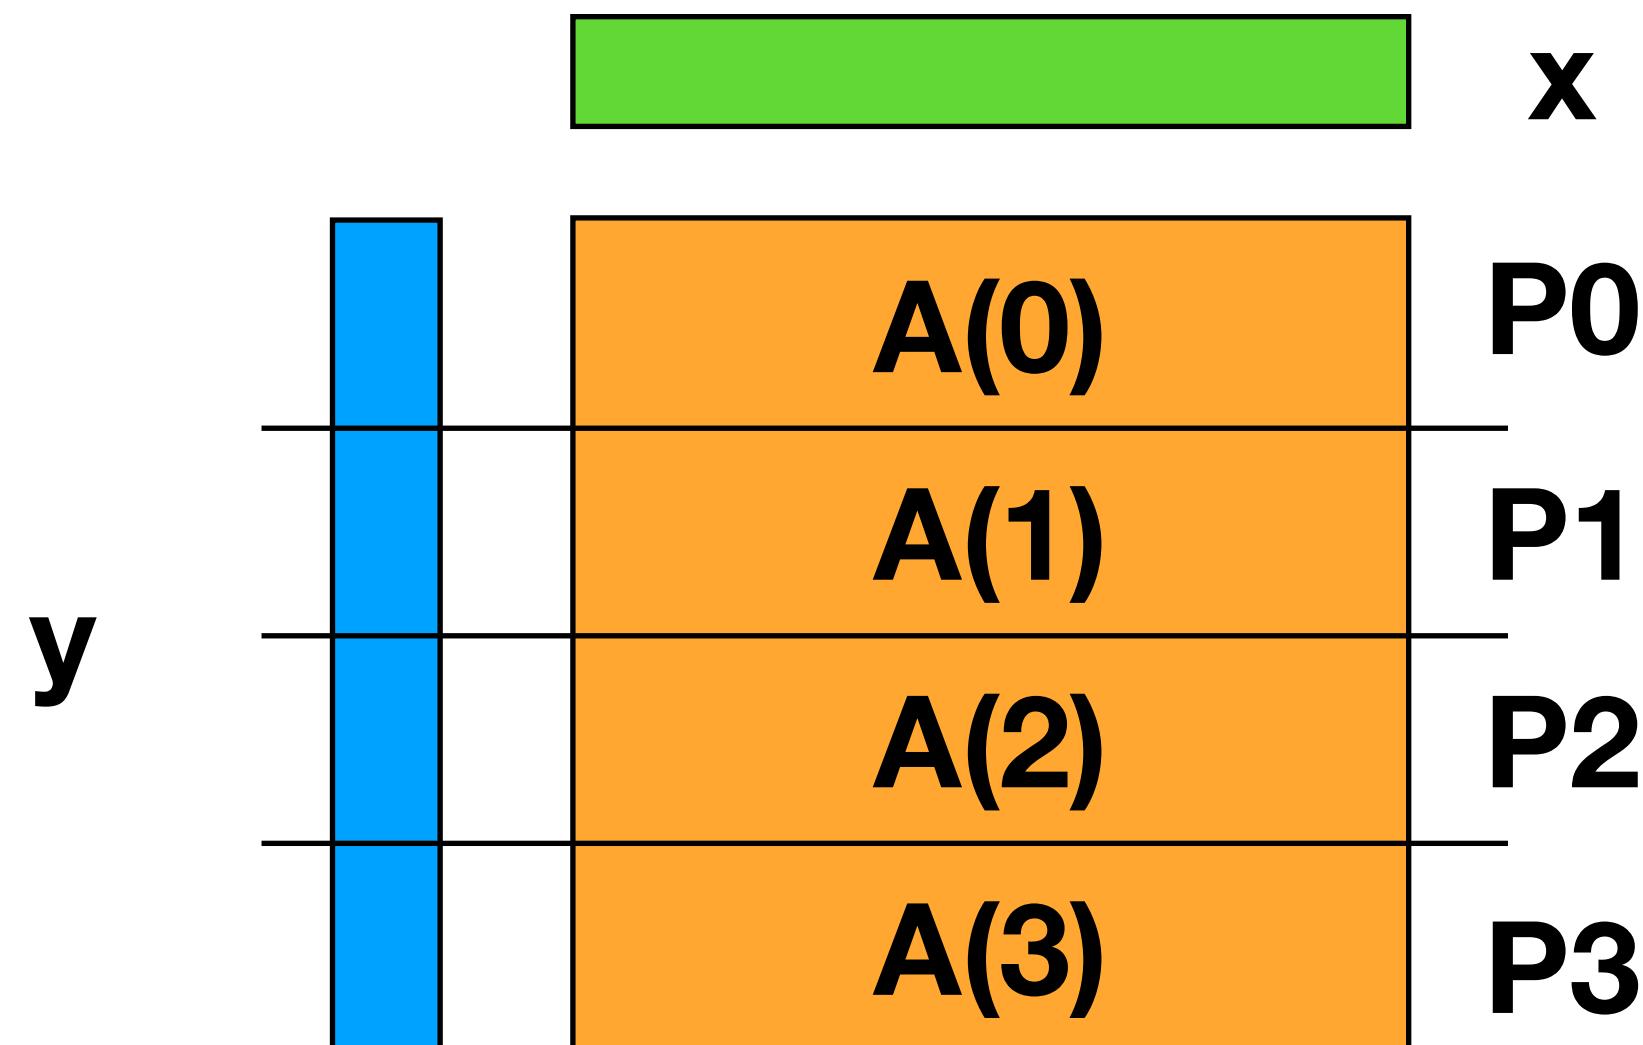
- $\sqrt{p}$  for square processor grid
- Can use other rectangular shapes  $p_{\text{row}} * p_{\text{col}} = p$



# Distributed SpMV

Row parallelism ( $y$  &  $A$  partitioned)

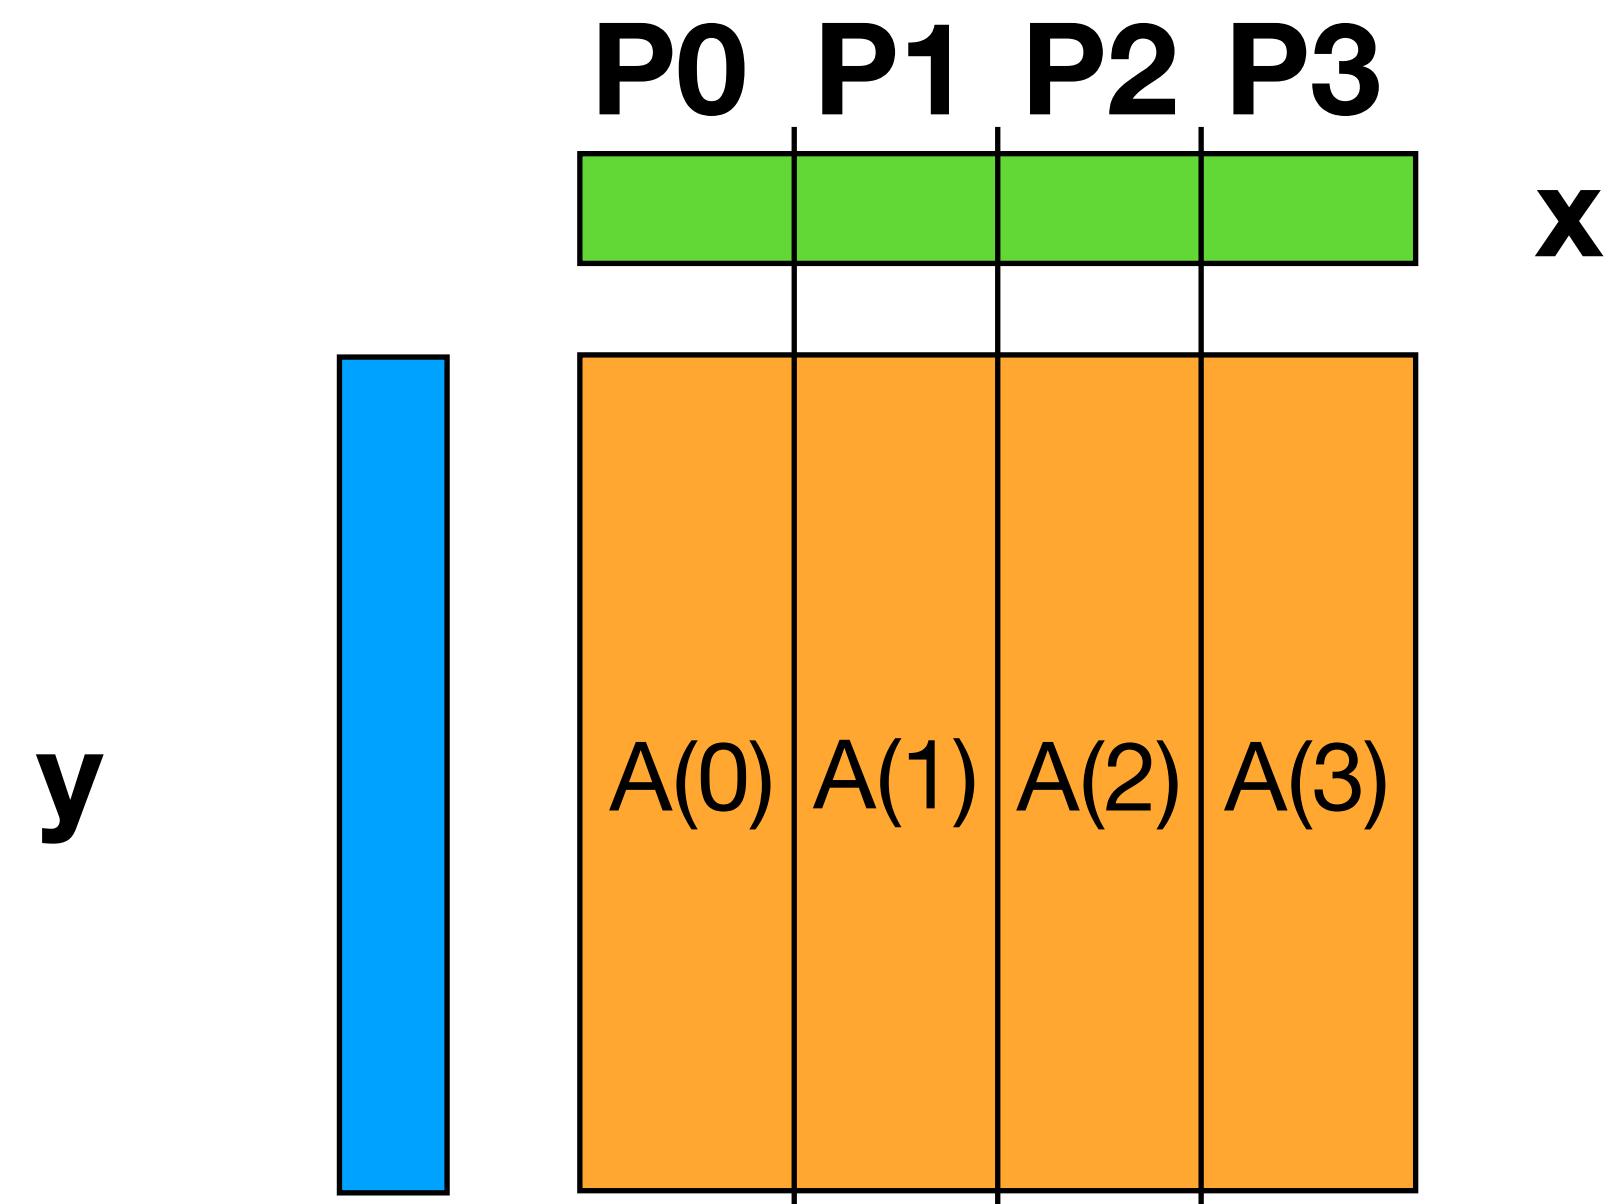
- Replicate  $x$  across processors
- Or exchange only necessary elements
- Are nonzeros clustered, e.g., near the diagonal?



# Distributed SpMV

Column parallelism ( $x$  &  $A$  partitioned)

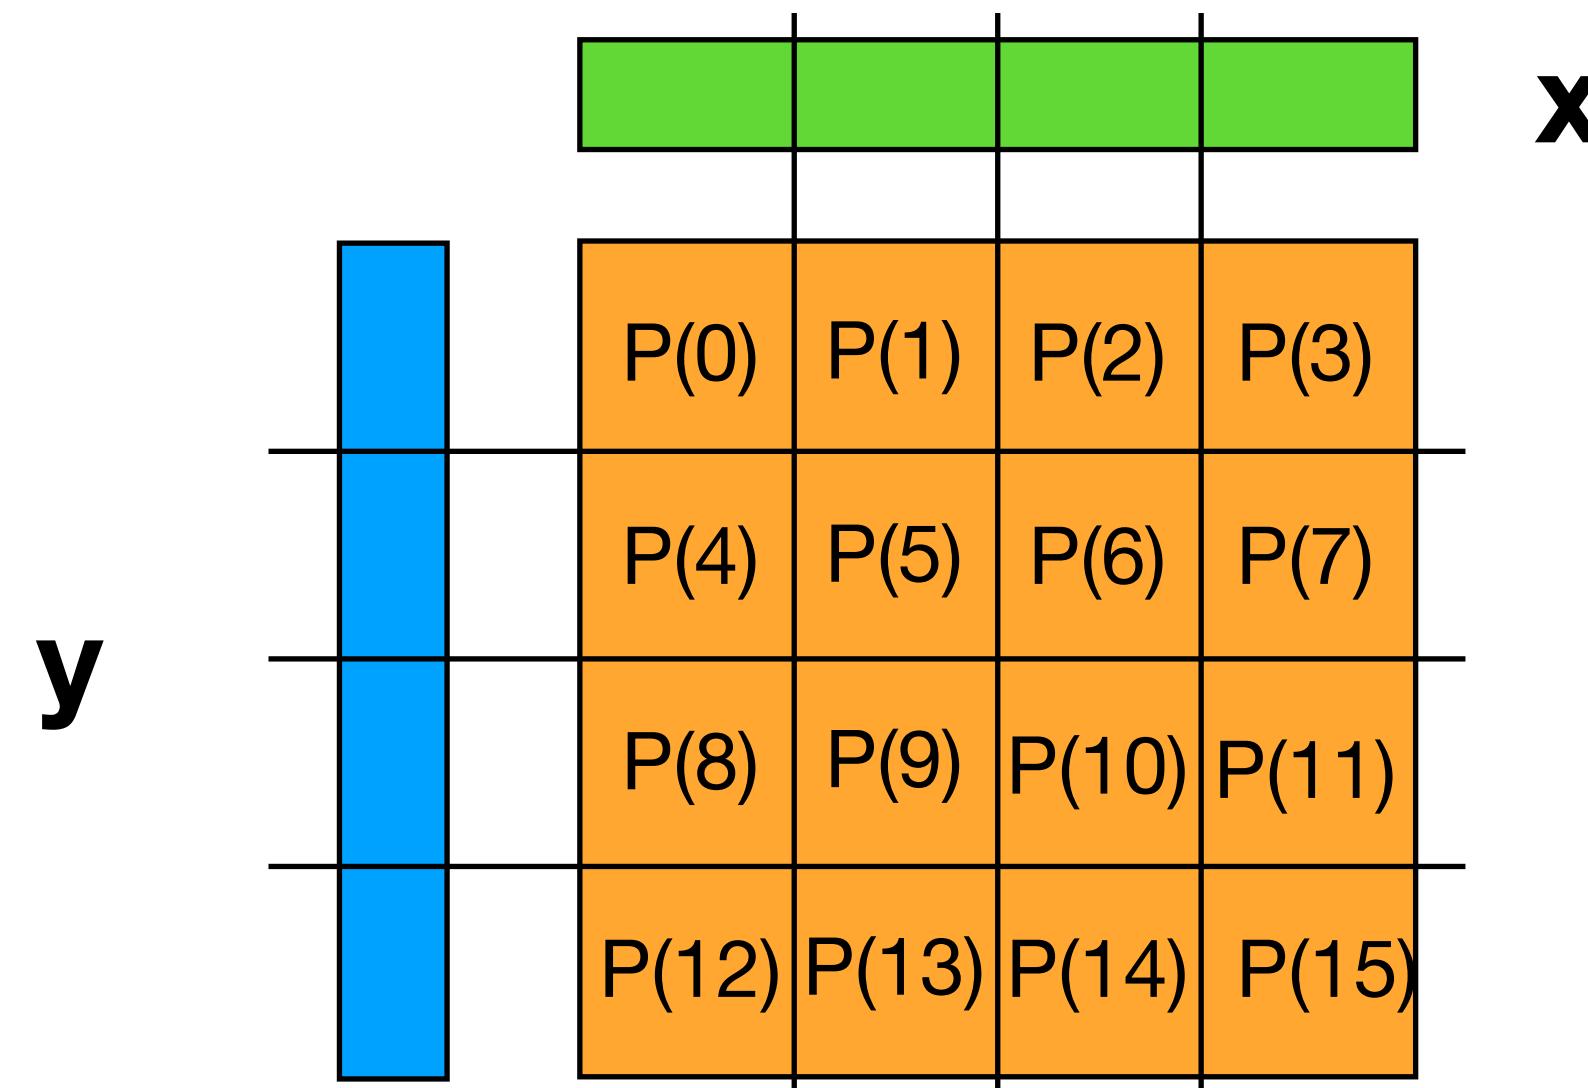
- Make temporary  $\text{temp\_y} = [0, \dots]$  on all processors;
- Update that; and (sparse?) sum-reduce over processors



# Distributed SpMV

2D parallelism for large  $p$  and **when nonzeros are uniform**

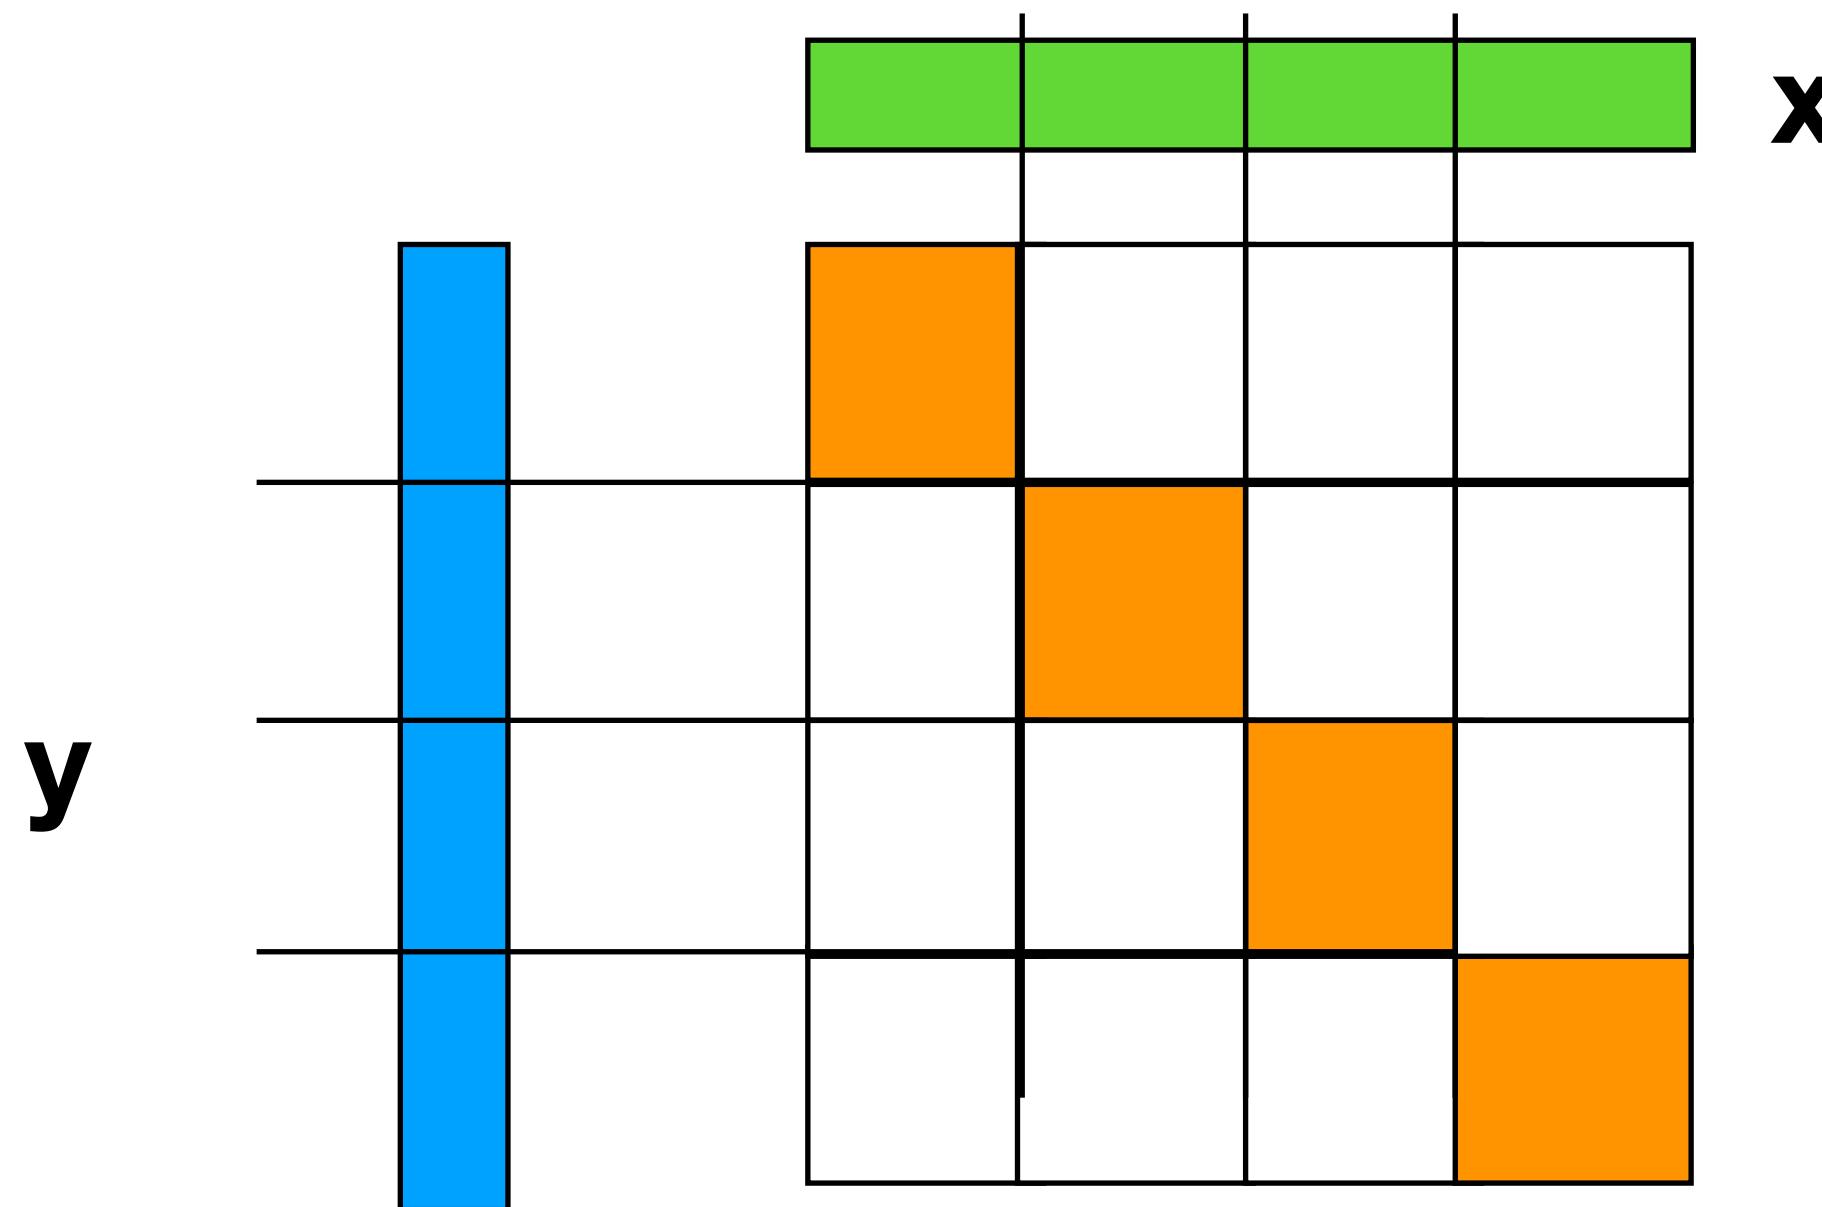
- Divide processors into  $p_1 \times p_2$  (e.g., square grid)
- Hybrid of Row and Column parallelism using teams
- NAS CG benchmark does this (random nonzero pattern)
- Bad load balance for clustered nonzeros



# Ideal Sparse Structure: P Diagonal Blocks

“Ideal” matrix structure for parallelism: block diagonal

- If  $p_i$  holds  $x_i$  and  $y_i$  blocks, no vectors to communicate
- If no nonzeroes outside these blocks, no communication is needed!



Dream scenario: reorder rows/columns to get close to this - most nonzeroes in diagonal blocks, few outside.

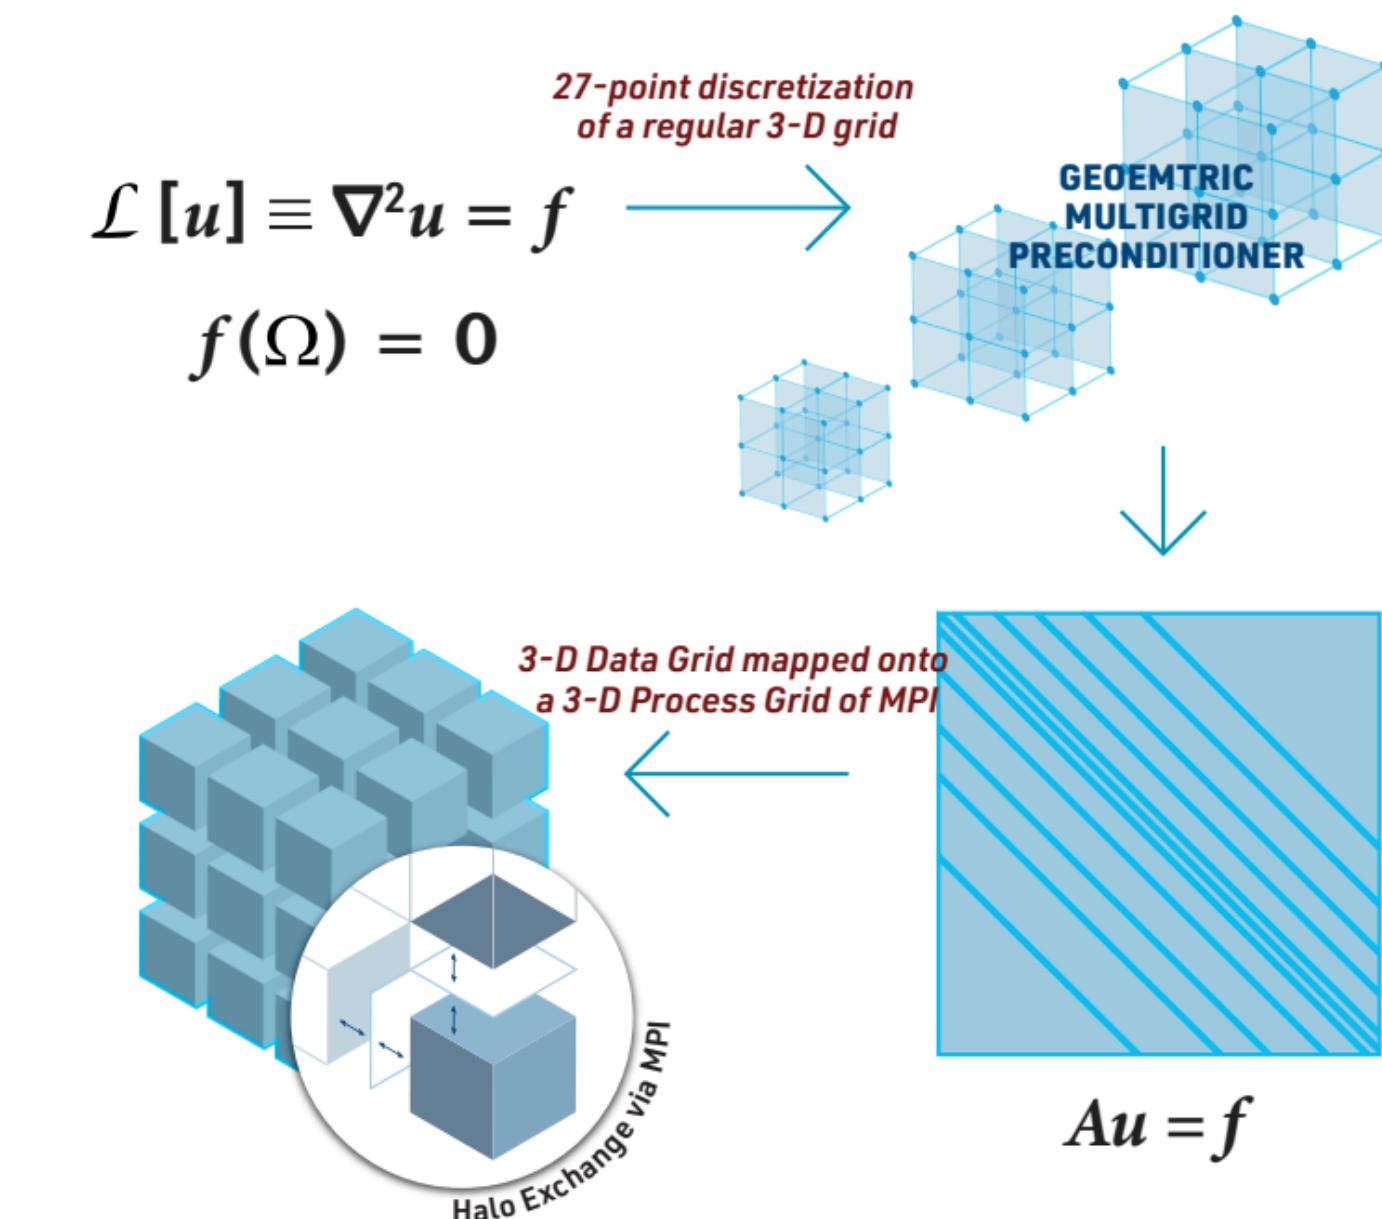
# High Performance Conjugate Gradients (HPCG) Benchmark

**Complement to LINPACK** (dense linear algebra), which is used for the TOP500.

Designed to exercise computational and data access patterns that more closely match a **different and broader set of applications**.

HPCG includes performance of the following basic operations:

- Conjugate gradient (27-point stencil)
- Sparse matrix-vector multiplication
- Global dot product
- Vector update
- and others



## November 2023 HPCG Results

New HPCG results were announced at SC23

Rank	Site	Computer	Cores	HPL Rmax (Pflop/s)	TOP500 Rank	HPCG (Pflop/s)	Fraction of Peak
1	RIKEN Center for Computational Science <b>Japan</b>	<b>Supercomputer Fugaku</b> — A64FX 48C 2.2GHz, Tofu interconnect D	7,630,848	442.01	4	16.00	3.0%
2	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Frontier</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	8,699,904	1194.00	1	14.05	0.8%
3	EuroHPC/CSC <b>Finland</b>	<b>LUMI</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	2,752,704	379.70	5	4.587	0.9%
4	EuroHPC/CINECA <b>Italy</b>	<b>Leonardo</b> — Xeon Platinum 8358 32C 2.6GHz, Quad-rail NVIDIA HDR100 Infiniband, NVIDIA A100 SXM4 64 GB	1,824,768	238.70	6	3.114	1.0%
5	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Summit</b> — IBM POWER9 22C 3.07GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	2,414,592	148.60	7	2.926	1.5%
6	DOE/SC/LBNL/NERSC <b>United States</b>	<b>Perlmutter</b> — AMD EPYC 7763 64C 2.45GHz, Slingshot-11, NVIDIA A100 SXM4 40 GB	888,832	79.23	12	1.905	1.7%
7	DOE/NNSA/LLNL <b>United States</b>	<b>Sierra</b> — IBM POWER9 22C 3.1GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	1,572,480	94.64	10	1.796	1.4%
8	NVIDIA Corporation <b>United States</b>	<b>Selene</b> — AMD EPYC 7742 64C 2.25GHz, Mellanox HDR Infiniband, NVIDIA A100	555,520	63.46	13	1.623	2.0%
9	Forschungszentrum Juelich (FZJ) <b>Germany</b>	<b>JUWELS Booster Module</b> — AMD EPYC 7402 24C 2.8GHz, Mellanox HDR InfiniBand/ParTec ParaStation ClusterSuite, NVIDIA A100	449,280	44.12	18	1.275	1.8%
10	Cyberscience Center, Tohoku University <b>Japan</b>	<b>AOBA-S</b> — Vector Engine Type 30A 16C 1.6GHz, Infiniband NDR200	64,512	17.22	50	1.089	5.5%

Not the same order as TOP500

## November 2023 HPCG Results

New HPCG results were announced at SC23

<https://www.netlib.org/benchmark/hpl/>

Rank	Site	Computer	Cores	HPL Rmax (Pflop/s)	TOP500 Rank	HPCG (Pflop/s)	Fraction of Peak
1	RIKEN Center for Computational Science <b>Japan</b>	<b>Supercomputer Fugaku</b> — A64FX 48C 2.2GHz, Tofu interconnect D	7,630,848	442.01	4	16.00	3.0%
2	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Frontier</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	8,699,904	1194.00	1	14.05	0.8%
3	EuroHPC/CSC <b>Finland</b>	<b>LUMI</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	2,752,704	379.70	5	4.587	0.9%
4	EuroHPC/CINECA <b>Italy</b>	<b>Leonardo</b> — Xeon Platinum 8358 32C 2.6GHz, Quad-rail NVIDIA HDR100 Infiniband, NVIDIA A100 SXM4 64 GB	1,824,768	238.70	6	3.114	1.0%
5	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Summit</b> — IBM POWER9 22C 3.07GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	2,414,592	148.60	7	2.926	1.5%
6	DOE/SC/LBNL/NERSC <b>United States</b>	<b>Perlmutter</b> — AMD EPYC 7763 64C 2.45GHz, Slingshot-11, NVIDIA A100 SXM4 40 GB	888,832	79.23	12	1.905	1.7%
7	DOE/NNSA/LLNL <b>United States</b>	<b>Sierra</b> — IBM POWER9 22C 3.1GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	1,572,480	94.64	10	1.796	1.4%
8	NVIDIA Corporation <b>United States</b>	<b>Selene</b> — AMD EPYC 7742 64C 2.25GHz, Mellanox HDR Infiniband, NVIDIA A100	555,520	63.46	13	1.623	2.0%
9	Forschungszentrum Juelich (FZJ) <b>Germany</b>	<b>JUWELS Booster Module</b> — AMD EPYC 7402 24C 2.8GHz, Mellanox HDR InfiniBand/ParTec ParaStation ClusterSuite, NVIDIA A100	449,280	44.12	18	1.275	1.8%
10	Cyberscience Center, Tohoku University <b>Japan</b>	<b>AOBA-S</b> — Vector Engine Type 30A 16C 1.6GHz, Infiniband NDR200	64,512	17.22	50	1.089	5.5%

Much lower Pflop/s compared to Linpack

## November 2023 HPCG Results

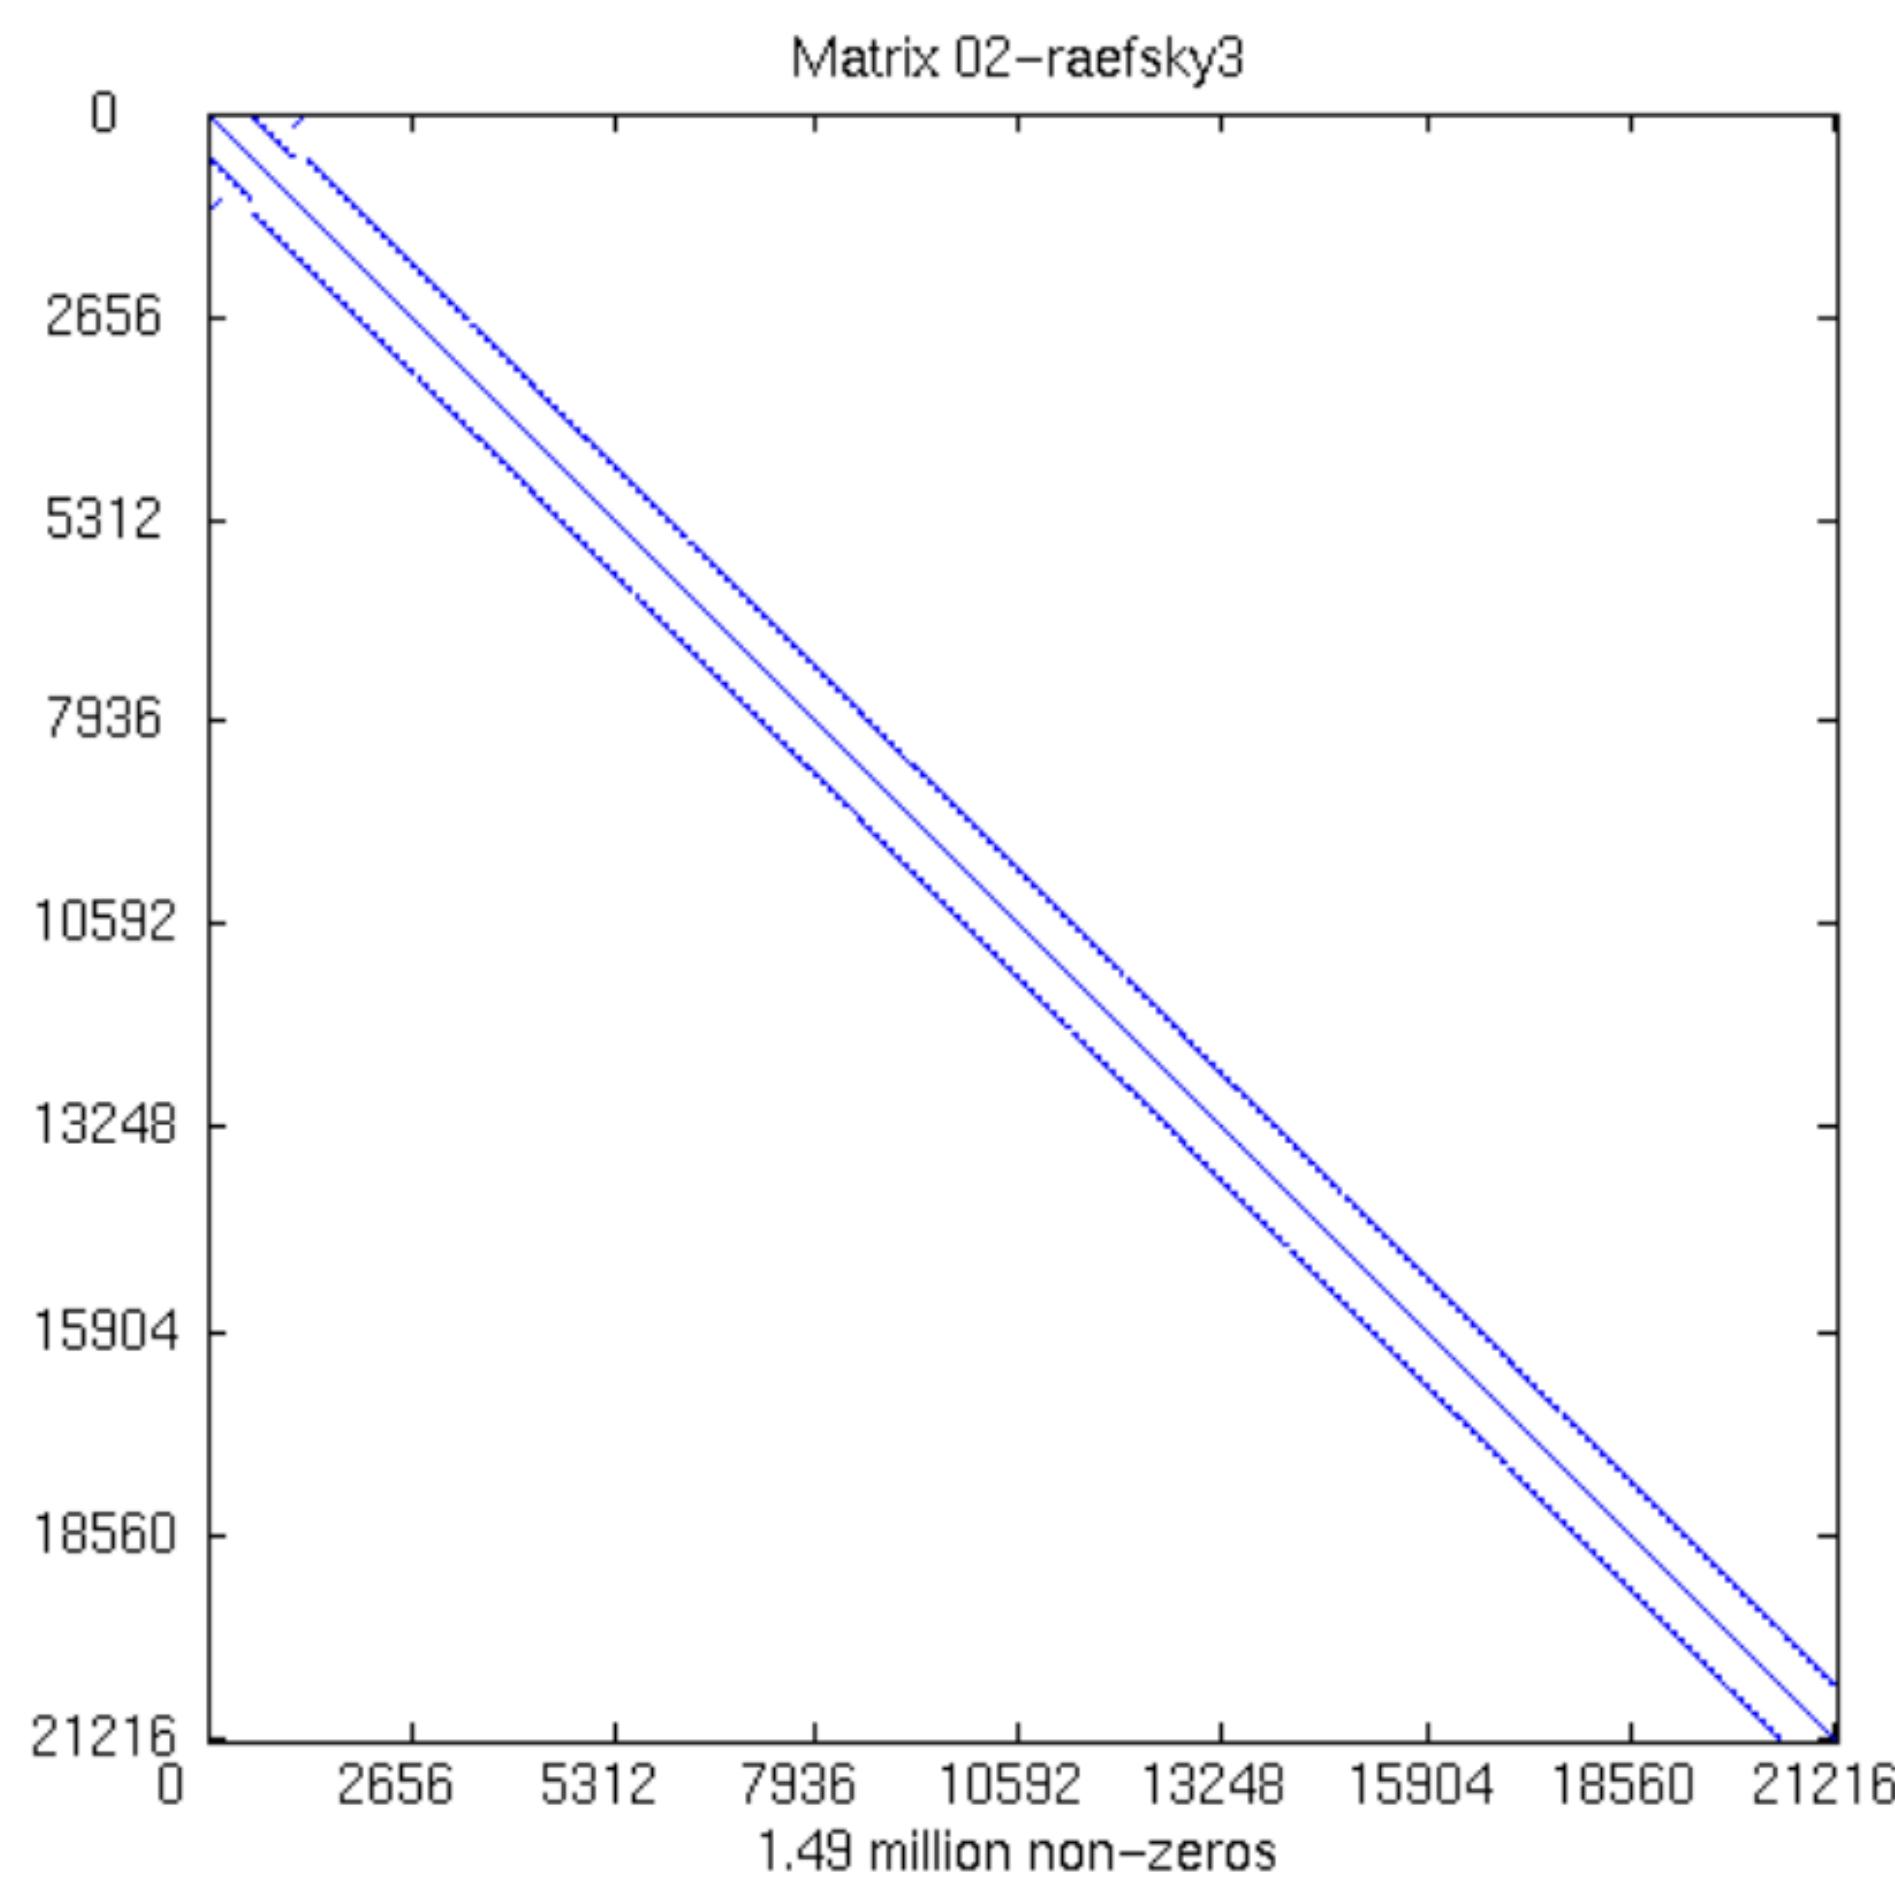
New HPCG results were announced at SC23

Rank	Site	Computer	Cores	HPL Rmax (Pflop/s)	TOP500 Rank	HPCG (Pflop/s)	Fraction of Peak
1	RIKEN Center for Computational Science <b>Japan</b>	<b>Supercomputer Fugaku</b> — A64FX 48C 2.2GHz, Tofu interconnect D	7,630,848	442.01	4	16.00	3.0%
2	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Frontier</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	8,699,904	1194.00	1	14.05	0.8%
3	EuroHPC/CSC <b>Finland</b>	<b>LUMI</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	2,752,704	379.70	5	4.587	0.9%
4	EuroHPC/CINECA <b>Italy</b>	<b>Leonardo</b> — Xeon Platinum 8358 32C 2.6GHz, Quad-rail NVIDIA HDR100 Infiniband, NVIDIA A100 SXM4 64 GB	1,824,768	238.70	6	3.114	1.0%
5	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Summit</b> — IBM POWER9 22C 3.07GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	2,414,592	148.60	7	2.926	1.5%
6	DOE/SC/LBNL/NERSC <b>United States</b>	<b>Perlmutter</b> — AMD EPYC 7763 64C 2.45GHz, Slingshot-11, NVIDIA A100 SXM4 40 GB	888,832	79.23	12	1.905	1.7%
7	DOE/NNSA/LLNL <b>United States</b>	<b>Sierra</b> — IBM POWER9 22C 3.1GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	1,572,480	94.64	10	1.796	1.4%
8	NVIDIA Corporation <b>United States</b>	<b>Selene</b> — AMD EPYC 7742 64C 2.25GHz, Mellanox HDR Infiniband, NVIDIA A100	555,520	63.46	13	1.623	2.0%
9	Forschungszentrum Juelich (FZJ) <b>Germany</b>	<b>JUWELS Booster Module</b> — AMD EPYC 7402 24C 2.8GHz, Mellanox HDR InfiniBand/ParTec ParaStation ClusterSuite, NVIDIA A100	449,280	44.12	18	1.275	1.8%
10	Cyberscience Center, Tohoku University <b>Japan</b>	<b>AOBA-S</b> — Vector Engine Type 30A 16C 1.6GHz, Infiniband NDR200	64,512	17.22	50	1.089	5.5%

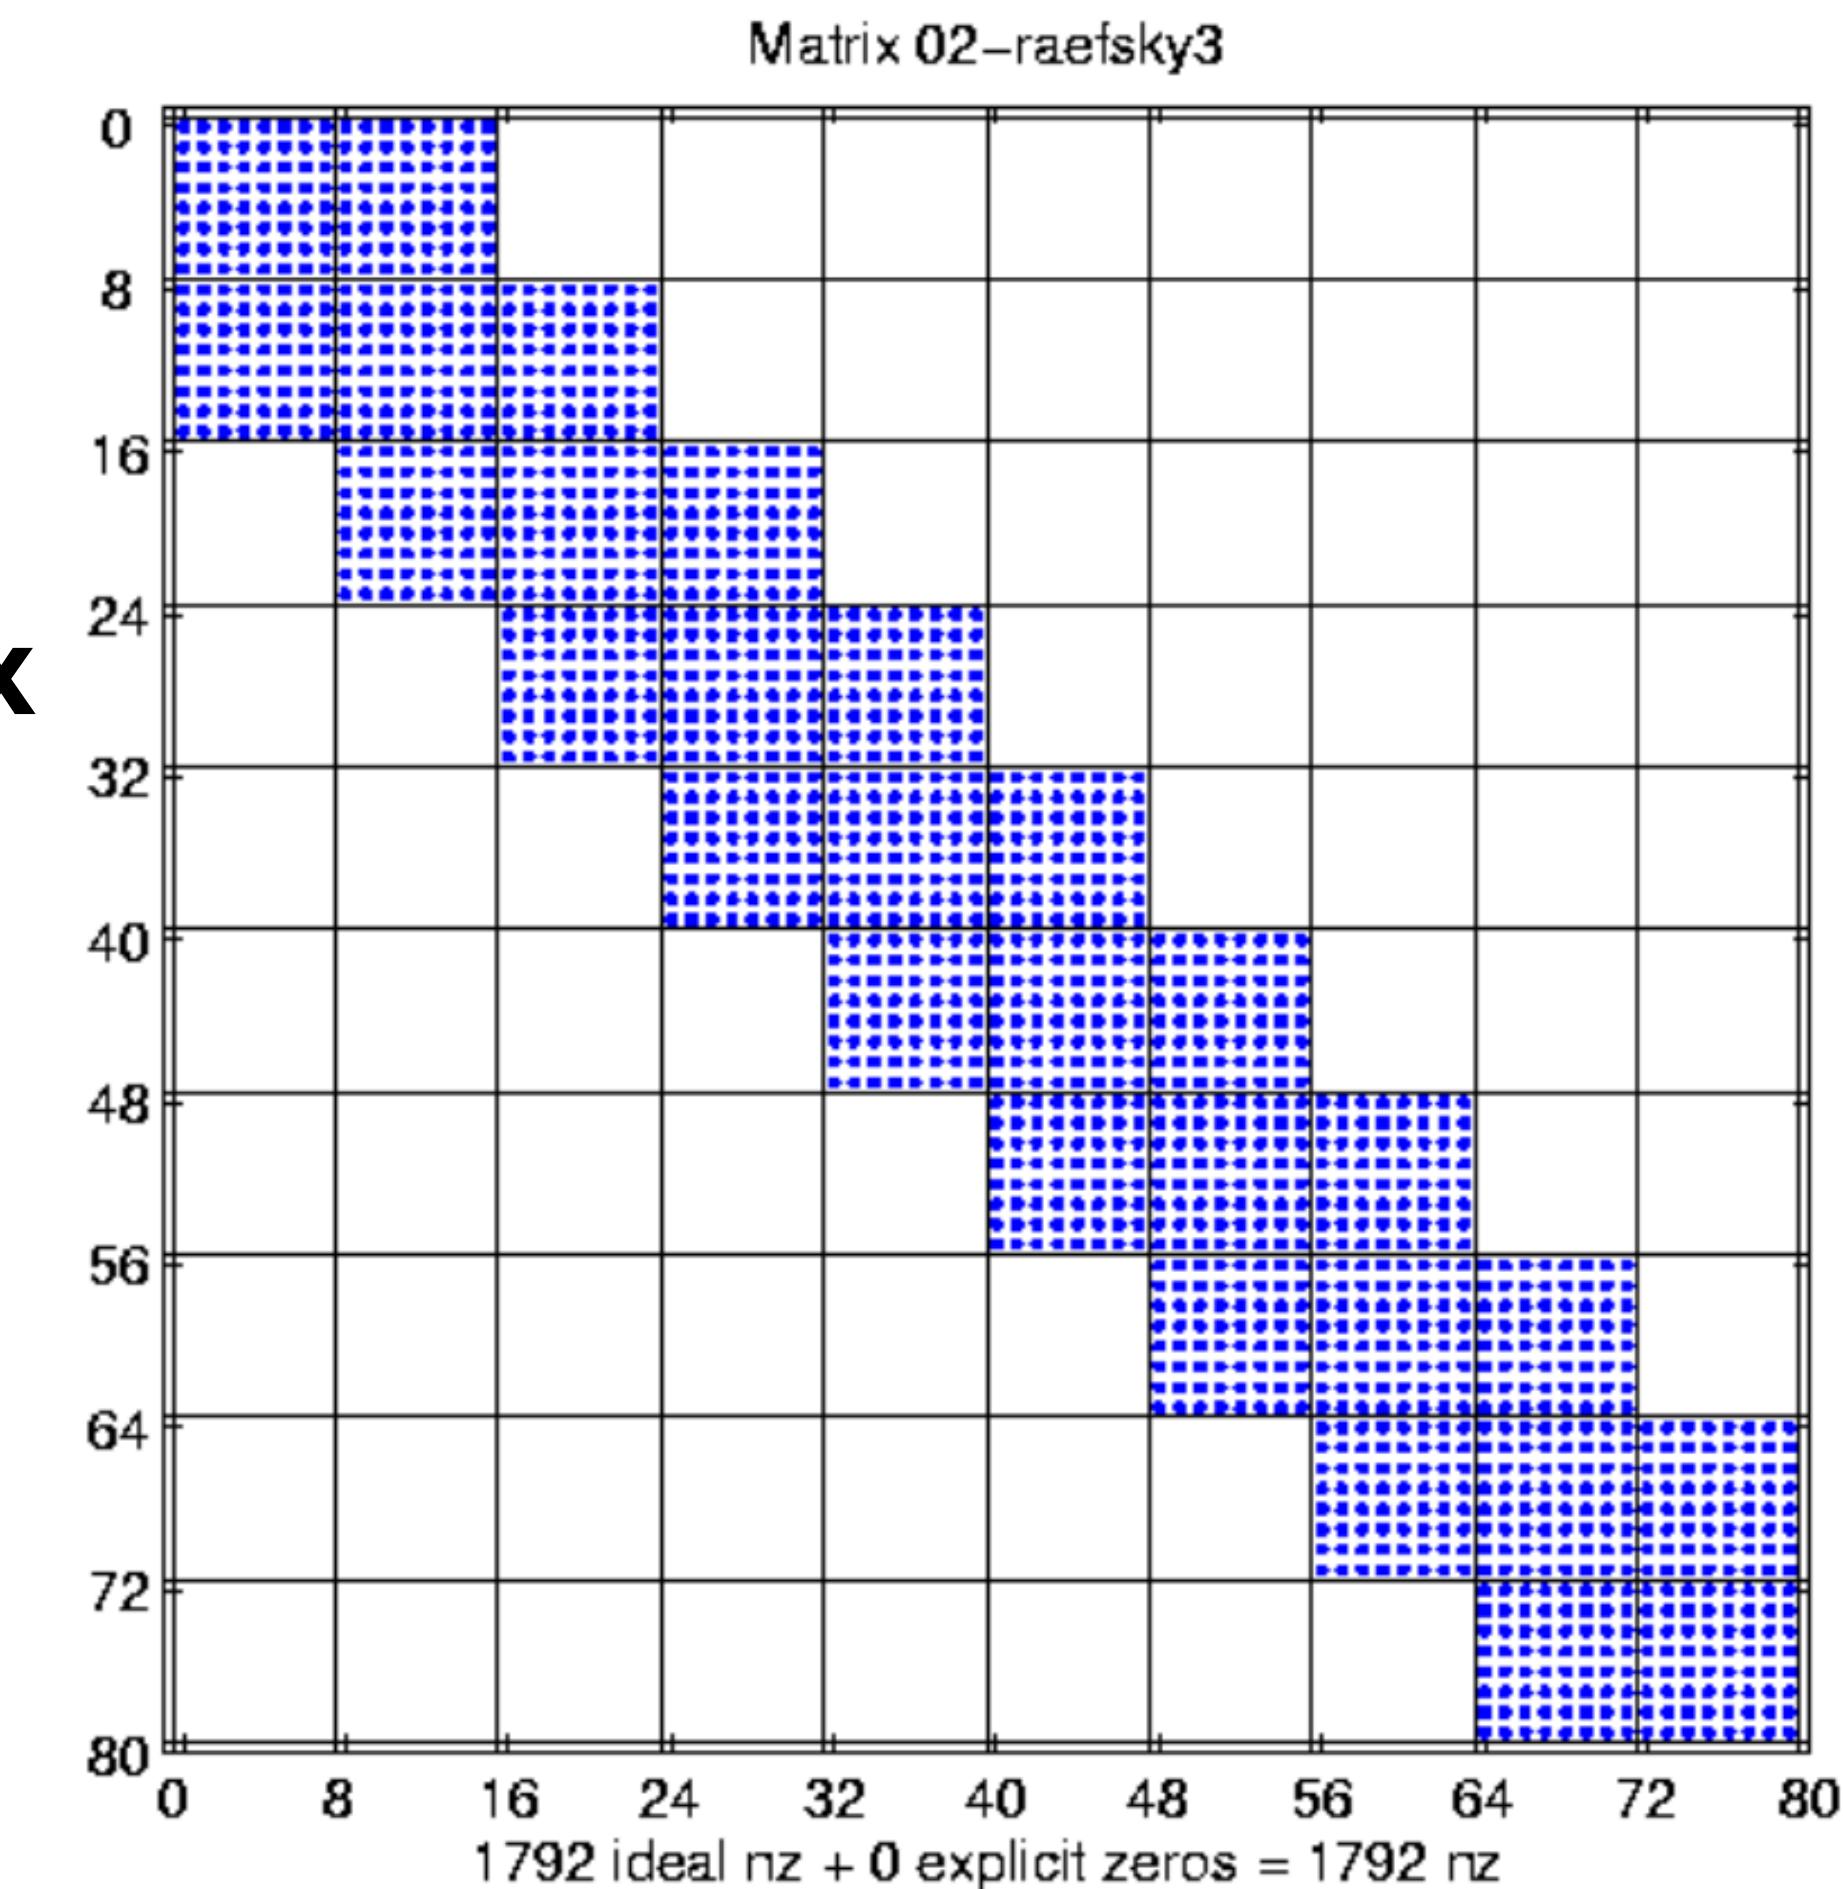
Small fraction  
of peak

# **Register/cache blocking and autotuning SpMV**

# Changing Matrix Format: Register Blocking



**Submatrix**



“Fast sparse matrix-vector multiplication by exploiting variable block structure,” Vuduc and Moon 2005.

# Using block structure in SpMV

The **bottleneck is the time to fetch** the matrix from memory

- Only 2 flops for each nz in matrix
- Fetching at  $\sim 1$  int (col idx) + 1 float (value) for 2 flops

1														
3	5													
		1												
		2												
		5	2	1										
2					4			3	1					
					5				2					

Non-zero Values    Column index    Row pointer  
# of stored elements: 14 values    # of fetches for a SpMV: 59 fetches  
# of fetches for a SpMV: 59 fetches    Filling Ratio: 100 %

CRS

Blocked Compressed Sparse Row (BCSR)

- Don't store each nonzero - instead, **store each nonzero r x c block** with 1 column index
- Time to fetch matrix from memory decreases

1														
3	5													
		1												
		2												
		5	2	1										
2					4			3	1					
					5				2					

Non-zero Values    Column index    Row pointer  
# of stored elements: 28 values    # of fetches for a SpMV: 30 fetches  
# of fetches for a SpMV: 30 fetches    Filling Ratio: 50 %

BCRS (2x2)

**Change both data structure and algorithm** - need to pick r and c and change algorithm accordingly

1														
3	5													
		1												
		2												
		5	2	1										
2					4			3	1					
					5				2					

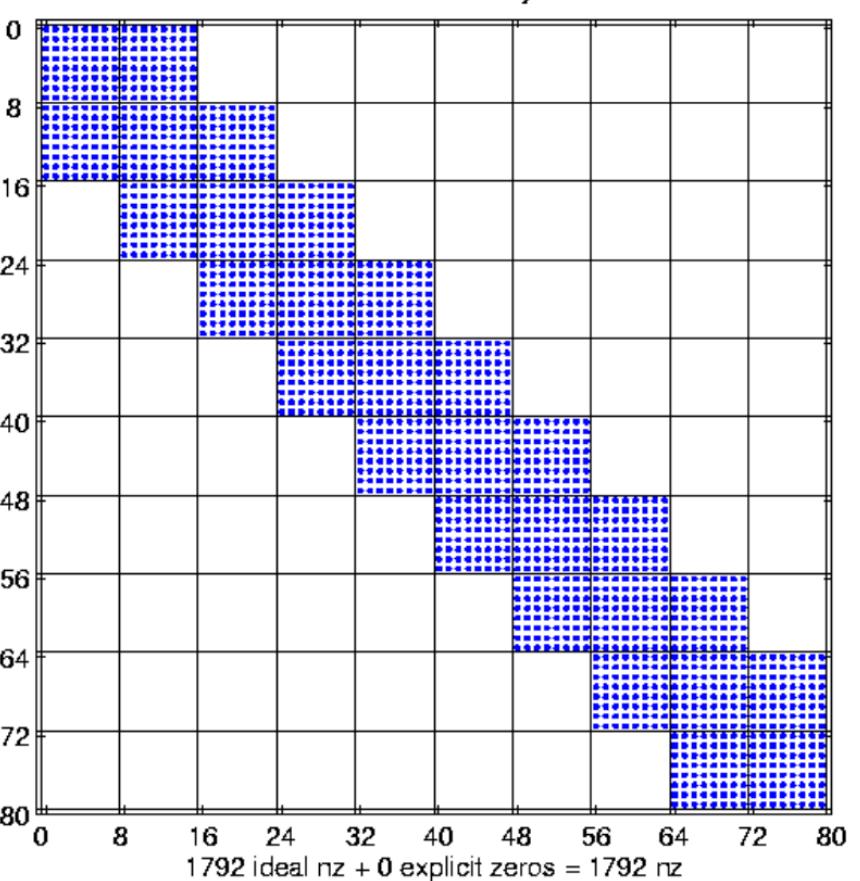
Non-zero Values    Column index    Row pointer  
# of stored elements: 48 values    # of fetches for a SpMV: 23 fetches  
# of fetches for a SpMV: 23 fetches    Filling Ratio: 29 %

BCRS (4x4)

# The need for search

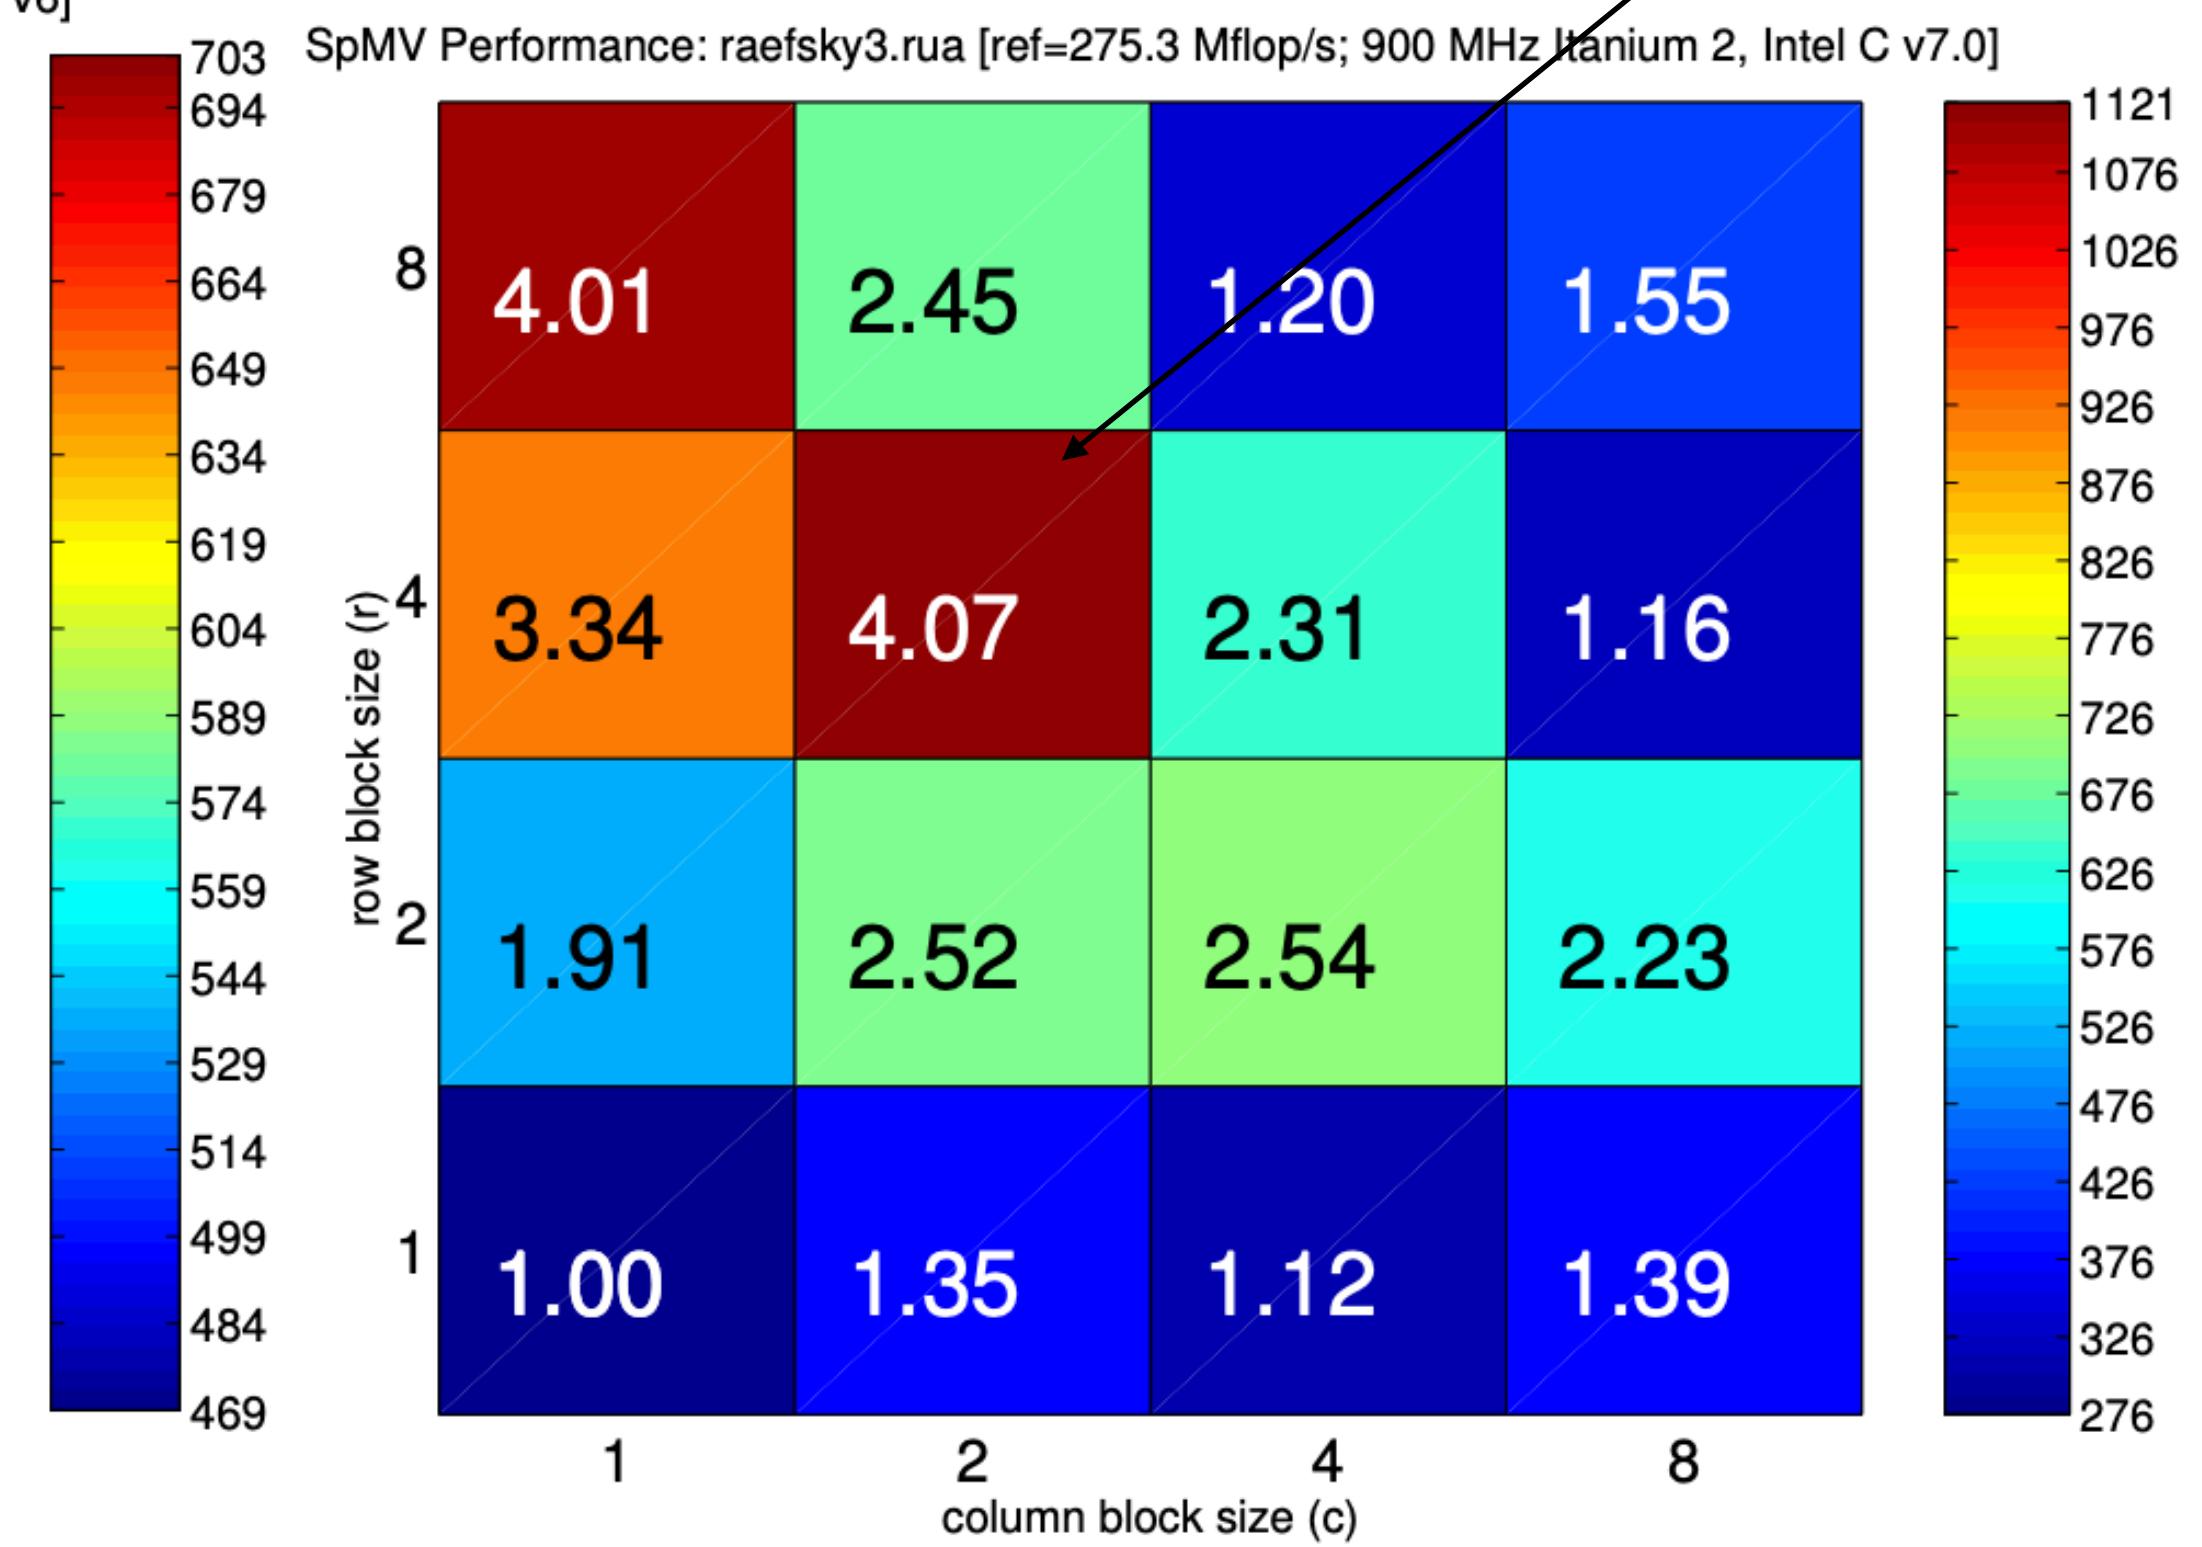
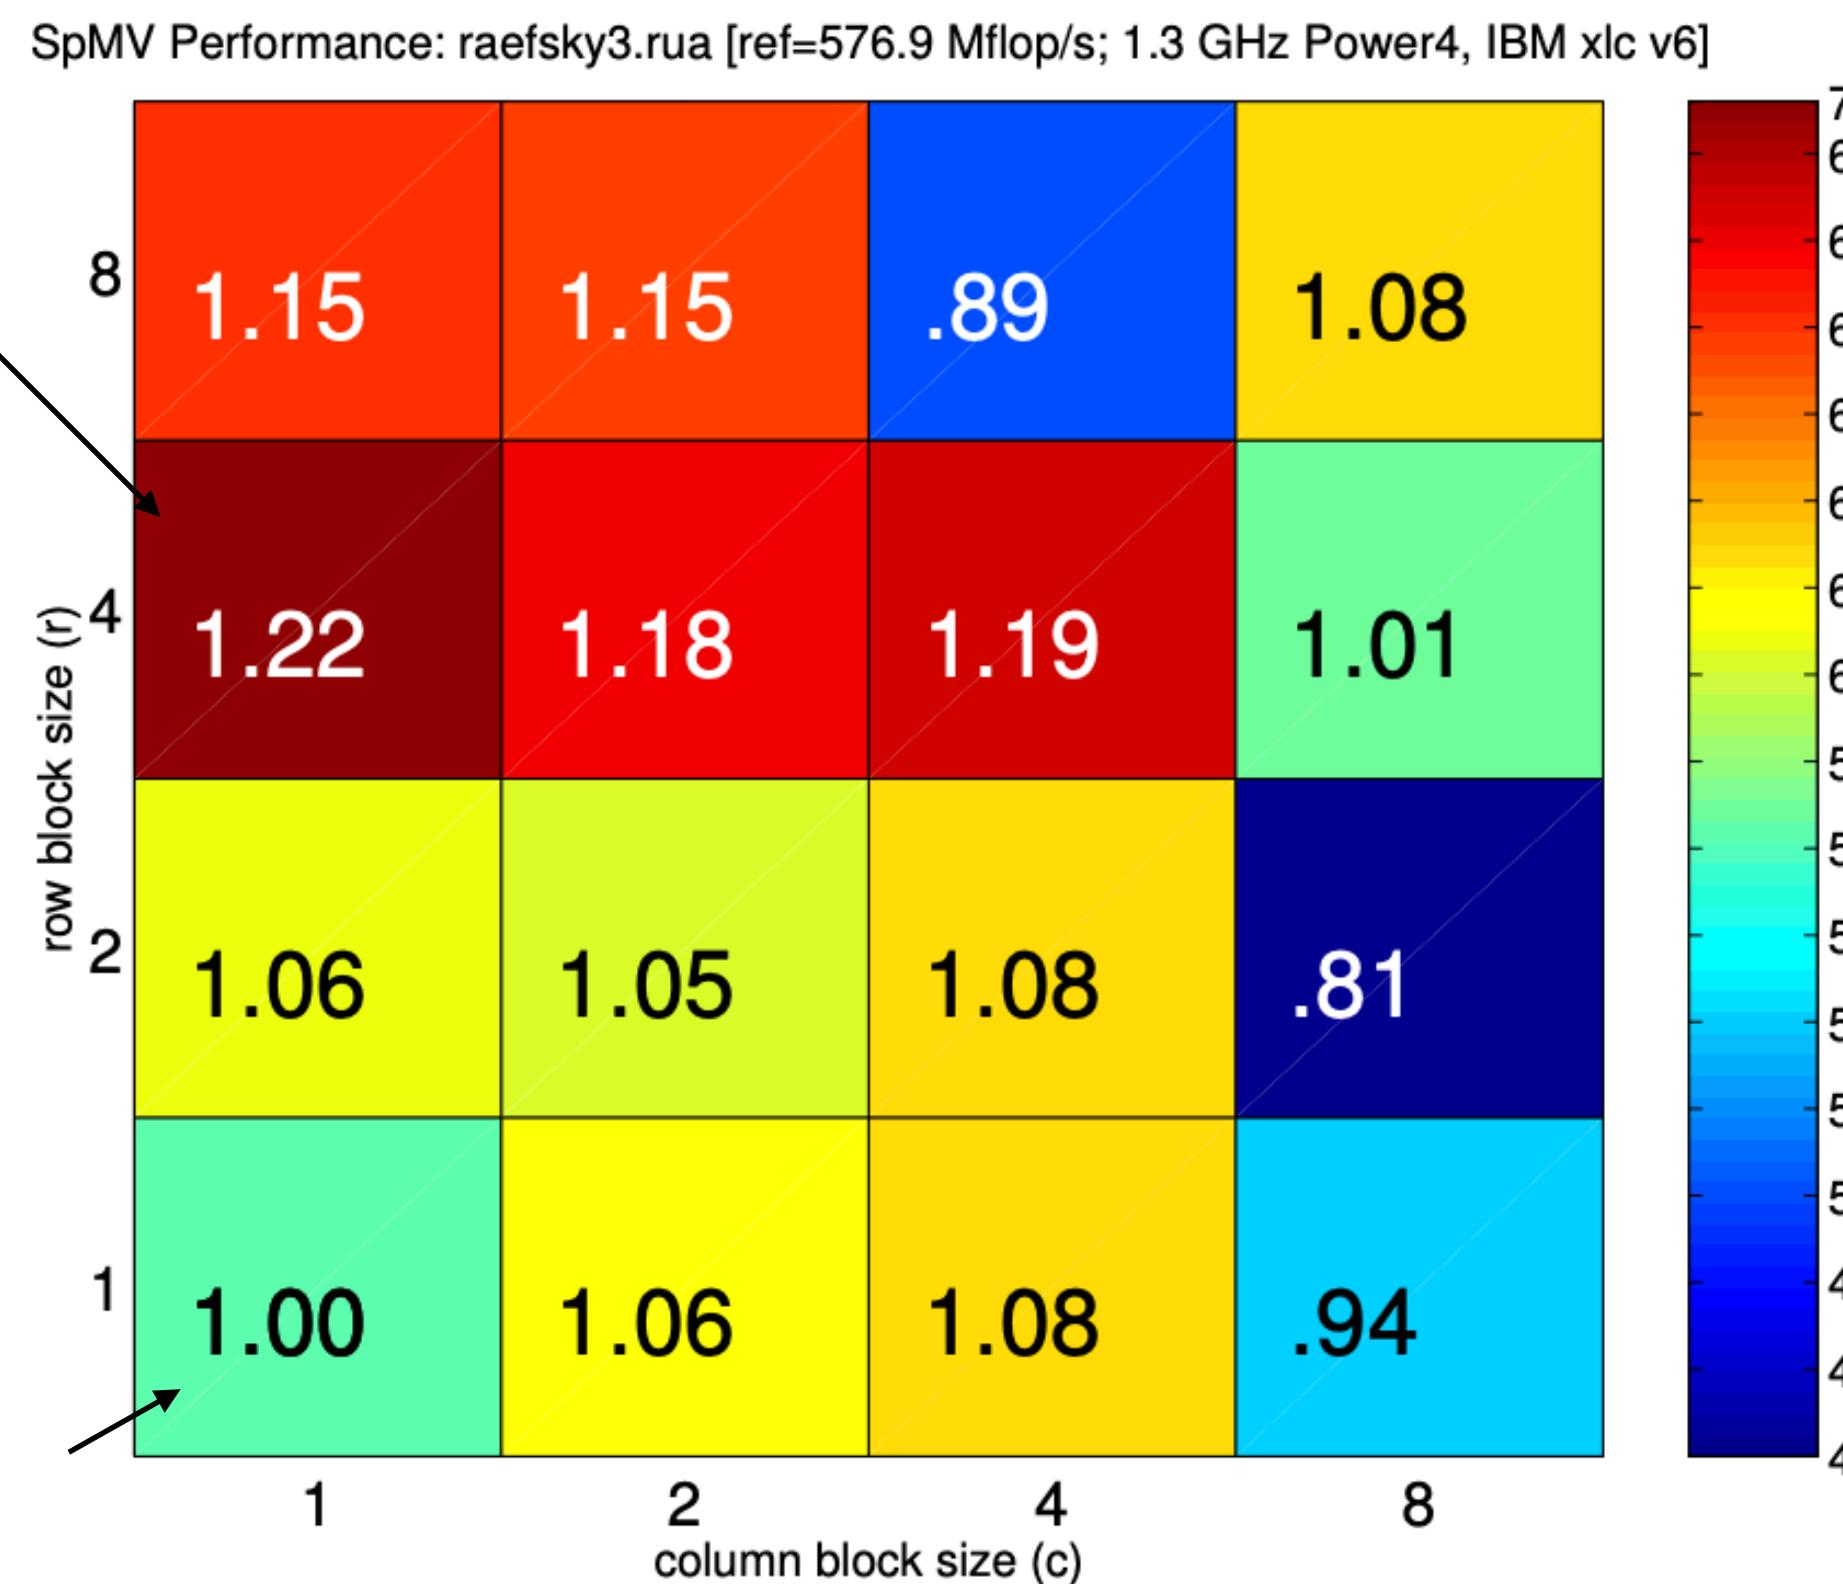
In the previous example, it seems like 8x8 is a natural choice.

Out of 6 platforms tested, 8x8 is the best on one (coming up)



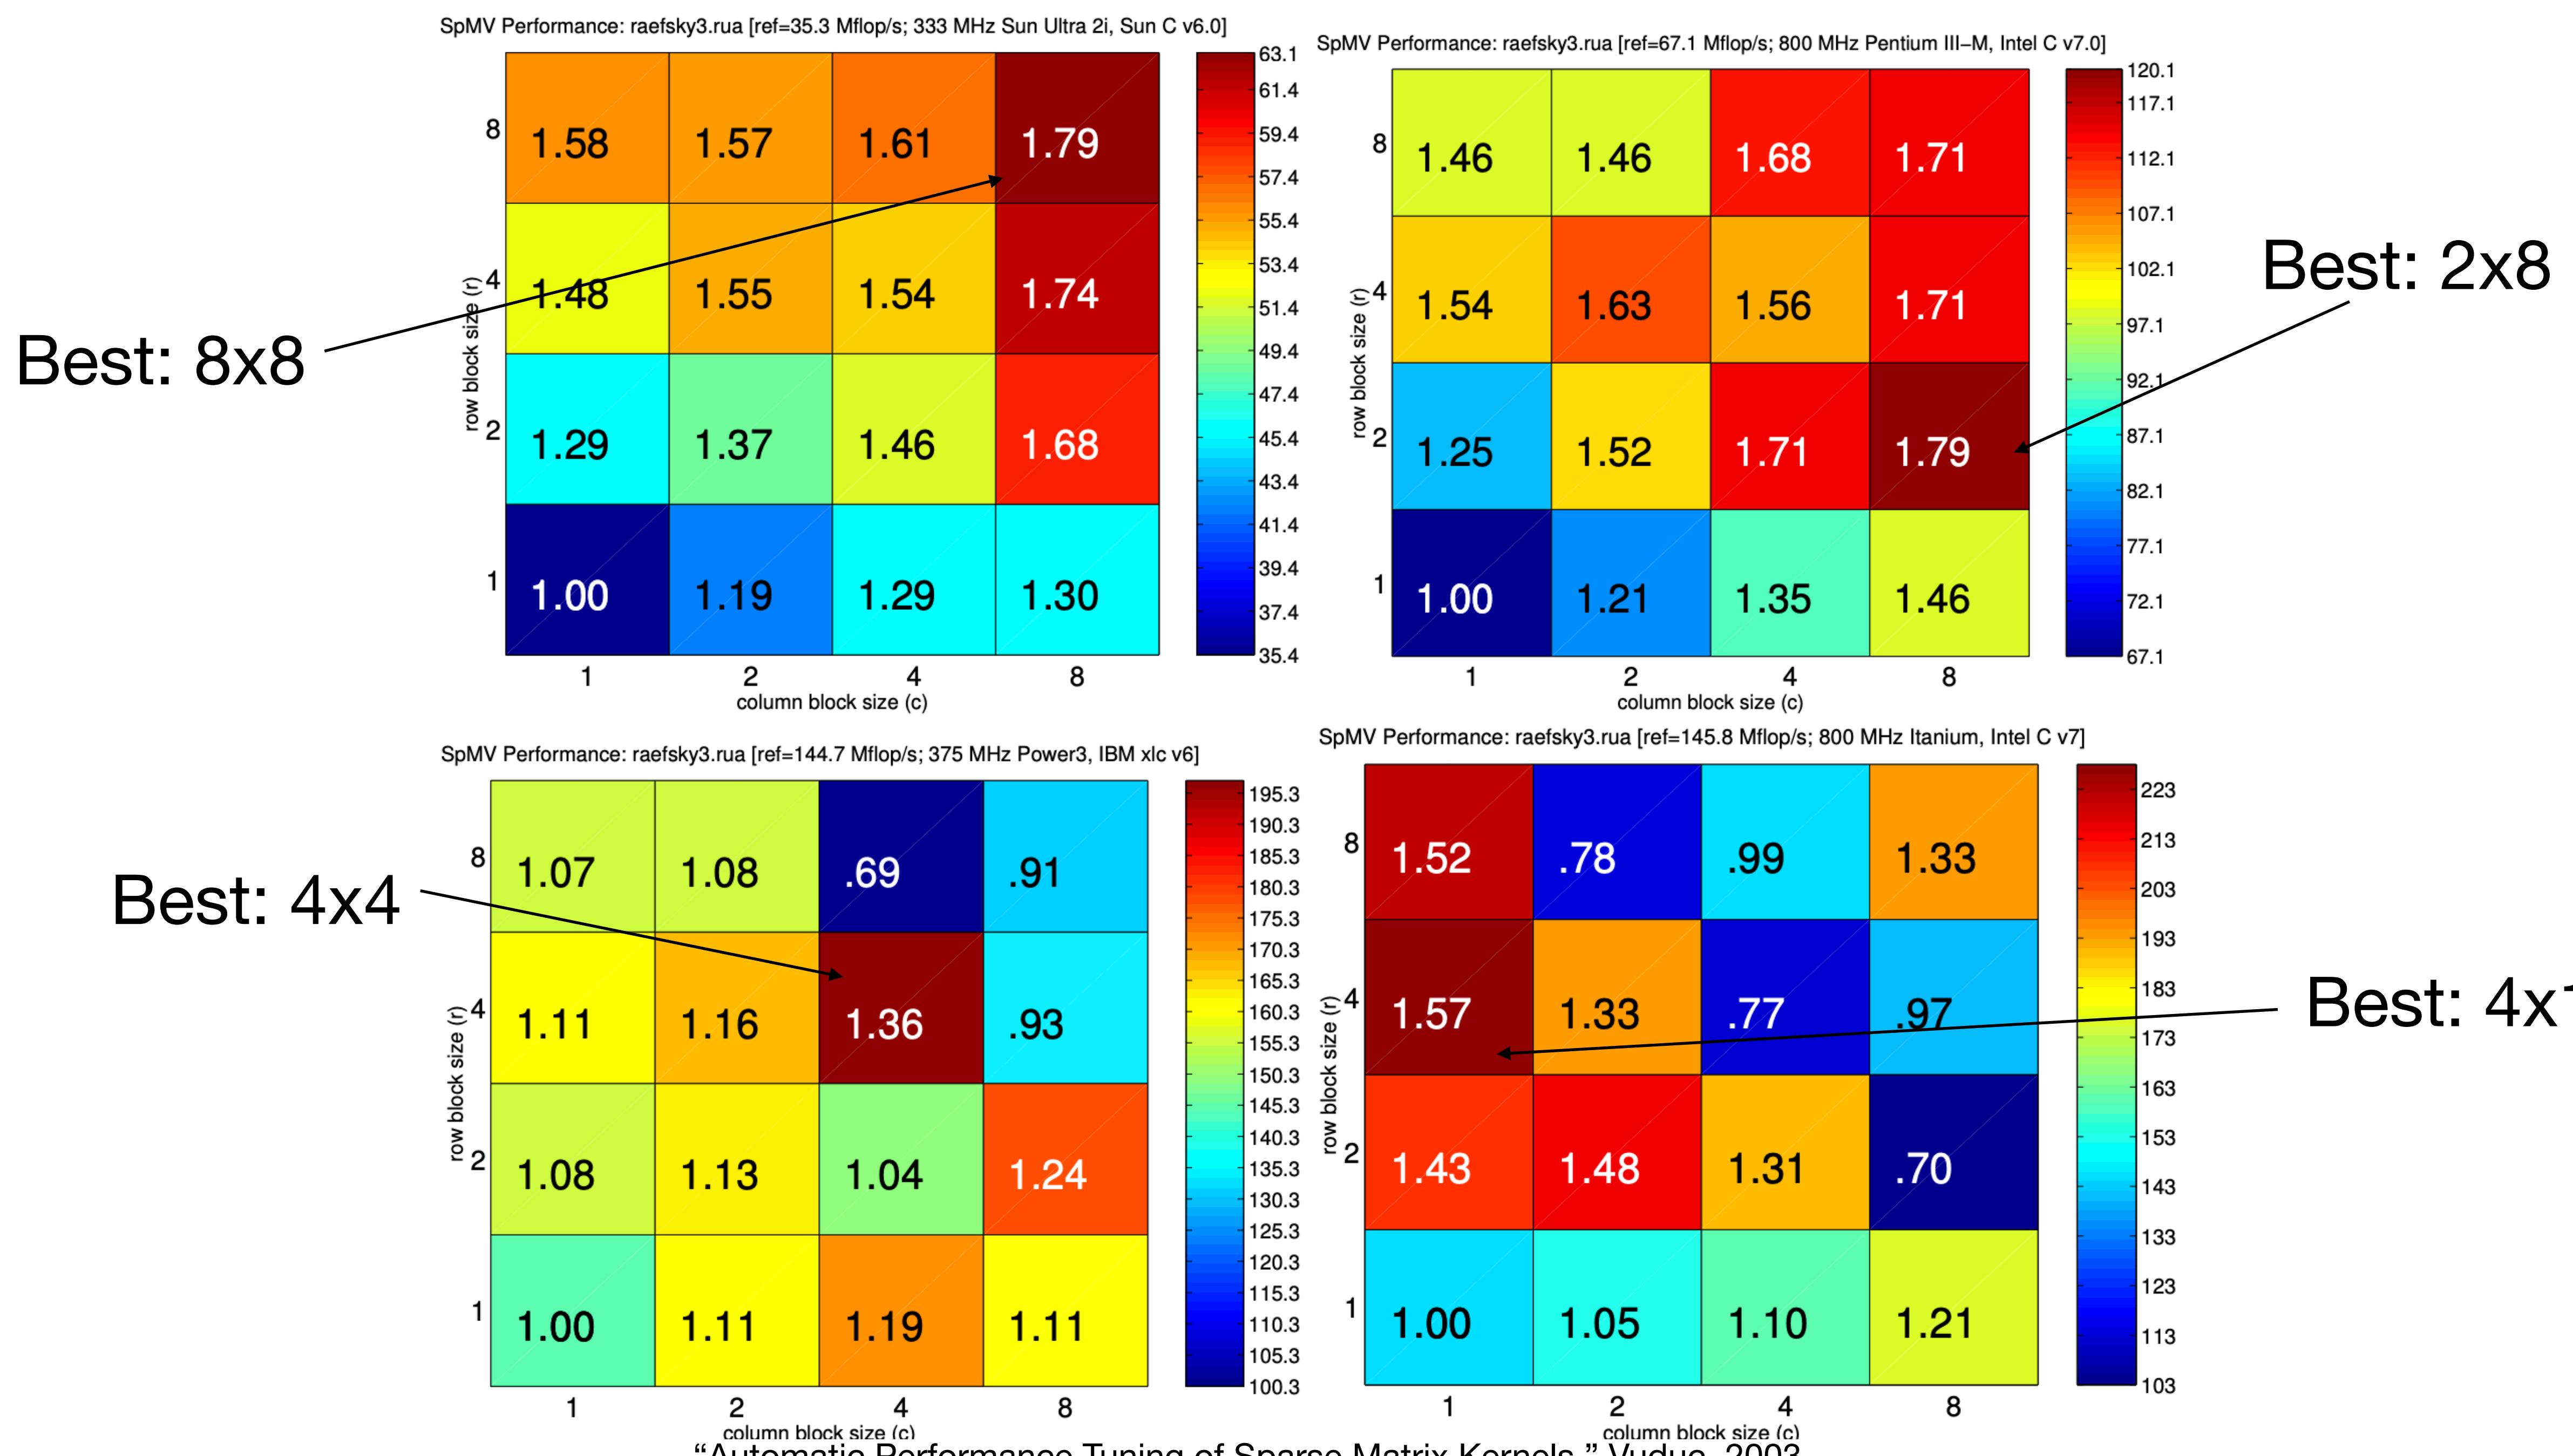
Best: 4x2

Best: 4x1

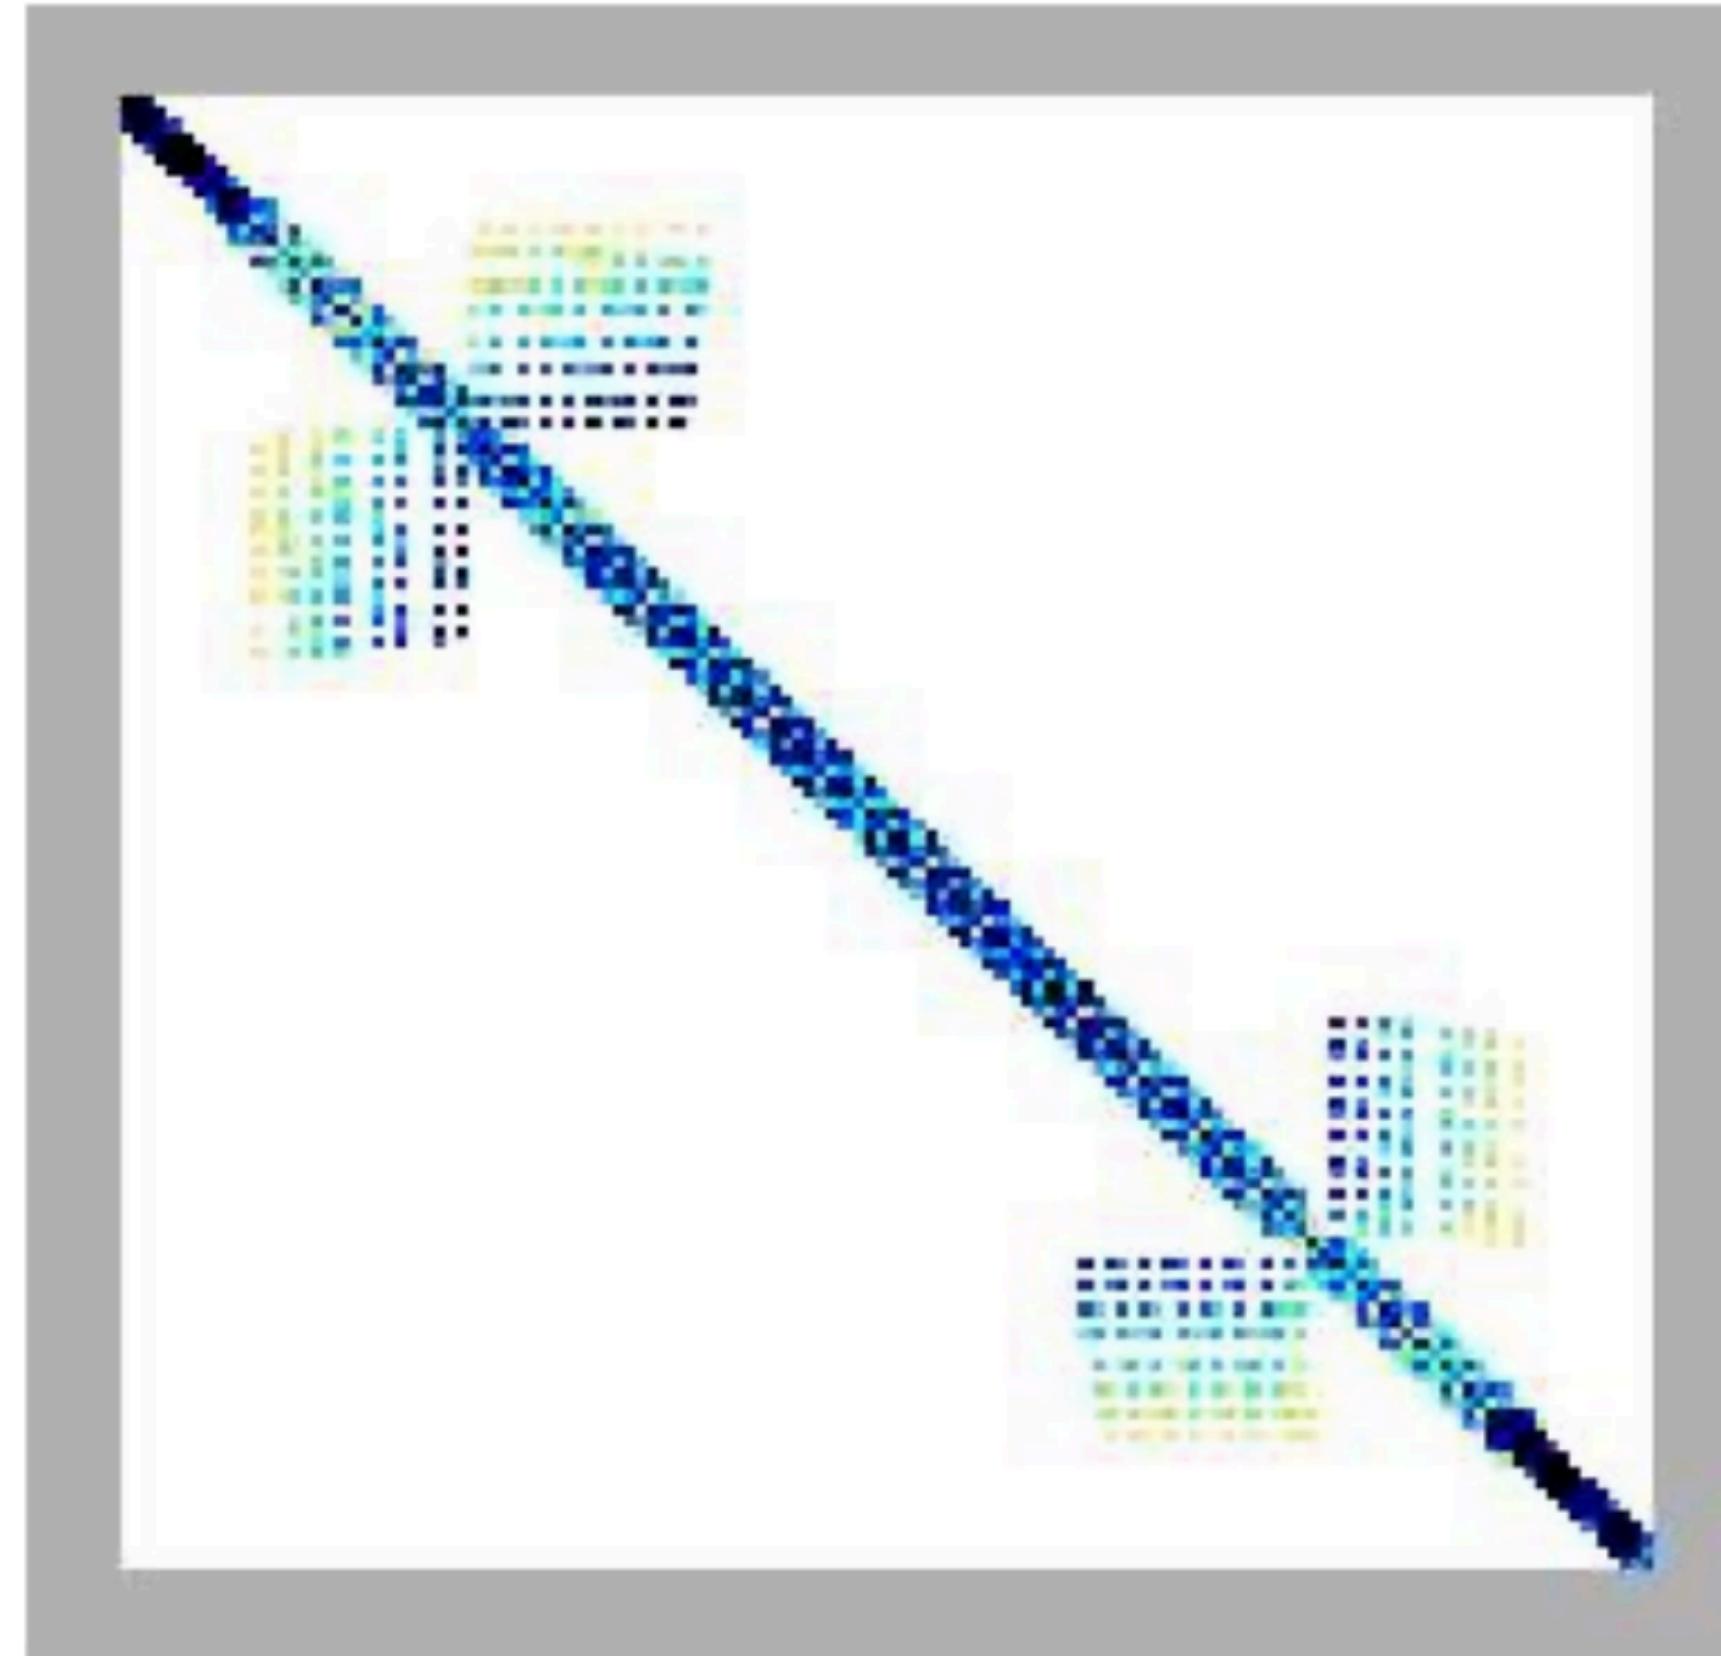


Reference

"Automatic Performance Tuning of Sparse Matrix Kernels," Vuduc, 2003.



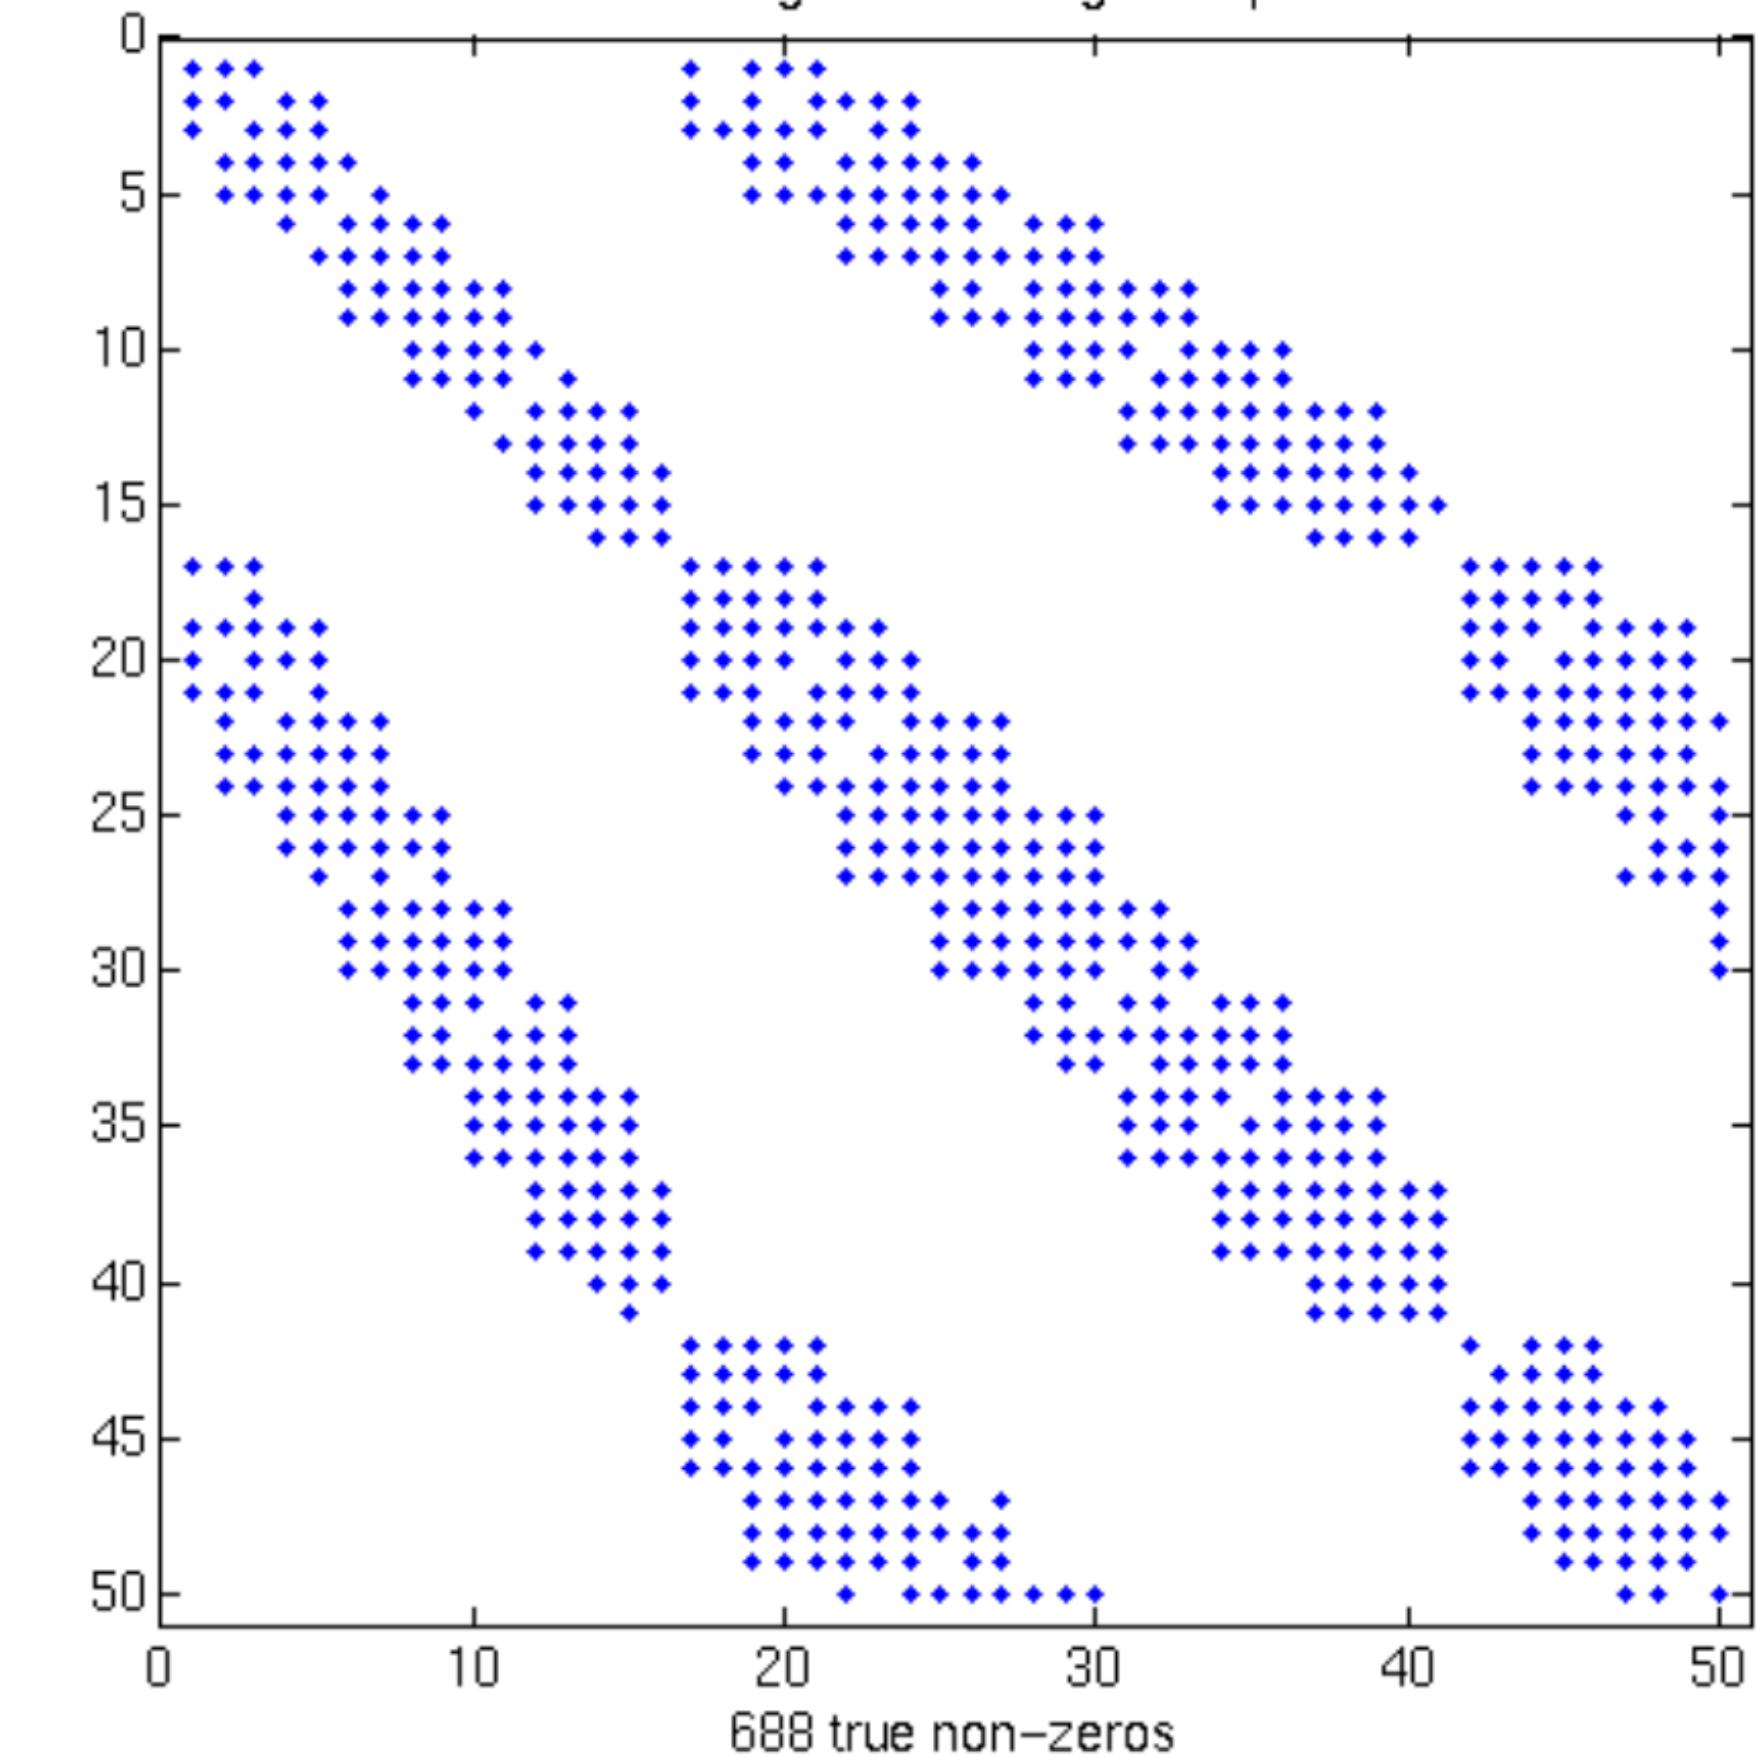
# But most matrices don't block so easily



Zoom into  
top corner



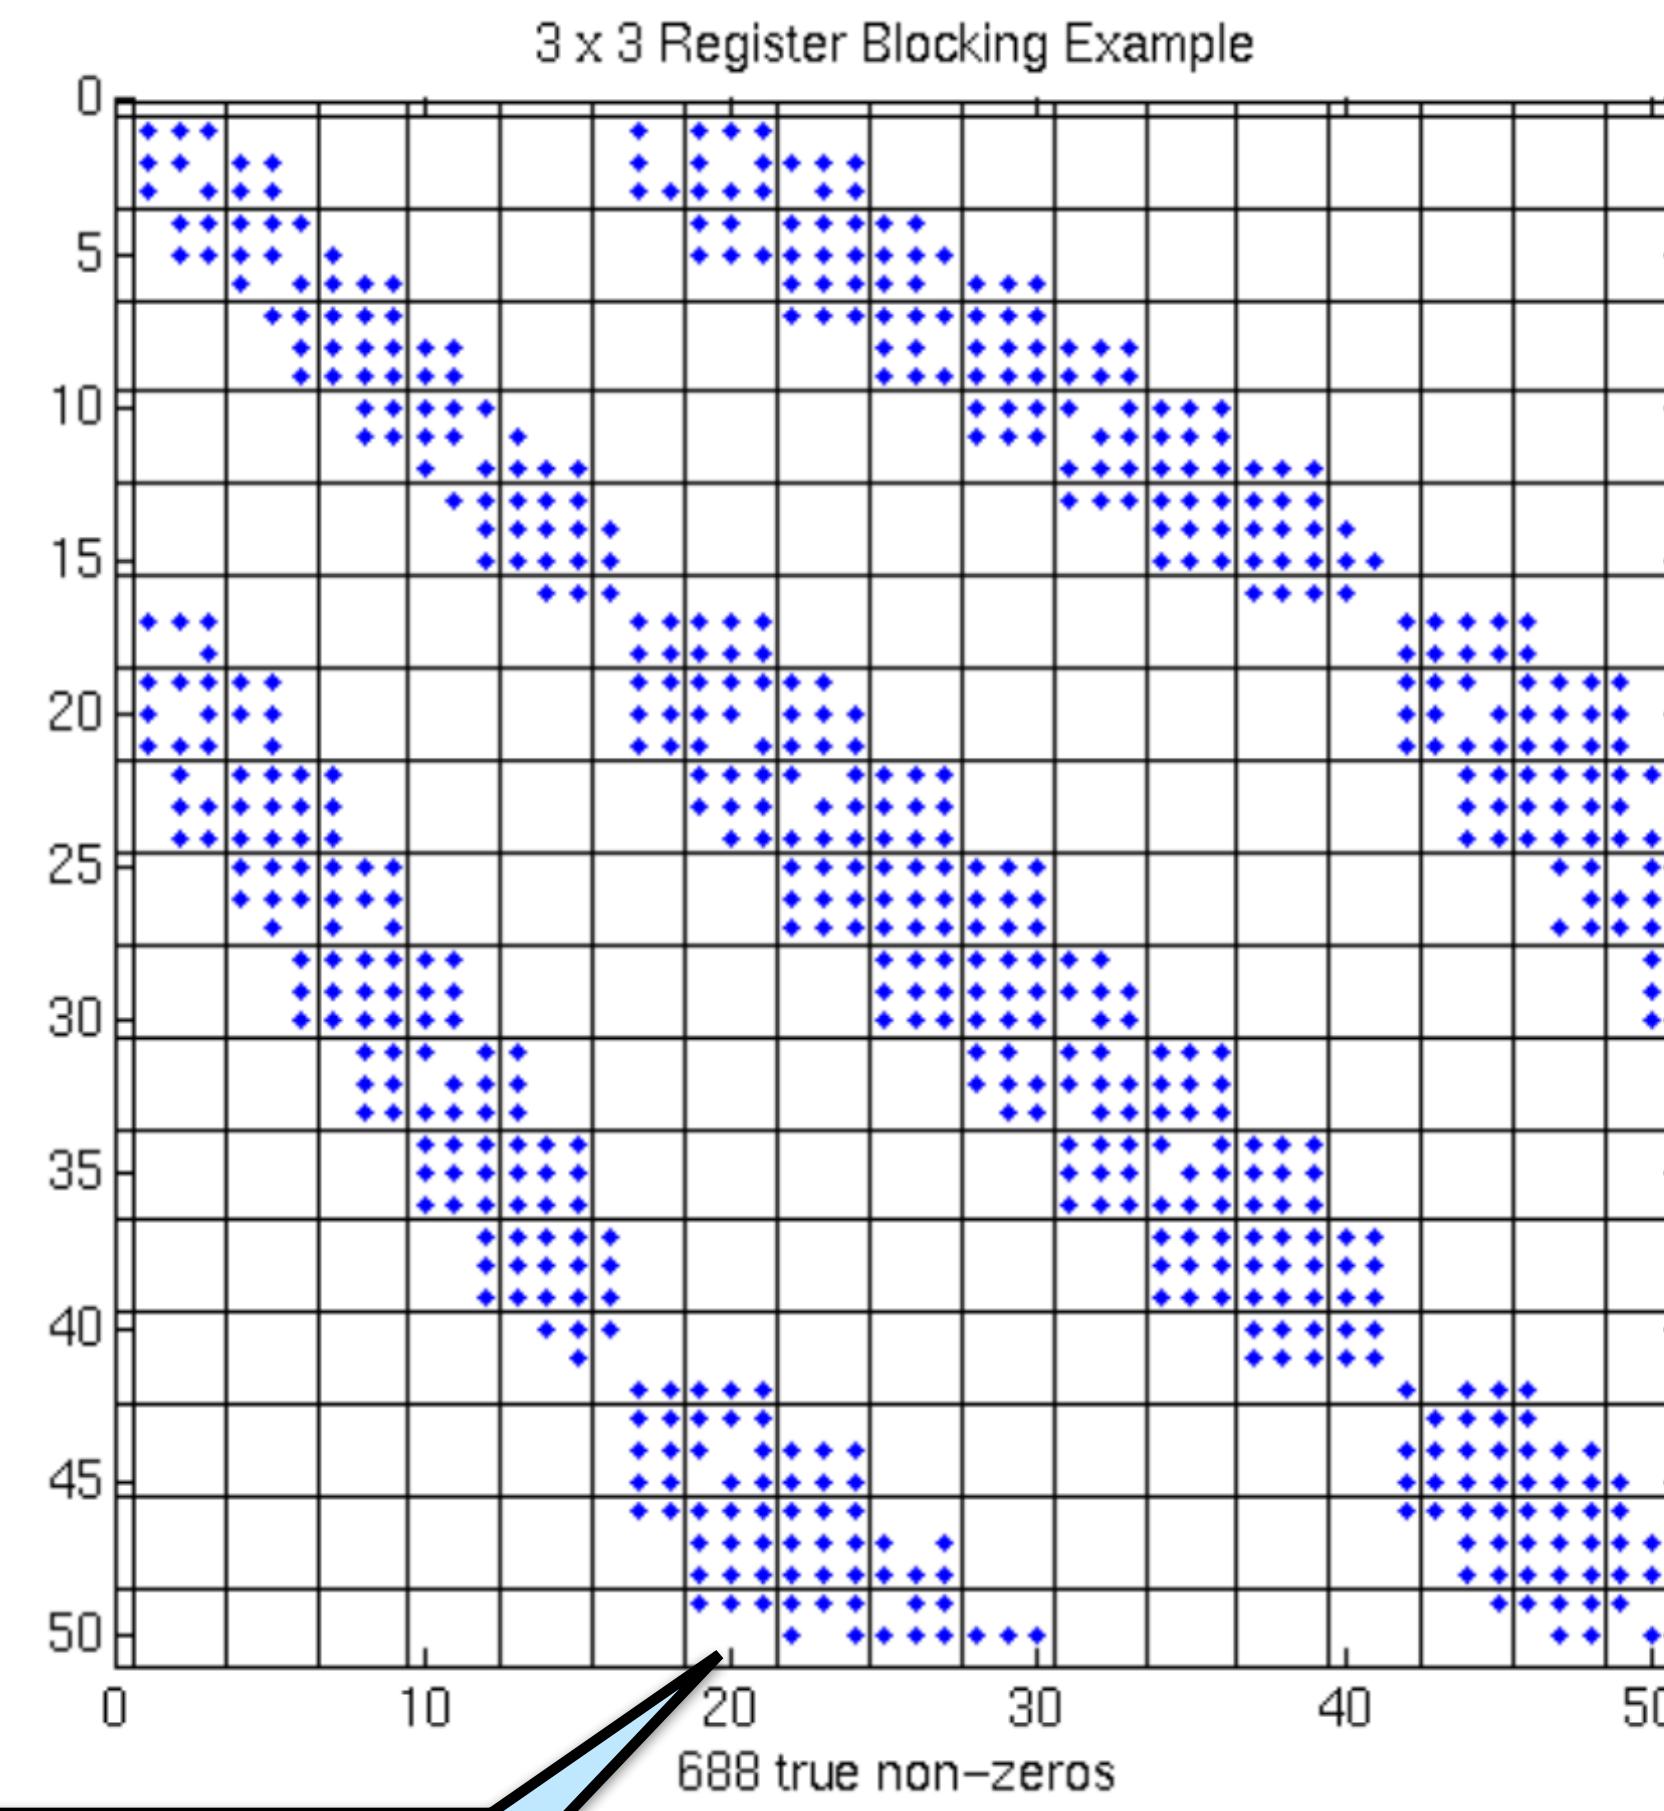
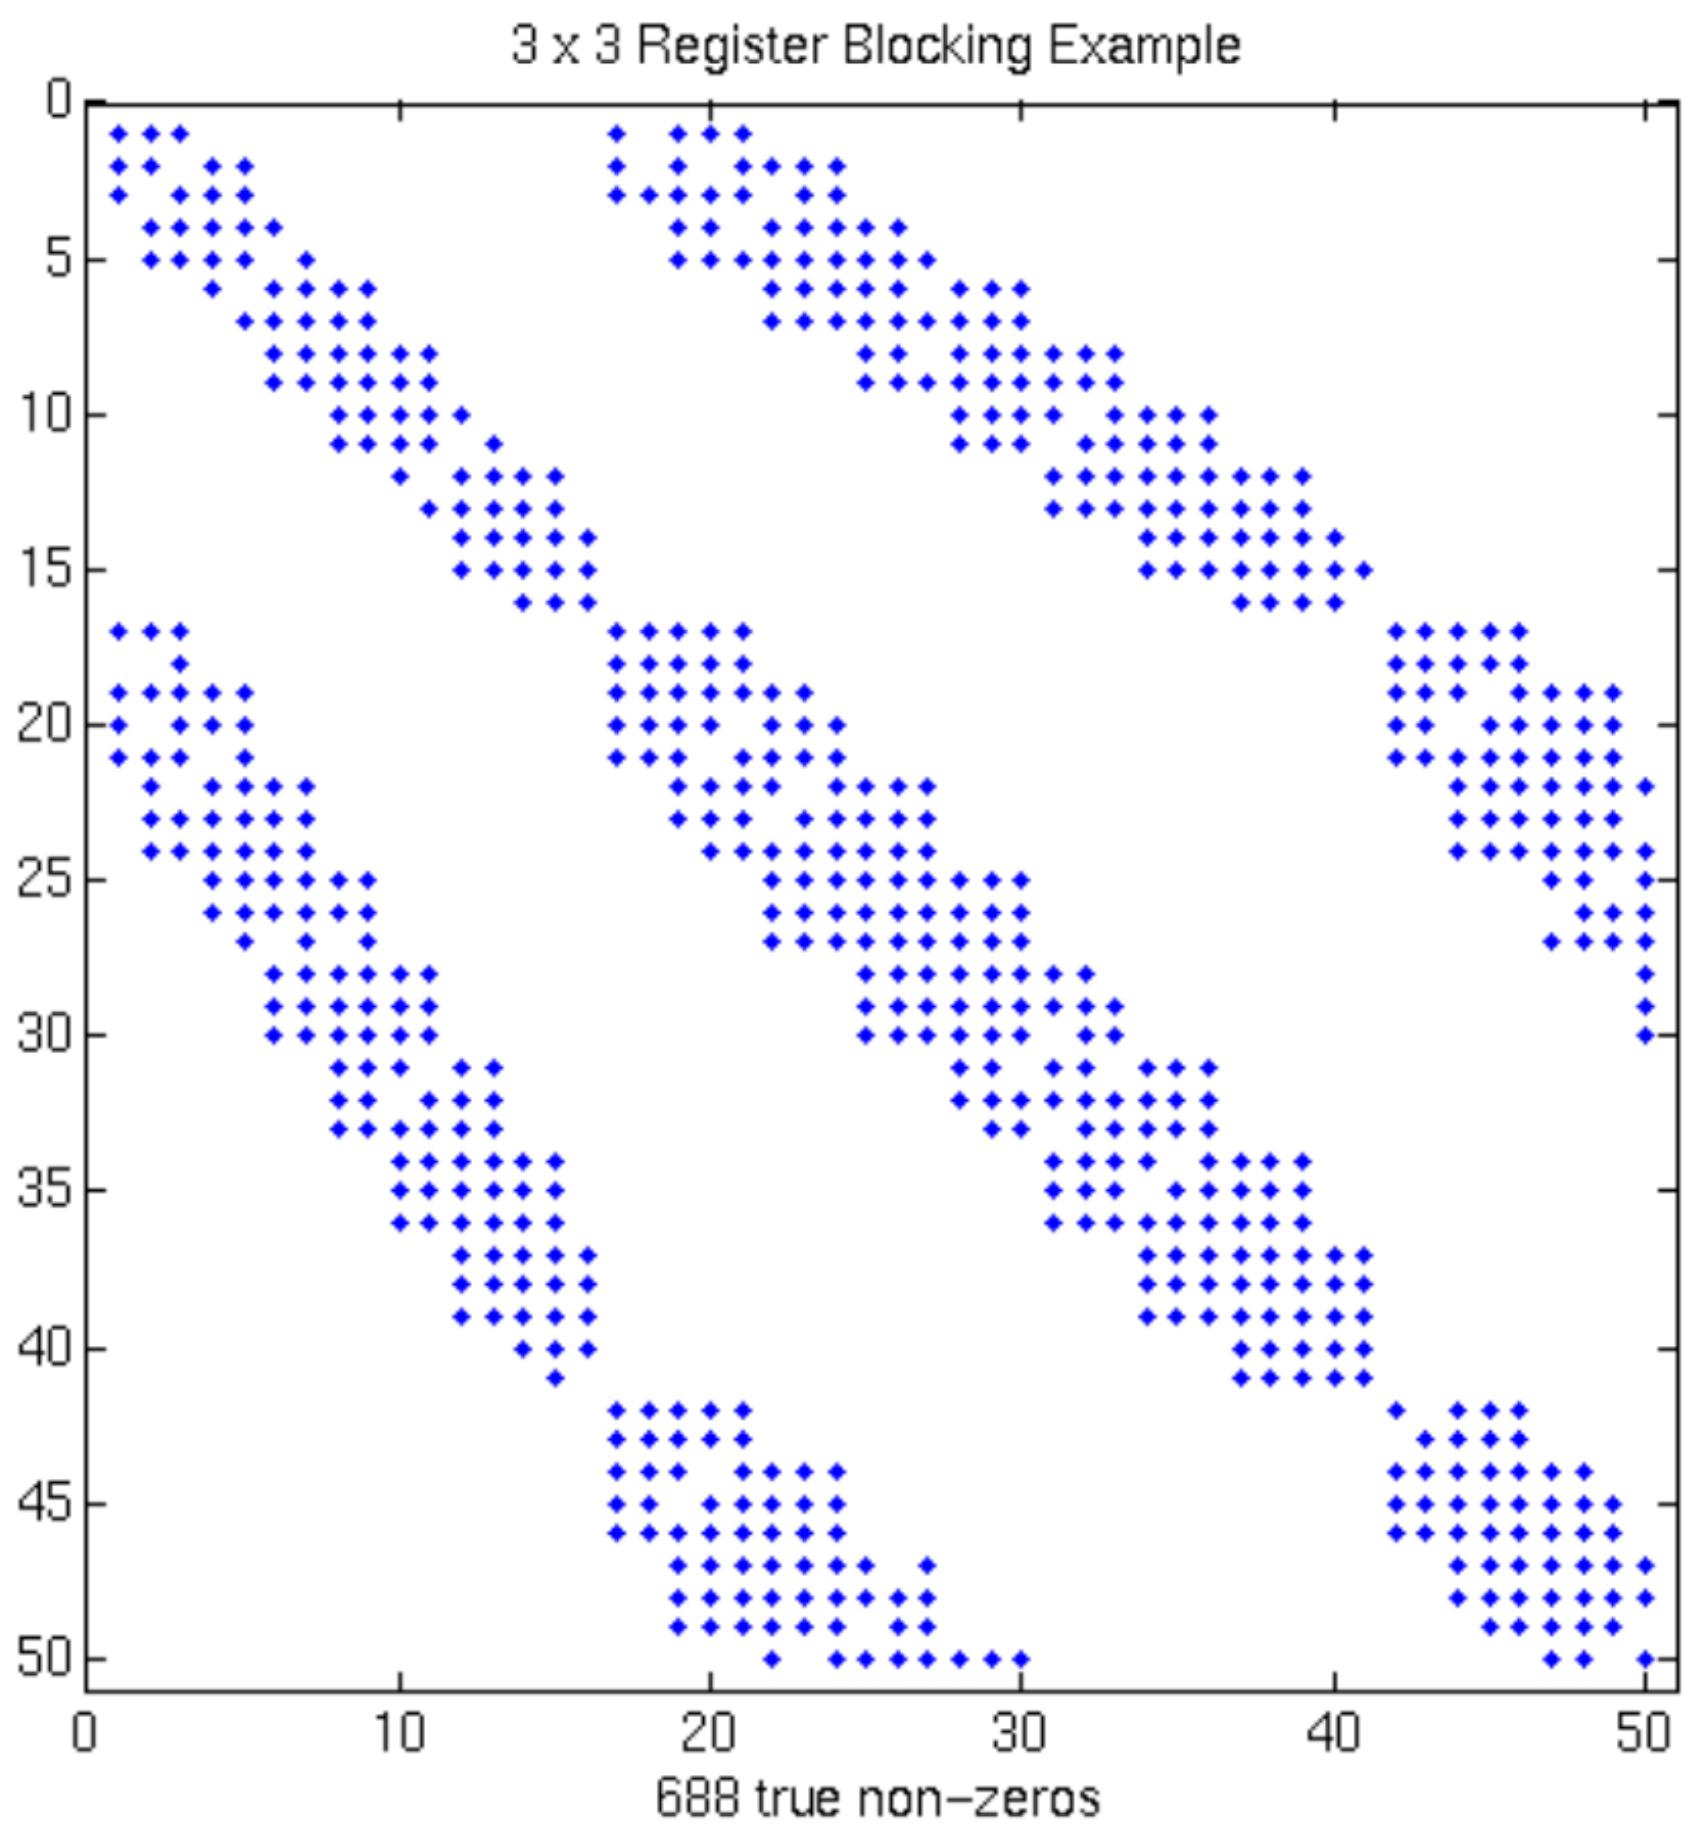
3 x 3 Register Blocking Example



Fluid dynamics problem - more complicated nonzero structure

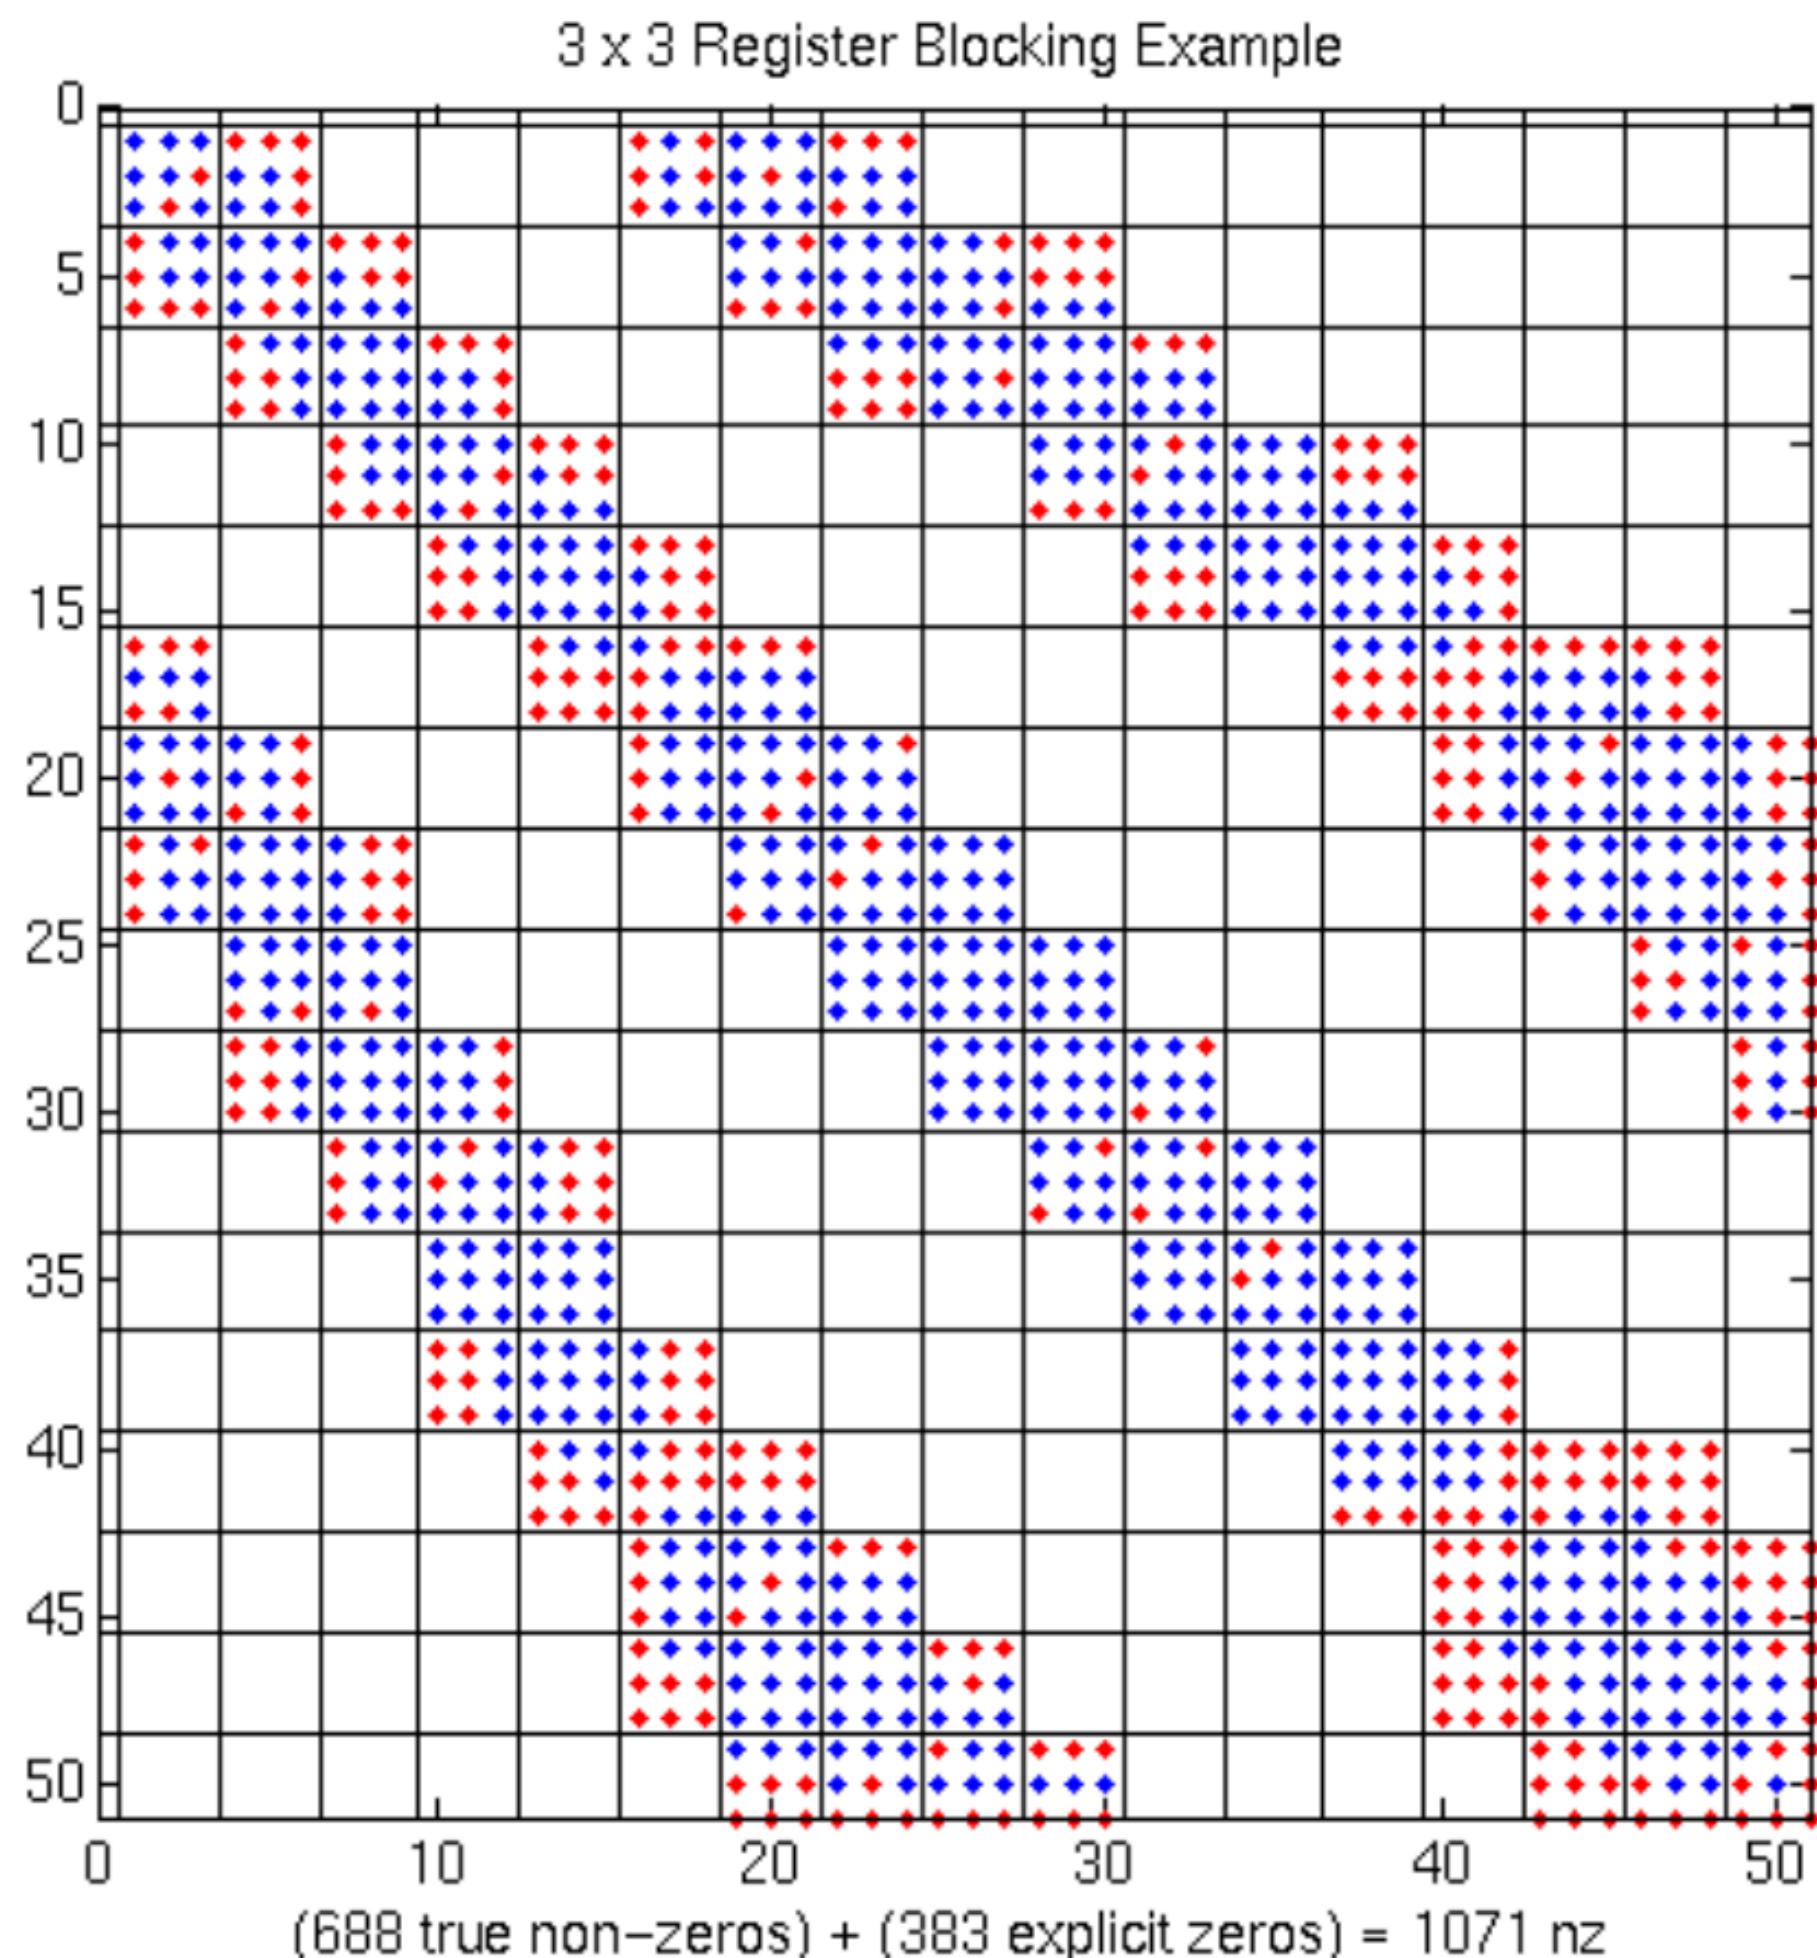
Total nnz = 1.1M

# 3x3 blocks look natural, but...



Many blocks  
are not full

# Extra work can improve efficiency



- Add explicit zeroes: 1.5x “fill overhead”
- Unroll loops
- More work but faster - 1.5x faster on PIII

# Libraries for sparse matrices

How to build optimized library when:

- Formats are not known? Libraries like PETSc and Trilinos will let the user provide format and SpMV

How to build optimized matrix kernel library (BLAS)?

- Nonzero structure is key to optimization

OSKI = Optimized Sparse Kernel Interface

- pOSKI for multicore



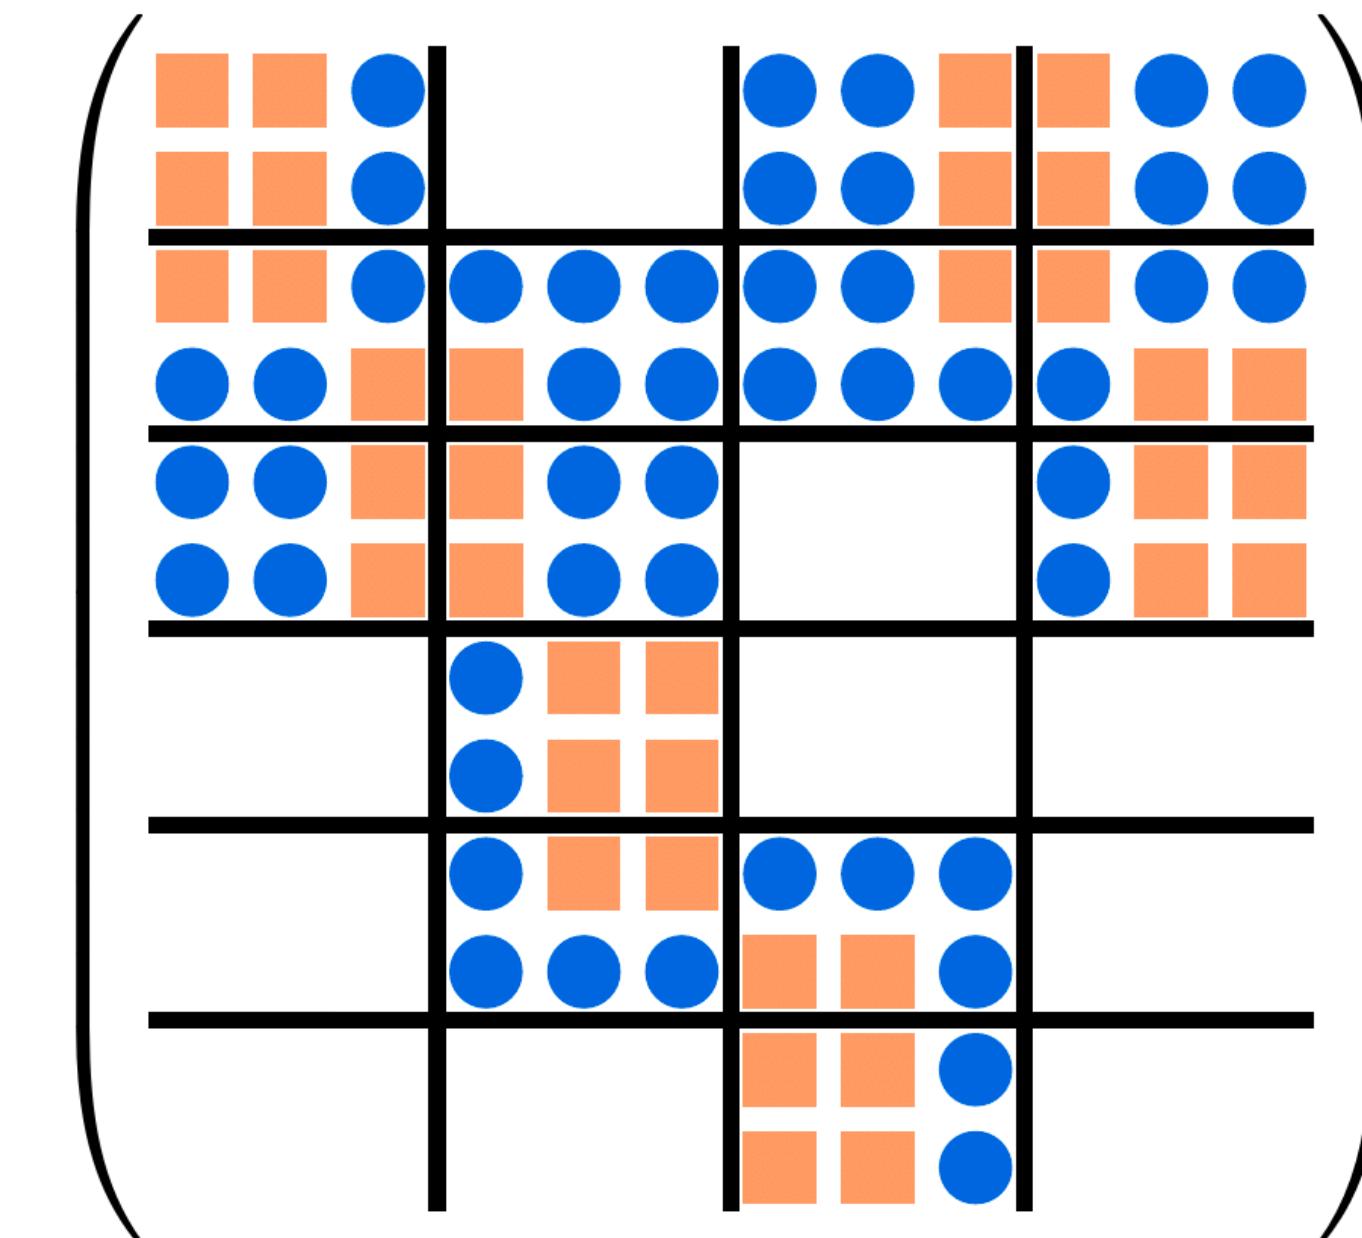
# Using the Fill to Formalize Blocking Scheme

## Quality [Im & Yelick, 01]

The fill of a  $N$ -dimensional tensor  $A$  is defined with respect to

- the number of nonzeros  $k(A)$
- a blocking  $\mathbf{b} = (b_1, \dots, b_N)$
- the number of nonempty blocks  $k_{\mathbf{b}}(A)$

$$f_{\mathbf{b}}(A) = \frac{b_1 b_2 \dots b_N k_{\mathbf{b}}(A)}{k(A)}$$



# Using the Fill to Formalize Blocking Scheme

## Quality [Im & Yelick, 01]

The fill of a  $N$ -dimensional tensor  $A$  is defined with respect to

- the number of nonzeros  $k(A)$
- a blocking  $\mathbf{b} = (b_1, \dots, b_N)$
- the number of nonempty blocks  $k_{\mathbf{b}}(A)$

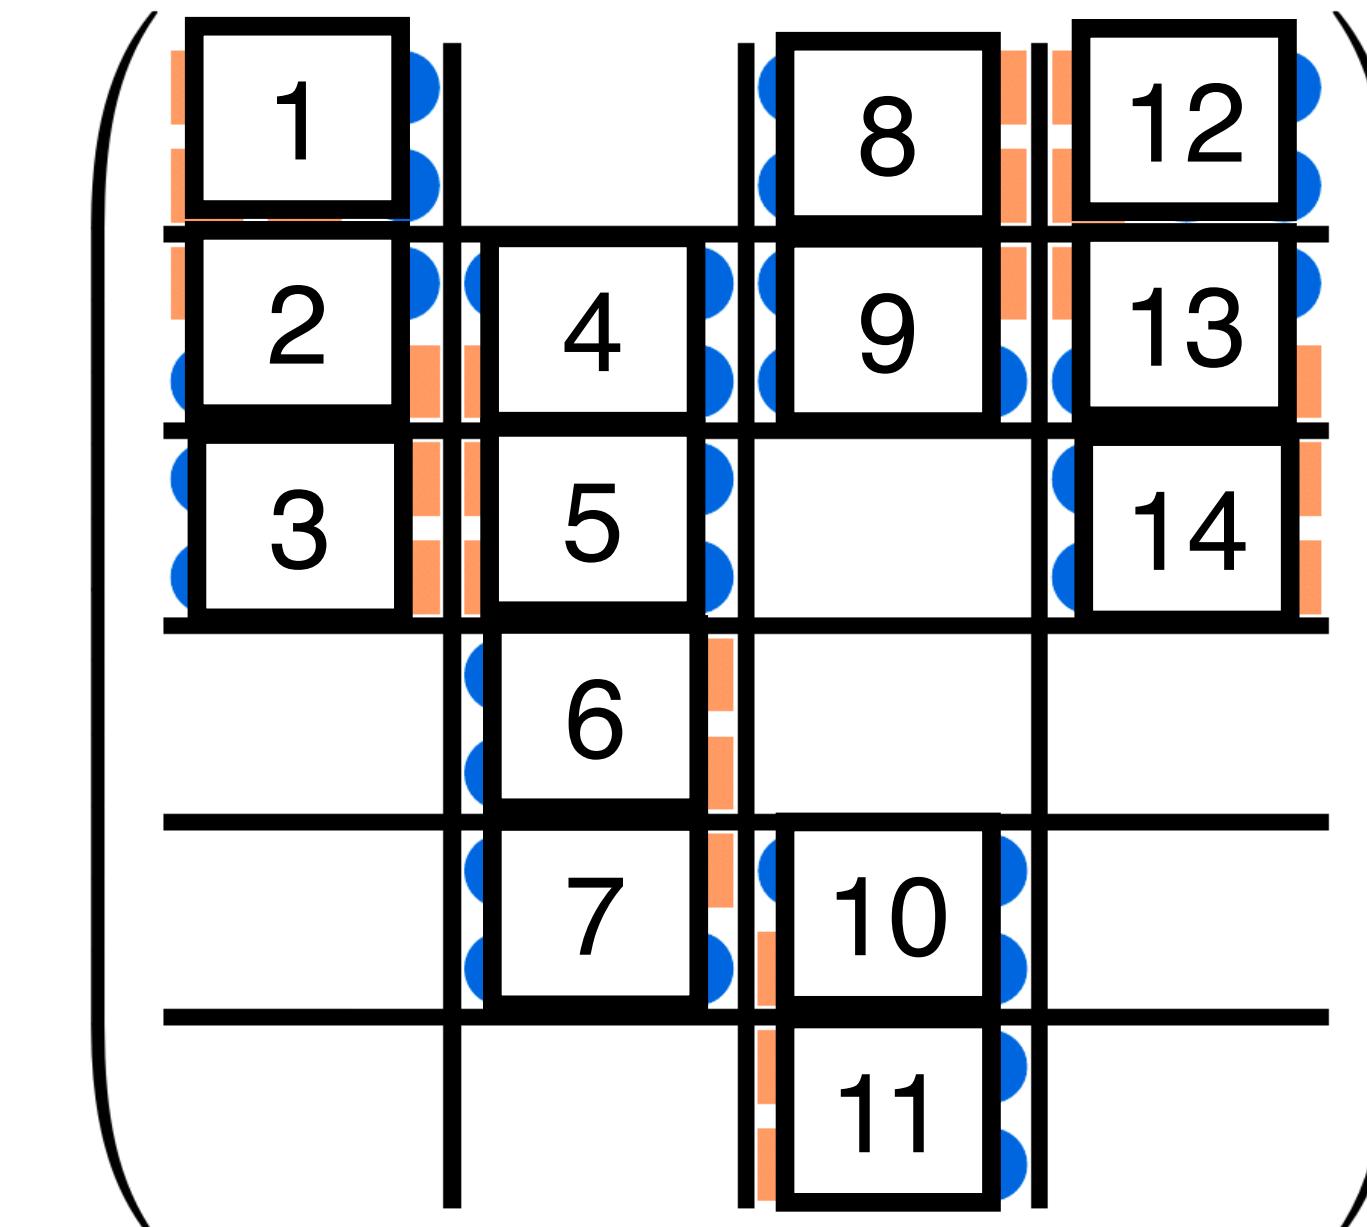
$$f_{\mathbf{b}}(A) = \frac{b_1 b_2 \dots b_N k_{\mathbf{b}}(A)}{k(A)}$$

Greater ~  
more wasted  
space

$$= \frac{2 \times 3 \times 14}{36} \approx 2.33$$

Nonempty blocks

Nonzeros

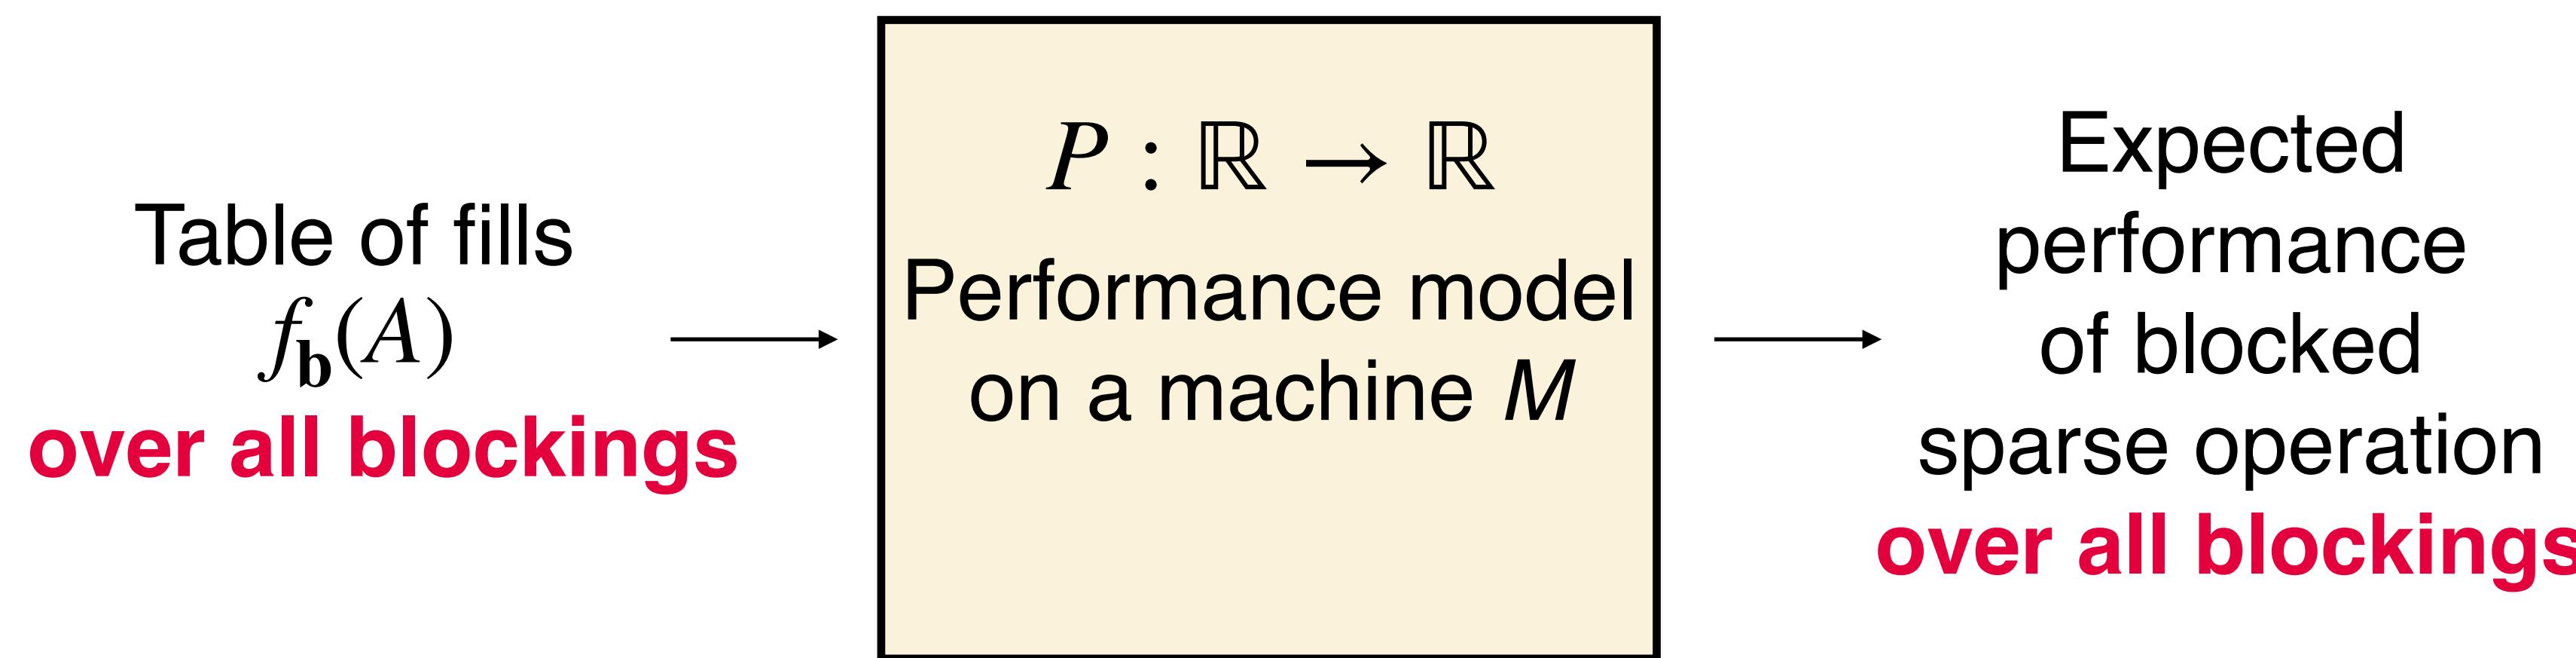


**Blocking  $\mathbf{b}$**

$b_1 = 2$
$b_2 = 3$
$k_{\mathbf{b}} = 14$
$k(A) = 36$

# Performance Modeling Using the Fill

A **performance model** is a function that maps the fill under a blocking  $\mathbf{b}$  to expected performance in FLOP/s.



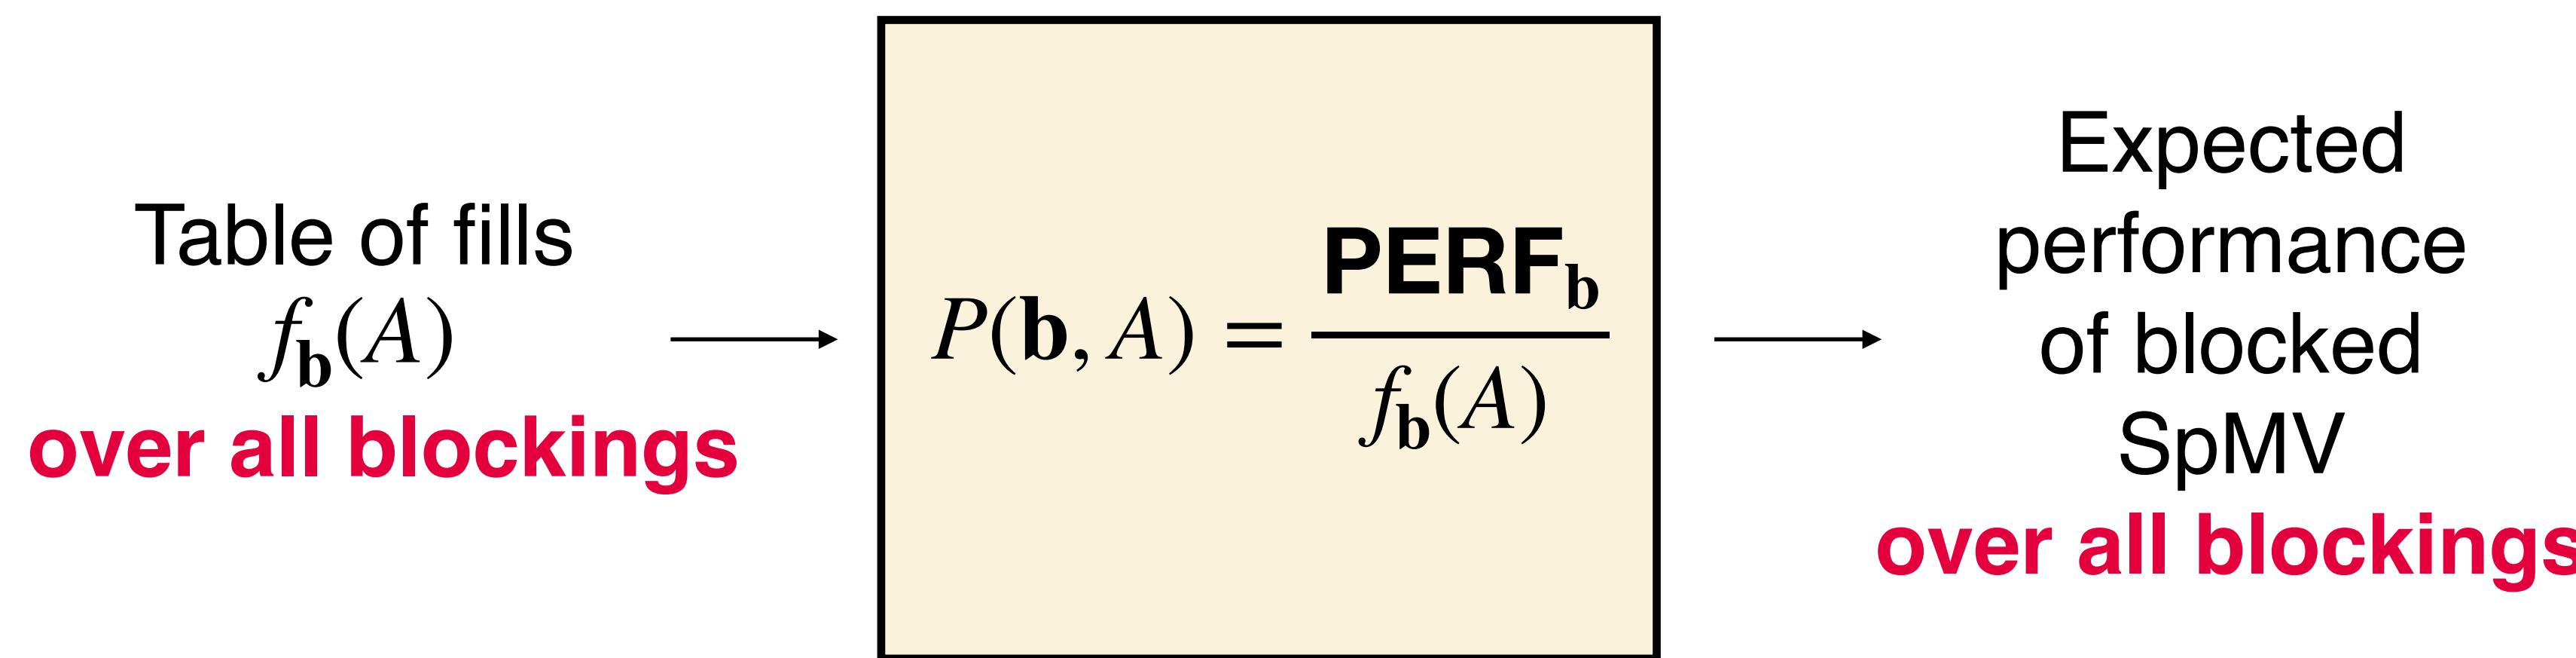
**Block size selection** chooses the block size that maximizes performance based on a performance model.

“Optimizing the Performance of Sparse Matrix-Vector Multiplication,” Im, 2000.

“Automatic Performance Tuning of Sparse Matrix Kernels,” Vuduc, 2003.

# Example: SPARSITY Performance Model for Blocked SpMV [Vuduc et al., 02]

The SPARSITY performance model uses a matrix **PERF(b)** of the performance of a machine  $M$  (in FLOP/s) on a dense matrix stored with blocking scheme  $b$ .

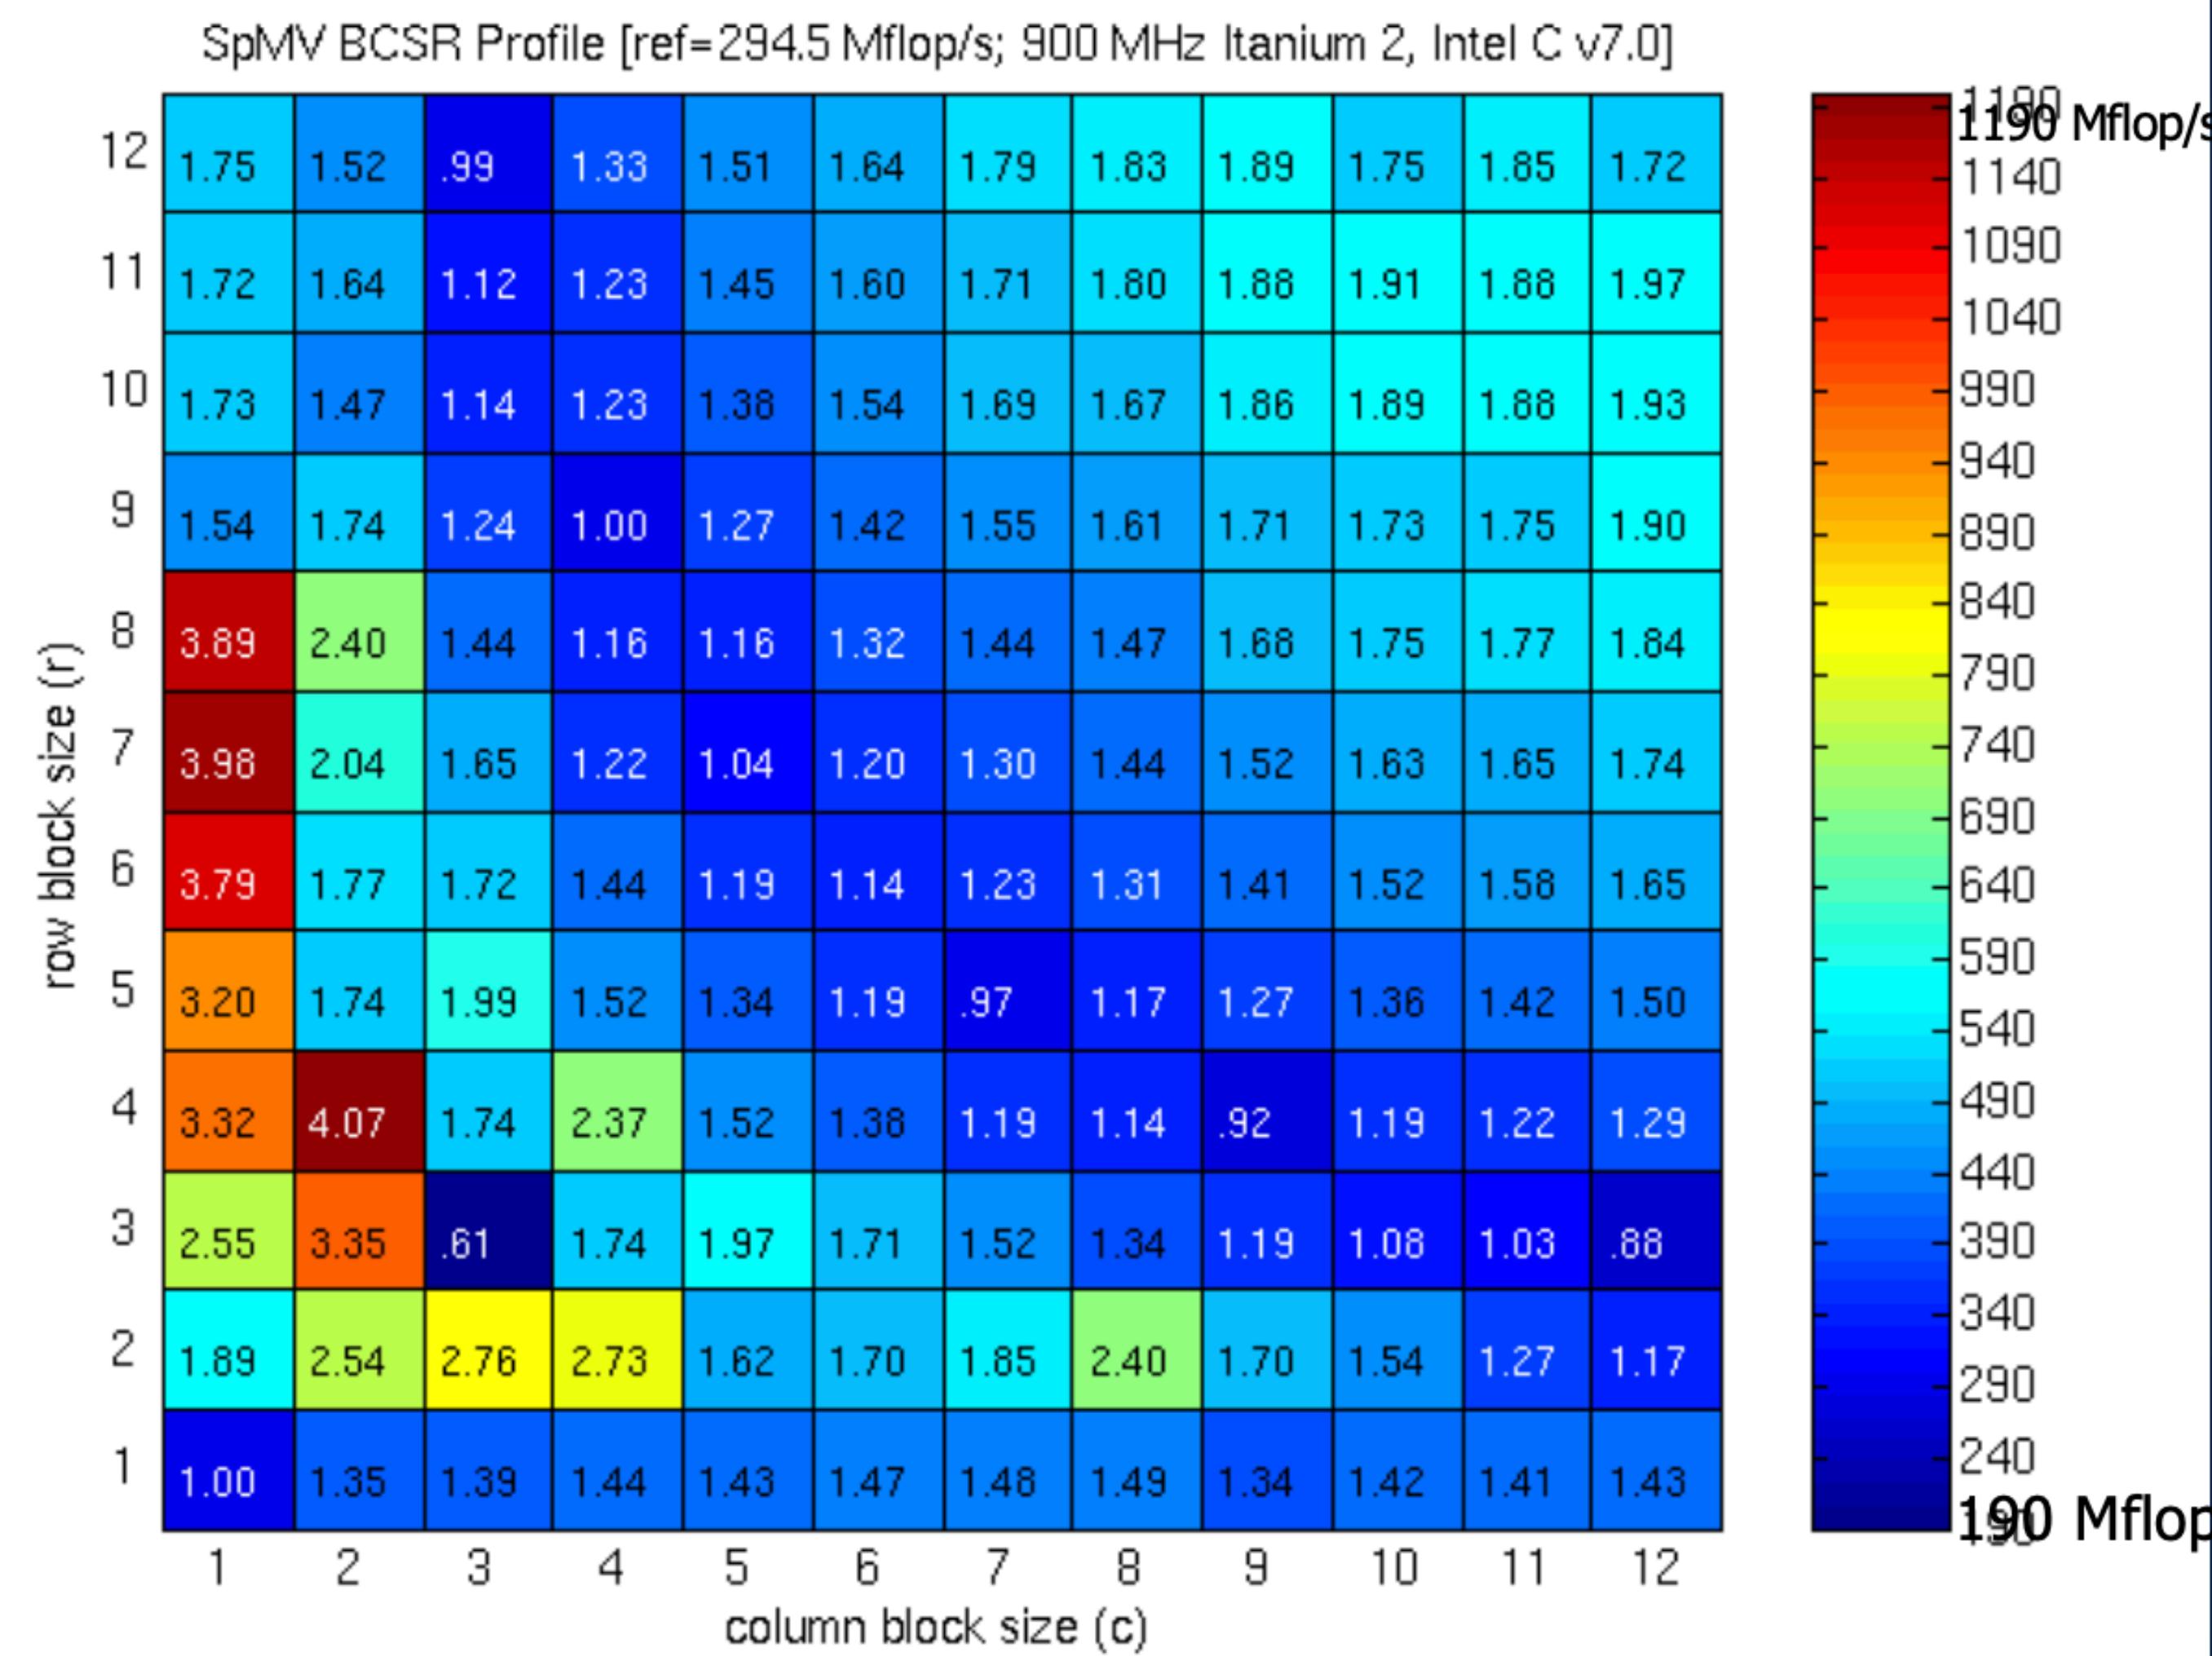


Vuduc et al. show that when the fill was known exactly, performance of the resulting blocking scheme was **optimal or near optimal** (within 5%) [V04]

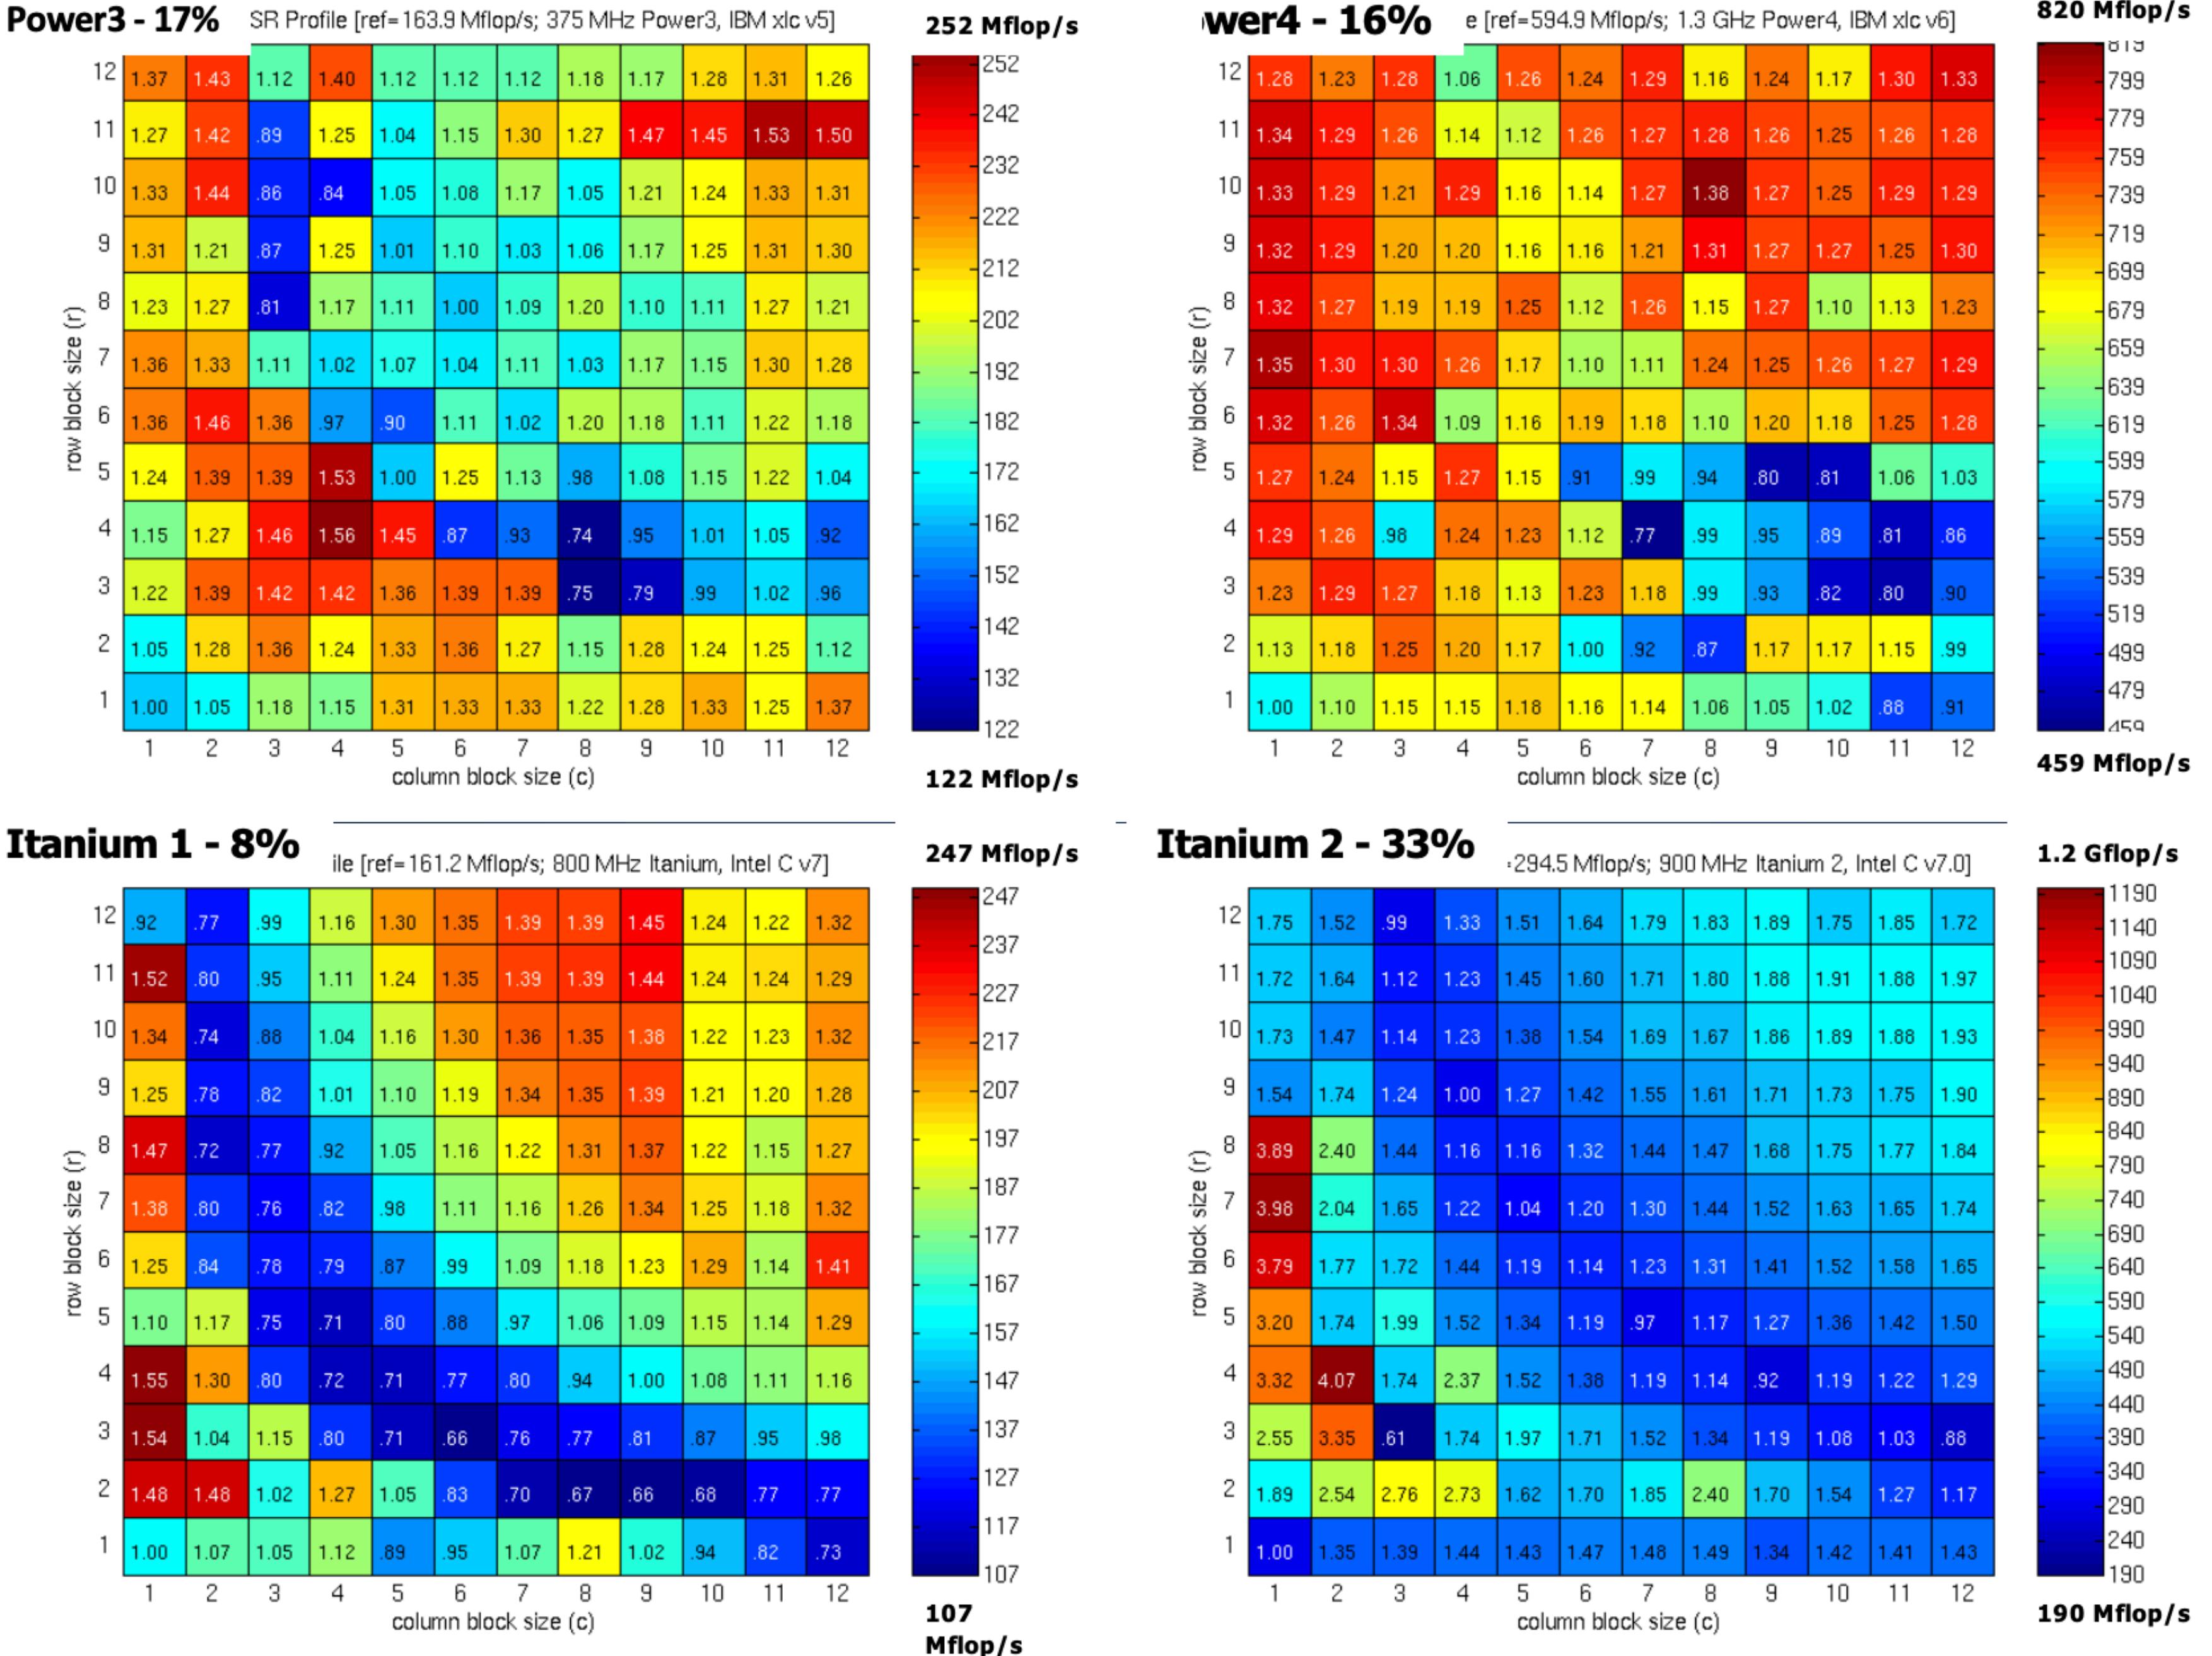
“Optimizing the Performance of Sparse Matrix-Vector Multiplication,” Im, 2000.

“Automatic Performance Tuning of Sparse Matrix Kernels,” Vuduc, 2003.

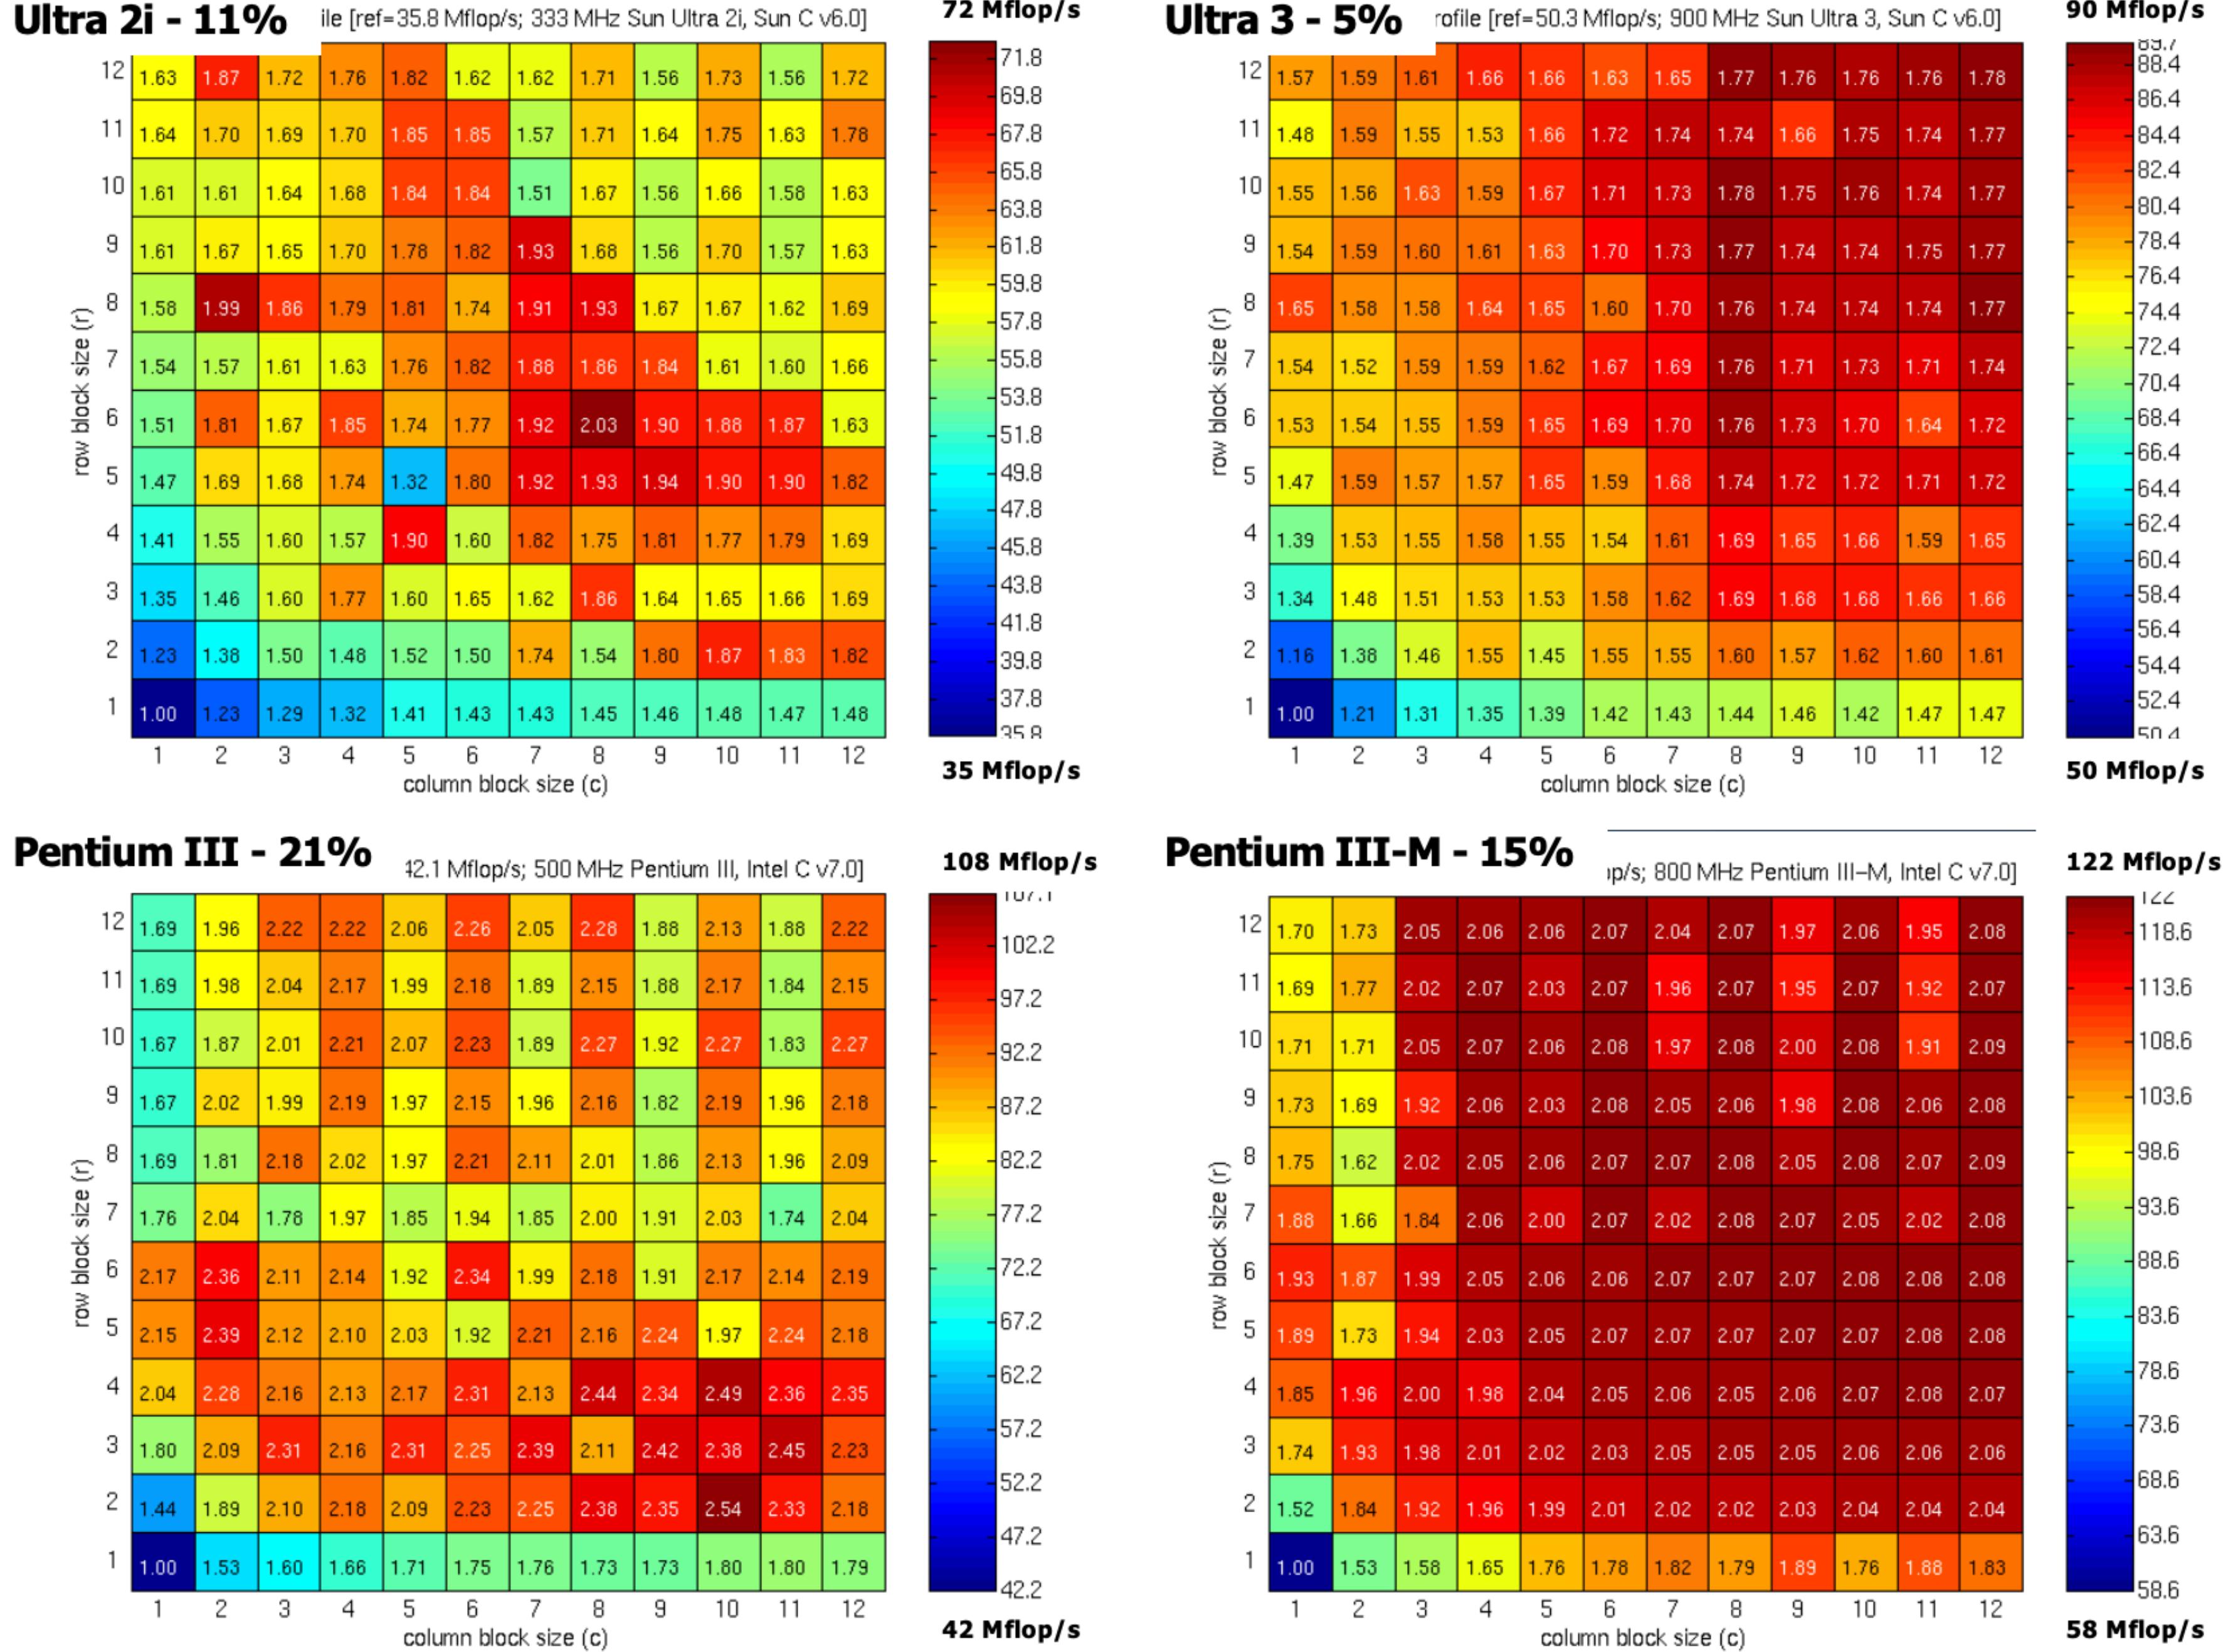
# Register Profile: dense matrix in sparse format



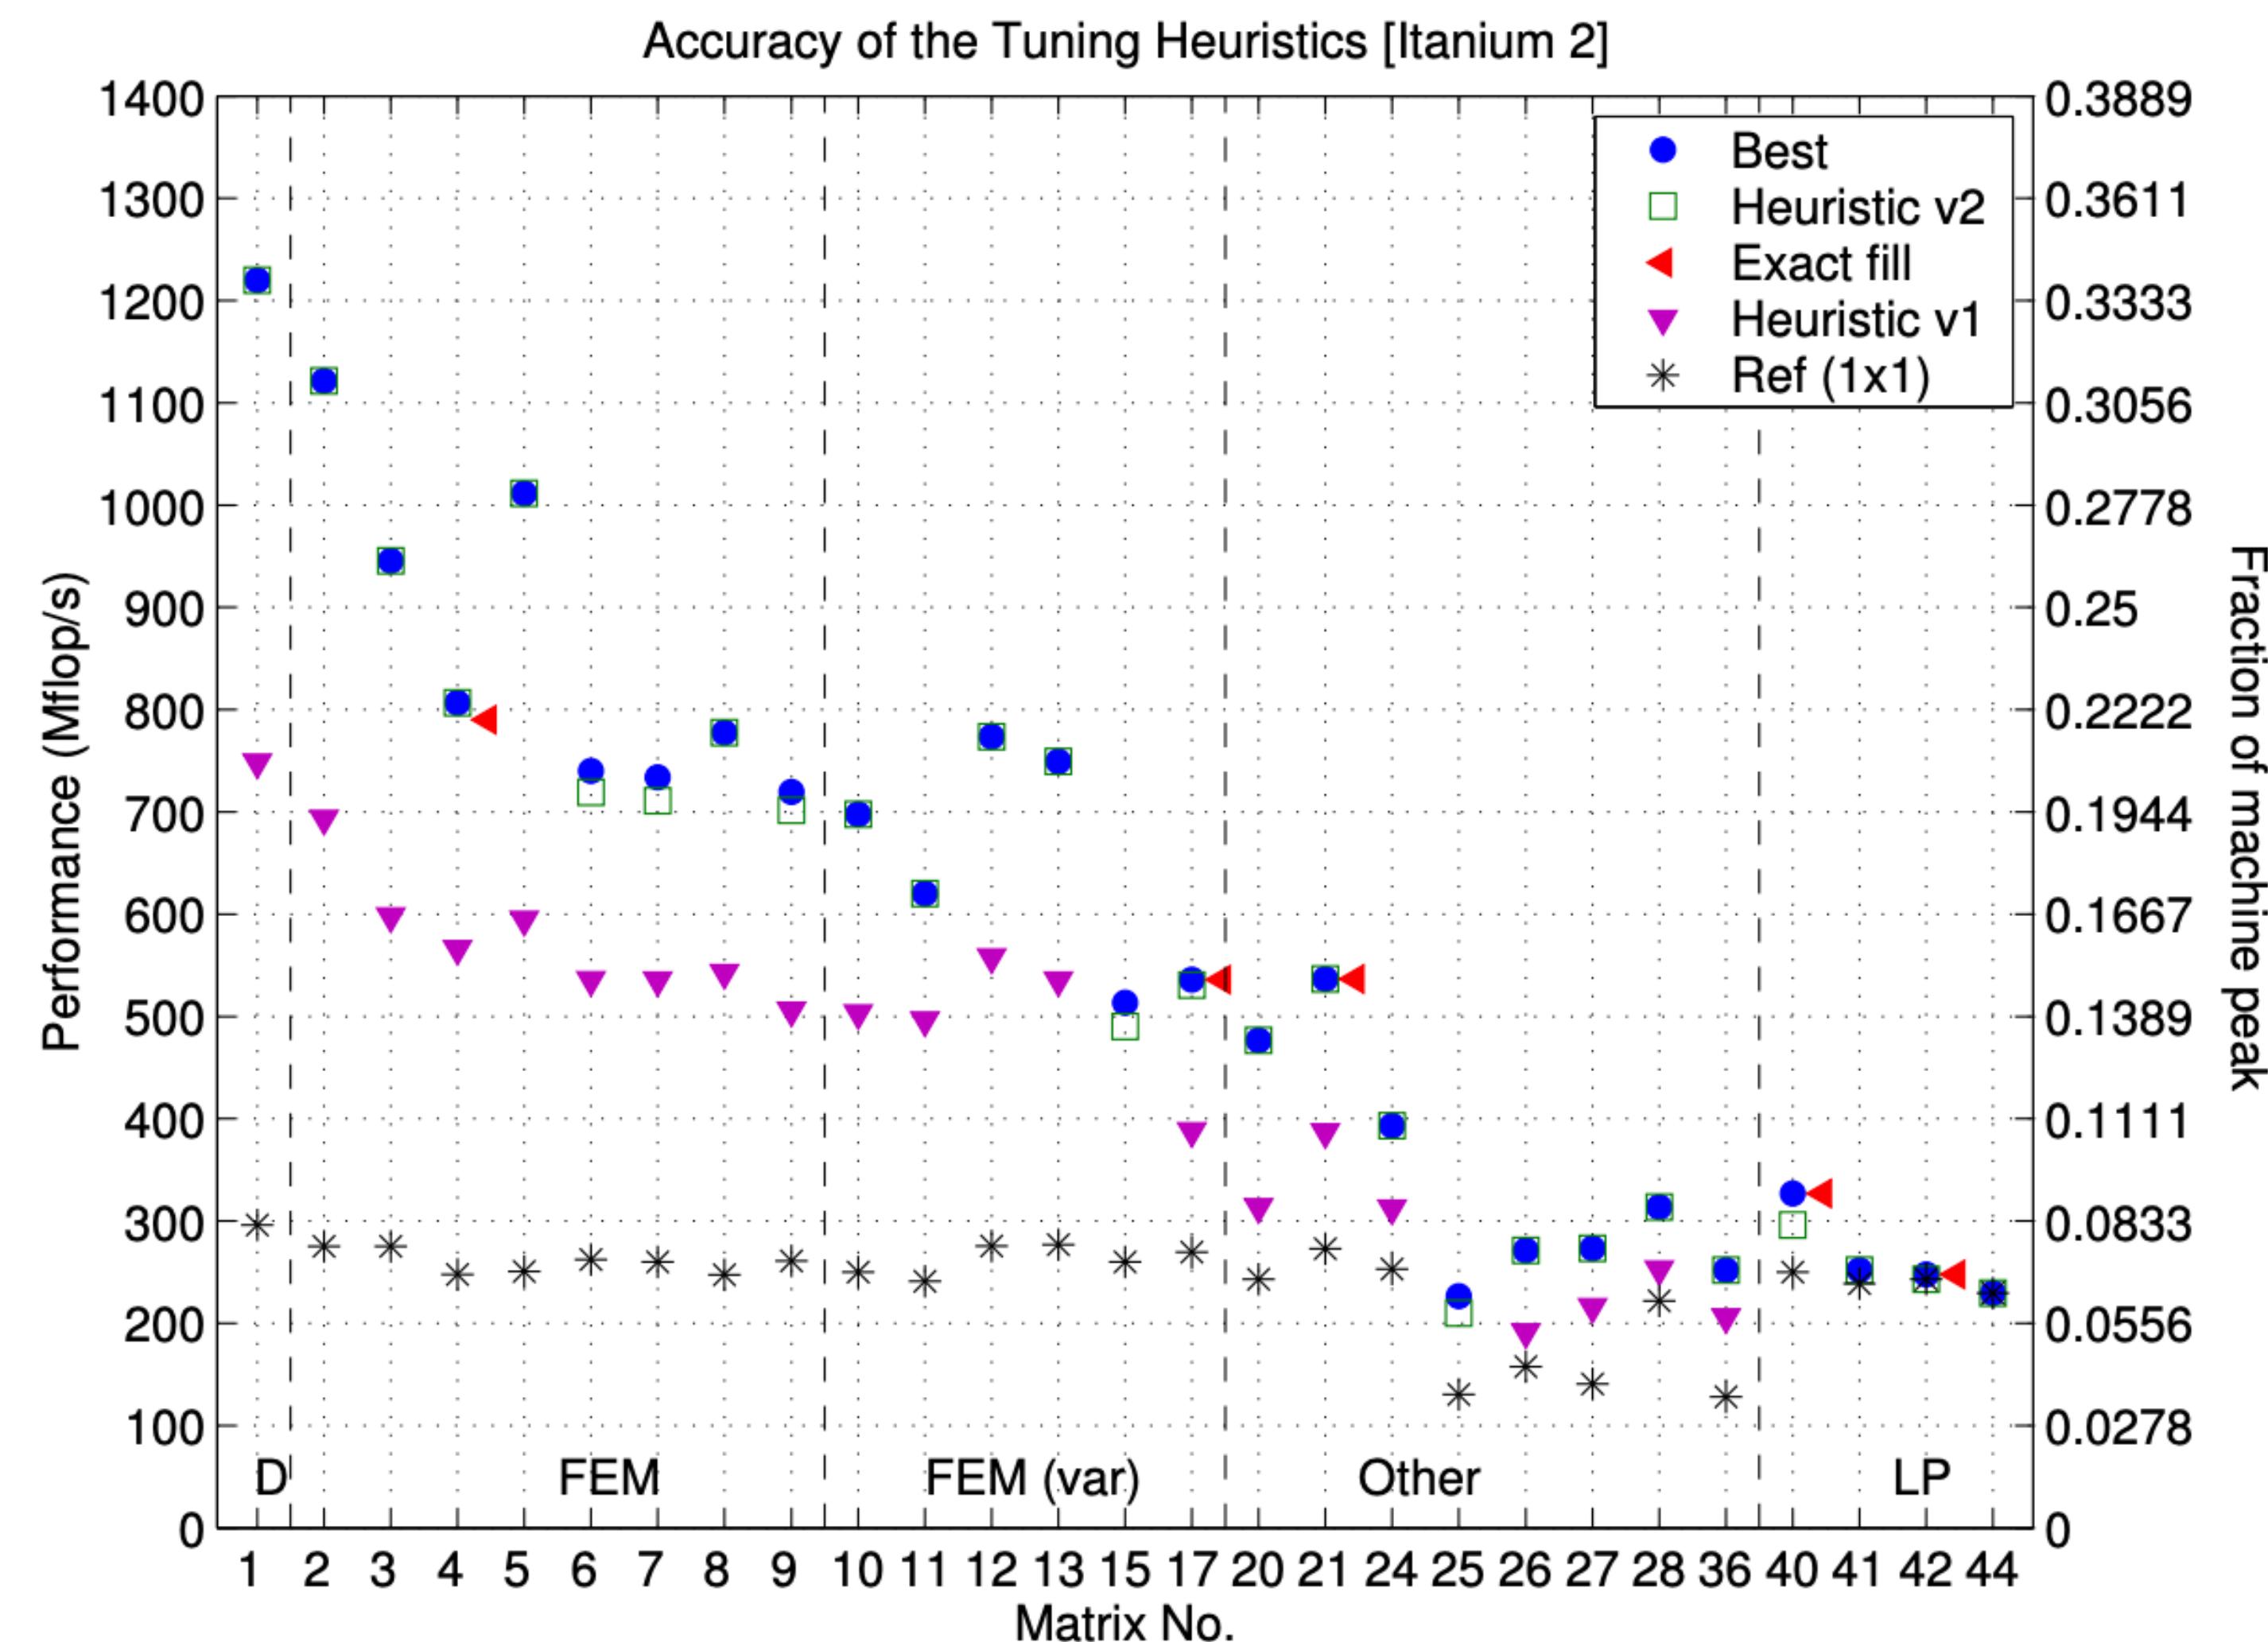
# Register Profiles



# Register Profiles

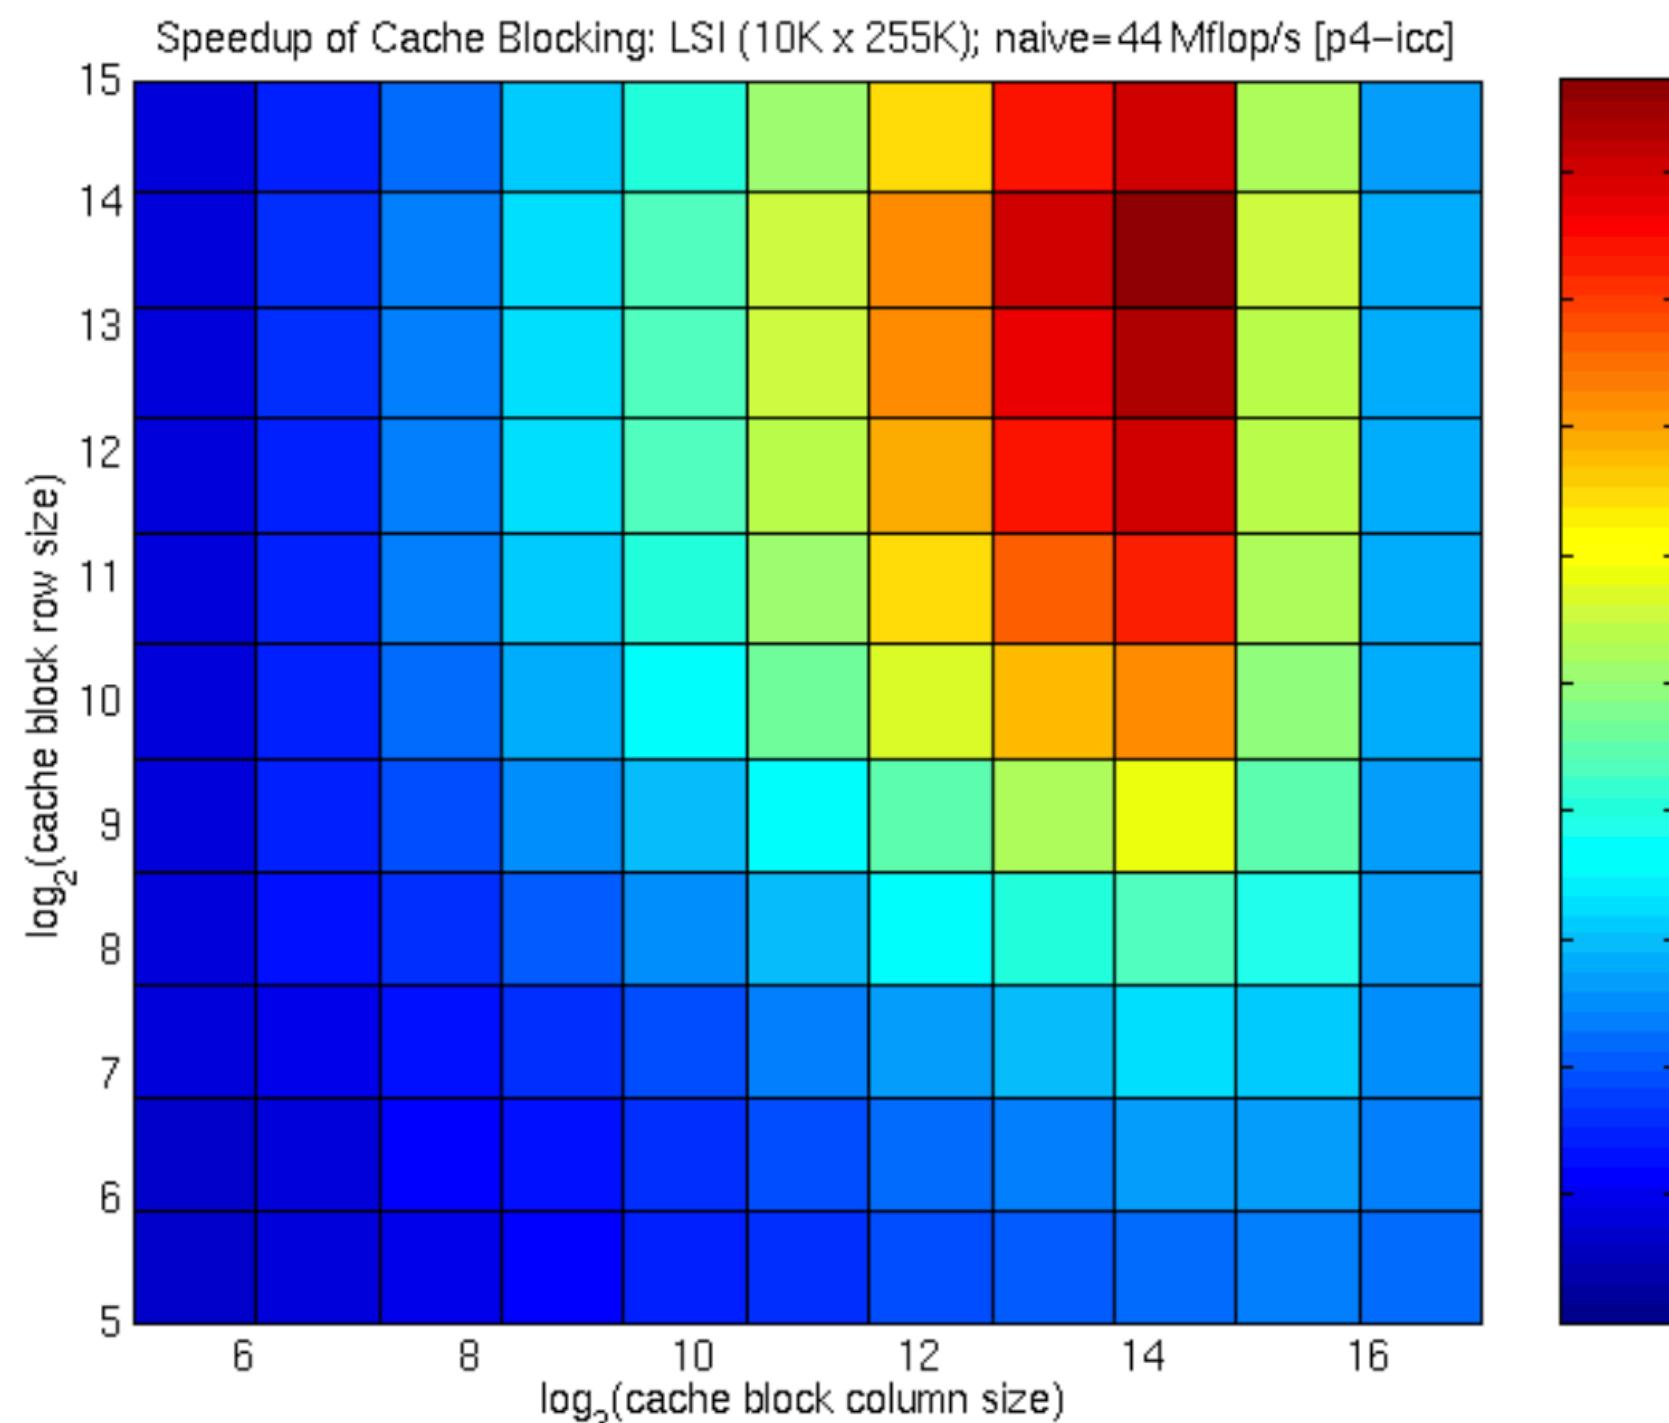


# Heuristic Tuning vs Best



# What about cache blocking?

- For CSR, dot-product, re-use opportunity is only in the x vector
- Matrices that are have some locality and “well ordered” e.g., near diagonal have good re-use
- Cache blocking can help for “short wide” matrices both on serial and SMPs



Matrix A

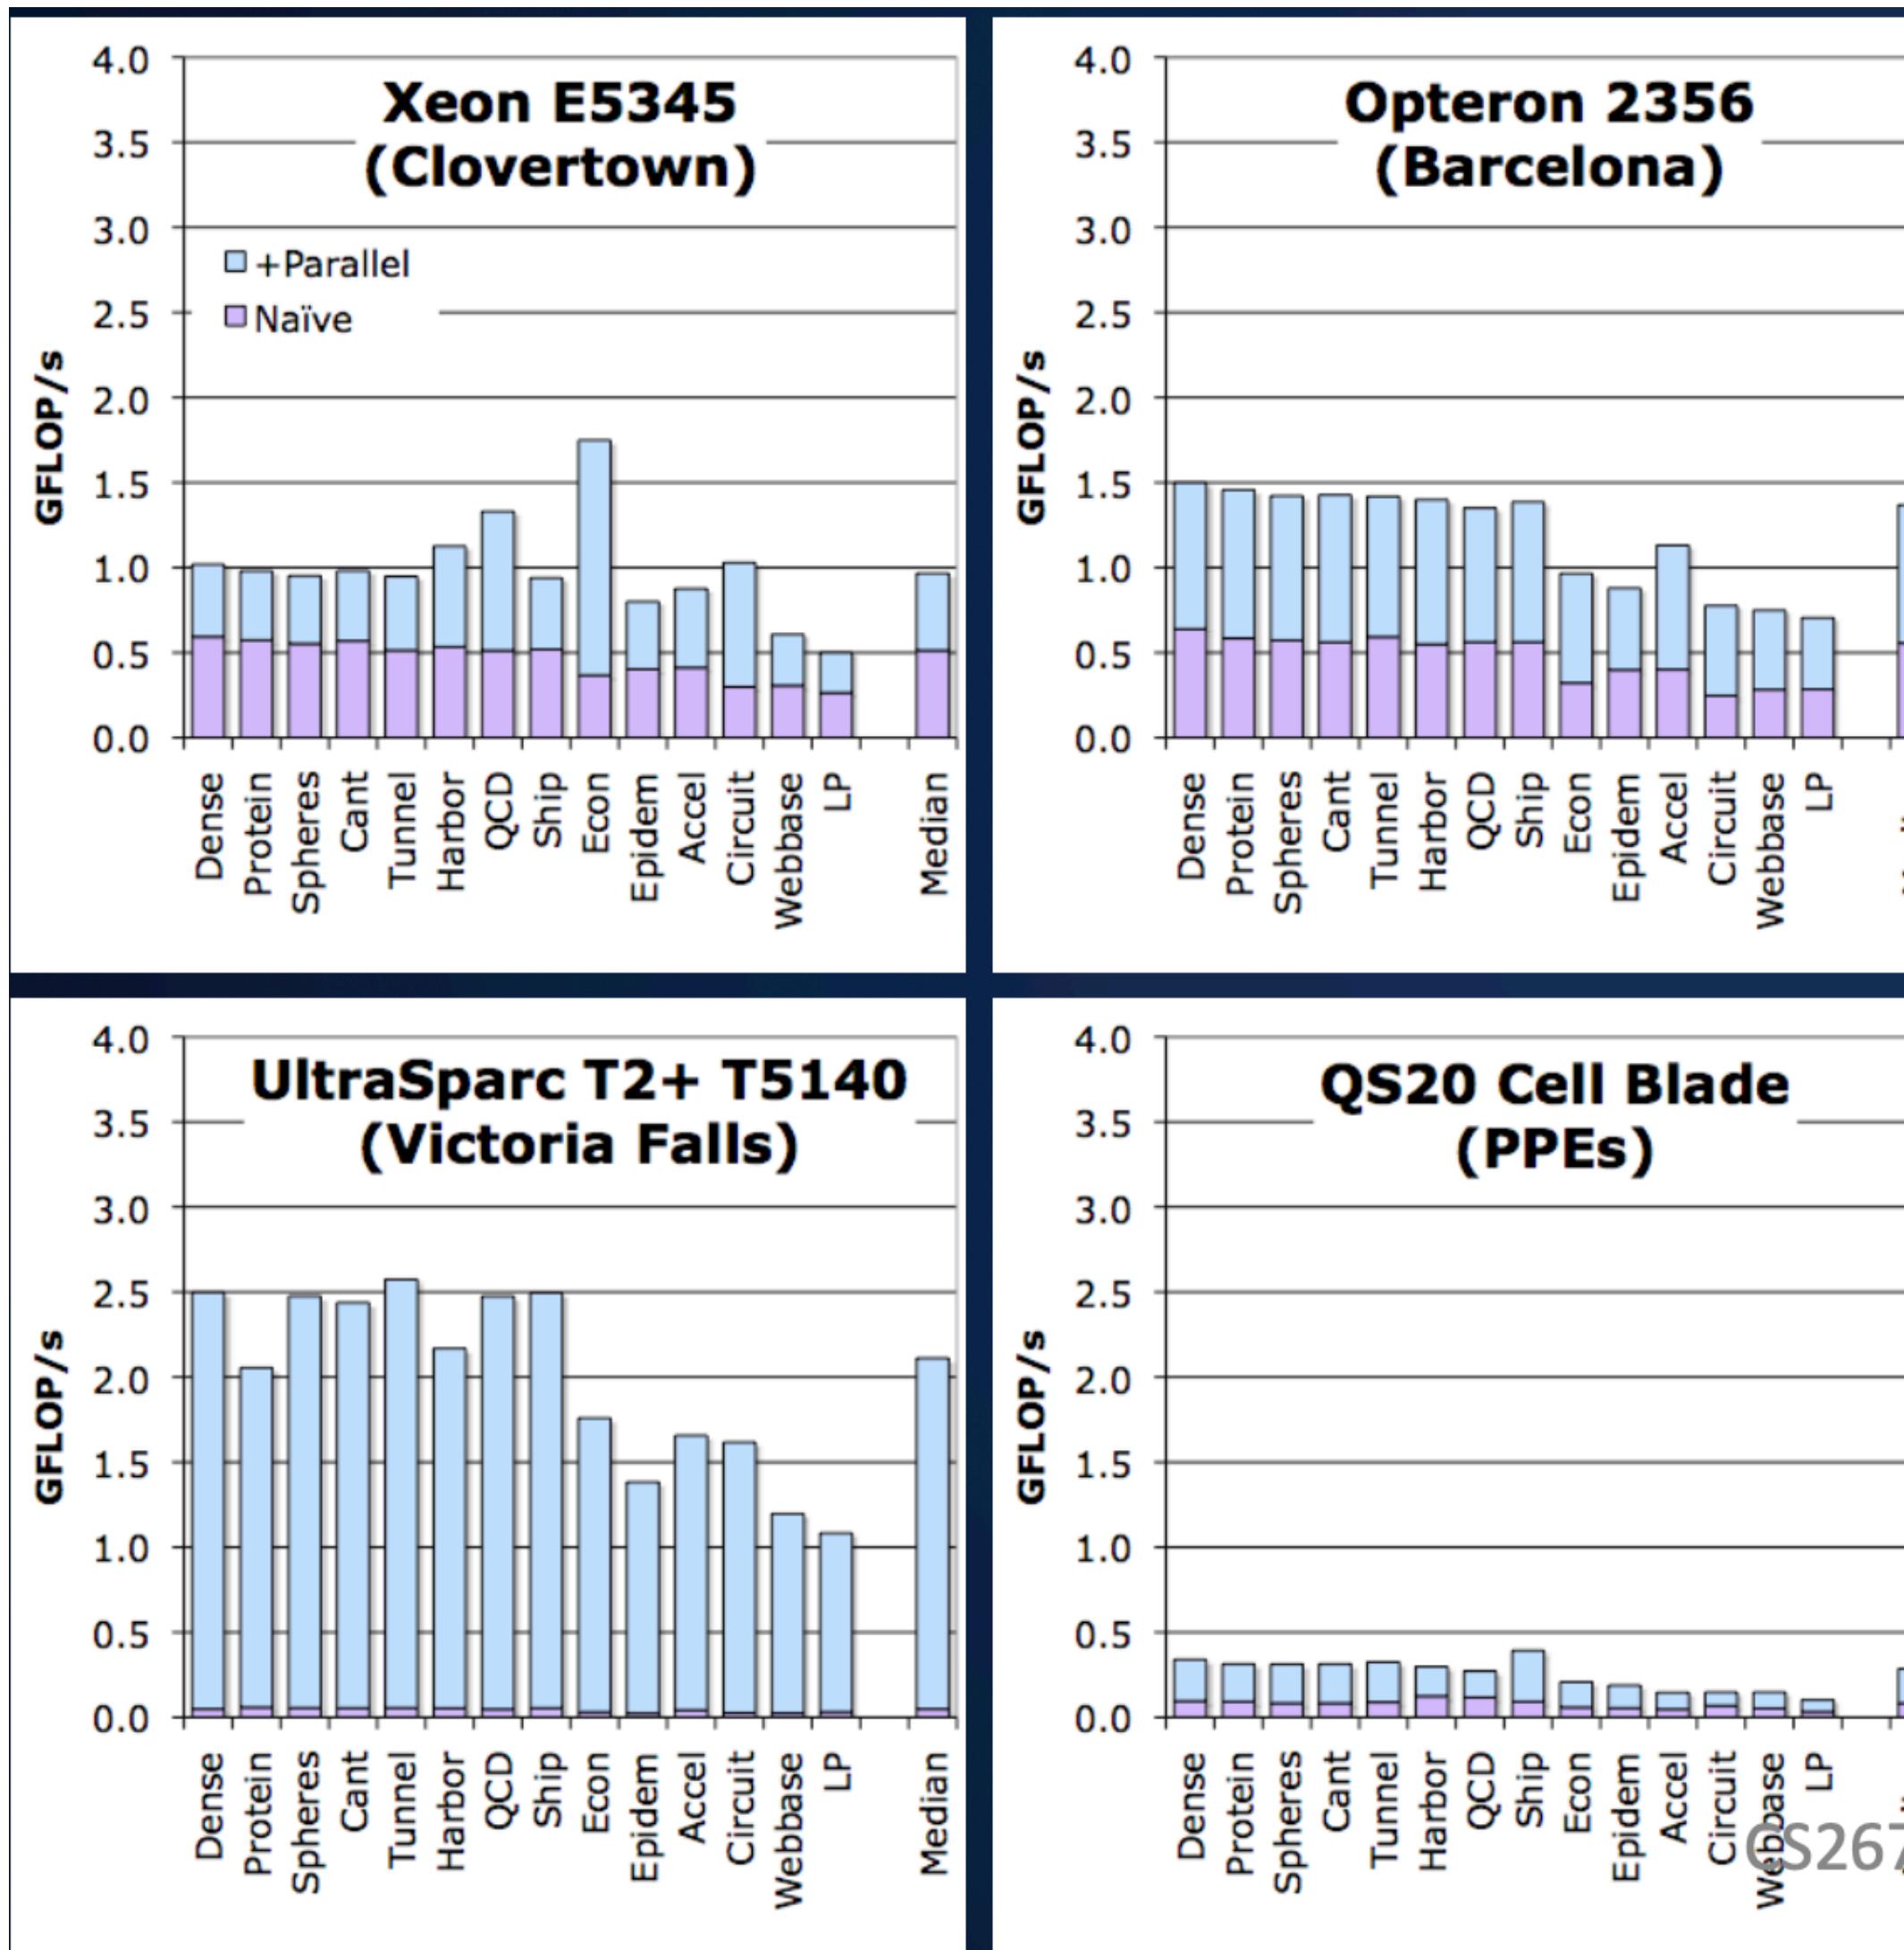
- 100k x 255k
- 3.7M non-zeros

Baseline: 44 Mflop/s

Best block size & performance:

- 16k x 16k
- 210 Mflop/s

# SpMV Performance (simple parallelization)

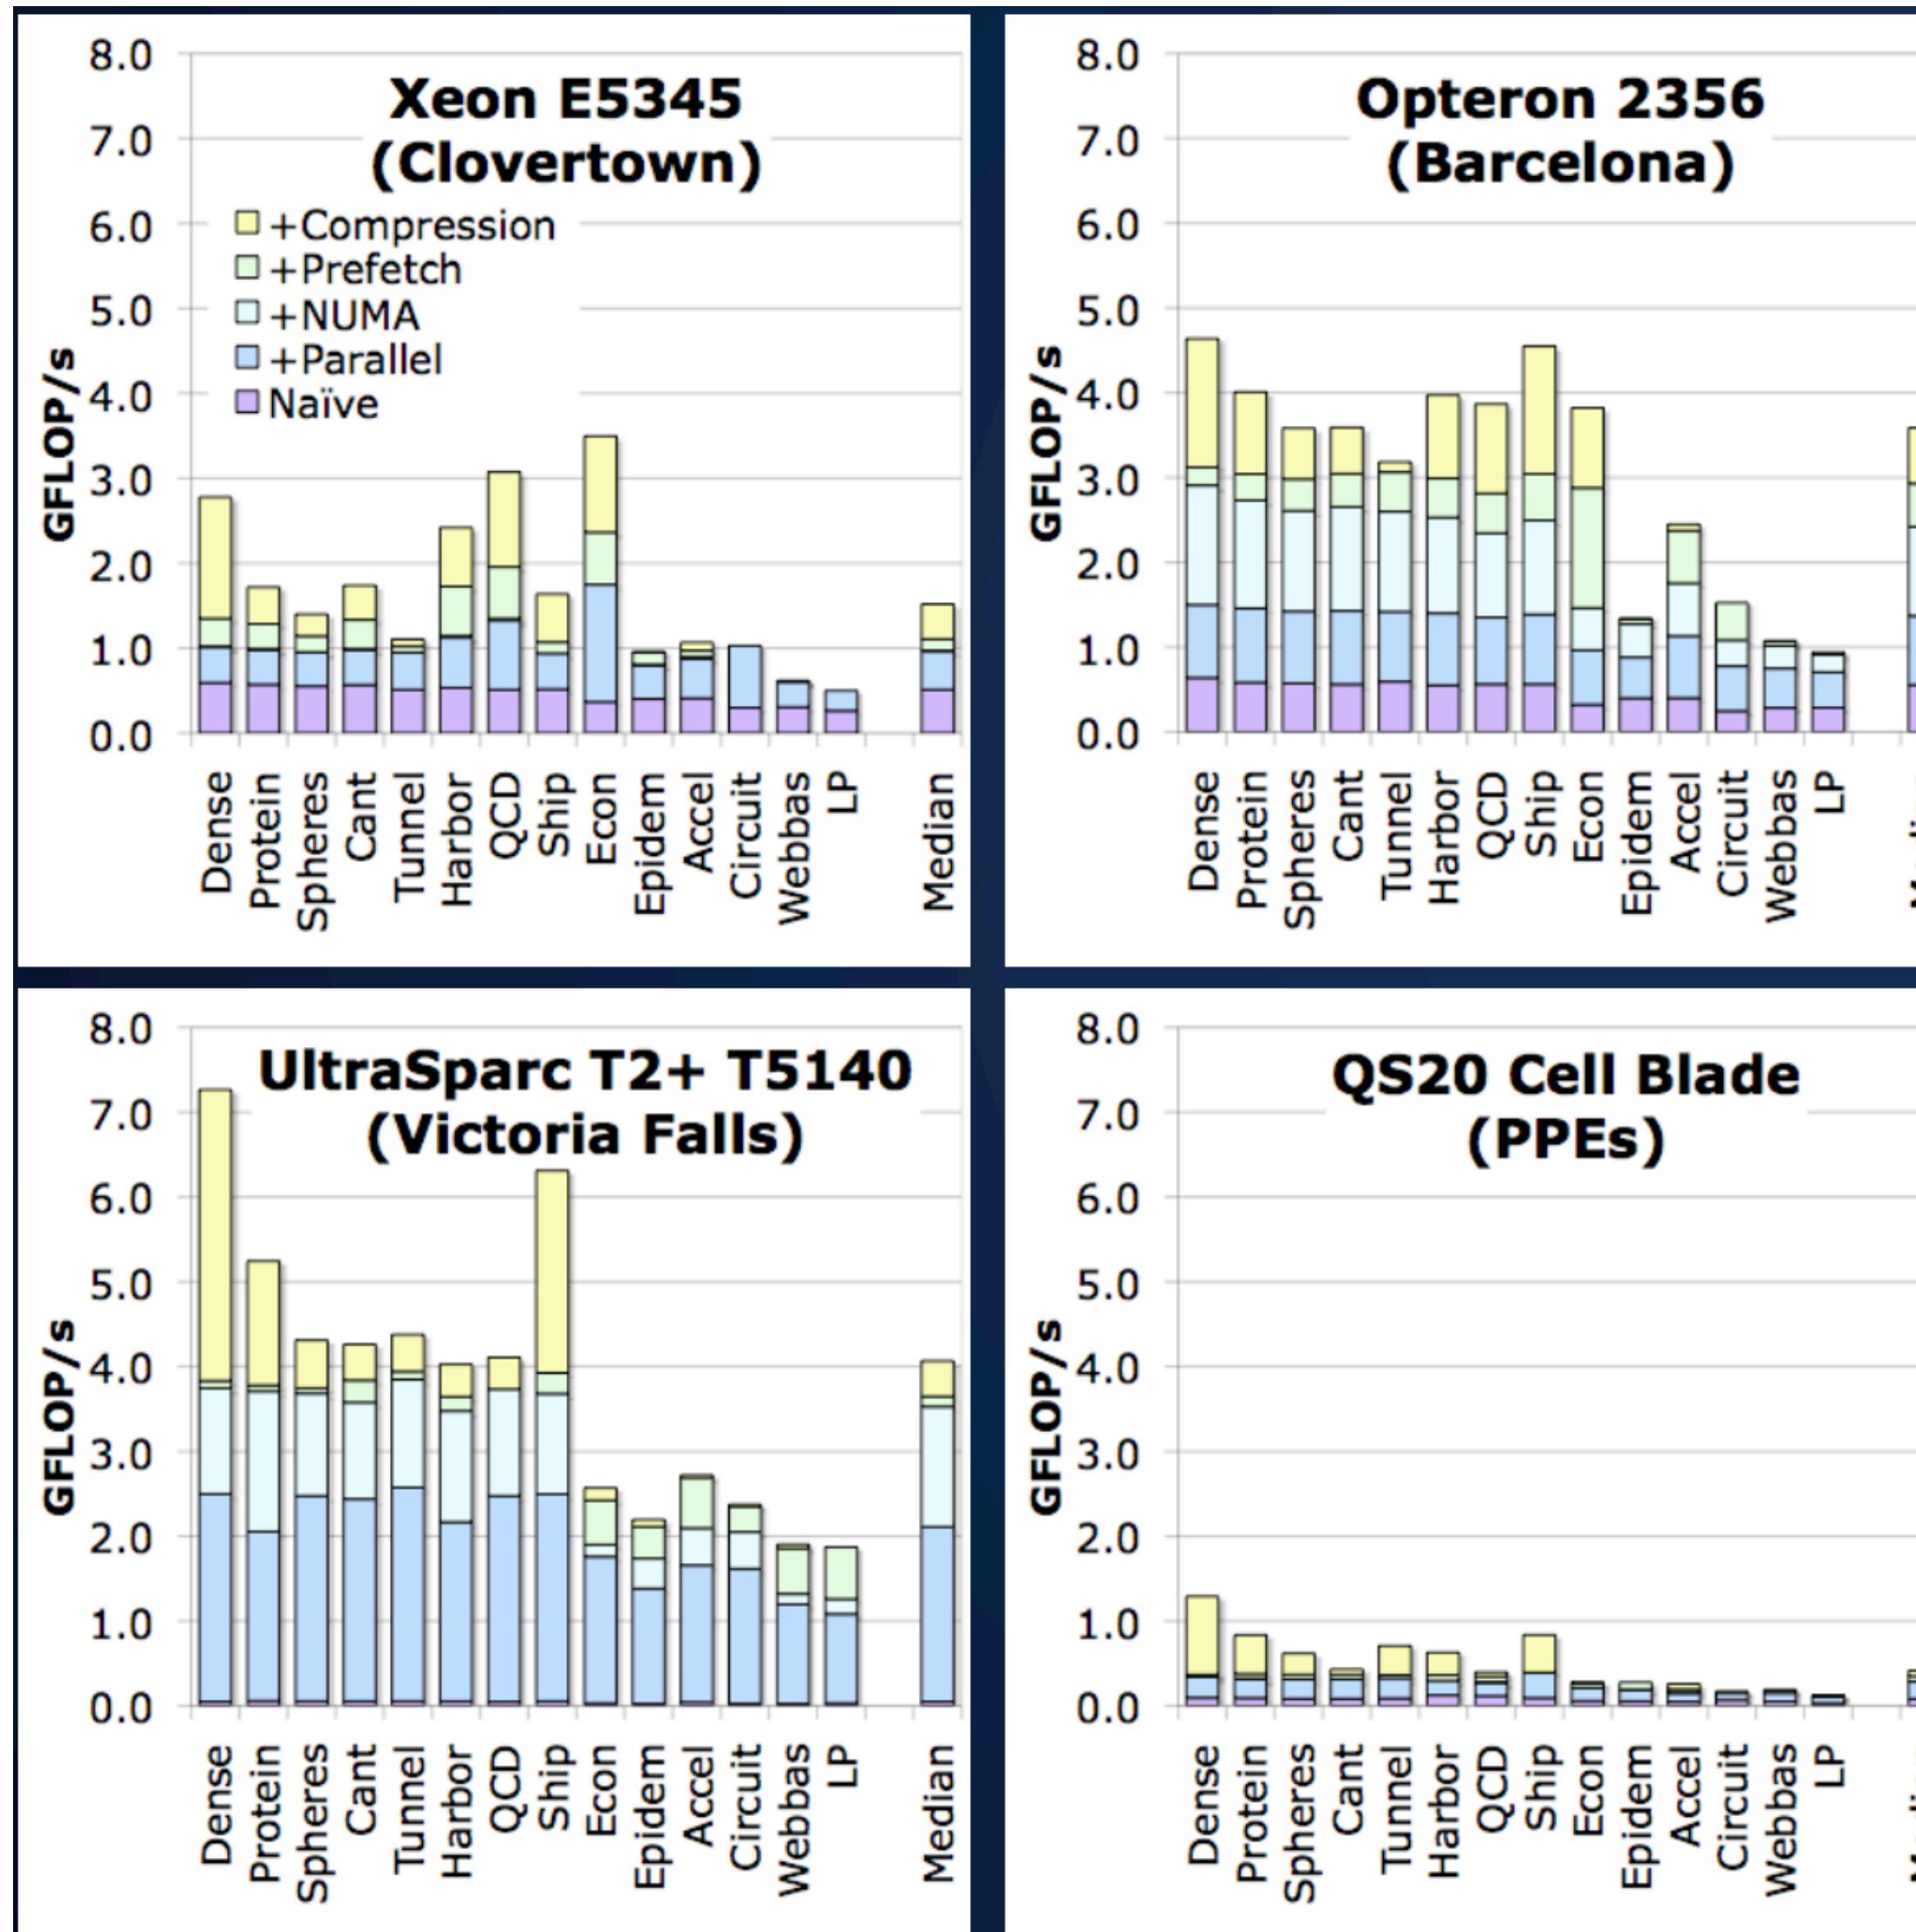


Out-of-the box SpMV performance  
on a suite of 14 matrices

Simplest solution = parallelization by  
rows

Scalability isn't great - Can we do  
better?

# SpMV Performance (matrix compression)



After maximizing memory bandwidth, the only hope is to minimize memory traffic.

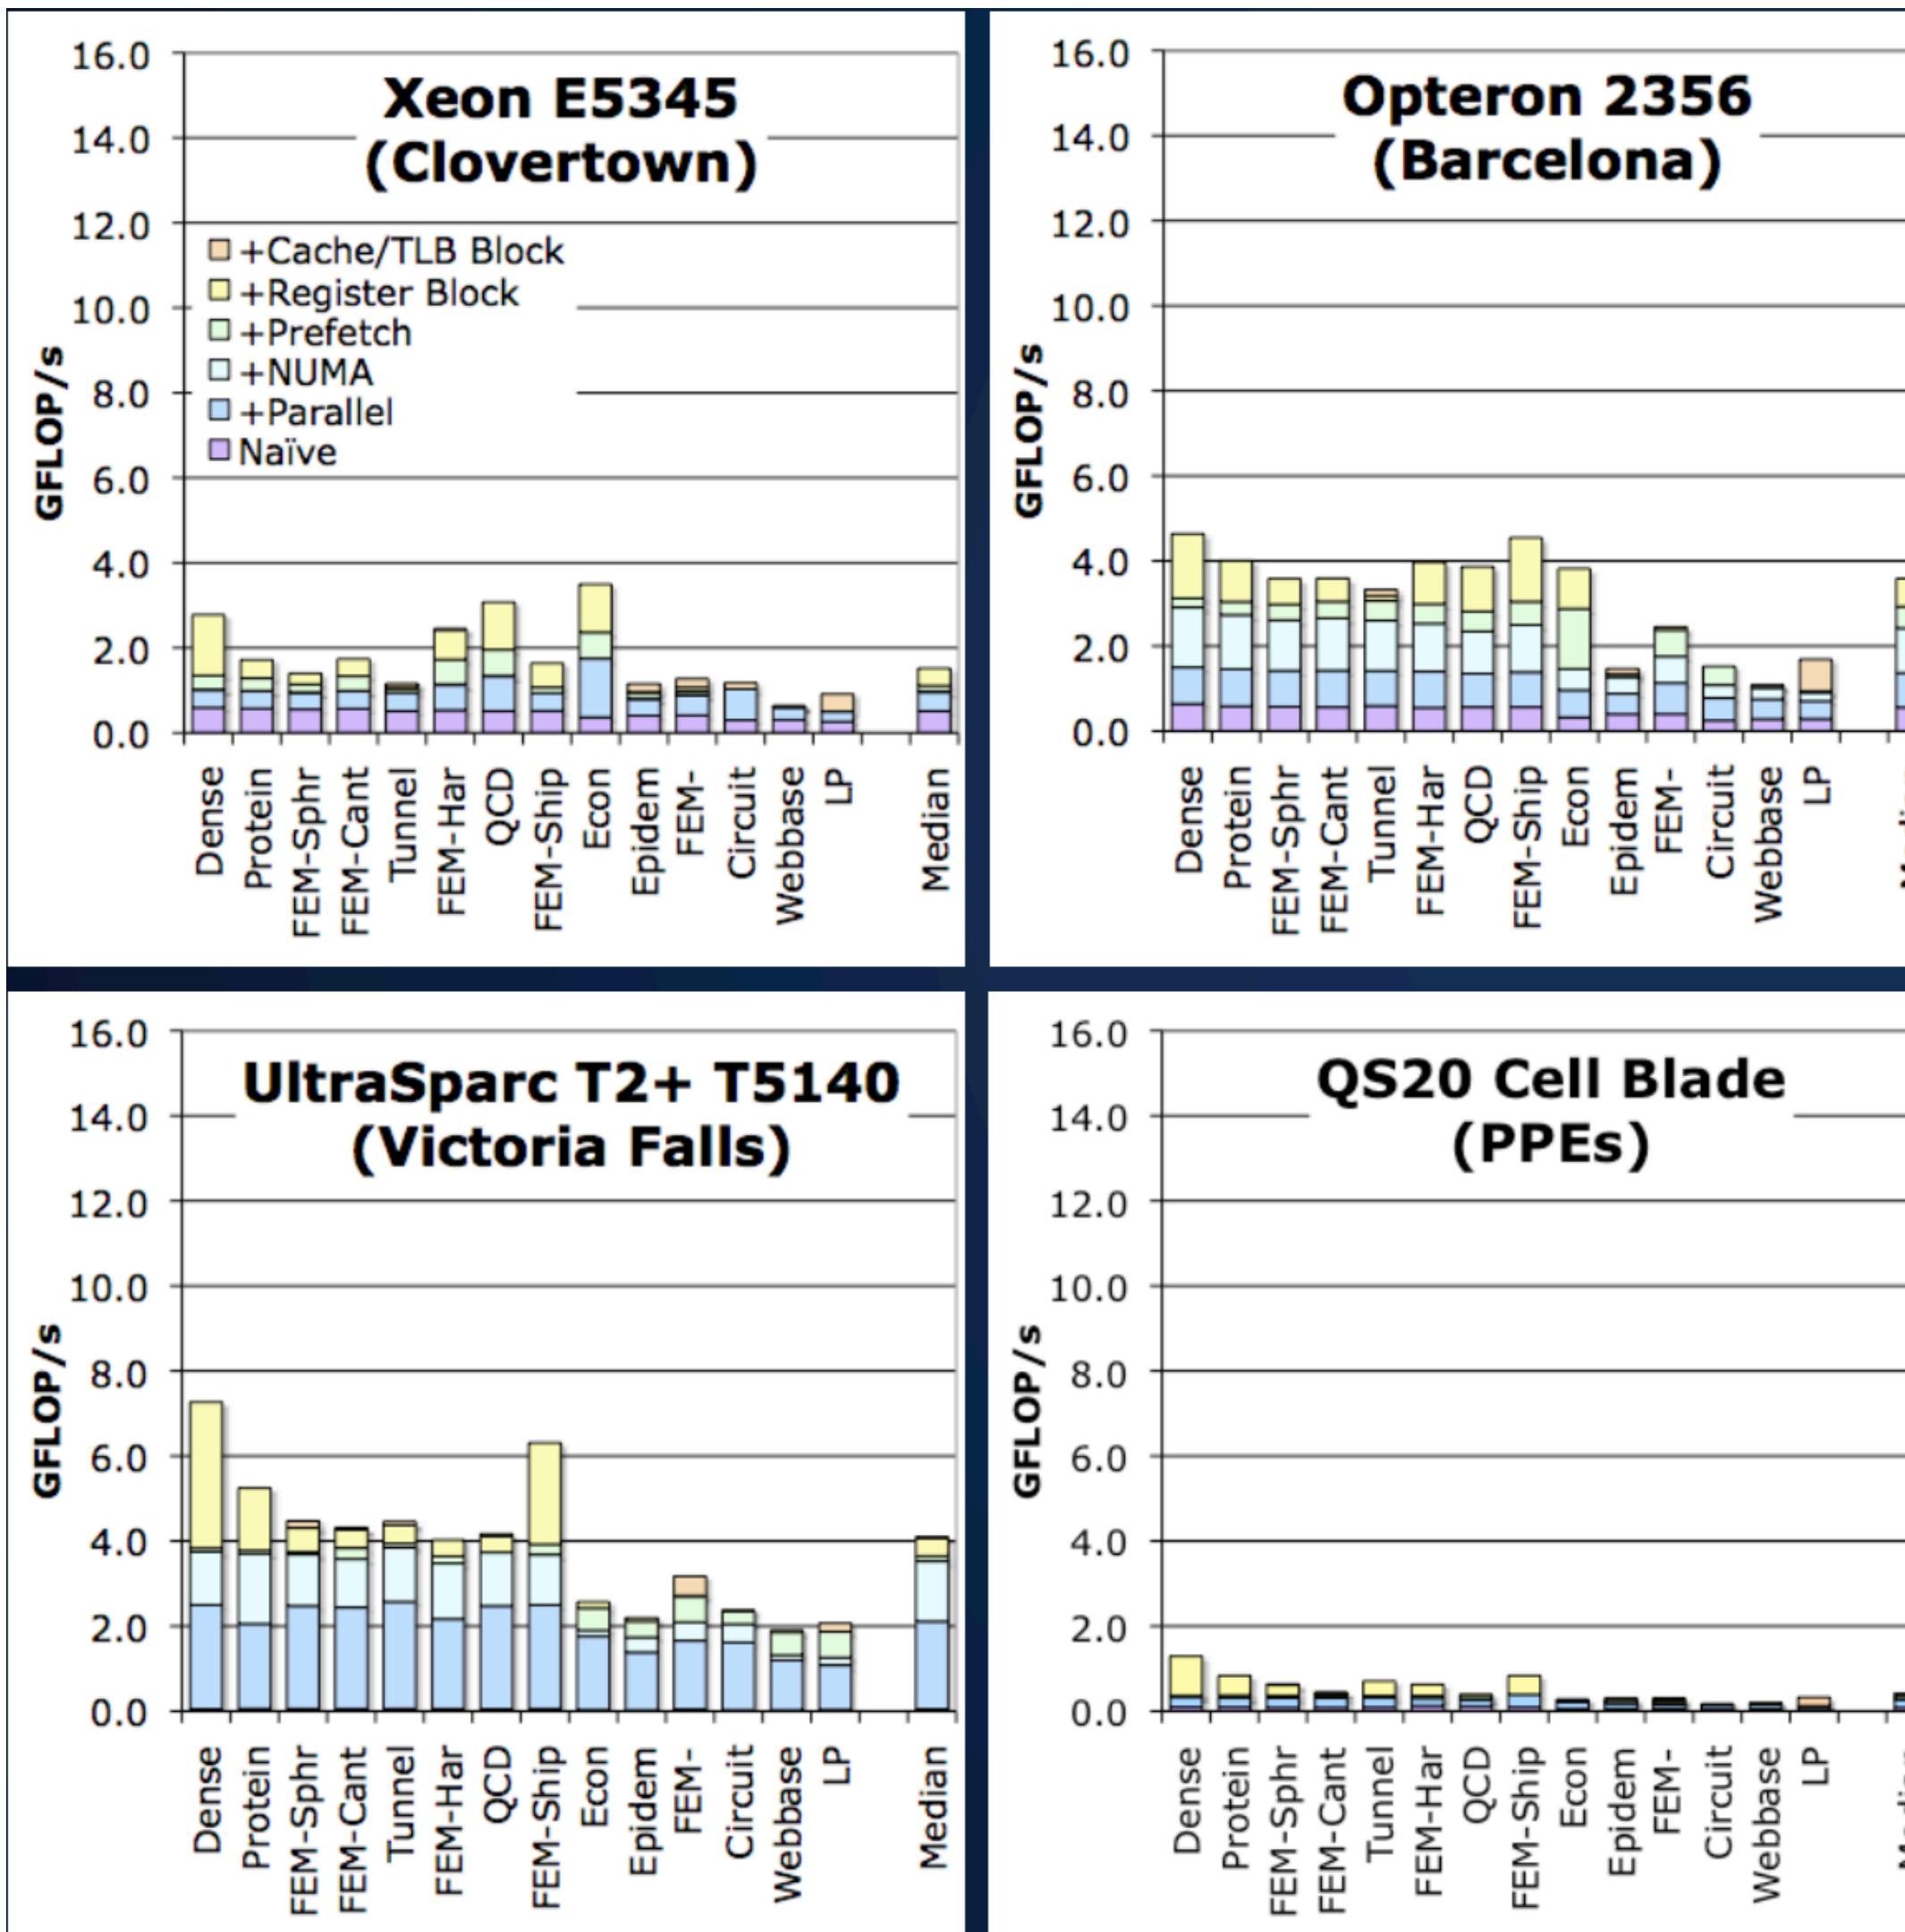
Compression: exploit

- register blocking
- other formats
- smaller indices

Benefit is matrix-dependent.

Register blocking enables efficient software prefetching (one per cache line)

# SpMV Performance (cache and TLB blocking)

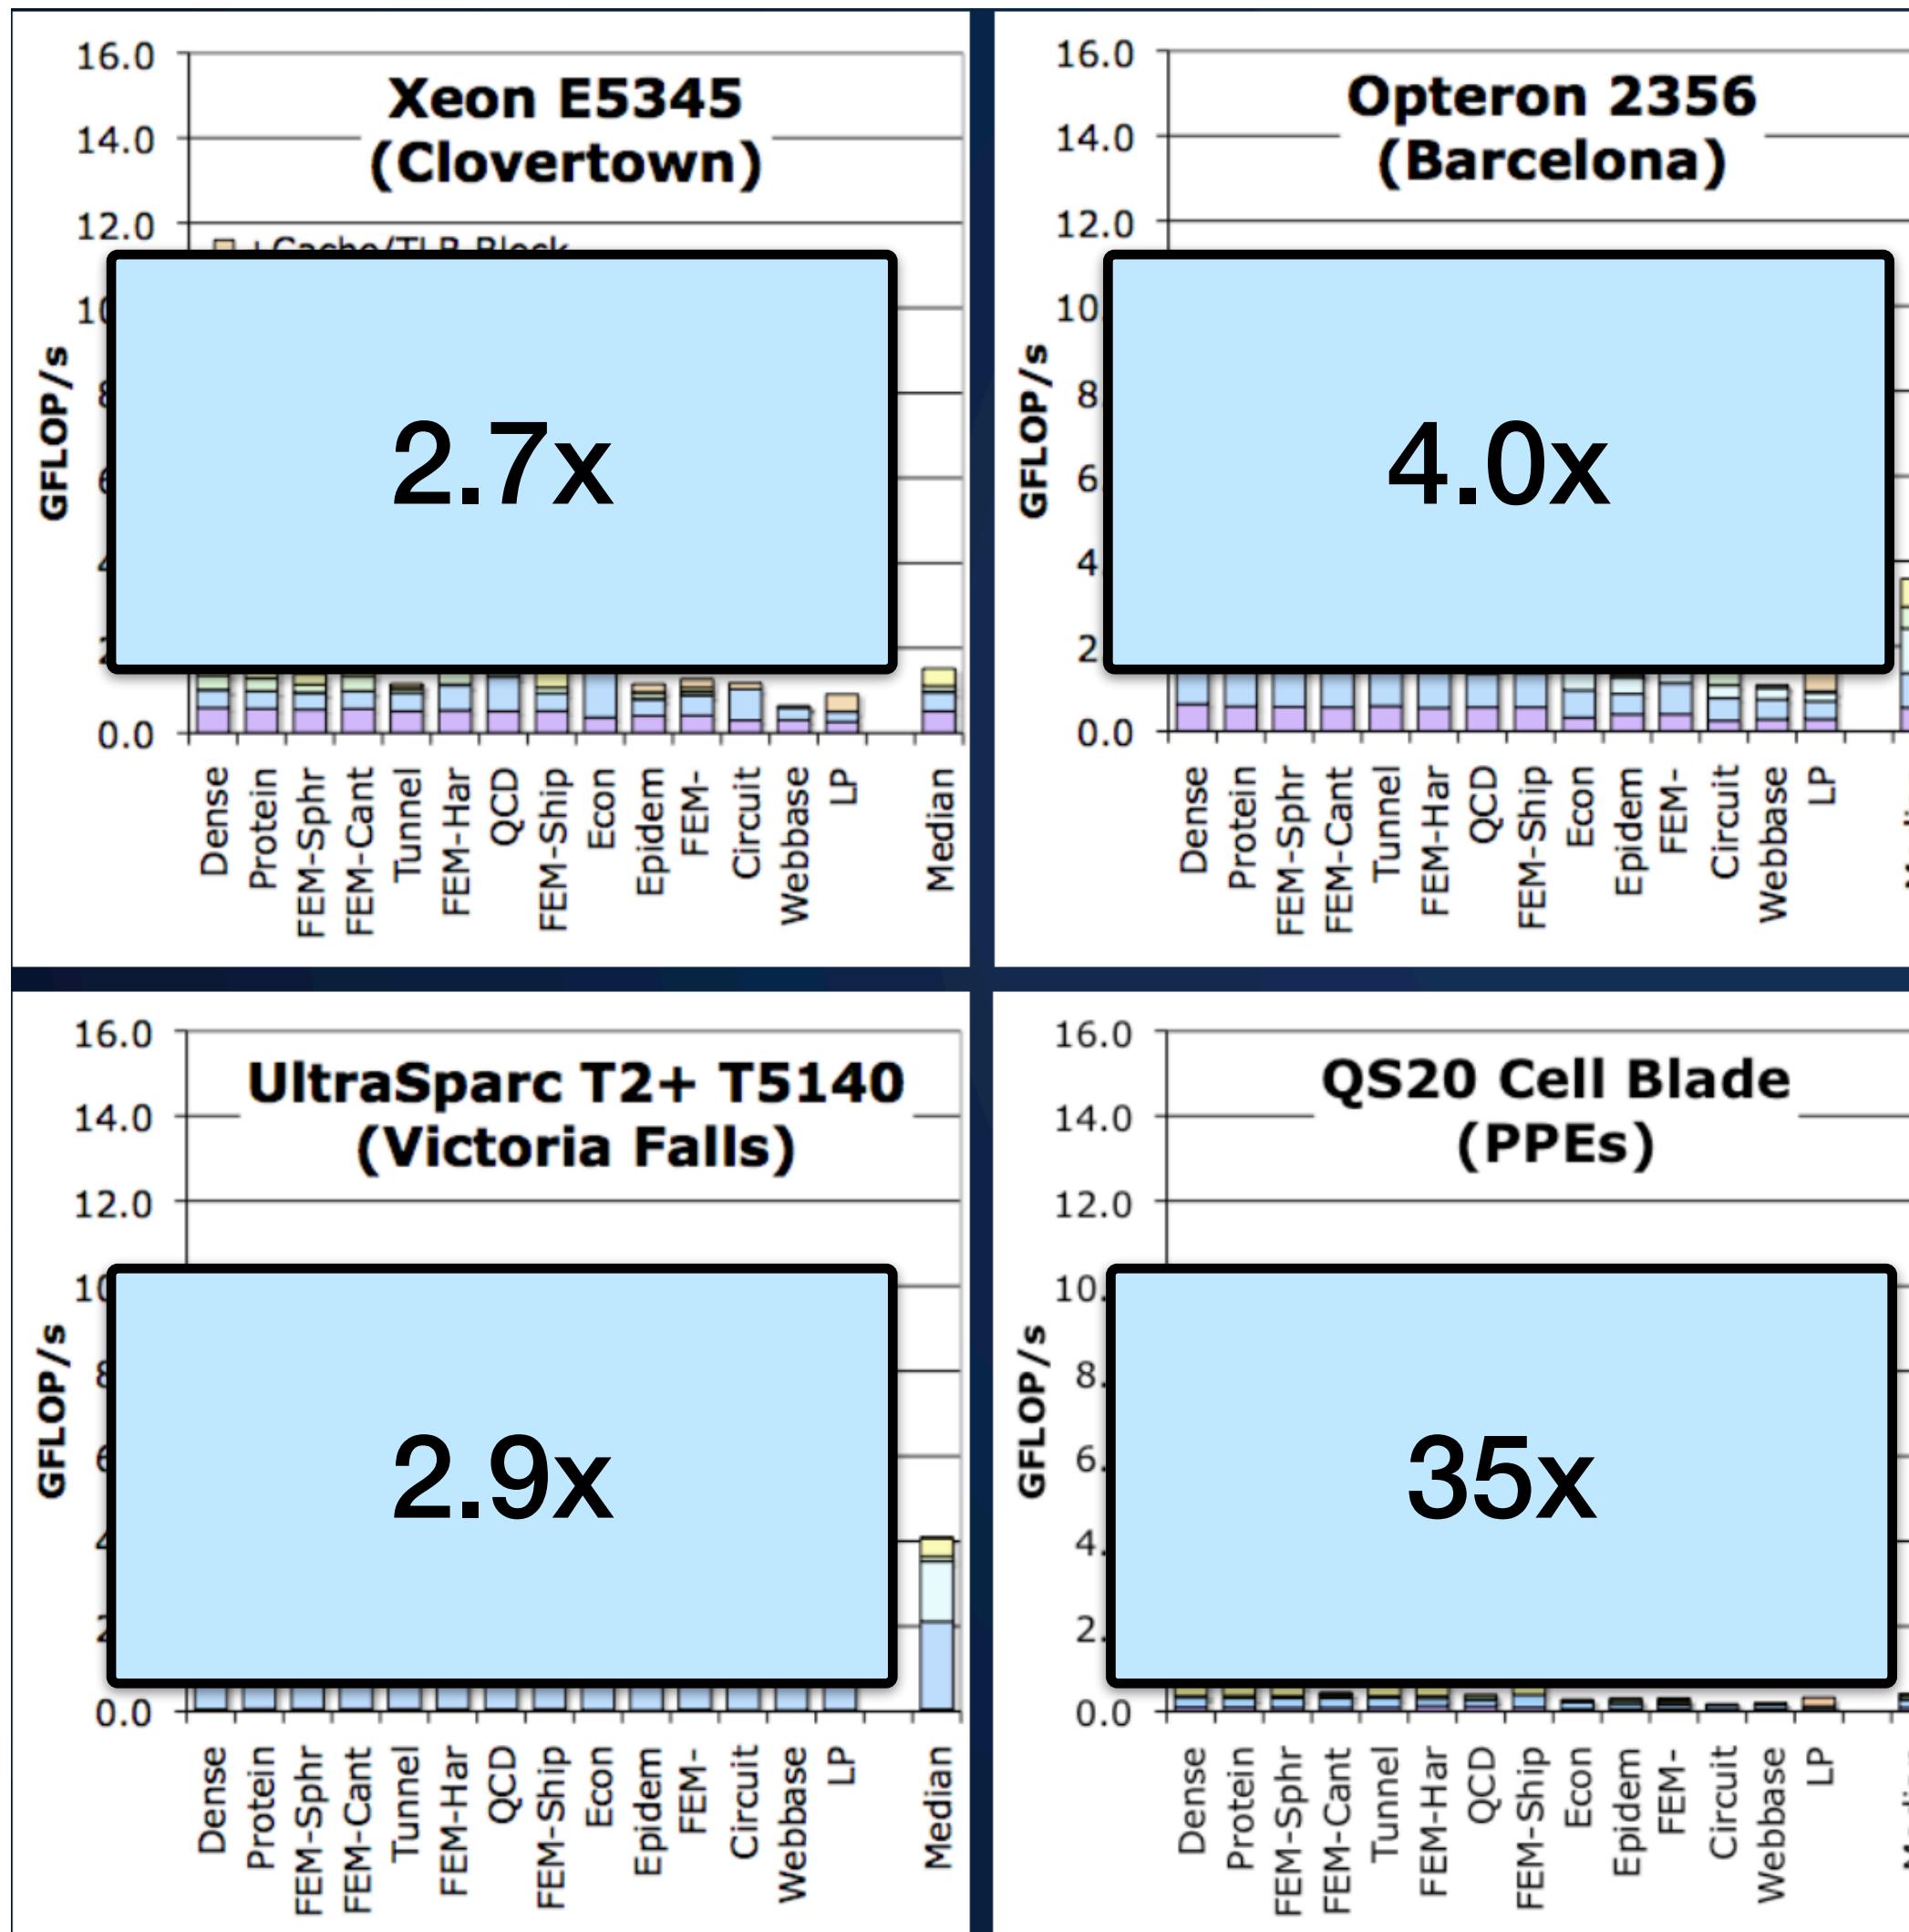


Fully auto-tuned SpMV performance across the suite of matrices

Why do some optimizations work better on some architectures?

- Matrices with naturally small working sets
- Architectures with giant caches

# SpMV Performance (cache and TLB blocking)



Fully auto-tuned SpMV performance across the suite of matrices

Why do some optimizations work better on some architectures?

- Matrices with naturally small working sets
- Architectures with giant caches

# Summary

- Tuning for sparse matrices: harder than dense ones
- SpMV: benefits lower due to **low Computational Intensity** (read the matrix)
- **Register blocking** and other “compression” can be a big win
- Cache blocking less so; other low level tuning (e.g., prefetch) some
- For distributed memory, **reordering** (e.g., graph partitioning) important
- **Autotuning** possible, but depends on sparsity structure; hybrid offline / online tuning
- After tuning SpMV should be **memory bandwidth limited**
- Optimizing at **higher-level algorithms** (across iterations) can improve reuse.