

CSE 6220/CX 4220

Introduction to HPC

Lecture 14: Dense Matrix Algorithms

Helen Xu


hxu615@gatech.edu



Georgia Tech College of Computing
School of Computational
Science and Engineering


Slide credits to Prof. Srinivas Aluru

Recap: Communication Primitives

- Broadcast
 - Reduce
 - AllReduce
 - Scan
 - Gather
 - AllGather
 - Scatter
 - All_to_All
 - Arbitrary Permutations: $O(\tau \cdot p + \mu \cdot m \cdot p)$
 - Hypercubic Permutations: $O(\tau \cdot \log p + \mu \cdot m \cdot p \cdot \log p)$
- 
- $\Theta(\tau \log p + \mu m \log p)$
- $\Theta(\tau \log p + \mu m p)$

Matrix Algorithms

- Regular structure
- Well-suited to **static partitioning** of computation
- Computation typically partitioned based on the data
 - input, output, intermediate
- Typical data partitioning strategies for matrix algorithms
 - 1D and 2D block, cyclic, block-cyclic
- Easier to think logically about the algorithms on a 2D mesh of processors:



Matrix-Vector Multiplication


Matrix-Vector Multiplication

Problem:

- multiply a dense $n \times n$ matrix A with an $n \times 1$ vector x
- yield the $n \times 1$ result vector y


What is the serial time complexity?

- Serial algorithm requires n^2 multiplications and additions




1D Partitioning for Matrix-Vector Multiply

- 1D partitioning splits up the matrix into **block rows** with at least one row on each processor.
- Given n rows, 1D partitioning can use at most n processors.




2D Partitioning for Matrix-Vector Multiply

- 2D partitioning splits up the matrix into **square sub-matrices** - at least 1 element per processor
- Theoretically, 2D partitioning can use up to n^2 processors (equal to number of elements)



Analysis of 1D Partitioning

- First, consider simple case: $p = n$
- $n \times n$ matrix is partitioned among p processors
 - *Each processor stores a complete row of the matrix*
- $n \times 1$ vector x is also partitioned in a conformal manner
 - *Each process owns one element of x*



Analysis of 1D Partitioning


- Step 1: Use AllGather to distribute all of x to each processor

$$O(\tau \log n + \mu \cdot n)$$

$\mu n > n$, so time is basically the cost of allgather

- Step 2: Processor P_i computes $y[i] = \sum_{j=0}^{n-1} A[i, j] \cdot x[j]$


$$O(n)$$



Analysis of 1D Partitioning

For $n > p$

- Initially, each processor stores
 - n/p complete rows of the matrix A
 - n/p elements of the vector x
- Distribute all of x vector to each processor
 - AllGather among p processors, messages of size n/p
- On each processor: n/p local dot products on vectors (rows) of length n




• Step 1: Comm. Time = What is the communication time?

• Step 2: Computation time = What is the computation time?

Analysis of 1D Partitioning

For $n > p$

- Initially, each processor stores
 - n/p complete rows of the matrix A
 - n/p elements of the vector x
- Distribute all of x vector to each processor
 - AllGather among p processors, messages of size n/p
- On each processor: n/p local dot products on vectors (rows) of length n




- Step 1: Comm. Time = $O\left(\tau \log p + \mu \cdot \frac{n}{p} \cdot \cancel{p}\right) = O(\tau \log p + \mu \cdot n)$

- Step 2: Computation time = What is the computation time?

Analysis of 1D Partitioning

For $n > p$

- Initially, each processor stores
 - n/p complete rows of the matrix A
 - n/p elements of the vector x
- Distribute all of x vector to each processor
 - AllGather among p processors, messages of size n/p
- On each processor: n/p local dot products on vectors (rows) of length n



- Step 1: Comm. Time = $O\left(\tau \log p + \mu \cdot \frac{n}{p} \cdot p\right) = O(\tau \log p + \mu \cdot n)$

- Step 2: Computation time = $O\left(\frac{n^2}{p}\right)$

Can only use up to n/μ processors because computation time needs to dominate for efficiency


Announcements

- HW4 due today EOD
- No class next week, PA2 due next Fri 3/7 EOD, late deadline 3/9 EOD
- Please be mindful about what resources you request on PACE ICE - see Canvas announcement
- Class on 3/10 in East Architecture 123 (same time as usual)


Analysis of 2D Partitioning

- Consider : $p = n^2$
- $n \times n$ matrix is partitioned among $p = n^2$ processors (virtual grid topology)
 - *Each processor owns a single element*
- $n \times 1$ vector x is distributed in the last column of n processors


What is a good algorithm given this distribution?




(a) Initial data distribution and communication steps to align the vector along the diagonal



(b) One-to-all broadcast of portions of the vector along process columns



(c) All-to-one reduction of partial results



(d) Final distribution of the result vector

Analysis of 2D Partitioning

$$p = n^2$$

- Algorithm

- align vector along the main diagonal: one-to-one communication $(\tau + \mu)\lg n$
- Broadcast vector element to n processors in column: one-to-all broadcast $(\tau + \mu)\lg n$
- Local multiplication 1
- Sum partial y values in each row: all-to-one reduction $(\tau + \mu)\lg n + \lg n$

Analysis of 2D Partitioning

$$p = n^2$$

- Algorithm

- align vector along the main diagonal: one-to-one communication $(\tau + \mu)\lg n$

- Broadcast vector element to n processors in column: one-to-all broadcast $(\tau + \mu)\lg n$

- Local multiplication 1

- Sum partial y values in each row: all-to-one reduction $(\tau + \mu)\lg n + \lg n$

- Communication time = $O(\tau \log n + \mu \cdot \log n)$

- Computation time = $O(\log n)$ From reduction

Is the algorithm efficient?

Analysis of 2D Partitioning

$$p = n^2$$

- Algorithm

- align vector along the main diagonal: one-to-one communication $(\tau + \mu)\lg n$
- Broadcast vector element to n processors in column: one-to-all broadcast $(\tau + \mu)\lg n$
- Local multiplication 1
- Sum partial y values in each row: all-to-one reduction $(\tau + \mu)\lg n + \lg n$

- Communication time = $O(\tau \log n + \mu \cdot \log n)$

- Computation time = $O(\log n)$ From reduction


- Is the algorithm efficient?

- **No** – we need $O(1)$ run-time!

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- When using fewer than n^2 processors, each process owns an $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ block of the matrix.
- The vector is distributed in portions of $\frac{n}{\sqrt{p}}$ elements in the last processor-column only.




2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- When using fewer than n^2 processors, each process owns an $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ block of the matrix.
- The vector is distributed in portions of $\frac{n}{\sqrt{p}}$ elements in the last processor-column only.
- In this case, the message sizes for the alignment, broadcast, and reduction are all $\frac{n}{\sqrt{p}}$.


Part of the vector



2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- When using fewer than n^2 processors, each process owns an $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ block of the matrix.
- The vector is distributed in portions of $\frac{n}{\sqrt{p}}$ elements in the last processor-column only.
- In this case, the message sizes for the alignment, broadcast, and reduction are all $\frac{n}{\sqrt{p}}$.
- The computation is a product of an $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ submatrix with a vector of length $\frac{n}{\sqrt{p}}$.



Run serial algorithm on the submatrix on each processor

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm
 - align vector along the main diagonal:

How many
vector elts?

What is the
communication
time?

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$

message size
 $O(n/\sqrt{p})$

assuming direct
connection,
otherwise $\times \lg(\sqrt{p})$

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
- Broadcast vector elements to processors in column:

assuming direct connection,
otherwise $\times \lg(\sqrt{p})$

How many
procs in each
column?

What is the
message size
per proc?

What is the
communication
time?

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
- Broadcast vector elements to processors in column:
 $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$

assuming direct connection,
otherwise $\times \lg(\sqrt{p})$

Each column
has $\sqrt{p}$ procs

message size
 $O(n/\sqrt{p})$

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
- Broadcast vector elements to processors in column:
 $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
- Local multiplication:

assuming direct connection,
otherwise $\times \lg(\sqrt{p})$

How much of the matrix is on each processor?

What is the computation time?

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
- Broadcast vector elements to processors in column:
 $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
- Local multiplication: $O\left(\frac{n^2}{p}\right)$
- Sum partial y values in each row:

assuming direct connection,
otherwise $\times \lg(\sqrt{p})$

What is the communication time?

What is the message size per proc?

How many procs in each row?

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
- Broadcast vector elements to processors in column:
 $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
- Local multiplication: $O\left(\frac{n^2}{p}\right)$
- Sum partial y values in each row: $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$

assuming direct connection,
otherwise $\times \lg(\sqrt{p})$

Each row has
 $\sqrt{p}$ procs

message size
 $O(n/\sqrt{p})$

2D Partitioning for Matrix-Vector Multiply

$$p < n^2$$

- Algorithm

- align vector along the main diagonal: $O\left(\tau + \mu \cdot \frac{n}{\sqrt{p}}\right)$
- Broadcast vector elements to processors in column:
 $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
- Local multiplication: $O\left(\frac{n^2}{p}\right)$
- Sum partial y values in each row: $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
- Communication time = $O\left(\tau \log \sqrt{p} + \mu \frac{n}{\sqrt{p}} \log \sqrt{p}\right)$
- Computation time = $O\left(\frac{n^2}{p}\right)$
- Efficient up to $O\left(\frac{n^2}{\log^2 n}\right)$ processors.

assuming direct
connection,
otherwise $\times \lg(\sqrt{p})$


More scalable than 1D
on the hypercube

Matrix-Matrix Multiplication

Matrix-Matrix Multiplication

- Consider the problem of multiplying two $n \times n$ dense, square matrices A and B to yield the product matrix $C = A \times B$.
- The serial complexity is $O(n^3)$.


```
for (i = 0; i < n; ++i)
  for (j = 0; j < n; ++j) {
    double sum = 0;
    for (k = 0; k < n; ++k)
      sum += A[i][k] * B[k][j];
    C[i][j] = sum;
  }
```



Initial Algorithm for Matrix-Matrix Multiplication

Assume we have n^2 processors, where every processor (i, j) has elements $A[i, j]$ and $B[i, j]$ in the initial distribution.

For processor (i, j) to compute $C[i, j]$, it needs the i -th row of A and j -th col of B .




Idea:

- Allgather i -th row of A onto each processor for each row i
- Allgather j -th col of B onto each processor for each col j
- Perform dot product of row and col to get $C[i, j]$


Data Requirements for Matrix-Matrix Multiply

Consider data needed for output matrix block shown in purple



Data Requirements for Matrix-Matrix Multiply

Consider data needed for output matrix block shown in purple



Matrix-Matrix Multiplication


- Consider the problem of multiplying two $n \times n$ dense, square matrices A and B to yield the product matrix $C = A \times B$.
- The serial complexity is $O(n^3)$.
 - We do not consider better serial algorithms (Strassen's method)
Note: these can still be used as serial kernels in the parallel algorithms.

Matrix-Matrix Multiplication

- Consider the problem of multiplying two $n \times n$ dense, square matrices A and B to yield the product matrix $C = A \times B$.
- The serial complexity is $O(n^3)$.
 - We do not consider better serial algorithms (Strassen's method)
Note: these can still be used as serial kernels in the parallel algorithms.
- A useful concept in this case is called *block* operations
 - an $n \times n$ matrix A can be regarded as a $q \times q$ array of blocks
 - $A_{i,j}$ ($0 \leq i, j < q$) such that each block is an $(n/q) \times (n/q)$ submatrix.
- In this view, we perform q^3 matrix multiplications, each involving $(n/q) \times (n/q)$ matrices.

Blocked Matrix-Matrix Multiply

- Consider two $n \times n$ matrices A and B partitioned into p blocks $A_{i,j}$ and $B_{i,j}$ ($0 \leq i, j < \sqrt{p}$) of size $(n/\sqrt{p}) \times (n/\sqrt{p})$ each.
- Process $P_{i,j}$ initially stores $A_{i,j}$ and $B_{i,j}$ and computes block $C_{i,j}$ of the result matrix.
- Computing submatrix $C_{i,j}$ requires all submatrices $A_{i,k}$ and $B_{k,j}$ for $0 \leq k < \sqrt{p}$.
 - AllGather blocks of A along rows and B along columns.
- Perform local submatrix multiplication.



Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:

What is the
communication time?

What is the message
size per proc? How
many procs participate
in each allgather?

Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
 - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$

Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
 - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices:

What is the computation time?

Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
 - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices:
 - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
- Total Time: $O\left(\frac{n^3}{p} + \tau \log p + \mu \frac{n^2}{\sqrt{p}}\right)$
- Efficient for $p = O(n^2)$

Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
 - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices:
 - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
- Total Time: $O\left(\frac{n^3}{p} + \tau \log p + \mu \frac{n^2}{\sqrt{p}}\right)$
- Efficient for $p = O(n^2)$

Any issues with this algorithm?

Analysis of Blocked Matrix-Matrix Multiply

- The two AllGather operations take time:
 - $O(\tau \log \sqrt{p} + \mu \frac{n^2}{p} \sqrt{p})$
- The computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices:
 - $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
- Total Time: $O\left(\frac{n^3}{p} + \tau \log p + \mu \frac{n^2}{\sqrt{p}}\right)$
- Efficient for $p = O(n^2)$
- Major drawback of the algorithm is that it is not memory optimal.


Extra $\sqrt{p}$ factor of
memory consumption!

Ideal: Memory is
serial_memory / p

Restriction - one instance of each element in A, B across all processors

Send a elements to the left one hop at a time

Send b elements to the north one hop at a time




$$c_{0,0} = a_{0,0}b_{0,0} + a_{0,1}b_{1,0} + a_{0,2}b_{2,0} + a_{0,3}b_{3,0}$$

$$c_{0,1} = \boxed{a_{0,0}}b_{0,1} + \boxed{a_{0,1}b_{1,1}} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

Goal - all processors are computing in each round

We need the corresponding a and b on each processor in each round




$$c_{0,0} = a_{0,0}b_{0,0} + a_{0,1}b_{1,0} + a_{0,2}b_{2,0} + a_{0,3}b_{3,0}$$

$$c_{0,1} = a_{0,0}b_{0,1} + a_{0,1}b_{1,1} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

Can compute the terms in any order

Goal - all processors are computing in each round


We need the corresponding a and b on each processor in each round



$$c_{0,0} = a_{0,0}b_{0,0} + a_{0,1}b_{1,0} + a_{0,2}b_{2,0} + a_{0,3}b_{3,0}$$


$$c_{0,1} = a_{0,0}b_{0,1} + a_{0,1}b_{1,1} + a_{0,2}b_{2,1} + a_{0,3}b_{3,1}$$

Can compute the terms in any order




Cannon's Algorithm: Rotations

Consider processor $P_{0,0}$: Step 1 - Multiply the **Red** blocks




Cannon's Algorithm: Rotations

$P_{0,0}$: Step 2 – Left shift A, Up shift B, Multiply the **Blue** blocks




Cannon's Algorithm: Rotations

$P_{0,0}$: Step 3 – Left shift A, Up shift B, Multiply the **Green** blocks




Cannon's Algorithm: Rotations

$P_{0,0}$: Step 4 – Left shift A, Up shift B, Multiply the **Yellow** blocks




Cannon's Algorithm: Initialization

Consider processor $P_{0,1}$




Cannon's Algorithm: Initialization

Consider processor $P_{0,1}$: Align blocks by shifting B up




Cannon's Algorithm

- Consider processor $P_{2,1}$
- We need to align the blocks!
 - In row 2 shift A_{ij} left by 2
 - In column 1 shift B_{ij} up by 1




Cannon's Algorithm

- Consider processor $P_{2,1}$
- Alignment Step 1




Cannon's Algorithm

- Consider processor $P_{2,1}$
- Alignment Step 2




Cannon's Algorithm

- Consider processor $P_{2,1}$
- Multiplication Step 1




Cannon's Algorithm

- Consider processor $P_{2,1}$
- Multiplication Step 2




Cannon's Algorithm

- Consider processor $P_{2,1}$
- Multiplication Step 3



Cannon's Algorithm

- Consider processor $P_{2,1}$
- Multiplication Step 4




Cannon's Algorithm: Initial Alignment

Consider processor $P_{i,j}$

- Aligning blocks:
 - Shift $A_{i,j}$ left by i
 - $A_{i,(j+i)}$
 - Shift $B_{i,j}$ up by j
 - $B_{(i+j),j}$
- $P_{i,j}$ will have $A_{i,(j+i) \bmod \sqrt{p}}$ and $B_{(i+j) \bmod \sqrt{p},j}$

Line up inner dimension



Analysis of Cannon's Algorithm

- In the alignment step,
 - maximum distance over which a block shifts is $\sqrt{p} - 1$,
 - two shift operations require a total of $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$ time.

Analysis of Cannon's Algorithm

- In the alignment step,
 - maximum distance over which a block shifts is $\sqrt{p} - 1$,
 - two shift operations require a total of $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$ time.
- Compute-and-shift phase has $\sqrt{p}$ steps
 - Computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices: $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
 - Communication requires $\sqrt{p}$ shifts $O(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p})$

Analysis of Cannon's Algorithm

- In the alignment step,
 - maximum distance over which a block shifts is $\sqrt{p} - 1$,
 - two shift operations require a total of $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$ time.
- Compute-and-shift phase has $\sqrt{p}$ steps
 - Computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices: $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
 - Communication requires $\sqrt{p}$ shifts $O(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p})$
- Total Time: $O\left(\frac{n^3}{p} + \tau\sqrt{p} + \mu\frac{n^2}{\sqrt{p}}\right)$

Analysis of Cannon's Algorithm

- In the alignment step,
 - maximum distance over which a block shifts is $\sqrt{p} - 1$,
 - two shift operations require a total of $O\left(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p}\right)$ time.
- Compute-and-shift phase has $\sqrt{p}$ steps
 - Computation requires $\sqrt{p}$ multiplications of $(n/\sqrt{p}) \times (n/\sqrt{p})$ sized submatrices: $O\left(\sqrt{p} \times \left(\frac{n}{\sqrt{p}}\right)^3\right) = O\left(\frac{n^3}{p}\right)$
 - Communication requires $\sqrt{p}$ shifts $O(\tau\sqrt{p} + \mu\frac{n^2}{p}\sqrt{p})$
- Total Time: $O\left(\frac{n^3}{p} + \tau\sqrt{p} + \mu\frac{n^2}{\sqrt{p}}\right)$
- **More messages, but Cannon's Algorithm is memory optimal.**


```

1  MatrixMatrixMultiply(int n, double *a, double *b, double *c,
2                        MPI_Comm comm)
3  {
4      int i;
5      int nlocal;
6      int npes, dims[2], periods[2];
7      int myrank, my2drank, mycoords[2];
8      int uprank, downrank, leftrank, rightrank, coords[2];
9      int shiftsource, shiftdest;
10     MPI_Status status;
11     MPI_Comm comm_2d;
12
13     /* Get the communicator related information */
14     MPI_Comm_size(comm, &npes);
15     MPI_Comm_rank(comm, &myrank);
16
17     /* Set up the Cartesian topology */
18     dims[0] = dims[1] = sqrt(npes);
19
20     /* Set the periods for wraparound connections */
21     periods[0] = periods[1] = 1;
22
23     /* Create the Cartesian topology, with rank reordering */
24     MPI_Cart_create(comm, 2, dims, periods, 1, &comm_2d);
25
26     /* Get the rank and coordinates with respect to the new topology */
27     MPI_Comm_rank(comm_2d, &my2drank);
28     MPI_Cart_coords(comm_2d, my2drank, 2, mycoords);
29
30     /* Compute ranks of the up and left shifts */
31     MPI_Cart_shift(comm_2d, 0, -1, &rightrank, &leftrank);
32     MPI_Cart_shift(comm_2d, 1, -1, &downrank, &uprank);
33
34     /* Determine the dimension of the local matrix block */
35     nlocal = n/dims[0];
36
37     /* Perform the initial matrix alignment. First for A and then for B */
38     MPI_Cart_shift(comm_2d, 0, -mycoords[0], &shiftsource, &shiftdest);
39     MPI_Sendrecv_replace(a, nlocal*nlocal, MPI_DOUBLE, shiftdest,
40                          1, shiftsource, 1, comm_2d, &status);

```




```

41
42 MPI_Cart_shift(comm_2d, 1, -mycoords[1], &shiftsource, &shiftdest);
43 MPI_Sendrecv_replace(b, nlocal*nlocal, MPI_DOUBLE,
44     shiftdest, 1, shiftsource, 1, comm_2d, &status);
45
46 /* Get into the main computation loop */
47 for (i=0; i<dims[0]; i++) {
48     MatrixMultiply(nlocal, a, b, c); /* c = c + a*b */
49
50     /* Shift matrix a left by one */
51     MPI_Sendrecv_replace(a, nlocal*nlocal, MPI_DOUBLE,
52         leftrank, 1, rightrank, 1, comm_2d, &status);
53
54     /* Shift matrix b up by one */
55     MPI_Sendrecv_replace(b, nlocal*nlocal, MPI_DOUBLE,
56         uprank, 1, downrank, 1, comm_2d, &status);
57 }
58
59 /* Restore the original distribution of a and b */
60 MPI_Cart_shift(comm_2d, 0, +mycoords[0], &shiftsource, &shiftdest);
61 MPI_Sendrecv_replace(a, nlocal*nlocal, MPI_DOUBLE,
62     shiftdest, 1, shiftsource, 1, comm_2d, &status);
63
64 MPI_Cart_shift(comm_2d, 1, +mycoords[1], &shiftsource, &shiftdest);
65 MPI_Sendrecv_replace(b, nlocal*nlocal, MPI_DOUBLE,
66     shiftdest, 1, shiftsource, 1, comm_2d, &status);
67
68 MPI_Comm_free(&comm_2d); /* Free up communicator */
69 }
70
71 /* This function performs a serial matrix-matrix multiplication c = a*b */
72 MatrixMultiply(int n, double *a, double *b, double *c)
73 {
74     int i, j, k;
75
76     for (i=0; i<n; i++)
77         for (j=0; j<n; j++)
78             for (k=0; k<n; k++)
79                 c[i*n+j] += a[i*n+k]*b[k*n+j];
80 }

```

Limitations of Cannon's algorithm



Requires a square 2D grid of processors

Also, A and B need to be square

Outer Product Matrix Multiplication

Inner-Product Matrix Multiplication

```
for(int i = 0; i < n; ++i) {  
    for(int j = 0; j < n; ++j) {  
        for(int k = 0; k < n; ++k) {  
            C[i][j] += A[i][k] * B[k][j];  
        }  
    }  
}
```

So far, we have been doing inner product matrix multiply, which calculates n^2 dot products (inner products).

$$C[i,j] = A[i, :] \times B[:, j]$$

Outer product definition

Multiplying two vectors u ($n \times 1$) and v ($1 \times n$) gives us a matrix of size $n \times n$ where each entry is the product of one element from u and v .

$$\begin{bmatrix} a \\ b \\ c \end{bmatrix} * \begin{bmatrix} e & f & g \end{bmatrix} = \begin{bmatrix} a*e & a*f & a*g \\ b*e & b*f & b*g \\ c*e & c*f & c*g \end{bmatrix}$$

Outer-Product Matrix Multiplication

```
for(int k = 0; k < n; ++k) {  
    for(int i = 0; i < n; ++i) {  
        for(int j = 0; j < n; ++j) {  
            C[i][j] += A[i][k] * B[k][j];  
        }  
    }  
}
```

The product of a column vector by a row vector is an $n \times n$ matrix, or the “multiplication table” of the two elements.

Running over the outer k loop in matrix multiply, we can express C as the **sum of k of these “multiplication table” matrices.**

$$C = \sum_k A[:, k] \times B[k, :]$$

SUMMA algorithm

Scalable Universal Matrix Multiplication Algorithm (SUMMA) [\[van de Geijn and Watts, 97\]](#)

- Used in ScaLAPACK, PLAPACK, and PBLAS libraries
- Based on outer product

Avoids need for square matrices and processor grid

```
for(int k = 0; k < n; ++k) {  
    C[:, :] += A[:, k] + B[k, :]  
}
```


Parallelize all (i, j)

SUMMA algorithm


On each process $P(i, j)$:

```
for(int k = 0; k < n; ++k) {  
  Bcast column k of A ( $A[i, k]$ ) within proc row i  
  Bcast row k of B ( $B[k, j]$ ) within proc column j  
  Do  $C += A[i, k] * B[k, j]$   
}
```

Compute the sum
of all outer
products




SUMMA Example: Initial Configuration



SUMMA Example: $k = 0$

Broadcast $A_{i,0}$
within processor
row




Inner dimension
needs to match

keep the original data as well, just not pictured

SUMMA Example: $k = 0$


Broadcast $B_{0,j}$
within processor
column



Inner dimension
matches now:
multiply and add
to C

SUMMA Example: $k = 0$


Broadcast $B_{0,j}$
within processor
column



Inner dimension
matches now


keep the original data as well, just not pictured

Next round: starting with initial configuration again




SUMMA Example: $k = 1$

Broadcast $A_{i,1}$
within processor
row




SUMMA Example: $k = 1$

Broadcast $B_{1,j}$
within processor
column




Inner dimension
matches now:
multiply and add
to C

Next round: starting with initial configuration again




SUMMA Example: $k = 2$

Broadcast $A_{i,2}$
within processor
row



SUMMA Example: $k = 2$

Broadcast $B_{2,j}$
within processor
column



Inner dimension
matches now:
multiply and add
to C

Blocking matrices in SUMMA

Suppose we have $n \times n$ square matrices A , B , C (for simplicity).

Rather than doing n iterations of the outer loop, pick some **block size b** (b columns of A and b rows of B) so that we can do n/b iterations of the outer loop.

```
for(int k = 0; k < n; k+=b) {  
    C[:,] += A[:, k:k+b-1] + B[k:k+b-1, :]  
}
```

Multiply $n \times b$ piece of A
with $b \times n$ piece of B


Communication cost

- The simplest choice of block size is $b = O(n/\sqrt{p})$ (assuming p processors in a square grid for simplicity).
- So each $n \times b$ submatrix is size $O(n^2/\sqrt{p})$.
- Each broadcast costs $O((\tau + \mu m)\lg(\sqrt{p})) = O((\tau + \mu n^2/\sqrt{p})(\lg p))$
- There are $n/b = O(\sqrt{p})$ rounds
- Total cost is (num rounds) $\times$ (cost of broadcast per round)

Substitute $m = n^2/\sqrt{p}$

2.5D Matrix Multiplication


- Assume can fit cn^2/P data per processor, $c > 1$
- Processors form $(P/c)^{1/2} \times (P/c)^{1/2} \times c$ grid



Example: $P = 32$, $c = 2$

2.5D Matrix Multiplication

- Assume can fit cn^2/P data per processor, $c > 1$
- Processors form $(P/c)^{1/2} \times (P/c)^{1/2} \times c$ grid



Initially $P(i,j,0)$ owns $A(i,j)$ and $B(i,j)$
each of size $n(c/P)^{1/2} \times n(c/P)^{1/2}$

- (1) $P(i,j,0)$ broadcasts $A(i,j)$ and $B(i,j)$ to $P(i,j,k)$
- (2) Processors at level k perform $1/c$ -th of SUMMA, i.e. $1/c$ -th of $\sum_m A(i,m) * B(m,j)$
- (3) Sum-reduce partial sums $\sum_m A(i,m) * B(m,j)$ along k -axis so $P(i,j,0)$ owns $C(i,j)$

2.5D Matrix Multiplication on IBM BG/P, 16k nodes / 64k cores

