

CSE 6220/CX 4220
Introduction to HPC

Lecture 18: Hierarchical Methods for the N-Body Problem

Helen Xu

hxu615@gatech.edu




Georgia Tech College of Computing
School of Computational
Science and Engineering

Slide credits mostly to Prof. James Demmel

Motivation

- **Computational simulation** is becoming an accepted paradigm for scientific discovery - but can involve several million variables
- Most large problems boil down to solution of linear systems or matrix-vector product
- Regular product requires $O(n^2)$ time and memory

Applying the idea recursively: cost drops to $O(n \log n)$ or even $O(n)$




https://www.cscamm.umd.edu/programs/fam04/dg_lecture1.pdf

Diagram from https://amath.colorado.edu/faculty/martinss/2014_CBMS/Lectures/lecture02.pdf

Motivation

- Intuition - **compress dependence on “distant” data**
- Compressing data of groups at nearby points can cut cost (work, communication) at distant points
- Examples:
 - multigrid for elliptic PDE
 - multilevel graph partitioning (METIS, Chaco, ...)
 - Barnes-Hut or Fast Multipole Method (FMM) for electrostatics/gravity/...

Applying the idea recursively: cost drops to $O(n \log n)$ or even $O(n)$



https://www.cscamm.umd.edu/programs/fam04/dg_lecture1.pdf

Diagram from https://amath.colorado.edu/faculty/martinss/2014_CBMS/Lectures/lecture02.pdf

Particle Simulation

```
t=0
while t < t_final:
  for i = 1 to n
    compute f(i) = force on particle i
  for i = 1 to n
    move particle i under force f(i) for time dt
  compute interesting properties of particles (energy, etc.)
  t = t + dt
end while
```

Most complicated component in computing $f(i)$ is **n-body force**, e.g., gravity or electrostatics

$$f(i) = \sum_{k \neq i} f(i, k)$$

$f(i, k)$ = force on i
from k

Obvious algorithm requires computing all-to-all interactions for $O(n^2)$, but we can do better

Motivation

- Listed as one of the top 10 algorithms of the 20th century (IEEE magazine)
- Frequently used in scientific computing





diagram from UC
Berkeley
CS267

Applications

- **Astrophysics and Celestial Mechanics - 1992**
 - Intel Delta = 1992 supercomputer, 512 Intel i860s
 - **17 million particles**, 600 time steps, 24 hours elapsed time
 - M. Warren and J. Salmon
 - Gordon Bell Prize at Supercomputing **1992**
 - Sustained **5.2 Gigafllops** = 44K Flops/particle/time step
 - 1% accuracy
 - Direct method (17 Flops/particle/time step) at 5.2 Gflops would have taken 18 years, 6570 times longer
- **Vortex particle simulation of turbulence – 2009**
 - Cluster of 256 NVIDIA GeForce 8800 GPUs
 - **16.8 million particles**
 - T. Hamada, R. Yokota, K. Nitadori. T. Narumi, K. Yasoki et al
 - Gordon Bell Prize for Price/Performance at Supercomputing **2009**
 - Sustained **20 Teraflops**, or **\$8/Gigaflop**

Reducing the number of particles in the force sum

- All later divide and conquer algorithms use same intuition
- Consider computing force on earth due to all celestial bodies
 - Look at night sky, # terms in force sum $\geq$ number of visible stars
 - Oops! One “star” is really the Andromeda galaxy, which contains billions of real stars
 - Seems like a lot more work than we thought ...
- Don't worry, ok to approximate all stars in Andromeda by a single point at its center of mass (CM) with same total mass (TM)
 - D = size of box containing Andromeda , r = distance of CM to Earth
 - Require that D/r be “small enough”




- Idea not new: Newton approximated earth and falling apple by CMs

What is new: Using points at CM recursively

- From Andromeda's point of view, Milky Way is also a point mass
- Within Andromeda, picture repeats itself
 - As long as $D1/r1$ is small enough, stars inside smaller box can be replaced by their CM to compute the force on Vulcan
 - Boxes nest in boxes recursively

Replacing Clusters by their Centers of Mass Recursively





Basic Data Structures: Quad Trees and Oct Trees

Quad Trees

- **Data structure to subdivide the plane**
 - Nodes can contain coordinates of center of box, side length
 - Eventually also coordinates of CM, total mass, etc.
- In a **complete** quad tree, each nonleaf node has 4 children



A Complete Quadtree with 4 Levels



Oct Trees



- **Similar Data Structure to subdivide space**

2 Levels of an Octree





Using Quad Trees and Oct Trees

- All approximate n-body simulation algorithms begin by constructing **a tree to hold all the particles**
- Interesting cases have nonuniformly distributed particles
 - In a complete tree most nodes would be empty, wasting space and time
- **Adaptive** Quad (Oct) Tree only subdivides space where particles are located



Example of an Adaptive Quad Tree

Adaptive quadtree where no square contains more than 1 particle



Child nodes enumerated counterclockwise from SW corner, empty ones excluded

In practice, have $q>1$ particles/square; tuning parameter

Cost of Adaptive Quad Tree Construction

- **Cost $\leq N * \text{maximum cost of Quad_Tree_Insert}$**
 $= O(N * \text{maximum depth of Quad_Tree})$
- **Uniform Distribution of particles**
 - **Depth of Quad_Tree = $O(\log N)$**
 - **Cost $\leq O(N * \log N)$**
- **Arbitrary distribution of particles**
 - **Depth of Quad_Tree = $O(\# \text{ bits in particle coords }) = O(b)$**
 - **Cost $\leq O(b N)$**

Barnes-Hut Algorithm (BH):
An $O(n \log n)$ approximation algorithm for the
N-Body Problem

Barnes-Hut Algorithm


- **High Level Algorithm (in 2D, for simplicity)**

- 1) **Build the QuadTree using QuadTreeBuild**
... **already described, cost = $O(N \log N)$ or $O(b N)$**
- 2) **For each node = subsquare in the QuadTree, compute the CM and total mass (TM) of all the particles it contains**
... **“post order traversal” of QuadTree, cost = $O(N \log N)$ or $O(b N)$**
- 3) **For each particle, traverse the QuadTree to compute the force on it, using the CM and TM of “distant” subsquares**
... **core of algorithm**
... **cost depends on accuracy desired but still $O(N \log N)$ or $O(bN)$**

Center of mass of a cell is the weighted average of centers of its children

Step 3 of BH: Compute force on each particle


- For each node = square, can approximate force on particles outside the node due to particles inside node by using the node's CM and TM
- This will be accurate enough if the node is "far away enough" from the particle
- For each particle, use as few nodes as possible to compute force, subject to accuracy constraint
- Need criterion to decide if a node is far enough from a particle
 - D = side length of node
 - r = distance from particle to CM of node
 - θ = user supplied error tolerance < 1
 - Use CM and TM to approximate force of node on box if $D/r < \theta$



Compute force on this particle


2D example

In total, Barnes-Hut required only
5 force calculations instead of 23!




--- far-field approximation based on center of mass
— direct force calculation

Example Performance Analysis



Example Error Analysis

Average vector distance between naive method and Barnes-Hut



In practice, 1.0 is a good default value

Fast Multipole Method (FMM): An $O(N)$ approximate algorithm for the N-Body Problem

Most slides by Prof. Lexing Ying (talk here: https://www.youtube.com/watch?v=qMLlyZi8Sz0&ab_channel=IBMResearch)

Some diagrams from <https://web.stanford.edu/~lexing/pfmm.pdf>

Fast Multipole Method (FMM)


- “A fast algorithm for particle simulation”, L. Greengard and V. Rokhlin, J. Comp. Phys. V. 73, 1987, many later papers
 - Many awards
- Differences from Barnes-Hut
 - FMM computes the *potential* at every point, not just the force
 - FMM uses more information in each box than the CM and TM, so it is both more accurate and more expensive
 - In compensation, FMM accesses a fixed set of boxes at every level, independent of D/r
 - BH uses fixed information (CM and TM) in every box, but # boxes increases with accuracy. FMM uses a fixed # boxes, but the amount of information per box increases with accuracy.
- FMM uses two kinds of expansions
 - **Outer expansions** represent potential outside node due to particles inside, analogous to (CM,TM)
 - **Inner expansions** represent potential inside node due to particles outside;
Computing this for every leaf node is the computational goal of FMM
- First review potential, then return to FMM

N-body Problem Formulation

- Points $\{x_i, i = 1, \dots, N\}$
- Charges $\{f_i, i = 1, \dots, N\}$
- Kernel $G(x, y)$, for example
 - $G(x, y) = \ln(|x - y|)$ in 2D
 - $G(x, y) = 1/|x - y|$ in 3D
- Compute potentials for x_i for $i = 1, \dots, N$

$$u_i = \sum_{j=1}^N G(x_i, x_j) f_j$$

Can be expressed
with matrix-vector
multiplication



Computational Cost


$$u_i = \sum_{j=1}^N G(x_i, x_j) f_j$$

- Direct computation is $O(N^2)$ steps : too costly
- Goal: lower the complexity
- Fast multipole method - Greengard and Rokhlin in 1987
 - For $G(x, y) = \ln |x - y|$ or $1/|x - y|$
 - Bring the complexity to $O(N)$ for some fixed accuracy ϵ
- Based on four key ideas


Idea 1: Well-separated Regions

- Suppose boxes B and A are well-separated
- Consider the influence from B to A: at each x_i in A, evaluate

$$u_i = \sum_{x_j \in B} G(x_i, x_j) f_j$$



Cost of 3 step approximation



- Reduce $O(N_A N_B)$ steps to $O(N_A + N_B)$ steps
- Good accuracy when **A, B are far away**
- We will use it whenever A and B are well separated

Low-rank matrix approximation





Algebraic Interpretation

$$\begin{array}{ccccccc}
 (u_i)_{x_i \in A} & & (G(x_i, x_j))_{x_i \in A, x_j \in B} & (f_j)_{x_j \in B} & & G(c_A, c_B) & (f_j)_{x_j \in B} \\
 \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} & = & \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} & \approx & \begin{array}{|c|} \hline 1 \\ 1 \\ \dots \\ 1 \\ \hline \end{array} & \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} & \begin{array}{|c|} \hline 1 & 1 & \dots & 1 \\ \hline \end{array} \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} \\
 & & & & & & \underbrace{\hspace{10em}}_{f_B} \\
 & & & & & & \underbrace{\hspace{10em}}_{u_A}
 \end{array}$$

- f_B = source representation
- u_A = target representation
- Problem: but each point x_i serves as **both a source and target**


Idea 2: Hierarchical Decomposition

- Partition the domain hierarchically until **each leaf box contains $O(1)$ number of points**





Idea 2: Hierarchical Decomposition

- Partition the domain hierarchically until **each leaf box contains $O(1)$ number of points**
- $N = \# \text{ points}$
- $\# \text{ levels} = O(\log_4 N)$
- $\# \text{ boxes on level } \ell = O(4^\ell)$
- $\text{Total } \# \text{ boxes} = O(N)$




Levels 0 and 1



At levels 0 and 1, the boxes are adjacent to each other, so **we cannot use the 3 step approximation** as the boxes are not well separated

Level 2




B and A are well separated, so use 3 step approximation for influence from B to A

Q: What about interaction with B its neighbors?

- **Go down one level in hierarchical decomposition**

Level 3




B' and A' are well-separated

Use the 3-step approx. for the influence from B' to A'

Interaction list of all B' = all such A', at most $6^2 - 3^2 = 27$ boxes

Level 4



- B'' and A'' are well-separated, so use 3-step approximation for influence from B'' to A''
- Q: influence from B'' to neighbors?
 - Suppose we are at the leaf level, so we cannot keep going down
 - **Evaluate neighbor influence directly**

Algorithm 1

1. For each box B in the tree, compute $f_B = \sum_{x_j \in B} f_j$

2. For $L = 2$ to last level

for each box B at level L

for each A in B 's interaction list

$O(1)$ boxes

$$u_A = u_A + G(c_A, c_B)f_B$$

3. for each box A in the tree

$$u_i = u_i + u_A \text{ for } x_i \in A$$

4. For each box B on the last level

$$u_i = u_i + \sum_{x_j \in \text{nghts}(B)} G(x_i, x_j)f_j, \text{ for } x_i \in B$$

Step 1

1. For each box B in the tree, compute $f_B = \sum_{x_j \in B} f_j$

For a given level ℓ , there 4^ℓ boxes and $O(N/4^\ell)$ points in each of the boxes at that level (assuming uniformity)

So the work to calculate the sum of all the masses at each level is $O(N)$

for $O(N \log N)$ total work

Step 2

- 2. For $L = 2$ to last level
 - for each box B at level L
 - for each A in B 's interaction list

$$u_A = u_A + G(c_A, c_B)f_B$$

For a given level ℓ , there 4^ℓ boxes and each has $O(1)$ neighbors in the interaction list. So the work to do this approximation, assuming you have

already calculated f_B , is $\sum_{\ell=0}^{\log_4 N} 4^\ell = O(N)$

Algorithm 1 Analysis

1. For each box B in the tree, compute $f_B = \sum_{x_j \in B} f_j$

$O(N \log N)$

2. For $L = 2$ to last level

for each box B at level L

for each A in B 's interaction list

$O(1)$ boxes

$O(N)$

$$u_A = u_A + G(c_A, c_B) f_B$$

3. for each box A in the tree

$$u_i = u_i + u_A \text{ for } x_i \in A$$

$O(N \log N)$

4. For each box B on the last level

$$u_i = u_i + \sum_{x_j \in \text{nghts}(B)} G(x_i, x_j) f_j, \text{ for } x_i \in B$$

$O(N)$

can we do better?

Total = $O(N \log N)$


Idea 3: reuse the computation

$$f_B = \sum_{x_j \in B} f_j = \sum_{x_j \in B_1} f_j + \sum_{x_j \in B_2} f_j + \sum_{x_j \in B_3} f_j + \sum_{x_j \in B_4} f_j = f_{B_1} + f_{B_2} + f_{B_3} + f_{B_4}$$

Idea: Compute f_B from its children

O(1) cost


But you have to traverse the quadtree bottom-up



Idea 3: reuse the computation

$$u_i = u_i + u_A, \text{ for } x_i \in A$$

- Similarly, add u_A to only its children - gradually pass down to the lower generations
- $O(1)$ cost
- Traverse quadtree top down



Algorithm 2 (FMM)

1. Go up the tree, for each box B

if B is a leaf box, $f_B = \sum_{x_j \in B} f_j$

else, $f_B = f_{B_1} + f_{B_2} + f_{B_3} + f_{B_4}$

2. Same as step 2 of algorithm 1

3. Go down the tree, for each box A

if A is a leaf box, $u_i = u_i + u_A$, for $x_i \in A$

else, $u_{A_1} = u_{A_1} + u_A$, same for A_2, A_3, A_4

4. Same as step 4 of algorithm 1

Algorithm 2 (FMM)

1. Go up the tree, for each box B

$$\text{if B is a leaf box, } f_B = \sum_{x_j \in B} f_j$$

$$\text{else, } f_B = f_{B_1} + f_{B_2} + f_{B_3} + f_{B_4}$$

$$O(N)$$

2. Same as step 2 of algorithm 1

$$O(N)$$

3. Go down the tree, for each box A

$$\text{if A is a leaf box, } u_i = u_i + u_A, \text{ for } x_i \in A$$

$$\text{else, } u_{A_1} = u_{A_1} + u_A, \text{ same for } A_2, A_3, A_4$$

$$O(N)$$

4. Same as step 4 of algorithm 1

$$O(N)$$

$$\text{Total} = O(N)$$

Three Main Ideas So Far

- When things are far away, use 3-step procedure (low-rank approximation)
- Use hierarchical decomposition to identify what is far away and what is nearby
- Reuse computation to do calculation of the total mass

what about accuracy? not just the total mass, but the mass distribution

Idea 4: Better Source/Target Representations

So far

- f_B is the sum of charges in B
- u_A is the potential at c_A from well-separated sources
- Low accuracy if A and B are not far

More generally

- f_B is a **compact representation** of the charge for B in well-separated targets
- Similarly, u_A is a **compact representation** of the potential in A by well-separated targets
- Better accuracy requires better compact reps.
 - In 2D, treat x, y as points in the complex plane
 - $G(x, y) = \text{RE}(\ln(x-y))$

More Fine-grained Representations

- In practice, to get good accuracy, we can use low-rank approximations rather than rank-1 approximations.
- Instead of representing f_B, u_A as a single value, represent them as a **series of coefficients rather than a single term** - first one is the initial approximation
- Keep $p = \lg(1/\epsilon)$ for accuracy ϵ
- Calculating the total mass now goes to $O(pn_B + pn_A)$. Still linear if $p = O(1)$

Summary of FMM

- The FMM relies on 4 main ideas:
 - Low-rank interaction
 - Hierarchical decomposition
 - Reuse of computation
 - More accurate compact representation for f_B and u_A
- So far we assumed the points are uniform, if they are not, adaptive algorithm is also available

Parallelizing BH, FMM, and related algorithms

Parallelizing Hierarchical N-Body Codes

Barnes-Hut, FMM and related algorithms have similar computational structure:

- 1) Build the QuadTree
- 2) Traverse QuadTree from leaves to root and build outer expansions
(just (TM,CM) for Barnes-Hut)
- 3) Traverse QuadTree from root to leaves and build any inner expansions
(FMM only)
- 4) Traverse QuadTree to accumulate forces for each particle

Parallelizing Hierarchical N-Body Codes

One parallelization scheme will work for them all [\[Blackston and Suel, SC '97\]](#)


- Assign regions of space to each processor
- Regions may have different shapes, to get load balance
 - Each region will have about N/p particles
- Each processor will store part of Quadtree containing all particles (=leaves) in its region, and their ancestors in Quadtree
 - Top of tree stored by all processors, lower nodes may also be shared
- Each processor will also store adjoining parts of Quadtree needed to compute forces for particles it owns
 - Subset of Quadtree needed by a processor called the **Locally Essential Tree (LET)**
- Given the LET, all force accumulations (step 4)) are done in parallel, without communication

Load Balancing Scheme 1: Orthogonal Recursive Bisection (ORB) [\[Warren and Salmon, SC '92\]](#)

- Recursively split region along axes into regions containing equal numbers of particles


Orthogonal Recursive Bisection

**Partitioning
for 16 procs:**



Load Balancing Scheme 2: Costzones

- **Called Costzones for Shared Memory**
 - PhD thesis, J.P. Singh, Stanford, 1993
- **Called “Hashed Oct Tree” for Distributed Memory**
 - Warren and Salmon, Supercomputing 93
- **We will use the name Costzones for both; also in Pbody**
- **Idea: partition QuadTree instead of space**
 - Estimate work for each node, call total work W



Example Performance Results

- **512 Proc Intel Delta**

- **Warren and Salmon, Supercomputing 92, Gordon Bell Prize**
- **8.8 M particles, uniformly distributed**
- **.1% to 1% RMS error, Barnes-Hut**
- **114 seconds = 5.8 Gflops**
 - **Decomposing domain** **7 secs**
 - **Building the OctTree** **7 secs**
 - **Tree Traversal** **33 secs**
 - **Communication during traversal** **6 secs**
 - **Force evaluation** **54 secs**
 - **Load imbalance** **7 secs**
- **Rises to 160 secs as distribution becomes nonuniform**